



Modelling Financial Markets

Chapter 1: Stylised Facts and Returns



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material

- ACF: AutoCorrelation Function
- AI: Artificial Intelligence
- ARCH: AutoRegressive Conditional Heteroskedasticity
- ARCH-LM: ARCH Lagrange Multiplier (test)
- BET: Bucharest Exchange Trading (index)
- BET-TR: BET Total Return (index)
- BNR: Banca Națională a României (the National Bank of Romania)
- BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- CAGR: Compound Annual Growth Rate
- CLT: Central Limit Theorem
- CRIX: CRyptO IndeX
- EGARCH: Exponential GARCH
- ETF: Exchange-Traded Fund
- GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- GEO: Government Emergency Ordinance
- GJR-GARCH: Glosten–Jagannathan–Runkle GARCH
- HAR-RV: Heterogeneous AutoRegressive model of Realised Volatility
- JB: Jarque–Bera (test)
- LB: Ljung–Box (test)
- LM: Lagrange Multiplier
- MDD: Maximum DrawDown
- OHLC: Open, High, Low, Close
- SR: Sharpe Ratio
- US: United States

Outline

Motivation

Prices and Returns

Descriptive Statistics and Tests

Stylised Facts

Markets Compared

Volatility Estimators

Risk-Adjusted Performance

Implications for Modelling

AI for Scientific Discovery

Conclusions

References

What Is a Stylised Fact?

Definition

- ▣ A **stylised fact** is a statistical regularity observed across many assets, markets and periods
- ▣ It is qualitative, not quantitative: “returns have heavy tails”, not “kurtosis = 7.3”
- ▣ It describes **returns**, not price levels; the standard survey is [Cont \(2001\)](#)

Key message of this chapter

- ▣ The **direction** of returns is almost unpredictable, but their **size** (absolute value, a proxy for risk) is highly structured

Why Stylised Facts Matter

- They are the **benchmark** every model of returns must reproduce
- A model that violates them misprices risk: margins, capital and option prices are wrong
- The textbook model of normally distributed, i.i.d. returns fails almost all of them
 - ▶ i.i.d. = **independent and identically distributed**: every day is a new, independent draw from the same distribution
- Stylised facts tell us **which** features a better model needs
- Later chapters build on them: volatility (5, 6, 9), risk (7, 8), machine learning and bubbles (13, 17)

Data Used in This Chapter

Series	Market	Period
S&P 500 index	US equities (price index)	1990–2026
BET-TR	Romanian equities (total return)	2014–2026
Bitcoin	crypto-asset, traded 24/7	2014–2026
EUR/RON	exchange rate, managed float	Jul 2005–2026
Gold	spot price, USD per ounce	1990–2026

- Daily closing values up to 18 September 2026 (repository folder data/market)
- Unless stated otherwise, returns are daily log returns (defined in Section 2)

The Five Series

- ▣ **S&P 500**: Standard & Poor's index of 500 large US companies; a **price index**, dividends are not included
- ▣ **BVB**: the Romanian stock exchange
 - ▶ **BET**: the benchmark index of the BVB
 - ▶ **BET-TR**: its version with dividends re-invested
- ▣ **EUR/RON**: price of one euro in Romanian lei; a **managed float** (market rate with central-bank interventions)
 - ▶ Source: the daily reference rate of the **BNR**
- ▣ **Gold**: spot price in US dollars (USD) per troy ounce; weekday quotes only (weekend prints removed)
- ▣ **Bitcoin**: price in USD; trades 24 hours a day, 7 days a week (24/7)

Data Quality Notes

- EUR/RON: official BNR reference rate instead of exchange quotes
 - ▶ Exchange quote series showed lag-1 autocorrelation (correlation of consecutive daily changes) of -0.4 to -0.7 in several years
 - ▶ Such values come from quote errors, not from the market
- S&P 500 **OHLC** data: Open, High, Low and Close prices of each day
 - ▶ Open prices before 2007 equal the previous close: unusable for range-based estimators
 - ▶ Range estimators therefore start in 2008 (Section 6)
- Lesson: always inspect the data before computing statistics

Prices Before Computers: the Stock Ticker



Stock ticker designed by Thomas Edison
(Smithsonian collection)

Photo: Jaclyn Nash, Smithsonian NMAH; public
domain; Wikimedia Commons

- Before computers, trades reached brokers as prices printed on a paper **ticker tape**
- A daily close is one snapshot of this stream of prices; returns are built from such snapshots
- The choice of snapshot (close, fixing, 00:00 UTC) matters, as Chapter 0 showed
- Prices at the level of individual trades: market microstructure, Chapter 10

Why Returns, Not Prices?

Returns instead of prices

- ▣ **Scale-free:** a 1% move means the same for a \$10 and a \$1,000 stock
- ▣ Prices are **non-stationary**: they trend and wander without a fixed mean (a unit root)
- ▣ **Stationary:** the distribution does not change over time; returns are much closer to it
- ▣ Investors care about relative gains

Adjusted and total-return prices

- ▣ **Dividends:** the price drops on the ex-dividend date, wealth does not
- ▣ **Splits:** mechanical price jumps
- ▣ **Adjusted prices** correct past prices for both
- ▣ **Total return** series re-invest dividends (e.g. BET-TR); the gap to a price index grows each year by the dividend yield

Simple Returns

□ Notation: P_t price at the end of day t , D_t dividend paid in $(t-1, t]$

□ **One-period simple (net) return**

$$R_t = \frac{P_t + D_t - P_{t-1}}{P_{t-1}} = \frac{P_t}{P_{t-1}} - 1 \quad (\text{no dividends})$$

□ **Gross return:** $1 + R_t = P_t/P_{t-1}$, the factor by which wealth is multiplied

□ Simple returns are bounded below: $R_t \geq -100\%$ (limited liability)

□ Reference textbooks: [Campbell, Lo & MacKinlay \(1997\)](#), ch. 1 and [Franke, Härdle & Hafner \(2019\)](#)

Multi-Period and Portfolio Simple Returns

- ▣ **Multi-period return** over k days: gross returns **multiply**

$$1 + R_t(k) = \prod_{j=0}^{k-1} (1 + R_{t-j})$$

- ▣ **Portfolio return**: simple returns aggregate **linearly** across assets

$$R_{p,t} = \sum_{i=1}^N w_i R_{i,t}, \quad \sum_i w_i = 1$$

► w_i : share of wealth invested in asset i at $t - 1$

- ▣ Advantage: exact portfolio accounting
- ▣ Disadvantage: products over time are awkward for statistics

Log Returns

- **Continuously compounded (log) return**

$$r_t = \ln(1 + R_t) = \ln P_t - \ln P_{t-1}$$

- Interpretation: the constant rate that, compounded continuously over one day, turns P_{t-1} into P_t
- Range: $r_t \in (-\infty, +\infty)$, convenient for Normal-distribution or Student- t models
- Back-transformation: $R_t = e^{r_t} - 1$

Properties of Log Returns

- **Time additivity:** multi-period log returns are sums

$$r_t(k) = \sum_{j=0}^{k-1} r_{t-j}$$

- ▶ Sums of many terms allow CLT arguments
- ▶ CLT: a sum of many weakly dependent terms approximately follows the Normal distribution
- **Not additive across assets:** $r_{p,t} \neq \sum_i w_i r_{i,t}$; the approximation is good only for small returns
- Rule of thumb
 - ▶ Log returns for time-series modelling of one asset
 - ▶ Simple returns for portfolios and performance reporting

Worked Example: Small vs Large Moves

Simple return R	Log return $r = \ln(1 + R)$	Difference
+1%	+0.995%	0.005 pp
-1%	-1.005%	0.005 pp
+20%	+18.2%	1.8 pp
-20%	-22.3%	2.3 pp
+50%	+40.5%	9.5 pp
-50%	-69.3%	19.3 pp

- pp = percentage points: the difference between two percentages
- Daily moves are small, so $r \approx R$; for crashes and long horizons the two differ a lot
- For every $R \neq 0$: $\ln(1 + R) < R$ (equality only at $R = 0$): log returns are lower than simple returns

Worked Example: Adding Returns

Over time: +50% then -50%



- ▣ Wealth: $1 \rightarrow 1.5 \rightarrow 0.75$, a loss of 25%
- ▣ Sum of simple returns: $0.5 - 0.5 = 0$ (misleading)
- ▣ Sum of log returns: $\ln 1.5 + \ln 0.5 = -0.288$
- ▣ $e^{-0.288} = 0.75$: exact

Across assets: 50/50 portfolio

- ▣ Asset A: $R = +10\%$; asset B: $R = -10\%$
- ▣ Portfolio simple return: $0.5(0.1) + 0.5(-0.1) = 0$: exact
- ▣ Weighted log returns:
 $0.5 \ln 1.1 + 0.5 \ln 0.9 = -0.005$
- ▣ True portfolio log return: $\ln 1 = 0$

- ▣ Log returns add over **time**; simple returns add over **assets**

Simple vs Log Returns in Practice



- Left: $r \approx R$ only for small moves; the gap grows with $|R|$
- Right: Bitcoin 2014–2026 multiplied wealth by 176.9, i.e. $\sum_t r_t = \ln 176.9 = 5.18$
- Adding daily simple returns gives $1 + \sum_t R_t = 8.83$: wrong by a factor of 20

Annualisation

- ▣ **Annualisation** converts daily statistics into yearly ones
- ▣ Let P be the number of periods per year
 - ▶ $P = 252$ trading days for stocks; $P = 365$ for crypto (traded every day); gold: its actual ≈ 260 weekday quotes a year
- ▣ If returns are i.i.d. with daily mean $\bar{r}$ and daily standard deviation $\hat{\sigma}$

$$\mu_{\text{ann}} = P \bar{r}, \quad \sigma_{\text{ann}} = \sqrt{P} \hat{\sigma}$$

- ▣ **Volatility** = standard deviation of returns, usually reported annualised
- ▣ The $\sqrt{P}$ rule fails under autocorrelation or volatility clustering (see Fact 6)

Worked Example: Annualising a Daily Volatility

S&P 500, 1990–2026

- Daily standard deviation $\hat{\sigma} = 1.133\%$
- $\sqrt{252} = 15.87$
- $\sigma_{\text{ann}} = 1.133\% \times 15.87 = 18.0\%$

Bitcoin, 2014–2026

- Daily standard deviation $\hat{\sigma} = 3.485\%$
- $\sqrt{365} = 19.10$
- $\sigma_{\text{ann}} = 3.485\% \times 19.10 = 66.6\%$

- Bitcoin is about 3.7 times as volatile as the S&P 500
- Using 252 instead of 365 for Bitcoin would understate its volatility by 17%

Arithmetic Mean, Geometric Mean and CAGR

- ▣ **Arithmetic mean:** the average one-period return, $\bar{R} = \frac{1}{T} \sum_t R_t$
- ▣ **Geometric mean:** the constant return that produces the same final wealth
- ▣ **CAGR** = Compound Annual Growth Rate: the geometric mean per year

$$\text{CAGR} = \left(\frac{P_T}{P_0} \right)^{1/Y} - 1, \quad Y = \text{number of years}$$

- ▣ Link between the two means

$$\mu_{\text{geo}} \approx \mu_{\text{arith}} - \frac{1}{2}\sigma^2$$

- ▣ The gap $\frac{1}{2}\sigma^2$ is the **volatility drag**: volatility lowers compound growth


Worked Example: CAGR of the S&P 500

- Close on 2 January 1990: $P_0 = 359.69$; close on 18 September 2026: $P_T = 7,650.50$
- Wealth multiple: $P_T/P_0 = 21.27$
- Length: $Y = 36.71$ years
- $\text{CAGR} = 21.27^{1/36.71} - 1 = 8.69\%$ per year
- Check with log returns: $\ln 21.27/36.71 = 0.0833$, and $e^{0.0833} - 1 = 8.69\%$
- Price index only: with re-invested dividends the CAGR would be higher

Volatility Drag: Bitcoin

- Bitcoin 2014–2026, $P = 365$
 - ▶ Mean of daily simple returns $\times 365$: 65.2% per year (arithmetic)
 - ▶ Realised CAGR: 53.9% per year (geometric)
- The difference of 11 pp is the volatility drag of a very volatile asset
- Which one to report?
 - ▶ CAGR for realised, multi-year performance
 - ▶ Arithmetic mean for the expected return of the next period

Five Markets: Prices and Returns



- ▣ **Left:** prices on log scale, trending and non-stationary
- ▣ **Right:** daily log returns
 - ▶ Mean-reverting around zero
 - ▶ Calm and turbulent periods alternate
 - ▶ Extreme days cluster: 2008, March 2020
- ▣ EUR/RON moves more than 1% on only 1.4% of days
- ▣ Bitcoin, 12 March 2020: log return -46.5% , i.e. a price fall of $1 - e^{-0.465} = 37.2\%$

Recap: Prices and Returns

- Work with returns: scale-free and close to stationary
- Simple returns aggregate across assets; log returns aggregate over time
- Annualise with P and $\sqrt{P}$; use 365 for crypto
- CAGR (geometric) is below the arithmetic mean by about $\sigma^2/2$
- Use total-return or adjusted series when dividends matter

Mean, Variance and Volatility

- ▣ Sample of daily returns $r_1, \dots, r_T$
- ▣ **Sample mean**

$$\bar{r} = \frac{1}{T} \sum_{t=1}^T r_t$$

- ▣ **Sample variance and standard deviation**

$$\hat{\sigma}^2 = \frac{1}{T-1} \sum_{t=1}^T (r_t - \bar{r})^2, \quad \hat{\sigma} = \sqrt{\hat{\sigma}^2}$$

- ▣ Daily means are tiny (hundredths of a percent); daily standard deviations are 10–60 times larger in our data
- ▣ Consequence: the mean is estimated very imprecisely, the volatility much better

Skewness

- ▣ **Skewness** measures the asymmetry of a distribution

$$\hat{S} = \frac{1}{T} \sum_{t=1}^T \left(\frac{r_t - \bar{r}}{\hat{\sigma}} \right)^3$$

- ▣ Normal distribution: $S = 0$ (symmetric)
- ▣ $S < 0$ (negative skew): large losses are more frequent than equally large gains
- ▣ $S > 0$ (positive skew): large gains dominate, as for a lottery ticket
- ▣ Equity indices typically show negative skewness: crashes are sharper than rallies

Kurtosis and Excess Kurtosis

- ▣ **Kurtosis** measures the weight of the tails

$$\hat{K} = \frac{1}{T} \sum_{t=1}^T \left(\frac{r_t - \bar{r}}{\hat{\sigma}} \right)^4$$

- ▣ Normal distribution: $K = 3$; **excess kurtosis** = $K - 3$ (zero for the Normal distribution)
- ▣ $K > 3$ (**leptokurtosis**): more mass in the centre **and** in the tails than the Normal distribution
- ▣ Caution: $\hat{K}$ depends on the fourth power of returns
 - ▶ It is dominated by a few extreme days and converges very slowly
 - ▶ **Tail index** α : the exponent in $P(|X| > x) \approx c x^{-\alpha}$; if $\alpha < 4$ the population kurtosis does not exist

Autocorrelation and the ACF

- ▣ **Autocorrelation** at lag k : the correlation between r_t and r_{t-k}

$$\hat{\rho}_k = \frac{\sum_{t=k+1}^T (r_t - \bar{r})(r_{t-k} - \bar{r})}{\sum_{t=1}^T (r_t - \bar{r})^2}$$

- ▣ **ACF** = AutoCorrelation Function: $\hat{\rho}_k$ plotted against $k = 1, 2, \dots$
- ▣ Under i.i.d. returns, $\hat{\rho}_k \approx N(0, 1/T)$
 - ▶ 95% band: $\pm 1.96/\sqrt{T}$; bars outside the band are significant
 - ▶ The band assumes i.i.d. returns: with heavy tails and volatility clustering it is too narrow (same issue as for Ljung-Box)
- ▣ We compute the ACF of r_t (linear predictability) and of $|r_t|$ or r_t^2 (predictability of volatility)

Hypothesis Tests: A Reminder

- H_0 (null hypothesis): the property we test, e.g. “returns follow the Normal distribution”
- **Test statistic**: a number computed from the data, large when H_0 looks false
- **Critical value**: the threshold that the statistic exceeds with probability 5% if H_0 is true
- **p-value**: the probability, under H_0 , of a statistic at least as large as the observed one
 - ▶ Reject H_0 at the 5% level if p-value < 0.05
- χ_m^2 (chi-square with m degrees of freedom): the reference distribution of the three tests below
 - ▶ 5% critical values: $\chi_2^2 = 5.99$, $\chi_5^2 = 11.1$, $\chi_{10}^2 = 18.3$

The Jarque–Bera Test

- ▣ **JB** = Jarque–Bera test of normality, [Jarque & Bera \(1980\)](#)
- ▣ H_0 : returns follow the Normal distribution, i.e. skewness 0 and excess kurtosis 0

$$JB = \frac{T}{6} \left(\hat{S}^2 + \frac{(\hat{K} - 3)^2}{4} \right) \stackrel{H_0}{\sim} \chi^2_2$$

- ▣ Large skewness or excess kurtosis give a large JB
- ▣ With thousands of observations, even small deviations from normality are detected

Worked Example: JB for the S&P 500

- ▣ S&P 500, 1990–2026: $T = 9,245$, $\hat{S} = -0.364$, $\hat{K} - 3 = 10.88$
- ▣ Skewness term: $\hat{S}^2 = 0.132$
- ▣ Kurtosis term: $10.88^2/4 = 29.6$
- ▣ $JB = \frac{9245}{6} (0.132 + 29.6) = 1540.8 \times 29.73 \approx 45,800$
- ▣ Critical value 5.99: normality is rejected by a factor of about 7,600
- ▣ Almost all of the statistic comes from the kurtosis term: the problem is the tails, not the asymmetry

The Ljung–Box Test

- ▣ **Ljung–Box (LB)** test of no autocorrelation up to lag m , [Ljung & Box \(1978\)](#)
- ▣ $H_0: \rho_1 = \dots = \rho_m = 0$

$$Q(m) = T(T+2) \sum_{k=1}^m \frac{\hat{\rho}_k^2}{T-k} \stackrel{H_0}{\sim} \chi_m^2$$

- ▣ Applied to r_t : linear predictability of returns
- ▣ Applied to r_t^2 : predictability of volatility
- ▣ We use $m = 10$ lags (critical value 18.3)

The ARCH-LM Test

- ▣ **ARCH** = AutoRegressive Conditional Heteroskedasticity, Engle (1982)
 - ▶ Conditional heteroskedasticity: the variance of today's return depends on past returns
- ▣ **LM** = Lagrange Multiplier: the testing principle used
- ▣ Procedure with q lags (we use $q = 5$)
 - ▶ Regress $\hat{\varepsilon}_t^2$ on $\hat{\varepsilon}_{t-1}^2, \dots, \hat{\varepsilon}_{t-q}^2$, where $\hat{\varepsilon}_t = r_t - \bar{r}$
 - ▶ Statistic: $TR^2 \sim \chi_q^2$ under H_0 of no ARCH effects
- ▣ A large statistic means that volatility is predictable from past squared returns

Summary Statistics: Five Markets

	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Observations	9245	2992	4384	5352	9559
Mean (% per day)	0.033	0.077	0.118	0.007	0.025
Std. dev. (% per day)	1.13	0.98	3.49	0.28	1.01
Annualised volatility (%)	18.0	15.5	66.6	4.4	16.3
Skewness	-0.36	-1.40	-0.70	0.73	-0.23
Excess kurtosis	10.9	17.8	12.0	15.1	9.3
Min / max log return (%)	-12.8 / 11.0	-11.9 / 6.7	-46.5 / 22.5	-2.5 / 2.9	-11.1 / 10.4
Jarque-Bera	45805	40433	26482	51575	34732
Ljung-Box $Q(10), r_t$	91.6	70.0	20.2	221.7	32.1
Ljung-Box $Q(10), r_t^2$	8018	714	213	2420	1049
ARCH-LM (5 lags)	2229	371	115	871	483

□ Min and max are daily **log** returns; EUR/RON from 1 July 2005 (new leu, RON)

□ 5% critical values: $\chi^2_2 = 5.99$, $\chi^2_{10} = 18.3$, $\chi^2_5 = 11.1$; all p-values < 0.03

Reading the Table

- ▣ **Normality is rejected everywhere:** JB statistics in the tens of thousands versus a critical value of 6
- ▣ **Heavy tails:** excess kurtosis between 9.3 (gold) and 17.8 (BET-TR)
- ▣ **Negative skewness** for equity, crypto and gold; strongest for BET-TR (-1.40)
- ▣ **Volatility is predictable:** $Q(10)$ on r_t^2 is 10–90 times larger than on r_t ; ARCH-LM rejects everywhere
- ▣ **Linear autocorrelation** of r_t : statistically significant but economically small
 - ▶ Large samples make tiny correlations significant; heavy tails and clustering distort the standard test
- ▣ EUR/RON is special: tiny volatility, high kurtosis, positive skewness, strong autocorrelation (Section 5)

Recap: Descriptive Statistics and Tests

- ▣ Skewness and excess kurtosis summarise asymmetry and tail weight
- ▣ JB tests normality; Ljung–Box tests autocorrelation; ARCH-LM tests predictable volatility
- ▣ All five markets reject normality and constant variance
- ▣ Statistical significance is not economic significance in samples of thousands of days

The People Behind the Stylised Facts (1/2)



Carl Friedrich Gauss (1777–1855)

Painting: C.A. Jensen (1840); public domain;
Wikimedia Commons

- The Normal distribution: the benchmark that returns violate



Louis Bachelier (1870–1946)

Photo: unknown author; public domain;
Wikimedia Commons

- Prices as a random walk, in his thesis Bachelier (1900)




William S. Gosset, “Student” (1876–1937)

Photo: unknown author; public domain;
Wikimedia Commons

- The Student- t distribution Student (1908), today a model for heavy tails

The People Behind the Stylised Facts (2/2)



Vilfredo Pareto (1848–1923)

Photo: unknown author (1870s); public domain; [Wikimedia Commons](#)



Benoît Mandelbrot (1924–2010)

Photo: Rama (2007); CC BY-SA 2.0 FR; [Wikimedia Commons](#)

- Pareto: **power laws** for large incomes, the ancestor of heavy-tail models
- Mandelbrot: cotton prices have heavy tails and the Normal-distribution model fails [Mandelbrot \(1963\)](#)
- Their tail index α is what the Hill estimator in this chapter measures


The Stylised Facts of Cont (2001)

1. Absence of linear autocorrelation
 2. Heavy tails
 3. Gain/loss asymmetry
 4. Aggregational Gaussianity
 5. Intermittency
 6. Volatility clustering
 7. Conditional heavy tails
 8. Slow decay of autocorrelation in absolute returns
 9. Leverage effect
 10. Volume/volatility correlation
 11. Asymmetry in time scales
- We verify most of them on the five markets, with data up to September 2026
 - Early evidence: Mandelbrot (1963), Fama (1965); textbook treatment: Taylor (2008), Franke, Härdle & Hafner (2019)

Facts 1 and 6: Uncorrelated but Not Independent

- **Fact 1 — absence of autocorrelation:** $\text{corr}(r_t, r_{t-k}) \approx 0$ for $k \geq 1$
 - ▶ Past returns do not predict the direction of future returns linearly
- **Fact 6 — volatility clustering:** large moves tend to be followed by large moves, of either sign
 - ▶ $\text{corr}(|r_t|, |r_{t-k}|) > 0$ for many lags
- Together: returns are **uncorrelated but not independent**
- Independence would require zero correlation of **every** function of returns, including $|r_t|$

Facts 1 and 6 on the S&P 500



- ▣ S&P 500 1990–2026, lags 1–100; grey band: $\pm 1.96/\sqrt{T}$
- ▣ Returns: almost all autocorrelations inside the band (lag 1: -0.08)
- ▣ Absolute and squared returns: large, positive and persistent (lag 1: 0.27 for $|r_t|$)


Why Facts 1 and 6 Matter for Modelling

- Fact 1: simple linear forecasts of returns (e.g. AR models) have little value
 - ▶ AR = autoregressive model: r_t regressed on its own past values
- Fact 6: the **variance** must be modelled as time-varying
 - ▶ $r_t = \sigma_t \varepsilon_t$ with σ_t depending on past returns
 - ▶ This is the idea behind GARCH models (Chapter 5)
- Risk forecasts (margins, capital requirements) must be updated as volatility changes

Fact 2: Heavy Tails

- ▣ **Heavy (fat) tails:** extreme returns occur far more often than under a Normal distribution with the same variance
- ▣ Diagnostic tools
 - ▶ Excess kurtosis above zero
 - ▶ Histogram compared with the density of the Normal distribution
 - ▶ **QQ-plot** (quantile–quantile plot): empirical quantiles against quantiles of the Normal distribution
- ▣ Reading a QQ-plot: points on the diagonal = Normal distribution; points bending away in the ends = heavy tails
- ▣ **Quantile** q_α : the value below which a fraction α of the data falls

Heavy Tails: Histograms and QQ-Plots



- Standardised returns $(r_t - \bar{r})/\hat{\sigma}$: higher peak and fatter tails than the Normal distribution
- A Student- t fits much better (next slides); QQ-plots bend away from the line in both tails

How Heavy Are the Tails?

	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Days beyond $3\hat{\sigma}$ (%)	1.55	1.44	1.82	2.06	1.42
Days beyond $5\hat{\sigma}$ (count)	30	12	10	29	26
Expected under the Normal distribution, $> 5\hat{\sigma}$	0.005	0.002	0.003	0.003	0.005
Student- t degrees of freedom $\hat{\nu}$ (MLE fit)	2.8	3.1	2.2	1.2	2.9

- Normal distribution: 0.27% of days beyond 3σ ; observed 5–8 times more
- Normal distribution: one 5σ day every 6,900 years of trading; observed 10–30 in 12–36 years
- $\hat{\nu}$: Student- t degrees of freedom (next slide), estimated by **maximum likelihood**: the value that makes the observed returns most probable

The Student- t Distribution

- Density with ν degrees of freedom, location μ and scale s

$$f(x) = \frac{\Gamma(\frac{\nu+1}{2})}{\Gamma(\frac{\nu}{2}) \sqrt{\nu\pi} s} \left(1 + \frac{(x - \mu)^2}{\nu s^2}\right)^{-\frac{\nu+1}{2}}$$

- ▶ $\Gamma(\cdot)$: the gamma function, a continuous extension of the factorial
- ν controls the tails: small ν = heavy tails
- $\nu \rightarrow \infty$ gives the Normal distribution; $\nu = 1$ gives the Cauchy distribution
- Student- t as a model of returns: [Praetz \(1972\)](#)

Student- t : Tails, Moments and Our Estimates

- ▣ Tails decay as a **power law**: $P(X > x) \propto x^{-\nu}$ for large x
 - ▶ Much slower than the tail of the Normal distribution, which decays like $e^{-x^2/2}$
- ▣ Moments exist only up to order $< \nu$
 - ▶ Variance $s^2\nu/(\nu - 2)$ for $\nu > 2$; excess kurtosis $6/(\nu - 4)$ for $\nu > 4$
- ▣ Our fits give $\hat{\nu} \approx 1.2\text{--}3.1$, but $\hat{\nu}$ is **not** a tail-index estimator
 - ▶ The likelihood is dominated by the many small returns in the centre, not by the tail
 - ▶ EUR/RON $\hat{\nu} = 1.2$ is an artefact of the fixing: 57% of daily changes are below 0.1%, plus rare jumps
- ▣ The tail itself is estimated with the **Hill estimator** (next slides)
- ▣ Power-law tails were first argued by [Mandelbrot \(1963\)](#) and [Fama \(1965\)](#)


The Hill Estimator of the Tail Index

- ▣ **Tail index** α : $P(|X| > x) \approx c x^{-\alpha}$ for large x ; moments of order $< \alpha$ exist
 - ▶ $\alpha > 2$: finite variance; $\alpha > 4$: finite kurtosis
- ▣ **Hill estimator**, [Hill \(1975\)](#): order the absolute returns $X_{(1)} \geq X_{(2)} \geq \dots \geq X_{(n)}$ and use the k largest

$$\hat{\alpha}_k = \left[\frac{1}{k} \sum_{i=1}^k \ln X_{(i)} - \ln X_{(k+1)} \right]^{-1}$$

- ▣ **Choice of k** : a bias–variance trade-off
 - ▶ Small k : few observations, noisy; large k : the centre of the distribution enters, biased
 - ▶ **Hill plot**: $\hat{\alpha}_k$ against k ; read the level where the curve is flat (here $k = 1\text{--}5\%$ of n)

Hill Plots for Five Markets



- $\hat{\alpha}$ for $|r_t|$ at $k = 2.5\%$ of n : S&P 500 3.1, BET-TR 2.7, Bitcoin 4.2, EUR/RON 2.8, gold 3.3
- For $k = 1-5\%$: $\hat{\alpha}$ between 2.3 and 3.5 except Bitcoin (3.1–4.2); loss tails heavier for BET-TR (2.2 vs 3.1) and Bitcoin (2.7 vs 3.4)
- Variance exists, but the **fourth moment very likely does not**: sample kurtosis never settles; tails are modelled with extreme value theory (Chapter 7)


Why Heavy Tails Matter for Modelling

- ▣ **VaR** = Value-at-Risk: the loss exceeded with a small probability (e.g. 1%) over a horizon
 - ▶ It is a quantile of the loss distribution (Chapter 7)
- ▣ A Normal-distribution model with the right variance **underestimates** extreme quantiles
- ▣ Stress scenarios based on “ 3σ events” are far too optimistic
- ▣ Estimators based on high moments (kurtosis) are unstable: prefer quantiles and tail counts
- ▣ Remedies: Student- t or skewed distributions, extreme value theory (Chapter 7)

Fact 3: Gain/Loss Asymmetry

- ▣ **Statement:** large losses are more frequent or larger than gains of the same probability
- ▣ Tools
 - ▶ **Survival function:** $P(X > x)$, the probability of a move larger than x
 - ▶ Compare it for losses ($-r_t$) and gains (r_t) on a log-log scale
 - ▶ Compare the 1% quantile ($q_{0.01}$, losses) with the 99% quantile ($q_{0.99}$, gains)
- ▣ Holds for equity indices; not for all asset classes

Gain/Loss Asymmetry in the Data




- The 1% loss quantile exceeds the 99% gain quantile for 4 of 5 markets
 - S&P 500: 3.15% vs 2.97%; BET-TR: 2.73% vs 2.40%; Bitcoin: 10.5% vs 9.9%
- EUR/RON is the exception: leu depreciations (up-moves) are larger (0.94% vs 0.76%)

Fact 4: Aggregational Gaussianity

- **Statement:** as the horizon grows, returns get closer and closer to the Normal distribution
- h -day log return: $r_t(h) = r_t + r_{t-1} + \cdots + r_{t-h+1}$
 - ▶ We use non-overlapping blocks, so that the h -day returns do not share days
- Mechanism: the CLT applies to sums of many daily returns
- Convergence is slow because daily returns are dependent (volatility clustering)

Aggregational Gaussianity in the Data



- Excess kurtosis of h -day returns, $h = 1, \dots, 60$
- Tails thin out: S&P 500 from 10.9 (daily) to 1.6 ($h = 60$), Bitcoin from 12.0 to 0.1
- $h = 60$ is a quarter for the S&P 500 (252 days a year) but two months for Bitcoin (365 days a year)
- At $h = 60$ only 154 (S&P 500) and 73 (Bitcoin) non-overlapping returns remain: approximate standard error $\sqrt{24/n} = 0.39$ and 0.57 under normality, so these values are very imprecise

Why Heavy Tails? Mixtures of Normals

- ▣ Suppose $r_t = \sigma_t \varepsilon_t$ with $\varepsilon_t \sim N(0, 1)$ independent of a random volatility σ_t

$$K = \frac{\mathbb{E}[r_t^4]}{\mathbb{E}[r_t^2]^2} = 3 \frac{\mathbb{E}[\sigma_t^4]}{\mathbb{E}[\sigma_t^2]^2} \geq 3$$

- ▣ Jensen's inequality ($\mathbb{E}[X^2] \geq \mathbb{E}[X]^2$) with $X = \sigma^2$: any variation in volatility creates excess kurtosis
- ▣ Example: $\sigma = 1\%$ on 90% of days and 3% on 10% of days

$$K = 3 \cdot \frac{0.9 \cdot 1 + 0.1 \cdot 81}{(0.9 + 0.1 \cdot 9)^2} = 3 \cdot \frac{9.0}{3.24} = 8.33$$

- ▣ This links heavy tails (Fact 2) to volatility clustering (Fact 6)


Fact 5: Intermittency

- ▣ **Statement:** returns show bursts of high volatility separated by calm periods (regimes)
- ▣ **Tool:** **rolling volatility** over a window of n days

$$\hat{\sigma}_t^{(n)} = \sqrt{P} \sqrt{\frac{1}{n-1} \sum_{j=0}^{n-1} \left(r_{t-j} - \bar{r}_t^{(n)} \right)^2}$$

- ▶ $\bar{r}_t^{(n)}$: mean over the same window; we use $n = 63$ days (one quarter)
- ▣ Constant-volatility models cannot produce such bursts

Intermittency in the Data




- 63-day rolling volatility, annualised, log scale
- Common spikes across markets: 2008–2009, March 2020
- EUR/RON volatility collapsed after 2012: an exchange-rate regime managed by the central bank

Fact 8: Slow Decay — Long Memory in Volatility

- ▣ **Statement:** the ACF of $|r_t|$ stays positive for months
- ▣ **Short memory:** autocorrelations decay **exponentially**, $\rho_k \approx c \phi^k$ with $|\phi| < 1$
 - ▶ Example: GARCH(1,1) volatility (Chapter 5)
- ▣ **Long memory:** autocorrelations decay **hyperbolically**, $\rho_k \approx c k^{-\beta}$ with $0 < \beta < 1$
 - ▶ Shocks to volatility die out very slowly
- ▣ Evidence: [Ding, Granger & Engle \(1993\)](#)

Long Memory in the Data



- ACF of $|r_t|$ up to one year (250 lags), log scale on the lag axis
- S&P 500: 0.27 at lag 1, 0.22 at lag 20, still 0.09 at lag 100
- BET-TR decays fastest (shorter sample); decay is slower than the exponential decay of GARCH(1,1)

Infinite Fourth Moments and Statistics on r_t^2

- ▣ The ACF of r_t^2 is a ratio of fourth-order moments: it needs $E[r_t^4] < \infty$, i.e. tail index $\alpha > 4$
- ▣ Our Hill estimates are about 2.5–3.5: the fourth moment is very likely infinite
 - ▶ The sample ACF of r_t^2 then converges slowly, to a limit that is not the Normal distribution; the $\pm 1.96/\sqrt{T}$ band is not valid, [Davis & Mikosch \(1998\)](#)
 - ▶ Ljung–Box on r_t^2 and ARCH-LM inherit the problem: the statistics are huge, but their p-values are not reliable
- ▣ The conclusion “volatility clusters” is robust; the exact magnitudes are not
- ▣ Remedy: use $|r_t|$, whose ACF needs only $\alpha > 2$ — this is one reason behind the Taylor effect (next slide)

The Taylor Effect

- **Definition:** compute the autocorrelation of $|r_t|^\delta$ for different powers $\delta > 0$
 - ▶ $\delta = 1$: absolute returns; $\delta = 2$: squared returns
- **Taylor effect:** the autocorrelation is highest for δ around 1, not for $\delta = 2$
- Implication: absolute returns are a more persistent, more robust measure of volatility than squared returns
- Documented by Taylor (1986); the name “Taylor effect” was given by Granger & Ding (1995); see also Ding, Granger & Engle (1993)

The Taylor Effect in the Data



- S&P 500: autocorrelation of $|r_t|^\delta$ as a function of δ
- Maximum at $\delta = 1.75$ (lag 1), 1.5 (lag 5) and 1.25 (lag 20): the optimum moves towards 1 at longer lags
- Motivates power-GARCH models and absolute-value volatility measures


Fact 9: The Leverage Effect

- ▣ **Statement:** falling prices raise **future** volatility more than rising prices
- ▣ **Leverage correlation** at lag k

$$L(k) = \text{corr}(r_t, |r_{t+k}|)$$

- ▶ $L(k) < 0$ for $k > 0$: a negative return today is followed by larger moves
- ▶ $L(k) \approx 0$ for $k < 0$: past volatility does not predict today's sign
- ▣ Explanations: financial leverage of firms, [Christie \(1982\)](#); the retarded volatility model, [Bouchaud, Matacz & Potters \(2001\)](#)

The Leverage Effect in the Data




- S&P 500: $L(1) = -0.12$, still -0.07 at $k = 10$; no effect for $k < 0$
- BET-TR -0.12 ; weak for Bitcoin (-0.07), negligible for gold (-0.02), reversed for EUR/RON ($+0.10$)
- Modelling: asymmetric volatility models, GJR-GARCH (Glosten–Jagannathan–Runkle) and EGARCH (Exponential GARCH), Chapter 5

Fact 10: Volume and Volatility Move Together

- **Statement:** days with large price moves are days with high trading volume
- **Abnormal log volume:** $\ln V_t$ minus its 250-day moving average
- **Spearman rank correlation:** the correlation between the ranks of two variables
 - ▶ Robust to outliers and to non-linear but monotone relations
- Mixture-of-distributions hypothesis, [Clark \(1973\)](#): prices move in a random “business time” driven by information flow, which also drives volume
- Survey of the price–volume relation: [Karpoff \(1987\)](#)

Volume and Volatility in the Data



- Abnormal log volume against $|r_t|$
- Spearman rank correlation: 0.25 for the S&P 500 and 0.39 for Bitcoin
- Volume is a useful additional input for volatility forecasts
- Bitcoin caveat: volume is exchange-reported and aggregated over many venues; part of it is **wash trading** (an exchange trading with itself to inflate volume), [Cong, Li, Tang & Yang \(2023\)](#)

Facts 7 and 11: Conditional Tails and Time Scales

Fact 7: conditional heavy tails

- After removing volatility clustering, standardised returns are still heavy-tailed
- Fit a volatility model, e.g. GARCH (Chapter 5), get $\hat{\sigma}_t$, compute $\hat{z}_t = r_t / \hat{\sigma}_t$
- Kurtosis of $\hat{z}_t$: lower than for r_t , still above 3
- Models need Student- t or skewed innovations

- Both facts are documented in [Cont \(2001\)](#)

Fact 11: asymmetry in time scales

- Weekly volatility predicts daily volatility better than the reverse
- Traders with different horizons react at different frequencies
- **HAR-RV**: Heterogeneous AutoRegressive model of Realised Volatility (Chapter 9)

Recap: The Stylised Facts

- ▣ Uncorrelated returns, strongly dependent volatility (Facts 1, 6, 8)
- ▣ Heavy tails, also conditionally, that thin out slowly with aggregation (Facts 2, 4, 7)
- ▣ Crashes sharper than rallies in equity markets (Facts 3, 9)
- ▣ Volatility in regimes, linked to volume and information flow (Facts 5, 10)
- ▣ Different horizons interact asymmetrically (Fact 11)

Crypto-Assets: The Same Facts, Amplified

Bitcoin vs S&P 500

- ▣ Volatility 66.6% vs 18.0% per year
- ▣ Excess kurtosis 12.0 vs 10.9
- ▣ Skewness -0.70 vs -0.36
- ▣ Worst day, log return: -46.5% (12 March 2020) vs -12.8% (16 March 2020)
- ▣ As price changes: -37.2% vs -12.0%

But some facts are weaker

- ▣ Leverage effect: weak, short-lived
- ▣ Volatility memory: ACF of $|r_t|$ decays faster (lag 100: 0.02 vs 0.09)
- ▣ Trading 24/7: 365 periods per year

Crypto-Assets: An Asset Class of Their Own

- Crypto as an asset class with its own statistical profile: [Härdle, Harvey & Reule \(2020\)](#)
- Tail, memory and moment factors classify cryptos as a separate asset class, mainly through the tail factor: [Pele et al. \(2023\)](#)
- Crypto modelling and the CRIX index in quantitative finance: [Härdle, Chen & Overbeck \(2017\)](#)
 - ▶ CRIX = CRypto IndeX, a benchmark index of the crypto market
- Deeper treatment in Chapter 16 (digital assets)

EUR/RON: A Managed Exchange Rate

- ▣ **Managed float:** the exchange rate is market-determined, but the central bank intervenes to limit fluctuations
- ▣ BNR reference rate, 1 July 2005–2026: from the new leu (RON, 1 RON = 10,000 old lei), as in Chapter 0
 - ▶ **Low volatility** (4.4% per year) but **high kurtosis** (15.1)
 - ▶ Long calm periods interrupted by sudden adjustments
 - ▶ Largest daily moves: October 2008 (+2.9% on 3 October, −2.5% on 9 October)
- ▣ **Positive** lag-1 autocorrelation (0.17), unlike equities
 - ▶ The central bank smooths adjustments over several days; a daily fixing, not a traded price

EUR/RON: Lessons

- Stylised facts are **market-specific**: always check the institutional setting
- Low volatility does not mean low tail risk: kurtosis is the second highest of the five series
- Data source matters
 - ▶ Exchange quote series showed spurious negative autocorrelation (quote errors)
 - ▶ Official reference rates avoid these errors

Romania: The BET Total Return Index

- ▣ BET-TR, 2014–2026, 2,992 days
- ▣ Facts confirmed: excess kurtosis 17.8, clustering ($Q(10)$ on r_t^2 : 714), 1% loss quantile 2.7%
- ▣ Strongest negative skewness of the five series (-1.40)
 - ▶ Worst day (log return) -11.9% on 19 December 2018: reaction to the bank-asset tax announced by the government on 18 December
 - ▶ The tax was adopted on 28 December 2018 as Government Emergency Ordinance (GEO) 114/2018
 - ▶ Policy shocks produce single-day crashes without comparable rallies

Romania: Autocorrelation in a Smaller Market

- ▣ Positive lag-1 autocorrelation (0.06), in contrast with the S&P 500 (-0.08)
- ▣ **Non-synchronous trading:** index constituents do not all trade at every moment
 - ▶ The index reflects news with a delay, creating positive autocorrelation, [Campbell, Lo & MacKinlay \(1997\)](#), ch. 3
- ▣ Frontier and emerging markets (below developed status in index-provider classifications)
 - ▶ Lower liquidity (it is costlier to trade quickly)
 - ▶ Stronger autocorrelation and policy-driven jumps

Why Use Open, High, Low and Close?

- **Close-to-close** estimator: the standard deviation of daily close-to-close log returns
 - ▶ Uses one number per day and ignores intraday information
- **Range-based** estimators use the OHLC prices
 - ▶ Notation: o_t, h_t, l_t, c_t are the log open, high, low and close of day t
 - ▶ The range $h_t - l_t$ is large on volatile days even if the close equals the open
- **Efficiency**: the ratio of variances of two estimators; higher efficiency = less noise for the same sample

Parkinson and Garman–Klass Estimators

- ▣ **Parkinson**, [Parkinson \(1980\)](#): uses only the daily range

$$\hat{\sigma}_P^2 = \frac{1}{4 \ln 2} \frac{1}{n} \sum_{t=1}^n (h_t - l_t)^2$$

- ▣ **Garman–Klass**, [Garman & Klass \(1980\)](#): combines the range with the open-to-close move

$$\hat{\sigma}_{GK}^2 = \frac{1}{n} \sum_{t=1}^n \left[\frac{1}{2} (h_t - l_t)^2 - (2 \ln 2 - 1) (c_t - o_t)^2 \right]$$

- ▣ Derived for a Brownian motion without drift and without overnight jumps
 - ▶ Brownian motion: the continuous-time random walk (Chapter 11)
 - ▶ Efficiency relative to close-to-close: about 5 (Parkinson) and 7–8 (Garman–Klass)

Rogers–Satchell and Yang–Zhang Estimators

- **Rogers–Satchell**, Rogers & Satchell (1991): unbiased under a non-zero **drift** (expected trend of the log price)



$$\hat{\sigma}_{RS}^2 = \frac{1}{n} \sum_{t=1}^n [(h_t - o_t)(h_t - c_t) + (l_t - o_t)(l_t - c_t)]$$

- **Yang–Zhang**, Yang & Zhang (2000): adds the overnight (close-to-open) jump

$$\hat{\sigma}_{YZ}^2 = \hat{\sigma}_{\text{overnight}}^2 + k \hat{\sigma}_{\text{open-close}}^2 + (1 - k) \hat{\sigma}_{RS}^2, \quad k = \frac{0.34}{1.34 + \frac{n+1}{n-1}}$$

- ▶ $\hat{\sigma}_{\text{overnight}}^2$: variance of $o_t - c_{t-1}$; $\hat{\sigma}_{\text{open-close}}^2$: variance of $c_t - o_t$
- ▶ Minimum-variance combination among drift-independent estimators robust to opening jumps
- Range estimators are the bridge to realised volatility from intraday data (Chapter 9)

Estimators on the S&P 500



- 21-day windows, 2008–2026; averages: close-to-close 16.4%, Yang–Zhang 13.9%, Parkinson 13.0%, Garman–Klass 12.4%, Rogers–Satchell 12.3%
- Pure range estimators are lower: they miss overnight gaps
- Yang–Zhang includes the overnight gap but stays 2.5 pp below close-to-close: the index open is partly stale (next slide)
- Range-based estimates are 3–4 times smoother (less noise in daily updates)

Data Quality: Stale Opens of the S&P 500 Index

- ▣ The index **open** is computed from the constituents' first prices; at 9:30 many stocks have not traded yet, so their previous close enters the open
- ▣ Open = previous close: about 96% of days until 2005, 3.5% over 2008–2026 (up to 26% in 2012–2013); on other days part of the open is still stale
- ▣ Comparison 2008–2026 with the SPY ETF, whose open is a traded price (annualised, %)

	Close-to-close	Yang–Zhang	Overnight	Open-to-close
S&P 500 index	16.4	13.9	6.7	17.3
SPY ETF	16.2	16.8	12.0	15.2

- ▣ A stale open moves part of the overnight move into the day: the index overnight volatility is understated and Yang–Zhang falls below close-to-close
- ▣ Lesson: for OHLC estimators use a traded instrument (ETF, futures), not the index


Drawdowns: Reminder and Focus

- Drawdown DD_t and maximum drawdown MDD were defined in Chapter 0 (section Key Measures)
- Here we add two things
 - ▶ A comparison of depth across the five markets
 - ▶ **Recovery duration**: the time from the peak until the price regains it
- Investors feel drawdowns more than volatility: a -50% loss needs $+100\%$ to recover

Worked Example: The S&P 500 in 2007–2009

- ▣ Peak: 9 October 2007, close 1,565.15
- ▣ Trough: 9 March 2009, close 676.53
- ▣ Maximum drawdown: $676.53/1565.15 - 1 = -56.8\%$
- ▣ Gain needed to recover: $1565.15/676.53 - 1 = +131\%$
- ▣ The price index regained the 2007 peak only on 28 March 2013, five and a half years after the peak

Drawdowns Across Markets



- Maximum drawdowns: Bitcoin -83%, S&P 500 -57%, gold -45%, BET-TR -31% (March 2020), EUR/RON -16%
- Depth and duration matter together

How Long Does Recovery Take?

	MDD	Peak	Trough	Peak regained
S&P 500	−57%	Oct 2007	Mar 2009	Mar 2013 (5.5 years)
BET-TR	−31%	Jan 2020	Mar 2020	Dec 2020 (0.9 years)
Bitcoin	−83%	Dec 2017	Dec 2018	Nov 2020 (3.0 years)
EUR/RON	−16%	Dec 2005	Jul 2007	Jan 2008 (2.0 years)
Gold	−45%	Sep 2011	Dec 2015	Jul 2020 (8.9 years)

- Depth and duration are different risks: gold's drawdown was shallower than Bitcoin's but lasted three times longer
- Fast V-shaped recoveries (BET-TR 2020) versus long underwater periods (gold 2011–2020)
- For EUR/RON a “drawdown” is a fall of the euro against the leu

The Sharpe Ratio

- ▣ **Sharpe ratio** (SR), [Sharpe \(1966\)](#); [Sharpe \(1994\)](#): excess return per unit of volatility

$$SR = \frac{\mu_{\text{ann}} - r_f}{\sigma_{\text{ann}}}$$

- ▶ r_f : risk-free rate; here $r_f = 0$ for comparability across markets
- ▣ **Worked example**, S&P 500 1990–2026
 - ▶ $\mu_{\text{ann}} = 252 \times 0.0331\% = 8.33\%$ (mean log return)
 - ▶ $\sigma_{\text{ann}} = 17.99\%$
 - ▶ $SR = 8.33/17.99 = 0.46$

The Sortino and Calmar Ratios

- **Sortino ratio**, [Sortino & van der Meer \(1991\)](#): penalises only downside deviation

$$SoR = \frac{\mu_{\text{ann}} - r_f}{\sigma_{\text{ann}}^-}, \quad \sigma^- = \sqrt{\mathbb{E}[\min(r_t, 0)^2]}$$

- ▶ σ^- : **downside deviation**, computed from negative returns only
- ▶ S&P 500: $\sigma_{\text{ann}}^- = 12.90\%$, so $SoR = 8.33/12.90 = 0.65$

- **Calmar ratio**: CAGR divided by the absolute maximum drawdown


$$\text{Calmar} = \frac{\text{CAGR}}{|MDD|}$$

- ▶ S&P 500: $8.69\%/56.78\% = 0.15$; very sensitive to one extreme episode

Caveats for Performance Measures

- All three ratios are estimates with large sampling error
- Heavy tails and negative skewness make the Sharpe ratio flattering
 - ▶ The Deflated Sharpe Ratio corrects for this and for multiple testing (Chapter 13)
- Different sample periods make comparisons across markets fragile
- Currency and interest rates: BET-TR is in RON, where short-term rates were well above USD rates; with $r_f = 0$ its Sharpe ratio is inflated
- Always report the ratios together with drawdowns and tail statistics

Risk and Return Across Markets



	CAGR	Sortino	Calmar
S&P 500	8.7%	0.65	0.15
BET-TR	21.2%	1.72	0.68
Bitcoin	53.9%	0.91	0.65
EUR/RON	1.8%	0.62	0.12
Gold	6.7%	0.56	0.15

- Different sample periods: not a horse race
- S&P 500 without dividends; BET-TR with dividends
- Higher volatility did not always mean higher Sharpe

Stylised Facts in One Table

Fact	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Heavy tails (JB rejects)	yes	yes	yes	yes	yes
Negative skewness	yes	strong	yes	no	yes
Volatility clustering (ARCH-LM)	yes	yes	yes	yes	yes
Slow decay of ACF $ r_t $	yes	weaker	weaker	yes	yes
Near-zero ACF of r_t	yes	small +	yes	no (+0.17)	yes
Leverage effect ($k = 1$)	-0.12	-0.12	-0.07	+0.10	-0.02

- The core facts (heavy tails, clustering, long memory in volatility) are universal
- Autocorrelation and asymmetry depend on market microstructure and institutions

What a Good Model Must Reproduce

Requirements

- ▣ Time-varying, persistent volatility
- ▣ Heavy-tailed innovations, also after conditioning on volatility
- ▣ Asymmetric response to bad news
- ▣ Approximately uncorrelated returns

- ▣ The random walk with i.i.d. increments following the Normal distribution (Bachelier–Black–Scholes world) violates Facts 2, 3 and 5–11; Fact 4 holds trivially (sums of Normally distributed variables follow the Normal distribution)

Where in this course

- ▣ Ch. 5: GARCH; Ch. 6: multivariate volatility and dependence
- ▣ Ch. 7–8: VaR, Expected Shortfall, extreme value theory; backtesting
- ▣ Ch. 9: realised volatility from intraday data
- ▣ Ch. 13: machine learning must respect the low signal-to-noise ratio

Practical Checklist for Empirical Work

- ▣ **Know your data**
 - ▶ Price vs total-return index; adjusted vs raw prices; fixing vs traded price
 - ▶ Official reference rates where possible (e.g. EUR/RON)
- ▣ **Clean before you model**
 - ▶ Non-trading days (zero volume), holidays, stale quotes, obvious errors
 - ▶ Check extremes: are the largest moves real events?
- ▣ **Use log returns** for single-asset dynamics, simple returns for portfolios
- ▣ **Report robust statistics:** quantiles and tail counts next to kurtosis
- ▣ **Annualise carefully:** 252 vs 365 periods; $\sqrt{P}$ fails under dependence
- ▣ **Reproducibility:** code and data on GitHub, as Quantlets

Key Formulas

Returns

- ▣ $R_t = P_t/P_{t-1} - 1, \quad r_t = \ln(1 + R_t)$
- ▣ $r_t(k) = \sum_{j=0}^{k-1} r_{t-j}$
- ▣ $R_{p,t} = \sum_i w_i R_{i,t}$
- ▣ $\text{CAGR} = (P_T/P_0)^{1/Y} - 1$

Tests

- ▣ $JB = \frac{T}{6} (\hat{S}^2 + (\hat{K} - 3)^2/4)$
- ▣ $Q(m) = T(T+2) \sum_{k=1}^m \hat{\rho}_k^2 / (T-k)$

Volatility

- ▣ $\sigma_{\text{ann}} = \sqrt{P} \hat{\sigma}$
- ▣ Parkinson: $\frac{1}{4 \ln 2} \overline{(h-l)^2}$
- ▣ Mixture: $K = 3 \mathbb{E}[\sigma^4] / \mathbb{E}[\sigma^2]^2$

Performance

- ▣ $SR = (\mu - r_f)/\sigma, \quad \text{SoR} = (\mu - r_f)/\sigma^-$
- ▣ $DD_t = P_t / \max_{s \leq t} P_s - 1$
- ▣ Calmar = $\text{CAGR} / |\text{MDD}|$

An Open Question: Are Tail Indices Stable over Time?

- The stylised facts are usually stated for a whole sample; the tail index α is treated as a constant
- Crypto-assets changed in structure: futures (2017), spot Bitcoin ETFs (January 2024), institutional holders
 - ▶ [Pele et al. \(2023\)](#): the tail factor is what separates cryptos from other asset classes
- Question: did institutionalisation make Bitcoin's tails thinner, i.e. raise α ?
- Inference is delicate: $\hat{\alpha}$ uses only the k largest observations, and returns are dependent
 - ▶ [Hoga \(2017\)](#): change-point tests for the tail index valid under β -mixing (e.g. GARCH-type) dependence

The Discovery Loop, Step by Step (1/2)

□ Literature

- ▶ Find peer-reviewed work since 2015 on time variation or structural breaks in tail indices of asset returns, with tests valid under GARCH-type dependence. For each: DOI, estimator, test, data. Flag entries you are unsure of.
- ▶ Resolve every DOI on Crossref; the title and authors must match

□ Hypothesis

- ▶ State competing hypotheses for a change in Bitcoin's tail after the 2024 spot ETFs: a change in tail shape, a fall in volatility only, sampling noise. What does each predict for the Hill index of raw returns and of GARCH-standardised returns?

□ Data and code

- ▶ Read BTC-USD.CC from data/market (close); log returns on its own calendar; Hill index of $|r|$ with $k = 2.5\%$ of n before and after 11 January 2024; moving-block bootstrap standard errors.
- ▶ Replicate a known number first: the chapter's Hill estimates, Bitcoin 4.2 and S&P 500 3.1

The Discovery Loop, Step by Step (2/2)

▣ Robustness

- ▶ Write a pre-registered robustness plan: k from 1% to 5% of n , loss and gain tails separately, GARCH-standardised residuals, block lengths from 10 to 50 days, a placebo date and a placebo asset.

▣ Critique

- ▶ Referee this design. With 24 tail observations after the break, what is the power against a change of 1 in the tail index? Which dependence assumption does each standard error need?

▣ Report

- ▶ $\hat{\alpha}$ in each window, both standard errors, the placebo, the minimum detectable change
- ▶ Prompts that shaped a decision go into AI_USE.md

- ▣ Scale versus shape: a lower volatility leaves α unchanged; only a change in the shape of the tail moves it

What the Human Must Check

- **Wrong return or price series:** prices, simple returns or a different calendar
 - ▶ test: reproduce the chapter's Hill table before anything else
- **Choice of k after seeing the result:** fix the rule ($k = 2.5\%$ of n) in advance; show the Hill plot for 1–5%
- **i.i.d. standard errors:** $\hat{\alpha}/\sqrt{k}$ ignores volatility clustering
 - ▶ test: compare with a block bootstrap; for the S&P 500 before 2024 it is 2.2 times larger
- **Break date chosen from the data:** use a date fixed by the event, or a sup-type change-point test with its own critical values ([Hoga \(2017\)](#))
- **Bad ticks create fake tails:** list the ten largest $|r_t|$ and match each to a dated event
- **Low power:** report the minimum detectable change, ≈ 2.8 standard errors of the difference

Mini-Case: Did Bitcoin's Tail Get Thinner after the ETFs?

- A typical assistant proposal: “institutional money since the spot ETFs has tamed Bitcoin's tails”
- Test: Hill index of $|r_t|$, $k = 2.5\%$ of n ; $H_0 : \alpha_{\text{post}} = \alpha_{\text{pre}}$ against $H_1 : \alpha_{\text{post}} > \alpha_{\text{pre}}$
 - ▶ $z = (\hat{\alpha}_{\text{post}} - \hat{\alpha}_{\text{pre}}) / \sqrt{s_{\text{pre}}^2 + s_{\text{post}}^2}$, s from a moving-block bootstrap (blocks of 20 days)
- Before (Sep 2014 – 10 Jan 2024, $k = 85$): $\hat{\alpha} = 4.16$ ($s = 0.52$)
- After (11 Jan 2024 – Sep 2026, $k = 24$): $\hat{\alpha} = 4.31$ ($s = 0.91$); $z = 0.15$
- Placebo, S&P 500 at the same date: $2.82 \rightarrow 2.86$, $z = 0.03$
- Interpretation: no evidence of a change, but the minimum detectable change is 2.9
 - ▶ the data cannot distinguish “unchanged” from a large thinning of the tail; wait for more data or pool assets

 ai_discovery_case.py

Your Turn: A Project Seed

- **Question:** did the tail index of Bitcoin and Ethereum change with institutionalisation (futures 2017, spot ETFs 2024)?
 - ▶ raw and GARCH-standardised returns: shape versus scale; loss and gain tails separately
- **Replicate first:** the chapter's Hill table ($k = 2.5\%$ of n) and the mini-case numbers
- **Deliverables**
 - ▶ a pre-registration note (H_0 , H_1 , rule for k , break dates, bootstrap) committed before the results
 - ▶ the change-point test of [Hoga \(2017\)](#) as a Quantlet; a report with power calculations
- **Logs:** prompts in `AI_USE.md`; every caught error (an i.i.d. standard error, a date chosen from the data) in `AI_ERRORS.md`

Conclusions

- ▣ Returns are **uncorrelated but not independent**: their size is predictable, their sign almost not
- ▣ Heavy tails are the rule: 5σ days are 1000+ times more frequent than the Normal distribution predicts
- ▣ Volatility clusters, has long memory and reacts asymmetrically to losses in equity markets
- ▣ Tails thin out with aggregation, but slowly
- ▣ Crypto shows the same facts with larger amplitude; managed exchange rates and thin markets add autocorrelation
- ▣ Range-based estimators extract more information from OHLC data
- ▣ Risk-adjusted measures must be read together with drawdowns and tail statistics

Further Study on Quantinar

Advanced courses on quantinar.com

- ▣  Statistics of Financial Markets
- ▣  Tukey's g- and h- transformations
- ▣  Applied Time Series Analysis with Python

Code and data

- ▣ All charts and tables: github.com/danpele/MFM/tree/main/Quantlets/Ch_01

References (1/3)

- ▣ Bouchaud, J.-P., Matacz, A., Potters, M. (2001). Leverage effect in financial markets: the retarded volatility model. *Physical Review Letters*, 87(22), 228701.
- ▣ Campbell, J.Y., Lo, A.W., MacKinlay, A.C. (1997). *The Econometrics of Financial Markets*. Princeton University Press.
- ▣ Christie, A.A. (1982). The stochastic behavior of common stock variances: value, leverage and interest rate effects. *Journal of Financial Economics*, 10(4), 407–432.
- ▣ Clark, P.K. (1973). A subordinated stochastic process model with finite variance for speculative prices. *Econometrica*, 41(1), 135–155.
- ▣ Cong, L.W., Li, X., Tang, K., Yang, Y. (2023). Crypto wash trading. *Management Science*, 69(11), 6427–6454.
- ▣ Cont, R. (2001). Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2), 223–236.
- ▣ Davis, R.A., Mikosch, T. (1998). The sample autocorrelations of heavy-tailed processes with applications to ARCH. *Annals of Statistics*, 26(5), 2049–2080.
- ▣ Ding, Z., Granger, C.W.J., Engle, R.F. (1993). A long memory property of stock market returns and a new model. *Journal of Empirical Finance*, 1(1), 83–106.
- ▣ Engle, R.F. (1982). Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4), 987–1007.
- ▣ Fama, E.F. (1965). The behavior of stock-market prices. *Journal of Business*, 38(1), 34–105.

References (2/3)

- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- ▣ Garman, M.B., Klass, M.J. (1980). *On the estimation of security price volatilities from historical data*. *Journal of Business*, 53(1), 67–78.
- ▣ Granger, C.W.J., Ding, Z. (1995). *Some properties of absolute return: an alternative measure of risk*. *Annales d'Économie et de Statistique*, 40, 67–91.
- ▣ Härdle, W.K., Chen, C.Y.-H., Overbeck, L. (Eds.) (2017). *Applied Quantitative Finance*, 3rd ed. Springer.
- ▣ Härdle, W.K., Harvey, C.R., Reule, R.C.G. (2020). *Understanding cryptocurrencies*. *Journal of Financial Econometrics*, 18(2), 181–208.
- ▣ Hill, B.M. (1975). *A simple general approach to inference about the tail of a distribution*. *Annals of Statistics*, 3(5), 1163–1174.
- ▣ Hoga, Y. (2017). *Change point tests for the tail index of β -mixing random variables*. *Econometric Theory*, 33(4), 915–954.
- ▣ Jarque, C.M., Bera, A.K. (1980). *Efficient tests for normality, homoscedasticity and serial independence of regression residuals*. *Economics Letters*, 6(3), 255–259.
- ▣ Karpoff, J.M. (1987). *The relation between price changes and trading volume: a survey*. *Journal of Financial and Quantitative Analysis*, 22(1), 109–126.
- ▣ Ljung, G.M., Box, G.E.P. (1978). *On a measure of lack of fit in time series models*. *Biometrika*, 65(2), 297–303.
- ▣ Mandelbrot, B. (1963). *The variation of certain speculative prices*. *Journal of Business*, 36(4), 394–419.

References (3/3)

- ▣ Parkinson, M. (1980). The extreme value method for estimating the variance of the rate of return. *Journal of Business*, 53(1), 61–65.
- ▣ Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets?. *The European Journal of Finance*, 29(10), 1064–1105.
- ▣ Praetz, P.D. (1972). The distribution of share price changes. *Journal of Business*, 45(1), 49–55.
- ▣ Rogers, L.C.G., Satchell, S.E. (1991). Estimating variance from high, low and closing prices. *Annals of Applied Probability*, 1(4), 504–512.
- ▣ Sharpe, W.F. (1966). Mutual fund performance. *Journal of Business*, 39(1), 119–138.
- ▣ Sharpe, W.F. (1994). The Sharpe ratio. *Journal of Portfolio Management*, 21(1), 49–58.
- ▣ Sortino, F.A., van der Meer, R. (1991). Downside risk. *Journal of Portfolio Management*, 17(4), 27–31.
- ▣ Taylor, S.J. (1986). Modelling Financial Time Series. Wiley; 2nd ed. (2008), World Scientific, doi:10.1142/6578.
- ▣ Yang, D., Zhang, Q. (2000). Drift-independent volatility estimation based on high, low, open, and close prices. *Journal of Business*, 73(3), 477–491.
- ▣ Bachelier, L. (1900). Théorie de la spéculation. *Annales scientifiques de l'École Normale Supérieure*, 17, 21–86.
- ▣ Student (1908). The probable error of a mean. *Biometrika*, 6(1), 1–25.