



Modelling Financial Markets

Chapter 1: Stylised Facts and Returns



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ **ACF**: AutoCorrelation Function
- ▣ **AI**: Artificial Intelligence
- ▣ **AR**: AutoRegressive (model); in event studies: Abnormal Return
- ▣ **ARCH**: AutoRegressive Conditional Heteroskedasticity
- ▣ **ARCH-LM**: ARCH Lagrange Multiplier (test)
- ▣ **BET-TR**: BET Total Return (index)
- ▣ **BNR**: Banca Națională a României — *the National Bank of Romania*
- ▣ **BVB**: Bursa de Valori București — *the Bucharest Stock Exchange*
- ▣ **CAGR**: Compound Annual Growth Rate
- ▣ **CI**: Confidence Interval
- ▣ **CLT**: Central Limit Theorem
- ▣ **CRIX**: CRypto IndeX
- ▣ **DAX**: Deutscher Aktienindex — *the German stock index*
- ▣ **EGARCH**: Exponential GARCH
- ▣ **ETF**: Exchange-Traded Fund
- ▣ **FTSE**: Financial Times Stock Exchange (FTSE Russell, index provider)
- ▣ **GARCH**: Generalised AutoRegressive Conditional Heteroskedasticity

Acronyms used in this material (2/2)

- ▣ **GEO**: Government Emergency Ordinance
- ▣ **GJR-GARCH**: Glosten–Jagannathan–Runkle GARCH
- ▣ **HAC**: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ **HAR-RV**: Heterogeneous AutoRegressive model of Realised Volatility
- ▣ **JB**: Jarque–Bera (test)
- ▣ **LB**: Ljung–Box (test)
- ▣ **LM**: Lagrange Multiplier
- ▣ **MDD**: Maximum DrawDown
- ▣ **OHLC**: Open, High, Low, Close
- ▣ **PDV**: Path-Dependent Volatility
- ▣ **QQ**: Quantile–Quantile (plot)
- ▣ **RMSE**: Root Mean Squared Error
- ▣ **SE**: Standard Error
- ▣ **TSPL**: Time-Shifted Power Law (kernel)
- ▣ **US**: United States
- ▣ **VIX**: CBOE Volatility Index
- ▣ **VIX9D**: Cboe 9-Day Volatility Index

Chapter Route and the Definition of a Stylised Fact

Route of this chapter

- Part I: returns, distribution and tails, dependence and its mechanisms
- Part II
 - ▶ persistence and the limits of inference, market exceptions
 - ▶ volatility from open, high, low and close prices
 - ▶ performance, a case study
 - ▶ one pass of the AI-assisted research loop
- Technical detail sits in the [Appendix](#) at the end

Definition

- A **stylised fact** is a qualitative statistical regularity of **returns**
 - ▶ seen across assets, markets and periods [Cont \(2001\)](#)
 - It remains a hypothesis [Ratliff-Crain, Van Oort, Koehler & Tivnan \(2025\)](#)
 - ▶ to be re-tested on each market and frequency
 - Key message: returns show little linear autocorrelation
 - ▶ but their **size** is highly predictable
-
- Stylised facts are the benchmark every model of returns must reproduce
 - ▶ the i.i.d. (independent and identically distributed) Normal-distribution model fails most of them

Data Used in This Chapter

Series	Price used	Type, currency	Calendar	Sample	T
S&P 500	daily close (plus open, high, low)	price index, USD	US trading days	1990–2026	9 245
BET-TR	daily close	total return, RON	BVB trading days	Sep 2014 – 2026	2 992
Bitcoin	daily close	price, USD	every day, 24/7	Sep 2014 – 2026	4 384
EUR/RON	BNR reference rate	RON per EUR, managed float	Romanian working days	Jul 2005 – 2026	5 352
Gold	daily close	spot, USD per ounce	weekdays only	1990–2026	9 559

- Daily closes up to 18 September 2026; BVB: Bucharest Stock Exchange
 - ▶ daily log returns on each series' own calendar
 - ▶ T : the number of daily returns in the series
- **Managed float:** a market-determined exchange rate with central-bank interventions
 - ▶ the official fixing of the National Bank of Romania (BNR) replaces the EUR/RON market (interbank) quotes
 - ▶ in several years these had a lag-1 autocorrelation of -0.4 to -0.7 , which came from quote errors
- In 1996–2005 the S&P 500 open equals the previous close on about 96% of days
 - ▶ still 16% in 2007
 - ▶ range-based estimators (built on the daily high–low range) therefore use data from 2008 on

Returns: Definitions and Aggregation (1/2)

	Simple return R_t	Log return r_t
Definition	$(P_t + D_t)/P_{t-1} - 1$	$\ln(1 + R_t) = \ln P_t - \ln P_{t-1}$ (no dividend)
Over k periods	$\prod_{j=0}^{k-1} (1 + R_{t-j}) - 1$	$\sum_{j=0}^{k-1} r_{t-j}$: additive over time
Across assets	$R_{p,t} = \sum_i w_i R_{i,t}$: exact	only approximately
Range	$R_t \geq -100\%$	$(-\infty, +\infty)$
Use	portfolios, performance reporting	single-asset dynamics, central limit theorem (CLT)

- ▣ The simple return is the relative change in the value of the investment; the log return is the change in the log price
- ▣ Notation
 - ▶ P_t : the closing price on day t ; D_t : the dividend paid in $(t-1, t]$
 - ▶ k : the number of periods aggregated; the index $j = 0, \dots, k-1$ runs over them
 - ▶ $R_{i,t}$: the simple return of asset i ; w_i : its weight in portfolio value, with $\sum_i w_i = 1$; $R_{p,t}$: the portfolio return
- ▣ With dividends use total-return or adjusted series (BET-TR)

Returns: Definitions and Aggregation (2/2)

- The log return is always below the simple return: $r = \ln(1 + R) < R$ for every $R \neq 0$
 - ▶ the reason: $\ln(1 + x) < x$ for every $x > -1$, $x \neq 0$, because the logarithm is concave
- The gap is negligible for small moves and large for big falls:
 - ▶ $R = +1\%$ gives $r = 0.995\%$
 - ▶ $R = -50\%$ gives $r = -69.3\%$
- Aggregation over time: the k -day log return is the sum of the k daily log returns
 - ▶ the k -day simple return needs the product of the growth factors $1 + R_{t-j}$
- Textbooks: [Campbell, Lo & MacKinlay \(1997\)](#), ch. 1; [Franke, Härdle & Hafner \(2019\)](#)

Annualisation, CAGR and Volatility Drag (1/2)

- ▣ Annualisation turns the daily mean and volatility into values per year
- ▣ Notation
 - ▶ A : the number of return intervals per year; 252 for stocks, 365 for crypto (traded every day), ≈ 260 for gold (the actual weekday quotes)
 - ▶ $\bar{r} = \frac{1}{T} \sum_{t=1}^T r_t$: the mean daily log return
 - ▶ $\hat{\sigma}$: the sample standard deviation of daily returns
- ▣ $\mu_{\text{ann}} = A \bar{r}$: the annual mean; needs only a constant mean
- ▣ $\sigma_{\text{ann}} = \sqrt{A} \hat{\sigma}$: the annual volatility; exact only for serially uncorrelated returns with constant variance
 - ▶ then daily variances add up, and the annual variance is $A \hat{\sigma}^2$
 - ▶ with autocorrelation the $\sqrt{A}$ factor is no longer correct
 - ▶ EUR/RON, lag-1 autocorrelation 0.17: annual volatility rises from 4.4% to 5.3% ([▶ Appendix](#); Seminar A2)


Annualisation, CAGR and Volatility Drag (2/2)

- **CAGR** (Compound Annual Growth Rate): the constant annual return that takes the price from P_0 to P_T

$$\text{CAGR} = (P_T/P_0)^{1/Y} - 1$$


- ▶ P_0, P_T : the price at the start and at the end of the sample; Y : the number of years
- ▶ CAGR is the annual geometric mean of the growth factors
- The geometric mean $\approx \mu_{\text{arith}} - \sigma^2/2$
 - ▶ μ_{arith} : the annual arithmetic mean of simple returns; σ^2 : their annual variance
 - ▶ **volatility drag**: the gap between the arithmetic and the geometric mean, $\approx \sigma^2/2$; it grows with the square of volatility
- Examples
 - ▶ S&P 500 1990–2026: $\hat{\sigma} = 1.133\%$ a day, 18.0% a year; CAGR 8.69% (price index, no dividends)
 - ▶ Bitcoin 2014–2026: arithmetic mean 65.2% a year ($A = 365$), CAGR 53.9%
 - ▶ a drag of 11 pp (percentage points)

Simple and Log Returns in Practice



- Left: $r \approx R$ only for small moves
- Right: Bitcoin 2014–2026 multiplied wealth by 176.9, i.e. $\sum_t r_t = \ln 176.9 = 5.18$
 - ▶ the sum of daily log returns is $\ln(P_T/P_0)$, the log of the multiple
 - ▶ adding daily simple returns gives $1 + \sum_t R_t = 8.83$, 20 times below the true multiple

Five Markets: Prices and Returns (1/2)



Five Markets: Prices and Returns (2/2)

- ▣ **Left:** prices on log scale, trending and non-stationary
- ▣ **Right:** daily log returns fluctuate around a small mean
 - ▶ calm and turbulent periods alternate
 - ▶ extremes cluster (2008, March 2020)
- ▣ EUR/RON moves more than 1% on only 1.4% of days
- ▣ Bitcoin, 12 March 2020: log return -46.5%
 - ▶ the price is multiplied by e^{r_t} , so it falls by $1 - e^{-0.465} = 37.2\%$

Sample Moments and Their Sampling Error (1/3)

- ▣ Skewness and kurtosis describe the shape of the distribution, after removing its mean and scale

$$z_t = \frac{r_t - \bar{r}}{\hat{\sigma}}, \quad \hat{S} = \frac{1}{T} \sum_{t=1}^T z_t^3, \quad \hat{K} = \frac{1}{T} \sum_{t=1}^T z_t^4$$

- ▶ z_t : the standardised return, with mean 0 and standard deviation 1
 - ▶ $\hat{S}$: sample skewness; $\hat{S} < 0$: the left tail is longer, large losses dominate
 - ▶ $\hat{K}$: sample kurtosis; $\hat{K} > 3$: more mass in the centre and in the tails than under the Normal distribution
 - ▶ S and K without a hat: the population values; the hat marks an estimated value
- ▣ Normal distribution: $S = 0$, $K = 3$, excess kurtosis $K - 3 = 0$

Sample Moments and Their Sampling Error (2/3)

- **Standard error (SE):** the standard deviation of an estimator across samples
- Mean: $SE(\bar{r}) = \sigma/\sqrt{T}$, with σ the standard deviation of daily returns, estimated by $\hat{\sigma}$
 - ▶ the 95% confidence interval: $\bar{r} \pm 1.96 SE(\bar{r})$
 - ▶ daily means are 13–40 times smaller than $\hat{\sigma}$
 - ▶ so the mean is the least precise moment
 - ▶ S&P 500: $\bar{r} = 0.033\% \pm 0.023\%$ a day (95%), $T = 9,245$
- i.i.d. data from the Normal distribution: $\sqrt{T} \hat{S} \rightarrow N(0, 6)$ and $\sqrt{T}(\hat{K} - 3) \rightarrow N(0, 24)$
 - ▶ $\rightarrow$: convergence in distribution as $T \rightarrow \infty$; $N(0, v)$: the Normal distribution with mean 0 and variance v
 - ▶ so $SE(\hat{S}) \approx \sqrt{6/T}$ and $SE(\hat{K}) \approx \sqrt{24/T}$; for the S&P 500: 0.025 and 0.051

Sample Moments and Their Sampling Error (3/3)

- The **delta method** gives the standard error of a function of an estimator

$$\sqrt{T}(\hat{\theta} - \theta) \rightarrow N(0, \nu) \implies \sqrt{T}(g(\hat{\theta}) - g(\theta)) \rightarrow N(0, g'(\theta)^2 \nu)$$

- ▶ θ : a parameter, e.g. a moment; $\hat{\theta}$: its estimator; ν : the asymptotic variance of $\hat{\theta}$
- ▶ g : a differentiable function, e.g. the ratio of moments that defines $\hat{S}$; g' : its derivative
- ▶ the idea: the first-order Taylor expansion $g(\hat{\theta}) \approx g(\theta) + g'(\theta)(\hat{\theta} - \theta)$
- Any i.i.d. distribution: the delta method gives a Normal limit
 - ▶ for $\hat{S}$ only if $\mathbb{E}r^6 < \infty$, for $\hat{K}$ only if $\mathbb{E}r^8 < \infty$
 - ▶ $\mathbb{E}$: the expectation; $\mathbb{E}r^6 < \infty$: the sixth moment is finite
 - ▶ the variance of $\hat{S}$ involves the sixth moment, that of $\hat{K}$ the eighth
 - ▶ whether these moments exist is checked empirically in the section on tails

Three Classical Tests and Their Assumptions (1/3)

- ▣ **ARCH** (AutoRegressive Conditional Heteroskedasticity): today's variance depends on past returns
- ▣ **Volatility clustering**: days with large moves tend to follow days with large moves
- ▣ The **tail index** α measures how fast the probability of extreme values decays

$$P(|r_t| > x) \approx c x^{-\alpha} \quad \text{for large } x$$

- ▶ x : a positive threshold; $c > 0$: a scale constant
- ▶ a small α means a heavy tail: the probability decays slowly as x grows
- ▶ so $\mathbb{E}|r_t|^p < \infty$ only for $p < \alpha$, where p is the order of the moment
- ▶ the estimates in this chapter give $\alpha \approx 3$

Three Classical Tests and Their Assumptions (2/3)

- ▣ **JB** (Jarque–Bera): $JB = T[\hat{S}^2/6 + (\hat{K} - 3)^2/24]$
 - ▶ the distance of skewness and kurtosis from their Normal-distribution values
- ▣ **LB** (Ljung–Box): $Q(m) = T(T+2) \sum_{k=1}^m \hat{\rho}_k^2 / (T-k)$
 - ▶ a weighted sum of the first m squared autocorrelations
 - ▶ $\hat{\rho}_k$: the sample autocorrelation at lag k , between r_t and r_{t-k} ; here $m = 10$ lags
- ▣ **ARCH-LM**: TR^2 from the regression of $\hat{\varepsilon}_t^2$ on $\hat{\varepsilon}_{t-1}^2, \dots, \hat{\varepsilon}_{t-q}^2$
 - ▶ LM: the Lagrange multiplier test; $\hat{\varepsilon}_t = r_t - \bar{r}$: the residuals; $q = 5$ lags
 - ▶ R^2 : the share of the variation in $\hat{\varepsilon}_t^2$ explained by past squared residuals
- ▣ **Reading**: all three statistics are ≥ 0 , and large values reject the null H_0
 - ▶ reference distributions: $\chi_2^2, \chi_m^2, \chi_q^2, \chi_n^2$: the chi-squared distribution with n degrees of freedom
 - ▶ 5% critical values: $\chi_2^2 = 5.99, \chi_5^2 = 11.1, \chi_{10}^2 = 18.3$

Three Classical Tests and Their Assumptions (3/3)

Test	Statistic	H_0	χ^2 limit needs	Problem here
JB, Jarque & Bera (1980)	$T \left[\frac{\hat{S}^2}{6} + \frac{(\hat{K}-3)^2}{24} \right]$	i.i.d., Normal distribution	H_0 itself (all moments exist)	a rejection does not say which part fails
LB, Ljung & Box (1978)	$T(T+2) \sum_{k=1}^m \frac{\hat{\rho}_k^2}{T-k}$	$\rho_1 = \dots = \rho_m = 0$	i.i.d., not only uncorrelated	volatility clustering: over-rejects
ARCH-LM, Engle (1982)	$TR^2, \hat{\varepsilon}_t^2$ on q lags	i.i.d., no ARCH effects	$\mathbb{E}r^4 < \infty$	$\alpha \approx 3 < 4$: size unreliable

- The question is not how to run the tests
 - ▶ but whether their reference distribution holds for returns

Worked Example: JB for the S&P 500 (1/2)

- Test statistic: $JB = T \left(\frac{\hat{S}^2}{6} + \frac{(\hat{K} - 3)^2}{24} \right)$
 - ▶ it measures how far $\hat{S}$ and $\hat{K}$ are from their Normal-distribution values, 0 and 3
 - ▶ distributed χ^2_2 under the null (5% critical value: 5.99)
- S&P 500, 1990–2026: $T = 9,245$, $\hat{S} = -0.364$, $\hat{K} - 3 = 10.88$
 - ▶ skewness term $\hat{S}^2/6 = 0.132/6 = 0.022$
 - ▶ kurtosis term $10.88^2/24 = 118.4/24 = 4.93$
 - ▶ $JB = 9,245 \times (0.022 + 4.93) \approx 45,800$, against the critical value 5.99
- Almost all of the statistic comes from the kurtosis term
 - ▶ the problem is the tails, not the asymmetry

Worked Example: JB for the S&P 500 (2/2)

- Under the null all moments exist
 - ▶ so the p-value is valid for the **joint** null “i.i.d. and Normal distribution”
 - ▶ heavy tails make JB grow with T
 - ▶ this is power, not a broken p-value
- A rejection does not say which part fails
 - ▶ it may be the marginal distribution or independence
 - ▶ tests of the marginal distribution with HAC (heteroskedasticity and autocorrelation consistent) variances:
 - ▶ [Appendix](#)

Robust Alternative: Quantile-Based Skewness and Kurtosis (1/3)

- Two quantile-based measures [Kim & White \(2004\)](#), which need no moments
 - ▶ q_p : the quantile of order p of returns, the value below which a share p of days lies; $q_{0.5}$: the median
- **Hinkley skewness** compares the lengths of the two tails around the median

$$\frac{q_{0.95} + q_{0.05} - 2q_{0.5}}{q_{0.95} - q_{0.05}}$$

- ▶ the numerator is $(q_{0.95} - q_{0.5}) - (q_{0.5} - q_{0.05})$: right tail minus left tail
 - ▶ it lies in $[-1, 1]$; a negative value indicates a longer left tail
- **Crow–Siddiqui excess kurtosis** compares the width of the tails with the width of the centre

$$\frac{q_{0.975} - q_{0.025}}{q_{0.75} - q_{0.25}} - 2.91$$

- ▶ 2.91: the value of the ratio for the Normal distribution; a positive value indicates heavier tails
 - ▶ both measures are 0 for the Normal distribution

Robust Alternative: Quantile-Based Skewness and Kurtosis (2/3)

- ▣ 95% confidence intervals (CI) from a **moving-block bootstrap**
 - ▶ it resamples blocks of 20 consecutive days (1,000 resamples)
 - ▶ the measure is recomputed on each resampled series
 - ▶ the blocks keep the dependence within them
- ▣ Valid for stationary, weakly dependent returns with a positive density at these quantiles
- ▣ A CI that excludes 0 indicates a significant departure from the Normal distribution

Robust Alternative: Quantile-Based Skewness and Kurtosis (3/3)

	Moment skew.	Hinkley [95% CI]	Excess kurt.	Crow–Siddiqui excess [95% CI]
S&P 500	−0.36	−0.064 [−0.092, −0.041]	10.9	1.54 [1.29, 1.79]
BET-TR	−1.40	−0.036 [−0.085, 0.025]	17.8	1.06 [0.73, 1.43]
Bitcoin	−0.70	−0.018 [−0.069, 0.032]	12.0	2.25 [1.88, 2.64]
EUR/RON	0.73	0.041 [−0.022, 0.093]	15.1	4.87 [3.74, 6.13]
Gold	−0.23	0.004 [−0.025, 0.030]	9.3	1.47 [1.29, 1.73]

- ▣ Quantile (Crow–Siddiqui) excess kurtosis is positive in all five markets
 - ▶ no CI contains 0
- ▣ Quantile-based negative skewness is established only for the S&P 500
 - ▶ BET-TR's moment skewness −1.40 comes from a few days (19 December 2018)

Summary Statistics: Five Markets

	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Observations	9245	2992	4384	5352	9559
Mean (% per day)	0.033	0.077	0.118	0.007	0.025
Std. dev. (% per day)	1.13	0.98	3.49	0.28	1.01
Annualised volatility (%)	18.0	15.5	66.6	4.4	16.3
Skewness	-0.36	-1.40	-0.70	0.73	-0.23
Excess kurtosis	10.9	17.8	12.0	15.1	9.3
Min / max log return (%)	-12.8 / 11.0	-11.9 / 6.7	-46.5 / 22.5	-2.5 / 2.9	-11.1 / 10.4
Jarque-Bera	45805	40433	26482	51575	34732
Ljung-Box $Q(10)$, r_t	91.6	70.0	20.2	221.7	32.1
Ljung-Box $Q(10)$, r_t^2	8018	714	213	2420	1049
ARCH-LM (5 lags)	2229	371	115	871	483





- Descriptive: heavy tails everywhere (excess kurtosis 9.3–17.8)
 - ▶ $Q(10)$ on r_t^2 is 10–90 times larger than on r_t
 - ▶ the size of returns is far more predictable than their direction
- Inferential: all i.i.d. p-values are below 0.03
 - ▶ but they are valid only under the assumptions in the table of the three classical tests

The Stylised Facts of Cont (2001)

Group	Facts
Distribution	2. heavy tails; 3. gain/loss asymmetry; 4. aggregational Gaussianity; 7. conditional heavy tails
Dependence of returns	1. absence of linear autocorrelation
Dependence of volatility	5. intermittency; 6. volatility clustering; 8. slow decay of the autocorrelation of $ r_t $; 11. asymmetry in time scales
Links to other variables	9. leverage effect; 10. volume/volatility correlation

- We re-test them on five markets, with data up to September 2026
 - we add the inference each test needs
- Early evidence: Mandelbrot (1963), Fama (1965); textbook treatment: Taylor (2008), Franke, Härdle & Hafner (2019)

Fact 2, Heavy Tails: Histograms and QQ-Plots (1/2)



Fact 2, Heavy Tails: Histograms and QQ-Plots (2/2)

- Extreme returns are far more frequent than under the Normal distribution
 - ▶ at the same variance
- Standardised returns $(r_t - \bar{r})/\hat{\sigma}$: higher peak and fatter tails
- QQ-plots compare empirical quantiles with those of the Normal distribution
 - ▶ they bend away from the line in both tails
- A Student- t fits much better

 MFM_ch1_heavy_tails

Tail Heaviness: Exceedances and the Student- t (1/3)

- The Student- t distribution, [Praetz \(1972\)](#): a heavy-tailed model for returns, with density

$$f(x) \propto \left(1 + \frac{(x - \mu)^2}{\nu s^2}\right)^{-(\nu+1)/2}$$

- Notation

- ▶ $f(x)$: the density at x ; $\propto$: “proportional to”, i.e. without the normalising constant
- ▶ μ : the centre of the distribution; $s > 0$: the scale parameter
- ▶ $\nu > 0$: the degrees of freedom; a small ν means heavy tails; as $\nu \rightarrow \infty$ we get the Normal distribution

- For a variable X with this distribution

- ▶ tail $P(X > x) \propto x^{-\nu}$: moments exist only of order $< \nu$
- ▶ variance $s^2 \nu / (\nu - 2)$ for $\nu > 2$
- ▶ excess kurtosis $6 / (\nu - 4)$ for $\nu > 4$

Tail Heaviness: Exceedances and the Student- t (2/3)

- $\hat{\nu}$, the estimated value of ν , is fitted by **maximum likelihood** to the whole distribution
 - ▶ it picks the μ , s and ν that make the observed data most probable
 - ▶ the many small returns dominate the fit
 - ▶ so a poor fit in the centre of the distribution moves $\hat{\nu}$
- EUR/RON: $\hat{\nu} = 1.2$ reflects that 57% of daily changes are below 0.1%, plus rare jumps

Tail Heaviness: Exceedances and the Student- t (3/3)

	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Days beyond $3\hat{\sigma}$ (%)	1.55	1.44	1.82	2.06	1.42
Days beyond $5\hat{\sigma}$ (count)	30	12	10	29	26
Expected under the Normal distribution, $> 5\hat{\sigma}$	0.005	0.002	0.003	0.003	0.005
Student- t degrees of freedom $\hat{\nu}$ (maximum likelihood)	2.8	3.1	2.2	1.2	2.9

□ Under the Normal distribution

- ▶ 0.27% of days lie beyond 3σ (observed 5–8 times more)
- ▶ one 5σ day occurs every 6,900 years of trading (observed 10–30)

The Hill Estimator of the Tail Index (1/3)

- The **tail index** α , defined with the classical tests: $P(|X| > x) \approx c x^{-\alpha}$ for large x
 - ▶ X : the variable studied, here $|r_t|$; $\mathbb{E}|X|^p < \infty$ only for $p < \alpha$
 - ▶ $\alpha > 2$: finite variance; $\alpha > 4$: finite kurtosis
 - ▶ power laws were first argued by [Mandelbrot \(1963\)](#) and [Fama \(1965\)](#)
- The **Hill estimator**, [Hill \(1975\)](#): estimates α from the k largest observations only
 - ▶ order them from largest to smallest: $X_{(1)} \geq X_{(2)} \geq \dots \geq X_{(n)}$
 - ▶ $X_{(i)}$: the i -th largest value; n : the number of observations
 - ▶ k : the number of tail observations; $X_{(k+1)}$ acts as the threshold where the tail starts
- **Pareto tail**: $P(X > x) = (x/u)^{-\alpha}$ for $x \geq u$, with threshold u ; the exact power law

The Hill Estimator of the Tail Index (2/3)

- The Hill estimator is the inverse of the mean log exceedance over the threshold

$$\hat{\alpha}_k = \left[\frac{1}{k} \sum_{i=1}^k \ln X_{(i)} - \ln X_{(k+1)} \right]^{-1} = \left[\frac{1}{k} \sum_{i=1}^k \ln \frac{X_{(i)}}{X_{(k+1)}} \right]^{-1}$$

- ▶ $\ln(X_{(i)}/X_{(k+1)})$: how far the i -th value exceeds the threshold, on a log scale
- ▶ for a Pareto tail these exceedances average $1/\alpha$
- ▶ large exceedances mean a heavy tail, hence a small $\hat{\alpha}_k$
- Standard error $SE(\hat{\alpha}_k) \approx \hat{\alpha}_k/\sqrt{k}$; 95% CI: $\hat{\alpha}_k \pm 1.96 SE$
 - ▶ it assumes independent tail observations and a Pareto tail
 - ▶ clustering of extremes makes it larger (▶ Appendix)

The Hill Estimator of the Tail Index (3/3)


□ Choice of k

- ▶ small k is noisy
- ▶ large k lets the centre in and biases $\hat{\alpha}$
- ▶ read the level where the **Hill plot** ($\hat{\alpha}_k$ against k) is flat, here $k = 1\text{--}5\%$ of n

□ Comparison with the Student- t , whose tail index is $\alpha = \nu$

- ▶ EUR/RON: the Hill estimator gives $\hat{\alpha} = 2.8$ while $\hat{\nu} = 1.2$
- ▶ a Student- t fitted to the whole distribution misdescribes the tails

Hill Plots for Five Markets



- $\hat{\alpha}$ for $|r_t|$ at $k = 2.5\%$ of n
 - ▶ S&P 500 3.1, BET-TR 2.7, Bitcoin 4.2, EUR/RON 2.8, gold 3.3
- For $k = 1-5\%$: between 2.3 and 3.5 except Bitcoin (3.1–4.2)
- At a common threshold, loss tails are heavier than gain tails
 - ▶ BET-TR: 2.2 versus 3.1
 - ▶ Bitcoin: 2.7 versus 3.4
- Tails are modelled with extreme value theory ([Chapter 7](#))

Tail Indices with Confidence Intervals (1/2)



$ r_t $, $k = 2.5\%$ of n	n	k	$\hat{\alpha}$	SE i.i.d.	95% CI (i.i.d.)	SE block bootstrap
S&P 500	9245	231	3.07	0.20	[2.68, 3.47]	0.33
BET-TR	2992	74	2.67	0.31	[2.06, 3.28]	0.45
Bitcoin	4384	109	4.22	0.40	[3.43, 5.01]	0.47
EUR/RON	5352	133	2.82	0.24	[2.34, 3.30]	0.28
Gold	9559	238	3.31	0.21	[2.89, 3.73]	0.24

- Moving-block bootstrap (blocks of 20 days, 500 resamples) keeps the clustering of extremes
 - SEs are 1.1–1.6 times the i.i.d. ones

Tail Indices with Confidence Intervals (2/2)

- $\alpha = 4$ is outside every interval except Bitcoin's
 - ▶ also when the interval uses the bootstrap SE
 - ▶ the fourth moment very likely does not exist
- $\alpha = 2$ is outside every i.i.d. interval
 - ▶ with the bootstrap SE only BET-TR reaches it: the lower limit $\hat{\alpha} - 1.96 \text{ SE} = 2.67 - 1.96 \times 0.45 = 1.79$
- Bitcoin's 4.2 uses a threshold where its Hill plot is not yet flat
 - ▶ at a common threshold its two tails give 2.71 and 3.36 (▶ [Appendix](#))

Fact 3: Gain/Loss Asymmetry in the Data (1/4)





Fact 3: Gain/Loss Asymmetry in the Data (2/4)

- Survival functions $P(X > x)$ of losses and gains, log-log scale
 - ▶ $P(X > x)$: the share of days on which X exceeds x ; $X = -r_t$ for losses, $X = r_t$ for gains
 - ▶ large losses are more frequent than gains of the same size
- **VaR** (Value-at-Risk) 1% = $-q_{0.01}$, the daily loss threshold exceeded with probability 1%; $q_{0.01}$: the 1% quantile of returns (Chapter 7)

 MFM_ch1_heavy_tails


Fact 3: Gain/Loss Asymmetry in the Data (3/4)



Fact 3: Gain/Loss Asymmetry in the Data (4/4)

- VaR 1% exceeds the 99% return quantile $q_{0.99}$ for 4 of 5 markets
 - ▶ e.g. for the S&P 500: 3.15% versus 2.97%
- EUR/RON is the exception: leu depreciations (up-moves) are larger
 - ▶ $q_{0.99} = 0.94\%$ versus VaR 1% = 0.76%
- A Normal-distribution model with the right variance underestimates such quantiles
 - ▶ stress tests based on “ 3σ events” are too optimistic

Facts 1 and 6: Uncorrelated but Not Independent (1/4)




Facts 1 and 6: Uncorrelated but Not Independent (2/4)

- **Fact 1:** $\text{corr}(r_t, r_{t-k}) \approx 0$, where k is the lag in days
 - ▶ S&P 500 1990–2026, lag 1: -0.08
 - ▶ i.i.d. reference band: $\pm 1.96/\sqrt{T}$; outside it the autocorrelation is significant at 5%
- **Fact 6, volatility clustering:** $\text{corr}(|r_t|, |r_{t-k}|) > 0$ for many lags (lag 1: 0.27)
 - ▶ large moves follow large moves of either sign

 MFM_ch1_volatility_clustering

Facts 1 and 6: Uncorrelated but Not Independent (3/4)



Facts 1 and 6: Uncorrelated but Not Independent (4/4)

- Independence would require zero correlation of **every** function of returns
 - ▶ returns are uncorrelated but not independent
- Model $r_t = \sigma_t \varepsilon_t$ with σ_t driven by past returns
 - ▶ σ_t : the conditional volatility on day t ; ε_t : i.i.d. shocks with mean 0 and variance 1
 - ▶ GARCH, generalised ARCH ([Chapter 5](#))

 MFM_ch1_volatility_clustering

The ACF under Conditional Heteroskedasticity (1/3)

- ▣ **ACF** (AutoCorrelation Function): the sample autocorrelations $\hat{\rho}_k$, plotted against the lag k

$$\hat{\rho}_k = \frac{\sum_t x_t x_{t-k}}{\sum_t x_t^2}, \quad x_t = r_t - \bar{r}$$

- ▶ x_t : the demeaned return; $\hat{\rho}_k \in [-1, 1]$
- ▶ $\hat{\rho}_k > 0$: a day above the mean tends to be followed, k days later, by a day above the mean
- ▣ Uncorrelated returns whose conditional variance moves: $\hat{\rho}_k \approx N(0, \tau_k)$ [Romano & Thombs \(1996\)](#)
 - ▶ τ_k : the variance of $\hat{\rho}_k$, estimated by $\hat{\tau}_k = \sum_t x_t^2 x_{t-k}^2 / (\sum_t x_t^2)^2$
 - ▶ i.i.d.: $T\tau_k = 1$, which gives the classical band $\pm 1.96/\sqrt{T}$
 - ▶ volatility clustering: $T\tau_k > 1$; the classical band is too narrow and flags autocorrelations as significant too often

The ACF under Conditional Heteroskedasticity (2/3)

- ▣ Robust t -statistic: $\hat{\rho}_k / \sqrt{\hat{\tau}_k}$, compared with ± 1.96
- ▣ Robust portmanteau test [Lobato, Nankervis & Savin \(2001\)](#): null hypothesis $\rho_1 = \dots = \rho_m = 0$

$$\tilde{Q}(m) = \sum_{k=1}^m \frac{\hat{\rho}_k^2}{\hat{\tau}_k}$$

- ▶ the sum of the squared robust t -statistics for the first m lags, so $\tilde{Q}(m) \geq 0$
- ▶ reference distribution χ_m^2 ; large values reject the null
- ▶ its assumptions: [Appendix](#)
- ▣ The difference from the Ljung–Box $Q(m)$
 - ▶ $Q(m)$ divides each $\hat{\rho}_k^2$, approximately, by the i.i.d. variance $1/T$
 - ▶ $\tilde{Q}(m)$ divides by the estimated variance $\hat{\tau}_k$

The ACF under Conditional Heteroskedasticity (3/3)

Lag 1	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
$\hat{\rho}_1$	-0.079	0.061	-0.022	0.170	-0.011
$T\hat{\rho}_1$	5.00	4.35	2.89	6.85	2.43
t -statistic, i.i.d. / robust	-7.6 / -3.4	3.3 / 1.6	-1.5 / -0.9	12.4 / 4.7	-1.1 / -0.7

- S&P 500: $Q(10) = 91.6$ but $\tilde{Q}(10) = 19.0$ ($p = 0.041$)
- With robust SEs only the S&P 500 ($t = -3.4$) and EUR/RON ($t = 4.7$) keep a significant lag-1 autocorrelation
 - it is economically small

Heavy Tails from Mixtures of Normal Distributions (1/2)

- The model: the return is a Normal shock times a random volatility

$$r_t = \sigma_t \varepsilon_t, \quad \varepsilon_t \sim N(0, 1) \text{ independent of } \sigma_t$$

- The kurtosis of returns is at least 3

$$K = \frac{\mathbb{E}[r_t^4]}{\mathbb{E}[r_t^2]^2} = 3 \frac{\mathbb{E}[\sigma_t^4]}{\mathbb{E}[\sigma_t^2]^2} \geq 3$$

- ▶ the middle equality: $\mathbb{E}[\varepsilon_t^4] = 3$ and $\mathbb{E}[\varepsilon_t^2] = 1$, and σ_t and ε_t are independent
- ▶ the inequality is Jensen's inequality ($\mathbb{E}[X^2] \geq \mathbb{E}[X]^2$) with $X = \sigma_t^2$
- ▶ equality only if σ_t is constant; any variation in volatility creates excess kurtosis


Heavy Tails from Mixtures of Normal Distributions (2/2)

- Example: $\sigma = 1\%$ on 90% of days and 3% on 10% of days

$$K = 3 \cdot \frac{0.9 \cdot 1 + 0.1 \cdot 81}{(0.9 + 0.1 \cdot 9)^2} = 3 \cdot \frac{9.0}{3.24} = 8.33$$


- ▶ numerator: $\mathbb{E}[\sigma^4] = 0.9 \cdot 1^4 + 0.1 \cdot 3^4$, in units of $(1\%)^4$
 - ▶ denominator: $\mathbb{E}[\sigma^2]^2$, with $\mathbb{E}[\sigma^2] = 0.9 \cdot 1^2 + 0.1 \cdot 3^2 = 1.8$
 - ▶ the excess kurtosis is 5.33, although each day's return, given σ , is Normally distributed
- This links heavy tails (Fact 2) to volatility clustering (Fact 6)
 - ▶ a mixture drawn independently every day has heavy tails but no clustering (Seminar A6, A8)

Fact 5: Intermittency in the Data




- Bursts of high volatility separated by calm periods: constant-volatility models cannot produce them
- 63-day rolling volatility $\sqrt{A} \times$ the standard deviation over the window, log scale
 - ▶ common spikes: 2008–2009, March 2020
- EUR/RON volatility collapsed after 2012, under the BNR's managed float

Fact 4: Aggregational Gaussianity



- h -day return: $r_t(h) = r_t + \dots + r_{t-h+1}$
 - ▶ computed in non-overlapping blocks
 - ▶ the CLT pulls sums towards the Normal distribution
 - ▶ slowly, because daily returns are dependent
- Excess kurtosis: S&P 500 $10.9 \rightarrow 1.6$ at $h = 60$, Bitcoin $12.0 \rightarrow 0.1$
- At $h = 60$ only 154 and 73 returns remain
 - ▶ SE of excess kurtosis (i.i.d., Normal): $\sqrt{24/n} = 0.39$ and 0.57
 - ▶ n : the number of h -day returns
 - ▶ so these values are very imprecise
- $h = 60$ is a quarter for the S&P 500 but two months for Bitcoin

Fact 9: The Leverage Effect (1/3)




■ i.i.d. reference band $\pm 1.96/\sqrt{T}$

Fact 9: The Leverage Effect (2/3)

- $L(k) = \text{corr}(r_t, |r_{t+k}|)$, $k > 0$: the correlation of today's return with the size of the return k days later
 - ▶ $L(k) < 0$: after falls, returns become larger in absolute value, i.e. volatility rises
- S&P 500: $L(1) = -0.12$, still -0.07 at $k = 10$
 - ▶ no effect for $k < 0$
- BET-TR -0.12 ; weak for Bitcoin (-0.07)
 - ▶ negligible for gold (-0.02), reversed for EUR/RON ($+0.10$)

 MFM_ch1_leverage_volume


Fact 9: The Leverage Effect (3/3)




i.i.d. reference band $\pm 1.96/\sqrt{T}$

- Possible explanations
 - ▶ financial leverage of firms [Christie \(1982\)](#)
 - ▶ the retarded volatility model [Bouchaud, Matacz & Potters \(2001\)](#)
- Asymmetric GARCH models: GJR-GARCH, EGARCH ([Chapter 5](#))

Fact 10: Volume and Volatility Move Together




- Abnormal log volume ($\ln V_t$ minus its 250-day moving average) against $|r_t|$; V_t : the volume traded on day t
 - ▶ Spearman rank correlation, the correlation between the ranks of the two series, in $[-1, 1]$: 0.25 (S&P 500) and 0.39 (Bitcoin)
- Mixture-of-distributions hypothesis [Clark \(1973\)](#): information flow drives both price moves and volume
 - ▶ survey of the empirical evidence: [Karpoff \(1987\)](#)
- Bitcoin volume is exchange-reported
 - ▶ it is partly **wash trading** (an exchange trading with itself) [Cong, Li, Tang & Yang \(2023\)](#)

Fact 8: Slow Decay of the ACF of $|r_t|$ (1/2)

- ρ_k : the autocorrelation of $|r_t|$ at lag k ; $c > 0$: a constant
- **Short memory:** $\sum_k |\rho_k| < \infty$
 - ▶ e.g. exponential decay $\rho_k \approx c \phi^k$, with $0 < \phi < 1$ (GARCH(1,1), Chapter 5)
- **Long memory:** hyperbolic decay $\rho_k \approx c k^{2d-1}$, $0 < d < \frac{1}{2}$ Ding, Granger & Engle (1993)
 - ▶ d : the memory parameter; a larger d means slower decay, and $\sum_k \rho_k = \infty$

Fact 8: Slow Decay of the ACF of $|r_t|$ (2/2)



- S&P 500: 0.27 at lag 1, 0.22 at lag 20, still 0.09 at lag 100
 - ▶ BET-TR decays fastest
- Slow empirical decay is not yet established long memory
 - ▶ structural breaks, i.e. abrupt shifts in the level of volatility, can imitate it

Estimating the Memory Parameter d (1/2)

- Long memory in the frequency domain: the spectral density grows without bound at low frequencies

$$f(\lambda) \approx G\lambda^{-2d} \quad \text{as } \lambda \rightarrow 0$$

- ▶ $f(\lambda)$: the spectral density, the contribution of cycles of frequency λ to the variance of the series; small λ : long cycles
- ▶ $G > 0$: a constant; d : the memory parameter
- ▶ semiparametric estimators use only the m lowest Fourier frequencies $\lambda_j = 2\pi j/T$, $j = 1, \dots, m$

- Two estimators (formulas in the [Appendix](#))

- ▶ local Whittle [Robinson \(1995\)](#): $\sqrt{m}(\hat{d} - d) \rightarrow N(0, \frac{1}{4})$, so $\text{SE} = 1/(2\sqrt{m})$
- ▶ log-periodogram regression [Geweke & Porter-Hudak \(1983\)](#): $\text{SE} = \pi/\sqrt{24m}$
- ▶ here $m = T^{0.65}$

Estimating the Memory Parameter d (2/2)

$ r_t , m = T^{0.65}$	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
m	378	181	232	265	386
Local Whittle $\hat{d}$ (SE)	0.48 (0.03)	0.37 (0.04)	0.35 (0.03)	0.39 (0.03)	0.40 (0.03)
Log-periodogram $\hat{d}$ (SE)	0.50 (0.03)	0.41 (0.05)	0.35 (0.04)	0.38 (0.04)	0.42 (0.03)

- $\hat{d} \approx 0.35\text{--}0.50$ everywhere
 - ▶ at least 8 SEs from $d = 0$ (short memory) with both estimators
- The S&P 500 is close to the stationarity boundary $\frac{1}{2}$
 - ▶ for $d \geq \frac{1}{2}$ the series is no longer stationary


Long Memory or Structural Breaks?

- Rare shifts in the level of volatility produce a slowly decaying ACF of $|r_t|$
 - ▶ and $\hat{d} > 0$ without long memory Diebold & Inoue (2001), Mikosch & Stărică (2004)
- Check: locate one variance break with the cumulative sum of squares Inclán & Tiao (1994)
 - ▶ then re-estimate d (local Whittle) on each side

$ r_t $	Break date	Volatility before / after (%)	$\hat{d}$ before (SE)	$\hat{d}$ after (SE)
S&P 500	20 Jul 1998	12.6 / 19.3	0.28 (0.04)	0.51 (0.03)
BET-TR	17 Dec 2018	12.5 / 16.9	0.19 (0.05)	0.37 (0.04)
Bitcoin	19 Jun 2022	74.9 / 47.7	0.39 (0.04)	0.15 (0.05)
EUR/RON	27 Jun 2013	6.6 / 2.2	0.35 (0.04)	0.35 (0.04)
Gold	29 Nov 2005	13.2 / 18.3	0.30 (0.03)	0.37 (0.03)


- $\hat{d}$ stays well above 0 in every sub-sample
 - ▶ one break does not explain the memory
- But $\hat{d}$ moves (Bitcoin $0.39 \rightarrow 0.15$)
 - ▶ several breaks or a time-varying d remain open
 - ▶ the break dates are descriptive

The Taylor Effect (1/2)



- **Taylor effect:** absolute returns $|r_t|$ are more strongly autocorrelated than squared returns r_t^2
 - ▶ $\text{corr}(|r_t|^\delta, |r_{t-k}|^\delta)$ peaks at a power δ close to 1
 - ▶ not at $\delta = 2$
- The chart: this autocorrelation as a function of $\delta > 0$
 - ▶ $\delta = 1$: absolute returns
 - ▶ $\delta = 2$: squared returns

The Taylor Effect (2/2)



- S&P 500: maximum at $\delta = 1.75$ (lag 1), 1.5 (lag 5) and 1.25 (lag 20)
 - ▶ the optimum moves towards 1 at longer lags
 - ▶ and is never at $\delta = 2$
- Documented by Taylor (1986), named by Granger & Ding (1995)
- The effect motivates
 - ▶ power-GARCH models, which model σ_t^δ with an estimated δ (Chapter 5)
 - ▶ absolute-value volatility measures

Moments, Tail Index and Valid Inference

Statistic	Needs	With Hill $\hat{\alpha} \approx 2.5\text{--}3.5$
Mean, variance of r_t	$\alpha > 2$	finite, slow convergence
ACF of $ r_t $, robust ACF bands, $\tilde{Q}$	$\alpha > 4$ for Normal limits	doubtful
Kurtosis, ACF of r_t^2 , ARCH-LM, Q on r_t^2	$\alpha > 4$ to exist, $\alpha > 8$ for Normal limits	no standard inference
Skewness	$\alpha > 6$ for a Normal limit	no standard inference
Quantiles, number of tail days, Hill $\hat{\alpha}$	no moments	valid, with dependence-robust SEs

- With an infinite fourth moment the sample ACF of r_t^2 converges slowly
 - ▶ its limit is not the Normal distribution [Davis & Mikosch \(1998\)](#)
 - ▶ the $\pm 1.96/\sqrt{T}$ band and the p-values of Q on r_t^2 and ARCH-LM are not reliable
- The conclusion “volatility clusters” is robust
 - ▶ the exact magnitudes are not robust
 - ▶ one reason behind the Taylor effect: $|r_t|$ needs fewer finite moments than r_t^2

Facts 7 and 11: Conditional Tails and Time Scales

Fact 7: conditional heavy tails

- Fit a volatility model, e.g. GARCH (Chapter 5)
 - ▶ and standardise: $\hat{z}_t = r_t / \hat{\sigma}_t$, with $\hat{\sigma}_t$ the model's volatility estimate for day t
- The kurtosis of $\hat{z}_t$ is lower than that of r_t , still above 3
- Models need heavy-tailed innovations (Student- t , possibly skewed)

Fact 11: asymmetry in time scales

- Weekly volatility predicts daily volatility better than the reverse
- Traders with different horizons react at different frequencies
- The heterogeneous autoregressive model of realised volatility (**HAR-RV**, Chapter 9)
 - ▶ realised volatility is the sum of squared intraday returns


- Both facts are documented in Cont (2001)
- Returns standardised by realised volatility from intraday data are close to the Normal distribution
 - ▶ Andersen, Bollerslev, Diebold & Ebens (2001) (Chapter 9)

Differences between the Five Markets

Market	Institutional feature	What differs	Candidate explanation
Bitcoin	trades 24/7; $A = 365$	volatility 66.6% versus 18.0% (S&P 500); worst day -46.5% ; weak leverage (-0.07); ACF of $ r_t $ at lag 100: 0.02 versus 0.09	no corporate leverage; volatility rises after large moves of either sign
EUR/RON	BNR daily fixing, managed float	volatility 4.4% but excess kurtosis 15.1; $\hat{\rho}_1 = +0.17$; largest moves October 2008 ($+2.9\%$, -2.5%)	adjustments smoothed over several days
BET-TR	total return, frontier/emerging market	skewness -1.40 ; worst day -11.9% on 19 Dec 2018 (bank-asset tax announced on 18 Dec, Government Emergency Ordinance 114/2018); $\hat{\rho}_1 = +0.06$	policy shocks; non-synchronous trading of constituents Campbell, Lo & MacKinlay (1997), ch. 3

- The explanations are hypotheses, not identified causes
- Low volatility is not low tail risk
- Heavy tails and volatility clustering hold in all five markets
 - ▶ autocorrelation and asymmetry depend on institutions

Open, High, Low and Close (OHLC) Prices (1/3)



Open, High, Low and Close (OHLC) Prices (2/3)

- **OHLC:** four prices of day t
 - ▶ O_t : the open
 - ▶ H_t : the high
 - ▶ L_t : the low
 - ▶ C_t : the close
- **Range-based** estimators use the whole path of the day
 - ▶ the range H_t/L_t is large on volatile days
 - ▶ even if the close equals the open

Open, High, Low and Close (OHLC) Prices (3/3)

- Logs of the four prices

$$o_t = \ln O_t, \quad h_t = \ln H_t, \quad l_t = \ln L_t, \quad c_t = \ln C_t$$

- ▶ the lower-case letter is the log of the price with the same initial

- Log range of day t

$$h_t - l_t = \ln(H_t/L_t) \geq 0$$

- ▶ approximately the percentage gap between the day's high and low

- Three log returns of day t

- ▶ close-to-close: $c_t - c_{t-1}$
- ▶ overnight: $o_t - c_{t-1}$
- ▶ intraday: $c_t - o_t$

- **Close-to-close estimator**: the standard deviation of $c_t - c_{t-1}$

- ▶ one observation per day

- **Efficiency** of an estimator X : $\text{Var}(\hat{\sigma}_{CC}^2)/\text{Var}(\hat{\sigma}_X^2)$

- ▶ an efficiency of 5 means the same precision with 5 times fewer days

Parkinson and Garman–Klass Estimators (1/2)

- ▣ **Parkinson, Parkinson (1980)**: estimates the daily variance from the daily range only

$$\hat{\sigma}_P^2 = \frac{1}{4 \ln 2} \frac{1}{n} \sum_{t=1}^n (h_t - l_t)^2$$

- ▶ n : the number of days; $h_t - l_t$: the log range of day t
 - ▶ for a Brownian motion with daily variance σ^2 , $\mathbb{E}[(h_t - l_t)^2] = 4 \ln 2 \sigma^2$; dividing by $4 \ln 2$ makes the estimator unbiased
 - ▶ $\hat{\sigma}_P^2$ is a daily variance; the annual volatility is $\sqrt{A \hat{\sigma}_P^2}$
- ▣ Assumption: a Brownian motion without drift and without overnight jumps
 - ▶ Brownian motion: the continuous-time random walk ([Chapter 11](#))
 - ▶ efficiency relative to close-to-close: about 5 (Seminar A5); the same precision from 5 times fewer days

Parkinson and Garman–Klass Estimators (2/2)

- ▣ **Garman–Klass**, Garman & Klass (1980): combines the range with the open-to-close move

$$\hat{\sigma}_{GK}^2 = \frac{1}{n} \sum_{t=1}^n \left[\frac{1}{2} (h_t - l_t)^2 - (2 \ln 2 - 1) (c_t - o_t)^2 \right]$$

- ▶ $c_t - o_t$: the open-to-close log return
 - ▶ the weights $\frac{1}{2}$ and $2 \ln 2 - 1 \approx 0.386$ come from the minimum-variance estimator derived by Garman and Klass
 - ▶ the same assumptions as Parkinson: no drift, no overnight jumps
- ▣ Efficiency relative to close-to-close: 7–8

Rogers–Satchell and Yang–Zhang Estimators (1/2)

- ▣ **Rogers–Satchell**, [Rogers & Satchell \(1991\)](#): unbiased under a non-zero **drift** (expected trend of the log price)

$$\hat{\sigma}_{RS}^2 = \frac{1}{n} \sum_{t=1}^n [(h_t - o_t)(h_t - c_t) + (l_t - o_t)(l_t - c_t)]$$

- ▶ the terms measure how far the high and the low moved away from the open and from the close
- ▶ both products are ≥ 0 , because $l_t \leq o_t, c_t \leq h_t$
- ▶ for a Brownian motion with drift their mean is exactly the daily variance
- ▶ it does not include the overnight jump


Rogers–Satchell and Yang–Zhang Estimators (2/2)

- ▣ **Yang–Zhang, Yang & Zhang (2000)**: adds the overnight (close-to-open) jump

$$\hat{\sigma}_{YZ}^2 = \hat{\sigma}_{\text{overnight}}^2 + k \hat{\sigma}_{\text{open-close}}^2 + (1 - k) \hat{\sigma}_{RS}^2, \quad k = \frac{0.34}{1.34 + \frac{n+1}{n-1}}$$

- ▶ $\hat{\sigma}_{\text{overnight}}^2$: the sample variance of $o_t - c_{t-1}$, the overnight return
 - ▶ $\hat{\sigma}_{\text{open-close}}^2$: the sample variance of $c_t - o_t$
 - ▶ k : the weight, chosen to make the combination minimum-variance; $k \approx 0.145$ for large n
 - ▶ minimum-variance combination among drift-independent estimators robust to opening jumps
- ▣ Range estimators are the bridge to realised volatility from intraday data ([Chapter 9](#))

Estimators on the S&P 500 (1/2)




Estimators on the S&P 500 (2/2)

- 21-day windows, 2008–2026; averages
 - ▶ close-to-close 16.4%, Yang–Zhang 13.9%, Parkinson 13.0%, Garman–Klass 12.4%, Rogers–Satchell 12.3%
- Pure range estimators are lower: they miss overnight gaps
- The range-based series are 3–4 times smoother: a property of the estimates, not by itself proof of efficiency


 MFM_ch1_vol_estimators

Stale Opens: the S&P 500 Index versus SPY (1/2)




- **Stale open:** the index open is computed from the constituents' first prices
 - ▶ at 9:30 many stocks have not traded yet
 - ▶ so their previous close enters the open
- SPY, an ETF (exchange-traded fund) that tracks the S&P 500, opens at a traded price
- Open = previous close
 - ▶ about 96% of days in 1996–2005
 - ▶ 3.5% over 2008–2026 (up to 26% in 2012–2013)

Stale Opens: the S&P 500 Index versus SPY (2/2)




- A stale open moves part of the overnight move into the day
 - ▶ index overnight volatility: 6.7%
 - ▶ SPY, whose open is a traded price: 12.0%
- Yang-Zhang falls below close-to-close on the index only
 - ▶ use a traded instrument for estimators that use the open

Drawdowns Across Markets (1/2)



- Drawdown $DD_t = P_t / \max_{s \leq t} P_s - 1$: the loss from the peak reached up to day t ([Chapter 0](#))
 - ▶ $\max_{s \leq t} P_s$: the highest price up to day t ;
 $DD_t \in [-1, 0]$, with 0 at a new peak
 - ▶ maximum drawdown $MDD = \min_t DD_t$: the deepest fall from a peak
- MDD in the five markets
 - ▶ Bitcoin -83%, S&P 500 -57%, gold -45%
 - ▶ BET-TR -31% (March 2020), EUR/RON -16%

Drawdowns Across Markets (2/2)



- S&P 500: 9 October 2007 (1,565.15) to 9 March 2009 (676.53): -56.8%
 - ▶ **recovery**: the price regaining its previous peak
 - ▶ after a fall of x it needs a gain of $1/(1-x) - 1$
 - ▶ here $1/(1 - 0.568) - 1 = +131\%$
- The MDD grows with the sample length
 - ▶ compare only equal horizons (▶ [Appendix](#))

Recovery Duration after the Maximum Drawdown

- ▣ **Recovery:** the price regaining the peak it had before the fall
- ▣ **Recovery duration:** the time from that peak until the price reaches it again

	MDD	Peak	Trough	Peak regained
S&P 500	−57%	Oct 2007	Mar 2009	Mar 2013 (5.5 years)
BET-TR	−31%	Jan 2020	Mar 2020	Dec 2020 (0.9 years)
Bitcoin	−83%	Dec 2017	Dec 2018	Nov 2020 (3.0 years)
EUR/RON	−16%	Dec 2005	Jul 2007	Jan 2008 (2.0 years)
Gold	−45%	Sep 2011	Dec 2015	Jul 2020 (8.9 years)

- ▣ Depth and duration are different risks
 - ▶ gold's drawdown was shallower than Bitcoin's but lasted three times longer
- ▣ Fast V-shaped recoveries (BET-TR 2020)
 - ▶ versus long underwater periods (gold 2011–2020)
- ▣ For EUR/RON a “drawdown” is a fall of the euro against the leu

Sharpe, Sortino and Calmar Ratios (1/3)

▣ **Sharpe ratio** Sharpe (1966), Sharpe (1994)

- ▶ the excess return earned per unit of volatility

$$SR = \frac{\mu_{\text{ann}} - r_f}{\sigma_{\text{ann}}}$$

- ▶ r_f : the risk-free rate; $r_f = 0$ here, for comparability
 - ▶ $\mu_{\text{ann}} = A\bar{R}$: the annual mean of **simple** returns
 - ▶ $\bar{R}$: the mean daily simple return
 - ▶ σ_{ann} : the annual volatility of simple returns
 - ▶ a higher SR means more return for the same risk
- ▣ S&P 500 1990–2026: $\mu_{\text{ann}} = 9.95\%$, $\sigma_{\text{ann}} = 17.97\%$, $SR = 0.55$
- ▶ $\mu_{\text{ann}} = \text{mean log return } 8.33\% \text{ plus } \approx \sigma^2/2 = 1.62 \text{ pp}$

Sharpe, Sortino and Calmar Ratios (2/3)

▣ **Sortino ratio** Sortino & van der Meer (1991)

- ▶ the same numerator as the Sharpe ratio
- ▶ in the denominator: only the volatility of losses

$$SoR = \frac{\mu_{\text{ann}} - r_f}{\sigma_{\text{ann}}^-}, \quad \sigma^- = \sqrt{\frac{1}{T} \sum_{t=1}^T \min(R_t, 0)^2}$$

- ▶ σ^- : the daily downside deviation
- ▶ $\min(R_t, 0)$ keeps only negative returns
- ▶ the average runs over all T days; positive days count as 0
- ▶ $\sigma_{\text{ann}}^- = \sqrt{A} \sigma^-$: the annual version

- ▣ S&P 500: $\sigma_{\text{ann}}^- = 12.68\%$, $SoR = 0.78$

Sharpe, Sortino and Calmar Ratios (3/3)

▣ **Calmar ratio**

- ▶ the annual growth earned per unit of maximum drawdown

$$\text{Calmar} = \frac{\text{CAGR}}{|MDD|}$$

- ▶ CAGR: the compound annual growth rate
- ▶ MDD: the maximum drawdown in the sample

▣ S&P 500: $8.69\%/56.78\% = 0.15$

- ▶ the value is driven by one episode

Sharpe Ratios with Standard Errors (1/2)

□ Delta method, i.i.d. returns from the Normal distribution [Lo \(2002\)](#)

- ▶ SR_d : the daily Sharpe ratio; $\widehat{SR}_d$: its estimate; $\text{Var}(\widehat{SR}_d) \approx (1 + SR_d^2/2)/T$
- ▶ annualised, $SR_{\text{ann}} = \sqrt{A} SR_d$: $\text{SE}(\widehat{SR}_{\text{ann}}) \approx \sqrt{A(1 + SR_d^2/2)/T} \approx 1/\sqrt{Y}$, with $Y = T/A$ years
- ▶ the span matters, not the frequency

	S&P 500	BET-TR	Bitcoin	EUR/RON	Gold
Years Y	36.7	12.0	12.0	21.2	36.7
Sharpe ratio (SE, Normal)	0.55 (0.17)	1.34 (0.29)	0.98 (0.29)	0.43 (0.22)	0.48 (0.17)


 MFM_ch1_risk_measures

Sharpe Ratios with Standard Errors (2/2)

- Clustering: HAC or studentised block-bootstrap tests [Ledoit & Wolf \(2008\)](#)
 - ▶ they still need finite fourth moments, which the Hill estimates put in doubt
- BET-TR versus S&P 500, common days 2014–2026
 - ▶ difference 0.61, $z = 1.75$ ($p = 0.08$)
 - ▶ z : the difference divided by its standard error, compared with $N(0, 1)$
 - ▶ block-bootstrap 95% CI $[-0.01, 1.21]$ (Seminar B13)
- Calmar has no usable sampling theory
 - ▶ its drawdown grows with the sample length


 MFM_ch1_risk_measures

Risk and Return Across Markets (1/2)



- Labels: the Sharpe ratio and, in brackets, its standard error (SE)
 - ▶ $SE = \sqrt{A(1 + SR_d^2)/T}$
 - ▶ valid for i.i.d. Normal simple returns
 - ▶ full table in the [Appendix](#)
- Different sample periods: not a ranking

Risk and Return Across Markets (2/2)



- S&P 500 without dividends; BET-TR with dividends and in RON
 - ▶ RON short-term rates were well above USD rates
 - ▶ with $r_f = 0$ the BET-TR Sharpe ratio is inflated
- Heavy tails and negative skewness flatter the Sharpe ratio
 - ▶ report it with drawdowns
 - ▶ and with tail-risk measures: VaR 1% and the tail index $\hat{\alpha}$

Case Study (1/14): Research Question and Dependent Variables

- Paper: Guyon & Lekeufack (2023), *Quantitative Finance*
 - ▶ the paper builds on Bollerslev, Hood, Huss & Pedersen (2018)
- Research question
 - ▶ what share of an equity index's volatility is explained by its past daily returns alone?
 - ▶ what functional form does this dependence take?
- Three dependent variables, modelled separately (Sections 3.4–3.5 of the paper)
 - ▶ VIX: the implied volatility of the S&P 500 over the next 30 days, computed from option prices
 - ▶ VIX9D: the same measure over a 9-day horizon
 - ▶ next-day realised volatility, computed from 5-minute returns

Case Study (2/14): Requirement, Relevance and Data

- Requirement: a single model must reproduce simultaneously
 - ▶ volatility clustering (Fact 6)
 - ▶ the slow decay of the autocorrelation of $|r_t|$ (Fact 8)
 - ▶ the leverage effect (Fact 9)
- Why it matters: if volatility is mostly a function of the price path, it is mostly endogenous
 - ▶ a **path-dependent volatility (PDV)** model needs little extra randomness
 - ▶ unlike stochastic volatility models ([Chapters 11–12](#))
- Data: daily closes of the S&P 500, Euro Stoxx 50, FTSE 100, DAX and Nikkei 225
 - ▶ training 1 Jan 2000 – 31 Dec 2018
 - ▶ test 1 Jan 2019 – 15 May 2022

Case Study (3/14): Model

- Equation (3.7) in the paper: today's volatility is a linear function of two summaries of past returns

$$\sigma_t = \beta_0 + \beta_1 R_{1,t} + \beta_2 \Sigma_t$$

- Notation

- ▶ σ_t : the dependent variable, for example the VIX at the close of day t
- ▶ S_{t_i} : the index level at the close of day t_i ; $t_i \leq t$ runs over the days up to and including t
- ▶ $r_{t_i} = S_{t_i}/S_{t_{i-1}} - 1$: the simple return of day t_i
- ▶ $R_{1,t}$: the trend component, built from past returns
- ▶ Σ_t : the volatility component, built from past squared returns
- ▶ $\beta_0, \beta_1, \beta_2$: coefficients estimated from the data

Case Study (4/14): The Two Components

- ▣ Trend component: a weighted average of past returns

$$R_{1,t} = \sum_{t_i \leq t} K_1(t - t_i) r_{t_i}$$

- ▶ $t - t_i$: the age of day t_i , in days; $K_1(t - t_i)$: the weight of that day
- ▶ $R_{1,t} < 0$ after a run of falls, $R_{1,t} > 0$ after a run of rises

- ▣ Volatility component: the square root of a weighted average of squared returns

$$\Sigma_t = \left(\sum_{t_i \leq t} K_2(t - t_i) r_{t_i}^2 \right)^{1/2}$$

- ▶ Σ_t is a weighted historical volatility, with weights K_2

- ▣ Interpreting the coefficients

- ▶ $\beta_1 < 0$: volatility rises after falls ($R_{1,t} < 0$), i.e. the leverage effect
- ▶ $\beta_2 \in (0, 1)$: high recent volatility raises today's volatility, i.e. volatility clustering, as in GARCH

Case Study (5/14): Weighting Kernels

- ▣ A **kernel** $K(\tau)$ gives the weight of a day that is τ days old; the weights are positive and sum to 1
- ▣ The paper uses a **time-shifted power law (TSPL)**, eq. (3.9)

$$K(\tau) = \frac{(\tau + \delta)^{-\alpha}}{Z_{\alpha, \delta}}, \quad \alpha > 1, \delta > 0, \tau = 0, 1, \dots, 999$$

- ▶ $Z_{\alpha, \delta}$: the normalising constant, the sum of $(\tau + \delta)^{-\alpha}$ over the last 1,000 days
- ▶ α : how fast the weights decay; a small α means long memory
- ▶ δ : the time shift; a small δ concentrates the weight on the last few days
- ▣ Resulting shape
 - ▶ large weight on the last few days (short memory)
 - ▶ a slowly decaying power-law tail (long memory, Fact 8)
- ▣ K_1 and K_2 each have their own parameters: (α_1, δ_1) and (α_2, δ_2)

Case Study (6/14): Estimation and Scores


- ▣ Seven parameters: $\beta_0, \beta_1, \beta_2, \alpha_1, \delta_1, \alpha_2, \delta_2$
 - ▶ estimated by nonlinear least squares on the training set: minimise $\sum_t (\sigma_t - \hat{\sigma}_t)^2$
 - ▶ $\hat{\sigma}_t$: the model value for day t
 - ▶ estimation details: [▶ Appendix](#)
- ▣ Scores, computed separately on the training and the test set
 - ▶ RMSE (root mean squared error) = $\sqrt{\frac{1}{T} \sum_t (\sigma_t - \hat{\sigma}_t)^2}$: the typical error, in volatility units
 - ▶ $R^2 = 1 - \sum_t (\sigma_t - \hat{\sigma}_t)^2 / \sum_t (\sigma_t - \bar{\sigma})^2$: the share of the variation of σ_t explained by the model
 - ▶ $\bar{\sigma}$: the mean of σ_t ; $R^2 = 1$ is a perfect fit, $R^2 = 0$ is no better than the mean

Case Study (7/14): R^2 on the Training and Test Samples

Index	Implied volatility (paper, Table 2)		Realised volatility (paper, Table 4)	
	R^2 train	R^2 test	R^2 train	R^2 test
S&P 500 (VIX)	0.946	0.855	0.738	0.654
S&P 500 (VIX), our replication	0.947	0.863	—	—
S&P 500 (9-day VIX)	0.876	0.914	—	—
Euro Stoxx 50	0.929	0.913	0.672	0.682
Nikkei 225	0.890	0.800	0.563	0.504

- Past index returns alone explain 80–92% of the variance of implied volatility out of sample
 - ▶ FTSE 100 and DAX in the [Appendix](#)
- On our data (S&P 500 and VIX): RMSE 0.020 training and 0.034 test
 - ▶ versus 0.020 and 0.035 in Table 2 of the paper
- Realised volatility scores are lower
 - ▶ 5-minute realised volatility is a noisy target (Remark 3 in the paper)
 - ▶ this course uses no 5-minute data, so Table 4 of the paper is not replicated

Case Study (8/14): Observed and Fitted VIX




Case Study (9/14): Observed and Fitted VIX

- Fitted on 3 Jan 2000 – 31 Dec 2018 (4,779 days)
 - ▶ started from Table 3 of the paper
 - ▶ then parameters frozen
- The test set is shaded

 MFM_ch1_path_dependent_vol

Case Study (10/14): The VIX after the Test Set




Case Study (11/14): The VIX after the Test Set

- After the test set, 16 May 2022 – 18 Sep 2026: $R^2 = 0.740$, RMSE 0.026
 - ▶ a lower RMSE, but the VIX varied less (standard deviation 5.0 versus 9.3 points)
 - ▶ so R^2 falls
- Ratio observed/fitted
 - ▶ mean 0.85 in April–May 2020
 - ▶ mean 1.16 from October 2020 to May 2022
 - ▶ 8 April 2025: VIX 52.3 versus 39.2 fitted

 MFM_ch1_path_dependent_vol

Case Study (12/14): Fitted Kernels



Case Study (13/14): Fitted Kernels

- Daily weights of K_2 , log-log scale
 - ▶ $\hat{\alpha}_2 = 1.58$, $\hat{\delta}_2 = 0.050$
 - ▶ versus 1.60, 0.052 in Table 3 of the paper
- The last 5 days carry 20% of the weight
 - ▶ the days beyond lag 250 still carry 10%
 - ▶ short and long memory in one kernel

 MFM_ch1_path_dependent_vol

Case Study (14/14): Limits and Extension

- ▣ Limits stated by the authors (Section 5 of the paper)
 - ▶ a small exogenous part remains
 - ▶ the ratio of observed to fitted VIX mean-reverts around 1
 - ▶ this ratio needs a stochastic factor
- ▣ Open extension (project seed)
 - ▶ does the same formula describe the realised volatility of Bitcoin or of BET-TR?

From Evidence to Model Requirements

Evidence (five markets)	Model requirement	Where in this course
quantile (Crow–Siddiqui) excess kurtosis > 0 in all five markets; Hill $\hat{\alpha} \approx 2.5\text{--}3.5$	heavy-tailed innovations, also conditionally	Ch. 5, Ch. 7 (extreme value theory)
ACF of $ r_t $ positive for months; $\hat{d} \approx 0.35\text{--}0.50$	persistent, time-varying volatility; long memory or breaks	Ch. 5, Ch. 9
$L(1) = -0.12$ (S&P 500, BET-TR); reversed for EUR/RON	asymmetric response to bad news where it exists	Ch. 5 (GJR-GARCH, EGARCH)
robust $\hat{\rho}_1$ economically small except EUR/RON	approximately uncorrelated returns; market microstructure matters	Ch. 2, Ch. 10
quantile skewness negative only for the S&P 500	asymmetry is market-specific	Ch. 7–8

- The random walk with i.i.d. increments from the Normal distribution (Bachelier–Black–Scholes world) violates Facts 2, 3 and 5–11
 - Fact 4 holds trivially
- ARCH-LM significance supports volatility clustering descriptively
 - with doubtful fourth moments it is not a formal proof

Open Question: Are Tail Indices Stable over Time? (1/2)

- The tail index α is usually treated as a constant
- Crypto-assets changed in structure
 - ▶ futures (2017), spot Bitcoin ETFs (January 2024), institutional holders
 - ▶ [Pele et al. \(2023\)](#): the tail factor separates cryptos from other asset classes
- Hypothesis proposed by an AI assistant
 - ▶ after the launch of the spot ETFs (January 2024), the entry of institutional investors has made the tails of Bitcoin's return distribution thinner
 - ▶ i.e. it has raised the tail index α

Open Question: Are Tail Indices Stable over Time? (2/2)

- ▣ Discovery loop in one pass
 - ▶ **literature**: ask for peer-reviewed work with DOIs, then resolve each DOI yourself
 - ▶ [Hoga \(2017\)](#): change-point tests for the tail index valid under β -mixing dependence
 - ▶ β -mixing: dependence that vanishes with the distance in time, e.g. in GARCH models
 - ▶ **hypotheses**: a change in tail shape, a fall in volatility only (scale does not move α), sampling noise
 - ▶ **data and code**: replicate the chapter's Hill estimates first (Bitcoin 4.2, S&P 500 3.1)
- ▣ Fixed in advance
 - ▶ $k = 2.5\%$ of n , the break date 11 January 2024
 - ▶ block-bootstrap SEs, a placebo asset

Mini-Case: Did Bitcoin's Tail Get Thinner after the ETFs? (1/2)

- Test: Hill estimate for $|r_t|$, $k = 2.5\%$ of n
 - ▶ $H_0 : \alpha_{\text{post}} = \alpha_{\text{pre}}$ against $H_1 : \alpha_{\text{post}} > \alpha_{\text{pre}}$
 - ▶ $\hat{\alpha}_{\text{pre}}, \hat{\alpha}_{\text{post}}$: the Hill estimates in the two subperiods; s : their standard errors, from a moving-block bootstrap (blocks of 20 days)
 - ▶ $z = (\hat{\alpha}_{\text{post}} - \hat{\alpha}_{\text{pre}}) / \sqrt{s_{\text{pre}}^2 + s_{\text{post}}^2}$: the difference divided by its standard error
 - ▶ under H_0 , $z \approx N(0, 1)$; reject H_0 at 5% if $z > 1.645$
- Results
 - ▶ before (Sep 2014 – 10 Jan 2024, $k = 85$): $\hat{\alpha} = 4.16$ ($s = 0.52$)
 - ▶ after (11 Jan 2024 – Sep 2026, $k = 24$): $\hat{\alpha} = 4.31$ ($s = 0.91$)
 - ▶ $z = 0.15$
- Placebo, S&P 500 at the same date: $2.82 \rightarrow 2.86$, $z = 0.03$

Mini-Case: Did Bitcoin's Tail Get Thinner after the ETFs? (2/2)

- ▣ **Minimum detectable change**, one-sided 5% test with 80% power
 - ▶ $(1.645 + 0.842) \times 1.05 = 2.6$
 - ▶ 1.645: the one-sided 5% critical value; 0.842: the 80% quantile of $N(0, 1)$, for 80% power
 - ▶ $1.05 = \sqrt{0.52^2 + 0.91^2}$ is the SE of the difference
 - ▶ the Normal-distribution approximation is rough with 24 tail observations
 - ▶ the data cannot distinguish an unchanged tail from a large thinning of the tail
- ▣ A non-significant result does not show the absence of an effect
 - ▶ wait for more data or pool assets

Project Seed

- **Question:** did the tail index of Bitcoin and Ethereum change with institutionalisation (futures 2017, spot ETFs 2024)?
 - ▶ raw and GARCH-standardised returns (shape versus scale)
 - ▶ loss and gain tails separately
- **Replicate first:** the chapter's Hill table and the mini-case numbers
- **Pre-register**, committed before the results
 - ▶ H_0 , H_1 , the rule for k , break dates and bootstrap
- **Report:** the change-point test of [Hoga \(2017\)](#) as a Quantlet
 - ▶ estimates, both SEs, the minimum detectable change
- **Log:** prompts in AI_USE.md
 - ▶ every caught error (an i.i.d. SE, a date chosen from the data) in AI_ERRORS.md
 - ▶ reusable prompts (▶ [Appendix](#)) and checks (▶ [Appendix](#))

Key Takeaways and the Link to Chapter 2

Takeaways

- ▣ Returns are uncorrelated but not independent: their size is predictable
- ▣ Heavy tails are the rule: tail indices near 3
 - ▶ the fourth moment very likely does not exist
 - ▶ kurtosis-based inference fails
- ▣ Volatility is persistent (long memory or breaks)
 - ▶ in equity markets it reacts asymmetrically to losses
- ▣ Every test needs a reference distribution that the data may violate

Check yourself

- ▣ Why does the $Q(10)$ on Bitcoin returns have $p = 0.028$ but the robust $\tilde{Q}(10)$ $p = 0.29$?
- ▣ Why can the tail index be estimated with the Hill estimator when the kurtosis does not exist?
- ▣ **Chapter 2:** are returns predictable at all?

Appendix: Formula Sheet with Assumptions

Returns and annualisation

- ▣ $R_t = P_t/P_{t-1} - 1$, $r_t = \ln(1 + R_t)$, $r_t(k) = \sum_{j=0}^{k-1} r_{t-j}$,
 $R_{p,t} = \sum_i w_i R_{i,t}$
- ▣ $\text{CAGR} = (P_T/P_0)^{1/Y} - 1$
- ▣ $\sigma_{\text{ann}} = \sqrt{A} \hat{\sigma}$: uncorrelated returns, constant variance

Tests

- ▣ $JB = T[\hat{S}^2/6 + (\hat{K} - 3)^2/24]$: χ_2^2 under i.i.d. Normal
- ▣ $Q(m) = T(T+2) \sum_{k=1}^m \hat{\rho}_k^2 / (T-k)$: χ_m^2 under i.i.d.
- ▣ $\tilde{Q}(m) = \sum_k \hat{\rho}_k^2 / \hat{\tau}_k$: $\mathbb{E}r^4 < \infty$, zero cross terms
- ▣ Hill $\text{SE}(\hat{\alpha}) \approx \hat{\alpha}/\sqrt{k}$: independent Pareto tail

Volatility

- ▣ Parkinson $\frac{1}{4 \ln 2} \overline{(h-l)^2}$, the bar a mean over days:
 driftless, no overnight jumps
- ▣ Mixture $K = 3 \mathbb{E}[\sigma^4] / \mathbb{E}[\sigma^2]^2$: σ_t independent of ε_t

Performance

- ▣ $SR = (\mu - r_f)/\sigma$ on simple returns
 ▶ $SE \approx 1/\sqrt{Y}$ for i.i.d. Normal returns
- ▣ $SoR = (\mu - r_f)/\sigma^-$
- ▣ $DD_t = P_t / \max_{s \leq t} P_s - 1$; Calmar = $\text{CAGR}/|MDD|$

Appendix: Prices Before Computers, the Stock Ticker



Stock ticker designed by Thomas Edison
(Smithsonian collection)

Photo: Jaclyn Nash, Smithsonian National Museum of
American History; public domain; Wikimedia Commons

- Before computers, trades reached brokers as prices printed on a paper **ticker tape**
- A daily close is one snapshot of this stream of prices
 - ▶ returns are built from such snapshots
- The choice of snapshot (close, fixing, a midnight snapshot) matters
 - ▶ as [Chapter 0](#) showed
- Prices at the level of individual trades: market microstructure, [Chapter 10](#)

Appendix: The People Behind the Stylised Facts (1/2)



Carl Friedrich Gauss (1777–1855)

Painting: C.A. Jensen (1840); public domain;
[Wikimedia Commons](#)

- The Normal distribution: the benchmark that returns violate



Louis Bachelier (1870–1946)

Photo: unknown author; public domain;
[Wikimedia Commons](#)

- Prices as a random walk, in his thesis [Bachelier \(1900\)](#)




William S. Gosset, "Student" (1876–1937)

Photo: unknown author; public domain;
[Wikimedia Commons](#)

- The Student-t distribution [Student \(1908\)](#), today a model for heavy tails



Appendix: The People Behind the Stylised Facts (2/2)



Vilfredo Pareto (1848–1923)

Photo: unknown author (1870s); public domain; [Wikimedia Commons](#)



Benoît Mandelbrot (1924–2010)

Photo: Rama (2007); CC BY-SA 2.0 France; [Wikimedia Commons](#)

- Pareto: **power laws** for large incomes, the ancestor of heavy-tail models
- Mandelbrot (1963): cotton prices
 - ▶ heavy tails, which the Normal distribution does not describe
- The exponent α of these power laws
 - ▶ is the tail index estimated with the Hill method

Appendix: Cont's Facts in Recent Studies

- Ratliff-Crain, Van Oort, Koehler & Tivnan (2025)
 - ▶ re-test all 11 facts on intraday returns of the Dow 30 stocks, October 2018 – March 2019
 - ▶ conclusive evidence for 8 facts, no support for the remaining 3
- Heavy tails beyond the Student- t : estimation of alpha-stable distributions, whose tail index is below 2 [Pele \(2014\)](#)
- Entropy-based measures of tail risk as an alternative to kurtosis [Pele, Lazar & Dufour \(2017\)](#)
- Crypto-assets: the same facts, with a tail factor that separates them from other asset classes [Pele et al. \(2023\)](#)

Appendix: Annual Variance under Autocorrelation (1/2)

- The variance of a sum of h daily returns depends on the autocorrelations

$$\text{Var}\left(\sum_{j=1}^h r_{t+j}\right) = h\sigma^2 \left[1 + 2 \sum_{j=1}^{h-1} (1 - j/h) \rho_j\right]$$

- ▶ h : the horizon in days; σ^2 : the daily variance; ρ_j : the autocorrelation at lag j
- ▶ assumption: weakly stationary returns (constant mean and variance, autocovariances depending only on the lag)
- ▶ without autocorrelation ($\rho_j = 0$) it is $h\sigma^2$, i.e. the $\sqrt{h}$ rule for volatility
- **Variance ratio** $VR(h) = \text{Var}(h\text{-day sum})/(h\sigma^2)$
 - ▶ $VR(h) = 1$: no autocorrelation; $VR(h) > 1$: positive autocorrelation; $VR(h) < 1$: negative autocorrelation
 - ▶ volatility over horizon h is the $\sqrt{h}$ -rule value multiplied by $\sqrt{VR(h)}$

Appendix: Annual Variance under Autocorrelation (2/2)

- Autoregression of order 1, AR(1), with $\rho_j = \rho^j$: $\sqrt{VR(h)} \rightarrow \sqrt{(1+\rho)/(1-\rho)}$ as h grows
 - ▶ ρ : the lag-1 autocorrelation, $|\rho| < 1$
 - ▶ EUR/RON, $\rho_1 = 0.17$: factor 1.19, annual volatility 4.4% $\rightarrow$ 5.3%
 - ▶ S&P 500, $\rho_1 = -0.079$: factor 0.92
- Using the sample lag-1 autocorrelation as an AR(1) parameter is an assumption
 - ▶ check the whole ACF (Seminar A2)
 - ▶ variance-ratio tests: [Chapter 2](#)

Appendix: The Null Hypothesis of JB, and Tests for Time Series

- ▣ JB is a Lagrange Multiplier test [Jarque & Bera \(1980\)](#)
 - ▶ under i.i.d. returns from the Normal distribution, $\sqrt{T}(\hat{S}, \hat{K} - 3) \rightarrow N(0, \text{diag}(6, 24))$, hence χ^2_2
 - ▶ $\text{diag}(6, 24)$: a diagonal covariance matrix; $\hat{S}$ and $\hat{K}$ are asymptotically independent, with variances $6/T$ and $24/T$
 - ▶ JB is the sum of the two squared standardised statistics, hence χ^2 with 2 degrees of freedom
- ▣ GARCH-type dependence changes $\text{Var}(\hat{S})$ and $\text{Var}(\hat{K})$
 - ▶ a rejection mixes the marginal distribution and independence
- ▣ [Bai & Ng \(2005\)](#): tests of skewness, kurtosis and normality of the marginal distribution of a time series
 - ▶ studentised with HAC variances
 - ▶ their separate skewness and kurtosis tests (null $S = 0$ or $K = 3$ without normality) need $\mathbb{E}r^6 < \infty$ and $\mathbb{E}r^8 < \infty$

Appendix: The Robust Portmanteau Test and Its Assumptions

- $\tilde{Q}(m) = \sum_{k=1}^m \hat{\rho}_k^2 / \hat{\tau}_k \sim \chi_m^2$ Lobato, Nankervis & Savin (2001) needs $\mathbb{E}r^4 < \infty$
 - ▶ $\sim$: “is distributed as”, here asymptotically
 - ▶ and zero cross terms: $\mathbb{E}[x_t^2 x_{t-k} x_{t-l}] = 0$ for two different lags $k \neq l$
 - ▶ then the $\hat{\rho}_k$ at different lags are asymptotically uncorrelated and the sum is χ_m^2
- Dalla, Giraitis & Phillips (2022) keep only the significant cross terms in the covariance of the $\hat{\rho}_k$
 - ▶ none for the S&P 500, BET-TR and Bitcoin
- $\hat{\tau}_k$ is the variance under $\rho_k = 0$
 - ▶ to compare a non-zero ρ_1 across periods use the HAC SE of the slope of r_t on r_{t-1} (Seminar B6)
- ACF of $|r_t|$ and of r_t^2 : Normal limits need $\alpha > 4$ and $\alpha > 8$ Davis & Mikosch (1998)

Appendix: Hill Estimator, Asymptotic Distribution and Alternatives (1/2)

- ▣ Hill = maximum likelihood estimator of a Pareto tail above $u = X_{(k+1)}$, conditional on u (Seminar A11)
- ▣ If $k \rightarrow \infty$ and $k/n \rightarrow 0$, and the tail is Pareto up to a second-order term

$$\sqrt{k}(\hat{\alpha}_k - \alpha) \rightarrow N(0, \alpha^2)$$

- ▶ second-order term: a correction to the power law that vanishes as $x \rightarrow \infty$
- ▶ consequence: $SE(\hat{\alpha}_k) \approx \alpha/\sqrt{k}$, the formula in the main part
- ▶ if k grows too fast, a bias term appears: the tail is only approximately Pareto
- ▣ Dependent data (GARCH-type): $\hat{\alpha}$ stays consistent
 - ▶ but extremes cluster and the variance exceeds α^2/k [Hsing \(1991\)](#)

Appendix: Hill Estimator, Asymptotic Distribution and Alternatives (2/2)

- ▣ Choosing k : plot $\hat{\alpha}$ against $\log k$ ("alt-Hill" plot) Drees, Resnick & de Haan (2000)
 - ▶ the log scale stretches the region of small k , where the actual tail lies
- ▣ Regression alternative Gabaix & Ibragimov (2011)
 - ▶ log rank on log size, with rank shifted by $-\frac{1}{2}$: $\ln(i - \frac{1}{2}) = a - b \ln X_{(i)}$, $i = 1, \dots, k$
 - ▶ the slope $\hat{b}$ estimates α ; the shift by $\frac{1}{2}$ reduces the small-sample bias
 - ▶ $SE = \hat{\alpha} \sqrt{2/k}$

Appendix: Two Tails and the Tail of $|r_t|$

- $P(|X| > x) = P(X > x) + P(X < -x)$
 - ▶ with indices α_- (losses) and α_+ (gains), $|X|$ has index $\min(\alpha_-, \alpha_+)$
- At the same k the estimates are not comparable
 - ▶ the k largest $|r_t|$ come from both tails
 - ▶ common threshold u : the lower of the two one-sided thresholds at $k = 2.5\%$ of n

	u (%)	k for $ r_t $	$\hat{\alpha}, r_t $	$\hat{\alpha}_-, \text{losses}$	$\hat{\alpha}_+, \text{gains}$
S&P 500	2.17	508	2.90	2.85	2.96
BET-TR	1.79	153	2.54	2.20	3.06
Bitcoin	7.11	233	3.02	2.71	3.36
EUR/RON	0.54	295	2.28	2.63	2.04
Gold	1.98	513	3.11	2.97	3.28

- At a common u , the $|r_t|$ estimate lies between the two tails
 - ▶ Bitcoin's 4.2 at $k = 2.5\%$ reflects second-order bias
 - ▶ report both one-sided indices

Appendix: Estimating d , the Formulas (1/2)

- ▣ Fourier frequencies $\lambda_j = 2\pi j/T$, $j = 1, \dots, m$; the m lowest are used, here $m = T^{0.65}$
- ▣ The **periodogram**: the sample estimator of the spectral density

$$I(\lambda_j) = \frac{1}{2\pi T} \left| \sum_{t=1}^T x_t e^{-i\lambda_j t} \right|^2$$

- ▶ x_t : the demeaned series (here $|r_t|$ minus its mean); i : the imaginary unit
- ▶ $I(\lambda_j)$: how much of the variation in x_t comes from the cycle of frequency λ_j
- ▣ Log-periodogram regression [Geweke & Porter-Hudak \(1983\)](#): regress $\ln I(\lambda_j)$ on $-2 \ln(2 \sin(\lambda_j/2))$
 - ▶ the reason: $\ln f(\lambda) \approx \ln G - 2d \ln \lambda$, and $2 \sin(\lambda/2) \approx \lambda$ at low frequencies
 - ▶ the slope is $\hat{d}$, $SE = \pi/\sqrt{24m}$

Appendix: Estimating d , the Formulas (2/2)

- Local Whittle [Robinson \(1995\)](#): $\hat{d}$ minimises an objective built from the approximation $f(\lambda) \approx G\lambda^{-2d}$

$$\hat{d} = \arg \min_d \left[\ln \left(\frac{1}{m} \sum_{j=1}^m \lambda_j^{2d} I(\lambda_j) \right) - \frac{2d}{m} \sum_{j=1}^m \ln \lambda_j \right]$$

- ▶ $\arg \min_d$: the value of d that minimises the expression; the constant G has been eliminated analytically
- ▶ $\sqrt{m}(\hat{d} - d) \rightarrow N(0, \frac{1}{4})$, so $SE = 1/(2\sqrt{m})$
- Local Whittle is robust to non-Gaussianity
 - ▶ a larger m lowers the SE
 - ▶ but lets short-run (GARCH) dynamics bias $\hat{d}$ (Seminar B15)

Appendix: Markets in Detail (1/2)

- ▣ **Crypto-assets:** an asset class with its own statistical profile [Härdle, Harvey & Reule \(2020\)](#)
 - ▶ tail, memory and moment factors classify cryptos as a separate asset class [Pele et al. \(2023\)](#)
 - ▶ the separation comes mainly through the tail factor
 - ▶ crypto modelling and the CRIX (CRypto IndeX): [Härdle, Chen & Overbeck \(2017\)](#); Chapter 16
- ▣ **EUR/RON:** BNR reference rate from 1 July 2005
 - ▶ the new leu (RON, 1 RON = 10,000 old lei)
 - ▶ long calm periods interrupted by sudden adjustments
 - ▶ EUR/RON market (interbank) quote series showed spurious negative autocorrelation (quote errors)

Appendix: Markets in Detail (2/2)

- ▣ **BET-TR**: the bank-asset tax was announced on 18 December 2018
 - ▶ adopted on 28 December 2018 as Government Emergency Ordinance (GEO) 114/2018
 - ▶ policy shocks produce single-day crashes without comparable rallies
- ▣ **Non-synchronous trading**: index constituents do not all trade at every moment
 - ▶ so the index reflects news with a delay
 - ▶ creating positive autocorrelation [Campbell, Lo & MacKinlay \(1997\)](#), ch. 3

Appendix: The Maximum Drawdown of a Brownian Motion

- ▣ Driftless Brownian motion of the log price with annual volatility σ over T years
 - ▶ here T is the horizon in years, not the number of days
- ▣ The maximum **log** drawdown $D = -\ln(1 + MDD) \geq 0$, the largest fall in the log price, has

$$\mathbb{E}[D] = \sqrt{\pi/2} \sigma \sqrt{T} \quad \text{Magdon-Ismail, Atiya, Pratap \& Abu-Mostafa (2004)}$$

- ▣ $MDD = e^{-D} - 1$; $\sqrt{\pi/2} \approx 1.25$
 - ▶ without drift the expected maximum drawdown grows in proportion to $\sqrt{T}$
 - ▶ with a positive drift $\mathbb{E}[D]$ grows only like $\ln T$
- ▣ Gold 1990–2026: $\sigma = 16.3\%$, $T = 36.7$
 - ▶ benchmark $\mathbb{E}[D] = 1.24$ versus observed 0.592 (−44.7%) (Seminar A4)
- ▣ Drawdowns and Calmar ratios from samples of different length are not comparable

Appendix: Sharpe Ratios beyond the Normal Distribution

- i.i.d., any distribution with $\mathbb{E}r^4 < \infty$ [Opdyke \(2007\)](#)
 - ▶ $\text{Var}(\widehat{SR}_d) \approx (1 - S SR_d + \frac{K-1}{4} SR_d^2) / T$
 - ▶ S, K : the skewness and kurtosis of daily returns; for the Normal distribution ($S = 0, K = 3$) this gives Lo's formula
 - ▶ negative skewness and high kurtosis raise the variance of $\widehat{SR}_d$ when $SR_d > 0$
 - ▶ here the SEs change by at most 0.02
- The Deflated Sharpe Ratio corrects for non-normality and multiple testing ([Chapter 13](#))

	CAGR	Sharpe	Sortino	Calmar
S&P 500	8.7%	0.55	0.78	0.15
BET-TR	21.2%	1.34	1.87	0.68
Bitcoin	53.9%	0.98	1.45	0.65
EUR/RON	1.8%	0.43	0.65	0.12
Gold	6.7%	0.48	0.69	0.15

Appendix: Case Study, Estimation Details (1/2)

- Targets (Sections 3.4–3.5 of the paper)
 - ▶ the 30-day implied volatility index of each index and the 9-day VIX
 - ▶ realised volatility of day $t + 1$
 - ▶ volatility in decimals ($0.03 = 3$ volatility points)
- Kernels, eq. (3.9) in the paper: $K(\tau) = Z_{\alpha, \delta}^{-1}(\tau + \delta)^{-\alpha}$
 - ▶ normalised so that $\sum_i K(i\Delta t)\Delta t = 1$ on the 1,000-day window
 - ▶ i : the age of the day, in days; $\Delta t = 1/252$: one day, in years
- Estimation: $(\beta_0, \beta_1, \beta_2, \alpha_1, \delta_1, \alpha_2, \delta_2)$ by nonlinear least squares on the training set
 - ▶ initialised using a lasso regression on features built with exponential kernels

Appendix: Case Study, Estimation Details (2/2)

- Our replication, started from the values in Table 3 of the paper (in parentheses)
 - ▶ $\hat{\beta}_0 = 0.052$ (0.057), $\hat{\beta}_1 = -0.090$ (-0.095)
 - ▶ $\hat{\beta}_2 = 0.868$ (0.82), $\hat{\alpha}_1 = 1.13$ (1.06)
 - ▶ the parameters differ slightly, but R^2 and RMSE are practically the same
- The 9-day VIX exists only from 2011 (shorter training set)
- The restricted models ($\beta_1 = 0$, $\beta_2 = 0$) are Seminar C3

Appendix: Case Study, R^2 for All Indices (1/2)

Index	Implied volatility (paper, Table 2)		Realised volatility (paper, Table 4)	
	R^2 train	R^2 test	R^2 train	R^2 test
S&P 500 (VIX)	0.946	0.855	0.738	0.654
S&P 500 (9-day VIX)	0.876	0.914	–	–
Euro Stoxx 50	0.929	0.913	0.672	0.682
FTSE 100	0.925	0.870	0.650	0.617
DAX	0.934	0.918	0.722	0.557
Nikkei 225	0.890	0.800	0.563	0.504

- VIX parameters: $\beta_0 = 0.057$, $\beta_1 = -0.095$, $\beta_2 = 0.82$
- Across all targets $\alpha_2 \in [1.2, 2.3]$ (Tables 3 and 5 of the paper)
 - long memory near criticality (Section 3.2 of the paper)

Appendix: Case Study, R^2 for All Indices (2/2)

- Dropping Σ_t hurts the most
- The model beats two ARCH benchmarks and four published path-dependent models
 - ▶ Figures 3.8–3.9 and Appendix C of the paper
- Its Markovian 4-factor version, driven by Brownian motion (Section 4.6 and Table 9 of the paper)
 - ▶ gives estimated values of about 0.12–0.17 of the Hurst exponent H for daily realised volatility, as in data
 - ▶ $H < 0.5$: a path “rougher” than Brownian motion
 - ▶ a model without rough volatility reproduces the small values of H estimated from daily data

Appendix: The Discovery Loop, Reusable Prompts (1/2)

□ Literature

- ▶ Find peer-reviewed work since 2015 on time variation or structural breaks in tail indices of asset returns, with tests valid under GARCH-type dependence.
- ▶ For each: DOI, estimator, test, data.
- ▶ Flag entries you are unsure of.

□ Hypothesis

- ▶ State competing hypotheses for a change in Bitcoin's tail after the 2024 spot ETFs.
- ▶ What does each imply for the Hill estimate of the tail index of raw and of GARCH-standardised returns?

Appendix: The Discovery Loop, Reusable Prompts (2/2)

□ Data and code

- ▶ Use daily Bitcoin closing prices in USD; log returns on its own calendar;
- ▶ Hill estimate of the tail index for $|r|$, with $k = 2.5\%$ of n , before and after 11 January 2024; moving-block bootstrap standard errors.

□ Robustness

- ▶ Write a pre-registered plan: k from 1% to 5% of n , loss and gain tails, GARCH-standardised residuals, block lengths 10 to 50 days, a placebo date and a placebo asset.

□ Critique

- ▶ Referee this design.
- ▶ With 24 tail observations after the break, what is the power against a change of 1 in the tail index?

Appendix: Mandatory Checks for the Researcher (1/2)

- **Wrong return or price series:** reproduce the chapter's Hill table before anything else
- **Choice of k after seeing the result**
 - ▶ fix the rule ($k = 2.5\%$ of n) in advance
 - ▶ show the Hill plot for 1–5%
- **i.i.d. standard errors:** $\hat{\alpha}/\sqrt{k}$ ignores volatility clustering
 - ▶ compare with a block bootstrap
 - ▶ for the S&P 500 before 2024 the bootstrap SE is 2.2 times larger

Appendix: Mandatory Checks for the Researcher (2/2)

- **Break date chosen from the data**
 - ▶ use a date fixed by the event
 - ▶ or a sup-type change-point test with its own critical values [Hoga \(2017\)](#)
- **Bad ticks create fake tails**
 - ▶ list the ten largest $|r_t|$ and match each to a dated event
- **Low power:** report the minimum detectable change
 - ▶ 2.8 SEs of the difference for a two-sided 5% test with 80% power
 - ▶ 2.5 SEs for a one-sided test

Appendix: Practical Checklist for Empirical Work (1/2)

- ▣ **Know your data**
 - ▶ price vs total-return index
 - ▶ adjusted vs raw prices
 - ▶ fixing vs traded price
 - ▶ official reference rates where possible (EUR/RON)
- ▣ **Clean before you model**
 - ▶ non-trading days, holidays, stale quotes, obvious errors
 - ▶ match the largest moves to real events
- ▣ **Use log returns** for single-asset dynamics, simple returns for portfolios

Appendix: Practical Checklist for Empirical Work (2/2)

- ▣ **Report robust statistics:** quantiles, the number of tail days and the tail index $\hat{\alpha}$ next to kurtosis
- ▣ **Annualise carefully**
 - ▶ 252 vs 365 periods
 - ▶ the $\sqrt{A}$ rule needs uncorrelated returns with constant variance
- ▣ **Reproducibility:** code and data on GitHub, as Quantlets

Appendix: Further Study on Quantinar

Advanced courses on quantinar.com

- ▣  Statistics of Financial Markets
- ▣  Tukey's g- and h- transformations
- ▣  Applied Time Series Analysis with Python

Code, data and essential reading

- ▣ All charts and tables: github.com/danpele/MFM/tree/main/Quantlets/Ch_01
- ▣ Essential reading: Cont (2001); Franke, Härdle & Hafner (2019); Campbell, Lo & MacKinlay (1997), ch. 1–2

References (1/4)

- Andersen, T.G., Bollerslev, T., Diebold, F.X., Ebens, H. (2001). The distribution of realized stock return volatility. *Journal of Financial Economics*, 61(1), 43–76.
- Bachelier, L. (1900). *Théorie de la spéculation*. *Annales scientifiques de l'École Normale Supérieure*, 17, 21–86.
- Bai, J., Ng, S. (2005). Tests for skewness, kurtosis, and normality for time series data. *Journal of Business & Economic Statistics*, 23(1), 49–60.
- Bollerslev, T., Hood, B., Huss, J., Pedersen, L.H. (2018). Risk everywhere: modeling and managing volatility. *Review of Financial Studies*, 31(7), 2729–2773.
- Bouchaud, J.-P., Matacz, A., Potters, M. (2001). Leverage effect in financial markets: the retarded volatility model. *Physical Review Letters*, 87(22), 228701.
- Campbell, J.Y., Lo, A.W., MacKinlay, A.C. (1997). *The Econometrics of Financial Markets*. Princeton University Press.
- Christie, A.A. (1982). The stochastic behavior of common stock variances: value, leverage and interest rate effects. *Journal of Financial Economics*, 10(4), 407–432.
- Clark, P.K. (1973). A subordinated stochastic process model with finite variance for speculative prices. *Econometrica*, 41(1), 135–155.
- Cong, L.W., Li, X., Tang, K., Yang, Y. (2023). Crypto wash trading. *Management Science*, 69(11), 6427–6454.
- Cont, R. (2001). Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2), 223–236.
- Dalla, V., Giraitis, L., Phillips, P.C.B. (2022). Robust tests for white noise and cross-correlation. *Econometric Theory*, 38(5), 913–941.
- Davis, R.A., Mikosch, T. (1998). The sample autocorrelations of heavy-tailed processes with applications to ARCH. *Annals of Statistics*, 26(5), 2049–2080.
- Diebold, F.X., Inoue, A. (2001). Long memory and regime switching. *Journal of Econometrics*, 105(1), 131–159.
- Ding, Z., Granger, C.W.J., Engle, R.F. (1993). A long memory property of stock market returns and a new model. *Journal of Empirical Finance*, 1(1), 83–106.

References (2/4)

- Drees, H., Resnick, S., de Haan, L. (2000). How to make a Hill plot. *Annals of Statistics*, 28(1), 254–274.
- Engle, R.F. (1982). Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4), 987–1007.
- Fama, E.F. (1965). The behavior of stock-market prices. *Journal of Business*, 38(1), 34–105.
- Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- Gabaix, X., Ibragimov, R. (2011). Rank – 1/2: a simple way to improve the OLS estimation of tail exponents. *Journal of Business & Economic Statistics*, 29(1), 24–39.
- Garman, M.B., Klass, M.J. (1980). On the estimation of security price volatilities from historical data. *Journal of Business*, 53(1), 67–78.
- Geweke, J., Porter-Hudak, S. (1983). The estimation and application of long memory time series models. *Journal of Time Series Analysis*, 4(4), 221–238.
- Granger, C.W.J., Ding, Z. (1995). Some properties of absolute return: an alternative measure of risk. *Annales d'Économie et de Statistique*, 40, 67–91.
- Guyon, J., Lekeufack, J. (2023). Volatility is (mostly) path-dependent. *Quantitative Finance*, 23(9), 1221–1258.
- Härdle, W.K., Chen, C.Y.-H., Overbeck, L. (Eds.) (2017). *Applied Quantitative Finance*, 3rd ed. Springer.
- Härdle, W.K., Harvey, C.R., Reule, R.C.G. (2020). Understanding cryptocurrencies. *Journal of Financial Econometrics*, 18(2), 181–208.
- Hill, B.M. (1975). A simple general approach to inference about the tail of a distribution. *Annals of Statistics*, 3(5), 1163–1174.
- Hoga, Y. (2017). Change point tests for the tail index of β -mixing random variables. *Econometric Theory*, 33(4), 915–954.
- Hsing, T. (1991). On tail index estimation using dependent data. *Annals of Statistics*, 19(3), 1547–1569.

References (3/4)

- Inclán, C., Tiao, G.C. (1994). Use of cumulative sums of squares for retrospective detection of changes of variance. *Journal of the American Statistical Association*, 89(427), 913–923.
- Jarque, C.M., Bera, A.K. (1980). Efficient tests for normality, homoscedasticity and serial independence of regression residuals. *Economics Letters*, 6(3), 255–259.
- Karpoff, J.M. (1987). The relation between price changes and trading volume: a survey. *Journal of Financial and Quantitative Analysis*, 22(1), 109–126.
- Kim, T.-H., White, H. (2004). On more robust estimation of skewness and kurtosis. *Finance Research Letters*, 1(1), 56–73.
- Ledoit, O., Wolf, M. (2008). Robust performance hypothesis testing with the Sharpe ratio. *Journal of Empirical Finance*, 15(5), 850–859.
- Ljung, G.M., Box, G.E.P. (1978). On a measure of lack of fit in time series models. *Biometrika*, 65(2), 297–303.
- Lo, A.W. (2002). The statistics of Sharpe ratios. *Financial Analysts Journal*, 58(4), 36–52.
- Lobato, I., Nankervis, J.C., Savin, N.E. (2001). Testing for autocorrelation using a modified Box–Pierce Q test. *International Economic Review*, 42(1), 187–205.
- Magdon-Ismail, M., Atiya, A.F., Pratap, A., Abu-Mostafa, Y.S. (2004). On the maximum drawdown of a Brownian motion. *Journal of Applied Probability*, 41(1), 147–161.
- Mandelbrot, B. (1963). The variation of certain speculative prices. *Journal of Business*, 36(4), 394–419.
- Mikosch, T., Stărică, C. (2004). Nonstationarities in financial time series, the long-range dependence, and the IGARCH effects. *Review of Economics and Statistics*, 86(1), 378–390.
- Opdyke, J.D. (2007). Comparing Sharpe ratios: so where are the p-values?. *Journal of Asset Management*, 8(5), 308–336.
- Parkinson, M. (1980). The extreme value method for estimating the variance of the rate of return. *Journal of Business*, 53(1), 61–65.

References (4/4)

- ▣ Pele, D.T. (2014). [A SAS approach for estimating the parameters of an alpha-stable distribution](#). *Procedia Economics and Finance*, 10, 68–77.
- ▣ Pele, D.T., Lazar, E., Dufour, A. (2017). [Information entropy and measures of market risk](#). *Entropy*, 19(5), 226.
- ▣ Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). [Are cryptos becoming alternative assets?](#). *The European Journal of Finance*, 29(10), 1064–1105.
- ▣ Praetz, P.D. (1972). [The distribution of share price changes](#). *Journal of Business*, 45(1), 49–55.
- ▣ Ratliff-Crain, E., Van Oort, C.M., Koehler, M., Tivnan, B.F. (2025). [Revisiting Cont's stylized facts for modern stock markets](#). *Quantitative Finance*, 25(9), 1343–1373.
- ▣ Robinson, P.M. (1995). [Gaussian semiparametric estimation of long range dependence](#). *Annals of Statistics*, 23(5), 1630–1661.
- ▣ Rogers, L.C.G., Satchell, S.E. (1991). [Estimating variance from high, low and closing prices](#). *Annals of Applied Probability*, 1(4), 504–512.
- ▣ Romano, J.P., Thombs, L.A. (1996). [Inference for autocorrelations under weak assumptions](#). *Journal of the American Statistical Association*, 91(434), 590–600.
- ▣ Sharpe, W.F. (1966). [Mutual fund performance](#). *Journal of Business*, 39(1), 119–138.
- ▣ Sharpe, W.F. (1994). [The Sharpe ratio](#). *Journal of Portfolio Management*, 21(1), 49–58.
- ▣ Sortino, F.A., van der Meer, R. (1991). [Downside risk](#). *Journal of Portfolio Management*, 17(4), 27–31.
- ▣ Student (1908). [The probable error of a mean](#). *Biometrika*, 6(1), 1–25.
- ▣ Taylor, S.J. (1986). *Modelling Financial Time Series*. Wiley; 2nd ed. (2008), World Scientific, doi:10.1142/6578.
- ▣ Yang, D., Zhang, Q. (2000). [Drift-independent volatility estimation based on high, low, open, and close prices](#). *Journal of Business*, 73(3), 477–492.