

Modelling Financial Markets

Chapter 2: Market Efficiency — From EMH to Adaptive Markets



Daniel Traian PELE

Bucharest University of Economic Studies
IDA Institute Digital Assets
Blockchain Research Center
AI4EFin Artificial Intelligence for Energy Finance
Romanian Academy, Institute for Economic Forecasting
MSCA Digital Finance

Chapter outline

- Can past prices, news or other public information predict returns after adjusting for risk and costs?
- Sections 1–5 contain the core material and end with a short check.
- The further-study topics contain the derivations, the other tests and the project topics.

- 1. Returns and information
- 2. Autocorrelation and variance ratios
- 3. Out-of-sample predictability
- 4. Event studies
- 5. Multiple testing and anomalies
- Topics for further study

1. Returns and information

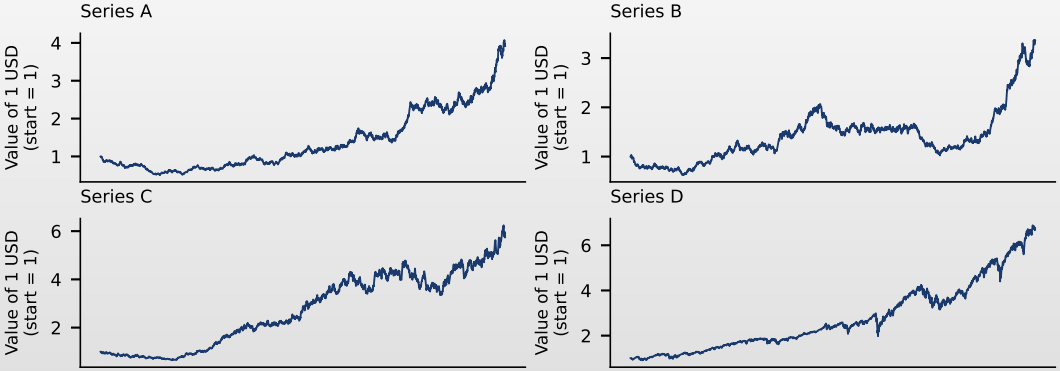
- **RW1**: Random Walk 1 — independent, identically distributed increments.
- **RW2**: Random Walk 2 — independent increments, possibly with different distributions.
- **RW3**: Random Walk 3 — uncorrelated increments, possibly dependent.

- Define the log return and the information available at time t.
- Distinguish RW1, RW2 and RW3 using their assumptions on returns.
- Explain why rejection also depends on the risk model.

A Game: Which Series Is the Real Market? (1/2)

S&P 500: Standard & Poor's 500.

One of these is the S&P 500 (2010-2026); three are simulated random walks with the same drift and volatility



MFM_ch2_random_walk_game

A Game: Which Series Is the Real Market? (2/2)

- Which panel is the S&P 500? Vote: A, B, C or D?

MFM_ch2_random_walk_game

A Game: Which Series Is the Real Market? (2/2)

- Which panel is the S&P 500? Vote: A, B, C or D?
- Answer: **D**.
- Roberts (1959): random walks produce "patterns" by chance
- We need statistical tests, not visual inspection
- Bachelier (1900) and Samuelson (1965): properly anticipated prices fluctuate randomly, so unpredictability is the mark of a well-functioning market

MFM_ch2_random_walk_game

Efficiency and the Information Set

- Fama (1970): a market is **efficient** (Efficient Market Hypothesis, EMH) if prices "fully reflect" all available information $\mathcal{I}_t$; portrait and surveys: Appendix; precursors: Appendix
  - formally: no strategy using $\mathcal{I}_t$ earns **expected** returns above the fair (risk-adjusted) return

Form	**Information set $\mathcal{I}_t$**	**Typical test (this chapter)**
Weak	Past prices and returns	Autocorrelation, runs, variance ratios, long memory
Semi-strong	All public information (news, reports)	Event studies
Strong	Public and private (insider) information	Returns of insiders and fund managers

- The strong form implies the semi-strong form, which implies the weak form
- Risk-adjusted returns with factor models follow in Chapter 3

The Joint-Hypothesis Problem

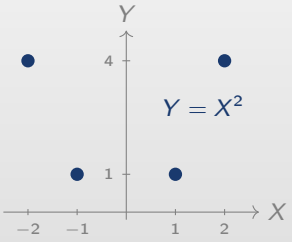
CAPM: Capital Asset Pricing Model.

- "Fair return" requires a model of expected returns (e.g. CAPM, Chapter 3)
- Every test of efficiency is a **joint** test of efficiency and of the asset-pricing model Fama (1991)
  - A rejection may mean an inefficient market **or** a wrong model of risk
  - Predictable **expected** returns are allowed if they compensate time-varying risk
- At daily horizons expected returns are tiny, so the choice of model matters little: weak-form tests are the cleanest
- Efficiency can never be "proved", only not rejected by a given test on given data

Correlation and dependence in returns

- Weak-form tests use only past prices. To say what it means for the past to contain no information, we need the difference between correlation and dependence.
- Lag-k autocorrelation is the correlation between a return and the return k periods earlier.
- Zero correlation excludes linear association, but may leave dependence in squared returns.

- Toy example: $X = -2, -1, 1, 2$, equally likely; $Y = X^2$.
- The mean of X is zero; the mean of XY is $(-8 - 1 + 1 + 8)/4 = 0$. Covariance, hence correlation, is zero.
- Yet knowing X determines Y exactly: they are dependent.



RW1, RW2, RW3: the restrictions on returns

- **i.i.d.**: independent and identically distributed.
- **GARCH**: Generalized AutoRegressive Conditional Heteroskedasticity.

$$r_t = \ln P_t - \ln P_{t-1} = \mu + \varepsilon_t$$

- μ: a constant mean; ε_t: the zero-mean innovation. The random walk concerns log prices; the restrictions concern their increments, the log returns.

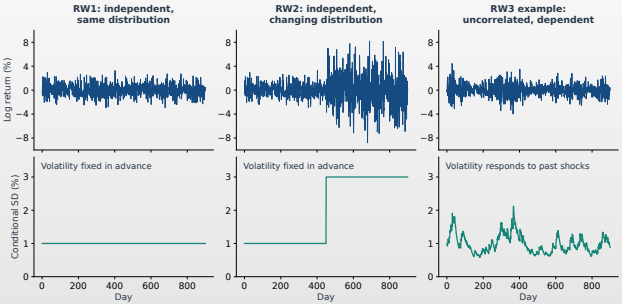
	Requirement	Illustrative return model
RW1	Independent, identical distribution	$r_t = \mu + 0.01 z_t$, z_t i.i.d. $N(0,1)$
RW2	Independent, distribution may vary	$r_t = \mu + \sigma_t z_t$, $\sigma_t = 1\%$ then 3%, fixed in advance
RW3	Uncorrelated, possibly dependent	GARCH: conditional variance responds to past squared shocks

How RW1, RW2 and RW3 are related

- RW1 implies RW2: i.i.d. returns are in particular independent.
- RW2 implies RW3: independent returns are in particular uncorrelated, provided their variance exists.
- A GARCH model keeps zero correlation but violates RW2: the size of today's move depends on past shocks.
- RW3 does not make tomorrow's mean unpredictable: a nonlinear function of the past can forecast it.

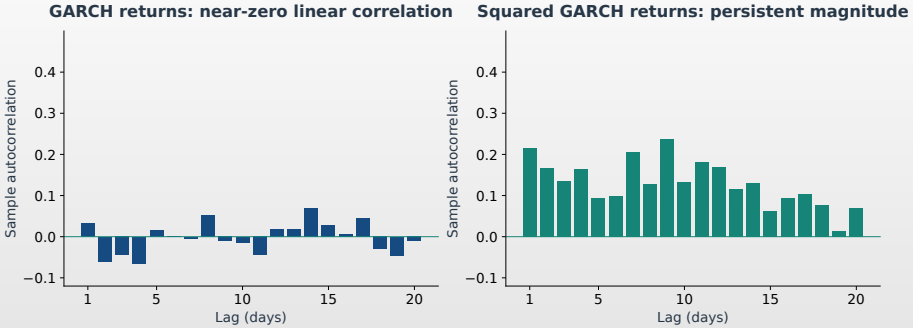
Three models for returns and volatility

SD: Standard Deviation.



- Three simulated series with mean zero: returns on top, the conditional standard deviation below; RW2 volatility is fixed in advance, GARCH volatility responds to shocks.
- Volatility clustering: large moves tend to be followed by large moves.
- The RW2 series also has a turbulent period although its returns are independent: the chart alone cannot tell a regime change fixed in advance from volatility clustering.

GARCH: returns versus squared returns



- Left: the autocorrelations of simulated GARCH returns are close to zero at all lags.
- Right: the autocorrelations of squared returns are positive up to lag 20: a day with a large move is followed by other days with large moves.
- These are illustrative sample estimates, not hypothesis tests.

Check: returns and information

- Can returns remain independent if their standard deviation moves from 1% to 3% on a date fixed in advance?
- Does zero autocorrelation establish independence?
- Why does rejecting an efficiency test not, on its own, prove that the market is inefficient?

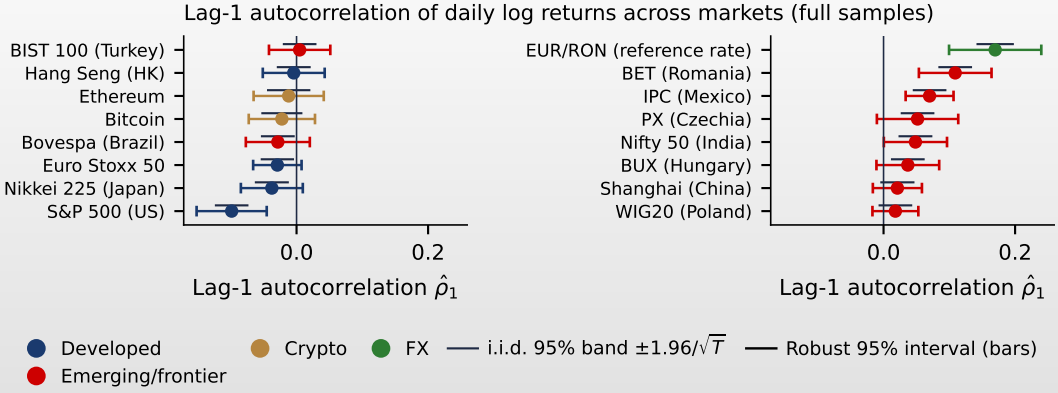
Answers: returns and information

- Yes. The series satisfies RW2 but violates RW1, because the distribution is no longer identical.
- No. In the GARCH example, return autocorrelation is zero, but squared returns are dependent.
- Every efficiency test also tests the model of expected returns.
- A rejection can come from a wrong risk model (the joint-hypothesis problem).

2. Autocorrelation and variance ratios

- Read the sign, uncertainty and economic size of an autocorrelation.
- Explain why volatility clustering changes inference.
- Read a variance-ratio profile without inferring a trading profit.

Lag-1 Autocorrelation Across Markets (1/3)



MFM_ch2_autocorrelation

Lag-1 Autocorrelation Across Markets (2/3)

SE: Standard Error.

- Each dot is the lag-1 autocorrelation $\hat{\rho}_1$ of one market's daily returns.
- The bar around the dot is the robust 95% confidence interval.
- The thin segments show the 95% band under the i.i.d. assumption, $\pm 1.96/\sqrt{T}$.
- Markets: BIST, WIG20, BUX, PX, IPC and BET are the stock-exchange indices of Istanbul, Warsaw, Budapest, Prague, Mexico and Bucharest; FX is foreign exchange (the EUR/RON rate).
- The robust interval allows for volatility clustering, so it is wider than the i.i.d. band.
- For the S&P 500, the robust standard error is 0.027 and the i.i.d. one is 0.012.

Lag-1 Autocorrelation Across Markets (3/3)

- BET: positive and significant autocorrelation, 0.109 (robust standard error 0.028).
- IPC (0.070) and EUR/RON (0.170): positive and significant autocorrelation.
- S&P 500: negative and significant autocorrelation, -0.099.
- Bitcoin (-0.022) and Ethereum (-0.012): autocorrelations not significant.
- For the BET, positive autocorrelation can come from thin trading in some stocks.
- For EUR/RON, it can come from the managed float of the exchange rate.
- For the S&P 500, the reversal can come from short-term overreaction and liquidity provision.
- The effect is small: yesterday's return explains $\hat{\rho}_1^2$ of the daily variance, 0.97% for the S&P 500 and 1.19% for the BET.

MFM_ch2_autocorrelation

The Variance Ratio: Idea

VR: Variance Ratio.

- Uncorrelated increments with constant variance: variance grows linearly with the horizon, $\text{Var}(r_t + \cdots + r_{t-q+1}) = q\,\text{Var}(r_t)$
  - q: the horizon in days (e.g. $q = 5$ is one trading week)
  - With changing variances (RW2, RW3) the variance of the sum is the sum of the variances; the sample $\widehat{\text{VR}}(q)$ still tends to 1 under the conditions of Lo & MacKinlay (1988)
- **Variance ratio**: $\text{VR}(q) = \dfrac{\text{Var}(r_t + \cdots + r_{t-q+1})}{q\,\text{Var}(r_t)}$
- For a covariance-stationary series, a weighted sum of autocorrelations: $\text{VR}(q) = 1 + 2\sum_{k=1}^{q-1}\left(1 - \frac{k}{q}\right)\rho_k$
- $\text{VR}(q) > 1$: persistence (positive autocorrelation); $\text{VR}(q) < 1$: reversal (mean reversion)
- For $q = 2$, the formula gives $\text{VR}(2) = 1 + \rho_1$
- For the S&P 500, $1 + \hat{\rho}_1 = 1 - 0.099 = 0.901$, the same value as $\widehat{\text{VR}}(2)$ estimated directly

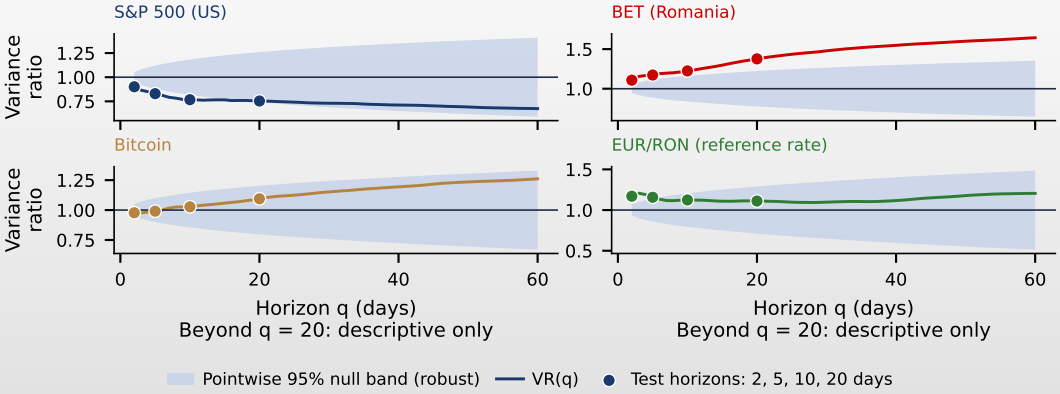
The robust variance-ratio band

- Without autocorrelation, the benchmark is $\mathrm{VR}(q) = 1$, even when volatility changes over time (the RW3 null).
- The robust 95% band is centred at 1, the value under the null: if $\widehat{\mathrm{VR}}(q)$ falls outside it, we reject the null at 5%.
- The band is not a confidence interval: a confidence interval would be centred at the estimate $\widehat{\mathrm{VR}}(q)$, not at 1.
- Its width depends on $\hat{\theta}(q)$, the estimated variance of $\sqrt{T}\,(\widehat{\mathrm{VR}}(q) - 1)$.
- When large moves follow large moves (volatility clustering), $\hat{\theta}(q)$ grows and the band widens.

Which variance-ratio horizons do we test?

- Lo & MacKinlay (1988) use $q = 2, 4, 8, 16$ on weekly data.
- Here the four test horizons are $q = 2, 5, 10, 20$ days: two days, about one week, two weeks and one trading month for equities.
- The robust null band has level 5% at each horizon taken separately, not at all horizons jointly.
- The four horizons are the analyst's choice; other horizons can give another result.
- The Chow–Denning test combines the chosen horizons into one joint rejection decision Chow & Denning (1993).
- The Chow–Denning test rejects RW3 for 5 of the 16 series, including the BET and the S&P 500; not for Bitcoin ($p = 0.835$).

Variance-Ratio Profiles (1/3)



MFM_ch2_variance_ratio

Variance-Ratio Profiles (2/3)

- Each curve shows the estimated variance ratio $\widehat{\mathrm{VR}}(q)$ for horizons q from 2 to 60 days.
- The shaded area is the robust 95% band around 1, the value without autocorrelation.
- S&P 500: $\widehat{\mathrm{VR}}(q)$ is below 1 (reversal) and leaves the band at 2, 5 and 10 days; at 20 days it is back inside
  - $\widehat{\mathrm{VR}}(2) = 0.901$, $\widehat{\mathrm{VR}}(5) = 0.830$, $\widehat{\mathrm{VR}}(10) = 0.767$, $\widehat{\mathrm{VR}}(20) = 0.753$
- BET: $\widehat{\mathrm{VR}}(q)$ is above 1 (persistence), rises with q and leaves the band at all four tested horizons
  - $\widehat{\mathrm{VR}}(2) = 1.108$, $\widehat{\mathrm{VR}}(5) = 1.173$, $\widehat{\mathrm{VR}}(10) = 1.225$, $\widehat{\mathrm{VR}}(20) = 1.377$

Variance-Ratio Profiles (3/3)

- The horizon $q = 60$, about three trading months, is only descriptive: it is not one of the four tested horizons
  - S&P 500: $\widehat{\mathrm{VR}}(60) = 0.674$; BET: $\widehat{\mathrm{VR}}(60) = 1.644$
- Bitcoin stays inside the band at all horizons: the null is not rejected.
- EUR/RON has $\widehat{\mathrm{VR}}(q)$ above 1 at short horizons, a persistence consistent with the managed float of the exchange rate.
- The band widens as q grows.
- A curve above or below 1 but inside the band is not statistical evidence of predictability.

MFM_ch2_variance_ratio

Why the robust band is wider: a numerical check

- For $q = 2$, the classical statistic is $z = \sqrt{T}\,(\widehat{\mathrm{VR}}(2) - 1)$; it assumes i.i.d. returns (RW1).
- The robust statistic z^* divides the same numerator by $\sqrt{\hat{\theta}(2)}$.
- Without volatility clustering, $\hat{\theta}(2) \approx 1$ and $z^* \approx z$; with volatility clustering, $\hat{\theta}(2) > 1$ and $|z^*| < |z|$.
- Toy example: $T = 400$ days and $\widehat{\mathrm{VR}}(2) = 0.90$ give $z = 20 \times (-0.10) = -2.0$: the null is rejected at 5%.
- If volatility clustering makes $\hat{\theta}(2) = 4$, then $z^* = -2.0/\sqrt{4} = -1.0$ and the null is no longer rejected.
- The robust band around 1 is then $\sqrt{4} = 2$ times wider.
- For the S&P 500, $z = -8.12$ and $z^* = -3.64$: the robust statistic is about 2.2 times smaller, but the null is still rejected.

Check: autocorrelation and variance ratios

- Why can an i.i.d. band be misleading under volatility clustering?
- Does positive autocorrelation imply profit after costs?
- What does a variance ratio below one suggest?
- Transfer: a test statistic has magnitude 2.4. If the VR estimate stays the same but its standard error doubles, what does the magnitude become?
- Is the null still rejected at 5% (two-sided)?

Answers: autocorrelation and variance ratios

- The i.i.d. standard error ignores dependence in squared returns.
- No: effect size, risk, implementation and costs still matter.
- A variance ratio below 1 suggests reversal over the chosen horizon.
- To call the deviation significant, we need a robust standard error.
- The statistic halves to 1.2.
- No: 1.2 does not cross 1.96; the estimated effect is unchanged, its evidence is weaker.

3. Out-of-sample predictability

- Keep estimation data separate from forecast evaluation.
- Compare squared errors with those of a stated benchmark.
- Separate predictive accuracy from profit after risk and costs.

What information enters a forecast?

- The predictor: dp, the log dividend–price ratio.
- Monthly dividends are interpolated from quarterly totals, so part of the information was not available exactly at the forecast date.
- The exercise is therefore pseudo-out-of-sample, not fully real-time.

How we compare out-of-sample forecasts

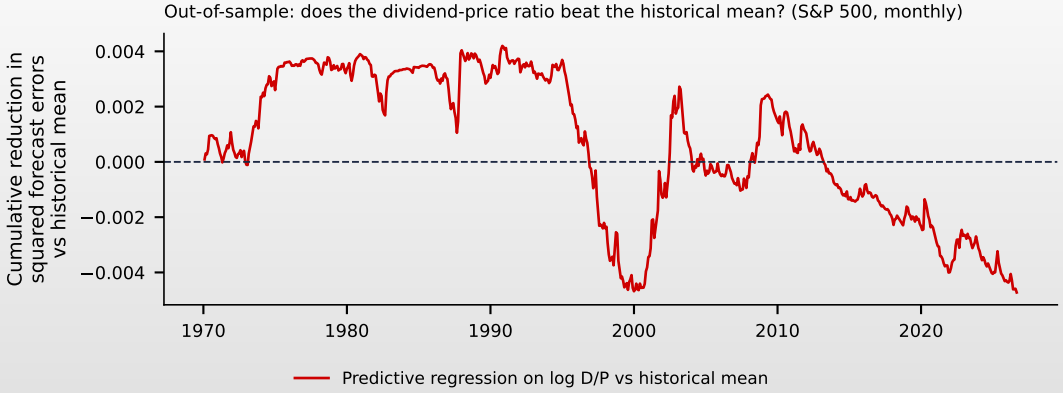
OS: Out-of-Sample.

- The model: each month we regress next month's return on this month's dp, $r_{t+1} = \alpha + \beta\, dp_t + u_{t+1}$, and use the estimate to forecast.
- An expanding window re-estimates the model each month using only observations before the month forecast.
- $R^2_{OS} = 1 - \sum_t (r_t - \hat{r}_t)^2 / \sum_t (r_t - \bar{r}_t)^2$.
- $\hat{r}_t$: the model forecast; $\bar{r}_t$: the historical-mean forecast from the same available data.
- A positive score means smaller squared errors than the benchmark on the evaluation sample.
- A positive score alone does not establish statistical significance or profit.
- A regression estimated on the full sample can have a positive in-sample R^2 and still a negative R^2_{OS}: in sample the model also fits noise.

Reading the cumulative squared-error chart

- Each month t from 1970 adds $(r_t - \bar{r}_t)^2 - (r_t - \hat{r}_t)^2$: the squared error of the historical mean minus that of the *dp* forecast.
- The curve is the running sum of these terms, in squared fractional log-return units.
- A rising segment marks months in which the *dp* forecast beat the historical mean; a falling segment, months in which it did worse.
- Example: errors of 2, 1 and 3 percentage points for the historical mean and 1, 2 and 3 for the *dp* forecast give the terms 3, -3 and 0; the curve rises to 3, falls back to 0, and $R^2_{OS} = 0$.
- The last value of the curve has the sign of R^2_{OS}: above zero, $R^2_{OS} > 0$; below zero, $R^2_{OS} < 0$.

The Dividend–Price Ratio: Out-of-Sample Forecasts



MFM_ch2_predictability

The Dividend–Price Ratio: Results for the S&P 500 (1/2)

- The S&P 500 curve ends below zero: over 680 monthly forecasts, $R^2_{OS} = -0.51\%$.
- The dividend–price ratio therefore has slightly larger squared errors than the historical mean.
- The Clark–West test Clark & West (2007) compares the squared errors of the two forecasts and corrects for the noise in estimating β; a large positive t favours the predictor.
- On the S&P 500, the Clark–West test gives $t = 0.42$ (one-sided $p = 0.34$): no significant out-of-sample predictability.
- The cumulative advantage of the *dp* forecast peaks in 1990 and is lost in the 1990s boom, when *dp* fell to record lows.
- After 2010, the *dp* forecast loses ground against the historical mean again.

The Dividend–Price Ratio: Results for the S&P 500 (2/2)

- The Campbell & Thompson (2008) restrictions (sign of β imposed, non-negative forecast) bind in 5.7% of months and leave R^2_{OS} practically unchanged, at -0.51%.
- Hypothetically, a predictor with $R^2_{OS} = 0.5\%$ would raise the expected return of a mean–variance investor who times the market by about 20% (formula: Appendix).
- The calculation uses a monthly market Sharpe ratio of 0.160: the mean excess return divided by its standard deviation.
- Here R^2_{OS} is negative, so this gain is absent.

Check: out-of-sample predictability

- What does a negative out-of-sample score say about the benchmark?
- Why can excellent in-sample fit fail on new data?
- Transfer: can a profitable strategy have a negative R^2_{OS}?
- Against which benchmark is a strategy's profit measured?
- The advantage of the *dp* forecast is lost in the 1990s: what does this say about the stability of the link between *dp* and returns?

Answers: out-of-sample predictability

- The model has larger squared errors than the benchmark on that sample.
- The model may fit noise or exploit information unavailable at the forecast origin.
- Yes. Profit can come from constant market exposure or from a few correct sign forecasts.
- Profit is measured against zero or the market, and on its own does not establish risk-adjusted profit after costs.
- The link is not stable: a predictor more accurate than the historical mean until 1990 loses all its advantage when *dp* falls to record lows.

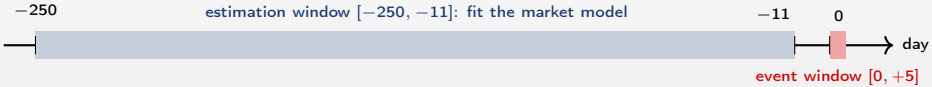
4. Event studies

CAR: Cumulative Abnormal Return.

- Separate observed, expected and abnormal returns.
- Interpret CAR over a stated event window.
- Distinguish a measured reaction from proof of its cause.

The Event-Study Method (1/2)

AR (event study): Abnormal Return.

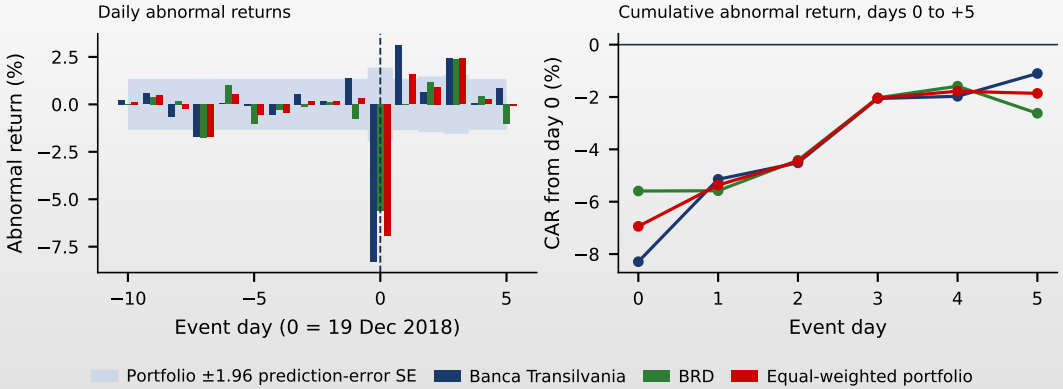


- Fama, Fisher, Jensen & Roll (1969): how fast do prices absorb public news (stock splits)?
- Market model: $r_{i,t} = a_i + b_i r_{m,t} + \varepsilon_{i,t}$, $r_{m,t}$ the market return; abnormal return $AR_{i,t} = r_{i,t} - \hat{a}_i - \hat{b}_i r_{m,t}$; cumulative abnormal return $CAR_i[\tau_1, \tau_2] = \sum_{t=\tau_1}^{\tau_2} AR_{i,t}$
  - a_i, b_i: intercept and market beta of firm i, fitted in the estimation window; $\hat{\sigma}_\varepsilon$: the residual standard deviation there; $[\tau_1, \tau_2]$: the days summed

The Event-Study Method (2/2)

- First-pass test: $t = CAR_i/(\hat{\sigma}_\varepsilon\sqrt{L})$, L the event-window length MacKinlay (1997)
  - AR: the return beyond what the market move predicts; the SD of a sum of L independent residuals is $\hat{\sigma}_\varepsilon\sqrt{L}$
- For unanticipated news, efficiency predicts a jump on day 0 and no predictable drift or reversal afterwards
- Daily data can only show absorption within a day; speed at the microsecond scale needs intraday data (Chapter 10; literature: Appendix)
- Application: the price response of Ethereum to network development milestones (Quantinar: Market Responses to ETH Development Milestones)

Worked Example: the Romanian Bank Tax of December 2018



MFM_ch2_seminar

The Bank Tax of December 2018: Results (1/2)

- The government announced the tax on bank assets on the evening of 18 December 2018; event day 0 is 19 December.
- We study Banca Transilvania (TLV), BRD and their equal-weighted portfolio; the market return is that of the BET index.
- For TLV, the estimated market beta is 1.18 and the residual standard deviation is 0.83% per day.
- On day 0, the abnormal return of TLV is $AR_0 = -8.3\%$, with $t = -10.0$ (first-pass formula).
- That day the log return of the BET was -11.9% (a fall of 11.2%).
- The prediction error also includes the uncertainty in $\hat{b}$ times this extreme market return; with it, t becomes -6.85.

The Bank Tax of December 2018: Results (2/2)

- Over days $+1, \ldots, +5$, the cumulative abnormal return of TLV is $+7.2\%$ ($t = 3.82$).
- Over days $0, \ldots, +5$, the cumulative abnormal return is only -1.1% ($t = -0.54$): the day-0 fall was almost recovered within a week.
- First possible explanation: an **overreaction** on day 0, against semi-strong efficiency.
- Second possible explanation: later news about the final form of the ordinance.
- A single event does not let us choose between the two explanations.

MFM_ch2_seminar

Check: event studies

- Why must you specify the event window?
- Could another announcement explain the same movement?

Answers: event studies

- CAR adds abnormal returns over the days of the window only.
- A window chosen after seeing the data can create or hide an effect.
- Yes: overlapping news is an alternative explanation, so the reaction alone is not causal proof.

5. Multiple testing and anomalies

- Account for the number of hypotheses tested.
- State what the evidence establishes and what it leaves open.
- Choose a check that could overturn your conclusion.

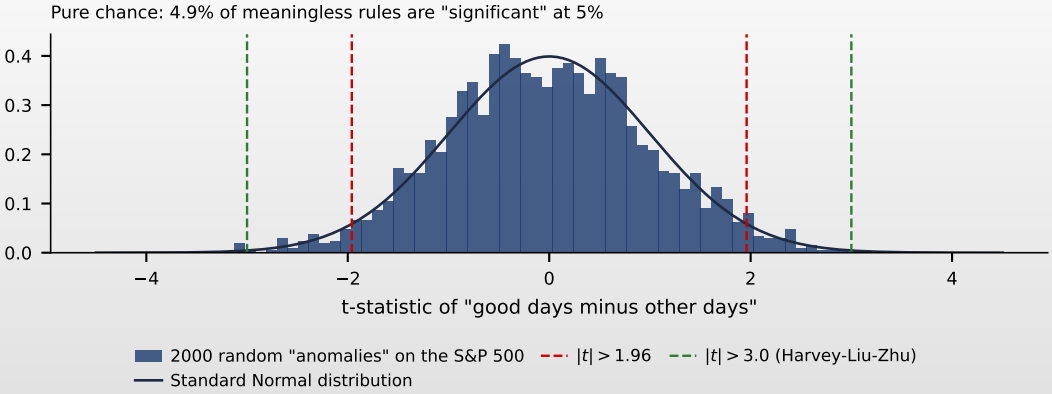
Calendar effects and data snooping

- A calendar effect is a difference in mean returns associated with a day or month.
- Turn of the month (TOM): the last trading day of a month and the first three of the next month.
- A January effect compares January returns with the other months.
- Data snooping means selecting a pattern after searching the same data for many alternatives.

How Many "Anomalies" Does Pure Chance Produce? (1/3)

- Experiment: 2 000 trading rules with no economic content, on daily S&P 500 returns.
- Each rule picks 20% of the days at random and calls them "good days".
- For each rule, a t test compares the mean return on "good days" with the mean on the other days.
- Every null hypothesis is true, so every rejection is false.

How Many "Anomalies" Does Pure Chance Produce? (2/3)



MFM_ch2_multiple_testing

How Many "Anomalies" Does Pure Chance Produce? (3/3)

- Result: 4.9% of the rules have $|t| > 1.96$, i.e. "significant" at 5%, close to what the 5% level implies.
- The most "spectacular" rule has $|t| = 3.19$: reported alone, it would pass even the strict bar $|t| > 3.0$, although it is pure chance.
- Harvey, Liu & Zhu (2016) propose the bar $|t| > 3.0$ for accepting a new factor; here 0.25% of the random rules pass it, i.e. 5 rules out of 2 000.
- How many rules were tried (the family of tests) matters as much as the bar chosen.
- Sullivan, Timmermann & White (1999) show how to correct for data snooping in trading rules.

MFM_ch2_multiple_testing

Reading the calendar-effects table

HAC: Heteroskedasticity and Autocorrelation Consistent.

- We apply the same logic to 15 real tests: five markets $\times$ three effects (weekday, January, TOM).
- January extra: the mean January return minus the mean of the other months, in percent per month.
- TOM extra: the mean return on TOM days minus the mean of the other days, in basis points per day.

January and Turn-of-the-Month Effects

Market	January extra (%/month)	t_{HAC}	TOM extra (bp/day)	t_{HAC}	p (TOM)
S&P 500 (US)	−0.66	−0.80	3.2	0.93	0.3539
BET (Romania)	3.48	1.49	17.3	3.62	0.0003
Euro Stoxx 50	−0.30	−0.29	5.1	1.21	0.2266
WIG20 (Poland)	0.53	0.40	8.2	1.73	0.0841
BUX (Hungary)	2.66	2.25	15.1	3.63	0.0003

- TOM: the extra return is +17.3 bp per day for the BET and +15.1 bp for the BUX, both with $|t| > 3.5$
- For the S&P 500, the TOM extra return is only +3.2 bp per day and not significant ($t = 0.93$)
- For the BET, January exceeds the other months by 3.48 percentage points, but the estimate is imprecise ($t = 1.49$)
- The BUX, with $t = 2.25$, is the only market with a January-effect p-value below 0.05
- The third effect, the day of the week, is not significant in any market: the smallest p-value is 0.2262 (WIG20)
- We ran 15 tests (5 markets $\times$ 3 effects): which survive a correction for multiple testing?

Multiple Testing: Bonferroni and Holm

- **FWER** (family-wise error rate): the probability of at least one false rejection; with M independent tests at level α, $\text{FWER} = 1 - (1-\alpha)^M$
  - $M = 15$, $\alpha = 5\%$: $\text{FWER} = 53.7\%$
- **Bonferroni**: reject only if $p_i < \alpha/M$
- Bonferroni guarantees FWER $\leq \alpha$ for any dependence between tests (bootstrap methods that use this dependence: Appendix)
- **Holm**: sort $p_{(1)} \leq \cdots \leq p_{(M)}$; reject while $p_{(j)} < \alpha/(M-j+1)$, stop at the first failure
  - Same guarantee, never less powerful than Bonferroni
- For our 15 calendar tests, the Bonferroni threshold is $0.05/15 = 0.0033$

Holm on the 15 calendar p-values (1/2)

- All p-values are two-sided; for January and TOM they come from the HAC t statistic compared with the standard Normal distribution.
- For the weekdays they come from a Wald test that the five daily means are equal (four restrictions, χ^2_4).

Market	Weekday, equal means	January	TOM
S&P 500 (US)	0.8263	0.4257	0.3539
BET (Romania)	0.3052	0.1368	0.0003
Euro Stoxx 50	0.6319	0.7738	0.2266
WIG20 (Poland)	0.2262	0.6865	0.0841
BUX (Hungary)	0.9327	0.0245	0.0003

Holm on the 15 calendar p-values (2/2)

- The two smallest p-values are TOM for the BET and TOM for the BUX, both 0.0003.
- They are below the Holm thresholds $0.05/15 = 0.0033$ and $0.05/14 = 0.0036$, so both nulls are rejected.
- The third p-value, January for the BUX (0.0245), exceeds the threshold $0.05/13 = 0.0038$, so Holm stops here.
- Without correction, January for the BUX would have been the third rejection, since $0.0245 < 0.05$.

Check: multiple testing and anomalies

- Why does testing many rules change the meaning of a small p-value?
- Does non-rejection prove efficiency?
- Which new sample or robustness check would you request?

Answers: multiple testing and anomalies

- Searching many rules creates more opportunities for false rejection.
- No. A test with low power may fail to reject even if the market is inefficient.
- The conclusion concerns only the alternative tested, not every form of predictability.
- For example, the same rule on post-publication data or on another market: a data-snooping artefact disappears out of sample; a real effect usually weakens after publication but does not vanish McLean & Pontiff (2016).

Main results of the chapter

- S&P 500: reversal at 2–10 days ($\widehat{\mathrm{VR}}(2) = 0.901$); BET: persistence at all four tested horizons.
- The dividend–price ratio does not beat the historical mean out of sample: $R^2_{OS} = -0.51\%$.
- Bank tax: TLV's abnormal return was -8.3% on day 0 and was almost recovered within five days.
- After the Holm correction only the TOM effects for the BET and the BUX remain significant.

Topics for further study

- ⊡ Data and conventions
- ⊡ Fair games, pricing and rational bubbles
- ⊡ Test construction and alternatives
- ⊡ Predictive regressions and bias
- ⊡ Long memory and changing efficiency
- ⊡ Event-study inference
- ⊡ Anomalies and behavioural explanations
- ⊡ Machine learning and research design
- ⊡ AI for scientific discovery
- ⊡ Synthesis and technical reference

Data Used in This Chapter

BNR: Banca Națională a României — the National Bank of Romania.

- 16 daily series up to 18 September 2026; sample sizes: ▸ Appendix
  - 13 equity indices from 2000: 4 developed (S&P 500, Euro Stoxx 50, Nikkei 225, Hang Seng), 9 emerging or frontier (BET, WIG20, BUX, PX, BIST 100, Bovespa, IPC, Nifty 50, Shanghai)
  - Bitcoin from 2014 in the full-sample and rolling tests; the yearly crypto analysis starts in 2011 to cover the market's first years
  - Ethereum from 2016 and EUR/RON, the BNR reference rate, from July 2005
- Daily log returns $r_t = \ln P_t - \ln P_{t-1}$, each series on its own calendar
  - P_t: the closing price (index level, exchange rate) on day t
- Equity indices use trading days only; a close equal to the previous close is removed, since it is usually a holiday filled with the previous price
  - the rule also removes the rare genuine zero-return day (S&P 500: 3 of about 6 700 days); the test statistics move by less than 0.05
- Crypto-assets trade 7 days a week: 365 observations per year
- Analyses that combine several series (e.g. event studies) align prices on common trading days before computing returns

Market abbreviations

- **US**: United States
- **HK**: Hong Kong
- **BIST**: Borsa İstanbul — the Istanbul Stock Exchange (BIST 100 index).
- **WIG20**: Warszawski Indeks Giełdowy 20 — the Warsaw Stock Exchange index of the 20 largest companies.
- **BUX**: Budapesti Értéktőzsde index — the Budapest Stock Exchange index.
- **PX**: Prague Stock Exchange index (index Burzy cenných papírů Praha).
- **IPC**: Índice de Precios y Cotizaciones — the Mexican Stock Exchange index.
- **BET**: Bucharest Exchange Trading (index).
- **FX**: Foreign eXchange.

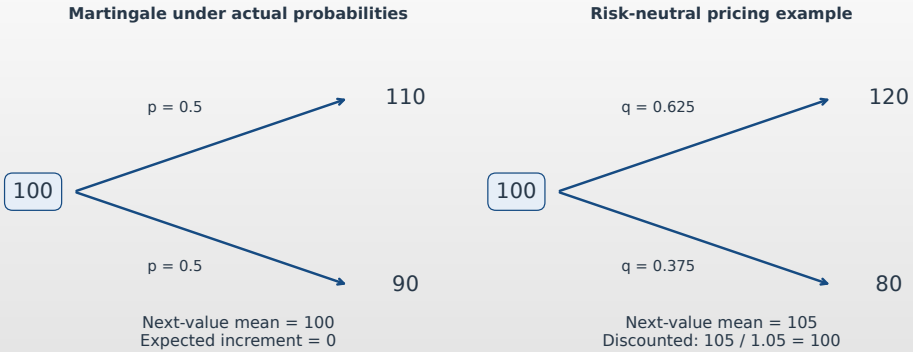
Pure arbitrage and relative-value trading

- Pure arbitrage is a strategy with no positive initial outlay, no possible loss and a positive gain in at least one state.
- Relative-value trading is risky when price convergence is uncertain.
- Example: a price gap of 20 points may widen to 30 before it closes; the position then needs extra collateral (margin) and may be liquidated.

Discounting and the risk-neutral measure

- A risk-neutral measure uses pricing probabilities under which discounted traded prices, including dividends, are martingales in an arbitrage-free model.
- Example: a stock without dividends costs 100 today; in one year it is worth 120 or 80; the risk-free rate is 5%.
  - The pricing probability $q = 0.625$ gives $q\,120 + (1-q)80 = 105$, and $105/1.05 = 100$, today's price.
- This q need not equal the actual probability of an increase: pricing probabilities are not forecasts of frequencies.

Probabilities and conditional means



Left: the actual conditional mean stays at 100. Right: pricing probabilities give 105, whose discounted value is 100.

Fair games and martingales

- **Fair game** for excess returns: $\mathbb{E}[r_{t+1} - \mu_t \mid \mathcal{I}_t] = 0$
  - $\mathcal{I}_t$: the information available at the end of day t; $\mathbb{E}[\cdot \mid \mathcal{I}_t]$: the expectation given that information
  - μ_t: the required (equilibrium) return, the compensation for risk; in words: the return beyond μ_t cannot be forecast from $\mathcal{I}_t$
- **Martingale**: a process whose expected next value, given the past, equals its current value
  - With a constant required simple return m, the total-return price discounted at m, $P_t/(1+m)^t$, is a martingale Campbell, Lo & MacKinlay (1997); discounting at the risk-free rate gives a martingale only under the risk-neutral measure
  - P_t itself is one only if $m = 0$: a zero expected **simple** return, not a zero expected log return

The Jensen term: why averages differ

- The logarithm is concave: averaging log outcomes gives less than taking the log of their average, unless outcomes are constant.
- Equally likely simple returns of $+10\%$ and -10% have mean zero.
- The mean log return is $\frac{1}{2}\ln(1.1) + \frac{1}{2}\ln(0.9) = \frac{1}{2}\ln(0.99) \approx -0.00503$.
- Along a path with one rise and one fall, wealth 100 becomes $100(1.1)(0.9) = 99$: a zero arithmetic mean does not preserve compounded wealth.

Prices, log prices and information

- "$\mathbb{E}[P_{t+1} \mid \mathcal{I}_t] = P_t$" and "$\mathbb{E}[\ln(P_{t+1}/P_t) \mid \mathcal{I}_t] = \mu$" are different restrictions
  - They differ by discounting and by the Jensen term: $\mathbb{E}[\ln(1+R)] \approx \mathbb{E}[R] - \mathrm{Var}(R)/2$ for small returns ($\mathbb{E}[R]^2$ neglected)
  - $R = P_{t+1}/P_t - 1$: the simple return; $\ln(1+R)$: the log return; the mean log return is lower than the mean simple return by about half the variance
- Consequence of the fair game: $\mathrm{corr}(\varepsilon_{t+1}, g(\mathcal{I}_t)) = 0$ for any function g of the past, $\varepsilon_{t+1} = r_{t+1} - \mu_t$
  - ε_{t+1}: the unexpected part of the return; in words, no past information is correlated with tomorrow's surprise
  - In particular, all autocorrelations $\rho_k = \mathrm{corr}(r_t, r_{t-k})$, $k \geq 1$, are zero when $\mu_t = \mu$; k: the lag in days

Rational bubble: definition

- Fundamental value: the present value of the cash flows expected from the asset.
  - For a share, these payments are dividends.
  - Discounting expresses future payments in today's money.
- A **rational bubble** is a positive price component above fundamental value.
  - It is sustained by the expectation of resale.
  - $P_t = P_t^* + B_t$: P_t is the price, P_t^* fundamental value, and B_t the bubble component.
- Rational expectations are consistent with the model, including crash risk.
  - They do not imply that a gain is certain.

Fundamental value, rational bubbles and fads

- With a constant required simple return m, the fair game is $\mathbb{E}[P_{t+1} + D_{t+1} \mid \mathcal{I}_t] = (1+m)\,P_t$; D_{t+1}: the cash flow (dividend) paid in $t+1$
  - Fundamental value $P_t^* = \mathbb{E}\big[\sum_{j\geq 1} D_{t+j}/(1+m)^j \mid \mathcal{I}_t\big]$; pricing error $B_t = P_t - P_t^*$
  - The fundamental value P_t^* satisfies the same fair-game equation as the price
  - Subtracting the two equations, the gap B_t satisfies $\mathbb{E}[B_{t+1} \mid \mathcal{I}_t] = (1+m)\,B_t$
- A fair game ties the price to the fundamental value only up to a **rational bubble**
  - A gap $B_t \neq 0$ is admissible only if it is expected to grow at the required rate m (Chapter 17)
  - A mean-reverting gap (fad) makes ε_{t+1} predictable and violates the fair game, although short-horizon tests have little power against it Summers (1986)
- The restriction is on the conditional mean only: the variance is not restricted, and volatility may be predictable (Chapter 1)

Rational bubble: a numerical example

- A share pays 10 RON at the end of every year, forever.
  - The first dividend is paid one year from today.
  - The constant required return is $m = 10\%$; fundamental value is $P_t^* = 10/0.10 = 100$ RON.
- Today $P_t = 120$ RON, so $B_t = 20$ RON.
  - Without a crash, $B_{t+1} = 1.1 \times 20 = 22$ RON.
  - After next year's dividend, fundamental value remains 100 RON; the price is therefore 122 RON.
- The total simple return is $(122 + 10 - 120)/120 = 10\%$.

Rational bubble: what happens if it can burst?

- Initial bubble $B_t = 20$ RON and required return $m = 10\%$.
  - Survival probability $\pi = 0.8$; crash probability $1 - \pi = 0.2$.
  - A crash makes the bubble zero; investors know the probabilities.
- The growth condition requires an expected bubble of $(1 + m)B_t = 22$ RON.
  - Conditional on survival: $B_{t+1} = 22/0.8 = 27.5$ RON.
  - Check: $0.8 \times 27.5 + 0.2 \times 0 = 22$.
  - Conditional growth is $27.5/20 - 1 = 37.5\%$.
- Crash mechanism: Blanchard & Watson (1982), Bubbles, Rational Expectations and Financial Markets.
- The purchase price is 120 RON; the dividend at year-end is 10 RON.
- Survival: the ex-dividend price is 127.5 RON.
  - Simple return: $(127.5 + 10 - 120)/120 \approx 14.6\%$.
- Crash: the price falls to fundamental value, 100 RON.
  - Simple return: $(100 + 10 - 120)/120 \approx -8.3\%$.
- Expected return: $0.8 \times (127.5 + 10 - 120)/120 + 0.2 \times (100 + 10 - 120)/120 = 10\%$.
  - The expected return satisfies the fair-game condition; either realised return can differ from it.

Rational bubble and the no-bubble condition

- The fair-game condition restricts the expected return over the next period.
- The additional **no-bubble condition** is $\lim_{k\to\infty} \frac{\mathbb{E}[P_{t+k} \mid \mathcal{I}_t]}{(1+m)^k} = 0$.
  - k: the number of periods until resale.
  - The discounted expected resale price vanishes as the horizon grows.
  - Then the current price is supported entirely by cash flows.
- The numerical examples verify the one-period pricing recursion.
  - Existence in an economic equilibrium needs further assumptions: Tirole (1985), Asset Bubbles and Overlapping Generations.

The Grossman–Stiglitz Paradox

- Gathering information is costly (analysts, data, computing)
- If prices reflected **all** information, nobody would be paid for collecting it Grossman & Stiglitz (1980)
  - Then nobody would collect it, and prices could not reflect it
- Resolution: an **equilibrium degree of inefficiency**, just large enough to reward informed traders for their costs
- Markets with cheap information and low costs (large, liquid) should be closer to efficiency: a prediction for our comparison of developed, emerging and crypto markets
- Efficiency does not rule out bubbles either: it says they are hard to time (Chapter 17)

Moments, tails and asymptotics

- Which assumptions make the test valid, and which alternatives can it detect?
- The kth absolute moment is $\mathbb{E}[|r|^k]$
  - a finite fourth moment requires $\mathbb{E}[r^4] < \infty$, not merely finite observed returns.
- Heavy tails allow extreme values more often than under the Normal distribution
  - Student t with four degrees of freedom has finite variance but infinite fourth moment.
- An asymptotic approximation describes behaviour as sample size grows
  - it need not be accurate in a short sample.
- Mixing conditions require sufficiently distant blocks to become nearly independent at a suitable rate
  - adjacent daily returns need not be independent.

Autocorrelation and Its Robust Standard Error

- Sample autocorrelation: $\hat{\rho}_k = \sum_{t=k+1}^{T} e_t e_{t-k} / \sum_{t=1}^{T} e_t^2$, with $e_t = r_t - \bar{r}$
  - T: number of daily returns; $\bar{r}$: their mean; k: the lag; $\hat{\rho}_k > 0$: persistence, $\hat{\rho}_k < 0$: reversal
- Under RW1 (i.i.d.): $\mathrm{SE}(\hat{\rho}_k) \approx 1/\sqrt{T}$; 95% band $\pm 1.96/\sqrt{T}$
- Unpredictable mean with volatility clustering (finite fourth moments): $\mathrm{SE}(\hat{\rho}_k) \approx \sqrt{\hat{\tau}_k}$
  - $\hat{\tau}_k = \sum_t e_t^2 e_{t-k}^2 / \left(\sum_t e_t^2\right)^2$ (robust variance)
  - When large moves cluster, $\hat{\tau}_k > 1/T$: the i.i.d. band is too narrow and classical tests reject too often; every test in this topic uses a variance that allows for this
- S&P 500: i.i.d. SE $= 0.012$, robust SE $= 0.027$, more than twice as large

Joint tests and reference distributions

- A portmanteau test combines several lag correlations into one statistic instead of selecting the most striking lag.
- A chi-squared distribution with m degrees of freedom is the distribution of a sum of squares of m independent standard Normal variables.
- For ten tested lags the reference is approximately χ^2_{10} under the stated assumptions
  - a statistic of 25 exceeds the 5% critical value of χ^2_{10}, 18.31.
- A Wald test similarly tests coefficient restrictions jointly
  - five equal weekday means impose four restrictions, $\beta_1 = \beta_2$, $\beta_1 = \beta_3$, $\beta_1 = \beta_4$, $\beta_1 = \beta_5$, so the reference is χ^2_4, with 5% critical value 9.49.
- Five restrictions, all $\beta_d = 0$, would test a different null: every weekday mean equal to zero.

Portmanteau Tests: Ljung–Box and Its Robust Version

- Ljung & Box (1978): $Q(m) = T(T+2)\sum_{k=1}^{m} \hat{\rho}_k^2/(T-k) \sim \chi_m^2$ under i.i.d. returns
  - Tests jointly H_0: $\rho_1 = \cdots = \rho_m = 0$
  - m: the number of lags
  - χ_m^2: the chi-squared distribution with m degrees of freedom
  - A large Q rejects H_0
- Robust version Lobato, Nankervis & Savin (2001): $\tilde{Q}(m) = \sum_{k=1}^{m} \hat{\rho}_k^2/\hat{\tau}_k \to \chi_m^2$ under RW3 when the $\hat{\rho}_k$ are asymptotically uncorrelated (condition on the slide The Heteroskedasticity-Robust Statistic z^*)
- Three markets at $m = 10$ (5% critical value 18.31)
  - S&P 500: $Q = 96.3$ but $\tilde{Q} = 19.9$ ($p = 0.031$)
  - BET: $Q = 108.0$, $\tilde{Q} = 24.4$ ($p = 0.007$)
  - Bitcoin: $Q = 20.2$ ($p = 0.028$) but $\tilde{Q} = 11.9$ ($p = 0.294$)
- The robust test keeps the rejections for S&P 500 and BET, but not for Bitcoin

The Runs Test

- Nonparametric: uses only the **signs** of returns (zero returns dropped), so it is immune to heavy tails Wald & Wolfowitz (1940)
- A **run** is a maximal block of equal signs: ++|-|+++ has 3 runs
- Moments of the number of runs R under independence, with n_1 positive and n_2 negative signs, $n = n_1 + n_2$
  - $\mathbb{E}[R] = \dfrac{2n_1n_2}{n} + 1, \qquad \mathrm{Var}[R] = \dfrac{2n_1n_2(2n_1n_2 - n)}{n^2(n-1)}, \qquad z = \dfrac{R - \mathbb{E}[R]}{\sqrt{\mathrm{Var}[R]}} \approx N(0,1)$
  - R: the observed number of runs; $\mathbb{E}[R]$, $\mathrm{Var}[R]$: its mean and variance if the signs were independent; $|z| > 1.96$ rejects at 5%
- Too **few** runs ($z < 0$): signs persist
- Too **many** runs ($z > 0$): signs reverse
- The test assumes a constant probability of a positive day
- S&P 500: $z = 4.34$, more runs than under independence; signs alternate, as $\hat{\rho}_1 < 0$
- BET: $z = -6.05$, fewer runs; signs persist, as $\hat{\rho}_1 > 0$
- Bitcoin: $z = 2.68$ rejects independent signs although $\hat{\rho}_1$ is not significant: signs and magnitudes carry different information (all markets: Appendix)

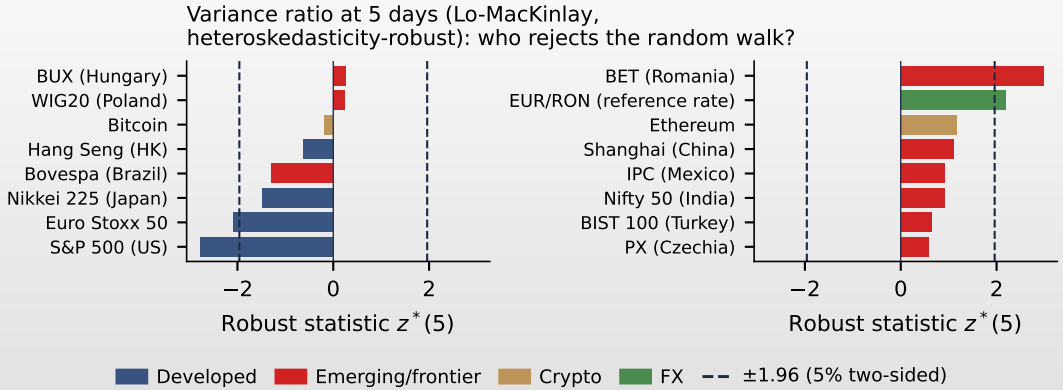
The Lo–MacKinlay Test

- Lo & MacKinlay (1988) estimate $\mathrm{VR}(q)$ from overlapping sums of q daily returns and correct the small-sample bias
  - $\hat{\sigma}_a^2 = \frac{1}{T-1}\sum_t (r_t - \hat{\mu})^2, \quad \hat{\sigma}_c^2(q) = \frac{1}{m}\sum_{t=q}^{T}\left(\sum_{j=0}^{q-1} r_{t-j} - q\hat{\mu}\right)^2$
  - $m = q(T-q+1)(1-q/T); \quad \widehat{\mathrm{VR}}(q) = \hat{\sigma}_c^2(q)/\hat{\sigma}_a^2$
  - $\hat{\mu}$: mean daily return; $\hat{\sigma}_a^2$: variance of one-day returns; $\hat{\sigma}_c^2(q)$: variance of the overlapping q-day sums divided by q; m: a divisor that removes the small-sample bias
- Test statistic under RW1 (homoskedastic returns, i.e. with constant variance)
  - $z(q) = \dfrac{\widehat{\mathrm{VR}}(q) - 1}{\sqrt{2(2q-1)(q-1)/(3qT)}} \approx N(0,1)$
  - the denominator is the standard error of $\widehat{\mathrm{VR}}(q)$ under RW1; $|z| > 1.96$ rejects $\mathrm{VR}(q) = 1$ at 5%
- This z rejects too often under volatility clustering (next slide)

The Heteroskedasticity-Robust Statistic z^*

- Under RW3 plus the conditions of Lo & MacKinlay (1988) (finite moments, weak dependence (mixing), $\mathbb{E}[\varepsilon_t^2 \varepsilon_{t-j}\varepsilon_{t-k}] = 0$ for $j \neq k$), the $\hat{\rho}_k$ are asymptotically uncorrelated with variances $\hat{\delta}_k / T$
  - $\hat{\delta}_k = T\hat{\tau}_k$, with $\hat{\tau}_k$ of the autocorrelation slide: the robust variance of $\sqrt{T}\hat{\rho}_k$, equal to 1 without volatility clustering
- Asymptotic variance of $\widehat{\mathrm{VR}}(q)$ and the test statistic
  - $\hat{\theta}(q) = \sum_{k=1}^{q-1} \left(\frac{2(q-k)}{q} \right)^2 \hat{\delta}_k, \qquad z^*(q) = \frac{\sqrt{T}\left(\widehat{\mathrm{VR}}(q) - 1\right)}{\sqrt{\hat{\theta}(q)}} \approx N(0,1)$
  - $2(q-k)/q$: the weight of $\hat{\rho}_k$ in $\widehat{\mathrm{VR}}(q)$; $|z^*| > 1.96$ rejects $\mathrm{VR}(q) = 1$ at 5%
- RW3 alone is not enough: the conditions hold for GARCH with symmetric innovations, but can fail with a leverage effect
  - Without the condition $\mathbb{E}[\varepsilon_t^2 \varepsilon_{t-j}\varepsilon_{t-k}] = 0$, $\hat{\theta}(q)$ needs the full covariance matrix of the $\hat{\rho}_k$
- S&P 500, $q = 2$: $\widehat{\mathrm{VR}} = 0.901$, $z = -8.12$ but $z^* = -3.64$: still significant, but about 2.2 times smaller in absolute value

The Variance Ratio at 5 Days on 16 Markets



MFM_ch2_variance_ratio

The 5-Day Variance Ratio and the Chow–Denning Test

- Each bar on the previous chart is the robust statistic $z^*(5)$ of the 5-day variance-ratio test for one of the 16 series
- The dashed lines mark ± 1.96, the bar of a two-sided 5% test
- At 5 days the null is rejected for the BET ($z^* = 2.98$), EUR/RON (2.19), the S&P 500 (-2.77) and the Euro Stoxx 50 (-2.08)
- Four horizons at 5% each: if the tests were independent, the probability of at least one false rejection would be $1 - 0.95^4 = 18.5\%$; the VR tests are correlated, so the true value is lower
  - Chow & Denning (1993): $\mathrm{CD} = \max_i |z^*(q_i)|$, the largest of the four $|z^*|$, over $q_i = 2, 5, 10, 20$; 5% critical value 2.49, above 1.96 to pay for the four looks
- The joint test rejects RW3 for 5 of the 16 series: BET, S&P 500, IPC, EUR/RON, Ethereum; Bitcoin: Chow–Denning p-value 0.835, no rejection
- No ranking "developed = efficient": the direction differs (reversal vs persistence)

MFM_ch2_variance_ratio

Beyond Linear Predictability: the Martingale Difference Hypothesis

- **MDH** (Martingale Difference Hypothesis): $\mathbb{E}[r_t \mid \mathcal{I}_{t-1}] = \mu$: no function of the past predicts the mean, linear or not
- RW3 (all $\rho_k = 0$) is weaker. Counterexample: $r_t = \varepsilon_t + \theta(\varepsilon_{t-1}^2 - \sigma^2)$, ε_t i.i.d. symmetric
  - $\sigma^2 = \mathrm{Var}(\varepsilon_t)$, so the extra term has mean 0; $\theta \neq 0$: how strongly yesterday's squared shock moves today's mean
  - all $\rho_k = 0$, yet $\mathbb{E}[r_t \mid \varepsilon_{t-1}] = \theta(\varepsilon_{t-1}^2 - \sigma^2)$ is predictable: autocorrelation tests cannot see it
  - The MDH tests on the next slide look for exactly this kind of dependence (results: Appendix)
- **BDS** (Brock–Dechert–Scheinkman) Brock, Dechert, Scheinkman & LeBaron (1996) compares the frequency of nearby return sequences with that implied by i.i.d. observations; it tests i.i.d., not the MDH: GARCH alone makes it reject
  - A BDS rejection is evidence against RW1, not against weak-form efficiency

What nonlinear tests examine

- MDH tests look for any dependence of the mean on the past, including nonlinear dependence.
- The generalised spectral test looks for mean dependence on nonlinear transformations of past returns
  - it can detect dependence missed by ordinary correlations. Escanciano & Velasco (2006)
- The automatic portmanteau test chooses the number of tested lags using a penalty
  - it can reject MDH through linear dependence but does not detect every nonlinear alternative. Escanciano & Lobato (2009)

Predetermined variables, bias and consistency

- Why does a persistent predictor distort inference in a finite sample?
- A predetermined predictor is known before the return shock it predicts
  - strict exogeneity additionally rules out correlation with shocks at every date.
- Today's dividend–price ratio can forecast next month's return, yet next month's price shock also changes the ratio: this feedback obstructs strict exogeneity.
- Bias is the estimator's repeated-sample mean minus the true parameter: mean slope 0.03 when the true slope is zero means bias 0.03.
- Consistency means convergence to the true parameter as data grow
  - it does not imply zero bias in a finite sample.

HAC standard errors and bandwidth

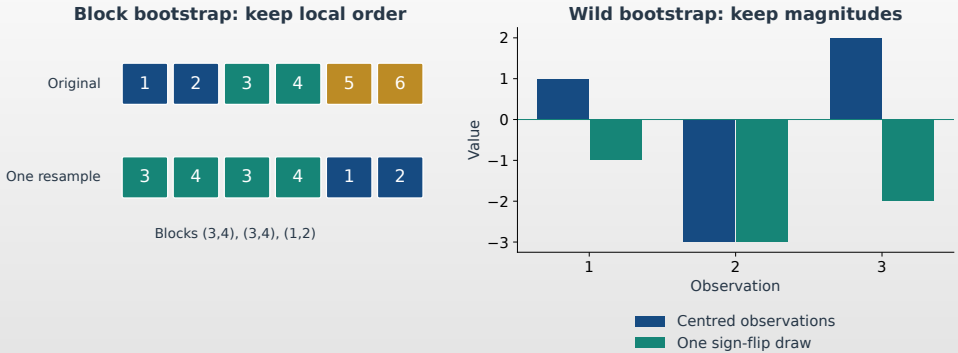
OLS: Ordinary Least Squares.

- HAC means heteroskedasticity and autocorrelation consistent: standard errors allow both changing variance and serial dependence under regularity conditions.
- Newey–West estimates long-run variance by adding weighted lag covariances to the variance
  - it changes uncertainty estimates, not the OLS coefficients.
- For a stationary series with variance 1, lag-1 covariance 0.2 and all other covariances zero, long-run variance is $1 + 2(0.2) = 1.4$.
- Bandwidth specifies how many lags enter the estimate
  - bandwidth 5 uses lags 1–5, with weights depending on the estimator.

Bootstrap: constructing repeated samples

- The bootstrap approximates the sampling distribution by generating many artificial samples and recomputing the statistic.
- An i.i.d. bootstrap draws observations with replacement: from $(1, 2, 3)$ one draw can be $(2, 2, 1)$
  - serial data need a dependence-aware design.
- Block bootstrap resamples consecutive blocks
  - drawing blocks $(1, 2)$ and $(3, 4)$ preserves within-block order, unlike shuffling individual days.
- Wild bootstrap multiplies centred returns by random signs: $(1, -3, 2)$ may become $(-1, -3, -2)$, keeping absolute magnitudes
  - validity depends on the null model.

What each bootstrap preserves



Blocks preserve local ordering; sign flips preserve absolute magnitudes. They impose different assumptions on the artificial samples.

Predictive Regressions and the Stambaugh Bias (1/2)

- Question: does a valuation ratio known today forecast next period's return?
- Model: $r_{t+1} = \alpha + \beta x_t + u_{t+1}, \quad x_{t+1} = \theta + \rho x_t + v_{t+1}$, with $\mathrm{Cov}(u, v) = \sigma_{uv}$
  - r: excess return; x: log dividend–price ratio $dp_t = \ln D_t - \ln P_t$, persistent ($\rho \approx 1$)
  - D_t: dividends paid over the last year; β: the predictive slope, 0 if returns are unpredictable; α, θ: intercepts
  - ρ: the persistence of x; u, v: the shocks of the two equations; σ_{uv}: their covariance; $\sigma_v^2 = \mathrm{Var}(v)$; T: number of periods

Predictive Regressions and the Stambaugh Bias (2/2)

- x_t is predetermined, not strictly exogenous: OLS is consistent but biased Stambaugh (1999)
  - Project u on v: $u = (\sigma_{uv}/\sigma_v^2)\, v + e$, so $\hat{\beta} - \beta = (\sigma_{uv}/\sigma_v^2)(\hat{\rho} - \rho) + \text{noise}$; e: the part of u unrelated to v
  - The step uses the same months and the same demeaned x_t in both regressions; the noise term has mean zero when e is independent of the whole path of x, e.g. (u, v) i.i.d. with a bivariate Normal distribution
  - $\hat{\rho}$ is biased downward in small samples Kendall (1954), and this error passes into $\hat{\beta}$
  - With $\mathbb{E}[\hat{\rho} - \rho] \approx -(1 + 3\rho)/T$: $\quad \mathbb{E}[\hat{\beta} - \beta] \approx -\dfrac{\sigma_{uv}}{\sigma_v^2}\,\dfrac{1 + 3\rho}{T}$
- Price up $\Rightarrow$ dp down: $\sigma_{uv} < 0$, so the bias is **upward**: predictability looks stronger than it is
- Near a unit root the t-statistic is not $N(0, 1)$ either: valid inference by a Bonferroni Q-test over a confidence interval for ρ Campbell & Yogo (2006)

Bonferroni Q: allowing for uncertain persistence

- Campbell–Yogo inference allows for uncertainty about predictor persistence instead of treating its estimate as known. Campbell & Yogo (2006)
- It combines a confidence set for persistence with tests of the predictive slope over that set, allocating the error probability across the steps.
- For example, if plausible persistence values range from 0.95 to 1, a conclusion must account for that range, not only the point estimate 0.98.
- This Q-test concerns predictive regressions
  - it is not the Ljung–Box Q statistic for autocorrelations.

Worked Example: the Stambaugh Bias for the S&P 500

- Monthly log excess return regressed on the lagged dp_t, 1934–2026, $T = 1\,111$ months; data construction: Appendix
- OLS: $\hat{\beta} = 0.267$ (% per month per unit of dp), Newey–West $t = 0.91$: not significant even before any correction
- Inputs: $\hat{\rho} = 0.9988$, $\text{corr}(\hat{u}, \hat{v}) = -0.97$, $\hat{\phi} = \hat{\sigma}_{uv}/\hat{\sigma}_v^2 = -0.925$ (as fractions, not in %)
- Substitution: bias $= -100\,\hat{\phi}\,(1 + 3\hat{\rho})/T = 0.333$ (in %, like $\hat{\beta}$)
  - corrected slope $= 0.267 - 0.333 = -0.066$: the whole estimate is bias
  - the correction is itself estimated: formal inference by Campbell & Yogo (2006) or a bootstrap
  - Kendall's formula is first order in $1/T$ and least accurate when $T(1 - \hat{\rho})$ is small, here 1.4 using the unrounded persistence estimate: read -0.066 as indicative of the sign and size, not as a precise slope
- Simulation with the same ρ and $\text{corr}(u, v)$, $\beta = 0$, $T = 360$ months: the nominal 5% t-test rejects $\beta = 0$ in 26.5% of samples

Out-of-Sample Predictability: R^2_{OS} and the Clark–West Test

- Most predictors in the literature have $R^2_{OS} < 0$ Welch & Goyal (2008); the 2022 update confirms it Goyal, Welch & Zafirov (2024).
- Nested models: under H_0 the larger model estimates a zero β, which inflates its MSPE (mean squared prediction error).
- The Diebold–Mariano test Diebold & Mariano (1995) then rejects less often than its nominal level.
- Clark & West (2007): $f_t = (r_t - \bar{r}_t)^2 - \left[(r_t - \hat{r}_t)^2 - (\bar{r}_t - \hat{r}_t)^2\right]$
  - f_t: the squared-error gain of the predictor, plus the term $(\bar{r}_t - \hat{r}_t)^2$ that removes the noise of estimating β
  - the t test on the mean of f_t is one-sided; a positive mean favours the predictor
- Campbell & Thompson (2008): impose the sign of β and a non-negative forecast.

Comparing forecasts and economic value

- Diebold–Mariano tests the mean loss difference between forecasts
- The Clark–West test is one-sided: it asks whether the mean reduction in squared error, adjusted for estimation noise, is positive
  - A two-sided test would also count a sufficiently negative gain as a departure from zero.
- The Sharpe ratio: mean excess return divided by its SD, at the same frequency.

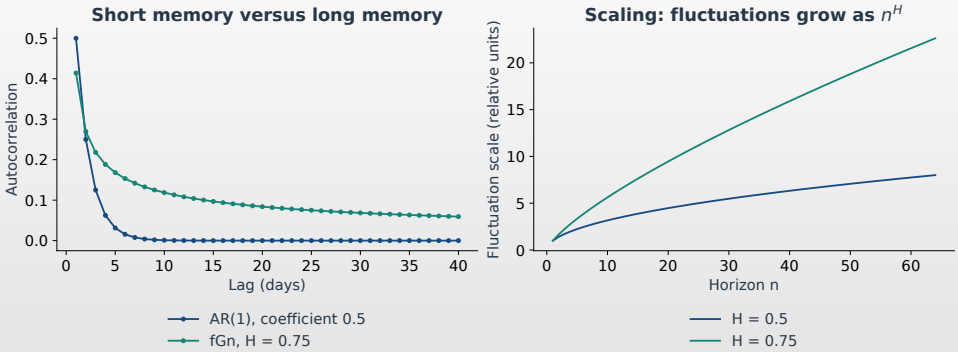
Short memory, long memory and scaling

- Does dependence persist across horizons and market regimes?

AR: AutoRegressive model.

- Short memory has absolutely summable autocovariances
  - an AR(1) with coefficient 0.5 has correlations 0.5^k at lag k.
- Long memory here means positive dependence decaying so slowly that the autocovariance sum diverges
  - power-law decay can produce this.
- Scaling $F(n) \propto n^H$ means multiplying the horizon by 2 multiplies fluctuations by 2^H: about 1.41 at $H = 0.5$, or 1.68 at $H = 0.75$.
- Anti-persistence describes offsetting increments, as in fractional Gaussian noise (fGn) with $H < 0.5$
  - fGn is a stationary Gaussian increment model with partial-sum variance proportional to n^{2H}.

Dependence across lags and horizons

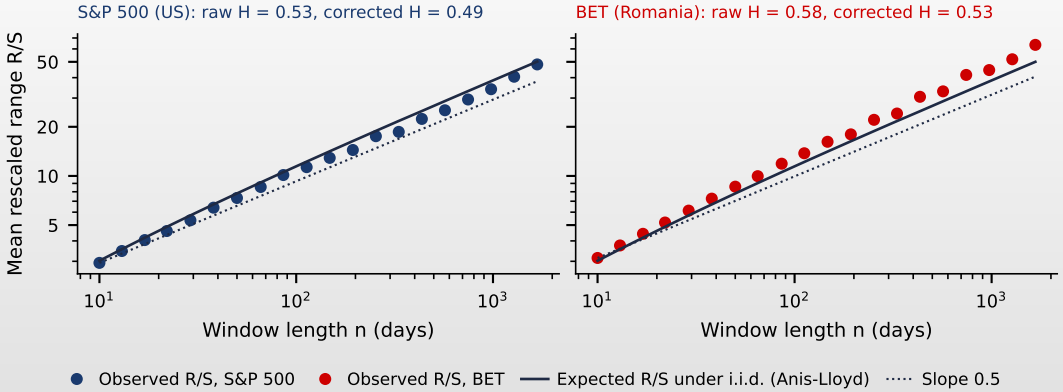


AR(1) correlations fade geometrically; the fGn example fades slowly. The right panel compares fluctuation growth across horizons.

The Rescaled Range (R/S) Statistic

- Using the rescaled range, Hurst (1951) found that years of high Nile floods followed other high-flood years more often than chance allows; history: Appendix
- Construction for a window of n returns with mean $\bar{r}_n$ and standard deviation s_n
  - Cumulative deviations $Y_k = \sum_{i=1}^{k}(r_i - \bar{r}_n)$, $k = 1, \ldots, n$
  - Range $R_n = \max_k Y_k - \min_k Y_k$; rescaled range R_n/s_n
- Scaling law: $\mathbb{E}[R_n/s_n] \approx c\, n^H$, where H is the **Hurst exponent**
  - c: a constant; taking logs, $\ln(R/S) \approx \ln c + H \ln n$, so H is a slope
  - Estimate: slope of $\ln(R/S)$ on $\ln n$, averaging R/S over non-overlapping windows
- $H = 0.5$: no long memory (short-run dependence allowed); $H > 0.5$: persistence, long memory with $\rho_k \sim k^{2H-2}$; $H < 0.5$: anti-persistence

R/S Analysis: S&P 500 and BET (1/2)



MFM_ch2_hurst

R/S Analysis: S&P 500 and BET (2/2)

- On the previous chart, the dots show the average R/S for each window length n.
- The solid line is the Anis & Lloyd (1976) expectation for i.i.d. returns from the Normal distribution (formula: ▸ Appendix).
- In small samples the raw slope exceeds 0.5 even for i.i.d. returns.
- The Anis–Lloyd correction: $H_{AL} = 0.5+$ the slope of the regression of $\ln(R/S) - \ln \mathbb{E}[R_n/s_n]$ on $\ln n$.
- For the S&P 500, the raw exponent is 0.530 and the corrected one 0.486: the apparent "persistence" came from small-sample bias.
- For the BET, the raw exponent is 0.578 and the corrected one 0.533: persistence remains after the correction.
- Open question: is BET's $H > 0.5$ long memory, or just its lag-1 autocorrelation $\hat{\rho}_1 = 0.109$?
- The classical R/S statistic cannot tell the two apart.

MFM_ch2_hurst

Variance Ratio and Hurst Exponent: Long or Short Memory?

- For an exactly self-similar process (fractional Gaussian noise), the variance of a sum of q returns is $\mathrm{Var}(r_t + \cdots + r_{t-q+1}) = q^{2H}\mathrm{Var}(r_t)$.
- Hence $\mathrm{VR}(q) = q^{2H-1}$, and an observed $\mathrm{VR}(q)$ implies a Hurst exponent $H(q) = 0.5 + \ln \mathrm{VR}(q)/(2 \ln q)$.
- At $q = 20$ days, the implied exponent is 0.553 for the BET and 0.453 for the S&P 500.
- The problem: a process without long memory can give $H(q) \neq 0.5$
  - Example: an AR(1) with $\rho_1 = -0.1$ has autocorrelations $\rho_k = (-0.1)^k$, so only short memory.
  - For this AR(1), $\mathrm{VR}(20) = 0.826$, so $H(20) = 0.468$: the process looks "anti-persistent" without any long memory.
  - At $q = 250$, the same process gives $H(250) = 0.482$: for short memory, $H(q)$ approaches 0.5 as q grows.
- An AR(1) with BET's autocorrelation, $\hat{\rho}_1 = 0.109$, gives $H(20) = 0.535$.
- So most of the gap between BET's $H(20) = 0.553$ and 0.5 is short memory.
- Separating long from short memory needs a statistic built for it: Lo's modified R/S (next slide).

Lo's Modified R/S Test

- Lo (1991): short-run autocorrelation inflates R/S and mimics long memory
- Modification: s_n is replaced by a long-run (HAC) standard deviation $\hat{\sigma}_\ell$ with ℓ lags
  - $\hat{\sigma}_\ell^2 = \hat{\gamma}_0 + 2\sum_{j=1}^{\ell}\left(1 - \frac{j}{\ell+1}\right)\hat{\gamma}_j$, $\hat{\gamma}_j$: autocovariance at lag j ($\hat{\gamma}_0$: the variance), with declining weights
  - $V = \dfrac{\max_k Y_k - \min_k Y_k}{\hat{\sigma}_\ell \sqrt{T}}$; Y_k: cumulative deviations over the full sample of T returns; bandwidth ℓ chosen as in Andrews (1991)
- Under short memory, V lies in $[0.809, 1.862]$ with probability 95%
- For the S&P 500 ($V = 1.559$), Bitcoin (1.547), PX (1.750) and EUR/RON (1.429), V is inside the interval: short memory is not rejected
- For the BET ($V = 1.961$, with $\ell = 7$ lags) and Ethereum ($V = 1.980$), V is just above 1.862
- This is weak evidence of long memory

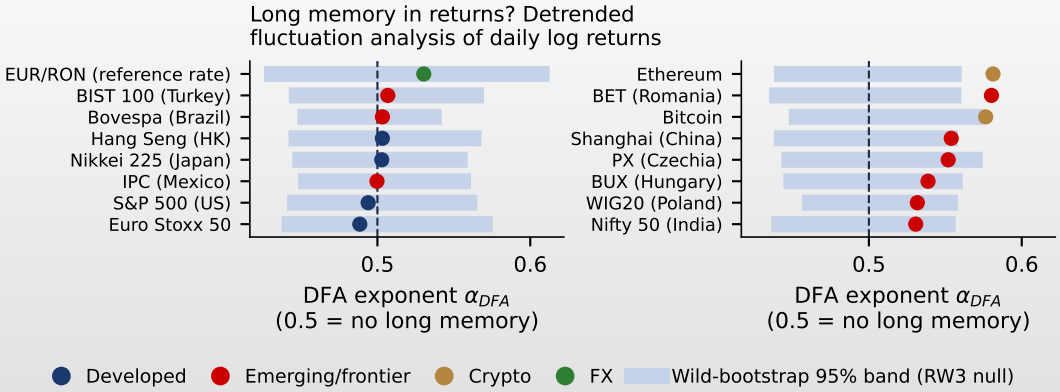
Detrended Fluctuation Analysis (DFA) (1/2)

- Peng et al. (1994): DFA measures how the fluctuation of cumulative returns grows with the window length.
- Removing a linear trend from the profile removes a constant mean of returns in each box.
- If the mean of returns itself trends, higher-order DFA is needed.
- The algorithm
  - The profile $y_k = \sum_{t \leq k} (r_t - \bar{r})$ is the running sum of deviations from the mean, over the full sample.
  - The profile is split into boxes of length n.
  - In each box a linear trend is fitted and the residuals are kept.
  - $F(n)$ is the root mean square of all residuals; the calculation is repeated for many n.
- $F(n) \propto n^{\alpha_{DFA}}$: the DFA exponent α_{DFA} plays the role of H ($\alpha_{DFA} = 0.5$: no long memory).

Detrended Fluctuation Analysis (DFA) (2/2)

- The null band comes from a **wild bootstrap**: each residual $e_t = r_t - \bar{r}$ is multiplied by a random sign $\eta_t = \pm 1$, with probability 1/2 (199 replications).
- The sign flip keeps the volatility $|e_t|$, allowed under RW3, but destroys serial correlation.
- A random shuffle would destroy both volatility and correlation.
- The implicit null is stricter than RW3: signs symmetric and independent, given volatility.
- The band is therefore a benchmark, not an exact test of RW3 and not a confidence interval for α_{DFA}.

DFA Exponent Across Markets (1/2)



MFM_ch2_hurst

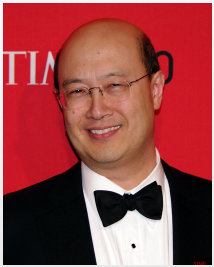
DFA Exponent Across Markets (2/2)

CI: Confidence Interval.

- On the previous chart, each dot is the DFA exponent α_{DFA} of one market.
- The bar is the 95% null band from random sign flips (volatility kept).
- For the S&P 500 (0.494), Euro Stoxx 50 (0.488), Nikkei (0.503), PX (0.552) and Shanghai (0.554), the exponent is inside the band.
- Three of the 16 series are above the band: Ethereum (0.581), BET (0.580, upper limit 0.561) and, borderline, Bitcoin (0.576, upper limit 0.575).
- A band from random shuffles also destroys volatility clustering; with it, five series seem to have memory.
- PX and Shanghai leave only the shuffled band: their apparent memory comes from volatility clustering.
- The local Whittle estimator of the long-memory parameter $d = H - 0.5$ gives for the BET $\hat{d} = 0.118$, with 95% CI $[0.062, 0.174]$ (Appendix).
- Conclusion: the evidence of long memory in returns is weak and depends on the method; memory in volatility (Chapter 1) is strong.

MFM_ch2_hurst

Adaptive Markets Hypothesis (AMH)



Andrew W. Lo, MIT

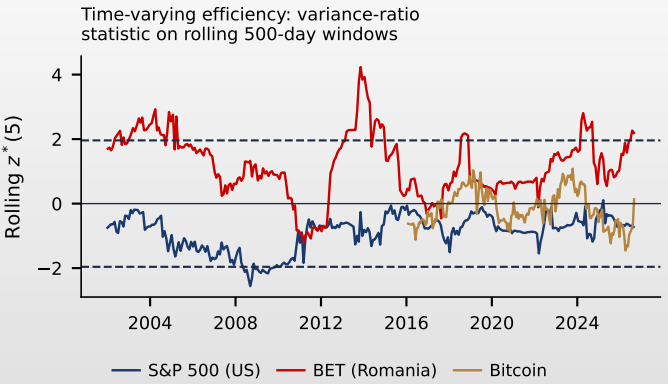
Photo: David Shankbone (2012); CC BY 3.0; Wikimedia Commons

- Full-sample tests give weak and unstable evidence; one possible reason: efficiency varies over time
- Andrew Lo proposed the Adaptive Markets Hypothesis (AMH) in Lo (2004) and developed it in the book Lo (2017)
- Markets are an ecology of competing species (traders, funds, algorithms) that adapt by trial and error, not by perfect optimisation
  - efficiency is a matter of degree and depends on the number of competitors, the profit opportunities and the environment
- Anomalies appear, weaken and may reappear
- Caution: time-varying statistics alone do not identify the AMH, since an Efficient Market Hypothesis (EMH) with time-varying risk premia also allows them; comparison table: Appendix
- Lim & Brooks (2011), a survey: stock-market efficiency varies over time, especially around crises
- Urquhart & McGroarty (2016): on four major indices (US, UK, Japan, euro area), serial dependence appears and disappears across subperiods with market conditions

Measuring Time-Varying Efficiency

- Compute a statistic on a **rolling window**: 500 trading days (about two years), moved forward by 21 days (one month)
  - Statistics: $z^*(5)$ of the variance-ratio test and the DFA exponent
- The DFA band is computed as before (random sign flips), but separately in each window, so it follows the window's volatility
- A fixed band simulated from i.i.d. Student-t returns with 4 degrees of freedom is $[0.380, 0.615]$
- The fixed band ignores volatility clustering, allowed under RW3; volatility clustering raises α_{DFA}, so the fixed band flags too many windows
- Two cautions for reading the results
  - Overlapping windows are strongly correlated: runs of rejections are not independent evidence
  - Even under the null about 5% of windows reject; a share near 5% is a description, not a calibrated test of stability

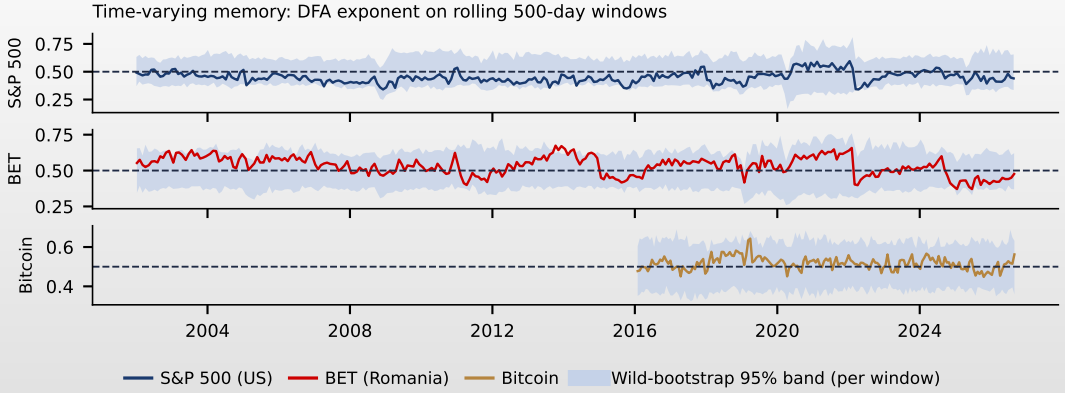
Rolling Variance Ratio: S&P 500, BET, Bitcoin



- The chart shows $z^*(5)$ on rolling 500-day windows; the dashed lines mark ± 1.96
- Share of windows with $|z^*(5)| > 1.96$: S&P 500 6.4%, BET 22.4%, Bitcoin 0.0%
- S&P 500: the full-sample rejection sits in a few turbulent windows
- BET: rejections come in episodes, then fade
- FTSE Russell upgraded Romania to Secondary Emerging from 21 Sep 2020, a possible change in the investor base (▸ Appendix)

MFM_ch2_rolling_efficiency

Rolling DFA Exponent: S&P 500, BET, Bitcoin (1/2)



MFM_ch2_rolling_efficiency

Rolling DFA Exponent: S&P 500, BET, Bitcoin (2/2)

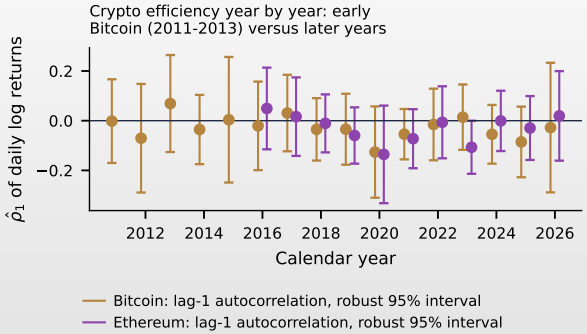
- On the previous chart, the shaded area is the wild-bootstrap 95% band, computed separately in each 500-day window.
- The exponent leaves its own band in 2.0% of windows for the S&P 500, 3.1% for the BET and 0% for Bitcoin.
- With the fixed Student-t_4 band, the shares would be 6.4%, 11.2% and 1.1%: the apparent BET "episodes" were volatility clustering.
- The test has low power: for the BET, the band averages $[0.36, 0.64]$.
- 500 observations cannot pin down a scaling exponent.
- Over the full sample, Bitcoin is only just above its band (0.576 against 0.575); in 500-day windows the test has too little power for so small an effect.

MFM_ch2_rolling_efficiency

The AMH, Tested

- Kim, Shamsuddin & Lim (2011) study the predictability of US stocks over a century of data.
- They show that predictability varies with market conditions (crises, bubbles) rather than along a trend; by linking predictability to state variables in a regression, they make the AMH testable.
- Our version: regress the rolling $z^*(5)$ on the window's annualised volatility (%) and a crisis indicator
  - The crisis indicator is 1 if the window overlaps Sep 2008–Jun 2009 or Feb–Jun 2020.
  - Newey–West errors with 24 lags, because consecutive windows share 95% of their days.
  - Volatility is the control variable: windows with similar volatility are compared.
- S&P 500: the crisis coefficient is -0.67 ($t = -2.37$) and the volatility coefficient 0.006 ($t = 0.27$); $R^2 = 0.20$.
- BET: the crisis coefficient is -0.82 ($t = -2.77$) and the volatility coefficient -0.008 ($t = -0.36$); $R^2 = 0.15$.
- In crisis windows $z^*(5)$ is lower, i.e. closer to reversal, even at equal volatility: a conditional association, not a causal effect.
- $z^*(5)$ depends both on $\widehat{\mathrm{VR}}(5) - 1$ and on its standard error; a change in z^* can come from either.
- State dependence is consistent with the AMH but does not identify it uniquely.

Crypto Efficiency Year by Year



- The AMH predicts that a new market becomes more efficient over time
- The chart: lag-1 autocorrelation for each calendar year, robust 95% intervals
- Only one of the 27 yearly autocorrelations is significant at 5%: Ethereum in 2023, $\hat{\rho}_1 = -0.107$
- By chance alone we would expect about 1.35 significant results
- Urquhart (2016) rejects efficiency in the early years, with other tests, another price source and classical standard errors (▸ Appendix)
- Power: one year of 365 returns gives a robust SE of about 0.06–0.13, so small effects are invisible

MFM_ch2_crypto_efficiency

Event-study inference

- The tests of the event-study section assume independent firms and constant variance; here we relax both.

Cross-sectional dependence and standardisation

- Two banks hit by the same news need not be independent observations: their abnormal returns can move together.
- Event-induced variance means the event itself raises dispersion
  - Using only quiet pre-event days can understate uncertainty.

Event Studies: the Prediction Error

- Prediction-error variance MacKinlay (1997): $\mathrm{Var}(AR_\tau) = \sigma_\varepsilon^2\Big[1 + \frac{1}{T_0} + \frac{(r_{m,\tau} - \hat{\mu}_m)^2}{(T_0 - 1)\hat{\sigma}_m^2}\Big]$, T_0: estimation days
  - $\hat{\mu}_m$, $\hat{\sigma}_m^2$: mean and variance of the market return in the estimation window.
  - The bracket grows when the event-day market move is unusual.
- Back to the bank tax of December 2018 (event-study section): on the event day the log return of the BET was -11.9%.
- The bracket is then 2.12, so the prediction error grows.
- The t statistic of Banca Transilvania (TLV) shrinks in absolute value from -10.0 to -6.85.
- For BRD, the day-0 abnormal return is -5.6%, with $t = -4.03$.
- For several firms, standardise each AR by its own prediction error and use the cross-sectional variance, robust to event-induced variance Boehmer, Musumeci & Poulsen (1991).

Event Studies: Correlation Across Firms

- When the event falls on the same day for all firms, abnormal returns are cross-correlated and cross-sectional tests reject too often.
- The equal-weighted portfolio test already allows for this correlation.
- Let $\bar{\varrho}$ be the average correlation between the firms' abnormal returns and N the number of firms.
- For TLV and BRD, $\bar{\varrho} = 0.153$.
- For the mean AR_0, the t computed under independence is -7.55; computed on the portfolio, t is -7.03.
- Kolari & Pynnönen (2010) multiply the cross-sectional statistic by $\sqrt{(1-\bar{\varrho})/(1+(N-1)\bar{\varrho})}$.
- For $N = 2$, the factor is 0.857, so the corrected $t \approx -6.47$, close to the portfolio t (-7.03).
- With ten firms and the same correlation, $(1+9\bar{\varrho})/(1-\bar{\varrho}) = 2.81$: the naive test would understate the variance of the mean AR by a factor of 2.81.

Anomalies and behavioural explanations

- Can behavioural biases explain the anomalies that survive?

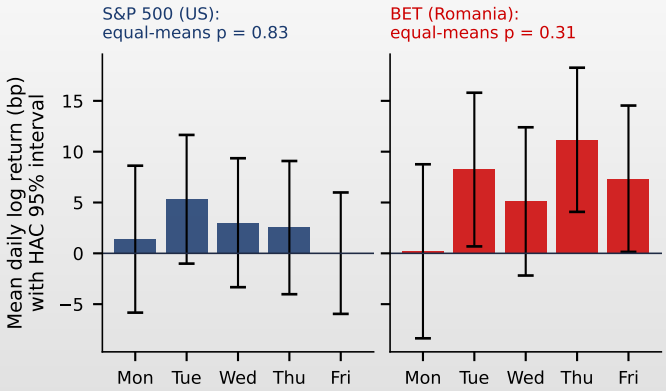
Momentum

- Cross-sectional momentum ranks assets: if A gained 10% and B gained 2%, a relative strategy may buy A and short B.
- Time-series momentum uses each asset's own past sign: both A and B would receive positive signals in that example.

Anomalies and Calendar Effects

- An **anomaly** is a return pattern the asset-pricing model cannot explain: an omitted risk factor (misspecified model), mispricing (behavioural bias) or data mining
  - Families: calendar, time-series (index momentum), cross-sectional (size, value, momentum: **Chapter 3**)
  - Returns 26% lower out of sample and 58% lower after publication McLean & Pontiff (2016)
- Calendar effects: Monday French (1980); January Rozeff & Kinney (1976); turn of the month (TOM) Lakonishok & Smidt (1988); monthly effect (first half of the month) Ariel (1987)
  - Day of week: $r_t = \sum_{d=1}^{5} \beta_d D_{d,t} + u_t$, Wald test of $\beta_1 = \cdots = \beta_5$ (four equality restrictions, reference χ^2_4); TOM and January: one dummy each
  - $D_{d,t} = 1$ if day t is weekday d (Monday, ..., Friday), else 0; β_d: the mean return on weekday d; u_t: the residual
  - Newey–West HAC errors Newey & West (1987), 5 daily or 3 monthly lags: residuals are heteroskedastic and autocorrelated

Day-of-the-Week Effect: S&P 500 and BET



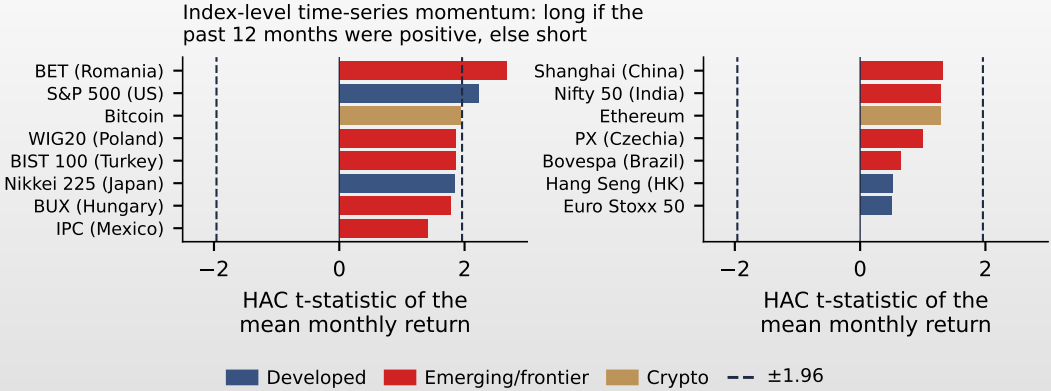
- Mean daily log return by weekday, in basis points (bp; 1 bp = 0.01%), 2000–2026, HAC 95% intervals
- S&P 500: from 0.0 bp (Friday) to 5.3 bp (Tuesday); equal-means $p = 0.83$: no Monday effect since 2000, unlike French (1980)
- BET: Thursday 11.2 bp ($t = 3.09$), yet the joint test gives $p = 0.31$
- Two different nulls:
  - a day's mean equal to zero (mixes the day effect with the positive average return)
  - all five means equal (the day effect itself)

MFM_ch2_calendar_anomalies

Time-Series Momentum on Indices (1/3)

- Time-series momentum (TSMOM) Moskowitz, Ooi & Pedersen (2012): signal $s_t = \text{sign}\left(\sum_{j=1}^{12} r_{t-j}\right)$
  - $+1$ (long) if the past 12 monthly log returns sum to a gain, -1 (short) otherwise
  - The strategy return is $s_t R_t$, R_t the index return of month t
- On the next chart, each bar is the HAC t (6 lags) of the strategy's mean monthly return in one market.

Time-Series Momentum on Indices (2/3)



MFM_ch2_tsmom

Time-Series Momentum on Indices (3/3)

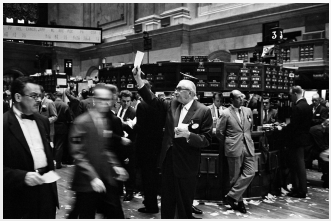
- The strategy has a significant mean return for the BET (1.40% a month, $t = 2.68$) and the S&P 500 ($t = 2.22$).
- For Bitcoin, $t = 1.94$, below 1.96: not significant at 5%.
- The TSMOM Sharpe ratio is below buy-and-hold: 0.65 vs 0.75 for the BET, 0.48 vs 0.52 for the S&P 500.
- Much of the TSMOM return is the equity premium earned while long.
- The Sharpe ratio is annualised ($\sqrt{12}\times$ monthly mean over monthly SD), on price returns, with no risk-free rate subtracted.
- Data limits: price indices without financing rates or dividends; 15 markets tested (full table: Appendix).

MFM_ch2_tsmom

Behavioural concepts

- Loss aversion: losses weigh more than equal gains, relative to a reference point.
- Prospect theory: gains and losses are judged relative to a reference point, with diminishing sensitivity and probability weighting.
- Disposition effect: the tendency to sell winners and keep losers.
- Anchoring: excess weight on an initial number, such as the purchase price.
- Overconfidence: overstating the precision of one's beliefs, which leads to too much trading.
- Limited attention: relevant information left unprocessed.
- Herding: following others' actions without an independent assessment.
- Underreaction and overreaction: a price adjusts too little or too much to news at first; the second is followed by reversal.

Psychology in Prices, from the Floor to the Feed



New York Stock Exchange trading floor, 1963

Photo: Thomas J. O'Halloran (1963); public domain; Wikimedia Commons

- **Excess volatility** Shiller (1981): prices move far more than the dividends that should justify them; **overreaction**: the source of the 3–5-year reversal De Bondt & Thaler (1985); portraits: ▸ Appendix
- **Prospect theory** Kahneman & Tversky (1979) explains the disposition effect
- **Herding**: in an illiquid market such as the BVB, Pele, Mazurencu-Marinescu & Nijkamp (2013) link herding to log-periodic bubbles; today its channels are order books, algorithms and social media (Chapter 10)
- Biases alone do not break efficiency: rational arbitrageurs could trade against them. Can they?

Limits to Arbitrage



Andrei Shleifer, Harvard

Photo: Manuel Castells (2018); CC BY-SA 2.0; Wikimedia Commons

- Shleifer & Vishny (1997): real arbitrage is risky and needs capital
- **Noise-trader risk** De Long, Shleifer, Summers & Waldmann (1990): mispricing can get worse before it corrects
  - Noise traders: investors who trade on signals without information content (sentiment, rumours)
- Funding constraints: managers lose clients after losses and must close positions at the worst time
- Costs: short-selling limits, margins, trading costs, no perfect substitutes
- Anomalies can persist where arbitrage is hard: in small, illiquid markets such as the Romanian one (BET) we find the strongest persistence among the stock indices
- The AMH reconciles both views: rational and biased behaviour coexist and evolve

Machine learning and research design

- How can a flexible model or research design be evaluated honestly?
- Anomalies and behavioural biases are usually found by searching among many predictors.

Can Machine Learning Beat Efficient Markets?

- ML (machine learning) learns a prediction rule from data and captures nonlinear patterns that autocorrelation tests miss (Chapter 13)
- Gu, Kelly & Xiu (2020): trees and neural networks give large economic gains over linear regressions in the US cross-section, yet their monthly out-of-sample R^2 is below 1%
- Large language models read news and filings: Tak & Pele (2025) forecast stock trends with LLaMA
- When many predictors are available, what does an in-sample rejection of efficiency mean?

Validation, leakage and shrinkage

- Flexible models, such as trees and neural networks, can also fit accidental sample patterns (overfitting).
- Walk-forward validation: the model is estimated only on data before the test period.
- Leakage: a forecast uses information unavailable at its date, such as an accounting figure revised later.
- Ridge adds to the OLS loss a penalty proportional to the squared coefficients; shrinkage pulls the coefficients towards zero and reduces estimation variance.
- With Normal errors and a Normal prior centred at zero, ridge is the posterior mean of the coefficients (Bayesian learning).

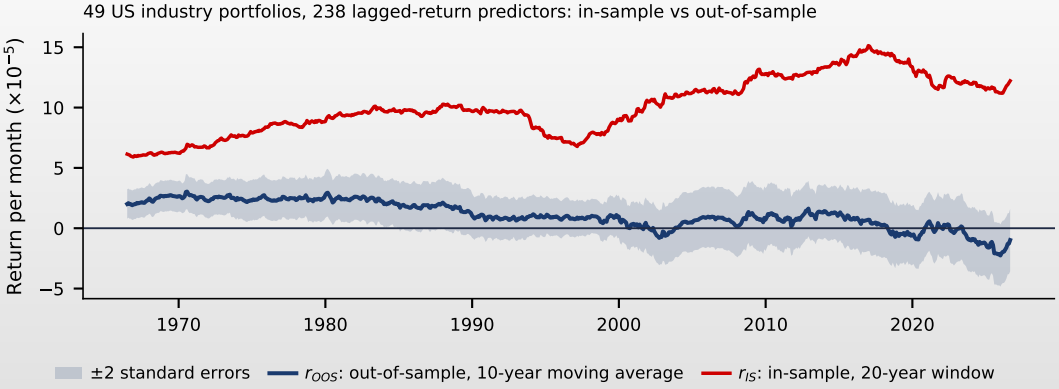
Case Study: Rational Learning with Many Predictors (1/2)

- Martin & Nagel (2022): investors price N assets whose cash flows depend on J characteristics with **unknown** coefficients and learn them optimally (Bayes) from past data (research question: Appendix)
  - The posterior mean of the coefficients is the ridge estimator of the previous slide.
  - Investors are risk-neutral, so predictability cannot be a risk premium.
- **In sample**: if the number of predictors per asset, J/N, tends to a constant $\psi > 0$, the test of no predictability rejects with probability tending to 1 (Martin & Nagel (2022), Prop. 4).
- A rejection is not evidence of irrationality: the rational-expectations null is the wrong benchmark when investors must learn the parameters (formal results: Appendix).

Case Study: Rational Learning with Many Predictors (2/2)

- **Out of sample**: a portfolio built with coefficients from other periods has mean return zero, $\mathbb{E}[r_{OOS}] = 0$ (Martin & Nagel (2022), Prop. 6; model and empirical design: Appendix).
- r_{OOS} is the return of that portfolio in a period not used to estimate its coefficients.
- So out-of-sample tests keep their economic meaning: in this model, they test efficiency.
- For US stocks, the in-sample R^2 is about 1%.
- The out-of-sample return r_{OOS} is much lower than the in-sample one and, after 2014, within two standard errors of zero (Martin & Nagel (2022), Fig. 8).

Replication on Course Data: In Sample vs Out of Sample (1/2)



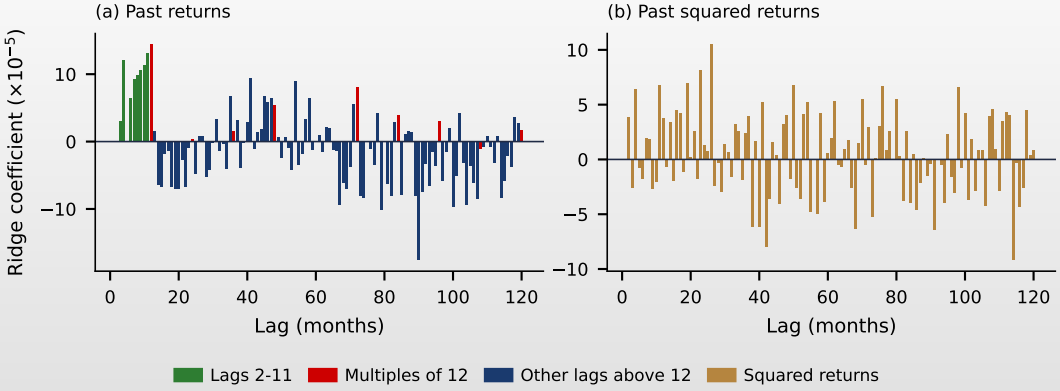
MFM_ch2_big_data

Replication on Course Data: In Sample vs Out of Sample (2/2)

- Data: 49 industry portfolios (Kenneth French Data Library), July 1926 – August 2026; 238 lagged-return predictors; rolling 20-year OLS, as in Martin & Nagel (2022), Sec. 7.
- r_{IS}, the in-sample return of the predicted-return portfolio, is between 5.90 and 15.14 ($\times 10^{-5}$ a month), positive by construction.
- The mean in-sample R^2 is 5.4%.
- r_{OOS}, the out-of-sample return of the same portfolio, has a mean of 1.05×10^{-5} over 842 months, about 11% of r_{IS}.
- Up to June 1986, the mean of r_{OOS} is 4.16 standard errors above zero; after June 1986, only 0.62 standard errors.
- The ± 2 standard-error bands are descriptive, not a test of $\mathbb{E}[r_{OOS}] = 0$.
- $J/N = 238/49 \approx 4.9$: each 20-year window has 49×240 monthly observations, so OLS remains possible, but the in-sample fit is inflated (mean R^2 of 5.4% against about 1% in Martin & Nagel (2022); details: Appendix).

MFM_ch2_big_data

Replication on Course Data: Ridge Coefficients by Lag (1/2)



MFM_ch2_big_data

Replication on Course Data: Ridge Coefficients by Lag (2/2)

- Full-sample ridge, January 1971 – August 2026; the penalty is chosen by leave-one-year-out cross-validation (56 years); compare Martin & Nagel (2022), Fig. 7.
- For lags 2–11, 8 of 10 coefficients are positive, matching **momentum** Jegadeesh & Titman (1993).
- For multiples of 12, 9 of 10 coefficients are positive, matching **seasonality** Heston & Sadka (2008).
- For the other lags above 12, 67% of the coefficients are negative, matching long-run **reversal** De Bondt & Thaler (1985).
- The coefficients on squared returns have no dominant sign.
- In cross-validation, ridge has $R^2 = 0.013\%$ and OLS -2.26%.
- Ridge recovers the known anomalies, yet they predict almost nothing in the years left out (limits and extensions: ▸ Appendix).

MFM_ch2_big_data

ETFs, placebo and price discovery

ETF: Exchange-Traded Fund.

- A spot ETF holds the underlying asset
  - one unit gives exposure through the fund, not through a futures contract.
- A placebo repeats the analysis at a date without the event: a similar apparent change in 2019 weakens a claim that the 2024 launch caused the change.
- Price discovery is the incorporation of information into prices: if exchange A reacts to news first and B follows, A may contribute more to price discovery.

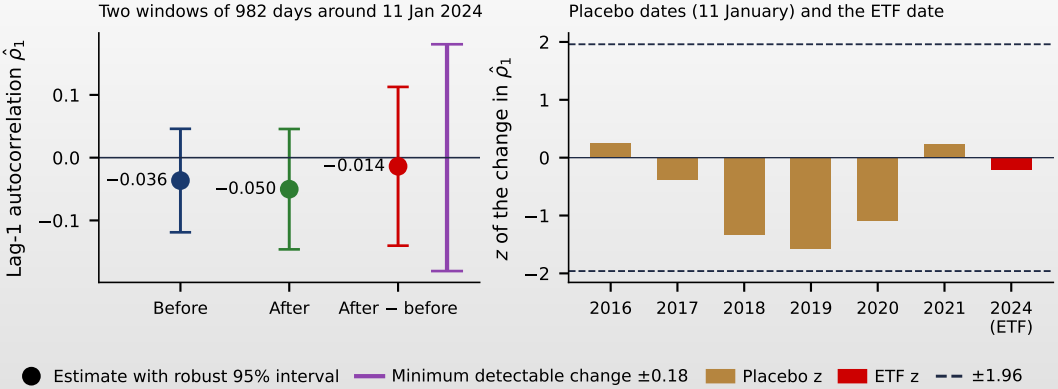
Mini-Case: Did Spot ETFs Make Bitcoin More Efficient?

- An artificial-intelligence (AI) assistant proposes a typical hypothesis: "US spot Bitcoin ETFs (trading since 11 January 2024) brought arbitrageurs, so Bitcoin became more efficient"
- The discovery loop; reusable prompts: Appendix
  - **literature**: Kia, Liu, Li, Song & Xu (2026) find that after the spot ETFs the spot market became the dominant venue for price discovery again; Tran & Leirvik (2020): crypto efficiency varies and has increased
  - **hypothesis**: $H_0 : \rho_{1,\text{post}} = \rho_{1,\text{pre}}$, two windows of equal length around the launch date (first trading day)
  - **data and code**: daily Bitcoin closing price (USD); first replicate the chapter's full-sample $\hat{\rho}_1 = -0.022$
- Checks fixed before seeing results; rationale for each: Appendix
  - robust standard errors (not $1/\sqrt{T}$); equal windows; the launch date only, plus placebo dates
  - a control asset for confounders (the April 2024 halving, rates, inflows)
  - a correction for every statistic tried; the minimum detectable effect

Estimation and Power (1/3)

- The test: two non-overlapping windows of 982 days, before and after 11 January 2024.
- $z = \Delta\hat{\rho}_1 / \sqrt{s^2_{\text{pre}} + s^2_{\text{post}}}$, with $\Delta\hat{\rho}_1 = \hat{\rho}_{1,\text{post}} - \hat{\rho}_{1,\text{pre}}$ and the robust standard errors s of the two estimates.
- **Minimum detectable effect (MDE)**: the smallest true change detected with 80% power at the 5% level.
- $\text{MDE} = (z_{0.975} + z_{0.80}) \times \text{SE}(\Delta\hat{\rho}_1) \approx 2.8 \times \text{SE}(\Delta\hat{\rho}_1)$
  - z_p: the p-quantile of the standard Normal distribution; $z_{0.975} = 1.96$ and $z_{0.80} = 0.84$
- Placebo dates: 11 January of each year 2016–2021, with windows of 982 returns.

Mini-Case: Estimate, Placebo and MDE



MFM_ch2_ai_discovery

Estimation and Power (2/3)

- The previous chart, left: $\hat{\rho}_1$ before and after the launch, with robust 95% intervals; the third point is the difference, with the $\pm$MDE bar.
- Right: the z statistic for the ETF date, among those of the six placebo dates.
- Before (May 2021 – Jan 2024): $\hat{\rho}_1 = -0.036$ ($s = 0.042$), VR(5) $= 1.000$ ($z^* = 0.00$).
- After (Jan 2024 – Sep 2026): $\hat{\rho}_1 = -0.050$ ($s = 0.049$), VR(5) $= 0.939$ ($z^* = -0.64$).
- $\Delta\hat{\rho}_1 = -0.014$, $z = -0.21$: an unchanged ρ_1 is not rejected.
- None of the six placebo dates gives a significant change in ρ_1 at 5%; the largest statistic is $|z| = 1.57$.

Estimation and Power (3/3)

- Here $\mathrm{SE}(\Delta\hat{\rho}_1) = \sqrt{0.042^2 + 0.049^2} \approx 0.065$, so the MDE is about 0.18.
- "Not significant" is not "no effect": the MDE (0.18) is about 13 times the estimated change (-0.014).
- If Bitcoin was already close to efficient before 2024, this design could not detect an improvement smaller than 0.18.
- A better outcome variable: intraday price discovery across venues, not daily ρ_1.

MFM_ch2_ai_discovery

Difference-in-differences: example and assumptions

- The before/after test cannot separate the effect of the ETFs from other changes in 2024.
- A causal conclusion needs a control group, for example Ethereum before its own ETFs (July 2024).
- Difference-in-differences (DiD) subtracts the control group's before–after change from the treated group's change.
- Parallel trends: without treatment, the treated and control groups would have changed in the same way.
- Similar pre-treatment trends support this assumption but cannot prove it.
- No anticipation: the treated group does not react before treatment.
- No spillover: treatment does not affect the control group; for example, a Bitcoin ETF must not change Ethereum trading.

Your Turn: A Project Seed

- **Question**: did the spot ETFs change the time-varying efficiency of Bitcoin, relative to a control asset?
  - rolling VR and DFA (this chapter) for Bitcoin and Ethereum; Ethereum is untreated until July 2024
  - if parallel trends, no anticipation and no spillovers are not credible, report descriptive, not causal, differences
- **Replicate first**: the chapter's yearly Bitcoin autocorrelations and the mini-case numbers; record which Bitcoin price series and first date each window uses, placebo windows included
- **Pre-register**: H_0, H_1, dates, windows and corrections, committed before the results
- **Report**: robust inference, placebo tests and the MDE
- Publish the code as a Quantlet
- **Log**: record the prompts used in `AI_USE.md`
- Record in `AI_ERRORS.md` every error of the AI assistant you caught (for example an i.i.d. standard error or an invented reference)

Synthesis, Check Yourself, and Next: Chapter 3

Five steps for every efficiency claim

- **Information set**: which information $\mathcal{I}_t$ do we assume prices reflect (past prices, public news)?
- **Null**: RW1, RW3 or the martingale difference, stated together with the risk model
- **Statistic**: its variance valid under volatility clustering
- **Uncertainty**: horizons, markets and windows tried; out-of-sample evidence
- **Economics**: a statistical rejection is not an exploitable profit after costs, and a profit does not prove an irrational market

Check yourself

- Why can RW3 be rejected for the S&P 500 while the market remains efficient in the economic sense?
- Why does BET's $H(20) > 0.5$ not prove long memory?
- Which of the two results, BET's persistence or Bitcoin's non-rejection, would survive a change of test, window and inference?
- Next, Chapter 3: risk-adjusted returns with factor models, and the factor zoo
- Textbooks: Campbell, Lo & MacKinlay (1997), ch. 2; Franke, Härdle & Hafner (2019); advanced courses: ▸ Appendix

Appendix: Pioneers, Bachelier and Samuelson



Louis Bachelier (1870–1946)

Photo: unknown author; public domain; Wikimedia Commons



Paul A. Samuelson (1915–2009)

Photo: Bernard Gotfryd (1970–1975); public domain; Wikimedia Commons

- Bachelier (1900): prices of French government bonds move like a random walk
  - "The mathematical expectation of the speculator is zero"
- Samuelson (1965): if prices are **properly anticipated**, their changes are unpredictable
- Both results are the ancestors of the Efficient Market Hypothesis

Back

Appendix: Eugene Fama



Eugene F. Fama, Nobel Prize 2013

Photo: Bengt Nyman (2013); CC BY 2.0; Wikimedia Commons

- Fama (1970) gave the definition used in this chapter and the three forms of efficiency
- Nobel Prize 2013, shared with Robert Shiller (a sceptic of efficiency) and Lars Peter Hansen
- Update and survey: Fama (1991); an accessible defence of the EMH against its critics: Malkiel (2003)

Back

Appendix: Summary Table, Weak-Form Tests 2000–2026

Market	T	$\hat{\rho}_1$	robust SE	$\widehat{\mathrm{VR}}(5)$	$z^*(5)$	CD p	Runs z
S&P 500 (US)	6 714	−0.099	0.027	0.830	−2.77	0.001	4.34
Euro Stoxx 50	6 836	−0.029	0.019	0.910	−2.08	0.097	3.41
Nikkei 225 (Japan)	6 545	−0.038	0.024	0.923	−1.48	0.317	1.97
Hang Seng (HK)	6 583	−0.004	0.024	0.967	−0.64	0.872	−0.17
BET (Romania)	6 670	0.109	0.028	1.173	2.98	0.001	−6.05
WIG20 (Poland)	6 685	0.018	0.018	1.009	0.24	0.777	2.25
BUX (Hungary)	6 666	0.037	0.024	1.013	0.26	0.427	−0.42
PX (Czechia)	6 682	0.052	0.032	1.041	0.59	0.350	−2.53
BIST 100 (Turkey)	6 692	0.005	0.024	1.033	0.64	0.949	−0.63
Bovespa (Brazil)	6 622	−0.028	0.025	0.928	−1.29	0.521	0.88
IPC (Mexico)	6 689	0.070	0.019	1.039	0.93	0.001	−3.27
Nifty 50 (India)	6 606	0.048	0.024	1.048	0.91	0.178	−5.46
Shanghai (China)	6 474	0.021	0.019	1.046	1.11	0.409	1.02
Bitcoin	4 384	−0.022	0.026	0.991	−0.19	0.835	2.68
Ethereum	3 913	−0.012	0.027	1.067	1.17	0.027	3.91
EUR/RON (reference rate)	5 352	0.170	0.036	1.157	2.19	0.000	−1.66

- T: number of daily returns; CD: Chow–Denning over $q = 2, 5, 10, 20$; EUR/RON from July 2005, crypto on 365 days per year

Back

Appendix: Runs Test on All Markets, Other VR Tests

- Runs z (table on the previous slide): Nifty 50 -5.46, IPC -3.27 (persistence in emerging markets); Ethereum 3.91, significant although $\hat{\rho}_1$ is not
- Choi (1999): an automatic VR test that chooses the horizon from the data (spectral estimate at frequency zero)
- Wild-bootstrap versions for small samples: Kim (2006); applications to Asian markets: Kim & Shamsuddin (2008)

Back

Appendix: Automatic Portmanteau Test of the MDH

- Escanciano & Lobato (2009): $AQ = T \sum_{j \leq \tilde{p}} \hat{\gamma}_j^2 / \hat{\kappa}_j \to \chi^2_1$, with $\hat{\gamma}_j = \frac{1}{T-j} \sum_t e_t e_{t-j}$ and $\hat{\kappa}_j = \frac{1}{T-j} \sum_t e_t^2 e_{t-j}^2$ (robust to conditional heteroskedasticity)
  - $\tilde{p} \leq 10$ maximises $Q_p = T \sum_{j \leq p} \hat{\gamma}_j^2 / \hat{\kappa}_j$ minus a penalty that switches from BIC ($p \ln T$, Bayesian Information Criterion) to AIC ($2p$, Akaike Information Criterion) when $\max_j \sqrt{T} |\hat{\rho}_j|$ exceeds $\sqrt{2.4 \ln T}$
- $e_t = r_t - \bar{r}$; $\hat{\gamma}_j$: autocovariance at lag j; $\hat{\kappa}_j$: its robust variance; a large AQ rejects the MDH
- AQ: S&P 500 13.14 ($p = 0.0003$), BET 14.92 ($p = 0.0001$), Bitcoin 0.75 ($p = 0.385$); the MDH is rejected at 5% for 5 of the 16 series

Back

Appendix: Predictive Regression, Data and Full Output

FRED: Federal Reserve Economic Data.

- Monthly log excess return (dividends included, minus the 3-month T-bill rate) regressed on the lagged dp_t
  - Data: Shiller (2026) (S&P Composite, to June 2026), extended with monthly averages of our daily S&P 500 closes (scaled over the last 12 common months) and the dividends of the SPDR S&P 500 ETF (SPY); 3-month US Treasury bill yield (FRED, series TB3MS; annual %), monthly log rate $\ln(1 + \text{TB3MS}/1200)$ lagged one month
- OLS: $\hat{\beta} = 0.267$, Newey–West SE 0.293 (12 lags), $t = 0.91$, OLS $t = 1.16$, $R^2 = 0.12\%$
- Corrected slope / unadjusted HAC SE $= -0.22$: only illustrative, it ignores the uncertainty of the correction
- Economic value Campbell & Thompson (2008) (number: slide The Dividend–Price Ratio: Results for the S&P 500): proportional rise in the expected return of the market timer
  - $\dfrac{R^2}{1-R^2} \cdot \dfrac{1+S^2}{S^2}$, S the monthly Sharpe ratio
- The Diebold & Mariano (1995) test: $t = -0.57$ for the same forecasts; the Clark–West test: $t = 0.42$

Back

Appendix: How Fast Is "Fast"?

CME: Chicago Mercantile Exchange.

- At high frequency, futures prices react to scheduled macro news within microseconds or less; Tak, Schlamp, Osterrieder & Pele (2026) measure this race at Eurex and CME
- Crypto events: Pele, Lin, Wang & Gaman (2025) study market responses to Ethereum development milestones
- Speed of adjustment is itself a measure of efficiency, and it depends on technology (Chapter 10, market microstructure)

Back

Appendix: Harold Hurst and the Nile



Harold Edwin Hurst (1880–1978)

Photo: Elliott & Fry (1953); public domain; Wikimedia Commons



Nilometer on Elephantine Island, Aswan

Photo: Olaf Tausch (2011); CC BY 3.0; Wikimedia Commons

- Hurst (1951) designed Nile reservoirs
- He found that the range of cumulative Nile flows grows faster than $\sqrt{n}$ with the horizon: years of high floods follow other years of high floods; hence the **rescaled range** statistic (R/S)

Back

Appendix: The Anis–Lloyd Expectation

- Anis & Lloyd (1976): exact $\mathbb{E}[R_n/s_n]$ under i.i.d. returns from the Normal distribution
  - $\mathbb{E}[R_n/s_n] = \dfrac{\Gamma((n-1)/2)}{\sqrt{\pi}\,\Gamma(n/2)} \sum_{i=1}^{n-1} \sqrt{\dfrac{n-i}{i}}$ (Γ: the gamma function)
- This $\mathbb{E}[R_n/s_n]$ is the reference subtracted in H_{AL} (R/S Analysis slide)
- Uncorrected estimates overstate memory by about 0.04–0.05 with daily data

Back

Appendix: Frequency, periodogram and local Whittle

- The spectral density allocates variance across frequencies
  - The periodogram is its raw sample analogue, built from squared Fourier amplitudes.
- Fourier components are sine and cosine waves; combining their amplitudes represents oscillations at different frequencies.
- Local Whittle estimates the long-memory parameter d from the spectral slope at low frequencies
  - $d = H - 0.5$; for example, $d = 0.1$ corresponds to $H = 0.6$.
  - Robinson (1995): the bandwidth m is the number of low frequencies used, not a number of lags.

Appendix: Long Memory with Inference, the Local Whittle Estimator

- $d = H - 0.5$; spectral density $f(\lambda) \sim G\lambda^{-2d}$ as $\lambda \to 0$ ($G > 0$ a constant; $d > 0$: the density explodes at low frequencies, long memory); fit only the first m Fourier frequencies $\lambda_j = 2\pi j/T$
  - $\hat{d} = \arg\min_d \left[\ln\left(\frac{1}{m} \sum_{j=1}^{m} \lambda_j^{2d} I(\lambda_j) \right) - \frac{2d}{m} \sum_{j=1}^{m} \ln \lambda_j \right]$, I: the periodogram
  - Robinson (1995): for $|d| < 1/2$, $m \to \infty$, $m/T \to 0$, a smooth f near 0 and a linear process with conditionally homoskedastic martingale-difference innovations, $\sqrt{m}(\hat{d} - d) \to N(0, 1/4)$ without Gaussianity: a confidence interval instead of a band
- Estimates with $m = T^{0.65}$ and the 95% CI $\hat{d} \pm 1.96/(2\sqrt{m})$
  - S&P 500 -0.013 $[-0.069, 0.043]$; Bitcoin 0.063 $[-0.002, 0.127]$
  - BET 0.118 $[0.062, 0.174]$; Ethereum 0.112 $[0.045, 0.179]$
- Bandwidth trade-off: a large m reaches frequencies where short-run autocorrelation (BET $\hat{\rho}_1 = 0.109$) biases $\hat{d}$ upward
  - with $m = T^{0.5}$ BET gives $\hat{d} = 0.135$ (SE 0.056): a sensitivity check supporting, not proving, that its $d > 0$ is not only short memory
  - GARCH returns also strain the conditions above

Back

Appendix: EMH versus AMH

	EMH	AMH
Investors	Rational, optimising	Bounded rationality, heuristics
Efficiency	An equilibrium property, tested as a restriction	A matter of degree that changes with conditions
Risk–return	Equilibrium risk premia, possibly time-varying	Changes with the population of investors
Anomalies	Omitted risk factors or data mining	Transient profit opportunities
Empirical tool	Tests of an equilibrium restriction	Tests of time and state variation (rolling windows)

Back

Appendix: Is Bitcoin Efficient? The Literature

- Urquhart (2016): Bitcoin inefficient over 2010–2016; some tests indicate efficiency in the later subperiod
  - His battery (Ljung–Box, runs, the rank test of Bartels (1982), automatic VR, BDS, R/S) includes BDS, which rejects under volatility clustering alone: not a test of RW3
- Nadarajah & Chu (2017): with a power transformation of returns, Bitcoin is weak-form efficient
- Kristoufek (2018): Bitcoin markets mostly inefficient in 2010–2017, with efficiency changing over time
- Tran & Leirvik (2020): time-varying efficiency across many crypto-assets, increasing over time
- Pele et al. (2023): crypto-assets have become a distinct asset class; their returns share the stylised facts of other markets
- Our own estimates, calendar year by calendar year (slide Crypto Efficiency Year by Year)
  - Bitcoin, lag-1 autocorrelation $\hat{\rho}_1$: -0.002 in 2011, -0.071 in 2012, 0.069 in 2013, all close to zero
  - Robust variance-ratio statistic $z^*(5)$: at most 1.71 in absolute value over all crypto-asset years, below the critical value 1.96
  - So in the early years we do not reject RW3, unlike Urquhart (2016)

Back

Appendix: Romania, Frontier then Emerging

BVB: Bursa de Valori București — the Bucharest Stock Exchange.

- FTSE Russell (2019): Romania reclassified from Frontier to Secondary Emerging, announced in September 2019, effective 21 September 2020
- Expected channels: more foreign investors, more analysts, more liquidity $\Rightarrow$ faster price discovery
- Earlier BVB evidence: crashes and log-periodic bubbles Pele & Mazurencu-Marinescu (2012); herding in an illiquid market Pele, Mazurencu-Marinescu & Nijkamp (2013)
- Did efficiency improve? Seminar C2 compares five years before and after, with WIG20 and BUX as a comparison group and a block bootstrap
  - Block bootstrap: resamples blocks of consecutive days to keep the dependence Künsch (1989); a before/after comparison is descriptive unless confounders are controlled

Back

Appendix: Reality Check, SPA and StepM explained

SPA: Superior Predictive Ability (test).

- Reality Check asks whether the best of all tried rules beats a benchmark, using a bootstrap distribution of their maximum performance. White (2000)
- SPA standardises performance by its uncertainty (studentisation) and reduces distortion from poor rules. Hansen (2005)
- StepM tests the strongest statistic first, then the remaining family using jointly resampled critical values
  - it controls FWER under its assumptions. Romano & Wolf (2005)
- Example: resample the same months for all 15 calendar rules, retaining common market shocks
  - separate resampling would destroy that dependence.

Back

Appendix: Dependent Tests, Reality Check, SPA and StepM

- Bonferroni and Holm ignore the correlation between tests (same markets, same days); Reality Check, SPA and StepM allow for it (explained on the previous slide)
- Our 15 calendar tests: StepM with a bootstrap of blocks of 3 calendar months, common to all markets, 999 replications
  - Critical p at the first step 0.0016 (Bonferroni: 0.0033); 2 rejections, BET and BUX TOM, as with Holm
  - Why stricter than Bonferroni? At nominal 5%, the bootstrap rejection rates of the HAC tests are 6.0% (day of week), 6.4% (TOM), 6.4% (January)
- Chordia, Goyal & Saretto (2020), from over two million strategies generated at random from real accounting data: thresholds $|t| > 3.8$ (time series) and $|t| > 3.4$ (cross-section); without them, about 45% of rejections are false
- Lowering false discoveries raises missed ones Harvey & Liu (2020)

Back

Appendix: Time-Series Momentum, All Markets

Market	Mean (%/month)	t_{HAC}	Sharpe TSMOM	Sharpe buy-and-hold	$\hat{\rho}_1$ monthly
S&P 500 (US)	0.61	2.22	0.48	0.52	0.03
Euro Stoxx 50	0.16	0.51	0.11	0.16	0.05
Nikkei 225 (Japan)	0.61	1.85	0.39	0.42	0.09
Hang Seng (HK)	0.18	0.53	0.10	0.19	0.04
BET (Romania)	1.40	2.68	0.65	0.75	0.19
WIG20 (Poland)	0.67	1.86	0.36	0.26	0.02
BUX (Hungary)	0.72	1.79	0.40	0.64	0.10
PX (Czechia)	0.38	0.99	0.24	0.45	0.10
BIST 100 (Turkey)	1.03	1.85	0.37	0.75	−0.08
Bovespa (Brazil)	0.28	0.65	0.14	0.51	0.11
IPC (Mexico)	0.44	1.41	0.32	0.63	0.03
Nifty 50 (India)	0.50	1.29	0.28	0.64	0.02
Shanghai (China)	0.59	1.32	0.30	0.23	0.08
Bitcoin	4.47	1.94	0.72	1.08	0.13
Ethereum	6.50	1.28	0.57	0.90	0.16

- 307 months for equities, fewer for crypto; Sharpe ratio $\sqrt{12}\,\bar{x}/s_x$ of price returns ($\bar{x}$, s_x: mean and SD of the monthly strategy returns), no risk-free rate subtracted; Moskowitz, Ooi & Pedersen (2012) use futures excess returns scaled to 40% ex-ante volatility

Back

Appendix: Robert Shiller and Richard Thaler



Robert J. Shiller, Nobel Prize 2013

Photo: Bengt Nyman (2013); CC BY 2.0; Wikimedia Commons



Richard H. Thaler, Nobel Prize 2017

Photo: Bengt Nyman (2017); CC BY 2.0; Wikimedia Commons

- Excess volatility Shiller (1981) and overreaction De Bondt & Thaler (1985): Psychology in Prices slide

Back

Appendix: Case Study, Research Question

- Martin & Nagel (2022): investors price N assets whose cash flows depend on J characteristics with unknown coefficients; what does a rejected no-predictability test mean when J is comparable to N?
  - the coefficients must be learned from past data: a high-dimensional prediction problem, as in machine learning
- Semi-strong efficiency: prices "fully reflect" public information Fama (1970); two readings
  - investors know the parameters of the cash-flow model (rational expectations: the null of most efficiency tests)
  - investors update their beliefs optimally (Bayes) with the data available in real time
- Why it matters: hundreds of in-sample return predictors, the "factor zoo" Harvey, Liu & Zhu (2016); text, satellite images and social media keep increasing J
  - does every significant predictor need a risk-based or a behavioural explanation?
- Link to this chapter: out-of-sample predictability Campbell & Thompson (2008) and the decay of anomalies after publication McLean & Pontiff (2016)

Back

Appendix: Case Study, Model and Method (1/2)

- **Model** (Martin & Nagel (2022), Sec. 2)
  - cash-flow growth $\Delta y_t = Xg + e_t$, with $e_t \sim N(0, I)$ and a zero interest rate
  - Δy_t: the N cash-flow growth rates; X: the $N \times J$ matrix of characteristics; g: the J unknown coefficients; I: identity matrix
- **Learning** (Prop. 1–2)
  - prior $g \sim N(0, (\theta/J)\, I)$, the true distribution from which g is drawn; θ: signal strength
  - realised return = underreaction caused by shrinkage + estimation noise + unpredictable shock
- **Asymptotics and simulations** (Sec. 4)
  - $N, J \to \infty$ with $J/N \to \psi > 0$
  - simulations with $N = 1000$, entries of X drawn from the standard Normal distribution and $\theta = 1$

Back

Appendix: Case Study, Model and Method (2/2)

- **CRSP**: Center for Research in Security Prices.
- **NYSE**: New York Stock Exchange.
- **Tests**
  - in sample: a cross-sectional OLS of r_{t+1} on X (requires $J < N$) with a joint χ^2 test
  - out of sample: a portfolio with weights $\frac{1}{N} X h_{s+1}$, from coefficients estimated in other periods
  - $r_{\text{OOS},t+1} = \frac{1}{N} h'_{s+1} X' r_{t+1}$, $s \neq t$ (eq. 19); r_{IS}: the same with $s = t$
  - h_{s+1}: the estimated coefficient vector; r_{t+1}: the N returns
- **Empirical design** (Sec. 7)
  - US stocks in CRSP, without those below the 20th NYSE size percentile or priced below \$1
  - predictors: simple and squared monthly returns at lags $t-2$ to $t-120$ (238 predictors), demeaned and standardised across stocks every month
  - full-sample ridge, 1971 – June 2019, penalty chosen by leave-one-year-out cross-validation
  - rolling 20-year OLS gives r_{OOS} from January 1956

Back

Appendix: Case Study, Formal Results

- In sample (Prop. 4 of Martin & Nagel (2022), main slide): for $\psi > 0.4$ ($\theta = 1$, one period of learning) the probability of no rejection is of the order $e^{-0.015N}$ (Fig. 2)
- Simulations ($N = 1000$, four periods of learning): a test of nominal size 5% rejects with probability above 90% once $J > N/2$ (Fig. 3b)
  - not a size distortion: the in-sample predictable component really is in the prices
- Single-predictor tests with k-FWER control (probability of k or more false rejections at most 5%): for $k = 20$ and $J/N = 0.05$, the probability of at least k rejections is 25% instead of 5%; about 80% for J/N close to 1 (Fig. 4)
  - multiple-testing corrections do not help: the problem is the null hypothesis, not the test
- Out of sample: $\mathbb{E}[r_{\mathrm{OOS},t+1}] = 0$ holds forward and backward in time (Prop. 6); positive only if investors shrink more than the data warrant (Fig. 6)
- US stocks: the ridge coefficients recover momentum, long-run reversal and momentum seasonality (Fig. 7); $\eta_S \approx 0.006\%$ per month (Fig. 8)

Back

Appendix: Case Study, In Sample vs Out of Sample in Detail

- $r_{\text{IS}} \geq 0$ in every 20-year window by construction: $r_{\text{IS}} = b' G^{-1} b/240$, with $G = \sum_s X_s' X_s / N_s$ and $b = \sum_s X_s' r_s / N_s$ over the 240 months s of the window
  - X_s: the $N_s \times 238$ predictor matrix of month s; r_s: the N_s portfolio returns; $G^{-1}b$: the pooled OLS coefficients
  - mean in-sample R^2 of 5.4%, against about 1% for individual stocks in Martin & Nagel (2022): 238 predictors against only 40–49 portfolios a month, pooled over 240 months, inflate the in-sample fit
- Out of sample (from July 1956), until June 1986: mean $r_{\text{OOS}} = 1.98 \times 10^{-5}$, 4.16 standard errors above 0; afterwards 0.36×10^{-5}, only 0.62 standard errors
  - the 10-year average is last above two standard errors in the window ending in December 1987; in the windows ending in 2014–2019 its largest ratio to the standard error is 1.25, as for stocks
- Interpretation suggested by Martin & Nagel (2022): the gap is consistent with investor learning, but does not establish it or rule out risk premia; early out-of-sample gains are also consistent with excess shrinkage before the computer era (Fig. 6)

Back

Appendix: Case Study, Replication, Limitations and Extensions

- **Replicate with the course data** (stock-level CRSP files are not part of them)
  - the 49 value-weighted industry portfolios of the main replication slides; same predictors, standardisation and windows
  - the size and price filters do not apply to portfolios
- **Limits stated by the authors**
  - lagged returns are equilibrium outcomes, not the exogenous X of the model: the application illustrates the wedge, it does not test the model
  - a t-test on the monthly r_{OOS} is not a valid test of $\mathbb{E}[r_{OOS}] = 0$: an expectation that also averages over the prior distribution of g
  - the backward result may rely on i.i.d. cash-flow growth, homogeneous investors and constant characteristics
- **Open extensions**
  - time series, as the authors suggest: does the dividend–price regression of this chapter show the same in-sample/out-of-sample wedge?
  - nonlinear models (trees, neural networks Gu, Kelly & Xiu (2020)) add functional flexibility without necessarily changing J: does the wedge grow?
  - AMH: is early out-of-sample predictability the excess shrinkage of investors with less computing power?

Back

Appendix: The Discovery Loop, Reusable Prompts (1/2)

- **Literature**
  - `List peer-reviewed studies since 2020 on how ETF introductions change the efficiency or price discovery of the underlying. For each: DOI, design (event study, difference-in-differences (DiD), structural break), efficiency measure. Flag any entry you are unsure exists.`
  - Resolve every DOI on Crossref; the title and authors must match
- **Hypothesis**
  - `Give three mechanisms by which spot ETFs could change Bitcoin's serial dependence (arbitrage, investor mix, trading hours). For each, the sign of the change in rho1 and VR(q), and a variable that separates it from the others.`
- **Data and code**
  - `Read the daily Bitcoin closing prices from the course data; log returns; rho1 with a heteroskedasticity-robust standard error and the Lo-MacKinlay VR(5) with z*, in two windows of equal length around 11 January 2024.`

Back

Appendix: The Discovery Loop, Reusable Prompts (2/2)

- **Robustness**
  - `Write a pre-registered robustness plan: window lengths, weekday-only closes, VR(q) for q = 2 to 20 with Chow-Denning, placebo dates, and Ethereum as a control until its own spot ETFs (July 2024).`
- **Critique**
  - `Act as a hostile referee: if Bitcoin was already efficient before 2024, what could this test ever detect? Compute the minimum detectable change in rho1 and name a better outcome variable.`
- **Report**
  - Estimates, robust standard errors, the placebo distribution, every test run (with a multiple-testing correction)
  - Prompts that shaped a decision go into `AI_USE.md`

Back

Appendix: What the Researcher Must Check

- **i.i.d. standard errors**: $1/\sqrt{T}$ rejects too often under volatility clustering
  - test: robust SE and z^*, as in the weak-form tests of this chapter
- **Unequal windows**: different power before and after; test: equal lengths, fixed in advance
- **Event date chosen from the data**
  - test: the launch date (first trading day) only, plus a placebo distribution at dates without an event
- **Confounding**: before/after is not causal; test: a control asset (Ethereum before July 2024) in a difference-in-differences
  - causal only under the conditions of the difference-in-differences slide; the common post-period (Jan–Jul 2024) is short
- **Many statistics**: ρ_1, VR(2–20), runs; test: Chow–Denning or Holm, not the smallest single p-value
- **Low power**: report the MDE of the Estimation and Power slide

Back

Appendix: Further Study on Quantinar

Advanced courses on quantinar.com

- Efficiency of cryptocurrency markets - A GMM-based analysis
- Generalized Method of Moments
- Market Responses to ETH Development Milestones

Code, data and textbooks

- All charts and tables: github.com/danpele/MFM/tree/main/Quantlets/Ch_02
- Textbooks: Campbell, Lo & MacKinlay (1997), Ch. 2; Franke, Härdle & Hafner (2019)

Back

Acronyms used in this material (1/2)

- **AIC**: Akaike Information Criterion
- **AMH**: Adaptive Markets Hypothesis
- **AR(p)**: AutoRegressive model of order p
- **AR**: Abnormal Return, in event studies.
- **BDS**: Brock–Dechert–Scheinkman (test)
- **BET**: Bucharest Exchange Trading (index)
- **BIC**: Bayesian Information Criterion
- **BIST**: Borsa İstanbul — *the Istanbul Stock Exchange (BIST 100 index)*
- **BNR**: Banca Națională a României — *the National Bank of Romania*
- **BUX**: Budapesti Értéktőzsde index — *the Budapest Stock Exchange index*
- **BVB**: Bursa de Valori București — *the Bucharest Stock Exchange*
- **CAPM**: Capital Asset Pricing Model
- **CAR**: Cumulative Abnormal Return
- **CD**: Chow–Denning (multiple variance-ratio test)
- **CI**: Confidence Interval
- **CME**: Chicago Mercantile Exchange
- **CRSP**: Center for Research in Security Prices
- **DFA**: Detrended Fluctuation Analysis
- **EMH**: Efficient Market Hypothesis
- **ETF**: Exchange-Traded Fund
- **fGn**: fractional Gaussian noise
- **FRED**: Federal Reserve Economic Data
- **FTSE**: Financial Times Stock Exchange (FTSE Russell, index provider)
- **FWER**: Family-Wise Error Rate
- **FX**: Foreign eXchange

Acronyms used in this material (2/2)

- **GARCH**: Generalized AutoRegressive Conditional Heteroskedasticity
- **GMM**: Generalized Method of Moments
- **HAC**: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- **HK**: Hong Kong
- **IPC**: Índice de Precios y Cotizaciones — *the Mexican Stock Exchange index*
- **MDE**: Minimum Detectable Effect
- **MDH**: Martingale Difference Hypothesis
- **MIT**: Massachusetts Institute of Technology
- **ML**: Machine Learning
- **MSPE**: Mean Squared Prediction Error
- **NYSE**: New York Stock Exchange
- **OLS**: Ordinary Least Squares
- **PX**: Prague Stock Exchange index (index Burzy cenných papírů Praha)
- **RW1**: Random Walk 1 (i.i.d. increments)
- **RW2**: Random Walk 2 (independent increments)
- **RW3**: Random Walk 3 (uncorrelated increments)
- **SD**: Standard Deviation
- **SE**: Standard Error
- **SPA**: Superior Predictive Ability (test)
- **TOM**: Turn Of the Month
- **TSMOM**: Time-Series MOMentum
- **US**: United States
- **VR**: Variance Ratio
- **WIG20**: Warszawski Indeks Giełdowy 20 — *the Warsaw Stock Exchange index of the 20 largest companies*

References (1/7)

- Andrews, D.W.K. (1991). Heteroskedasticity and autocorrelation consistent covariance matrix estimation. *Econometrica*, 59(3), 817–858.
- Anis, A.A., Lloyd, E.H. (1976). The expected value of the adjusted rescaled Hurst range of independent normal summands. *Biometrika*, 63(1), 111–116.
- Ariel, R.A. (1987). A monthly effect in stock returns. *Journal of Financial Economics*, 18(1), 161–174.
- Bachelier, L. (1900). Théorie de la spéculation. *Annales scientifiques de l'École Normale Supérieure*, 17, 21–86.
- Bartels, R. (1982). The rank version of von Neumann's ratio test for randomness. *Journal of the American Statistical Association*, 77(377), 40–46.
- Blanchard, O.J., Watson, M.W. (1982). Bubbles, rational expectations and financial markets. In P. Wachtel (ed.), *Crises in the Economic and International Financial System*, Lexington Books, 295–315.
- Boehmer, E., Musumeci, J., Poulsen, A.B. (1991). Event-study methodology under conditions of event-induced variance. *Journal of Financial Economics*, 30(2), 253–272.
- Brock, W.A., Dechert, W.D., Scheinkman, J.A., LeBaron, B. (1996). A test for independence based on the correlation dimension. *Econometric Reviews*, 15(3), 197–235.
- Campbell, J.Y., Lo, A.W., MacKinlay, A.C. (1997). The Econometrics of Financial Markets. Princeton University Press.
- Campbell, J.Y., Thompson, S.B. (2008). Predicting excess stock returns out of sample: can anything beat the historical average? *Review of Financial Studies*, 21(4), 1509–1531.
- Campbell, J.Y., Yogo, M. (2006). Efficient tests of stock return predictability. *Journal of Financial Economics*, 81(1), 27–60.
- Choi, I. (1999). Testing the random walk hypothesis for real exchange rates. *Journal of Applied Econometrics*, 14(3), 293–308.
- Chordia, T., Goyal, A., Saretto, A. (2020). Anomalies and false rejections. *Review of Financial Studies*, 33(5), 2134–2179.

References (2/7)

- Chow, K.V., Denning, K.C. (1993). A simple multiple variance ratio test. *Journal of Econometrics*, 58(3), 385–401.
- Clark, T.E., West, K.D. (2007). Approximately normal tests for equal predictive accuracy in nested models. *Journal of Econometrics*, 138(1), 291–311.
- De Bondt, W.F.M., Thaler, R. (1985). Does the stock market overreact? *Journal of Finance*, 40(3), 793–805.
- De Long, J.B., Shleifer, A., Summers, L.H., Waldmann, R.J. (1990). Noise trader risk in financial markets. *Journal of Political Economy*, 98(4), 703–738.
- Diebold, F.X., Mariano, R.S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263.
- Escanciano, J.C., Lobato, I.N. (2009). An automatic Portmanteau test for serial correlation. *Journal of Econometrics*, 151(2), 140–149.
- Escanciano, J.C., Velasco, C. (2006). Generalized spectral tests for the martingale difference hypothesis. *Journal of Econometrics*, 134(1), 151–185.
- Fama, E.F. (1970). Efficient capital markets: a review of theory and empirical work. *Journal of Finance*, 25(2), 383–417.
- Fama, E.F. (1991). Efficient capital markets: II. *Journal of Finance*, 46(5), 1575–1617.
- Fama, E.F., Fisher, L., Jensen, M.C., Roll, R. (1969). The adjustment of stock prices to new information. *International Economic Review*, 10(1), 1–21.
- Franke, J., Härdle, W.K., Hafner, C.M. (2019). Statistics of Financial Markets: An Introduction, 5th ed. Springer.
- French, K.R. (1980). Stock returns and the weekend effect. *Journal of Financial Economics*, 8(1), 55–69.
- French, K.R. (2026). Kenneth R. French Data Library. Tuck School of Business, Dartmouth College.

References (3/7)

- FTSE Russell (2019). FTSE Russell announces results of country classification review (press release, 26 September 2019).
- Goyal, A., Welch, I., Zafirov, A. (2024). A comprehensive 2022 look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 37(11), 3490–3557.
- Grossman, S.J., Stiglitz, J.E. (1980). On the impossibility of informationally efficient markets. *American Economic Review*, 70(3), 393–408.
- Gu, S., Kelly, B., Xiu, D. (2020). Empirical asset pricing via machine learning. *Review of Financial Studies*, 33(5), 2223–2273.
- Hansen, P.R. (2005). A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4), 365–380.
- Harvey, C.R., Liu, Y. (2020). False (and missed) discoveries in financial economics. *Journal of Finance*, 75(5), 2503–2553.
- Harvey, C.R., Liu, Y., Zhu, H. (2016). . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1), 5–68.
- Heston, S.L., Sadka, R. (2008). Seasonality in the cross-section of stock returns. *Journal of Financial Economics*, 87(2), 418–445.
- Hurst, H.E. (1951). Long-term storage capacity of reservoirs. *Transactions of the American Society of Civil Engineers*, 116(1), 770–799.
- Jegadeesh, N., Titman, S. (1993). Returns to buying winners and selling losers: implications for stock market efficiency. *Journal of Finance*, 48(1), 65–91.
- Kahneman, D., Tversky, A. (1979). Prospect theory: an analysis of decision under risk. *Econometrica*, 47(2), 263–291.
- Kendall, M.G. (1954). Note on bias in the estimation of autocorrelation. *Biometrika*, 41(3–4), 403–404.
- Kia, K., Liu, B., Li, Q., Song, V., Xu, K. (2026). Price discovery in Bitcoin ETF market. *Financial Review*, 61(2), 435–449.

References (4/7)

- Kim, J.H. (2006). Wild bootstrapping variance ratio tests. *Economics Letters*, 92(1), 38–43.
- Kim, J.H., Shamsuddin, A. (2008). Are Asian stock markets efficient? Evidence from new multiple variance ratio tests. *Journal of Empirical Finance*, 15(3), 518–532.
- Kim, J.H., Shamsuddin, A., Lim, K.-P. (2011). Stock return predictability and the adaptive markets hypothesis: evidence from century-long U.S. data. *Journal of Empirical Finance*, 18(5), 868–879.
- Kolari, J.W., Pynnönen, S. (2010). Event study testing with cross-sectional correlation of abnormal returns. *Review of Financial Studies*, 23(11), 3996–4025.
- Kristoufek, L. (2018). On Bitcoin markets (in)efficiency and its evolution. *Physica A*, 503, 257–262.
- Künsch, H.R. (1989). The jackknife and the bootstrap for general stationary observations. *Annals of Statistics*, 17(3), 1217–1241.
- Lakonishok, J., Smidt, S. (1988). Are seasonal anomalies real? A ninety-year perspective. *Review of Financial Studies*, 1(4), 403–425.
- Lim, K.-P., Brooks, R. (2011). The evolution of stock market efficiency over time: a survey of the empirical literature. *Journal of Economic Surveys*, 25(1), 69–108.
- Ljung, G.M., Box, G.E.P. (1978). On a measure of lack of fit in time series models. *Biometrika*, 65(2), 297–303.
- Lo, A.W. (1991). Long-term memory in stock market prices. *Econometrica*, 59(5), 1279–1313.
- Lo, A.W. (2004). The adaptive markets hypothesis. *Journal of Portfolio Management*, 30(5), 15–29.
- Lo, A.W. (2017). Adaptive Markets: Financial Evolution at the Speed of Thought. Princeton University Press.
- Lo, A.W., MacKinlay, A.C. (1988). Stock market prices do not follow random walks: evidence from a simple specification test. *Review of Financial Studies*, 1(1), 41–66.

References (5/7)

- Lobato, I., Nankervis, J.C., Savin, N.E. (2001). Testing for autocorrelation using a modified Box–Pierce Q test. *International Economic Review*, 42(1), 187–205.
- MacKinlay, A.C. (1997). Event studies in economics and finance. *Journal of Economic Literature*, 35(1), 13–39.
- Malkiel, B.G. (2003). The efficient market hypothesis and its critics. *Journal of Economic Perspectives*, 17(1), 59–82.
- Martin, I.W.R., Nagel, S. (2022). Market efficiency in the age of big data. *Journal of Financial Economics*, 145(1), 154–177.
- McLean, R.D., Pontiff, J. (2016). Does academic research destroy stock return predictability? *Journal of Finance*, 71(1), 5–32.
- Moskowitz, T.J., Ooi, Y.H., Pedersen, L.H. (2012). Time series momentum. *Journal of Financial Economics*, 104(2), 228–250.
- Nadarajah, S., Chu, J. (2017). On the inefficiency of Bitcoin. *Economics Letters*, 150, 6–9.
- Newey, W.K., West, K.D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.
- Pele, D.T., Lin, M.-B., Wang, R., Gaman, S. (2025). Market responses to ETH development milestones. IDA Institute Digital Assets, Bucharest University of Economic Studies.
- Pele, D.T., Mazurencu-Marinescu, M. (2012). Modelling stock market crashes: the case of Bucharest Stock Exchange. *Procedia – Social and Behavioral Sciences*, 58, 533–542.
- Pele, D.T., Mazurencu-Marinescu, M., Nijkamp, P. (2013). Herding behaviour, bubbles and log periodic power laws in illiquid stock markets: a case study on the Bucharest Stock Exchange. SSRN Working Paper.
- Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets? *European Journal of Finance*, 29(10), 1064–1105.
- Peng, C.-K., Buldyrev, S.V., Havlin, S., Simons, M., Stanley, H.E., Goldberger, A.L. (1994). Mosaic organization of DNA nucleotides. *Physical Review E*, 49(2), 1685–1689.

References (6/7)

- Roberts, H.V. (1959). Stock-market "patterns" and financial analysis: methodological suggestions. *Journal of Finance*, 14(1), 1–10.
- Robinson, P.M. (1995). Gaussian semiparametric estimation of long range dependence. *Annals of Statistics*, 23(5), 1630–1661.
- Romano, J.P., Wolf, M. (2005). Stepwise multiple testing as formalized data snooping. *Econometrica*, 73(4), 1237–1282.
- Rozeff, M.S., Kinney, W.R. (1976). Capital market seasonality: the case of stock returns. *Journal of Financial Economics*, 3(4), 379–402.
- Samuelson, P.A. (1965). Proof that properly anticipated prices fluctuate randomly. *Industrial Management Review*, 6(2), 41–49; reprinted 2015, World Scientific.
- Shiller, R.J. (1981). Do stock prices move too much to be justified by subsequent changes in dividends? *American Economic Review*, 71(3), 421–436.
- Shiller, R.J. (2026). U.S. stock markets 1871–present and CAPE ratio (monthly data, accessed September 2026).
- Shleifer, A., Vishny, R.W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1), 35–55.
- Stambaugh, R.F. (1999). Predictive regressions. *Journal of Financial Economics*, 54(3), 375–421.
- Sullivan, R., Timmermann, A., White, H. (1999). Data-snooping, technical trading rule performance, and the bootstrap. *Journal of Finance*, 54(5), 1647–1691.
- Summers, L.H. (1986). Does the stock market rationally reflect fundamental values? *Journal of Finance*, 41(3), 591–601.
- Tak, R., Pele, D.T. (2025). LLM-driven stock prediction: capturing market trends with LLaMA. *Proceedings of the International Conference on Business Excellence*, 19(1), 529–540.
- Tak, R., Schlamp, S., Osterrieder, J., Pele, D.T. (2026). Racing the news: nanosecond market efficiency at Eurex and CME. SSRN Working Paper.

References (7/7)

- Tirole, J. (1985). Asset bubbles and overlapping generations. *Econometrica*, 53(6), 1499–1528.
- Tran, V.L., Leirvik, T. (2020). Efficiency in the markets of crypto-currencies. *Finance Research Letters*, 35, 101382.
- Urquhart, A. (2016). The inefficiency of Bitcoin. *Economics Letters*, 148, 80–82.
- Urquhart, A., McGroarty, F. (2016). Are stock markets really efficient? Evidence of the adaptive market hypothesis. *International Review of Financial Analysis*, 47, 39–49.
- Wald, A., Wolfowitz, J. (1940). On a test whether two samples are from the same population. *Annals of Mathematical Statistics*, 11(2), 147–162.
- Welch, I., Goyal, A. (2008). A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 21(4), 1455–1508.
- White, H. (2000). A reality check for data snooping. *Econometrica*, 68(5), 1097–1126.