



Modelling Financial Markets

Chapter 9: Realised Volatility



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model)
- ▣ BV: Bipower Variation
- ▣ BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- ▣ CI: Confidence Interval
- ▣ CLT: Central Limit Theorem
- ▣ DJIA: Dow Jones Industrial Average
- ▣ DM: Diebold–Mariano (test)
- ▣ ETF: Exchange-Traded Fund
- ▣ EWMA: Exponentially Weighted Moving Average
- ▣ FDR: False Discovery Rate
- ▣ FOMC: Federal Open Market Committee
- ▣ FRED: Federal Reserve Economic Data
- ▣ FWER: Family-Wise Error Rate
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GMM: Generalized Method of Moments
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HAR: Heterogeneous AutoRegressive (model)
- ▣ HAR-CJ: HAR model with Continuous and Jump components
- ▣ HAR-RV: Heterogeneous AutoRegressive model of Realised Volatility
- ▣ HARQ: HAR model with Realised Quarticity
- ▣ HEAVY: High-frEQUENCY-bAsed VolatilitY (models)
- ▣ IQ: Integrated Quarticity
- ▣ LLM: Large Language Model
- ▣ MAE: Mean Absolute Error
- ▣ MCS: Model Confidence Set

Acronyms used in this material (2/2)

- ▣ ML: Machine Learning
- ▣ MSE: Mean Squared Error
- ▣ MZ: Mincer–Zarnowitz (regression)
- ▣ NYSE: New York Stock Exchange
- ▣ OLS: Ordinary Least Squares
- ▣ QLIKE: Quasi-LIKelihood (loss function)
- ▣ RF: Random Forest
- ▣ RK: Realised Kernel
- ▣ RQ: Realised Quarticity
- ▣ RV: Realised Variance
- ▣ SE: Standard Error
- ▣ SHAR: Semivariance HAR model
- ▣ TRV: Truncated Realised Variance
- ▣ TSRV: Two-Scale Realised Variance
- ▣ US: United States
- ▣ UTC: Coordinated Universal Time (Temps Universel Coordonné)
- ▣ VIX: CBOE Volatility IndeX
- ▣ VRP: Variance Risk Premium

Outline

Motivation

Realised Variance

Microstructure Noise

Jumps

Stylised Facts of Realised Volatility

The HAR-RV Model

Forecasting: HAR vs GARCH

Rough Volatility

Case Study: Machine Learning vs HAR

AI for Scientific Discovery

Summary

References

Volatility Is Latent: Can We Measure It?

The problem

- Volatility drives risk (Chapters 7–8), option prices (Chapter 12) and portfolio weights (Chapter 4)
- But it is never observed: one daily return is a single draw from a distribution whose dispersion we want to estimate
- GARCH (Chapter 5) **infers** today's variance from past daily returns through a model

The idea of this chapter

- Inside one day there are hundreds of price changes: add up their squares
- The sum **measures** the day's variance almost without a model: realised variance [Andersen & Bollerslev \(1998\)](#)
- Measured volatility can then be modelled and forecast like any time series
- Early idea: [Merton \(1980\)](#) showed that variance, unlike the mean, is estimated better with more frequent data

Realised Volatility: From Sceptics to Standard



Torben G. Andersen, 2012

Photo: Soralum (2012); CC BY-SA 4.0; Wikimedia

[Commons](#)

- Measure first, then model: Andersen, Bollerslev, Diebold & Labys (2001), Andersen, Bollerslev, Diebold & Ebens (2001) study the distribution of realised volatility; Andersen, Bollerslev, Diebold & Labys (2003) forecast it directly
- Why the sceptics were wrong (Section 7) and textbook treatment: Franke, Härdle & Hafner (2019)

Data Used in This Chapter

Series	Frequency	Period	Days
SPY (ETF on the S&P 500)	5-minute bars, 09:30–16:00 New York time	2020–2026	1 489
Bitcoin	5-minute bars, 24-hour UTC days	2024–2026	659
SPY, Bitcoin	daily closes	2000/2014–2026	for GARCH
VIX	daily close	2020–2026	implied volatility

- Up to 18 September 2026; returns $r_i = 100(\ln P_i - \ln P_{i-1})$ in %; a bar is labelled by its start time
- SPY: 78 returns per full day, 42 on half-days; bars off the 5-minute grid removed; missing bars inside a day carry the last price
- Bitcoin: 288 returns per day; days with fewer than 95% of the bars dropped (88 of 747 calendar days, mostly a data gap in October 2025)

A Model for Intraday Prices

- Notation: day t has M intervals of length $\Delta = 1/M$, returns $r_{t,i}$, daily return $r_t = \sum_i r_{t,i}$
 - ▶ Variances in %²; annualised volatility $\sqrt{252 \times \text{variance}}$ (SPY), $\sqrt{365 \times \text{variance}}$ (Bitcoin); the SPY overnight return (close to next open) is outside the bars
- Log price in continuous time: $dp(s) = \mu(s) ds + \sigma(s) dW(s)$, with W a Brownian motion (Chapter 11)
 - ▶ $\sigma(s)$ is the **spot volatility**: it changes within the day and from day to day
- **Integrated variance** of day t : the quantity we want to measure
 - ▶ $IV_t = \int_{t-1}^t \sigma^2(s) ds$
- If σ is independent of W (no leverage) and the drift is zero (or fixed by the path of σ), then given that path the daily return is Normal: $r_t \mid \sigma \sim N(\int \mu, IV_t)$
 - ▶ With leverage (Section 6) this conditional normality fails, but the limit theory of RV below still holds
- IV_t is the natural daily variance; GARCH's σ_t^2 is a model-based forecast of it from yesterday's information

Definition: Realised Variance

- **Realised variance:** the sum of squared intraday returns


- ▶
$$RV_t = \sum_{i=1}^M r_{t,i}^2, \text{ realised volatility} = \sqrt{RV_t}$$

- As $M \rightarrow \infty$, $RV_t \rightarrow IV_t$: the drift vanishes and the sum of squares converges to the **quadratic variation** Barndorff-Nielsen & Shephard (2002)

- ▶ The mean of SPY 5-minute returns is below 1% of their standard deviation: it can be ignored

- No parameters to estimate: RV_t is a **model-free** measurement of one day's variance
- The squared daily return r_t^2 is the special case $M = 1$: unbiased (with zero drift) but extremely noisy

One Day, 78 Returns: 7 April 2025




□ Left: SPY on 5-minute bars on the most volatile day of the sample; right: running sum of squared returns

Reading the Chart: 7 April 2025

- SPY opened -3.27% below the previous close, after tariff announcements; 78 five-minute returns followed
 - ▶ $RV = 45.1\%^2$: annualised realised volatility $\sqrt{252 \times 45.1} = 107\%$
- One 5-minute return of $+4.28\%$ (interval ending 10:15) contributed 41% of the day's realised variance
- Open-to-close the day ended $+3.09\%$: the daily return alone would call it a normal day, realised variance calls it one of the most turbulent days since 2020

Daily Realised Volatility: SPY and Bitcoin

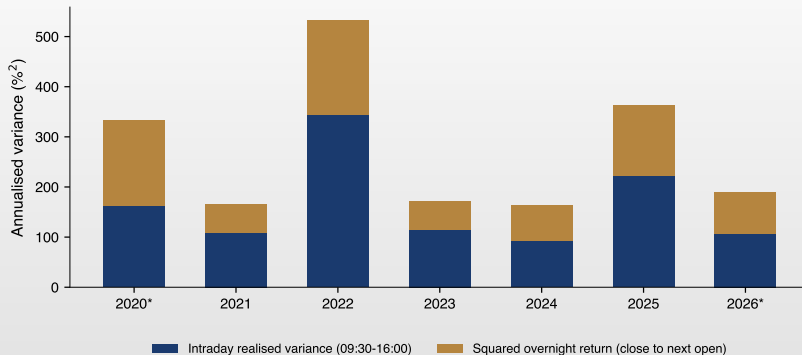


□ Annualised, from 5-minute returns; SPY: regular US session; Bitcoin: 24 hours; the gap in October 2025 is missing data

Reading the Realised Volatility Series

- SPY: volatility from mean RV $\sqrt{252 \overline{RV}} = 12.9\%$ p.a., median day 9.7%, range 2.4% to 106.7%
 - ▶ Volatility computed from mean variance exceeds the median: a few very volatile days raise the root-mean-square of daily volatilities
- Bitcoin: $\sqrt{365 \overline{RV}} = 45.1\%$ p.a., median 37.8%, maximum 136.1% on 2026-02-05
- Both series show **clustering**, as with GARCH (Chapter 5), but every point is a measurement, not a model output

Where Is the Risk: Day or Night?



- SPY, mean daily variance by year, annualised: intraday realised variance plus squared overnight return; 2020 and 2026 are partial years

Reading the Overnight Decomposition

- Overnight returns carry 38% of SPY's daily variance, although the market is closed for 17.5 hours
 - ▶ News (earnings, macro data, Asian and European sessions) accumulates and is priced at the open
- Annualised: close-to-close 16.6%; open-to-close 13.6%; overnight 10.1%
- Largest overnight move: -4.08% on 2024-08-05
- For daily risk we use **total** variance: $RV_t + r_{\text{overnight},t}^2$ (annualised mean 16.4%)
 - ▶ Alternatives: scale intraday RV up, or weight the two parts optimally [Hansen & Lunde \(2005\)](#)

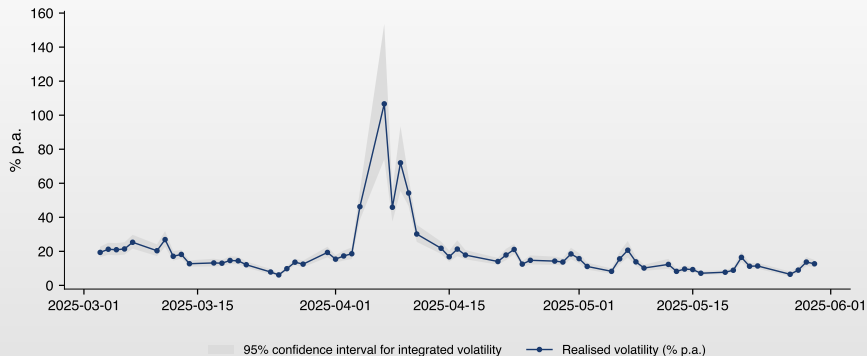
The Setting: Itô Semimartingales and Stable Convergence

- Itô semimartingale: $dp = \mu ds + \sigma dW + dJ$, with μ, σ adapted and σ càdlàg; σ may depend on W (leverage)
 - ▶ Without jumps: $\sqrt{M}(RV_t - IV_t) \xrightarrow{\mathcal{L}-s} MN(0, 2IQ_t)$, a **mixed** Normal distribution whose variance $2IQ_t$, with integrated quarticity IQ , $IQ_t = \int \sigma^4(s) ds$, is itself random [Barndorff-Nielsen & Shephard \(2002\)](#)
- **Stable** convergence ($\mathcal{L}$ -s): the limit holds jointly with every variable of the path's σ -field [Mykland & Zhang \(2009\)](#)
 - ▶ Hence $(RV_t - IV_t)/\sqrt{2RQ_t/M} \rightarrow N(0, 1)$: studentising by $\sqrt{2RQ_t/M}$, with $RQ_t \rightarrow IQ_t$ random, is legitimate only under stable convergence
- Finite samples: the t -statistic for $\ln RV_t$ has better coverage than the one in levels [Barndorff-Nielsen & Shephard \(2005\)](#); we use it on the next slide

How Precise Is Realised Variance?

- Feasible central limit theorem for RV [Barndorff-Nielsen & Shephard \(2002\)](#):
 $\sqrt{M}(\text{RV}_t - \text{IV}_t)/\sqrt{2\text{RQ}_t} \rightarrow N(0, 1)$
 - ▶ IQ_t is estimated by realised quarticity $\text{RQ}_t = \frac{M}{3} \sum_i r_{t,i}^4$
- Standard error: $\text{se}(\text{RV}_t) = \sqrt{\frac{2}{3} \sum_i r_{t,i}^4}$
- 95% CI on the log scale (always positive): $\text{RV}_t \exp(\pm 1.96 \text{se}(\text{RV}_t)/\text{RV}_t)$
 - ▶ With constant volatility, $\text{se}/\text{RV} \approx \sqrt{2/M}$: 16% for $M = 78$
- RV is a measurement **with error**: it matters when RV is the target of a forecast (Section 6)

Confidence Intervals for Daily Volatility

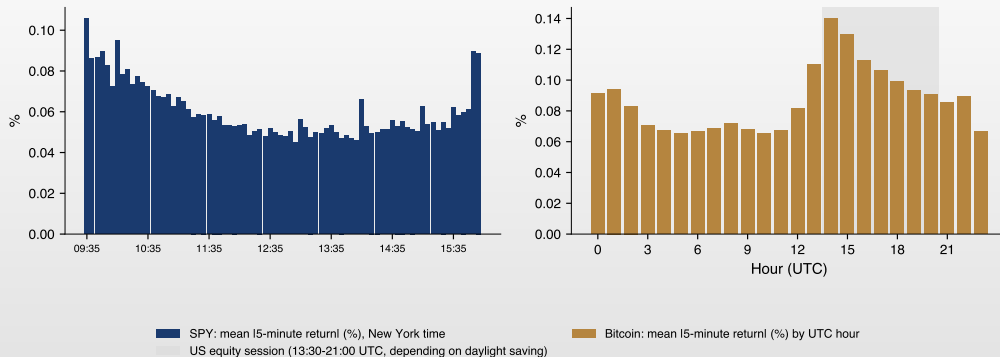


□ SPY, March–May 2025; band: 95% confidence interval for $\sqrt{252 IV_t}$ from realised quarticity

Reading the Confidence Intervals

- Median relative width of the interval (for volatility): 37% of its level
 - ▶ The band assumes a continuous price and negligible noise: on a jump day RV tends to $IV_t + \sum_j \kappa_j^2$, and $\sum r^4$ keeps the term κ^4 , so the band has no nominal coverage for either the total variation or IV_t (Section 4)
- 9 April 2025 (the day the tariff pause was announced): volatility 72%, CI [55%, 93%]
 - ▶ The interval widens when a few returns dominate the sum: $\sum r^4$ grows faster than $(\sum r^2)^2$
- Still far more precise than r_t^2 , whose relative error is of order $\sqrt{2} \approx 141\%$ under the Normal distribution
- Practical rule: a single day's volatility is known to about $\pm 20\%$ (95% band), averages over weeks much more precisely

Intraday Seasonality



Left: SPY, mean $|r_{t,i}|$ by 5-minute interval of the day (full days only); right: Bitcoin by UTC hour

Reading the Seasonality Pattern

- SPY: a **U-shape**: 0.106% in the first 5 minutes, 0.050% around midday, 0.089% in the last 5 minutes [Andersen & Bollerslev \(1997\)](#)
 - ▶ The open prices overnight news; the close concentrates index and fund rebalancing
- Bitcoin trades 24/7 but is most volatile at 14:00 UTC, near the US equity opening (09:30 New York: 13:30 UTC in summer, 14:30 UTC in winter); quietest at 10:00 UTC (ratio 2.2)
- RV sums over the day, so seasonality does not bias it; it does bias BV: $\mathbb{E}|r_i| |r_{i-1}| \propto \sigma_i \sigma_{i-1} \leq (\sigma_i^2 + \sigma_{i-1}^2)/2$, so fast changes of scale push RV – BV up and the jump test (Section 4) over-rejects
 - ▶ Remedy: divide $r_{t,i}$ by a periodicity factor f_i (standard-deviation estimator of [Boudt, Croux & Laurent \(2011\)](#)): $f = 1.78$ at the open, 0.73 at the minimum; SPY rejections at 5%: 235 $\rightarrow$ 167

Recap: Realised Variance

- ▣ $RV_t = \sum_i r_{t,i}^2$ measures integrated variance without a model; precision improves with M
- ▣ Realised quarticity gives its standard error: about $\sqrt{2/78} = 16\%$ of RV for one day of 5-minute data, i.e. a 95% band of roughly $\pm 20\%$ for volatility
- ▣ For SPY, overnight returns add 38% of daily variance: use total variance for daily risk
- ▣ Volatility follows a daily U-shape (SPY) and a UTC-hour pattern (Bitcoin)

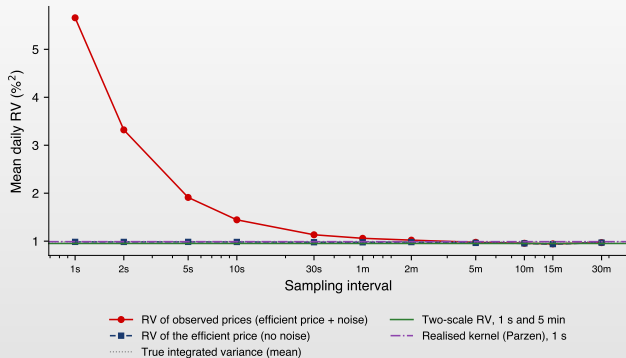
Why Not Use Every Tick?

- Observed prices are not the efficient price: $p_i^{\text{obs}} = p_i^* + \varepsilon_i$
 - ▶ Sources of ε : bid–ask bounce, tick size (price discreteness), stale quotes (Chapter 10)
- With i.i.d. noise of variance ω^2 and n returns per day [Zhang, Mykland & Aït-Sahalia \(2005\)](#), [Hansen & Lunde \(2006\)](#)
 - ▶ $\mathbb{E}[\text{RV}_t^{(n)} \mid \sigma] = \text{IV}_t + 2n\omega^2$ (zero drift, no leverage): the bias **grows** with the sampling frequency
 - ▶ Observed returns become negatively autocorrelated: $\text{cov}(r_i, r_{i-1}) = -\omega^2$
- Trade-off: more returns reduce sampling error but add noise bias

Deriving the Noise Bias and the Optimal Frequency

- Observed return $r_i = r_i^* + \varepsilon_i - \varepsilon_{i-1}$, with ε i.i.d. $(0, \omega^2)$ and independent of p^*
 - ▶ Given the path of σ (zero drift, no leverage: σ independent of W):
 $\mathbb{E}[\text{RV}^{(n)} \mid \sigma] = \sum_i \mathbb{E}[r_i^{*2} \mid \sigma] + \sum_i \mathbb{E}[(\varepsilon_i - \varepsilon_{i-1})^2] = \text{IV} + 2n\omega^2$; $\text{cov}(r_i, r_{i-1}) = -\mathbb{E}[\varepsilon_{i-1}^2] = -\omega^2$
 - ▶ By a law of large numbers for the noise terms, $\text{RV}^{(n)}/(2n) \xrightarrow{P} \omega^2$ as $n \rightarrow \infty$: the all-tick RV estimates the noise, not IV
[Bandi & Russell \(2008\)](#)
- MSE: $\mathbb{E}[(\text{RV}^{(n)} - \text{IV})^2] \approx \underbrace{2\text{IQ}/n}_{\text{discretisation}} + \underbrace{4n^2\omega^4}_{\text{squared bias}} + O(n)$
 - ▶ First-order condition $-2\text{IQ}/n^2 + 8n\omega^4 = 0$ gives $n^* = (\text{IQ}/(4\omega^4))^{1/3}$ [Bandi & Russell \(2008\)](#)
- Numbers: $\text{IV} = 1\%^2$, $\text{IQ} = 1$, $\omega = 0.01\%$: bias 0.016 at 5 minutes, 4.680 at 1 second; $n^* = 292$ (one return every 80 s)

The Volatility Signature Plot: Simulation



100 simulated days, one price per second; efficient price with random daily volatility plus i.i.d. noise, $\omega = 0.01\%$

Reading the Simulated Signature Plot

- True mean IV = 0.98; RV of observed prices: 5.66 at 1 second, 1.06 at 1 minute, 0.98 at 5 minutes
 - ▶ The bias matches $2n\omega^2$ from the previous slide
- Without noise, RV is flat in the sampling interval: the rise at high frequency is the fingerprint of noise
- Noise-robust estimators on 1-second data: two-scale RV 0.95, realised kernel 0.99
 - ▶ Error standard deviation across days: 5-minute RV 0.187, two-scale 0.150, kernel 0.045: using all data pays off only with a noise correction

Remedies: Sparse Sampling and Subsampling

- **Sparse sampling:** use one return every 5 minutes instead of every tick
 - ▶ Simple and robust; throws away most of the data [Liu, Patton & Sheppard \(2015\)](#)
- **Subsampling:** compute RV on K shifted grids (offsets $0, 1, \dots, K - 1$) and average
 - ▶ $RV_t^{\text{avg}} = \frac{1}{K} \sum_{k=0}^{K-1} RV_t^{(k)}$: same bias as one sparse grid, smaller variance
- **Two-scale RV (TSRV)** [Zhang, Mykland & Aït-Sahalia \(2005\)](#): remove the remaining bias with the all-data RV
 - ▶ $TSRV_t = RV_t^{\text{avg}} - \frac{\bar{n}}{n} RV_t^{\text{all}}$, $\bar{n} = (n - K + 1)/K$; divide by $1 - \bar{n}/n$ in small samples
 - ▶ $\bar{n}$ is the average number of returns per sparse grid; our grids also keep the open and the close, so $\bar{n}$ is their actual average count

Realised Kernels



Neil Shephard, 2004

Photo: Renate Schmid (2004); CC BY-SA
2.0 de; Wikimedia Commons

- ▣ **Realised kernel** Barndorff-Nielsen, Hansen, Lunde & Shephard (2008): add weighted autocovariances of returns
 - ▶ $RK_t = \gamma_0 + 2 \sum_{h=1}^H k\left(\frac{h}{H+1}\right) \gamma_h$, $\gamma_h = \sum_i r_{t,i} r_{t,i-h}$
 - ▶ Parzen weight function k : 1 at 0, falling smoothly to 0 at 1; keeps RK positive
- ▣ Noise makes γ_1 negative; the kernel adds it back and cancels the bias
- ▣ Bandwidth H grows with the noise-to-signal ratio: $H \approx 3.5 \xi^{4/5} n^{3/5}$, $\xi^2 = \omega^2 / \sqrt{IQ}$
- ▣ Shephard and co-authors developed kernels, bipower variation and the asymptotic theory of RV

Rates of Convergence under Noise

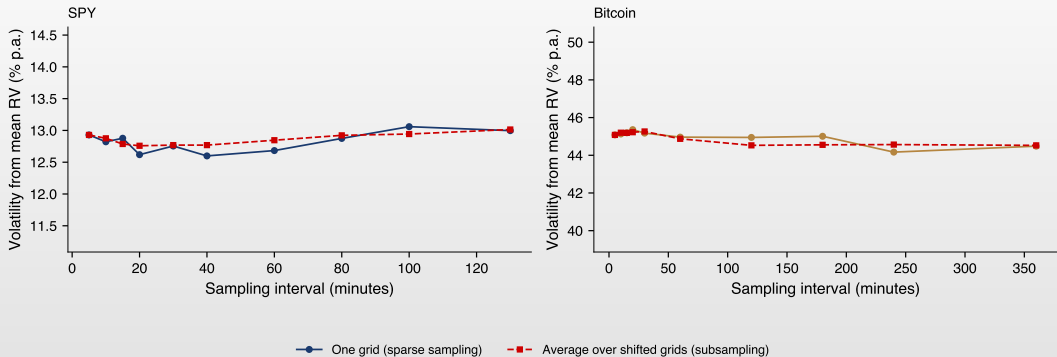
Estimator	Error of order	Positive	Source
RV on all ticks	inconsistent (bias $2n\omega^2$)	yes	Bandi & Russell (2008)
Sparse RV, fixed 5-minute grid	$M^{-1/2}$, M fixed; bias stays	yes	Liu, Patton & Sheppard (2015)
Two-scale RV	$n^{-1/6}$	no	Zhang, Mykland & Aït-Sahalia (2005)
Multi-scale RV	$n^{-1/4}$	no	Zhang (2006)
Realised kernel, flat-top / Parzen	$n^{-1/4}$ / $n^{-1/5}$	no / yes	Barndorff-Nielsen, Hansen, Lunde & Shephard (2008)
Pre-averaging	$n^{-1/4}$	no	Jacod, Li, Mykland, Podolskij & Vetter (2009)

- $n^{-1/4}$ is the optimal rate with i.i.d. noise: the price of noise is half the usual exponent
- With $n = 23\,400$ seconds: $n^{1/4} \approx 12.4$, against $\sqrt{78} \approx 8.8$ for 5-minute RV: a heuristic only (rates ignore constants, noise size and the bias of sparse RV), consistent with the modest gains in [Liu, Patton & Sheppard \(2015\)](#)

Pre-Averaging

- Local averages of $k_n \approx \theta\sqrt{n}$ returns with a weight g on $[0, 1]$, e.g. $g(x) = \min(x, 1 - x)$ Jacod, Li, Mykland, Podolskij & Vetter (2009)
 - ▶ $\bar{r}_i = \sum_{j=1}^{k_n-1} g(j/k_n) r_{i+j}$: noise variance shrinks like $1/k_n$, the signal keeps order k_n/n
- Estimator, with $\psi_1 = \int_0^1 g'(x)^2 dx$, $\psi_2 = \int_0^1 g(x)^2 dx$ (1 and 1/12 for the triangular g)
 - ▶ $\widehat{IV} = \frac{1}{\theta \psi_2 \sqrt{n}} \sum_i \bar{r}_i^2 - \frac{\psi_1}{2 \theta^2 \psi_2 n} \sum_i r_i^2$ (second term: noise bias correction)
- Rate $n^{-1/4}$; the same device gives jump-robust (pre-averaged bipower) and covariance estimators
- Built for tick data: on 5-minute bars little noise is left to average out (next slides)

The Signature Plot on Real Data



- Annualised volatility from mean RV by sampling interval: one grid vs the average over all shifted grids

Reading the Empirical Signature Plot

- SPY: 12.9% at 5 minutes, 12.8% at 30 minutes, 13.0% at 130 minutes: flat
 - ▶ First-order autocorrelation of 5-minute returns: -0.008 (SPY), 0.005 (Bitcoin): no visible noise at 5 minutes
- Bitcoin: 45.1% at 5 minutes, 44.5% at 6 hours: flat as well
- On 5-minute SPY data: two-scale $RV / RV = 0.97$, realised kernel / $RV = 0.95$; correlations with RV 0.89 and 0.85
 - ▶ The diagnostics show little noise bias at 5 minutes: the corrections move RV by a few per cent and add estimation variability
- Consistent with [Liu, Patton & Sheppard \(2015\)](#): almost 400 estimators on 31 assets rarely beat 5-minute RV; noise-robust estimators pay off when noise is large relative to the signal (tick or second-level data, illiquid assets)

Realised Covariance and Asynchronous Trading

- Realised covariance $\sum_i r_{t,i}^{(1)} r_{t,i}^{(2)} \rightarrow \int \sigma_{12}(s) ds$ needs prices of both assets on the **same** grid
 - ▶ Epps effect [Epps \(1979\)](#): measured correlation falls towards 0 as sampling gets finer, because trades are asynchronous and the last price is stale
- [Hayashi & Yoshida \(2005\)](#): sum the products of **all** pairs of returns whose time intervals overlap
 - ▶ $\sum_{i,j} r_i^{(1)} r_j^{(2)} 1\{(t_{i-1}, t_i] \cap (s_{j-1}, s_j] \neq \emptyset\}$: unbiased without synchronisation (in the absence of noise)
- With noise as well: pre-averaged or multivariate kernel versions; application: intraday portfolio risk (Chapter 4)

Recap: Microstructure Noise

- ▣ Noise adds $2n\omega^2$ to RV: sampling too often overstates volatility
- ▣ The signature plot is the diagnostic: RV should not depend on the sampling interval
- ▣ Remedies: 5-minute sampling, subsampling, two-scale RV ($n^{-1/6}$), flat-top realised kernels and pre-averaging ($n^{-1/4}$), Parzen kernel ($n^{-1/5}$)
- ▣ For liquid assets at 5 minutes (SPY, Bitcoin) plain RV is the benchmark

Prices Jump

- Add jumps to the model: $dp = \mu ds + \sigma dW + dJ$, with J a jump process (news, Chapter 11)
 - ▶ Quadratic variation now has two parts: $RV_t \rightarrow IV_t + \sum_j \kappa_j^2$, with κ_j the jump sizes in day t
- Why separate them? Jumps are less persistent than diffusive volatility, and they price differently (options, Chapter 12)
- We need an estimator of IV_t that ignores jumps

Bipower Variation



Ole E. Barndorff-Nielsen, 2007

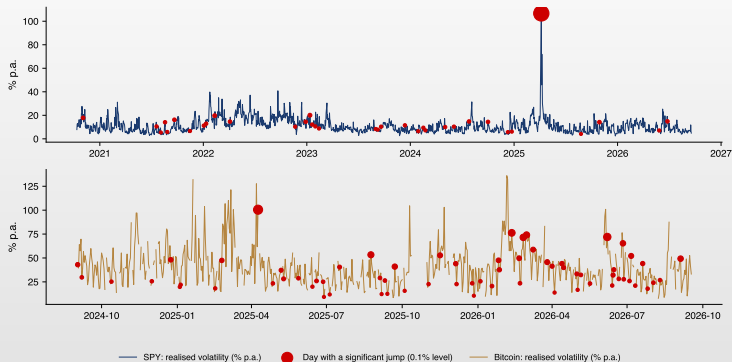
Photo: Thomas Steiner (2007); CC BY-SA 2.5;
Wikimedia Commons

- **Bipower variation** Barndorff-Nielsen & Shephard (2004): products of adjacent absolute returns
 - ▶ $BV_t = \mu_1^{-2} \frac{M}{M-1} \sum_{i=2}^M |r_{t,i}| |r_{t,i-1}| \rightarrow IV_t$
 - ▶ $\mu_1 = \mathbb{E}|Z| = \sqrt{2/\pi} = 0.798$ for $Z \sim N(0, 1)$, so $\mu_1^{-2} = \pi/2 = 1.571$
- A jump enters one return: it is multiplied by a small neighbour, so its effect vanishes as M grows
- Jump part of the day: $\max(RV_t - BV_t, 0)$
- Barndorff-Nielsen (1935–2022) also shaped the Lévy-process models used for jumps

A Test for Jumps

- ▣ Ratio statistic [Huang & Tauchen \(2005\)](#), [Barndorff-Nielsen & Shephard \(2006\)](#):
$$z_t = \frac{(RV_t - BV_t)/RV_t}{\sqrt{\theta M^{-1} \max(1, TQ_t/BV_t^2)}}$$
 - ▶ $\theta = \pi^2/4 + \pi - 5 = 0.609$; tripower quarticity $TQ_t = M\mu_{4/3}^{-3} \frac{M}{M-2} \sum_i |r_{t,i}|^{4/3} |r_{t,i-1}|^{4/3} |r_{t,i-2}|^{4/3}$, $\mu_{4/3} = \mathbb{E}|Z|^{4/3}$
- ▣ Under no jumps $z_t \rightarrow N(0, 1)$; a jump makes $RV > BV$ and z_t large and positive
- ▣ One test per day: at level α we expect $\alpha \times$ days false alarms
 - ▶ We use $\alpha = 0.1\%$ (critical value 3.09): about 1.5 false jump days for SPY
- ▣ Significant day: jump part $J_t = RV_t - BV_t$, continuous part $C_t = RV_t - J_t$; otherwise $J_t = 0$

Jump Days: SPY and Bitcoin



□ Red points: days with a significant jump at the 0.1% level; point size proportional to the jump part J_t

Reading the Jump Test

- SPY: 35 jump days of 1 489 (2.4%); jumps are 2.8% of total realised variance
 - ▶ On a jump day the jump is on average 35% of that day's RV; BV/RV averages 0.95 over all days
 - ▶ Largest statistic: 2021-08-27 ($z = 5.2$), jump in the 10:00–10:05 interval, when the Federal Reserve chair's Jackson Hole speech was released
- Bitcoin: 67 jump days of 659 (10.2%), but only 1.8% of variance
 - ▶ Many small jumps on quiet days (weekends): one isolated move dominates a calm day
- The turbulent days of April 2025 are mostly **not** jump days: many large returns in a row raise BV as well

Many Days, Many Tests: Spurious Jumps

- 1 489 daily tests at 5%: 235 rejections against 74 expected under no jumps; which of them are false?
- [Benjamini & Hochberg \(1995\)](#): sort the p -values, reject the k smallest, with k the largest index such that $p_{(k)} \leq kq/N$; the FDR, $\mathbb{E}[\text{false/rejections}]$, stays $\leq q$ (with false/rejections $:= 0$ when nothing is rejected)
 - ▶ The guarantee needs valid p -values that are independent or positively dependent [Benjamini & Yekutieli \(2001\)](#); daily jump tests are only approximately so
 - ▶ SPY: FDR at 5%: 42 jump days; at 1%: 9; Bonferroni (FWER 5%): 7; FDR 5% after the periodicity correction: 17
- [Bajgrowicz, Scaillet & Treccani \(2016\)](#): without multiplicity control most detected jumps are spurious; the true ones are rare and cluster on news
- [Christensen, Oomen & Podolskij \(2014\)](#): at tick frequency most 5-minute “jumps” are steep but continuous bursts of volatility (fact or friction)

Locating Jumps Within the Day: Lee–Mykland

- Lee & Mykland (2008): $L_{t,i} = r_{t,i} / \hat{\sigma}_{t,i}$, with $\hat{\sigma}_{t,i}^2$ the bipower average of the previous $K = 270$ returns (their value for 5-minute data)
 - ▶ Under no jumps, $\max_i |L_i|$ over n returns has a Gumbel limit: reject if $(|L| - C_n) / S_n > -\ln(-\ln(1 - \alpha))$, with $C_n = \sqrt{2 \ln n} - \frac{\ln \pi + \ln \ln n}{2\sqrt{2 \ln n}}$ and $S_n = 1 / \sqrt{2 \ln n}$ (our $\hat{\sigma}^2$ includes the factor $\pi/2$, so the constant $c = \sqrt{2/\pi}$ of Lee & Mykland (2008) becomes 1)
 - ▶ SPY, returns divided by f_i , $\alpha = 1\%$, $n = 115\,453$ (the whole sample): threshold $|L| > 5.41$
- 169 jumps on 105 days; 62 of them on Wednesdays between 14:00 and 15:00, when FOMC statements (14:00) and press conferences (14:30) are released
 - ▶ Largest: $L = -22.9$ on 2025-10-29; only 10% of these days are also flagged by the daily ratio test
- The two tests ask different questions: one extreme return relative to local volatility, versus a daily gap between RV and BV

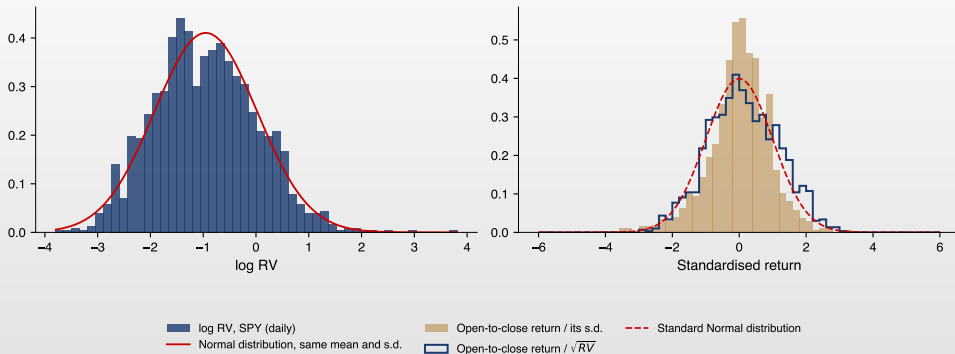
Jump-Robust Estimators of Integrated Variance

- Truncated RV [Mancini \(2009\)](#): $\text{TRV}_t = \sum_i r_{t,i}^2 \mathbf{1}\{|r_{t,i}| \leq c \hat{\sigma}_t f_i \Delta^\varpi\}$, $\varpi \in (0, 1/2)$
 - ▶ Diffusive increments are $O_p(\Delta^{1/2})$ and end up below the threshold; jumps are $O(1)$ and are cut off
- MedRV [Andersen, Dobrev & Schaumburg \(2012\)](#): $\frac{\pi}{6-4\sqrt{3}+\pi} \frac{M}{M-2} \sum_i \text{med}(|r_{i-1}|, |r_i|, |r_{i+1}|)^2$
 - ▶ An isolated return larger than both neighbours never enters the median (which shifts to a neighbour), so a large isolated jump is excluded, though its effect is not removed exactly in finite samples; MedRV is less sensitive than BV to isolated zero returns; two adjacent jumps can still enter the median
- Jump share of SPY RV: BV 6.3%, MedRV 7.2%, TRV 12.5% ($c = 3$, $\varpi = 0.49$; truncation bites on 44% of days: c is a tuning choice)
- [Aït-Sahalia & Jacod \(2009\)](#): compare power variations $\sum |r|^p$, $p > 2$, at Δ and $k\Delta$; the ratio tends to 1 with jumps and to $k^{p/2-1}$ without: known limits under both hypotheses

Recap: Jumps

- ▣ RV measures total quadratic variation: diffusion plus squared jumps
- ▣ Bipower variation is robust to jumps; $RV - BV$ estimates the jump part
- ▣ The ratio test flags about 2.4% of SPY days at the 0.1% level; jumps are a small share of total variance
- ▣ Test many days, control the FDR: 235 SPY days at 5% shrink to 42 with Benjamini–Hochberg; Lee–Mykland locates jumps in time (FOMC releases)
 - ▶ Jumps are separated from continuous variation in the HAR-CJ model (Section 5, Seminar 9)

Distributions: log RV and Standardised Returns

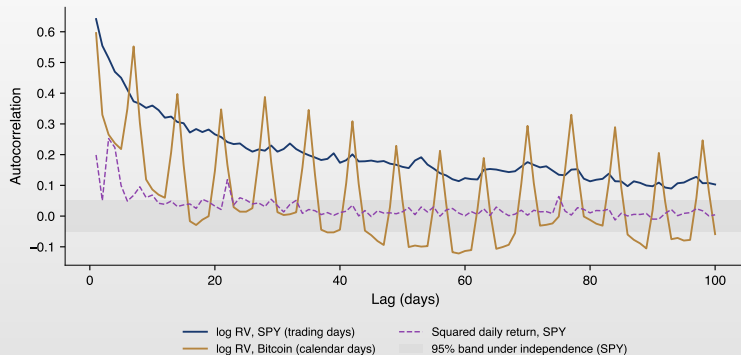


- Left: daily log RV of SPY and the Normal density; right: open-to-close returns scaled by their standard deviation (filled) and by $\sqrt{RV_t}$ (outline)

Reading the Distributions

- SPY RV: skewness 20.0, kurtosis 550; log RV: skewness 0.29, kurtosis 3.20 (the Normal distribution: 0 and 3)
 - ▶ Bitcoin log RV: skewness -0.27 , kurtosis 3.33: model volatility in logs
- Open-to-close returns: kurtosis 15.73; divided by $\sqrt{RV_t}$: 2.65, standard deviation 1.05
 - ▶ Bitcoin: from 7.57 to 2.77
- Heavy tails of daily returns (Chapter 1) come mostly, ex post, from changing volatility: approximately a **mixture** of Normal distributions with variances IV_t , exact only without leverage and jumps (Section 2)

Long Memory in Realised Volatility

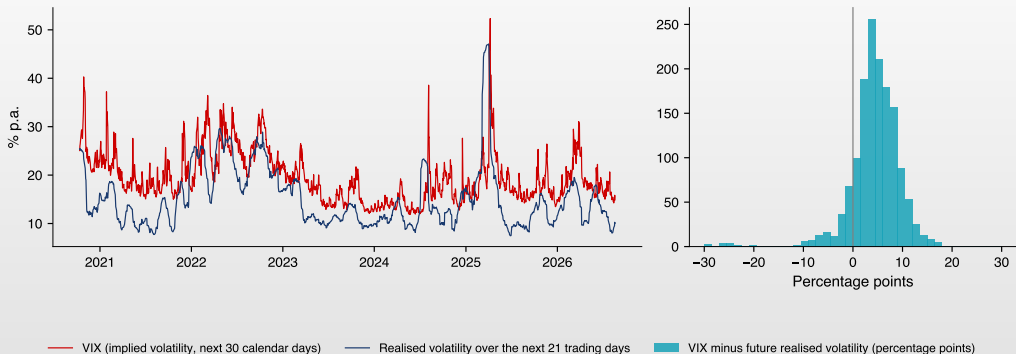


- Autocorrelations up to 100 days: log RV of SPY (total variance) and Bitcoin, squared daily SPY returns

Reading the Autocorrelations

- SPY log RV: 0.64 at lag 1, 0.24 at one month, still 0.10 at 100 days: **long memory**
 - ▶ An AR(1) would decay geometrically: $0.64^{22} \approx 0.00005$
- Squared daily returns show the same persistence, hidden by noise: 0.20 at lag 1
- Bitcoin: peaks at 7, 14, 21, ... days: a **weekly cycle**
 - ▶ Weekend RV is only 0.41 of weekday RV (Saturday 27% vs Wednesday 48% p.a.): no equity or macro news, fewer institutional traders

Implied vs Realised Volatility



- Left: VIX and SPY realised volatility over the next 21 trading days (including overnight returns); right: the gap between them

Reading the Variance Risk Premium

- VIX: the market's 30-day volatility, implied by S&P 500 options (Chapter 12); mean 19.5% vs 15.2% realised afterwards
 - ▶ VIX exceeded the subsequent realised volatility on 88% of 1 468 days; mean gap 4.4 percentage points
- **Variance risk premium (VRP)** Carr & Wu (2009), Bollerslev, Tauchen & Zhou (2009):
$$\text{VRP}_t = \mathbb{E}_t^Q[\text{RV}_{t,t+21}] - \mathbb{E}_t^P[\text{RV}_{t,t+21}]$$
, risk-neutral minus physical expectation of the same future variance; VIX^2 approximates the first term
 - ▶ Investors pay to insure against volatility spikes; the chart shows an **ex-post** spread (VIX minus the volatility that followed), a noisy proxy of this ex-ante premium
- As a forecast: regressing future variance on VIX^2 gives slope 0.56 (SE 0.05), $R^2 = 0.23$: VIX is informative but biased upwards
- The ex-post spread turns negative when realised volatility exceeds the earlier implied volatility (March–April 2025): a forecast error, not evidence of a negative ex-ante premium

Recap: Stylised Facts

- log RV is close to Normal Andersen, Bollerslev, Diebold & Labys (2001), Andersen, Bollerslev, Diebold & Ebens (2001): model and forecast volatility in logs
- Returns divided by same-day $\sqrt{RV}$ are close to $N(0, 1)$ ex post: fat tails come mostly from changing volatility
- Long memory: the model must weigh the last month, not only yesterday
- Implied volatility carries a risk premium: it is an upward-biased forecast

The HAR-RV Model

- Corsi (2009): $RV_{t+1} = \beta_0 + \beta_d RV_t + \beta_w RV_t^{(w)} + \beta_m RV_t^{(m)} + \varepsilon_{t+1}$
 - ▶ $RV_t^{(w)} = \frac{1}{5} \sum_{j=0}^4 RV_{t-j}$ (week), $RV_t^{(m)} = \frac{1}{22} \sum_{j=0}^{21} RV_{t-j}$ (month)
 - ▶ Bitcoin trades every day: week = 7 days, month = 30 days
- Heterogeneous market hypothesis Müller et al. (1997): daily, weekly and monthly traders; a few components with different persistence mimic long memory
- An AR(22) with only 3 free slopes: weight $\beta_d + \beta_w/5 + \beta_m/22$ on lag 1, $\beta_w/5 + \beta_m/22$ on lags 2–5, $\beta_m/22$ on lags 6–22
- Estimation by OLS; errors are heteroskedastic and autocorrelated: use Newey–West (HAC) standard errors
 - ▶ log-HAR: the same regression for $\ln RV$; forecast $\exp(\hat{y} + \hat{s}^2/2)$, $\hat{s}^2$ the residual variance: the conditional mean if the errors are $N(0, s^2)$ (lognormal approximation)
 - ▶ For non-Normal errors use the smearing factor $\frac{1}{T} \sum_t e^{\hat{\varepsilon}_t}$ Duan (1983); with heteroskedastic errors both are approximations

Extensions of HAR

- ▣ **HAR-CJ Andersen, Bollerslev & Diebold (2007)**: separate continuous and jump parts as regressors
 - ▶ $\ln RV_{t+1} = \beta_0 + \beta_{Cd} \ln C_t + \beta_{Cw} \ln C_t^{(w)} + \beta_{Cm} \ln C_t^{(m)}$
 - ▶ $+ \beta_{Jd} \ln(1 + J_t) + \beta_{Jw} \ln(1 + J_t^{(w)}) + \beta_{Jm} \ln(1 + J_t^{(m)}) + \varepsilon_{t+1}$
 - ▶ $C_t^{(w)}, J_t^{(w)}$ ($C_t^{(m)}, J_t^{(m)}$): arithmetic means over the last 5 (22) days, as $RV_t^{(w)}$; logs are taken after averaging
- ▣ **HARQ Bollerslev, Patton & Quaedvlieg (2016)**: less weight on yesterday's RV when it is imprecise (high realised quarticity)
 - ▶ β_d becomes $\beta_d + \beta_Q \sqrt{RQ_t}$ with $\beta_Q < 0$
- ▣ Weekend or holiday dummies, leverage (negative returns raise future RV), implied volatility as an extra regressor
- ▣ Multi-step: regress $\sum_{j=1}^h RV_{t+j}$ directly on the same regressors

HAR Estimates: SPY and Bitcoin


	β_0	β_d	β_w	β_m	R^2
SPY, HAR (levels)	0.273 (0.064)	0.399 (0.089)	0.222 (0.096)	0.114 (0.053)	0.31
SPY, log-HAR	-0.079 (0.027)	0.362 (0.037)	0.337 (0.048)	0.167 (0.044)	0.47
Bitcoin, HAR (levels)	1.496 (0.412)	0.412 (0.089)	0.409 (0.122)	-0.065 (0.097)	0.34
Bitcoin, log-HAR	0.223 (0.094)	0.380 (0.041)	0.465 (0.080)	-0.022 (0.098)	0.36

- Full samples; SPY: total daily variance (intraday RV plus squared overnight return), 1 467 days; Bitcoin: 24-hour RV, 547 days
- Newey–West standard errors in brackets

Reading the HAR Estimates

- SPY: all three components are significant; in logs $\hat{\beta}_d = 0.362$, $\hat{\beta}_w = 0.337$, $\hat{\beta}_m = 0.167$
 - ▶ Persistence $\hat{\beta}_d + \hat{\beta}_w + \hat{\beta}_m = 0.87$: shocks fade slowly, as in GARCH (Chapter 5)
- $R^2 = 0.47$ for log variance vs 0.31 for variance in levels: not comparable, the dependent variables differ; the out-of-sample losses on a common target (Section 7) decide
- Bitcoin: daily and weekly components matter; the monthly coefficient is close to 0 (-0.022 , SE 0.098)
 - ▶ Short sample (two years) and a strong weekly cycle: the week absorbs most of the memory

Implied Lag Weights and In-Sample Fit



- Left: weights of log-HAR on each past day (Bitcoin: complete 7- and 30-day windows); right: SPY, last 250 days, realised volatility and log-HAR fitted values

Reading the Lag Weights

- SPY: yesterday gets 0.437, each of days 2–5 gets 0.075, each of days 6–22 gets 0.008
- The staircase decays slowly: the month still holds weight $22 \times 0.008 \approx 0.167$
 - ▶ Three coefficients reproduce the long memory of the autocorrelation function
- Bitcoin: yesterday 0.446, days 2–7 about 0.066 each, almost nothing beyond a week
- In-sample fitted values follow the level of volatility, not the day-to-day noise

HAR as a Restricted AR(22): Testing the Restrictions

- Unrestricted: $\ln RV_{t+1} = \phi_0 + \sum_{k=1}^{22} \phi_k \ln RV_{t+1-k} + \varepsilon_{t+1}$
 - ▶ HAR imposes $\phi_2 = \dots = \phi_5$ (3 restrictions) and $\phi_6 = \dots = \phi_{22}$ (16): $R\phi = 0$ with 19 rows
- Wald: $W = (R\hat{\phi})^\top (R\hat{V}_{\text{NW}}R^\top)^{-1} R\hat{\phi} \rightarrow \chi^2_{19}$ under the HAR restrictions
- SPY log total variance: $W = 18.1$, $p = 0.51$; R^2 0.477 (AR(22)) vs 0.470 (HAR)
 - ▶ The staircase is not rejected: HAR saves 19 parameters at almost no in-sample cost, and less estimation noise helps out of sample

Measurement Error in the Regressors: Why HARQ

- $RV_t = IV_t + u_t$ with $\mathbb{E}[u_t | \sigma] = 0$ and $\text{Var}(u_t | \sigma) = 2IQ_t/M$ (the CLT of Section 2)
 - ▶ If $IV_{t+1} = \beta_0 + \beta IV_t + \eta_{t+1}$, OLS of RV_{t+1} on RV_t gives $\text{plim } \hat{\beta} = \beta \frac{\text{Var}(IV)}{\text{Var}(IV) + \mathbb{E}[2IQ/M]} < \beta$ for $\beta > 0$: attenuation
- If the signal variance does not rise with IQ_t (a working approximation), the signal-to-noise ratio of RV_t falls when IQ_t is high; HARQ approximates the falling weight by $\beta_d + \beta_Q \sqrt{RQ_t}$, $\beta_Q < 0$ [Bollerslev, Patton & Quaedvlieg \(2016\)](#)
- SPY, levels, Newey–West SE: intraday RV $\hat{\beta}_Q = -0.0088$ (0.0042); total variance -0.0022 (0.0024)
 - ▶ Daily slope (intraday RV) 1.13 at the median $\sqrt{RQ}$, 1.09 at its 99th percentile: the sign predicted by the theory, a small effect

Recap: HAR-RV

- HAR regresses tomorrow's RV on daily, weekly and monthly averages: OLS, three slopes
- It approximates long memory with a staircase of lag weights
- Estimate in logs; report Newey–West standard errors
- HAR is an AR(22) with 19 restrictions that the SPY data do not reject; measurement error in RV_t attenuates the daily slope, which HARQ corrects approximately
- Extensions: jumps (HAR-CJ), measurement error (HARQ), calendar effects, implied volatility

Forecasting Design

- Target: tomorrow's daily variance; SPY: total variance, 2022-01-03 to 18 September 2026 (1 180 days); Bitcoin: 2025-03-01 onwards (422 days)
 - ▶ Every forecast uses only data available the evening before
- Five models
 - ▶ log-HAR and HAR on realised variance, re-estimated daily on an expanding window
 - ▶ GARCH(1,1) with Student- t errors on daily returns since 2000 (SPY) or 2014 (Bitcoin), re-estimated every 21 days [Bollerslev \(1986\)](#)
 - ▶ EWMA with $\lambda = 0.94$ (Chapter 5) and "yesterday's RV" (random walk)
- HAR sees intraday data; GARCH and EWMA see only daily returns

Evaluating Against a Noisy Target

- True variance is unobserved; we compare forecasts F_t with the proxy RV_t
 - ▶ A loss is **robust** if the ranking by *expected* loss with a proxy that is conditionally unbiased (given the forecasts' information) equals the ranking with true variance [Patton \(2011\)](#)
- Two robust losses
 - ▶ $MSE = (RV_t - F_t)^2$: dominated by the most volatile days
 - ▶ $QLIKE = \frac{RV_t}{F_t} - \ln \frac{RV_t}{F_t} - 1 \geq 0$: depends only on the ratio; penalises under-prediction more
- Common mistake: MAE, or MSE on volatility instead of variance, is **not** robust
 - ▶ [Andersen & Bollerslev \(1998\)](#): against r_t^2 , GARCH forecasts explained a few % of the variation; against RV the same forecasts explain far more: the proxy was noisy, not the model

Tests: Diebold–Mariano and Mincer–Zarnowitz

- Diebold & Mariano (1995): loss difference $d_t = L(RV_t, F_t^A) - L(RV_t, F_t^B)$; $t = \bar{d}/\text{se}_{\text{NW}}(\bar{d}) \rightarrow N(0, 1)$
 - ▶ $t < 0$: model A has the lower loss; HAC standard error because d_t is autocorrelated (Chapter 8)
- Mincer & Zarnowitz (1969): $RV_t = a + b F_t + u_t$; an unbiased forecast has $a = 0$, $b = 1$
 - ▶ Joint Wald test with Newey–West covariance; R^2 measures how much of the variation the forecast explains
 - ▶ R^2 is capped by the noise in RV: even the true variance would not reach 1 Andersen & Bollerslev (1998)

Results: SPY, 2022–2026

Model	QLIKE	DM vs log-HAR (p)	MSE	MZ b	MZ R^2
log-HAR	0.310	–	4.15	1.16	0.30
HAR	0.326	–1.95 (0.052)	4.48	0.75	0.26
GARCH(1,1)- t	0.343	–2.30 (0.021)	4.87	0.70	0.20
EWMA	0.393	–4.02 (< 0.001)	5.14	0.74	0.14
Yesterday's RV	0.502	–5.95 (< 0.001)	5.41	0.54	0.29

□ 1180 one-day-ahead forecasts of total daily variance; DM: QLIKE differences, Newey–West; $t < 0$ favours log-HAR

Reading the SPY Results

- log-HAR has the lowest QLIKE and MSE: QLIKE 10% below GARCH(1,1)- t (DM $t = -2.30$, $p = 0.021$)
 - ▶ Against EWMA and yesterday's RV the gain is highly significant; against HAR in levels only marginally ($p = 0.052$)
- MSE differences are not significant: a few crisis days (April 2025) dominate the squared errors
- Mincer–Zarnowitz: log-HAR $b = 1.16$ ($p = 0.554$ for $a = 0$, $b = 1$), GARCH $b = 0.70$: no significant bias, but log-HAR explains more (R^2 0.30 vs 0.20)
- Yesterday's RV has a high R^2 but $b = 0.54$: right direction, far too extreme

Results: Bitcoin, 2025–2026

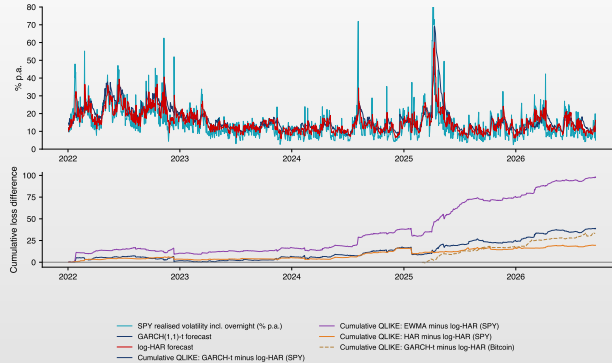
Model	QLIKE	DM vs log-HAR (p)	MSE	MZ b	MZ R^2
log-HAR	0.282	–	26.6	0.82	0.38
HAR	0.409	–1.22 (0.223)	27.6	0.89	0.34
GARCH(1,1)- t	0.360	–4.32 (< 0.001)	42.2	0.55	0.23
EWMA	0.346	–2.96 (0.003)	36.3	0.64	0.18
Yesterday's RV	0.466	–5.38 (< 0.001)	35.0	0.58	0.33

□ 422 one-day-ahead forecasts of 24-hour realised variance; same models and tests as for SPY

Reading the Bitcoin Results

- log-HAR wins again: QLIKE 22% below GARCH(1,1)- t (DM $t = -4.32$, $p < 0.001$)
 - ▶ GARCH on daily data is biased: MZ $b = 0.55$ ($p < 0.001$): it overreacts to large daily moves
- HAR in levels is worse than GARCH under QLIKE (0.409 vs 0.360): after calm weekends it predicts too low a variance, and QLIKE punishes under-prediction
- Intraday information can enter Bitcoin risk forecasts; entropy of high-frequency returns has been proposed as one such signal ([Pele & Mazurencu-Marinescu-Pele \(2019\)](#))
- Crypto-assets behave like a distinct asset class with their own volatility dynamics [Pele et al. \(2023\)](#)

Forecasts over Time




- Top: SPY realised volatility (with overnight) and the two forecasts, annualised; bottom: cumulative QLIKE loss of each model minus log-HAR (rising line = log-HAR better)

Reading the Forecast Paths

- After a shock both forecasts jump; GARCH stays high longer (its persistence is close to 1), log-HAR returns faster
 - ▶ April 2025: the GARCH forecast stayed above 40% p.a. until late April, while realised volatility was below 20% from mid-April
- The cumulative loss difference rises steadily, with jumps at turbulent episodes: the gain is not driven by one day
- Of the three models plotted against log-HAR, EWMA loses most: a fixed $\lambda = 0.94$ reacts too much to single days

Mincer–Zarnowitz Regressions



□ SPY, 2022–2026: realised variance against the forecast (log scales); dashed: 45-degree line; black: fitted regression in levels

Reading the Mincer–Zarnowitz Plots

- log-HAR forecasts span a wider range: they separate calm from turbulent days better
- Only 1% of GARCH forecasts are below $0.26\%^2$ per day (log-HAR: minimum $0.14\%^2$): on calm days realised variance falls far below GARCH
- The cloud is wide for both: one day's variance is hard to predict, and RV itself is measured with error
- Weekly horizon and model extensions: Seminar 9 (B6, B7)

Beyond HAR: Semivariances and Joint Return–RV Models

- Realised semivariances $RS_t^\pm = \sum_i r_{t,i}^2 \mathbf{1}\{r_{t,i} \gtrless 0\}$, $RS^+ + RS^- = RV$; SHAR [Patton & Sheppard \(2015\)](#): $\beta_d RV_t$ becomes $\beta^+ RS_t^+ + \beta^- RS_t^-$
 - ▶ SPY, levels (overnight return counted as one more return of the day): $\hat{\beta}^- = 0.56$ (0.23), $\hat{\beta}^+ = 0.21$ (0.18); Wald $\beta^+ = \beta^-$: $p = 0.35$
 - ▶ Bad volatility is more persistent, as in [Patton & Sheppard \(2015\)](#), but six years of data do not separate the two slopes
- Realized GARCH [Hansen, Huang & Shek \(2012\)](#): a GARCH equation driven by the realised measure x_t plus a **measurement equation**
 - ▶ $\ln h_t = \omega + \beta \ln h_{t-1} + \gamma \ln x_{t-1}$, $\ln x_t = \xi + \varphi \ln h_t + \tau(z_t) + u_t$, $r_t = \sqrt{h_t} z_t$: joint quasi-maximum likelihood
- HEAVY models [Shephard & Sheppard \(2010\)](#): one GARCH-type equation for the daily variance, driven by lagged RV, and one for RV itself: consistent multi-step forecasts

Which Models Survive? The Model Confidence Set

Model	QLIKE	MCS p -value	In MCS (90%)
log-HAR	0.310	1.000	✓
SHAR	0.320	0.224	✓
HAR	0.326	0.014	
HARQ	0.344	0.095	
GARCH(1,1)- t	0.343	0.095	
EWMA	0.393	0.012	
Yesterday's RV	0.502	0.000	

- Hansen, Lunde & Nason (2011): test equal predictive ability of all models (range statistic, block bootstrap), drop the worst, repeat; the surviving set contains the best model with asymptotic probability $\geq 90\%$ (under the method's regularity conditions)
- SPY, 1 180 days from 2022: the set is {log-HAR, SHAR}; HARQ and SHAR are estimated in levels with the filter of Bollerslev, Patton & Quaadvlieg (2016) (forecasts outside the in-sample range replaced by the mean)
- A multiple-comparison procedure: many pairwise DM tests without size control would overstate the evidence (Chapter 8)


Recap: Forecasting

- Evaluate variance forecasts with robust losses (QLIKE, MSE on variance) and test differences with Diebold–Mariano
- Intraday data help: log-HAR beats GARCH(1,1)- t by 10% (SPY) and 22% (Bitcoin) in QLIKE
- Work in logs: HAR in levels loses under QLIKE when it under-predicts
- MSE rankings rest on a few crisis days and are rarely significant; the MCS keeps only log-HAR and SHAR for SPY

Is Volatility Rough?

- Gatheral, Jaisson & Rosenbaum (2018): log volatility behaves like a **fractional Brownian motion** with Hurst exponent $H \approx 0.1$
 - ▶ $H = 0.5$: ordinary Brownian motion; $H < 0.5$: increments are negatively correlated, paths are rougher; $H > 0.5$: smoother, trending
- Scaling law: $m(q, \Delta) = \frac{1}{T} \sum_t |\ln \sigma_{t+\Delta} - \ln \sigma_t|^q \propto \Delta^{qH}$
 - ▶ Regress $\ln m(q, \Delta)$ on $\ln \Delta$: slope $\zeta_q = qH$; then H from the slopes across q
- Rough models reproduce the shape of the implied volatility smile for short maturities (Chapter 12)

Scaling of Log Volatility Increments



□ Left: $\ln m(q, \Delta)$ against $\ln \Delta$ for SPY, $\Delta = 1, \dots, 30$ days; right: slopes ζ_q against q

Reading the Roughness Estimates

- Straight lines on the log-log scale: the scaling law holds; slopes are linear in q
 - ▶ $\hat{H} = 0.11$ for SPY (block bootstrap 95% CI [0.08, 0.13]), 0.07 for Bitcoin (only observed pairs at each calendar lag; interpolating the missing days gives 0.08): far below 0.5
- Caution: RV measures volatility with error, and measurement error lowers $\hat{H}$
 - ▶ A simulated Brownian motion gives $\hat{H} = 0.46$; plus noise of a similar relative size, $\hat{H} = 0.18$ (Seminar 9, B8)
- The scaling estimator on noisy RV proxies is **inconsistent** for the Hurst exponent; consistency needs high-frequency asymptotics [Fukasawa, Takabatake & Westphal \(2022\)](#)
 - ▶ Rough-looking estimates can come from estimation error in a non-rough model [Cont & Das \(2024\)](#); short-run roughness can coexist with long memory [Bennedsen, Lunde & Pakkanen \(2022\)](#)

Recap: Rough Volatility

- The moments of log-volatility increments scale like Δ^{qH} with $H \approx 0.1$
- Our data: $\hat{H} = 0.11$ (SPY) and 0.07 (Bitcoin)
- Part of the roughness can be measurement error in RV: i.i.d. error biases $\hat{H}$ down, other error structures can bias it either way, so $\hat{H}$ is not a bound on the latent H
- Rough volatility models are an active research area in option pricing

Case Study: Research Question

- Christensen, Siggaard & Veliyev (2023): does off-the-shelf machine learning (ML) forecast daily realised variance better than the HAR family, and if so, why?
 - ▶ HAR Corsi (2009) is linear in three regressors built from RV (daily, weekly and monthly averages), while RV is right-skewed and has long memory (stylised facts above): a nonlinear map may fit better
 - ▶ Its extensions (log-HAR, SHAR Patton & Sheppard (2015), HARQ Bollerslev, Patton & Quaedvlieg (2016)) still condition only on past returns
- Why other predictors rarely enter: OLS overfits when regressors are many, correlated and noisy
 - ▶ Candidates: implied volatility, the VIX, trading volume, macro and uncertainty indices (Table 1 of the paper)
- Why it matters: variance forecasts feed VaR (Chapter 7) and option pricing (Chapter 12); ML must beat a strong benchmark with three slope coefficients and an intercept
 - ▶ Design choice: minimal hyperparameter tuning, so the test is of the methods “as is”, not of the tuning effort


Case Study: Data and Method

- Data (Section 3): 29 of the 30 DJIA stocks (Visa excluded), 29 January 2001 – 31 December 2017, $T = 4\,257$ days; 5-minute RV from cleaned NYSE trade prices, $n = 78$ returns a day
 - ▶ Split: training 70% (2 964 days), validation 10% (424), test 20% (847); robustness: training 1 000 or 2 000 days, validation 200 (Appendix A.1)
 - ▶ $\mathcal{M}_{\text{HAR}}$: daily, weekly, monthly RV; $\mathcal{M}_{\text{ALL}}$ adds 9 predictors (Table 1): implied volatility, earnings-day dummy, 1-week momentum, change in log dollar volume; VIX, squared Hang Seng return, business-conditions index, change in the 3-month T-bill rate, policy-uncertainty index
- Models (Section 2): HAR, log-HAR, HAR with leverage, SHAR, HARQ by OLS on a rolling window; ridge (L_2 penalty), lasso (L_1), elastic net (a mix), adaptive and post-lasso, penalties tuned on the validation set
 - ▶ Trees: bagging, random forest (RF) with Breiman's defaults, 500 trees, $J/3$ candidate predictors per split (J : number of predictors), minimum node size 5 (Table 15), gradient boosting
 - ▶ Neural networks NN_k^e : $k = 1, \dots, 4$ hidden layers (2; 4-2; 8-4-2; 16-8-4-2 neurons), leaky ReLU, Adam, dropout, early stopping; $e = 1$: best of 100 seeds on validation, $e = 10$: average of the best 10
- Evaluation (Section 2.6): out-of-sample MSE relative to HAR, averaged over stocks; one-sided DM test per stock [Diebold & Mariano \(1995\)](#); MCS at 90% [Hansen, Lunde & Nason \(2011\)](#)

Case Study: Key Results

- One day ahead, RV lags only (Table 2): MSE relative to HAR: log-HAR 0.995, HARQ 1.059, RF 1.020, neural networks 0.954–0.990
 - ▶ Past RV leaves little to learn: the networks give the clearest gains (1–5%), log-HAR barely beats HAR
- One day ahead, all 12 predictors (Table 3): OLS with all predictors 0.966, HARQ 1.289; elastic net 0.916, RF 0.901, log-HAR 0.901, NN_2^{10} 0.885
 - ▶ Extra predictors pay most with shrinkage, nonlinearity or the log transform; plain OLS gains less (0.966), and HARQ overfits
- One month ahead, all predictors (Table 7): RF 0.604, bagging 0.621, networks 0.786–0.964, log-HAR 0.919, OLS with all predictors 1.290
 - ▶ The gain grows with the horizon: the ML fits are more persistent (Figure 8) and approximate the long memory of RV better
- MCS at 90% (Figure 4): with all predictors, the networks stay in the set for about 80% of stocks, HAR for about 20–40%
 - ▶ Accumulated local effects (Figure 7): daily RV, weekly RV and implied volatility dominate every model; RF and networks rank implied volatility higher than the linear models

Case Study: Gains by Horizon




□ Source: Christensen, Siggaard & Veliyev (2023), Tables 2–7: out-of-sample MSE relative to HAR, averaged over 29 DJIA stocks (test set: the last 20% of 2001–2017); left: RV lags only, right: all 12 predictors

Case Study: Reading the Horizon Profile

- ▣ RV lags only: the six models shown stay between 0.937 and 1.111 times the HAR MSE at all three horizons
 - ▶ NN_2^{10} : 0.958, 0.937, 0.958; three RV averages leave little for a flexible model to learn
- ▣ All predictors: trees and networks improve with the horizon
 - ▶ Random forest $0.901 \rightarrow 0.796 \rightarrow 0.604$; NN_2^{10} $0.885 \rightarrow 0.846 \rightarrow 0.786$
 - ▶ The authors attribute this to the higher persistence of the ML fits, closer to the long memory of RV (Figure 8)
- ▣ The linear models given the same predictors deteriorate with the horizon
 - ▶ OLS with all predictors $0.966 \rightarrow 1.037 \rightarrow 1.290$, elastic net $0.916 \rightarrow 0.925 \rightarrow 1.107$, HARQ $1.289 \rightarrow 1.459 \rightarrow 1.569$
 - ▶ Only log-HAR stays below 1 (0.901, 0.970, 0.919): the log transform down-weights the high-variance days that dominate the MSE

Case Study: Ranking of 21 Models, One Month Ahead



Source: Christensen, Siggaard & Veliev (2023), Table 7: MSE relative to HAR, one month ahead, all 12 predictors, averaged over 29 DJIA stocks; colour: model family; NN_k^e: k hidden layers, e networks averaged

Case Study: Reading the Ranking

- The top two average many trees fitted on bootstrap samples: random forest 0.604, bagging 0.621; gradient boosting 0.829
 - ▶ All eight networks beat HAR (0.786–0.964), but depth does not help: the two four-layer networks are the weakest (0.902, 0.964)
- All five penalised regressions lose to HAR: ridge 1.111 and elastic net 1.107, the lasso variants 1.363–1.386
 - ▶ Linear HAR extensions with the extra predictors do worse still (from OLS with all predictors, 1.290, to HARQ, 1.569); log-HAR (0.919) is the only linear model below 1
- Each bar is an average of MSE ratios over stocks; Table 7 reports per-stock DM tests, not a test of the average ranking; the MCS of Figure 4 concerns one-day forecasts
 - ▶ Whether this ranking survives on SPY and under QLIKE is the question of Seminar 9, C3

Case Study: Replicate, Limits, Extend

- Replicate on course data: SPY 5-minute bars (data/market/intraday), open-to-close RV with 78 returns a day, as in Section 3; the authors report that SPY confirms the stock results (footnote 12)
 - ▶ Split of Appendix A.1: 1 000 training days, 200 validation days, the rest to 18 September 2026 as test set
 - ▶ Predictors: VIX (computed from S&P 500 index options; a proxy for SPY's own implied volatility, which would need SPY options), squared Hang Seng return, SPY momentum and dollar volume, T-bill and policy uncertainty (FRED); no earnings dummy for a fund (Seminar 9, C3)
- Limits stated by the authors
 - ▶ Minimal tuning: more tuning could improve ML further (Section 7); the networks are trained once and not re-estimated in the test window (Section 3)
 - ▶ Rankings depend on the split: with 1 000 training days the financial crisis falls in the test set and the HAR extensions beat HAR (Section 4.1, Appendix A.1)
 - ▶ HAR models with a separate jump component are not included (Section 2.2); accumulated local effects describe the average prediction and miss rare events such as earnings days (Section 4.2)
- Open extension: the paper ranks variance forecasts by MSE only; does the ML advantage survive QLIKE, which is robust to the noisy RV proxy [Patton \(2011\)](#), and does it hold for Bitcoin's 24-hour RV?

An Open Question: Is Volatility Really Rough?

- Gatheral, Jaisson & Rosenbaum (2018) report $H \approx 0.1$, but H is estimated from RV, a noisy proxy of the latent σ_t
 - ▶ Cont & Das (2024): the scaling estimator applied to realised volatility finds roughness even when spot volatility is smooth: an estimation artefact
 - ▶ Bolko, Christensen, Pakkanen & Veliyev (2023) (GMM that models the RV error) and Fukasawa, Takabatake & Westphal (2022) (consistent under high-frequency asymptotics) still find $H < 0.5$
- Formal question: $\ln \text{RV}_t^{1/2} = \ln \sigma_t + e_t$, with $\ln \sigma_t$ fractional of index H
 - ▶ H_0 : $H = 1/2$ (classical diffusive volatility) against H_1 : $H < 1/2$
 - ▶ Identification problem: a small H and a large $\text{Var}(e_t)$ both flatten $m(2, \Delta)$ at short lags
- Stakes: the short-maturity skew and volatility forecasts depend on H (Chapter 12)

The Discovery Loop, Step by Step (1/2)

- Tools: an AI assistant (an LLM such as Claude, ChatGPT, Gemini or GitHub Copilot); literature search: [Semantic Scholar](#), [Elicit](#), [Google Scholar](#)
- 1. **Literature**: tabulate the evidence, then check every DOI on [Crossref](#) (the title must match)
 - ▶ Tabulate the estimates of H published since 2018: estimator, sampling frequency, treatment of the RV measurement error, point estimate and standard error, DOI.
- 2. **Competing hypotheses**, each with its implied shape of $m(2, \Delta)$
 - ▶ Derive $m(2, \Delta)$ for three processes that give H near 0.1 with the scaling estimator: fractional log-volatility with $H = 0.1$; $H = 0.5$ plus i.i.d. error; a two-factor model with one fast mean-reverting factor.
- 3. **Data and code**: replicate the chapter's number first
 - ▶ Write Python that builds daily total variance from data/market/intraday/SPY.US_5m.csv (5-minute RV plus the squared overnight return, previous close to today's open), computes $m(q, \Delta)$ for q in $\{0.5, 1, 1.5, 2, 3\}$, $\Delta = 1..30$, and the slope H . Target: $\hat{H} = 0.11$

The Discovery Loop, Step by Step (2/2)

- 4. **Robustness**, pre-registered, with a power check
 - ▶ Write a pre-registered plan: lag windows [1,30], [5,30], [10,50]; RV, bipower variation and 30-minute subsampled RV; a fit $m(2, \Delta) = c \Delta^{2H} + k$; block-bootstrap intervals; size and power by simulation at $H = 0.1$ and 0.5 .
- 5. **Critique**: ask for the strongest counter-argument
 - ▶ Act as a referee: which non-i.i.d. error structures in log RV (serially correlated, level-dependent, driven by overnight returns or jumps) would survive the constant-k correction, and how would you detect each?
- 6. **Report**: all windows and measures, bootstrap intervals, the simulation design, code and the AI log

What the Researcher Must Check

- **Proxy error:** RV is an estimate with a known standard error; test: simulate $H = 0.5$ plus error of the same size and check that the procedure does not return $H \approx 0.1$
- **Overlapping increments:** $m(q, \Delta)$ reuses the same days at every Δ , so OLS standard errors of ζ_q are invalid; test: block bootstrap
- **Microstructure noise and jumps:** check the signature plot, then repeat with BV and with subsampled RV
- **Specification search:** lag windows and q grids chosen after seeing $\hat{H}$ inflate significance; fix them in advance and report all
- **Invented estimates:** assistants quote " $H \approx 0.1$ for all assets"; check each number against the table of the cited paper

Mini-Case: Is $\hat{H}$ an Artefact of i.i.d. Error?

- A typical assistant proposal: $\hat{H} \approx 0.1$ comes from i.i.d. measurement error in RV. Implication: $m(2, \Delta) = c \Delta^{2H} + k$ with $k = 2 \text{Var}(e_t) > 0$, so $\hat{H}$ rises when short lags are dropped

	$\Delta \in [1, 30]$	[5, 30]	[10, 50]	Fit with k , [1, 50]
Simulated $H = 0.5$ plus error	0.18	0.24	0.26	0.39
SPY, total RV	0.106	0.089	0.077	0.096
SPY, intraday RV	0.129	0.094	0.078	0.113

- First column replicates the chapter; the simulated row (one path, not a power study) shows the mechanism: $\hat{H}$ rises and the fit recovers most of the gap to 0.5
- SPY: $\hat{H}$ falls with the window and the fitted k is 0; $\hat{H} = 0.096$, block bootstrap 95% CI [0.058, 0.268]
- Interpretation: an illustrative diagnostic, inconsistent with the i.i.d.-error explanation; a formal rejection needs a simulation-calibrated test; the non-i.i.d. mechanism of [Cont & Das \(2024\)](#) is not tested by this design. Code: `Quantlets/Ch_09/ai_discovery_case.py`

Your Turn: A Project Seed

- Question: does the roughness of SPY and Bitcoin volatility survive an estimator that models the RV measurement error?
- Replicate first: $\hat{H} = 0.11$ (SPY) and 0.07 (Bitcoin) from the scaling law, and $\hat{H} = 0.18$ for simulated $H = 0.5$ plus error (Seminar 9, B8)
- Deliverables
 - ▶ A pre-registration (hypotheses, estimators, lag windows, tests, simulation design) dated before the first run
 - ▶ The GMM estimator of [Bolko, Christensen, Pakkanen & Veliyev \(2023\)](#) on the course RV series, with Monte Carlo size and power at $H \in \{0.1, 0.5\}$
 - ▶ Code, a two-page report with every specification tried, and a referee-style critique of your own result
- Log every prompt that shaped a design choice in AI_USE.md and every error the assistant made in AI_ERRORS.md: github.com/danpele/MFM/tree/main/project

Summary

- ▣ Realised variance turns intraday prices into a measurement of daily variance, with a known standard error
- ▣ Sample at 5 minutes for liquid assets; use noise-robust estimators when noise is material (tick data, illiquid assets); add the squared overnight return for daily risk
- ▣ Bipower variation, MedRV and truncation separate jumps from continuous variation; with FDR control jumps are rare and a small share of variance
- ▣ log RV is nearly Normal and has long memory: the log-HAR model captures it with three coefficients
- ▣ Out of sample, log-HAR beats GARCH for SPY and Bitcoin under QLIKE; MSE comparisons are inconclusive

Which Tool When?

Situation	Suggested tool
Liquid asset, 1–5 minute bars	RV at 5 minutes, subsampled
Tick data, or an illiquid stock (BVB)	Pre-averaging or flat-top realised kernel ($n^{-1/4}$; Parzen $n^{-1/5}$); check the signature plot
Daily risk of an exchange-traded asset	Total variance: RV plus squared overnight return
Separate news jumps from diffusion	BV or MedRV, ratio test with FDR control; Lee–Mykland for the jump time
Daily or weekly variance forecast	log-HAR (plus jumps, weekends, RQ); benchmark GARCH
Comparing forecasts	QLIKE and MSE on variance, Diebold–Mariano, Mincer–Zarnowitz, model confidence set

Coming Next: Market Microstructure (Chapter 10)

- Where does the noise in high-frequency prices come from? The limit order book, bid–ask spreads, tick sizes
- Liquidity and market impact: how prices react to trading
- High-frequency trading and crypto exchange order books
- Further reading: Andersen, Bollerslev, Diebold & Labys (2003), Corsi (2009), Franke, Härdle & Hafner (2019)

Further Study on Quantinar

Advanced courses on quantinar.com

-  Nonparametric and Semiparametric Models
-  Kalman Filter
-  Applied Time Series Analysis with Python

Code and data

- Quantlets: github.com/danpele/MFM/tree/main/Quantlets/Ch_09

References (1/5)

- ▣ Aït-Sahalia, Y., Jacod, J. (2009). Testing for jumps in a discretely observed process. *Annals of Statistics*, 37(1), 184–222.
- ▣ Andersen, T.G., Bollerslev, T. (1997). Intraday periodicity and volatility persistence in financial markets. *Journal of Empirical Finance*, 4(2–3), 115–158.
- ▣ Andersen, T.G., Bollerslev, T. (1998). Answering the skeptics: yes, standard volatility models do provide accurate forecasts. *International Economic Review*, 39(4), 885–905.
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X. (2007). Roughing it up: including jump components in the measurement, modeling, and forecasting of return volatility. *Review of Economics and Statistics*, 89(4), 701–720.
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X., Ebens, H. (2001). The distribution of realized stock return volatility. *Journal of Financial Economics*, 61(1), 43–76.
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X., Labys, P. (2001). The distribution of realized exchange rate volatility. *Journal of the American Statistical Association*, 96(453), 42–55.
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X., Labys, P. (2003). Modeling and forecasting realized volatility. *Econometrica*, 71(2), 579–625.
- ▣ Andersen, T.G., Dobrev, D., Schaumburg, E. (2012). Jump-robust volatility estimation using nearest neighbor truncation. *Journal of Econometrics*, 169(1), 75–93.
- ▣ Bajgrowicz, P., Scaillet, O., Treccani, A. (2016). Jumps in high-frequency data: spurious detections, dynamics, and news. *Management Science*, 62(8), 2198–2217.
- ▣ Bandi, F.M., Russell, J.R. (2008). Microstructure noise, realized variance, and optimal sampling. *Review of Economic Studies*, 75(2), 339–369.
- ▣ Barndorff-Nielsen, O.E., Hansen, P.R., Lunde, A., Shephard, N. (2008). Designing realized kernels to measure the ex post variation of equity prices in the presence of noise. *Econometrica*, 76(6), 1481–1536.

References (2/5)

- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2002). [Econometric analysis of realized volatility and its use in estimating stochastic volatility models](#). *Journal of the Royal Statistical Society, Series B*, 64(2), 253–280.
- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2004). [Power and bipower variation with stochastic volatility and jumps](#). *Journal of Financial Econometrics*, 2(1), 1–37.
- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2005). [How accurate is the asymptotic approximation to the distribution of realised variance?](#). In D.W.K. Andrews, J.H. Stock (Eds.), *Identification and Inference for Econometric Models*, Cambridge University Press, 306–331.
- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2006). [Econometrics of testing for jumps in financial economics using bipower variation](#). *Journal of Financial Econometrics*, 4(1), 1–30.
- ▣ Benjamini, Y., Hochberg, Y. (1995). [Controlling the false discovery rate: a practical and powerful approach to multiple testing](#). *Journal of the Royal Statistical Society, Series B*, 57(1), 289–300.
- ▣ Benjamini, Y., Yekutieli, D. (2001). [The control of the false discovery rate in multiple testing under dependency](#). *The Annals of Statistics*, 29(4), 1165–1188.
- ▣ Bennedsen, M., Lunde, A., Pakkanen, M.S. (2022). [Decoupling the short- and long-term behavior of stochastic volatility](#). *Journal of Financial Econometrics*, 20(5), 961–1006.
- ▣ Bolko, A.E., Christensen, K., Pakkanen, M.S., Veliyev, B. (2023). [A GMM approach to estimate the roughness of stochastic volatility](#). *Journal of Econometrics*, 235(2), 745–778.
- ▣ Bollerslev, T. (1986). [Generalized autoregressive conditional heteroskedasticity](#). *Journal of Econometrics*, 31(3), 307–327.
- ▣ Bollerslev, T., Patton, A.J., Quaedvlieg, R. (2016). [Exploiting the errors: a simple approach for improved volatility forecasting](#). *Journal of Econometrics*, 192(1), 1–18.
- ▣ Bollerslev, T., Tauchen, G., Zhou, H. (2009). [Expected stock returns and variance risk premia](#). *Review of Financial Studies*, 22(11), 4463–4492.

References (3/5)

- ▣ Boudt, K., Croux, C., Laurent, S. (2011). Robust estimation of intraweek periodicity in volatility and jump detection. *Journal of Empirical Finance*, 18(2), 353–367.
- ▣ Carr, P., Wu, L. (2009). Variance risk premiums. *Review of Financial Studies*, 22(3), 1311–1341.
- ▣ Christensen, K., Oomen, R.C.A., Podolskij, M. (2014). Fact or friction: jumps at ultra high frequency. *Journal of Financial Economics*, 114(3), 576–599.
- ▣ Christensen, K., Siggaard, M., Veliyev, B. (2023). A machine learning approach to volatility forecasting. *Journal of Financial Econometrics*, 21(5), 1680–1727.
- ▣ Cont, R., Das, P. (2024). Rough volatility: fact or artefact?. *Sankhya B*, 86(1), 191–223.
- ▣ Corsi, F. (2009). A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Diebold, F.X., Mariano, R.S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263.
- ▣ Duan, N. (1983). Smearing estimate: a nonparametric retransformation method. *Journal of the American Statistical Association*, 78(383), 605–610.
- ▣ Epps, T.W. (1979). Comovements in stock prices in the very short run. *Journal of the American Statistical Association*, 74(366), 291–298.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- ▣ Fukasawa, M., Takabatake, T., Westphal, R. (2022). Consistent estimation for fractional stochastic volatility model under high-frequency asymptotics. *Mathematical Finance*, 32(4), 1086–1132.

References (4/5)

- ▣ Gatheral, J., Jaisson, T., Rosenbaum, M. (2018). *Volatility is rough*. *Quantitative Finance*, 18(6), 933–949.
- ▣ Hansen, P.R., Huang, Z., Shek, H.H. (2012). *Realized GARCH: a joint model for returns and realized measures of volatility*. *Journal of Applied Econometrics*, 27(6), 877–906.
- ▣ Hansen, P.R., Lunde, A. (2005). *A forecast comparison of volatility models: does anything beat a GARCH(1,1)?*. *Journal of Applied Econometrics*, 20(7), 873–889.
- ▣ Hansen, P.R., Lunde, A. (2006). *Realized variance and market microstructure noise*. *Journal of Business & Economic Statistics*, 24(2), 127–161.
- ▣ Hansen, P.R., Lunde, A., Nason, J.M. (2011). *The model confidence set*. *Econometrica*, 79(2), 453–497.
- ▣ Hayashi, T., Yoshida, N. (2005). *On covariance estimation of non-synchronously observed diffusion processes*. *Bernoulli*, 11(2), 359–379.
- ▣ Huang, X., Tauchen, G. (2005). *The relative contribution of jumps to total price variance*. *Journal of Financial Econometrics*, 3(4), 456–499.
- ▣ Jacod, J., Li, Y., Mykland, P.A., Podolskij, M., Vetter, M. (2009). *Microstructure noise in the continuous case: the pre-averaging approach*. *Stochastic Processes and their Applications*, 119(7), 2249–2276.
- ▣ Lee, S.S., Mykland, P.A. (2008). *Jumps in financial markets: a new nonparametric test and jump dynamics*. *Review of Financial Studies*, 21(6), 2535–2563.
- ▣ Liu, L.Y., Patton, A.J., Sheppard, K. (2015). *Does anything beat 5-minute RV? A comparison of realized measures across multiple asset classes*. *Journal of Econometrics*, 187(1), 293–311.
- ▣ Mancini, C. (2009). *Non-parametric threshold estimation for models with stochastic diffusion coefficient and jumps*. *Scandinavian Journal of Statistics*, 36(2), 270–296.

References (5/5)

- ▣ Merton, R.C. (1980). On estimating the expected return on the market: an exploratory investigation. *Journal of Financial Economics*, 8(4), 323–361.
- ▣ Mincer, J., Zarnowitz, V. (1969). The evaluation of economic forecasts. In J. Mincer (Ed.), *Economic Forecasts and Expectations*, NBER, 3–46.
- ▣ Müller, U.A., Dacorogna, M.M., Davé, R.D., Olsen, R.B., Pictet, O.V., von Weizsäcker, J.E. (1997). Volatilities of different time resolutions: analyzing the dynamics of market components. *Journal of Empirical Finance*, 4(2–3), 213–239.
- ▣ Mykland, P.A., Zhang, L. (2009). Inference for continuous semimartingales observed at high frequency. *Econometrica*, 77(5), 1403–1445.
- ▣ Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- ▣ Patton, A.J., Sheppard, K. (2015). Good volatility, bad volatility: signed jumps and the persistence of volatility. *Review of Economics and Statistics*, 97(3), 683–697.
- ▣ Pele, D.T., Mazurencu-Marinescu-Pele, M. (2019). Using high-frequency entropy to forecast Bitcoin's daily value at risk. *Entropy*, 21(2), 102.
- ▣ Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets?. *European Journal of Finance*, 29(10), 1064–1105.
- ▣ Shephard, N., Sheppard, K. (2010). Realising the future: forecasting with high-frequency-based volatility (HEAVY) models. *Journal of Applied Econometrics*, 25(2), 197–231.
- ▣ Zhang, L. (2006). Efficient estimation of stochastic volatility using noisy observations: a multi-scale approach. *Bernoulli*, 12(6), 1019–1043.
- ▣ Zhang, L., Mykland, P.A., Aït-Sahalia, Y. (2005). A tale of two time scales: determining integrated volatility with noisy high-frequency data. *Journal of the American Statistical Association*, 100(472), 1394–1411.