



Modelling Financial Markets

Chapter 12: Options and the Volatility Surface



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model); in event studies: Abnormal Return
- ▣ ATM: At-The-Money
- ▣ COVID-19: COronaVirus Disease 2019
- ▣ CRR: Cox–Ross–Rubinstein (binomial tree)
- ▣ DVOL: Deribit Volatility Index
- ▣ ETF: Exchange-Traded Fund
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GBM: Gradient Boosting Machine
- ▣ GJR-GARCH: Glosten–Jagannathan–Runkle GARCH
- ▣ GMM: Generalized Method of Moments
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HAR: Heterogeneous AutoRegressive (model)
- ▣ HAR-RV: Heterogeneous AutoRegressive model of Realised Volatility
- ▣ ITM: In-The-Money
- ▣ MA: Moving Average
- ▣ MZ: Mincer–Zarnowitz (regression)
- ▣ OLS: Ordinary Least Squares
- ▣ OTM: Out-of-The-Money
- ▣ QLIKE: Quasi-LIKElihood (loss function)
- ▣ RV: Realised Variance
- ▣ SSVI: Surface SVI (Gatheral–Jacquier parameterisation of the whole surface)
- ▣ SVI: Stochastic Volatility Inspired (parameterisation)
- ▣ SVIX: Simple VIX: index of simple variance swaps (Martin, 2017)
- ▣ US: United States
- ▣ UTC: Coordinated Universal Time (Temps Universel Coordonné)

Acronyms used in this material (2/2)

- **VIX: CBOE Volatility Index**
- **VIX3M: Cboe 3-Month Volatility Index**
- **VIX9D: Cboe 9-Day Volatility Index**
- **VRP: Variance Risk Premium**
- **VVIX: Cboe VIX of VIX Index**
- **XFG: Quantlet prefix of Applied Quantitative Finance (Härdle, Chen & Overbeck)**

Chapter Route and Outcomes

- An **option** gives the right, not the obligation, to buy (call) or sell (put) an asset at a fixed strike price
 - ▶ It pays only in some states: a bet on the **distribution**, not only on the direction; option prices reveal what the market expects about volatility and crashes
- Part I: pricing by replication and what a hedge leaves behind
 - ▶ Parity and static spanning; the tree and the Black–Scholes equation; the Greeks; discrete-hedging error and the gamma-weighted volatility P&L
 - ▶ Implied volatility, the smile, and a full Bitcoin surface with its risk-neutral density
- Part II: implied volatility as information
 - ▶ Variance swaps and the VIX; the variance risk premium; is the VIX a good forecast? 0DTE options (zero days to expiry)
 - ▶ Case study: are option-implied bounds on the equity premium valid and tight? AI for scientific discovery
- Outcome: estimate a surface and a density under no-arbitrage constraints, with bands, and test option-implied forecasts with the right standard errors

Chicago, 1973: Listed Options Begin



Trading floor of the Chicago Board of Trade in 2006; its members founded the Cboe in 1973

Photo: [Richie Diesterheft \(2006\)](#); CC BY 2.0; [Wikimedia](#)

[Commons](#)

- The Chicago Board Options Exchange (Cboe) opened in 1973 as the first exchange for standardised, listed options [MacKenzie & Millo \(2003\)](#)
 - ▶ Before: options were negotiated one by one, over the counter
- The same year: the Black–Scholes and Merton pricing papers [Black & Scholes \(1973\)](#), [Merton \(1973\)](#)
- Traders soon used the formula on the floor: theory changed the market it described [MacKenzie & Millo \(2003\)](#)
- Textbooks: [Franke, Härdle & Hafner \(2019\)](#), [Härdle, Chen & Overbeck \(2017\)](#), [Gatheral \(2006\)](#)

The Formula and Its Authors




Fischer Black (1938–1995)

Photo: Dalmatine at English Wikipedia;
public domain; Wikimedia Commons



Myron Scholes (2008)

Photo: Nobellaureatesphotographer (2008);
CC BY-SA 3.0; Wikimedia Commons



Robert C. Merton (2006)

Photo: Digarnick (2006); public domain;
Wikimedia Commons

- Key idea: a portfolio of the stock and cash can **replicate** the option, so its price follows from no-arbitrage [Black & Scholes \(1973\)](#)
- Scholes and Merton received the [1997 Nobel Memorial Prize in Economic Sciences](#); Black had died in 1995

Data Used in This Chapter

Series	What it measures	Period	Owner
S&P 500	US equity index, daily close	1990–2026	S&P Dow Jones Indices
SPY	ETF on the S&P 500: open and close	1993–2026	State Street
VIX	30-day implied volatility of the S&P 500	1990–2026	Cboe
VIX9D, VIX3M	the same over 9 days and 3 months	2011 / 2007–2026	Cboe
VVIX	implied volatility of the VIX	2006–2026	Cboe
Bitcoin options	all listed contracts, one dated snapshot	22 Sep 2026	Deribit
DVOL	30-day implied volatility of Bitcoin	2021–2026	Deribit

- Daily data up to 18 September 2026; the option snapshot was taken on 22 September 2026, 07:51 UTC
- Implied volatility is quoted in annualised %; “vol points” = percentage points of volatility

Notation and Prerequisites

- European call and put pay off $(S_T - K)^+$ and $(K - S_T)^+$ at expiry T ; $\tau = T - t$ in years; American options can be exercised at any time up to T
 - ▶ Both pay-offs are convex in S_T : the value of this convexity is what the smile prices
- Forward $F = Se^{(r-q)\tau}$, log-moneyness $k = \ln(K/F)$; ATM: $k \approx 0$; OTM: puts with $K < F$ and calls with $K > F$; ITM: the opposite
 - ▶ Most of this chapter works in k and in total implied variance $w = \sigma_{imp}^2 \tau$
- Model-free bounds (no dividends): $\max(S - Ke^{-r\tau}, 0) \leq C \leq S$ and $\max(Ke^{-r\tau} - S, 0) \leq P \leq Ke^{-r\tau}$
 - ▶ With $r > 0$ an American call on a stock without dividends is never exercised early: it is worth at least $S - Ke^{-r\tau} > S - K$
- Assumed: Brownian motion, Itô's lemma and martingales (Chapter 11); regression with HAC standard errors (Chapter 8)

Put–Call Parity

- Two portfolios with the same pay-off at T (European options, same K and T , no dividends):
 - ▶ A: one call + cash Ke^{-rT} (worth K at T) $\Rightarrow \max(S_T, K)$
 - ▶ B: one put + one share $\Rightarrow \max(K, S_T)$
- Same pay-off in every state, so the same price today, otherwise there is an arbitrage:
 - ▶ $C_t - P_t = S_t - Ke^{-rT}$ (with dividend yield q : S_te^{-qT} instead of S_t)
- No model for S is needed: parity uses only no-arbitrage [Merton \(1973\)](#)
- Consequence: a call and a put with the same K and T have the **same** implied volatility



Static Replication and Its Corollaries

- Carr–Madan spanning: any twice-differentiable pay-off $f(S_T)$ and any point F (Taylor with integral remainder; proof in the Appendix):
 - ▶ $f(S_T) = f(F) + f'(F)(S_T - F) + \int_0^F f''(K)(K - S_T)^+ dK + \int_F^\infty f''(K)(S_T - K)^+ dK$
 - ▶ Priced today (F the forward): $V_0 = e^{-r\tau} f(F) + \int_0^F f''(K) P(K) dK + \int_F^\infty f''(K) C(K) dK$, whatever the model for S [Bakshi, Kapadia & Madan \(2003\)](#)
- **Variance:** $f(S) = -2 \ln(S/F)$ has $f''(K) = 2/K^2$: the strip with weights $1/K^2$ is the variance swap and the VIX (Section 8)
 - ▶ Low strikes get the largest weights, so an expensive put wing makes variance expensive
- **Density:** a butterfly $[C(K - h) - 2C(K) + C(K + h)]/h^2 \rightarrow \partial^2 C / \partial K^2 = e^{-r\tau} q(K)$ [Breedon & Litzenberger \(1978\)](#): the risk-neutral density is a price (Section 7)
- Straddles, spreads and protective puts are finite strips with f'' concentrated at one or two strikes; moments and SVIX: Appendix

No Arbitrage and the Pricing Measure

- One period, $S \rightarrow uS$ or dS , $d < R = e^{r\Delta t} < u$: $\Delta = \frac{C_u - C_d}{(u - d)S}$ shares plus cash replicate the call, so
 - ▶ $C = \frac{q C_u + (1 - q) C_d}{R}$, $q = \frac{R - d}{u - d}$: q is the unique probability under which S_t/R^t is a martingale
- The real probability of an up move never enters: risk preferences are already in S
- **Fundamental theorem:** no arbitrage $\Leftrightarrow$ an equivalent martingale measure Q exists; complete market $\Leftrightarrow Q$ is unique [Harrison & Pliska \(1981\)](#), [Delbaen & Schachermayer \(1994\)](#)
 - ▶ With jumps or stochastic volatility the market is incomplete: many Q , and option prices carry risk premia (Section 8)
- CRR tree [Cox, Ross & Rubinstein \(1979\)](#): $u = e^{\sigma\sqrt{\Delta t}}$, $d = 1/u$, value backwards from T ; American options: $\max(\text{continuation}, \text{exercise})$ at every node

From the Tree to Black-Scholes



- Left: CRR price of a European call for $N = 1, \dots, 200$ steps; right: American vs European put
- Converges to Black-Scholes (10.451): $N = 50$: 10.411; $N = 200$: 10.441; $N = 1000$: 10.449
 - The error is of order $1/N$ and oscillates: $N = 10$ gives 10.253, $N = 11$ gives 10.611 (strike on or between nodes; Seminar 12, A4)
- American put at $S = K = 100$: 6.090 vs European 5.574; early-exercise premium 0.516: deep in the money it pays to receive K now

The Black–Scholes Model and Equation


- Geometric Brownian motion $dS_t = \mu S_t dt + \sigma S_t dW_t$: log returns over τ follow the Normal distribution with variance $\sigma^2 \tau$; constant σ and r , continuous trading, no costs, short selling allowed
 - Replicate $V(S, t)$ with a **self-financing** portfolio of Δ_t shares and cash B_t at r : $d\Pi = \Delta dS + rB dt$
- Itô's lemma: $dV = (V_t + \mu S V_S + \frac{1}{2} \sigma^2 S^2 V_{SS}) dt + \sigma S V_S dW_t$
- Choose $\Delta = V_S$, start from $\Pi = V$ ($B = V - S V_S$): $d(V - \Pi) = (V_t + \frac{1}{2} \sigma^2 S^2 V_{SS} - rB) dt$; both dW_t and μ drop out
- A riskless position that costs nothing must earn nothing, hence the **Black–Scholes equation** Black & Scholes (1973), Merton (1973)
 - $V_t + \frac{1}{2} \sigma^2 S^2 V_{SS} + rS V_S - rV = 0$, with $V(S, T) = \text{pay-off}$
- That μ is irrelevant is a result, not an assumption; Feynman–Kac: $V = e^{-r\tau} E^Q[\text{pay-off}]$ with S drifting at r under Q

The Black–Scholes Formula

- European call and put, no dividends (with yield q : replace S by $Se^{-q\tau}$):
 - ▶ $C = S\Phi(d_1) - Ke^{-r\tau}\Phi(d_2), \quad P = Ke^{-r\tau}\Phi(-d_2) - S\Phi(-d_1)$
 - ▶ $d_1 = \frac{\ln(S/K) + (r + \sigma^2/2)\tau}{\sigma\sqrt{\tau}}, \quad d_2 = d_1 - \sigma\sqrt{\tau},$ with Φ the standard Normal distribution function
- Reading the two terms:
 - ▶ $\Phi(d_2)$ = risk-neutral probability that the call ends in the money
 - ▶ $\Phi(d_1)$ = number of shares in the replicating portfolio (the delta)
- Five inputs: S, K, τ, r observed; σ is the only one not observed: it must be estimated or implied



The Greeks: One Identity, All Sensitivities

- Delta $\Delta = \Phi(d_1)$ (put: $\Phi(d_1) - 1$); gamma $\Gamma = \varphi(d_1)/(S\sigma\sqrt{\tau})$, φ the standard Normal density; key identity $S\varphi(d_1) = Ke^{-r\tau}\varphi(d_2)$
- Vega $= S\varphi(d_1)\sqrt{\tau} = \Gamma S^2\sigma\tau$: for one expiry gamma and vega are proportional (Seminar 12, A5–A6); theta $\Theta = rC - rS\Delta - \frac{1}{2}\sigma^2 S^2\Gamma$; rho $= K\tau e^{-r\tau}\Phi(d_2)$



- Call with $K = 100$; short maturities concentrate gamma and theta at the strike, long maturities carry the most vega; vega per vol point, theta per calendar day

Discrete Delta Hedging: Algorithm and Error




- Sell one call at its Black–Scholes price, buy Δ_0 shares, keep the rest in cash; rebalance N times; **hedging error** = shares + cash – pay-off at expiry (GBM, $\mu = 8\%$, $\sigma = 20\%$, three months, 20 000 paths)
- Mean error near 0; s.d. monthly 40.8%, weekly 21.5%, daily 9.5%, four times a day 4.8% of the premium; worst 1%: -2.68 (weekly), -1.17 (daily)
- Log-log slope -0.48 ; theory (Appendix): s.d. $\approx \sqrt{\pi/4} \mathcal{V}\sigma/\sqrt{N}$, 1.01, 0.44, 0.22 vs simulated 0.96, 0.43, 0.22 ($N = 12, 63, 252$); costs grow like N

Derivation: the Gamma-Weighted Volatility P&L


- Sell the option at σ_{imp} and hold $\Delta = \partial_S V^{imp}$; the price moves with σ_{real} : $dS = \mu S dt + \sigma_{real} S dW_t$
- Itô on V^{imp} : $dV = V_t dt + \Delta dS + \frac{1}{2} \Gamma \sigma_{real}^2 S^2 dt$; the equation at σ_{imp} gives $V_t = rV - rS\Delta - \frac{1}{2} \sigma_{imp}^2 S^2 \Gamma$
- Short option + Δ shares + cash at r , self-financing: $d \text{P\&L} = r \text{P\&L} dt + \frac{1}{2} \Gamma_t S_t^2 (\sigma_{imp}^2 - \sigma_{real}^2) dt$, so over the life
 - ▶ $\text{P\&L} = \int_0^T e^{r(T-t)} \frac{1}{2} \Gamma_t S_t^2 (\sigma_{imp}^2 - \sigma_{real,t}^2) dt$: path-dependent, weighted by Γ , large near K and near T
- A short variance swap pays $\sigma_{imp}^2 - \sigma_{real}^2$ with constant weights; a delta-hedged option is a variance bet with random weights
- First-order size: vega $\times$ (implied – realised); here vega = 0.197 per vol point, so 5 vol points $\approx$ 0.99

Selling Volatility at 20%: Three Outcomes



- Three-month call sold at 20% implied volatility, hedged daily; realised volatility 15%, 20% or 25%
- Realised 15%: mean P&L +0.98, positive on 99.8% of paths
- Realised 20%: mean +0.00: only discrete-hedging noise remains (s.d. 0.43)
- Realised 25%: mean -0.99, positive on only 2.2% of paths
- Each mean is close to $\text{vega} \times 5 \text{ vol points}$ (± 0.99): a clean bet on volatility

Thirty-Six Years of Selling S&P 500 Volatility



- Every 21 trading days: sell a one-month ATM call priced with the VIX, hedge it daily at the same volatility; P&L in % of the index; interest and dividends ignored
- 440 months from 1990-01: mean 0.48% a month, Newey–West $t = 16.0$ [Newey & West \(1987\)](#); positive in 86% of months; VIX above the realised volatility of the month in 84% of months (19.3 vs 15.3)
- Worst month -4.25% (2020-03), skewness -1.46 : many small gains, a few large losses; consistent with a **volatility risk premium**, but actual seller returns need option quotes and costs [Bakshi & Kapadia \(2003\)](#)

Implied Volatility

- **Implied volatility** σ_{imp} : the σ that makes the Black–Scholes price equal the market price
 - ▶ $BS(S, K, \tau, r, \sigma_{imp}) = C^{market}$; unique because vega > 0 : the price increases in σ
- Newton's method: $\sigma_{k+1} = \sigma_k - \frac{BS(\sigma_k) - C^{market}}{\mathcal{V}(\sigma_k)}$
 - ▶ Example: $S = K = 100$, three months, $r = 2\%$, market price 4.50; start at $\sigma_0 = 20.00\%$
 - ▶ $BS(\sigma_0) = 4.2322$, $\mathcal{V} = 19.848 \Rightarrow \sigma_1 = 21.3495\%$; the next step changes only the 6th decimal: $\sigma_{imp} = 21.3495\%$
- Traders quote options in implied volatility, not in money: it is comparable across strikes, maturities and assets
- If Black–Scholes were right, all options on one asset would have the **same** implied volatility

After October 1987: the Smile Appears



Amsterdam stock exchange after the crash, 21 October
1987

Photo: [Bart Molendijk / Anefo \(1987\)](#); CC0; [Wikimedia Commons](#)

- Before 1987 index options had a nearly flat implied volatility across strikes [Rubinstein \(1994\)](#)
 - ▶ Since the crash, low-strike puts are much more expensive: the skew
- Option prices showed crash fears from October 1986 to August 1987, but not strongly in the two months immediately before the crash [Bates \(1991\)](#)
- **Smile**: implied volatility as a function of K ; **term structure**: as a function of T ; **surface**: both
- Crash risk on the Bucharest Stock Exchange: [Pele & Mazurencu-Marinescu \(2012\)](#)


Why Smiles Exist and How Models Produce Them

- **Fat tails:** large moves are more frequent than the Normal distribution allows (Chapter 1), so deep OTM options are worth more than Black–Scholes says
- **Leverage and volatility feedback:** when prices fall, volatility rises [Christie \(1982\)](#); negative skewness makes low-strike puts dearer
- **Crash fear and demand for insurance:** investors pay a premium for puts; intermediaries need compensation for crash risk they cannot hedge [Bates \(1991\)](#)

Model	Idea	Fits best
Jumps Merton (1976)	Poisson jumps added to the Brownian motion	steep short-dated smiles
Stochastic volatility Heston (1993) , Hull & White (1987)	variance follows its own mean-reverting process, correlated with S	skew and term structure of longer maturities
Local volatility Gatheral (2006)	$\sigma(S, t)$ read off the whole surface (Dupire's formula)	matches every quoted price today, exotic options
Parametric smile, SVI Gatheral & Jacquier (2014)	five parameters per expiry, no dynamics	interpolation and quoting (Section 7)


- **Processes:** Chapter 11; the implied-volatility surface in detail: [Franke, Härdle & Hafner \(2019\)](#), [Čížek, Härdle & Weron \(2011\)](#)

Three Sources of a Smile



- One-month implied volatility from the historical S&P 500 monthly returns (rescaled to a martingale, $r = 0$, no risk adjustment), a jump model and a stochastic-volatility model
- History: 9 225 overlapping returns, skewness -1.27 , excess kurtosis 6.8: 30.6% at $K/S = 0.8$, 13.9% ATM, 19.3% at 1.2
- Merton jumps (one crash of about -8% a year): 31.0% vs 16.3%; Heston, $\rho = -0.7$: 28.6% vs 16.6% vs 13.1%
- Risk-neutral index distributions are more left-skewed than realised ones [Jackwerth \(2000\)](#): priced crash risk steepens market smiles further

The Leverage Effect in the Data



- Daily S&P 500 log return vs the change of the VIX on the same day, 1990–2026
- Correlation -0.79 over 9 244 days: implied volatility rises when the market falls
- Slope -1.16 VIX points per 1% return, steeper on down days (-1.39) than on up days (-1.01); mean VIX change $+0.88$ vs -0.76
- The $\rho < 0$ of Heston and the asymmetry of GJR-GARCH (Chapter 5); low-strike puts gain from both the fall and the rise in volatility, which is why they are expensive


Bitcoin Options and the SVI Smile

- Deribit lists European options on Bitcoin, traded 24/7; **inverse** options: premium and pay-off in BTC; each expiry has its own forward F ; DVOL is its 30-day implied-volatility index [Deribit DVOL](#)
 - ▶ Our snapshot: 22 September 2026, 07:51 UTC; 9 expiries with at least two days left, 375 quoted OTM options; crypto risk: [Pele et al. \(2023\)](#); hedging Bitcoin options: [Matic, Packham & Härdle \(2023\)](#)
- Work with the **total implied variance** $w(k) = \sigma_{imp}^2(k) \tau$ at log-moneyness $k = \ln(K/F)$
 - ▶ SVI (stochastic volatility inspired) [Gatheral & Jacquier \(2014\)](#): $w(k) = a + b \left(\rho(k - m) + \sqrt{(k - m)^2 + s^2} \right)$
 - ▶ a level, b slope of the wings, ρ tilt (negative = put skew), m shift, s curvature at the bottom; wing slopes $b(1 \pm \rho)$
- One SVI fit per expiry by least squares on w ; the fit must then be checked against the no-arbitrage constraints (next slide)

No-Arbitrage Constraints on the Smile

- **Wings:** Lee (2004) bounds the asymptotic slope, $\limsup_{|k| \rightarrow \infty} w(k)/|k| \leq 2$; for SVI $b(1 + |\rho|) \leq 2$ is necessary, and $w \geq 0$ needs $a + bs\sqrt{1 - \rho^2} \geq 0$
 - ▶ A steeper wing means fatter tails of S_T under Q (moment formula: Appendix)
- **Butterfly** (call prices convex in K , density $q \geq 0$), with $w' = \partial_k w$ Gatheral & Jacquier (2014):
 - ▶ $g(k) = \left(1 - \frac{kw'}{2w}\right)^2 - \frac{w'^2}{4} \left(\frac{1}{w} + \frac{1}{4}\right) + \frac{w''}{2} \geq 0$ for all k , plus $d_+(k) = -k/\sqrt{w} + \sqrt{w}/2 \rightarrow -\infty$ as $k \rightarrow \infty$ (Gatheral & Jacquier (2014), Lemma 2.2)
- **Calendar:** $\partial_\tau w(k, \tau) \geq 0$ for every k : total variance cannot fall with maturity
- Slice-by-slice least squares imposes none of them; SSVI, the surface version of SVI, has simple sufficient conditions (Gatheral & Jacquier (2014), Theorems 4.1–4.2)

The Bitcoin Smile




- Mark implied volatilities (mark IV: the reference implied volatility computed by the exchange) of OTM options (dots) and SVI fits (lines) for five expiries
- A true **smile**: both wings rise, unlike the one-sided equity skew; Bitcoin can jump up as well as down [Zulfiqar & Gulzar \(2021\)](#)
- Short expiries have the steepest wings: a jump over a few days matters more relative to $\sigma\sqrt{T}$
- SVI root mean squared error at most 0.76 vol points per expiry; puts slightly dearer than calls: skew from -5.2 to -0.3 vol points

Are the Nine Bitcoin Fits Arbitrage-Free?

- ▣ Wings: largest slopes 0.31 (right) and 0.39 (left), far below 2: Lee's bound holds
- ▣ Butterfly: $\min_k g(k) > 0$ on the grid $k \in [-1.5, 1.5]$ for all 9 expiries (lowest 0.036, 27 November 2026)
- ▣ Calendar: 1 of 8 adjacent pairs crosses: $w(2 \text{ October } 2026) > w(9 \text{ October } 2026)$ on 24% of $k \in [-0.5, 0.5]$, by at most 0.0075
 - ▶ A calendar arbitrage created by the estimator, not quoted by the market: separate fits ignore the cross-expiry restriction
- ▣ Remedy: estimate the surface jointly under the constraints (SSVI, or least squares with $g \geq 0$ and $\partial_\tau w \geq 0$)

Term Structure and Skew




- Left: ATM implied volatility by expiry and the DVOL index; right: skew = implied volatility at $k = +0.15$ minus at $k = -0.15$
- ATM between 36.4% and 39.8%; longest expiry (276 days) 39.8%: upward sloping, calm now and more uncertainty later
- DVOL 38.1 vs ATM 37.0% and 36.4% around 30 days: a variance-strip index need not lie between ATM volatilities
- Skew negative for 9 of 9 expiries, closest to 0 at long maturities; one snapshot is one day

The Risk-Neutral Density: From Prices to a Density

- Breeden & Litzenberger (1978): the second derivative of call prices in the strike gives the risk-neutral density q of S_T
 - ▶ $q(K) = e^{r\tau} \frac{\partial^2 C(K)}{\partial K^2}$; convexity of C in K (no butterfly arbitrage) is exactly $q \geq 0$
- From noisy, discrete quotes this is an ill-posed inverse problem: differentiating twice amplifies the noise; fit a smooth smile (SVI), then differentiate (other estimators: Appendix)
- Uncertainty: pairs bootstrap of the quotes, refit, differentiate: pointwise bands for q and intervals for tail probabilities
- Beyond the last quoted strike the data identify nothing: the tail is an extrapolation; report where the quotes end
- q mixes probabilities and risk aversion: pricing kernel $\propto q/p$ with p the physical density Aït-Sahalia & Lo (1998), Grith, Härdle & Park (2013); its non-monotone shape is the pricing-kernel puzzle Jackwerth (2000)

What the Market Prices for Bitcoin



- Density from the SVI smile of the expiry closest to 30 days (38 days, forward 85 742 USD, ATM 36.4%), 95% pairs-bootstrap band, vs the log-normal density; shaded: outside the 56 quotes ($0.47F$ to $1.34F$)
- Fall of more than 30%: 1.21% [1.08, 1.24] vs 0.14% log-normal, about 8 times more; fall of more than 20%: 4.28% vs 3.27%; rise above 20%: 5.3% vs 5.3%
- Narrow band ($\pm 2.5\%$ at the centre, $\pm 8\%$ at $0.7F$): smoothed mark volatilities, so fitting noise, not bid-ask uncertainty; 0.13 of the 1.21 percentage points lie below the lowest quoted strike
- Prices of insurance, not forecasts: they include the premium for crash risk

 MFM_ch12_btc_rnd

Variance Swaps and the Option Strip

- A **variance swap** pays at T : notional $\times$ (realised variance – strike K_{var}), with realised variance $\frac{252}{n} \sum r_t^2$; the fair strike is the risk-neutral expected variance [Demeterfi, Derman, Kamal & Zou \(1999\)](#)
 - ▶ $K_{var} = \frac{2e^{r\tau}}{\tau} \left(\int_0^F \frac{P(K)}{K^2} dK + \int_F^\infty \frac{C(K)}{K^2} dK \right)$ (spanning with $f = -2 \ln(S/F)$, Section 2)
- Why it works: Itô on $\ln S$ (continuous paths, $r = q = 0$): $-2 \ln(S_T/F) = \int_0^T \sigma_t^2 dt - 2 \int_0^T dS_t/S_t$, and the last term is a Q -martingale [Britten-Jones & Neuberger \(2000\)](#)
 - ▶ “Model-free” means free of a model for σ_t , not of continuity: with crash jumps the strip understates expected realised variance [Carr & Wu \(2009\)](#) (Appendix)
- Buyer of variance = insurance against turbulence; seller = earns a premium, pays in crashes [Carr & Wu \(2009\)](#)


The VIX Index

- Cboe Volatility Index (VIX), since 1993; since 2003 computed from S&P 500 options with the variance-swap formula [Cboe \(2025\)](#)
 - ▶ $\sigma^2 = \frac{2}{T} \sum_i \frac{\Delta K_i}{K_i^2} e^{RT} Q(K_i) - \frac{1}{T} \left(\frac{F}{K_0} - 1 \right)^2$, K_0 = first strike at or below F ; Q = bid-ask mid of the put for $K_i < K_0$, of the call for $K_i > K_0$, the average of both at K_0
 - ▶ $\Delta K_i = (K_{i+1} - K_{i-1})/2$, one-sided at the two ends; options with a zero bid are dropped, and each wing stops after two consecutive zero bids
- Two expiries around 30 days, interpolated to exactly 30 days; $VIX = 100 \sigma$
- VIX futures since 2004, VIX options since 2006 [Cboe \(2025\)](#); the “fear gauge” [Whaley \(2000\)](#)
- Siblings: VIX9D (9 days), VIX3M (3 months), VVIX (the VIX of the VIX)

Worked Example: a VIX for Bitcoin


- Cboe rules on our Deribit snapshot: USD mid $Q = \text{BTC mid} \times \text{index } S; F$ and $e^{RT} = F/S$ from Deribit's forward (Cboe: F from parity, R from Treasury yields)
 - ▶ Near expiry: 17 days, 26 OTM options, $\sigma = 39.5\%$; next expiry: 38 days, 56 OTM options, $\sigma = 40.2\%$
- Interpolate the total variances $\sigma_i^2 T_i$ linearly in time to 30 days:
 - ▶ $\sigma_{30}^2 = \frac{1}{30} \left[T_1 \sigma_1^2 \frac{T_2 - 30}{T_2 - T_1} + T_2 \sigma_2^2 \frac{30 - T_1}{T_2 - T_1} \right]$ (days)
- Our Bitcoin VIX: 40.0; Deribit's DVOL at the same moment: 38.1 [Alexander & Imeraj \(2020\)](#); both exceed the ATM volatilities: the variance swap also pays for the wings
- Truncation and strike spacing move the index by at most 0.6 vol points, less than the gap to DVOL (1.9): an index is an estimator, so report its sensitivity (Appendix; Seminar 12, A8)

The VIX, 1990–2026



- Daily close with the six highest peaks at least 500 days apart
- Mean 19.4 (5.6% a month), range 9.1–82.7 (2020-03-16, COVID-19); above 30 on 7.9% of days: sudden spikes, slow decay, as in GARCH (Chapter 5)


The VIX Term Structure



- Top: implied volatility over 9 days, 30 days and 3 months; bottom: ratio VIX / VIX3M, 2011–2026
- Normally upward sloping (**contango**): mean $VIX / VIX3M = 0.89$; VIX9D below the VIX on 73% of days
- Inverted (**backwardation**) on 7.6% of days, always in crises: fear is concentrated in the near term
- Next 21 days after backwardation: return 2.88%, realised volatility 26.7%, vs 0.80% and 13.6% in contango: high risk, high expected return after panics [Bollerslev, Tauchen & Zhou \(2009\)](#)

 MFM_ch12_vix_term

The Volatility of Volatility




- VIX (top) and VVIX (bottom), daily closes, 2006–2026
- Mean VVIX 93.5: the VIX is expected to move about 5.3 times more than the index, in relative terms
- Highest value 207.6 on 2020-03-16, the day of the VIX peak; daily changes of VIX and VVIX correlate 0.71
- Volatility is itself random: the core of stochastic-volatility models such as Heston (1993)

The Variance Risk Premium


- Variance risk premium (VRP) Carr & Wu (2009), Bollerslev, Tauchen & Zhou (2009); three versions that must not be confused:
 - ▶ **Economic (ex ante):** $E_t^Q[RV_{t,t+21}] - E_t^P[RV_{t,t+21}]$; VIX_t^2 estimates the first term, the second needs a forecast
 - ▶ **Ex post** (used here): $VIX_t^2 - RV_{t,t+21}$, the realised P&L of the seller of a variance swap, known only at $t + 21$
 - ▶ **Ex-ante proxy:** $VIX_t^2 - RV_{t-21,t}$, past variance as the forecast of E_t^P ; known at t , so it can serve as a predictor (Section 11)
 - ▶ $RV_{t,t+21} = \frac{252}{21} \sum_{i=1}^{21} r_{t+i}^2$ (annualised, in $\%^2$)
- Positive on average if investors dislike states with high volatility, which are also bad states for wealth
- Overlapping 21-day windows are autocorrelated: Newey–West with at least 21 lags Newey & West (1987); persistent volatility can extend the dependence, so check longer bandwidths

Implied vs Realised Volatility, S&P 500



- Top: VIX and the volatility realised over the next 21 trading days; bottom: their difference, 1990–2026
- 9 203 days: mean VIX 19.4, mean realised volatility 15.3; premium 4.09 vol points (Newey–West s.e. 0.27); in variance 113 %², $t = 5.5$
- Positive on 86% of days; the exceptions are the crashes: worst -72.0 vol points, window starting 2020-02-20 (COVID-19)
- Same message as the delta-hedged option: sellers of volatility are paid on average and lose in crises

Is the VIX a Good Forecast?




- Mincer–Zarnowitz [Mincer & Zarnowitz \(1969\)](#):
 $RV_{t,t+21} = \alpha + \beta VIX_t^2 + \varepsilon_t$; unbiased means $\alpha = 0$, $\beta = 1$;
 Newey–West errors, 21 lags
- $\hat{\alpha} = -51.9$ (34.4), $\hat{\beta} = 0.86$ (0.10), $R^2 = 0.37$; separately
 neither differs from (0, 1): $t = -1.51$ and -1.34
- But the estimates correlate -0.90 : joint Wald $\chi^2_2 = 41.7$,
 $p = 8.9 \cdot 10^{-10}$: unbiasedness rejected; most points lie
 below the 45-degree line
- Past realised variance alone: $R^2 = 0.29$; with both, the VIX
 keeps its weight (0.77), past variance adds little (0.08, s.e.
 0.13)

Logs, Instruments and HAR-RV: Inference Done Right

Specification	$\hat{\beta}$	s.e.	$t(\beta = 1)$	Note
Levels, OLS	0.86	0.10	-1.34	joint Wald $p = 8.9 \cdot 10^{-10}$
Logs, OLS	1.06	0.035	1.65	intercept -0.94 , $R^2 = 0.57$
Logs, instrumental variables	1.02	0.042	0.60	first-stage $t = 37.7$

- Levels are dominated by 2008 and 2020; in logs the slope is close to 1 and the premium is a proportional shift (intercept < 0)
- VIX^2 measures expected variance with error, which attenuates OLS: instrument it with VIX^2 21 days earlier [Christensen & Prabhala \(1998\)](#); the slope barely moves (valid only if the lag is uncorrelated with the forecast error, untestable with one instrument)
- Against HAR-RV [Corsi \(2009\)](#): in sample the HAR terms add information ($\chi^2_3 = 14.2$, $p = 0.003$); out of sample from 2000 the VIX is about 6% better in QLIKE loss, Diebold–Mariano $p = 0.35$ [Diebold & Mariano \(1995\)](#): not significant
 - Test details: Appendix

The Variance Risk Premium of Bitcoin




- DVOL and the Bitcoin volatility realised over the next 30 calendar days, 2021–2026; bottom: their difference
- 1975 days: mean DVOL 60.3 vs realised 51.7; premium 8.5 vol points (s.e. 1.6, $t = 5.3$): twice the S&P 500 premium in vol points, similar relative to the level
- Positive on 71% of days, less often than for the S&P 500; worst -46.9 vol points (window starting 2026-01-28)
- A short sample: 66 non-overlapping 30-day periods; a Bitcoin VIX and its premium: Alexander & Imeraj (2020)

 MFM_ch12_vrp_btc

Options That Live One Day


- ▣ **0DTE** (zero days to expiry) options: traded on the day they expire [Cboe 0DTE](#)
 - ▶ S&P 500 options now expire every trading day, so a 0DTE contract exists every day
- ▣ More than 75% of retail trades in S&P 500 options involve 0DTE contracts, and retail investors lose money on them [Beckmeyer, Branger & Gayda \(2023\)](#)
 - ▶ 0DTE options were 48% of index-option trading in August 2023 [Brogaard, Han & Won \(2023\)](#)
- ▣ More 0DTE trading is followed by higher index volatility, mostly through speculative retail trades [Brogaard, Han & Won \(2023\)](#); ultra-short pricing needs skewness and kurtosis corrections [Bandi, Fusari & Renò \(2023\)](#)
- ▣ Theory from Section 4: with $\tau \rightarrow 0$ gamma and theta explode at the money

Gamma and Time Decay Near Expiry



- Left: gamma of a call with $K = 100$ one day, one week and one month before expiry; right: daily time decay of an ATM call
- ATM gamma one day before expiry 0.396 vs 0.086 one month before: 4.6 times larger ($= \sqrt{21}$); a 1% move in the last day changes delta by about 0.4
- Time decay 0.167 in the last day vs 0.030 with 30 days left, about 5.5 times more: the seller collects fast decay and carries large gamma risk

Selling a Same-Day Straddle: a Stylised Test



- SPY, 2011–2026: sell an ATM straddle at the open, settle at the close; premium = $\sqrt{2/\pi} \sigma_{day} S$, σ_{day} from the previous VIX9D scaled to the session (share 0.74 of daily variance, estimated on 1993–2010)
- 3 950 days: mean premium 0.76% vs mean absolute move 0.55%; the seller wins on 75% of days, 20.7 basis points a day
- Worst day –8.25% (2025-04-09), about 11 premia; stylised pricing measures the premium inside VIX9D, not traded 0DTE prices [Bandi, Fusari & Renò \(2023\)](#), [Beckmeyer, Branger & Gayda \(2023\)](#)

Case Study: the Research Question

- **Question Back, Crotty & Kazempour (2022):** option prices give lower bounds on the equity premium; are they valid and tight conditionally, and do they forecast returns out of sample?
 - ▶ Why: a tight bound is a real-time equity-premium forecast with no estimated parameter; **Martin (2017)** could not reject $\alpha = 0$, $\beta = 1$ on 1996–2012, a test with low power
- **Martin bound** (eq. 3), $R_{t,T}$ the gross market return, $R_{f,t,T}$ the gross riskless return, var_t^* the risk-neutral variance ($\text{SVIX}_{t,T}^2$ not annualised here):
 - ▶ $E_t[R_{t,T}] - R_{f,t,T} \geq R_{f,t,T} \text{SVIX}_{t,T}^2$, $\text{SVIX}_{t,T}^2 = \text{var}_t^*(R_{t,T}/R_{f,t,T})$
 - ▶ Valid if $\text{cov}_t(M_{t,T}R_{t,T}, R_{t,T}) \leq 0$ ($M_{t,T}$ the stochastic discount factor); tight if this covariance is zero, e.g. a log-utility investor who holds the market
- **Chabi-Yo–Loudis bound** **Chabi-Yo & Loudis (2020)**: built from the first four risk-neutral moments, free of preference parameters in its restricted form; almost always above the Martin bound (Figure 3.1)
- **Link to this chapter:** SVIX is the static-spanning strip with $f = S^2$ (Section 2 and Appendix)
 - ▶ Realised slackness $s_{t,T} = R_{t,T}^e - b_{t,T}$ (excess return minus bound): valid means $E_t[s_{t,T}] \geq 0$, tight means $E_t[s_{t,T}] = 0$

Case Study: Data and Method


- **Data** (Section 3, Appendix C): daily Martin and Chabi-Yo–Loudis bounds, January 1990 – December 2020, from S&P 500 index options, interpolated to 30, 90, 180 and 360 days; S&P 500 excess returns over 21, 63, 126 and 252 trading days, annualised, in % (filters: Appendix)
- **Moments** (eq. 10): for any nonnegative z_t known at t , $E[(R_{t,T}^e - b_{t,T}) z_t] \geq 0$; $z_t =$ a constant and nine Welch–Goyal predictors made positive [Welch & Goyal \(2008\)](#) (Table 3)
- **Kodde–Palm tests** [Kodde & Palm \(1986\)](#) (eqs. 11–14): $\bar{\lambda}$ the 10 sample means, Σ their Hansen–Hodrick covariance [Hansen & Hodrick \(1980\)](#), with lags equal to the days in the horizon
 - ▶ $D_1 = \min_{\lambda \geq 0} (\lambda - \bar{\lambda})^\top \Sigma^{-1} (\lambda - \bar{\lambda})$ tests validity ($H_0: \lambda_0 \geq 0$); $D_2 = \bar{\lambda}^\top \Sigma^{-1} \bar{\lambda} - D_1$ tests tightness ($H_0: \lambda_0 = 0$ vs $\lambda_0 \geq 0$); both follow mixtures of χ^2 laws
- **Forecasting**: bias-corrected predictive regressions [Amihud & Hurvich \(2004\)](#) (Tables 5, 6) and monthly out-of-sample forecasts against the historical mean (Table 8; R_{OS}^2 : out-of-sample R^2 against the historical mean; protocol: Appendix)

Case Study: Results (Tables 1, 3, 5, 6 and 8)

Horizon	Mean R^e (%)	Mean b (%)	D_1	D_2	$\hat{\beta}$ (s.e.)	R_{OS}^2 (%)
1 month	8.76	3.81	0.00	29.93	1.36 (1.23)	-1.2
3 months	8.73	3.86	0.00	39.51	1.39 (1.47)	-1.2
6 months	8.67	3.83	0.00	49.28	1.78 (0.85)	-0.2
12 months	8.86	3.68	0.00	71.21	1.04 (0.88)	-3.0

- Source: [Back, Crotty & Kazempour \(2022\)](#), *Journal of Financial Economics*, Tables 1, 3, 5, 6 and 8
- **Validity not rejected:** all 10 sample moments are positive, so $D_1 = 0$ and $p = 100\%$ at every horizon; this does not prove validity in every state
- **Not tight:** D_2 exceeds 27.13, the upper 0.1% Kodde–Palm critical value, at every horizon (finite-sample p : 3.6, 0.4, 2.6, 8.3%); mean slackness 4.8 to 5.2 points a year: the bound is less than half of the realised premium
- **Forecasting:** slopes positive but imprecise (Table 5); with the nine predictors as controls a one-s.d. rise in the 12-month bound goes with a 6.6-point higher 12-month excess return, but bootstrap p -values are 20.7, 0.0, 0.2, 9.5% (Table 6)
- **Out of sample:** the Martin bound loses to the historical mean at every horizon; the best case, Chabi-Yo–Loudis at 6 months, reaches 4.3%, not significant (Table 8): low power is one possible explanation


Case Study: the Bound and the Realised Premium



- Top: the authors' daily 1-month bounds, 1990–2020 ([replication package](#)); bottom, course data: 12-month excess return of the S&P 500 price index vs the 12-month Martin bound
- 1-month bound: median 2.7%, peak 54.4% (16 March 2020); 12-month peak 21.1% (20 November 2008): a risk-neutral variance, it jumps with the VIX in crises
- Chabi-Yo–Loudis at or above Martin on 99.99% of days: negative skewness and fat tails raise it
- Realised 12-month excess return below the bound on 33% of dates (2000–2002, 2007–2008): the bound restricts $E_t[R_{t,T}^e]$, not each realisation
- Correlation bound–realised 0.15 (overlapping, descriptive): return noise dwarfs the bound, hence imprecise slopes and negative R_{OS}^2

 MFM_ch12_bck_bounds

Case Study: Mean Premium and Mean Bound



- Dark blue: Table 1 of the paper (S&P 500 with dividends); teal: course data (price index, the same formation days: 7789 at 1 month, 7598 at 12 months); red: mean of the authors' Martin bound
- Mean slackness on the course data 2.79, 2.69, 2.63, 2.91 points a year vs 4.8–5.2 in the paper: the gap of 2.2–2.3 points is mostly the missing dividend yield
- The mean bound covers 41–44% of the premium with dividends, 56–59% of the price-index premium; single realisations fall below it on 41% (1 month) to 33% (12 months) of days, the average never (Seminar 12, C3)

Case Study: Replication, Limitations and an Open Extension

- **Replicate** (Seminar 12, C3): the authors' daily bound series ([replication package](#)), S&P 500 daily closes of the course data, the daily risk-free rate ([Kenneth French Data Library](#)), the Welch–Goyal predictors ([Amit Goyal's data page](#))
 - ▶ Rebuild Table 1 and D_1 , D_2 of Table 3; the course index has no dividends, so excess returns and slackness are lower: a data limitation
- **Limitations stated by the authors**
 - ▶ A short series, limited to the years after 1990 by the need for option prices: the out-of-sample tests have little power (Section 6)
 - ▶ Validity holds even when the bound is replaced by its time-series mean, so it says nothing about predictive power (Section 4.3)
 - ▶ More recent bounds and formulas for risk premia are not tested (Introduction)
- **Open extension:** is the Martin bound still valid on 2021–2026 (the 2022 bear market, the 0DTE era of Section 9)? The course data have VIX, not SVIX: VIX weights OTM options by $1/K^2$ and exceeds SVIX when the left tail is priced ([Martin \(2017\)](#), Result 3); is a VIX-based bound still valid, and can the Kodde–Palm test detect it if not?

An Open Question: Does the VRP Still Predict Returns?

- [Bollerslev, Tauchen & Zhou \(2009\)](#), Table 2: monthly data 1990–2007, $y_t^{(h)} = \frac{12}{h} \sum_{j=1}^h (r_{m,t+j} - r_{f,t+j})$ on $\text{VRP}_t = \text{VIX}_t^2/12 - \text{RV}_t$
 - ▶ RV_t : realised variance of month t ; at $h = 3$: slope 0.47, Hodrick $t = 2.86$, adjusted $R^2 = 6.82\%$, the peak of a hump-shaped horizon profile
- [Goyal, Welch & Zafirov \(2024\)](#): 29 equity-premium predictors proposed after 2008, re-examined with data to 2021: over one third lose in-sample significance, half of the rest fail out of sample
- Open: is VRP predictability a stable price of variance and jump risk, or a property of 1990–2007? The 2008–2026 sample adds two crashes and the 0DTE era
- Formally: $H_0 : \beta_3 = 0$ vs $H_1 : \beta_3 > 0$ in $y_t^{(3)} = \alpha + \beta_3 \text{VRP}_t + u_t$
 - ▶ Out of sample: H_0 : the VRP forecast does not lower the mean squared prediction error of the historical mean; Clark–West statistic [Clark & West \(2007\)](#)

The Discovery Loop and What the Researcher Must Check

- **Literature:** ask an assistant for peer-reviewed tests of VRP predictability since 2009 (DOI, sample, table)
 - ▶ Check: resolve every DOI on Crossref and compare the cited table with the paper, digit by digit
- **Hypotheses:** time-varying risk aversion, jump-risk compensation, or a small-sample artefact of persistent regressors and overlapping returns
 - ▶ Check: each must imply something testable for the horizon profile, for 2008–2026 and for non-US markets
- **Data and code:** build $VIX^2/12$ and monthly RV from the daily VIX and S&P 500 closes of the course data, reproduce Table 2 ($h = 3$) before extending anything
 - ▶ Check timing (VIX at the end of month t , returns from $t + 1$) and match the published means (implied 33.23, realised 14.93)
- **Robustness and critique:** a dated pre-analysis plan (horizons, start date, benchmark, Holm correction); a hostile-referee prompt
 - ▶ Check overlap (Hodrick errors [Hodrick \(1992\)](#)), horizon and start-date mining (R_{OS}^2 against the start date), the weight of 2008 and 2020; full prompts: Appendix

Mini-Case: Replicate, Then Extend

- A typical assistant proposal: "the VRP predicts the quarterly equity premium; extend the 2009 result to today". We test it on the course data
- Replication, February 1990 to September 2007, $h = 3$, daily data: $\hat{\beta}_3 = 0.42$, $t_{\text{Hodrick}} = 2.25$, adjusted $R^2 = 4.6\%$ (published: 0.47, 2.86, 6.82%)
 - ▶ Mean implied variance 33.1 is close to 33.23; mean realised 20.6 vs 14.93: daily close-to-close data, a stated limitation
- Extension to May 2026 (436 months): $\hat{\beta}_3 = 0.09$, $t_{\text{Hodrick}} = 1.33$, adjusted $R^2 = 1.0\%$; all horizons with four kinds of standard errors: Appendix
- Out of sample from January 2008, 221 expanding-window forecasts: $R_{OS}^2 = 1 - \sum_t (y_t - \hat{y}_t)^2 / \sum_t (y_t - \bar{y}_t)^2 = -22.2\%$ [Campbell & Thompson \(2008\)](#); Clark–West $t = 0.64$ ($p = 0.26$): H_0 not rejected
- Reading: sign and size replicate in-sample; the longer sample does not support stable predictability. Outliers, measurement or a vanished premium: that is the open part

Your Turn: a Project Seed

- Question: does the VRP predict the equity premium after 2008, and does the Bitcoin premium of Section 8 (DVOL) predict Bitcoin returns?
- Replicate first: [Bollerslev, Tauchen & Zhou \(2009\)](#), Table 2, $h = 3$: slope 0.47, $t = 2.86$, adjusted $R^2 = 6.82\%$; explain every gap
- Deliverables
 - ▶ A dated pre-analysis plan; the replication table; out-of-sample results for all horizons with a Holm correction
 - ▶ Robustness (5-minute variance, subsamples, without 2008 and 2020); a referee-style critique and your reply
- Log each prompt that shaped a decision in `AI_USE.md`; each invented reference, wrong number or look-ahead caught in `AI_ERRORS.md`, with the test that caught it ([templates](#))

Key Takeaways

- Parity, bounds and static spanning need no model; the tree and the Black–Scholes equation price by dynamic replication
- The Greeks measure exposures; discrete hedging leaves an error of order $1/\sqrt{N}$
- A delta-hedged option is a bet on realised vs implied variance; implied variance exceeded realised variance on average on the S&P 500 and on Bitcoin
- The surface and the risk-neutral density are estimates: impose no-arbitrage constraints and report bands
- The VIX is a strip estimator with jump and truncation biases; as a forecast it needs joint tests, logs and HAC loss comparisons
- Option-implied bounds on the equity premium are valid but not tight in 1990–2020: evidence rests on average slackness, not on forecasting power

Frontier, Next Chapter and Further Study

▣ Frontier

- ▶ Andersen, Fusari & Todorov (2015): a left-tail factor in index-option surfaces, not spanned by volatility, predicts future equity and variance premia
- ▶ Bali et al. (2023): nonlinear machine learning on option and stock characteristics predicts equity-option returns out of sample
- ▶ Open: jointly estimated, arbitrage-free surfaces with valid confidence bands; the pricing of 0DTE options

▣ Next: Chapter 13, machine learning for financial prediction, with volatility and option-implied features among the inputs

▣ Further reading: Franke, Härdle & Hafner (2019), Härdle, Chen & Overbeck (2017), Čížek, Härdle & Weron (2011), Gatheral (2006); pricing kernels: Ait-Sahalia & Lo (1998), Grith, Härdle & Park (2013); crypto options: Matic, Packham & Härdle (2023), Alexander & Imeraj (2020), Pele et al. (2023)

- ▶ Quantinar:  XFG Advanced Methods in Quantitative Finance,  Pricing Kernels and Risk Premia,  Hedging Cryptocurrency Options
- ▶ Quantlets: github.com/danpele/MFM/tree/main/Quantlets/Ch_12

Appendix: Proof of Spanning, Moments and SVIX

- Proof for $S_T > F$: $\int_F^{S_T} f''(K)(S_T - K) dK = f(S_T) - f(F) - f'(F)(S_T - F)$ by parts; the put integral vanishes; symmetric for $S_T < F$
- Take $e^{-r\tau} E^Q[\cdot]$ with F the forward (a forward contract costs 0): a bond, a forward and a strip of OTM options with weights $f''(K)$ replicate any smooth pay-off
- **Moments:** $f = [\ln(S/F)]^n$, $n = 2, 3, 4$, price the raw risk-neutral moments of the log return; centring and standardising gives volatility, skewness and kurtosis [Bakshi, Kapadia & Madan \(2003\)](#)
- **SVIX** [Martin \(2017\)](#), from $f = S^2$: $SVIX_t^2 = \frac{2}{\tau R_f S_t^2} \left[\int_0^F P dK + \int_F^\infty C dK \right] = \text{var}_t^Q(R_T/R_f)/\tau$, with R_T, R_f gross returns
 - If $\text{cov}_t(M_T R_T, R_T) \leq 0$ (M_T the stochastic discount factor): $\frac{1}{\tau}(E_t R_T - R_f) \geq R_f SVIX_t^2$, a lower bound on the equity premium (case study, Section 10)

Appendix: Why $1/\sqrt{N}$? The Discrete Hedging Error

- Over one interval $\Delta t = T/N$ the hedged book earns $\frac{1}{2}\Gamma S^2 [\sigma^2 \Delta t - (\Delta S/S)^2]$: time decay against realised squared return
- Mean 0 under the model to leading order in Δt ; $\text{Var}[(\Delta S/S)^2] \approx 2\sigma^4 \Delta t^2$ and the terms are approximately uncorrelated, so $\text{Var}(\text{error}) \approx \frac{1}{2}\sigma^4 \Delta t E \int_0^T \Gamma_t^2 S_t^4 dt$
- ATM call: $E[\Gamma_t^2 S_t^4] \approx \frac{K^2}{2\pi\sigma^2\sqrt{T^2 - t^2}}$ integrates to $K^2/(4\sigma^2)$, so s.d. $\approx \sigma K \sqrt{T/(8N)} = \sqrt{\pi/4} \mathcal{V}\sigma/\sqrt{N}$
 - With vega $\mathcal{V} = 19.72$ per unit: theory 1.01, 0.44, 0.22 vs simulated 0.96, 0.43, 0.22 for $N = 12, 63, 252$
- Asymptotic theory of the replication error: [Bertsimas, Kogan & Lo \(2000\)](#); it fails with jumps (the error does not vanish as $N \rightarrow \infty$), with stochastic volatility (vega risk) and with costs

Appendix: the Moment Formula for the Wings

- Moment formula of [Lee \(2004\)](#): $\beta_R = \limsup_{k \rightarrow \infty} w(k)/k \in [0, 2]$ and $\tilde{p} = \frac{1}{2\beta_R} + \frac{\beta_R}{8} - \frac{1}{2}$
 - ▶ $\tilde{p} = \sup\{p : E^Q[S_T^{1+p}] < \infty\}$: a steeper wing means fatter tails; the left wing mirrors it with negative moments
- Bitcoin: the largest right slope 0.31 gives $\tilde{p} = 1.18$; a property of SVI's linear extrapolation beyond the last strike, not of the data
- For SVI a right-wing slope below 2 also gives the tail condition $d_+(k) \rightarrow -\infty$ of [Gatheral & Jacquier \(2014\)](#), Lemma 2.2

Appendix: Estimators of the Risk-Neutral Density

- Fit a smile (SVI) and differentiate the implied call prices twice, as in the lecture
- Kernel regression of C on (K, τ) Aït-Sahalia & Lo (1998); the same under monotonicity and convexity in K Aït-Sahalia & Duarte (2003)
- Smoothness-penalised implied trees Jackwerth & Rubinstein (1996); shape-invariant pricing kernels Grith, Härdle & Park (2013)
- Tails beyond the quotes: SVI wings or an extreme-value tail; either way an assumption, not an estimate

Appendix: What the Strip Measures When Prices Jump

- With jumps J : the strip prices $\int \sigma_t^2 dt + 2 \sum (e^J - 1 - J)$, but realised variance is $\int \sigma_t^2 dt + \sum J^2$
 - ▶ Gap $2E^Q \sum (e^J - 1 - J - J^2/2) \approx E^Q \sum J^3/3$: negative for crashes, so the strip understates expected realised variance Carr & Wu (2009)
- Merton model of Section 6 (one jump a year, mean -8%): strip 17.82 vs $\sqrt{E^Q[\text{QV}]} = 18.00$ vol points, -2.0% in variance

Appendix: Truncation and Discretisation of the Bitcoin Index

- Only a finite set of strikes is quoted: truncating the strip biases the index down, a coarse grid biases it up [Jiang & Tian \(2005\)](#)
- Bitcoin, same two expiries, 30-day index: quoted strikes 40.03; SVI prices at the same strikes 39.99
 - ▶ SVI on a dense grid inside the quoted range 39.78; extended to $k \in [-4, 4]$ with the SVI wings 40.34
- The near expiry is quoted only from $0.76F$ to $1.17F$: its volatility rises from 39.05 to 40.00 when the wings are added
- The rest of the gap to DVOL reflects its quote selection and methodology, not quantified here

Appendix: Encompassing and Out-of-Sample Tests

- Heterogeneous autoregressive model (HAR-RV) [Corsi \(2009\)](#): past daily, weekly and monthly realised variance; encompassing regression in logs [Busch, Christensen & Nielsen \(2011\)](#)
 - ▶ VIX slope 0.90 (0.056); HAR terms jointly $\chi^2_3 = 14.2$, $p = 0.003$; R^2 : VIX 0.570, HAR 0.474, both 0.575
- Out of sample from 2000 (6 697 days; coefficients re-estimated every 21 days on targets already observed)
 - ▶ QLIKE loss $L = RV/\hat{\sigma}^2 - \ln(RV/\hat{\sigma}^2) - 1$, robust to noise in the RV proxy [Patton \(2011\)](#)
 - ▶ Loss ratio VIX^2 / HAR : 0.94, Diebold–Mariano $t = -0.93$ ($p = 0.35$) [Diebold & Mariano \(1995\)](#); MZ-corrected VIX: 0.95 ($p = 0.35$)
- Targets overlap by 20 of 21 days: the Diebold–Mariano variance is HAC with 21 lags; instrumental variables by just-identified GMM with a HAC covariance

Appendix: Does the Ex-Ante VRP Predict Returns? Overlapping-Horizon Inference

Horizon	$\hat{b}$	t_{NW}	t_{HH}	$t_{Hodrick}$	p_{boot}	R^2 (%)
$h = 1$	0.07	0.33	0.36	0.41	0.35	0.2
$h = 3$	0.09	0.69	0.67	1.33	0.07	1.2
$h = 6$	0.03	0.49	0.48	0.65	0.26	0.3
$h = 12$	-0.02	-0.64	-0.79	-0.54	0.69	0.2

- Monthly 1990–2026: $\frac{12}{h} \sum_{j=1}^h r_{t+j}^e = a + b \text{VRP}_t + u_t$, $\text{VRP}_t = \text{VIX}_t^2/12 - \text{RV}_t$, ex-ante proxy with past RV_t as the forecast [Bollerslev, Tauchen & Zhou \(2009\)](#)
- If one-month errors are serially uncorrelated, overlap makes u_t an $\text{MA}(h-1)$: Newey–West (h lags), Hansen–Hodrick [Hansen & Hodrick \(1980\)](#), Hodrick 1B [Hodrick \(1992\)](#) (better size); these bandwidths are choices, so check their sensitivity
- p_{boot} : one-sided, VRP simulated as $\text{AR}(1)$ ($\hat{\phi} = 0.13$) under $b = 0$, residuals resampled jointly [Stambaugh \(1999\)](#); no horizon is significant on the full sample (Section 11 revisits 1990–2007)

Appendix: Case Study, Filters and Forecasting Protocol

- ▣ Option data: Cboe for 1990–1995 and 2020–2021, OptionMetrics for 1996–2019
- ▣ Filters: standard options only (no weekly or quarterly), bid > 0 , the OTM option at each strike, 7 to 550 days, at least 10 strikes per expiry
- ▣ Returns: S&P 500 with dividends minus the Fama–French risk-free rate, up to March 2021
- ▣ Out of sample (Table 8): monthly forecasts against the expanding post-1926 market mean, 60-month initial window
 - ▶ $R_{OS}^2 = 1 - \sum \varepsilon_t^2 / \sum \nu_t^2$, ε_t and ν_t the errors of the forecast and of the benchmark; Diebold–Mariano tests Diebold & Mariano (1995)

Appendix: The Discovery Loop, Reusable Prompts (1/2)

- **Literature:** "List peer-reviewed tests of VRP return predictability since 2009: DOI, sample, horizon, table, headline estimate."
 - ▶ Resolve every DOI on Crossref and read the cited table yourself; search with Semantic Scholar, Google Scholar, Elicit
- **Hypotheses:** "Give three competing explanations of VRP predictability and the testable implication of each for the horizon profile, for 2008-2026 and for non-US markets."
- **Data and code:** "Write Python that builds VIX squared/12 and monthly realised variance from the daily VIX and S&P 500 closes of the course data and reproduces Table 2 of BTZ (2009), $h = 3$, with Hodrick (1992) standard errors."
 - ▶ Any assistant (Claude, ChatGPT, Gemini, GitHub Copilot); replicate the published number before extending anything

Appendix: The Discovery Loop, Reusable Prompts (2/2)

- **Robustness**, fixed before the results: "Draft a pre-analysis plan: horizons, out-of-sample start date, benchmark, test statistic and the multiple-testing correction across horizons."
 - ▶ Placebo: VRP lagged 12 months or block-shuffled; subsamples; 5-minute realised variance of SPY (from 2020) instead of daily squared returns
- **Critique**: "Act as a hostile referee: why could this out-of-sample R^2 be spurious or biased, and which test settles each objection?"
 - ▶ Expect: bias with persistent regressors, overlap, the choice of start date, the weight of 2008 and 2020
- **Report**: every horizon and specification tried, in one table; the plan dated before the results; AI use logged

References (1/4)

- ▣ Aït-Sahalia, Y., Duarte, J. (2003). Nonparametric option pricing under shape restrictions. *Journal of Econometrics*, 116(1–2), 9–47.
- ▣ Aït-Sahalia, Y., Lo, A.W. (1998). Nonparametric estimation of state-price densities implicit in financial asset prices. *Journal of Finance*, 53(2), 499–547.
- ▣ Alexander, C., Imeraj, A. (2020). The Bitcoin VIX and its variance risk premium. *Journal of Alternative Investments*, 23(4), 84–109.
- ▣ Amihud, Y., Hurvich, C.M. (2004). Predictive regressions: a reduced-bias estimation method. *Journal of Financial and Quantitative Analysis*, 39(4), 813–841.
- ▣ Andersen, T.G., Fusari, N., Todorov, V. (2015). The risk premia embedded in index options. *Journal of Financial Economics*, 117(3), 558–584.
- ▣ Back, K., Crotty, K., Kazempour, S.M. (2022). Validity, tightness, and forecasting power of risk premium bounds. *Journal of Financial Economics*, 144(3), 732–760.
- ▣ Bakshi, G., Kapadia, N. (2003). Delta-hedged gains and the negative market volatility risk premium. *Review of Financial Studies*, 16(2), 527–566.
- ▣ Bakshi, G., Kapadia, N., Madan, D. (2003). Stock return characteristics, skew laws, and the differential pricing of individual equity options. *Review of Financial Studies*, 16(1), 101–143.
- ▣ Bali, T.G., Beckmeyer, H., Mörke, M., Weigert, F. (2023). Option return predictability with machine learning and big data. *Review of Financial Studies*, 36(9), 3548–3602.
- ▣ Bandi, F.M., Fusari, N., Renò, R. (2023). ODTE option pricing. SSRN Working Paper 4503344.
- ▣ Bates, D.S. (1991). The crash of '87: was it expected? The evidence from options markets. *Journal of Finance*, 46(3), 1009–1044.
- ▣ Beckmeyer, H., Branger, N., Gayda, L. (2023). Retail traders love ODTE options... but should they?. SSRN Working Paper 4404704.
- ▣ Bertsimas, D., Kogan, L., Lo, A.W. (2000). When is time continuous?. *Journal of Financial Economics*, 55(2), 173–204.
- ▣ Black, F., Scholes, M. (1973). The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3), 637–654.
- ▣ Bollerslev, T., Tauchen, G., Zhou, H. (2009). Expected stock returns and variance risk premia. *Review of Financial Studies*, 22(11), 4463–4492.
- ▣ Breeden, D.T., Litzenberger, R.H. (1978). Prices of state-contingent claims implicit in option prices. *Journal of Business*, 51(4), 621–651.
- ▣ Britten-Jones, M., Neuberger, A. (2000). Option prices, implied price processes, and stochastic volatility. *Journal of Finance*, 55(2), 839–866.
- ▣ Brogaard, J., Han, J., Won, P.Y. (2023). Does ODTE options trading increase volatility?. SSRN Working Paper 4426358.
- ▣ Busch, T., Christensen, B.J., Nielsen, M.Ø. (2011). The role of implied volatility in forecasting future realized volatility and jumps in foreign exchange, stock, and bond markets. *Journal of Econometrics*, 160(1), 48–57.
- ▣ Campbell, J.Y., Thompson, S.B. (2008). Predicting excess stock returns out of sample: can anything beat the historical average?. *Review of Financial Studies*, 21(4), 1509–1531.

References (2/4)

- ▣ Carr, P., Wu, L. (2009). [Variance risk premiums](#). *Review of Financial Studies*, 22(3), 1311–1341.
- ▣ Cboe Global Indices (2025). [Cboe Volatility Index: Methodology](#). Cboe Global Markets.
- ▣ Cboe Global Markets. [ODTE options](#). Product page.
- ▣ Chabi-Yo, F., Loudis, J. (2020). [The conditional expected market return](#). *Journal of Financial Economics*, 137(3), 752–786.
- ▣ Christensen, B.J., Prabhala, N.R. (1998). [The relation between implied and realized volatility](#). *Journal of Financial Economics*, 50(2), 125–150.
- ▣ Christie, A.A. (1982). [The stochastic behavior of common stock variances: value, leverage and interest rate effects](#). *Journal of Financial Economics*, 10(4), 407–432.
- ▣ Čížek, P., Härdle, W.K., Weron, R. (Eds.) (2011). [Statistical Tools for Finance and Insurance](#), 2nd ed. Springer.
- ▣ Clark, T.E., West, K.D. (2007). [Approximately normal tests for equal predictive accuracy in nested models](#). *Journal of Econometrics*, 138(1), 291–311.
- ▣ Corsi, F. (2009). [A simple approximate long-memory model of realized volatility](#). *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Cox, J.C., Ross, S.A., Rubinstein, M. (1979). [Option pricing: a simplified approach](#). *Journal of Financial Economics*, 7(3), 229–263.
- ▣ Delbaen, F., Schachermayer, W. (1994). [A general version of the fundamental theorem of asset pricing](#). *Mathematische Annalen*, 300(1), 463–520.
- ▣ Demeterfi, K., Derman, E., Kamal, M., Zou, J. (1999). [A guide to volatility and variance swaps](#). *Journal of Derivatives*, 6(4), 9–32.
- ▣ Deribit. [DVOL: Bitcoin volatility index](#). Statistics page.
- ▣ Diebold, F.X., Mariano, R.S. (1995). [Comparing predictive accuracy](#). *Journal of Business & Economic Statistics*, 13(3), 253–263.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.
- ▣ Gatheral, J. (2006). [The Volatility Surface: A Practitioner's Guide](#). Wiley.
- ▣ Gatheral, J., Jacquier, A. (2014). [Arbitrage-free SVI volatility surfaces](#). *Quantitative Finance*, 14(1), 59–71.
- ▣ Goyal, A., Welch, I., Zafirov, A. (2024). [A comprehensive 2022 look at the empirical performance of equity premium prediction](#). *Review of Financial Studies*, 37(11), 3490–3557.
- ▣ Grith, M., Härdle, W.K., Park, J. (2013). [Shape invariant modeling of pricing kernels and risk aversion](#). *Journal of Financial Econometrics*, 11(2), 370–399.
- ▣ Hansen, L.P., Hodrick, R.J. (1980). [Forward exchange rates as optimal predictors of future spot rates: an econometric analysis](#). *Journal of Political Economy*, 88(5), 829–853.

References (3/4)

- ▣ Härdle, W.K., Chen, C.Y.-H., Overbeck, L. (Eds.) (2017). *Applied Quantitative Finance*, 3rd ed. Springer.
- ▣ Harrison, J.M., Pliska, S.R. (1981). Martingales and stochastic integrals in the theory of continuous trading. *Stochastic Processes and their Applications*, 11(3), 215–260.
- ▣ Heston, S.L. (1993). A closed-form solution for options with stochastic volatility with applications to bond and currency options. *Review of Financial Studies*, 6(2), 327–343.
- ▣ Hodrick, R.J. (1992). Dividend yields and expected stock returns: alternative procedures for inference and measurement. *Review of Financial Studies*, 5(3), 357–386.
- ▣ Hull, J., White, A. (1987). The pricing of options on assets with stochastic volatilities. *Journal of Finance*, 42(2), 281–300.
- ▣ Jackwerth, J.C. (2000). Recovering risk aversion from option prices and realized returns. *Review of Financial Studies*, 13(2), 433–451.
- ▣ Jackwerth, J.C., Rubinstein, M. (1996). Recovering probability distributions from option prices. *Journal of Finance*, 51(5), 1611–1631.
- ▣ Jiang, G.J., Tian, Y.S. (2005). The model-free implied volatility and its information content. *Review of Financial Studies*, 18(4), 1305–1342.
- ▣ Kodde, D.A., Palm, F.C. (1986). Wald criteria for jointly testing equality and inequality restrictions. *Econometrica*, 54(5), 1243–1248.
- ▣ Lee, R.W. (2004). The moment formula for implied volatility at extreme strikes. *Mathematical Finance*, 14(3), 469–480.
- ▣ MacKenzie, D., Millo, Y. (2003). Constructing a market, performing theory: the historical sociology of a financial derivatives exchange. *American Journal of Sociology*, 109(1), 107–145.
- ▣ Martin, I. (2017). What is the expected return on the market?. *Quarterly Journal of Economics*, 132(1), 367–433.
- ▣ Matic, J.L., Packham, N., Härdle, W.K. (2023). Hedging cryptocurrency options. *Review of Derivatives Research*, 26(1), 91–133.
- ▣ Merton, R.C. (1973). Theory of rational option pricing. *Bell Journal of Economics and Management Science*, 4(1), 141–183.
- ▣ Merton, R.C. (1976). Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1–2), 125–144.
- ▣ Mincer, J., Zarnowitz, V. (1969). The evaluation of economic forecasts. In *Economic Forecasts and Expectations*, NBER, 3–46.
- ▣ Newey, W.K., West, K.D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.
- ▣ Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- ▣ Pele, D.T., Mazurencu-Marinescu, M. (2012). Modelling stock market crashes: the case of Bucharest Stock Exchange. *Procedia – Social and Behavioral Sciences*, 58, 533–542.
- ▣ Pele, D.T., Wesselhöft, N., Härdle, W.K., Kolossatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets?. *European Journal of Finance*, 29(10), 1064–1105.

References (4/4)

- ▣ Rubinstein, M. (1994). Implied binomial trees. *Journal of Finance*, 49(3), 771–818.
- ▣ Stambaugh, R.F. (1999). Predictive regressions. *Journal of Financial Economics*, 54(3), 375–421.
- ▣ Welch, I., Goyal, A. (2008). A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 21(4), 1455–1508.
- ▣ Whaley, R.E. (2000). The investor fear gauge. *Journal of Portfolio Management*, 26(3), 12–17.
- ▣ Zulfikar, N., Gulzar, S. (2021). Implied volatility estimation of bitcoin options and the stylized facts of option pricing. *Financial Innovation*, 7, 67.