



Modelling Financial Markets

Chapter 14: Deep Learning and Time-Series Foundation Models



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material

- AI: Artificial Intelligence
- AR: AutoRegressive (model)
- BET: Bucharest Exchange Trading (index)
- CI: Confidence Interval
- COVID-19: COroNaVirus Disease 2019
- CPU: Central Processing Unit
- DM: Diebold–Mariano (test)
- ES: Expected Shortfall
- FHS: Filtered Historical Simulation
- FZ0: Fissler–Ziegel loss, zero-homogeneous version
- GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- GIFT: General Time Series Forecasting Model Evaluation (benchmark GIFT-Eval)
- GRU: Gated Recurrent Unit
- HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- HAR: Heterogeneous AutoRegressive (model)
- HS: Historical Simulation
- LLM: Large Language Model
- LSTM: Long Short-Term Memory (network)
- MCS: Model Confidence Set
- MDE: Minimum Detectable Effect
- MLP: MultiLayer Perceptron
- MSE: Mean Squared Error
- QLIKE: Quasi-LIKElihood (loss function)
- RNN: Recurrent Neural Network
- RV: Realised Variance
- US: United States

Outline

Why Deep Learning for Financial Time Series?

From the Perceptron to Deep Networks

Recurrent Networks and the LSTM

Attention and Transformers

Time-Series Foundation Models

Test 1: Forecasting Returns

Test 2: Forecasting Realised Volatility

Test 3: VaR 1% and ES 2.5%


AI for Scientific Discovery

Practice and Conclusions

References

The Promise and the Question

- **Deep learning:** models built from many layers of simple non-linear units, trained on data
 - ▶ It transformed image recognition and language processing
 - ▶ Since 2024: **foundation models** for time series, pretrained on billions of observations
- The question of this chapter
 - ▶ Do these models forecast **returns**, **volatility** and **tail risk** better than the classical models of Chapters 5–9?
 - ▶ We answer with the evaluation tools of Chapter 8




A data centre: deep networks are trained on racks of such servers

Photo: [BalticServers.com](#) (2013); CC BY-SA 3.0; [Wikimedia Commons](#)

Three Forecasting Targets

- ▣ **Returns** r_t (conditional mean)
 - ▶ Very low signal-to-noise ratio; predictable patterns are traded away (Chapter 2)
- ▣ **Volatility** (conditional variance)
 - ▶ Persistent and clustered (Chapters 1 and 5): there is a signal to learn
- ▣ **Tail quantiles**: VaR 1% and ES 2.5%
 - ▶ Need the shape of the lower tail, not just the centre of the distribution (Chapter 7)
- ▣ Each target gets its own loss function and its own benchmark model

Where Is the Signal?



- S&P 500 daily returns since 2000 (6 718 days): autocorrelation at lags 1 to 50
- Grey band: 95% range under independence, $\pm 1.96/\sqrt{T} = \pm 0.024$

Interpretation: Mean versus Size of Moves

- ▣ Returns: first-order autocorrelation -0.099
 - ▶ 15 of 50 lags outside the band, but small and of changing sign: little to learn about the mean
- ▣ Absolute returns: autocorrelation 0.29 at lag 1, still 0.13 at lag 50
 - ▶ All 50 lags outside the band: volatility clustering with long memory
- ▣ Prediction for the rest of the chapter: any flexible model can find the volatility signal; none should find much in the mean

What We Do in This Chapter


- ▣ From the perceptron to deep networks: neurons, layers, training
- ▣ Recurrent networks and the LSTM: why memory is hard to learn
- ▣ Attention and Transformers: patches, positions, quantile outputs
- ▣ Foundation models: Chronos, TimesFM, Moirai; zero-shot use and pretraining leakage
- ▣ Three empirical tests: returns, realised volatility, VaR 1% and ES 2.5%

Section Summary

- ▣ Three targets with very different signal-to-noise ratios: mean, volatility, tail
- ▣ Absolute returns are strongly autocorrelated, returns are not
- ▣ Every new model is judged against a simple benchmark with the tests of Chapter 8

The Perceptron (1958)

- **Artificial neuron:** weighted sum of inputs plus a bias, passed through an activation function
 - ▶ $a = \varphi(w^T x + b)$
 - ▶ w : weights, b : bias, φ : activation (non-linear function)
- **Rosenblatt (1958):** the perceptron learns w from labelled examples
 - ▶ Built in hardware as the Mark I Perceptron (photo)
 - ▶ One layer only separates classes with a straight line (hyperplane)



The Mark I Perceptron, Cornell Aeronautical Laboratory

Photo: J.C. Hay, A.E. Murray (1960); public domain; [Wikimedia](#)

[Commons](#)

The Multilayer Perceptron (MLP)

- Layers applied one after another: $h^{(\ell)} = \varphi(W^{(\ell)}h^{(\ell-1)} + b^{(\ell)})$, $h^{(0)} = x$
 - ▶ $h^{(\ell)}$: hidden state of layer ℓ ; the last layer is linear for regression
- Common activations
 - ▶ ReLU (rectified linear unit): $\max(0, z)$
 - ▶ $\tanh(z) \in (-1, 1)$; sigmoid $\sigma(z) = 1/(1 + e^{-z}) \in (0, 1)$
- **Universal approximation** (Hornik, Stinchcombe & White (1989)): one hidden layer with enough units approximates any continuous function
 - ▶ It says nothing about how much data is needed: the key issue in finance

Training a Network

- ▣ Loss: mean squared error (MSE) $\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2$ for point forecasts; pinball loss for quantiles
- ▣ **Backpropagation** (Rumelhart, Hinton & Williams (1986)): the chain rule gives the gradient of the loss with respect to every weight
- ▣ **Gradient descent**: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$, learning rate η
 - ▶ Mini-batches: gradient on a random subset of examples; **epoch**: one pass through the data
 - ▶ Adam (Kingma & Ba (2015)): step sizes adapted per parameter from running moments of the gradient

Overfitting and How Deep Learning Fights It

- ▣ Networks have many parameters relative to financial samples
 - ▶ Our small LSTM has 1 233 parameters; Chronos-2 has about 119 million
- ▣ **Early stopping**: stop training when the loss on a validation set stops falling
 - ▶ In time series the validation set must come after the training set (Chapter 13: purged cross-validation)
- ▣ **Weight decay**: penalty $\lambda \|\theta\|^2$; **dropout**: random switching-off of units during training
- ▣ **Ensembles**: average several networks trained from different random seeds

Deep Learning Comes of Age

- ▣ [LeCun, Bengio & Hinton \(2015\)](#): review of the methods behind the breakthroughs in vision and speech
- ▣ 2024: Nobel Prize in Physics to Hopfield and Hinton for foundational work on neural networks
- ▣ In finance: [Gu, Kelly & Xiu \(2020\)](#) find that neural networks improve monthly stock-return predictions through non-linear interactions (Chapter 13)
 - ▶ Cross-section of thousands of stocks, not one index: much more data



Geoffrey Hinton, Nobel laureate 2024


Photo: [Arthur Petron \(2024\)](#); CC BY-SA 4.0; [Wikimedia Commons](#)

Section Summary

- Neuron: weighted sum plus non-linear activation; MLP: layers of neurons
- Training: backpropagation plus gradient descent (Adam) on mini-batches
- Flexibility is cheap, data are not: early stopping, weight decay and ensembles are essential

Recurrent Neural Networks (RNN)

- A hidden state carries information forward in time
 - ▶ $h_t = \tanh(W h_{t-1} + U x_t + b)$, $\hat{y}_{t+1} = v^\top h_t + c$
 - ▶ The same weights W, U, b at every step: the number of parameters does not grow with the sequence length
- Unrolled over time, an RNN is a very deep network with shared weights



Why Long Memory Is Hard: the Vanishing Gradient

- Backpropagation through time multiplies one Jacobian matrix per step

- ▶ $\frac{\partial h_T}{\partial h_{T-k}} = \prod_{j=0}^{k-1} \text{diag}(1 - h_{T-j}^2) W$

- ▶ Scalar case: $|w| \tanh'(\cdot) \leq |w|$, so the gradient is at most $|w|^k$

- With $|w| = 0.9$: $0.9^{20} = 0.12$, $0.9^{50} = 0.005$

- ▶ Bengio, Simard & Frasconi (1994): gradient-based learning of long-range dependencies is difficult
- ▶ Exploding gradients (norms above 1) are the mirror problem; clipping the gradient norm is the usual fix

The LSTM Cell

- ▣ **Long short-term memory** (Hochreiter & Schmidhuber (1997)): a cell state c_t with gates that decide what to keep, write and output
 - ▶ Forget gate: $f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$
 - ▶ Input gate: $i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$; candidate $g_t = \tanh(W_g x_t + U_g h_{t-1} + b_g)$
 - ▶ Output gate: $o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$
- ▣ Updates: $c_t = f_t \odot c_{t-1} + i_t \odot g_t$, $h_t = o_t \odot \tanh(c_t)$
 - ▶ $\odot$: element-by-element product; gates take values in $(0, 1)$

Why the LSTM Remembers

- Along the cell state the only multiplier is the forget gate
 - ▶ $\frac{\partial c_T}{\partial c_{T-k}} = \prod_{j=0}^{k-1} f_{T-j}$ (direct path)
 - ▶ With $f \approx 1$ the gradient survives many steps
- Constant forget gate $f = \sigma(b_f)$
 - ▶ $b_f = 0$: $f = 0.5$, $0.5^{50} \approx 9 \times 10^{-16}$
 - ▶ $b_f = 3$: $f = 0.953$, $0.953^{50} \approx 0.09$; half-life 14 steps
- GRU (gated recurrent unit): a lighter variant with two gates




Sepp Hochreiter, co-inventor of the LSTM

Photo: Eulenreich (2015); CC BY-SA 4.0; Wikimedia

[Commons](#)

The Vanishing Gradient, Measured



- ▣ $|\partial h_T / \partial x_{T-k}|$ by automatic differentiation; 16 hidden units, 50 random initialisations, log scale
- ▣ The same random inputs for the three networks; only the architecture and the forget-gate bias change

Interpretation: Architecture Is Not Enough

- Simple RNN: gradient share 8.3×10^{-5} at lag 20 and 7.4×10^{-9} at lag 50
 - ▶ Practically nothing is learned from inputs more than a few weeks back
- LSTM with forget-gate bias 0: 1.3×10^{-9} at lag 50, no better than the RNN
 - ▶ At initialisation $f \approx 0.5$: the gate halves the memory at every step
- LSTM with forget-gate bias 3: 8.4% at lag 50, 9.0% at lag 100
 - ▶ Matches the constant-gate calculation $0.953^{50} \approx 0.09$
 - ▶ Practical rule: initialise the forget-gate bias at a positive value

How Big Is a Small LSTM?

- Parameters of one LSTM layer with input size d and hidden size h
 - ▶ $4 \times (hd + h^2 + h)$ (PyTorch uses two bias vectors: $4 \times (hd + h^2 + 2h)$)
 - ▶ Four blocks: three gates and the candidate
- Our return model: $d = 1$, $h = 16$, linear output: 1 233 parameters
 - ▶ Trained on at most 4,000 daily returns: about three observations per parameter
- Application in finance: [Fischer & Krauss \(2018\)](#) use LSTMs to predict whether S&P 500 constituents beat the cross-sectional median the next day

Section Summary

- RNN: one hidden state, shared weights; long memory fails because gradients vanish
- LSTM: additive cell state and gates; the forget gate controls how long information lives
- Even small LSTMs have thousands of parameters: a lot for one financial series


Attention: Looking Back Directly

- Each time step asks a question (query) and compares it with every past step (key)
 - ▶ $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V$ (Vaswani et al. (2017))
 - ▶ $Q = XW_Q$, $K = XW_K$, $V = XW_V$: linear projections of the inputs X ; d_k : key dimension
 - ▶ softmax turns scores into weights that sum to one; division by $\sqrt{d_k}$ keeps the scores from saturating it
- No product of Jacobians over time: any past step is one hop away
 - ▶ Cost grows with the square of the number of inputs n : $O(n^2)$

Positions and Patches

- **Positional encoding:** attention treats its inputs as a set, so the order must be added
 - ▶ Sinusoids of different frequencies added to the inputs (original Transformer)
 - ▶ Rotary embeddings: rotate queries and keys by an angle proportional to the position (used by Chronos-2)
- **Patching** (Nie et al. (2023)): a token is a block of consecutive observations, not a single one
 - ▶ 512 daily returns in patches of 16 days: 32 tokens instead of 512, about 256 times fewer attention scores
 - ▶ Each token summarises local dynamics (a few weeks of market)

Patching a 512-Day Context



- The last 512 daily returns of the S&P 500, split as Chronos-2 (16 days) and TimesFM-2.5 (32 days) read them
- Each shaded block becomes one embedding vector for the Transformer

Other Neural Architectures for Forecasting

- DeepAR ([Salinas et al. \(2020\)](#)): an RNN that outputs the parameters of a predictive distribution, trained on many related series
- N-BEATS ([Oreshkin et al. \(2020\)](#)): stacks of fully connected blocks that forecast and subtract, without recurrence or attention
- PatchTST ([Nie et al. \(2023\)](#)): Transformer on patches, each series processed separately
- Common thread: learn across many series, because one financial series is short
 - ▶ Foundation models push this idea to its limit: learn across all available series

Section Summary

- ▣ Attention compares every step with every other step: no vanishing gradient through time, but quadratic cost
- ▣ Positions must be encoded; patches reduce the number of tokens and carry local context
- ▣ Modern forecasting networks learn across many series


What Is a Foundation Model?

- ▣ A large network pretrained once on a very broad corpus, then used on new tasks
 - ▶ Idea borrowed from large language models (LLM): pretrain on text, then use on any question
- ▣ **Zero-shot** forecasting: no training on the target series; only its recent history (the **context**) is passed at forecast time
 - ▶ Output: a set of quantiles of the future values, not only a point
 - ▶ Fine-tuning (further training on the target series) is possible but not used here
- ▣ Appeal for finance: one model for thousands of series, no estimation per series

Chronos: Time Series as a Language

- ▣ Ansari et al. (2024): turn values into tokens and train a language model (T5) with cross-entropy
 - ▶ Mean scaling: divide the context by $s = \frac{1}{C} \sum_t |x_t|$
 - ▶ Quantisation: 4,093 equally spaced bin centres in $[-15, 15]$ (width 0.0073), plus special tokens
 - ▶ Forecast: sample future tokens, map back to values, read quantiles from the samples
- ▣ Pretraining corpus: public datasets (energy, transport, retail, weather, web, finance and economics) plus synthetic series (TSMixup, KernelSynth)
 - ▶ Model sizes 20 to 710 million parameters; the small version has 46 million

What Chronos Sees: Scaled Returns and Tokens



- Context: the last 512 S&P 500 daily returns; scale $s = 0.672\%$ (mean absolute return)
- Right: where the scaled values fall within the token range $[-15, 15]$

Interpretation: a Vocabulary Mostly Unused

- ▣ Scaled returns range from -9.2 to 13.5; 96% lie within ± 3
 - ▶ The 512 returns use 336 distinct tokens out of 4,093; ± 3 alone spans 818 bins
- ▣ Scaling makes the model free of units: 1% in a calm year and 3% in a crisis can map to the same token
 - ▶ Good for transfer across series; the level of volatility must be recovered from the context itself
- ▣ A crash day inflates s and compresses all other days towards zero (Seminar, A6)

Chronos-Bolt and Chronos-2

- ▣ **Chronos-Bolt**: patches of 16 values and a direct quantile output instead of token sampling
 - ▶ Quantile levels 10%, 20%, ..., 90% only; small version 48 million parameters
- ▣ **Chronos-2** ([Ansari et al. \(2025\)](#)): robust scaling, patches of 16, attention across time and across groups of related series
 - ▶ 21 quantile levels: 1%, 5%, 10%, ..., 95%, 99%; about 119 million parameters
 - ▶ Training corpus: selected datasets of the Chronos and [GIFT-Eval](#) corpora plus synthetic multivariate series
 - ▶ Benchmarks: [GIFT-Eval](#) ([Aksu et al. \(2024\)](#)) and fev-bench

TimesFM, Moirai and Lag-Llama

- **TimesFM** (Das et al. (2024)): decoder-only Transformer on patches of 32 values
 - ▶ Pretrained on about 10^{11} time points: Google Trends, Wikipedia page views, synthetic series, M4, electricity, traffic, weather
 - ▶ TimesFM-2.5: mean plus deciles; the released checkpoint has about 231 million parameters
- **Moirai** (Woo et al. (2024)): any number of variables, mixture-distribution output; Moirai 2.0 (Liu et al. (2025))
 - ▶ Released under a non-commercial licence; not used in our tests
- **Lag-Llama** (Rasul et al. (2023)): decoder-only model with lagged values as features and a Student- t output

The Three Models We Test

Model	Parameters	Patch	Quantile levels	Weights published
Chronos-Bolt (small)	48 M	16	10%–90%	25.11.2024
TimesFM-2.5	231 M	32	mean, 10%–90%	01.10.2025
Chronos-2	119 M	16	1%, 5%, . . . , 99%	30.10.2025

- All three run on an ordinary CPU; context of 512 daily observations, one-day-ahead forecasts
- Publication dates of the weights from the Hugging Face model repositories; they fix what the models could have seen
- Only Chronos-2 reaches the 1% level needed for VaR 1%

Pretraining Leakage

- A test is out of sample only if the test period was never seen during training
 - ▶ The Chronos corpus includes financial series (e.g. daily exchange rates 1990–2016); later corpora are broader
 - ▶ We cannot rule out that our series, or series driven by the same shocks, were in the pretraining data
- Our remedy: a **post-release window** from 3 November 2025, after all three sets of weights were published
 - ▶ Clean, but short (about 220 trading days; 320 calendar days for Bitcoin): low power, especially for VaR 1%
- Look-ahead in our own pipeline is avoided too: every scaling, window and estimate uses only past data


From a Quantile Grid to What We Need

- The model returns quantiles q_u at levels $u_1 < \dots < u_K$
 - ▶ Point forecast: the median $q_{0.5}$, or the mean $\int_0^1 q_u du$ approximated on the grid
 - ▶ Mean of a transformed variable: $E[g(Y)] = \int_0^1 g(q_u) du$ for increasing g (e.g. $g = \exp$ for log variance)
- Tail risk (Chapter 7): $\text{VaR}_\alpha = -q_\alpha$, $\text{ES}_\alpha = -\frac{1}{\alpha} \int_0^\alpha q_u du$
 - ▶ Below the lowest level the grid is silent: we hold q_u at its last value, which understates ES
 - ▶ Chronos-Bolt and TimesFM stop at 10%: a "VaR 1%" request returns the 10% quantile

Worked Example: the Last Chronos-2 Forecast

- S&P 500, forecast for the day after 18.09.2026: $q_{0.01} = -2.70\%$, $q_{0.05} = -1.40\%$
 - ▶ VaR 1% = $-q_{0.01} = 2.70\%$
 - ▶ Linear interpolation: $q_{0.025} = q_{0.01} + \frac{0.015}{0.04}(q_{0.05} - q_{0.01}) = -2.21\%$, so VaR 2.5% = 2.21%
- ES 2.5%: $-\frac{1}{0.025} \left[0.01 q_{0.01} + 0.015 \frac{q_{0.01} + q_{0.025}}{2} \right] = 2.56\%$
 - ▶ A lower bound: the tail below 1% is cut at $q_{0.01}$

A Zero-Shot Fan Chart for Volatility



- Chronos-2 on log realised variance of SPY (512-day context), 22 trading days ahead from 18.09.2026; annualised volatility
- Bands: quantiles 5%–95% and 25%–75% of the predictive distribution

Interpretation: Mean Reversion Learned Without Training

- Last observed volatility 4.8% p.a.; average of the last 20 days 6.6%
 - ▶ Median forecast at 22 days: 6.6%, with a 90% band from 3.8% to 15.4%
- The band widens with the horizon and is skewed upwards: volatility can jump, not collapse
 - ▶ Features of GARCH and HAR forecasts (Chapters 5 and 9), obtained here with no estimation on SPY data
- Whether these forecasts are **better** is an empirical question: next sections

Section Summary

- ▣ Foundation models: pretrained once on large corpora, used zero-shot with a context window
- ▣ They read a series through scaling, tokens or patches, and return quantiles
- ▣ Quantile grids limit tail estimates; pretraining overlap requires a post-release test window


Design of the Return Test

- Daily log returns (%): S&P 500 and BET from 04.01.2016 (2 693 and 2 675 days), Bitcoin from 01.01.2019 (2 818 days)
 - ▶ Every forecast uses only data up to the previous close
- Models: historical mean, AR(1) on a 1,000-day window, LSTM, Chronos-2 and Chronos-Bolt (median), TimesFM-2.5 (mean)
- Benchmark: the zero forecast; out-of-sample $R^2 = 1 - \sum_t (r_t - \hat{r}_t)^2 / \sum_t r_t^2$
 - ▶ Zero rather than the historical mean, as in [Gu, Kelly & Xiu \(2020\)](#); compare [Campbell & Thompson \(2008\)](#)
 - ▶ Inference: block-bootstrap CI for R^2 (blocks of 20 days), Diebold–Mariano (DM) test on squared errors with Newey–West variance ([Diebold & Mariano \(1995\)](#))

The LSTM for Returns

- ▣ Input: the last 20 returns, standardised with the mean and standard deviation of the training period
 - ▶ Output: next-day return; loss MSE; Adam; weight decay 10^{-4}
- ▣ Trained once on up to 4,000 days before the test period; last 20% of them for early stopping
 - ▶ Five random seeds; the forecast is the average of the five networks
- ▣ One hidden layer of 16 units: 1 233 parameters

Training the LSTM: Early Stopping




- S&P 500, seed 0: MSE of standardised returns by epoch, on the training and validation parts
- Training stops 26 epochs in; the weights of epoch 1 are kept

Interpretation: Little to Learn

- Both losses barely move; the validation loss is lowest after the first epoch, so early stopping keeps epoch 1
 - ▶ The validation MSE is below the training MSE because the validation years were calmer: the level reflects volatility, not skill
- Early stopping picks a network close to its starting point: forecasts near the mean
 - ▶ Standard deviation of the LSTM forecasts for the S&P 500: 0.047%, against 1.13% for returns
- A healthy sign: a network that forecast large moves would be fitting noise

Out-of-Sample R^2 for Returns




□ R^2 against the zero forecast (%), with 95% block-bootstrap CIs

Interpretation: No Model Beats Zero, Except on the BET

R^2 (%) [DM p]	Historical mean	AR(1)	LSTM (5 seeds)	Chronos-2	Chronos-Bolt	TimesFM-2.5
S&P 500	+0.15 [0.095]	+0.94 [0.639]	+0.82 [0.116]	-0.26 [0.418]	-1.57 [0.078]	-3.35 [0.001]
Bitcoin	+0.07 [0.691]	-0.28 [0.326]	-0.08 [0.135]	-0.35 [0.413]	-0.61 [0.070]	-2.29 [0.001]
BET	+0.30 [0.081]	-0.00 [1.000]	+1.18 [0.015]	-0.39 [0.562]	+0.12 [0.831]	-3.27 [0.070]

- S&P 500 and Bitcoin: no model beats zero significantly; TimesFM-2.5 is significantly **worse** than zero
 - ▶ Foundation models forecast too much movement: the s.d. of TimesFM forecasts on the S&P 500 is 0.141%, three times that of the LSTM (0.047%)
- BET: LSTM R^2 +1.18% (DM $p = 0.015$): a small but real signal in a less efficient market (Chapter 2)

Cumulative Squared-Error Gains on the S&P 500




□ $\sum_{s \leq t} [r_s^2 - (r_s - \hat{r}_s)^2]$: rises when the model beats the zero forecast, falls when it loses

Interpretation: Gains Come from a Few Days

- All curves stay near zero in calm years and jump in March 2020
 - ▶ Squared errors are dominated by crisis days: one month decides the ranking
- AR(1) gains in 2020 from the negative autocorrelation of crash days, then gives part back
 - ▶ TimesFM-2.5 loses steadily: its forecasts are too volatile
- Lesson: report the path, not only the final R^2 ; a single episode can create or destroy "predictability"

The Seed Lottery



- Out-of-sample R^2 of five LSTMs that differ only in the random seed, and of their average

Interpretation: Report All Seeds

- ▣ Bitcoin: single seeds range from -1.47% to +0.21%; S&P 500 from +0.35% to +1.40%
 - ▶ Seed-to-seed variation is as large as the effect itself
- ▣ BET: all 5 seeds positive, between +1.03% and +1.25%: the one robust signal
- ▣ Choosing the best seed after looking at test results is backtest overfitting (Chapter 13)

Section Summary

- ▣ Daily returns of large markets: no deep or foundation model beats the zero forecast
- ▣ Zero-shot point forecasts can be worse than zero: too much movement is predicted
- ▣ A small LSTM finds a robust signal in the BET, consistent with its lower efficiency


Target, Benchmark and Loss

- ▣ **Realised variance** (RV, [Andersen et al. \(2003\)](#)): sum of squared 5-minute log returns within the trading day
 - ▶ $RV_t = \sum_{j=1}^{78} r_{t,j}^2$ for SPY (regular session, 6.5 hours); average volatility 11.9% p.a.
- ▣ **HAR** ([Corsi \(2009\)](#)): $RV_t = \beta_0 + \beta_d RV_{t-1} + \beta_w \overline{RV}_{t-5:t-1} + \beta_m \overline{RV}_{t-22:t-1} + \varepsilon_t$
 - ▶ log-HAR: the same on log RV, forecast $\exp(\hat{y} + \hat{s}^2/2)$
- ▣ **QLIKE** loss ([Patton \(2011\)](#)): $L = RV_t/\hat{h}_t - \log(RV_t/\hat{h}_t) - 1$
 - ▶ Ranks forecasts correctly even though RV is a noisy proxy; minimised by the conditional mean

Design of the Volatility Test

- ▣ SPY realised variance, one-day-ahead forecasts from 11.10.2022 to 18.09.2026 (979 days)
 - ▶ HAR and log-HAR on a rolling 500-day window; random walk: $\hat{h}_t = \text{RV}_{t-1}$
 - ▶ GARCH-t on daily SPY returns, rescaled to RV units by the ratio of average RV to average squared return over 500 days
 - ▶ LSTM on the last 22 values of log RV, retrained every 125 days on 500 days (hidden size 8, three seeds)
- ▣ Foundation models on log RV (context 512): forecast $\hat{h}_t = \int_0^1 \exp(q_u) du$ on the quantile grid
 - ▶ Post-release window: 218 days from 3 November 2025

Forecasts Through the April 2025 Shock




- SPY realised volatility (annualised) and one-day-ahead forecasts, February–July 2025
- Peak: 107% p.a. on 07.04.2025, after the US tariff announcements

Interpretation: Everybody Is One Day Late

- No model anticipates the jump; all of them react the next day
 - ▶ One-day forecasts can only extrapolate the past; shocks come from news
- The foundation models and log-HAR decay back at similar speeds
 - ▶ GARCH-t reacts to the size of daily returns and decays more slowly
- The ranking is decided by the average over about four years: next slide

QLIKE Losses of Eight Volatility Models



Left: full test period; right: window after the release of all model weights

Interpretation: Foundation Models Match log-HAR, Not More

Model	QLIKE	DM vs log-HAR	p	MCS p	QLIKE after 3.11.2025
Random walk	0.326	+4.53	<0.001	0.00	0.384
HAR	0.247	+1.30	0.194	0.26	0.270
log-HAR	0.234	—	—	0.95	0.238
GARCH-t	0.283	+3.43	<0.001	0.03	0.295
LSTM	0.290	+1.87	0.062	0.19	0.231
Chronos-2	0.231	-0.62	0.537	1.00	0.242
Chronos-Bolt	0.232	-0.18	0.857	0.95	0.236
TimesFM-2.5	0.233	-0.08	0.938	0.95	0.244


- Chronos-2, Chronos-Bolt, TimesFM-2.5 and log-HAR: DM tests not significant, all in the 90% MCS
- Random walk and GARCH-t are eliminated from the MCS; HAR in levels over-forecasts (mean forecast / mean RV = 1.36)
- After 3 November 2025 the small LSTM has the lowest QLIKE (0.231): over 218 days, rankings are unstable

Median or Mean? A Costly Detail

Chronos-2 and TimesFM-2.5 on log RV	QLIKE	Mean forecast / mean RV
Chronos-2, $\int_0^1 \exp(q_u) du$	0.231	1.02
Chronos-2, $\exp(q_{0.5})$	0.259	0.79
TimesFM-2.5, $\int_0^1 \exp(q_u) du$	0.233	0.94
TimesFM-2.5, $\exp(q_{0.5})$	0.259	0.79

- RV is right-skewed: its median is about 20% below its mean
 - ▶ $\exp(\text{median of log RV})$ is the median of RV, not its mean: a biased volatility forecast
 - ▶ The median version is significantly worse than the mean version (DM +3.60 for Chronos-2) and worse than log-HAR (0.234)

How Much Context Does Chronos-2 Need?



□ QLIKE of Chronos-2 on log RV with contexts of 32 to 512 days (same forecast days)

 MFM_ch14_realised_vol

Interpretation: Long Memory Needs Long Context

- ▣ QLIKE falls from 0.290 with 32 days to 0.231 with 512 days
 - ▶ With less than about 256 days, Chronos-2 is worse than log-HAR (0.234)
- ▣ Volatility has long memory: months of history tell the model the slow-moving level
 - ▶ HAR encodes the same idea explicitly, through its 22-day average
- ▣ Context length is a modelling choice to report and to test, like a window length


Section Summary

- ▣ Zero-shot foundation models forecast realised variance as well as a well-specified log-HAR
- ▣ They are not significantly better, before or after the release of their weights
- ▣ Use the mean, not the median, and a long context; always compare with HAR

Design of the Tail-Risk Test

- $\text{VaR}_\alpha = -q_\alpha(r_{t+1} \mid \mathcal{F}_t)$, the loss exceeded with probability α ; $\text{ES}_\alpha = -\frac{1}{\alpha} \int_0^\alpha q_u du$
 - ▶ Levels: VaR 1% (coverage tests) and the pair (VaR 2.5%, ES 2.5%) for the FZ0 loss (Fissler & Ziegel (2016), Patton, Ziegel & Chen (2019))
- Classical models on a 1,000-day window: HS, GARCH(1,1)-t, FHS (McNeil & Frey (2000)), re-estimated every 20 days
- Foundation models, two ways
 - ▶ **Raw**: Chronos-2 quantiles of tomorrow's return (context: 512 returns)
 - ▶ **Hybrid**: volatility $\hat{\sigma}_t$ = forecast mean of $|r_t|$ from Chronos-2 or TimesFM; FHS on $r_t/\hat{\sigma}_t$

VaR 1% Breach Rates




- Share of days with a loss above VaR 1% (log scale); grey band: 95% range of the breach rate if the model were correct
- Chronos-Bolt: its lowest quantile (10%) used as "VaR 1%"

Interpretation: Raw Quantiles Are Too Narrow

Breaches [Kupiec p]	HS	GARCH-t	FHS	Chronos-2 quantiles	FHS + Chronos-2 vol.	FHS + TimesFM vol.
S&P 500 (26.9)	37 [0.065]	44 [0.002]	30 [0.559]	45 [0.001]	25 [0.705]	27 [0.989]
Bitcoin (28.2)	19 [0.065]	27 [0.822]	25 [0.539]	33 [0.374]	27 [0.822]	27 [0.822]
BET (26.8)	32 [0.322]	41 [0.010]	30 [0.536]	40 [0.016]	26 [0.884]	32 [0.322]

- In brackets after the market: the expected number of breaches, 1% of the days
- Raw Chronos-2 quantiles: too many breaches on the S&P 500 (1.67%); Chronos-Bolt "VaR 1%": 13.5%
 - ▶ Hybrid FHS with foundation-model volatility: breach rates close to 1% in all three markets

The COVID-19 Crash: VaR 1% of Three Models




□ S&P 500, February–June 2020 (104 days): daily losses and VaR 1% forecasts; dots: breaches

Interpretation: Speed of Reaction

- Breaches in the 104 days: FHS 2, raw Chronos-2 4, FHS with Chronos-2 volatility 2, HS 12
 - ▶ All models are hit in the first days of the crash: none anticipates it
- Chronos-2 reads the context in the first days, so its quantiles widen fast
 - ▶ Scaling by the mean absolute value of the context: one crash day widens the whole distribution
- The hybrid overshoots: its VaR 1% peaks at 58.9%, against 26.4% for FHS and a largest loss of 12.8%
 - ▶ Few breaches, but a high capital cost: this is what the FZ0 loss will penalise

Where Deciles Suffice: VaR 10%



- Breach rate of the 10% quantile (target 10%) for HS, FHS and the three foundation models

Interpretation: Even the 10% Quantile Is Too High

- ▣ S&P 500: HS 9.8%, FHS 9.6%; Chronos-2 11.7%, Chronos-Bolt 13.5%, TimesFM-2.5 12.0%
 - ▶ Bitcoin: Chronos-2 10.6%, Chronos-Bolt 12.8%, TimesFM-2.5 11.6%; BET: 11.6%, 14.1%, 11.8%
- ▣ All three foundation models breach their 10% quantile too often: predictive distributions too narrow for returns
 - ▶ A plausible reason: pretraining corpora are dominated by smoother series (energy, traffic, web) with lighter tails than returns
- ▣ Used as volatility filters, the same models inherit the tails from the data, through FHS

Comparing (VaR 2.5%, ES 2.5%) Forecasts: the MCS

HS	0.140	0.343	0.015
GARCH-t	0.187	0.597	0.341
FHS	1.000	0.597	1.000
Chronos-2 quantiles	0.140	0.641	0.239
FHS + Chronos-2 volatility	0.140	1.000	0.341
FHS + TimesFM volatility	0.183	0.704	0.202
	S&P 500	Bitcoin	BET

□ p -values of the 90% model confidence set ([Hansen, Lunde & Nason \(2011\)](#)) for the FZ0 loss; green: kept, red: eliminated

Interpretation: Hybrids Pass, FHS Remains the Reference

FZ0 [MCS p]	HS	GARCH-t	FHS	Chronos-2 quantiles	FHS + Chronos-2 vol.	FHS + TimesFM vol.
S&P 500	1.309 [0.14]	1.026 [0.19]	0.999 [1.00]	1.147 [0.14]	1.084 [0.14]	1.059 [0.18]
Bitcoin	2.237 [0.34]	2.192 [0.60]	2.192 [0.60]	2.181 [0.64]	2.151 [1.00]	2.160 [0.70]
BET	1.230 [0.01]	1.001 [0.34]	0.963 [1.00]	1.032 [0.24]	1.007 [0.34]	1.027 [0.20]

- S&P 500: FHS has the lowest FZ0 loss; hybrids cover correctly but with larger, more conservative VaR (3.20% vs 2.78% on average)
- Bitcoin: the hybrid FHS with Chronos-2 volatility has the lowest FZ0 loss, but the differences are not significant
- The MCS eliminates no model on the S&P 500 and Bitcoin, only HS on the BET; pairwise DM tests are sharper: FHS beats the Chronos-2 hybrid on the S&P 500 (DM +2.65, $p = 0.008$) and on the BET (DM +2.36, $p = 0.018$)

After the Release of the Weights

Breaches after 3.11.2025	HS	GARCH-t	FHS	Chronos-2 quantiles	FHS + Chronos-2 vol.	FHS + TimesFM vol.
S&P 500 (220 days)	0	2	2	2	2	2
Bitcoin (320 days)	3	3	4	4	3	4
BET (217 days)	4	4	4	5	1	2

- Expected number of breaches: 2.2 (S&P 500), 3.2 (Bitcoin), 2.2 (BET)
 - ▶ No leakage is possible here, but with two or three expected breaches no test has power
- For tail risk, the post-release window cannot confirm or reject anything yet: it must grow for years

Section Summary

- ▣ Raw foundation-model quantiles are too narrow for returns, even at 10%; Chronos-Bolt and TimesFM cannot reach 1% at all
- ▣ As volatility filters inside FHS they give correct coverage, but a significantly higher FZ0 loss than GARCH-based FHS on the S&P 500 and BET
- ▣ None of them replaces GARCH-based FHS on the evidence so far

An Open Question: Skill or Memory?

- Test 2: zero-shot Chronos-2 matches log-HAR on SPY realised variance (QLIKE 0.231 vs 0.234, DM $p = 0.537$)
- Rahimikia, Ni & Wang (2025): on daily excess returns across global markets, off-the-shelf foundation models do poorly zero-shot and after fine-tuning; models pretrained from scratch on financial data gain
- Lopez-Lira, Tang & Zhu (2025): once a model has seen the realised values, its counterfactual forecasting skill is not identified; detected memorisation is only a lower bound (shown for LLMs; the logic applies to any pretrained forecaster)
- Open: is the parity learned structure (long memory, mean reversion) or recall of the test period, and can a short post-release window tell?
 - $d_t = L(\hat{h}_t^{FM}, RV_t) - L(\hat{h}_t^{\log HAR}, RV_t)$; memorisation implies $E[d_t | \text{post}] > E[d_t | \text{pre}]$; H_0 : equality

The Discovery Loop, Step by Step (1/2)

- **Literature:** "Find evaluations of time-series foundation models on financial returns or volatility since 2024: models, data, loss, benchmark, and whether the test period overlaps the pretraining corpus."
 - ▶ Check every arXiv identifier and read the corpus documentation yourself; Semantic Scholar, Google Scholar, Elicit help with the search
- **Hypotheses:** "List predictions that separate memorisation from skill for a zero-shot volatility forecaster: pre vs post release, never-public series, simulated series with the same dynamics."
 - ▶ Skill predicts equal gains on a simulated log-HAR path; memory predicts gains only on public series before the cutoff
- **Data and code:** "From `ch14_rv.csv`, compute the QLIKE differential Chronos-2 minus log-HAR, test pre vs post 3 November 2025 with HAC errors, and report the minimum detectable effect at 80% power."
 - ▶ Any assistant (Claude, ChatGPT, Gemini, GitHub Copilot); first reproduce the chapter's QLIKE of log-HAR, 0.234

The Discovery Loop, Step by Step (2/2)

- **Robustness**, fixed in advance: "Design a pre-registered test with more power: which cross-section of series, which variance proxy, which panel DM test, and how many post-release days."
 - ▶ Placebo: a fake release date one year earlier must show no break; simulated series cannot be memorised
- **Critique**: "Attack this claim: no memorisation, because the post-release DM test is insignificant. What is wrong with the inference?"
 - ▶ Expect: low power, absence of evidence, and a regime change (the 2025 tariff shock) that coincides with the release
- **Report**: model versions and weight release dates, context length, mean or median, every seed and window tried

What the Human Must Check

- **Pretraining overlap:** a public series may sit in the corpus
 - ▶ Test: post-release window, simulated or private series; record the release date of the weights
- **Low power sold as a null result:** "not significant" over a few hundred days
 - ▶ Test: report the minimum detectable effect (MDE) and the CI of $E[d_t]$, not only p
- **Wrong functional:** the median of log RV used as a mean forecast
 - ▶ Test: mean forecast / mean RV should be close to 1 (Chronos-2 median: 0.79)
- **Leakage in your own code:** HAR averages or scalings that include day t (Seminar B8)
 - ▶ Test: recompute the forecast for t from data truncated at $t - 1$; it must not change
- **Selection:** contexts, seeds and models chosen after seeing the test
 - ▶ Test: log every trial; compare the whole set with the MCS ([Hansen, Lunde & Nason \(2011\)](#))

Mini-Case: a Test Without Power

- A typical assistant proposal: "the parity of Chronos-2 comes from memorised data, so its edge must vanish after release". We test it on the saved forecasts
- Mean QLIKE differential d_t : -0.0050 before 3 November 2025 (761 days) vs +0.0038 after (218 days)
- Change +0.0088, HAC $t = 1.00$ ($p = 0.32$): the direction memorisation predicts, but H_0 is not rejected
- Power: with 218 days, MDE (5% two-sided, 80% power) = $(z_{0.975} + z_{0.8}) \hat{s}_{HAC} / \sqrt{n} = 0.0180$
 - ▶ 7.7% of the log-HAR QLIKE, 3.6 times the pre-release gap; detecting that gap needs about 7 400 days, some 29 years
- Reading: the window cannot separate memory from skill; "no memorisation" is not a finding. Power must come from the cross-section, not from waiting

 ai_discovery_case.py

Your Turn: a Project Seed

- Question: does zero-shot skill for daily variance survive where memorisation is impossible: after release, on simulated series, on less-covered markets?
- Replicate first: the chapter's QLIKE of log-HAR (0.234) and Chronos-2 (0.231), DM $p = 0.537$, and the 218-day post-release window
- Deliverables
 - ▶ A pre-registered power calculation; a panel of series from the course data with a range-based variance proxy and a proxy-robust loss ([Patton \(2011\)](#))
 - ▶ Pre/post tests with placebo release dates; the MCS across models; a referee-style critique and your reply
- Log each prompt that shaped a decision in AI_USE.md; each invented reference, leak or wrong number caught in AI_ERRORS.md, with the test that caught it ([templates](#))

Pitfalls Checklist

- ▣ **Leakage:** pretraining overlap, scaling with full-sample statistics, shuffled validation windows
 - ▶ Seminar B8: HAR averages that include the forecast day cut QLIKE by a quarter, a gain no honest model achieves
- ▣ **Weak benchmarks:** compare with zero (returns), log-HAR (volatility), FHS (tails), not with a random walk
- ▣ **Selection:** seeds, context lengths, architectures chosen after seeing the test period
- ▣ **Wrong functional:** medians used as means; deciles used as 1% quantiles

When to Use What

Target	Reference model	Where deep learning can help
Daily index returns	Zero forecast	Less efficient markets (BET); many assets, many features
Realised volatility	log-HAR	Zero-shot forecasts for thousands of series without estimation
VaR 1%, ES 2.5%	FHS (GARCH filter)	Volatility filter inside FHS, not raw quantiles

- Foundation models are cheap to try: no estimation, one line of code, a CPU is enough
- Their value is breadth (many series, no tuning), not accuracy on a single well-studied series

Related Research

- ▣ Agakishiev, Härdle & Pele (2025): forecasting cryptocurrency liquidity, Springer Handbook of Computational Statistics
- ▣ Będowska-Sójka, Wójcik & Pele (2026): machine-learning early-warning systems for "zombie" crypto-assets
- ▣ Pele et al. (2023): cryptocurrencies classified with statistical learning across many series
- ▣ Textbooks: Franke, Härdle & Hafner (2019) and Härdle, Chen & Overbeck (2017) for the statistics of financial markets behind the benchmarks

Key Takeaways

- LSTMs solve the vanishing gradient through the forget gate; Transformers avoid it through attention
- Foundation models forecast zero-shot from a context window and return quantiles
- Returns: nothing beats zero on large markets; volatility: foundation models equal log-HAR; tails: use them only as volatility filters
- Evaluate every new model with the tools of Chapter 8, against a strong benchmark and after the release of its weights

References (1/6)

- ▣ Agakishiev, I., Härdle, W.K., Pele, D.T. (2025). [Cryptocurrency liquidity forecasting](#). In *Springer Handbooks of Computational Statistics*, 41–60. Springer.
- ▣ Aksu, T., Woo, G., Liu, J., Liu, X., Liu, C., Savarese, S., Xiong, C., Sahoo, D. (2024). [GIFT-Eval: a benchmark for general time series forecasting model evaluation](#). [arXiv:2410.10393](#).
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X., Labys, P. (2003). [Modeling and forecasting realized volatility](#). *Econometrica*, 71(2), 579–625.
- ▣ Ansari, A.F., Stella, L., Turkmen, C., Zhang, X., et al. (2024). [Chronos: learning the language of time series](#). *Transactions on Machine Learning Research*; [arXiv:2403.07815](#).
- ▣ Ansari, A.F., Shchur, O., Küken, J., Auer, A., et al. (2025). [Chronos-2: from univariate to universal forecasting](#). [arXiv:2510.15821](#).
- ▣ Będowska-Sójka, B., Wójcik, P., Pele, D.T. (2026). [Early warning systems for cryptocurrency markets: predicting 'zombie' assets using machine learning](#). *North American Journal of Economics and Finance*, 81, 102543.
- ▣ Bengio, Y., Simard, P., Frasconi, P. (1994). [Learning long-term dependencies with gradient descent is difficult](#). *IEEE Transactions on Neural Networks*, 5(2), 157–166.

References (2/6)

- ▣ Campbell, J.Y., Thompson, S.B. (2008). [Predicting excess stock returns out of sample: can anything beat the historical average?](#) *Review of Financial Studies*, 21(4), 1509–1531.
- ▣ Corsi, F. (2009). [A simple approximate long-memory model of realized volatility.](#) *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Das, A., Kong, W., Sen, R., Zhou, Y. (2024). [A decoder-only foundation model for time-series forecasting.](#) *Proceedings of ICML*, PMLR 235; arXiv:2310.10688.
- ▣ Diebold, F.X., Mariano, R.S. (1995). [Comparing predictive accuracy.](#) *Journal of Business & Economic Statistics*, 13(3), 253–263.
- ▣ Fischer, T., Krauss, C. (2018). [Deep learning with long short-term memory networks for financial market predictions.](#) *European Journal of Operational Research*, 270(2), 654–669.
- ▣ Fissler, T., Ziegel, J.F. (2016). [Higher order elicibility and Osband's principle.](#) *Annals of Statistics*, 44(4), 1680–1707.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.

References (3/6)

- ▣ Gu, S., Kelly, B., Xiu, D. (2020). [Empirical asset pricing via machine learning](#). *Review of Financial Studies*, 33(5), 2223–2273.
- ▣ Hansen, P.R., Lunde, A., Nason, J.M. (2011). [The model confidence set](#). *Econometrica*, 79(2), 453–497.
- ▣ Härdle, W.K., Chen, C.Y.-H., Overbeck, L. (eds.) (2017). [Applied Quantitative Finance](#), 3rd ed. Springer.
- ▣ Hochreiter, S., Schmidhuber, J. (1997). [Long short-term memory](#). *Neural Computation*, 9(8), 1735–1780.
- ▣ Hornik, K., Stinchcombe, M., White, H. (1989). [Multilayer feedforward networks are universal approximators](#). *Neural Networks*, 2(5), 359–366.
- ▣ Kingma, D.P., Ba, J. (2015). [Adam: a method for stochastic optimization](#). *ICLR*; arXiv:1412.6980.
- ▣ LeCun, Y., Bengio, Y., Hinton, G. (2015). [Deep learning](#). *Nature*, 521(7553), 436–444.

References (4/6)

- ▣ Liu, C., Aksu, T., Liu, J., Liu, X., et al. (2025). [Moirai 2.0: when less is more for time series forecasting](#). [arXiv:2511.11698](#).
- ▣ Lopez-Lira, A., Tang, Y., Zhu, M. (2025). [The memorization problem: can we trust LLMs' economic forecasts?](#) [arXiv:2504.14765](#).
- ▣ McNeil, A.J., Frey, R. (2000). [Estimation of tail-related risk measures for heteroscedastic financial time series: an extreme value approach](#). *Journal of Empirical Finance*, 7(3–4), 271–300.
- ▣ Nie, Y., Nguyen, N.H., Sinthong, P., Kalagnanam, J. (2023). [A time series is worth 64 words: long-term forecasting with Transformers](#). *ICLR*; [arXiv:2211.14730](#).
- ▣ Oreshkin, B.N., Carpo, D., Chapados, N., Bengio, Y. (2020). [N-BEATS: neural basis expansion analysis for interpretable time series forecasting](#). *ICLR*; [arXiv:1905.10437](#).
- ▣ Patton, A.J. (2011). [Volatility forecast comparison using imperfect volatility proxies](#). *Journal of Econometrics*, 160(1), 246–256.
- ▣ Patton, A.J., Ziegel, J.F., Chen, R. (2019). [Dynamic semiparametric models for expected shortfall \(and value-at-risk\)](#). *Journal of Econometrics*, 211(2), 388–413.

References (5/6)

- ▣ Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets? *The European Journal of Finance*, 29(10), 1064–1105.
- ▣ Rahimikia, E., Ni, H., Wang, W. (2025). Re(Visiting) time series foundation models in finance. arXiv:2511.18578.
- ▣ Rasul, K., Ashok, A., Williams, A.R., Ghonia, H., et al. (2023). Lag-Llama: towards foundation models for probabilistic time series forecasting. arXiv:2310.08278.
- ▣ Rosenblatt, F. (1958). The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6), 386–408.
- ▣ Rumelhart, D.E., Hinton, G.E., Williams, R.J. (1986). Learning representations by back-propagating errors. *Nature*, 323(6088), 533–536.
- ▣ Salinas, D., Flunkert, V., Gasthaus, J., Januschowski, T. (2020). DeepAR: probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting*, 36(3), 1181–1191.
- ▣ Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., et al. (2017). Attention is all you need. *NeurIPS*; arXiv:1706.03762.

References (6/6)

- Woo, G., Liu, C., Kumar, A., Xiong, C., Savarese, S., Sahoo, D. (2024). Unified training of universal time series forecasting Transformers. *ICML*; arXiv:2402.02592.