








Modelling Financial Markets

Chapter 16: Digital Assets and DeFi



Daniel Traian PELE

Bucharest University of Economic Studies
IDA Institute Digital Assets
Blockchain Research Center
AI4EFin Artificial Intelligence for Energy Finance
Romanian Academy, Institute for Economic Forecasting
MSCA Digital Finance

Digital Assets: What Is Different and How to Measure It

Digital Assets, Pricing and Liquidity

- The question: do cryptos, stablecoins and DeFi pools behave like familiar assets, and which measurements must change?
- Part I: the crypto market, stylised facts of returns, crypto indices (CRIX), stablecoins
- Part II: DeFi (decentralised finance) and automated market makers, spot ETFs (exchange-traded funds that hold the coin) and crypto equities
  - Cryptos as an asset class, a case study, one pass of the research loop assisted by artificial intelligence (AI)
- Technical detail sits in the Appendix at the end

What you should be able to do

- measure a 7-day market on the right calendar and attach intervals to its tail index, volatility and correlations
- build a market-value index with a divisor and model a stablecoin peg as a threshold autoregression (threshold AR)
- price a swap in a constant-product pool and explain why providing liquidity is selling volatility
- judge ETF tracking, crypto betas and asset-class claims
  - with tests that respect closing times and unknown break dates

Core Path and Optional Depth

First reading: the core path

- Market and returns: market value and concentration, crypto vs shares and gold, the 2020 break
- Indices and stablecoins: the index and its divisor, the threshold AR of the peg, March 2023
- DeFi: a swap step by step, impermanent loss and LVR
- ETFs and factors: tracking, the Dimson beta, the GRS test, eight coins
- Research with an assistant: the ETF mini-case

Second reading: optional depth

- The CRIX rule, Fama–MacBeth, the asset-class map, the case study of Liu, Tsyvinski & Wu (2022)
- Appendix: conditions for the Hill interval, the sup-Wald limit, the LVR derivation, robust factor tests
- Worked exercises: Seminar 16, A1 (index divisor), Seminar 16, A2 (swap and arbitrage band), Seminar 16, B5 (threshold AR)

Blockchains, Coins and Tokens



Bust of Satoshi Nakamoto, Graphisoft Park, Budapest

Photo: Fekist (2021); CC BY-SA 4.0; Wikimedia Commons

- **Blockchain** Nakamoto (2008): a shared ledger of blocks, each carrying the hash of the previous one
  - Secured by proof of work (Bitcoin) or proof of stake (Ethereum since 15 September 2022, ethereum.org: The Merge)
- Bitcoin issuance halves over time up to 21 million coins Bitcoin.org FAQ; 20.09 million existed on 18 Sep 2026
- **Coins** are native assets of a blockchain
  - **tokens** (stablecoins, governance tokens, tokenised funds) are issued by **smart contracts** Buterin (2014)
  - DeFi is built from such contracts Schär (2021)
- Surveys: technology Lin, Pele & Ren (2025); economics Härdle, Harvey & Reule (2020)

Seven-Day Crypto Data Next to Five-Day Markets

Series	Content	Period
Crypto prices	Bitcoin, Ethereum, Solana, Dogecoin, Cardano, Litecoin; daily, 7 days a week	2018–2026
Stablecoins	Tether (USDT), USD Coin (USDC), Dai (DAI) prices; supply by coin (DefiLlama)	2019–2026
Coin supply	Coins in circulation, daily (Coin Metrics Community Data)	2018–2026
DeFi	Total value locked, all chains and by chain (DefiLlama)	2020–2026
ETFs and shares	spot ETFs on Bitcoin (IBIT) and Ether (ETHA), Coinbase, Strategy; S&P 500, Nasdaq 100 (ETF QQQ), gold (XAU/USD, USD per ounce)	2018–2026
Asset universe	26 assets: cryptos, equity indices, currencies, commodities, bonds	2019–2026

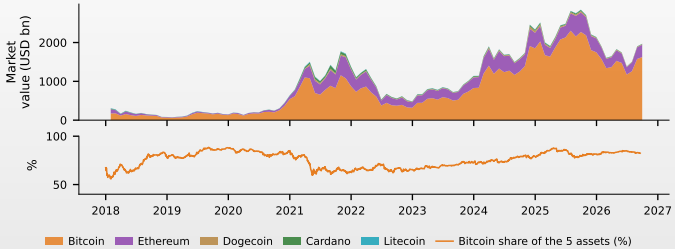
- Daily closes up to 18 September 2026; crypto on its own 7-day calendar; analyses that combine assets use common days (prices aligned on common trading days first, then returns); some tests use longer histories and extra series (next slide)

Longer Histories and Extra Series

Series	Period	Used for
Bitcoin and S&P 500, daily	from January 2016	2017–2019 correlations, the Forbes–Rigobon base period, the break test (2 692 common days)
Bitcoin and XRP; Ether	from January 2014; from 6 August 2015 (week 32)	the case study: Table 1 of Liu, Tsyvinski & Wu (2022) rebuilt, 52-week volatility
XRP, BNB, Chainlink, daily	weekly returns from 9 Feb 2018	the GRS test on eight coins (Section 7), the case-study replication
4-week Treasury bill, DTB4WK (FRED)	the same weeks as the GRS test	excess returns in the GRS test
One-month Treasury bill rate (Kenneth French Data Library)	the case-study weeks	the case study, as in the paper's definition of CMKT
Spot ETFs IBIT and ETHA	from 11 Jan 2024 and 23 Jul 2024	ETF tracking and the Dimson beta
Coinbase and Strategy shares	weekly, April 2021 – September 2026	the two-factor regressions of crypto equities

- The table on the previous slide gives the default window; a slide that uses another window or series says so next to its numbers

Five Coins, One Dominant Asset



- **Market value** of coin i at month end t: $M_{i,t} = P_{i,t} \times Q_{i,t}$, with $P_{i,t}$ the price in USD and $Q_{i,t}$ the coins in circulation
  - Bitcoin: 1.63 trillion USD on 18 Sep 2026
- The five coins: 1 970 billion USD today, peak 3 135 billion (6 Oct 2025); the value moves with prices, supply changes slowly
- Bitcoin's share: 82.5% today, between 55.8% (1 Feb 2018) and 88.5% (5 Sep 2019)

MFM_ch16_market_value

Five Coins, One Dominant Asset: Concentration

- Concentration: **Herfindahl–Hirschman index** $\mathrm{HHI} = \sum_i w_i^2$, $w_i = M_{i,t} / \sum_j M_{j,t}$ the value share of coin i
  - From $1/5$ (equal shares) to 1 (one coin)
  - Today 0.71, an effective number of coins $1/\mathrm{HHI} = 1.42$
  - Check: two coins with shares 80% and 20% give $\mathrm{HHI} = 0.8^2 + 0.2^2 = 0.68$ and $1/\mathrm{HHI} = 1.47$, about one and a half "equal" coins
- The five coins: Bitcoin, Ethereum, Dogecoin, Cardano and Litecoin; Solana is in the price data but has no public circulating-supply series, so it has no market value here (Section 3)
- XRP (the native token of the XRP Ledger) and Chainlink are left out: their reported supply includes issuer-held tokens (free float matters)

Returns, volatility and asymmetry

- Daily log return in %: $r_t = 100\,(\ln P_t - \ln P_{t-1})$
  - P_t: closing price on day t; each series on its own calendar (crypto 7 days a week, shares 5)
  - Not the simple percentage change $R_t = 100\,(P_t/P_{t-1} - 1)$: $r_t = 100\ln(1 + R_t/100)$, so a 20% fall ($R_t = -20$) is $r_t = -22.3$; log returns add over time; a portfolio's simple return is the weighted mean of its assets' simple returns
- Annualised volatility $\sigma_a = s\sqrt{A}$
  - s: sample standard deviation of the daily r_t; A: return periods per year, about 365 for crypto and 252 for shares
- Skewness $S = \overline{(r - \bar{r})^3}/s^3$: $S < 0$ means large losses are more frequent than large gains
  - A bar denotes the sample mean: $\bar{r}$ is the mean return, $\overline{(r - \bar{r})^3}$ the mean cubed deviation

Drawdown, clustering and tail losses

- Maximum drawdown $\mathrm{MDD} = \min_t(P_t / \max_{u \le t} P_u - 1)$: the deepest fall from a previous peak, between -100% and 0
  - $\max_{u \le t} P_u$: the running peak, the highest price on the days u up to t
- Volatility clustering $\hat{\rho}_1(|r|)$: sample autocorrelation of $|r_t|$ at lag 1 (Chapter 1); > 0: a large move tends to follow a large move
- Historical daily $\mathrm{VaR}_{1\%} = -q_{1\%}(r)$ and $\mathrm{ES}_{2.5\%} = -\overline{r \mid r \le q_{2.5\%}(r)}$ (Chapter 7)
  - $q_p(r)$: empirical p-quantile of the returns; VaR (value at risk): the loss threshold exceeded with probability 1%; ES (expected shortfall): the mean loss on the worst 2.5% of the days
  - Why two levels: the Basel market-risk rules replaced VaR at 1% by ES at 2.5% (Basel Committee, MAR33, para. 33.3); under the Normal distribution both are close, 2.326σ and 2.338σ; $\mathrm{ES}_p \ge \mathrm{VaR}_p$ holds at the same p, not across levels

Estimating the Tail Index: the Hill Estimator

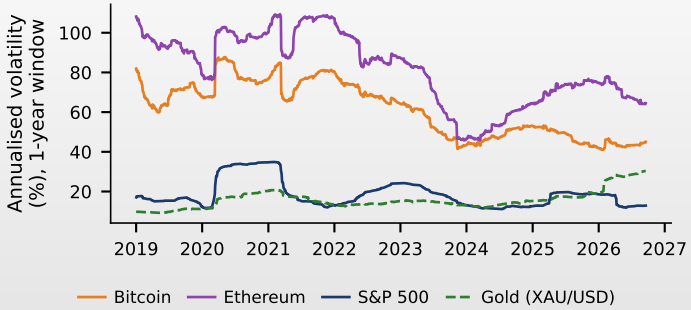
- Power-law tail of the daily loss $L = -r$: $P(L > x) \sim Cx^{-\alpha}$ for large x
  - $C > 0$: a scale constant; $\sim$: the ratio of the two sides tends to 1 as $x \to \infty$
  - α: the **tail index**; smaller α = heavier tail; the moment $E|r|^q$ is finite only for $q < \alpha$
- Hill estimator Hill (1975) from the k largest losses: $\hat{\alpha}_H = \left(\frac{1}{k}\sum_{j=1}^{k} \ln(L_{(j)}/L_{(k+1)})\right)^{-1}$
  - $L_{(1)} \geq L_{(2)} \geq \ldots$: the losses sorted from the largest; $k = 5\%$ of the losses; $L_{(k+1)}$ acts as the threshold
  - Each log ratio says how far a large loss lies above the threshold: large ratios on average give a small $\hat{\alpha}_H$, a heavy tail
  - Standard error (SE) $\mathrm{SE} = \hat{\alpha}_H/\sqrt{k}$ for independent and identically distributed (i.i.d.) data; 95% confidence interval (CI) $\hat{\alpha}_H \pm 1.96\,\mathrm{SE}$ (conditions and dependence in the Appendix)
- All six i.i.d. 95% CIs lie below 3.7: $\alpha < 4$, hence $E|r|^4 = \infty$, is plausible but not proven
  - If $\alpha < 4$, the sample excess kurtosis $K = \overline{(r - \bar{r})^4}/s^4 - 3$ has no finite population value to estimate: report $\hat{\alpha}_H$ with its CI instead
- Sample 2018–2026 (Solana from April 2020); comparison assets: S&P 500 and gold

Crypto vs Shares and Gold, 2018–2026

	n	σ_a (%)	S	MDD (%)	$\hat{\rho}_1(\lvert r\rvert)$	$\hat{\alpha}_H$	95% CI	$\mathrm{VaR}_{1\%}$	$\mathrm{ES}_{2.5\%}$
Bitcoin	3 182	64	−0.95	−82	0.17	2.73	[2.12, 3.33]	10.2	10.3
Ethereum	3 182	84	−0.82	−94	0.16	2.67	[2.08, 3.26]	13.5	13.7
Solana	2 351	117	−0.22	−96	0.28	2.85	[2.12, 3.58]	14.5	17.1
Dogecoin	3 182	124	4.88	−92	0.37	2.83	[2.22, 3.45]	13.8	16.4
S&P 500	2 189	19	−0.65	−34	0.36	2.87	[2.07, 3.67]	3.4	3.8
Gold	2 269	17	−0.55	−28	0.13	2.75	[1.99, 3.51]	3.0	3.2

- Daily log returns in %; VaR and ES are losses of the daily log return; $\hat{\alpha}_H$ from the largest 5% of losses, CI $\hat{\alpha}_H \pm 1.96\,\hat{\alpha}_H/\sqrt{k}$ (1.96: the 97.5% quantile of the Normal distribution)
- Bitcoin: 3.3 times the S&P 500's volatility; daily VaR 1% 10.2% in log terms, a loss of 9.7% of the position; worst day 12 Mar 2020 (COVID-19 crash)
- Drawdowns of 82–96% for the coins, 34% for the S&P 500; Hill indices 2.67–2.87 for all six assets: crypto is more volatile, but the data give no strong evidence of different tail indices (tests in the ▸ Appendix)

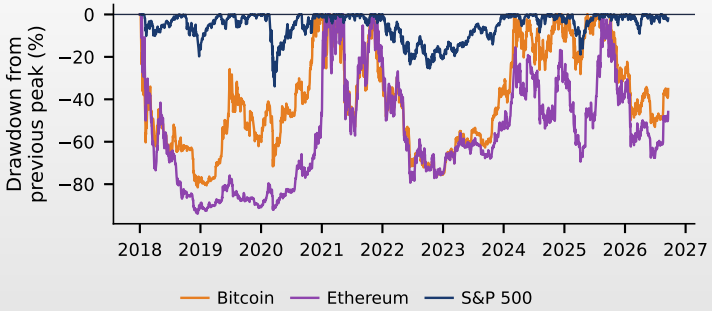
Volatility Over Time



- Annualised standard deviation of daily log returns over a moving one-year window
- Bitcoin fell from 88% to 45% today, Ethereum from 109% to 64%: consistent with a maturing market, but the chart does not identify the cause
- Bitcoin is still 3.5 times as volatile as the S&P 500 (13%); gold reached 30% in 2026, a record for the sample

MFM_ch16_stylised

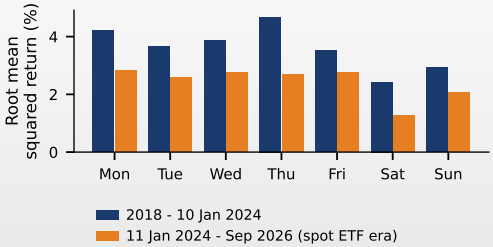
Drawdowns: How Deep and How Long



- Fall from the previous maximum, in %: Bitcoin was more than 50% below its peak on 36% of the days since 2018, Ethereum on 61%; the S&P 500 at most 34% (March 2020)
- Today: Bitcoin 35%, Ethereum 46% below their peaks; long "crypto winters" (2018–2020, 2022–2023) follow each boom (bubbles: Chapter 17)

MFM_ch16_stylised

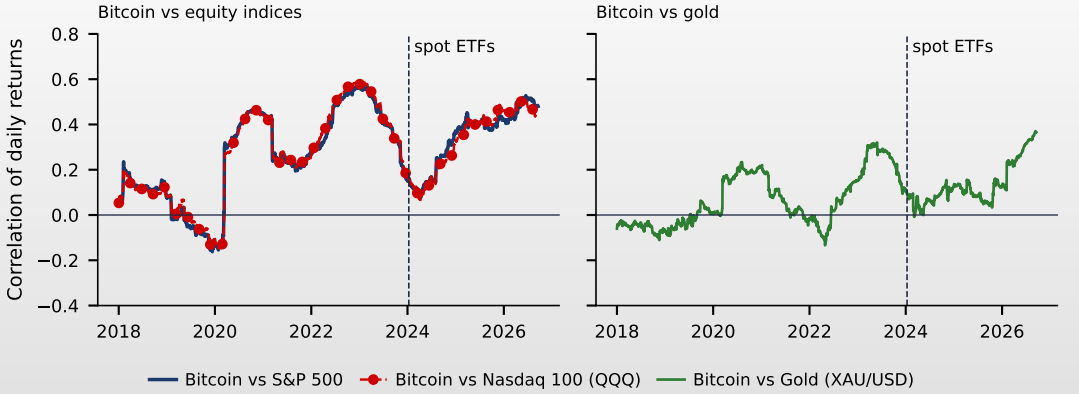
A 7-Day Market with a Weekend Lull



- Root mean squared daily Bitcoin return by day of the week; weekend variance / weekday variance: 0.45 before and 0.39 after 11 January 2024
- Bitcoin trades on Saturdays, but macro news, US equity trading and ETF flows do not; all days calmer after the ETF launch (weekday root mean square 4.01% $\rightarrow$ 2.73%)
- Annualise with 365 days, but do not treat all days as identical; did the ratio change? Seminar 16, B1

MFM_ch16_weekday

Is Bitcoin Moving with Shares? (1/2)



MFM_ch16_correlation

Is Bitcoin Moving with Shares? (2/2)

- Chart: correlation of daily log returns over a moving window of 250 common days
- Whole-period correlations, Bitcoin vs S&P 500: 0.02 (2017–2019, $n = 753$), 0.39 (2020–2023, $n = 1\,006$), 0.38 since the ETFs ($n = 674$ days)
- Nasdaq 100 like the S&P 500 (0.38); gold -0.02, 0.12, 0.22: a smaller diversification benefit than before 2020 Bouri et al. (2017), Liu & Tsyvinski (2021)
- The chart suggests a break in early 2020, not in January 2024: tested on the next two slides

MFM_ch16_correlation

Did the Bitcoin–Equity Link Break in 2020? (1/2): the Test

- Model on 2 692 common days since January 2016: $r_t^{\mathrm{BTC}} = a + b\, r_t^{\mathrm{SPX}} + u_t$
  - r_t^{BTC}, r_t^{SPX}: daily log returns of Bitcoin and the S&P 500; a: intercept; b: Bitcoin's beta to the S&P 500; u_t: error
  - (a, b) may shift at an unknown date τ: one pair before τ, another after; T: number of days
- For each candidate τ: $W(\tau)$, the Wald statistic of H_0: the same (a, b) before and after τ; a large $W(\tau)$ signals a large change at τ
  - Newey–West covariances (Bartlett, 8 lags) in each regime: heteroskedasticity and autocorrelation consistent (HAC)
- sup-Wald Andrews (1993): $\sup_{\tau \in [0.15T, 0.85T]} W(\tau)$, the largest W over the central 70% of the sample
  - Under H_0 the sup is not χ^2_2, because τ is chosen to make W large
  - 5% critical value 11.55, simulated from the limit distribution (Appendix); Seminar 16 quotes the tabulated value 11.70 of Andrews (1993): two approximations of the same limit, and sup $W = 20.21$ exceeds both

Did the Bitcoin–Equity Link Break in 2020? (2/2): the Result

- sup $W = 20.21$ ($p = 0.0008$): above the 5% critical value 11.55, so the no-break hypothesis is rejected
  - Least-squares break date 9 Mar 2020 ($W = 18.81$): the τ that minimises the sum of squared residuals (SSR) of the two regimes
  - At the ETF launch $W = 2.31$
  - sup W itself is reached on 14 Aug 2017: $W(\tau)$ uses a separate HAC covariance in each regime, so its maximiser need not minimise the sum of squared residuals; the test rejects the no-break hypothesis; its maximiser is only two trading days from the trimmed lower boundary (10 August 2017), so the least-squares date and its interval are the dating evidence
- 95% CI for the date Bai (1997): 27 Nov 2018 – 15 Jun 2021; it excludes January 2024
  - Beta b to the S&P 500: 0.07 before, 1.22 after; correlation 0.01 and 0.40
- Several breaks at unknown dates: the sequential tests and CIs of Bai & Perron (1998)

Correlation or Volatility? The Forbes–Rigobon Adjustment (1/2)

- Pearson correlation rises mechanically when the volatility of the driving market rises, with the same underlying link Forbes & Rigobon (2002)
- Adjusted correlation: the correlation the volatile period would show at the volatility of the calm period
  - $\rho^* = \dfrac{\rho}{\sqrt{1+\delta(1-\rho^2)}}, \qquad \delta = \dfrac{\sigma^2_{\mathrm{SPX,high}}}{\sigma^2_{\mathrm{SPX,low}}} - 1$
  - ρ: correlation of Bitcoin and S&P 500 returns in the volatile period; ρ^*: the adjusted correlation
  - $\sigma^2_{\mathrm{SPX,high}}$, $\sigma^2_{\mathrm{SPX,low}}$: variance of S&P 500 returns in the volatile period and in the calm base period (2017–2019)
  - δ: relative rise of that variance; $\delta = 0$ gives $\rho^* = \rho$, $\delta > 0$ gives $\rho^* < \rho$
- Caveat: the adjustment assumes no endogeneity and no omitted common shocks, which a pandemic hardly satisfies

Correlation or Volatility? The Forbes–Rigobon Adjustment (2/2)

- 2020–2023: S&P 500 volatility 23.1% vs 12.8% ($\delta = 2.24$); $\hat{\rho} = 0.39$ becomes $\rho^* = 0.23$
  - Since the ETFs: $\hat{\rho} = 0.38$, $\rho^* = 0.33$ ($\delta = 0.42$)
- Change against 2017–2019 ($\hat{\rho} = 0.02$): $\rho^*_{\text{2020–23}} - \hat{\rho}_{\text{2017–19}} = 0.21$, 95% block-bootstrap CI $[0.09, 0.34]$
  - The CI excludes 0: the adjusted correlation did rise
  - Under the adjustment's assumptions, 44% of the raw rise is attributed to volatility; the adjusted increase remains positive: stronger co-movement, not an identified mechanism
  - The share: $1 - \dfrac{\rho^* - \hat{\rho}_{\text{2017–19}}}{\hat{\rho} - \hat{\rho}_{\text{2017–19}}} = 1 - \dfrac{0.23 - 0.02}{0.39 - 0.02} \approx 43\%$ with the rounded values; 44% with unrounded correlations

A Crypto Index and Its Divisor



Humboldt-Universität zu Berlin, where CRIX was developed

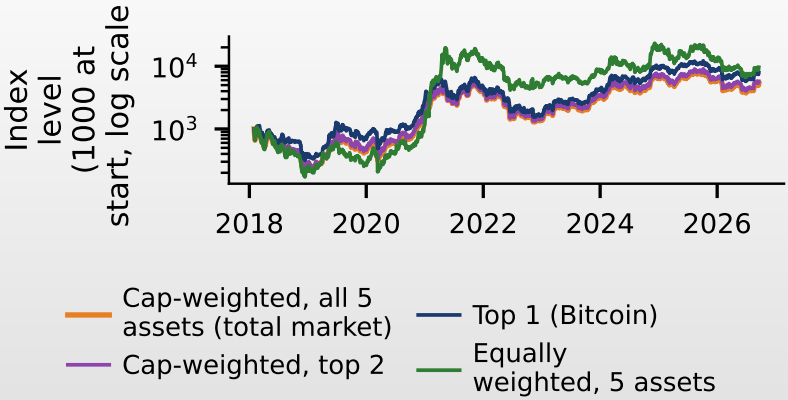
Photo: Diego Delso (2016); CC BY-SA 4.0; Wikimedia Commons

- **CRIX** Trimborn & Härdle (2018): a market-value weighted benchmark of the crypto market, published at thecrix.de
  - Review of crypto indices: Găman, Pele, Härdle & Będowska-Sójka (2026)
- Level on day t of reallocation period τ: $I_t = \sum_{i=1}^{k} P_{i,t}\, Q_{i,\tau} / D_\tau$
  - k: number of constituents; $P_{i,t}$: price of coin i on day t; $Q_{i,\tau}$: its coins in circulation, fixed within the period; D_τ: the **divisor**
  - Index return = weighted average of the constituents' returns, weights = value shares at the start of the period; a Laspeyres-type index (quantities fixed within the period)
- At a reallocation (new Q, new members) on day t the divisor is reset so that the level does not jump: $D_{\tau+1} = \sum_i P_{i,t} Q_{i,\tau+1} / I_t$
  - Here: monthly reallocation, weights from the previous month end; worked examples: Seminar 16, A1, A4

How Many Constituents? The CRIX Rule

- Candidate index CRIX(k, β), fitted to track the total market index TMI (all coins) Trimborn & Härdle (2018)
  - k: number of constituents; the k_1 largest coins enter with weight 1, the next $s = k - k_1$ with free weights $\beta = (\beta_1, \ldots, \beta_s)$, fitted to the data
- Tracking residual: the daily gap between the two index returns, $\hat{\varepsilon}_t(k, \beta) = \Delta \ln \mathrm{TMI}_t - \Delta \ln \mathrm{CRIX}_t(k, \beta)$
  - $\Delta \ln X_t = \ln X_t - \ln X_{t-1}$: the daily log return of index X
- Akaike information criterion (AIC) $\mathrm{AIC}(k) = -2 \log L\{\hat{\varepsilon}(k, \hat{\beta})\} + 2s$; the rule picks the k with the smallest AIC
  - L: likelihood of the residuals, larger when the candidate tracks TMI closely; $\hat{\beta}$: the fitted weights; $2s$: penalty for the free weights
  - More coins track better but cost parameters; k reviewed quarterly on the last three months of data, members reallocated monthly by market value
  - Two stopping rules, ECRIX (first rise of the AIC) and EFCRIX (global minimum); likelihood details in the Appendix
- Data limitation: circulating supply is public for five of our coins only (Bitcoin, Ethereum, Dogecoin, Cardano, Litecoin), so $k \leq 4$ and TMI has five members, not hundreds

Crypto Indices Since 2018 (1/2)



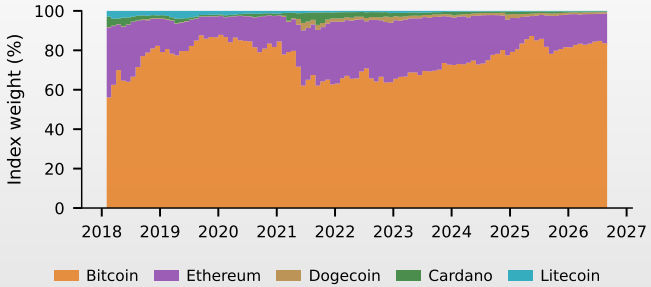
MFM_ch16_crix

Crypto Indices Since 2018 (2/2)

- Log scale; total-market index 21.4% a year over 8.6 years (compound annual growth $(I_T/I_0)^{1/\text{years}} - 1$, I_0 and I_T the first and last levels), volatility 65.7%; Bitcoin alone did better: the other coins lost weight and value
- Equal weights: 30.0% a year at 82.5% volatility, driven by 2021 (+998% vs +113%): a few extreme months, not a stable premium
- An index is only as good as its universe: today's survivors are not the leaders of 2018 Będowska-Sójka, Wójcik & Pele (2026)

MFM_ch16_crix

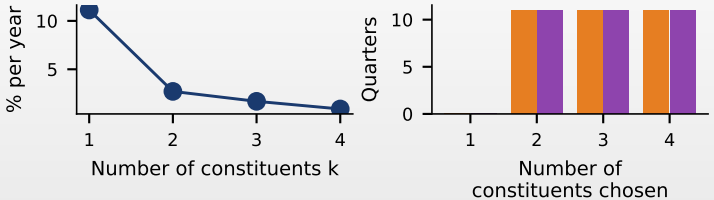
Index Weights Over Time



- Value weights at each monthly reallocation, %: Bitcoin 75.7% on average, never below 56.1%, today 82.9%; Ethereum today 15.8%; the other three together today below 2%
- Index weights are fixed at the last reallocation, from the previous month end: hence 82.9% here and the 82.5% market-value share of 18 Sep 2026
- Altcoins (coins other than Bitcoin) gain weight late in booms (2021) and lose it in busts: a value-weighted crypto index is mostly Bitcoin

MFM_ch16_crix

Tracking Error and the CRIX Rule



- Left: tracking error of the top-k index (weights 1) against the five-coin index: 11.1% a year for Bitcoin alone, 2.7% for two coins, 1.7% for three, 0.9% for four
- Right: the rule picks 2 coins in 11, 3 in 11 and 4 in 11 of 33 quarters (2 in the last); ECRIX and EFCRIX agree in every quarter

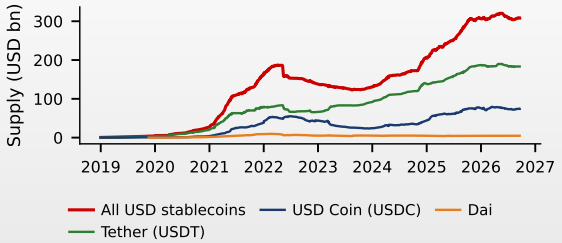
MFM_ch16_crix

How a Stablecoin Holds 1 USD

Design	Backing	Examples (USD bn, 18 Sep 2026)
Fiat-backed	Bank deposits and short US Treasury bills held by the issuer (fiat: government-issued currency, here USD)	Tether (USDT) 183.33, USD Coin (USDC) 73.92, PayPal USD (PYUSD) 2.81
Crypto-backed	Over-collateralised crypto deposits or hedged positions, managed by smart contracts	Sky Dollar (USDS) 6.49, Ethena USDe 4.74, Dai (DAI) 4.77
Algorithmic	No external collateral: supply expands and contracts against a second token	TerraUSD (collapsed in May 2022)

- Authorised traders **mint** (pay 1 USD, receive 1 coin) and **redeem** (return 1 coin, receive 1 USD) with the issuer: arbitrage keeps the market price P (USD) in the band $1 - c_r \leq P \leq 1 + c_m$, with c_r, c_m the redemption and minting costs per coin, in USD
  - Below the band: buy and redeem; above: mint and sell (evidence: Lyons & Viswanath-Natraj (2023); Tether issuance and Bitcoin prices: Griffin & Shams (2020))
- The band breaks when redemption stops or the backing is doubted; reserve and redemption rules: MiCA, Regulation (EU) 2023/1114, GENIUS Act, Public Law 119-27

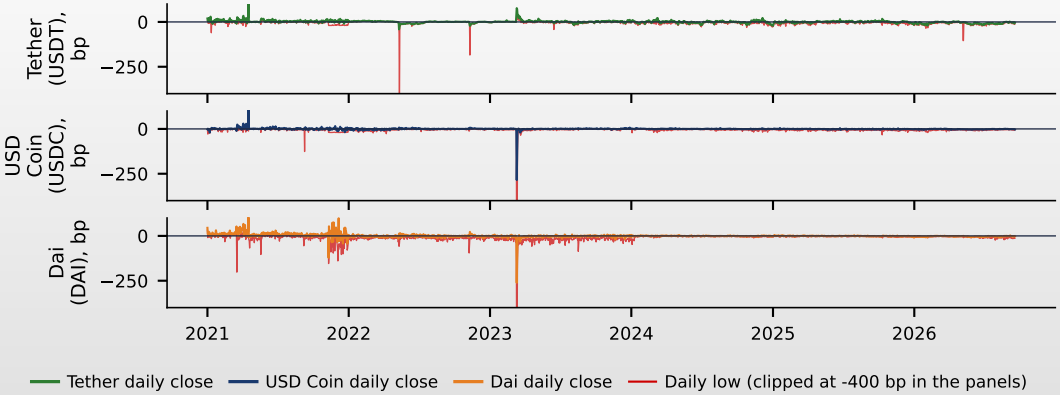
The Rise of Stablecoins



- Supply in circulation (DefiLlama): from 4.2 billion USD in January 2020 to 308 billion on 18 Sep 2026, about 74 times: the cash leg of crypto trading and DeFi
- Tether 183 billion (60%), USD Coin 74 billion (24%); fiat-backed 283, crypto-backed 22.4, algorithmic 0.15 billion for the 45 coins above 100 million USD; the smaller coins, not classified, add 2.6 billion (together the 308)
- 2022–2023 dip: 187 billion (2 Apr 2022) to 123 billion (19 Aug 2023); TerraUSD vanished, USD Coin fell from 55 to 23 billion

MFM_ch16_stablecoins

Deviations from 1 USD (1/2)



MFM_ch16_stablecoins

Deviations from 1 USD (2/2)

- Daily close and daily low minus 1 USD, in basis points (1 bp = 0.01%), 2021–2026
  - Lows below −400 bp are cut at the panel edge
- Normal times: standard deviation of the close 7.3 bp (Tether), 7.5 (USD Coin), 10.8 (Dai)
  - Beyond 50 bp on less than 1% of the days
- Stress is rare but deep: USD Coin's low −1 226 bp (11 Mar 2023), Tether's −515 bp (12 May 2022)
- A stablecoin is a claim with a short put on the reserves
  - Nothing happens most days, a large loss is possible on a few (the payoff)

MFM_ch16_stablecoins

Arbitrage Bands Imply a Threshold AR (1/2): the Model

- Deviation from the peg in basis points: $d_t = 10^4(P_t - 1)$
  - P_t: daily close in USD; e.g. $P_t = 0.997$ gives $d_t = -30$ bp
  - Inside the band $|d| \le c$ arbitrage does not pay: weak pull; outside it does: fast reversion
- Threshold autoregression (TAR), the equilibrium form of Balke & Fomby (1997): one AR(1) inside the band, another outside
  - $d_t = \begin{cases} \phi_{\text{in}}\, d_{t-1} + \varepsilon_t, & |d_{t-1}| \le c \\ \phi_{\text{out}}\, d_{t-1} + \varepsilon_t, & |d_{t-1}| > c \end{cases}$
  - c: the threshold, the half-width of the band in bp; ε_t: a shock with mean 0
  - ϕ_{in}, ϕ_{out}: persistence inside and outside the band; $|\phi| < 1$: deviations shrink, faster the closer ϕ is to 0
  - Half-life $\ln 0.5 / \ln|\phi|$: the number of days in which a deviation halves
- $\hat{c}$ by least squares over a grid of $|d_{t-1}|$ values, keeping at least 15% of the days in each regime
  - $\hat{c}$ (in bp) is an empirical regime boundary: it matches the costs only under a symmetric band $c_r = c_m$ (USD per token), as $10^4 c_r$ bp, and other frictions also move it

Arbitrage Bands Imply a Threshold AR (2/2): the Linearity Test

- H_0: $\phi_{\text{in}} = \phi_{\text{out}}$: one linear AR(1), no band; H_1: the two persistences differ
  - For a given c: $W(c)$, the Wald statistic of H_0 with White standard errors (robust to heteroskedasticity); the test uses $\sup_c W(c)$, the largest over the grid
- Under H_0 the threshold c does not enter the model: it is not identified, so $\sup_c W(c)$ is not χ^2_1 (the Davies problem)
  - χ^2_1 critical values are too small for a maximum over many c: the test would reject too often
- p-value from the fixed-regressor bootstrap of Hansen (1996), valid with heteroskedasticity
  - Artificial series $d^*_t = \hat{\varepsilon}_t \eta_t$: $\hat{\varepsilon}_t$ the residuals of the linear AR(1) (the model under H_0), $\eta_t \sim N(0,1)$ independent draws
  - The regressors d_{t-1} are held fixed; $\sup_c W^*(c)$ is recomputed on each artificial series; p = share of replications with $\sup_c W^* \geq \sup_c W$
- Sample 2021–2026: earlier quotes are thin and noisy

Threshold AR Estimates for the Three Stablecoins

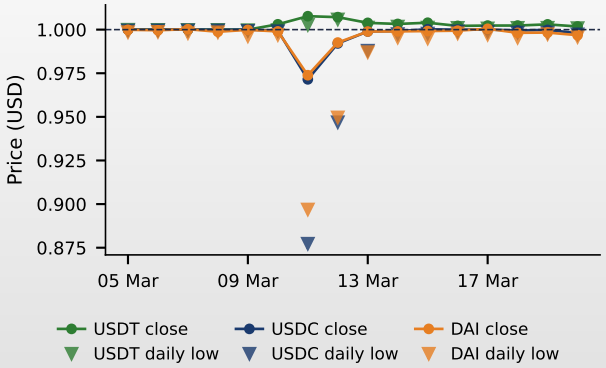
	$\hat{c}$ (bp)	$\hat{\phi}_{\text{in}}$ (SE)	$\hat{\phi}_{\text{out}}$ (SE)	sup W	p bootstrap	p (χ^2_1)	$\hat{\phi}$ linear
Tether	8.6	0.86 (0.02)	0.69 (0.11)	7.4	0.084	0.006	0.73
USD Coin	2.9	0.59 (0.04)	0.28 (0.03)	32.5	0.001	< 0.001	0.29
Dai	6.0	0.72 (0.08)	0.32 (0.07)	13.8	0.005	< 0.001	0.33

- USD Coin and Dai: linearity rejected; reversion is much faster outside the band (half-life 0.5 vs 1.3 days for USD Coin)
  - Without 9–16 March 2023, USD Coin gives $p = 0.29$ (bootstrap) but 0.019 (χ^2_1): one stress week carries the evidence, and χ^2 over-rejects (step by step: Seminar 16, B5)
- Tether: the widest band ($\hat{c} = 8.6$ bp) and slow reversion inside it, but linearity is not rejected at 5%: consistent with fewer direct redeemers Ma, Zeng & Zhang (2025)

How Precisely Is the Band Width Estimated?

- USD Coin: $\hat{c} = 2.87$ bp, while the search grid ends at 2.93 bp; 15.7% of the days lie outside, just above the 15% minimum (Tether: $\hat{c} = 8.57$ of 8.77 bp, 15.4% outside)
  - USD Coin has an interior SSR minimum, close to the upper boundary; the 15% rule does not force the estimate to equal the largest admissible threshold
  - Check sensitivity to trimming and to March 2023 before interpreting the band: proximity to the boundary alone does not make this estimate a lower bound
  - Dai, with 19.6% of the days outside, is farther from the trimming boundary
- What depends on the split and what does not
  - The half-lives (1.3 vs 0.5 days for USD Coin) move with $\hat{c}$
  - The linearity test does not: $\sup_c W(c)$ and its bootstrap already search the whole grid
- Check, with the USD Coin estimates: outside the band a deviation of -40 bp is followed on average by $0.28 \times (-40) \approx -11$ bp the next day
  - Inside, -2 bp becomes $0.59 \times (-2) \approx -1.2$ bp; half-life $\ln 0.5 / \ln 0.59 = 1.3$ days

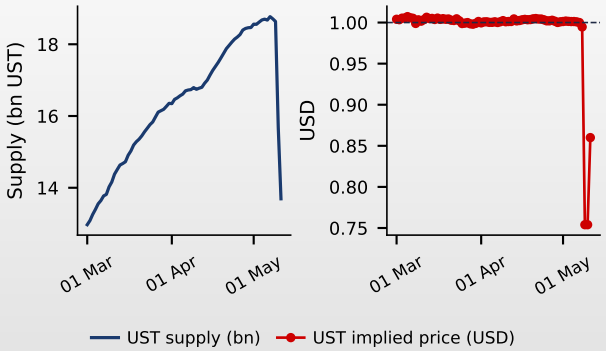
March 2023: a Bank Failure Hits USD Coin



- Daily close (lines) and daily low (triangles), 5–20 March 2023
- Friday 10 March: Silicon Valley Bank (SVB) is closed, the FDIC (Federal Deposit Insurance Corporation) is appointed receiver FDIC (2023)
  - USD Coin held part of its reserves there Diop, Chevallier & Sanhaji (2024)
- Weekend: redemptions could not settle, the band disappeared
  - USD Coin's low 0.8774 USD on Saturday 11 March (close 0.9715)
  - Dai, crypto-backed, fell to 0.8970

MFM_ch16_stablecoins

TerraUSD: an Algorithmic Stablecoin Collapses



- Supply and implied price (market value / supply), March–May 2022 (DefiLlama)
- TerraUSD (UST) held its peg by exchange with the token LUNA, not with reserves
  - Demand came from the high yields of the Anchor lending protocol Liu, Makarov & Schoar (2023)
- Supply 18.8 billion on 7 May 2022, 13.7 billion four days later (27% redeemed)
  - Implied price 0.754 USD on 10 May 2022

MFM_ch16_stablecoins

Two Stress Episodes: Doubted Backing, Failed Design

USD Coin, March 2023: the backing

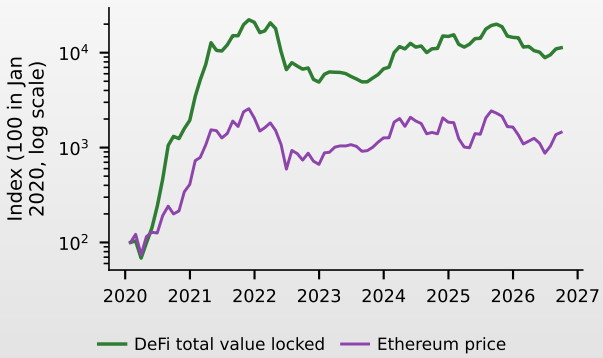
- Tether rose to 1.0077 (high 1.0296): a flight to the other large stablecoin
- Sunday 12 March: all SVB depositors protected Treasury, Federal Reserve & FDIC (2023)
  - The peg returned within two days: the event tested the backing, not the design

TerraUSD, May 2022: the design

- Redemption minted more LUNA and pushed its price down: a death spiral that wiped out 50 billion USD in three days
  - Sophisticated holders ran first and lost less Liu, Makarov & Schoar (2023)
- Without external collateral the peg rests on confidence alone: a run is self-fulfilling

- The short put in one line: per coin the holder gets $\min(1, R) = 1 - \max(1 - R, 0)$, with R the reserve value per coin in USD
  - A sure 1 USD minus a put on R with strike 1: the holder has sold that put to the issuer; $R = 0.9$ pays $1 - 0.1 = 0.9$ USD, $R \geq 1$ pays 1

Total Value Locked in DeFi (1/2)



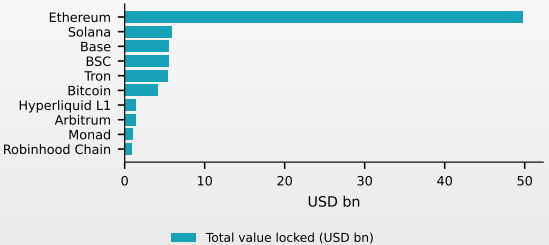
- Total value locked (TVL): the value of tokens deposited in DeFi protocols, all chains (DefiLlama), and the Ether price
  - Month ends, 100 in January 2020, log scale
- TVL: 0.6 billion USD in January 2020, 177 billion at the peak (9 Nov 2021), 88 billion today (50% below the peak)

MFM_ch16_defi

Total Value Locked in DeFi (2/2)

- Monthly $\Delta \ln \mathrm{TVL}_t = a + b\,\Delta \ln P_t^{\mathrm{ETH}} + u_t$, Newey–West SE
  - b: the % change in TVL per 1% change in the Ether price P_t^{ETH}: $\hat{b} = 0.79$ (SE 0.07), $R^2 = 0.62$, $n = 80$ months
- TVL mixes token prices with net deposits: measure activity with fixed-price quantities, net flows or fees
  - DeFi risks: leverage, liquidity mismatch, interconnection Aramonte, Huang & Schrimpf (2021)

Where the Value Is Locked



- TVL by blockchain, ten largest, 18 September 2026 (DefiLlama): Ethereum 49.8 billion USD, 57% of 88 billion; Solana 5.9, Base 5.6, BSC (BNB Smart Chain) 5.6, Tron 5.4
- Base and Arbitrum are layer-2 networks that settle on Ethereum: cheaper transactions, the same security anchor
- L1 in the chart (Hyperliquid L1): a layer-1 network, a blockchain that settles its own transactions

MFM_ch16_defi

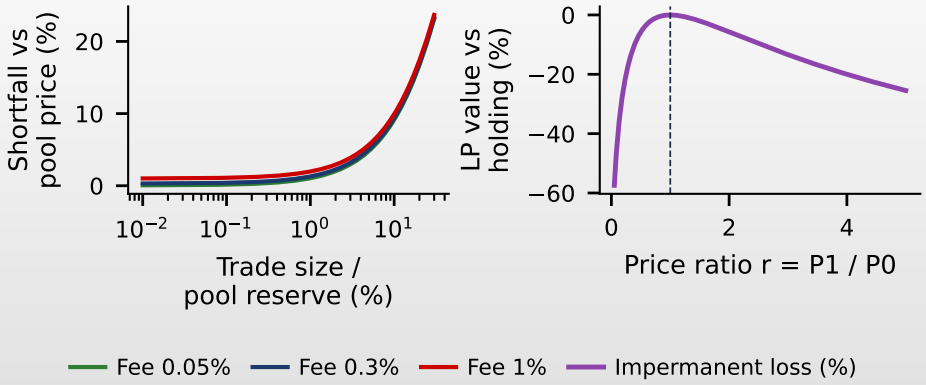
The Constant-Product Market Maker

- A pool holds reserves x of token X (e.g. Ether) and y of token Y (e.g. USD Coin); it accepts any trade that keeps $x \cdot y = k$, with the input counted net of the fee (below) Adams, Zinsmeister & Robinson (2020)
  - k: the invariant, set by the deposits; pool price (marginal): $p = y/x$ units of Y per X
- A trader sells Δx of X and receives Δy of Y; the pool keeps a fee f, so only $\gamma \Delta x$ enters the invariant, $\gamma = 1 - f$
  - The fee-adjusted input preserves xy in the pricing equation; the actual input reserve is $x + \Delta x$, so the actual reserve product increases when the fee stays in the pool: $(x + \gamma\Delta x)(y - \Delta y) = xy \;\Rightarrow\; \Delta y = \dfrac{y\,\gamma\Delta x}{x + \gamma\Delta x}$
- Execution price $\Delta y/\Delta x = p\,\dfrac{\gamma}{1+\gamma\delta}$, with $\delta = \Delta x/x$ the trade size relative to the reserve
  - **Shortfall** against the pool price: $1 - \gamma/(1+\gamma\delta)$: fee plus price impact; it grows with f and with δ
- Arbitrageurs keep p in line with other exchanges: the pool price follows the market Angeris et al. (2019)

Worked Example: a Swap in an Ether Pool

- Pool: $x = 2\,000$ Ether, $y = 5\,000\,000$ USD Coin, so $p = 2\,500$ and $k = 10\,000 \times 10^6$; fee $f = 0.3\%$
  - Sell $\Delta x = 20$ Ether ($\delta = 1\%$); after the fee, $\gamma \Delta x = 19.94$
- $\Delta y = \dfrac{5\,000\,000 \times 19.94}{2\,000 + 19.94} = 49\,358$ USD Coin, an execution price of 2 467.9 per Ether
- Shortfall: $1 - 2\,467.9/2\,500 = 1.28\%$ (0.3% fee plus about 1% price impact)
- New pool price: $(5\,000\,000 - \Delta y)/2\,020 = 2\,450.8$, i.e. -1.97%
  - Why 2 020 and not 2 019.94: the fee of 0.06 Ether stays in the reserve (Uniswap v2); the product of the reserves rises to $2\,020 \times 4\,950\,642 = 10\,000.3 \times 10^6$, slightly above k: the invariant binds for the fee-adjusted input only, and the extra is the liquidity providers' income
  - An arbitrageur now buys Ether in the pool and sells it elsewhere until p is back inside the no-arbitrage band that the fee sets around the market price (Seminar 16, A2)

Slippage and Impermanent Loss (1/2)



MFM_ch16_amm

Slippage and Impermanent Loss (2/2)

- Left: shortfall by trade size and fee; with a 0.3% fee 0.40% for 0.1% of the reserve, 1.28% for 1%, 9.34% for 10% (0.05% fee: 1.04% at 1%)
- Small trades pay the fee, large ones the price impact (about δ): liquidity is the depth of the reserve, not an order book (Chapter 10); concentrated liquidity deepens the pool near the price (Uniswap v3)
- Right: value of a liquidity provider (LP) position relative to holding the two tokens, against the price ratio $r = P_1/P_0$ (next slide)

MFM_ch16_amm

Impermanent Loss and Loss Versus Rebalancing (1/2): Impermanent Loss

- A **liquidity provider** (LP) deposits x_0 of X and $y_0 = p_0 x_0$ of Y (half and half in value) and receives a share of the pool
  - p_0: the price at the deposit; p_1: the market price later; $r = p_1/p_0$: the price relative (written P_1/P_0 in the chart and in the seminar)
  - Without fees, arbitrage moves the pool to p_1: from $xy = k$ and $y/x = p_1$, $x_1 = \sqrt{k/p_1}$, $y_1 = \sqrt{k\,p_1}$
- Two values in units of Y: the pool position $V_{\mathrm{LP}} = p_1 x_1 + y_1 = 2\sqrt{k\,p_1}$ and simply holding the deposit $V_{\mathrm{H}} = p_1 x_0 + y_0$
- **Impermanent loss**, the relative shortfall of the pool against holding, using $k = p_0 x_0^2$:

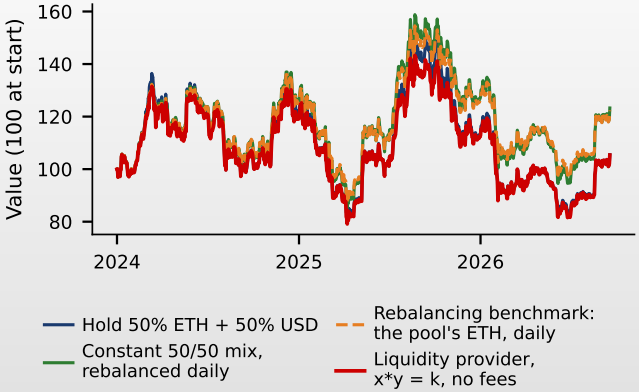
  $$\mathrm{IL}(r) = \frac{V_{\mathrm{LP}}}{V_{\mathrm{H}}} - 1 = \frac{2\sqrt{r}}{1+r} - 1 \leq 0$$

  - $\mathrm{IL} = 0$ only at $r = 1$; $r = 2$ or $r = 0.5$: -5.72%; $r = 4$ or 0.25: -20.00%
  - Symmetric in $\ln r$: a doubling and a halving of the price cost the same
  - Check: a 21% rise, $r = 1.21$, gives $2 \times 1.1/2.21 - 1 = -0.45\%$: small moves cost little, IL grows roughly with $(\ln r)^2$
- "Impermanent": it vanishes if the price returns to p_0; it becomes permanent when the LP withdraws

Impermanent Loss and Loss Versus Rebalancing (2/2): LVR

- A fairer benchmark than holding: a **rebalancing** portfolio that holds the pool's quantity of Ether at every instant and trades at market prices Milionis et al. (2022)
  - The pool rebalances too, but at stale prices set by arbitrageurs: it always trades at a loss to them
- For $x \cdot y = k$, the loss versus rebalancing (LVR) grows at the rate $\mathrm{LVR}/V = \sigma^2/8$ per unit of time Milionis et al. (2022)
  - LVR: the gap between the rebalancing portfolio and the pool; V: the pool value; σ: annualised volatility of the price
  - Example: $\sigma = 60\%$ gives $0.36/8 = 4.5\%$ of the pool value a year
- Fees must exceed $\sigma^2/8$ a year for a liquidity provider to beat this benchmark (not to earn a positive return): a volatility-selling business
- Derivation by Itô in the Appendix
  - A weighted pool $x^w y^{1-w} = k$, with $w \in (0,1)$ the value weight of X ($w = 1/2$: constant product), loses at the rate $\sigma^2 w(1-w)/2$ (Seminar 16, A7)

A Liquidity Provider in an Ether Pool (1/2)



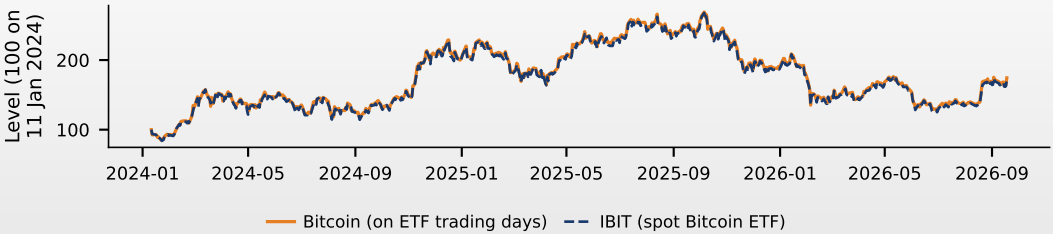
MFM_ch16_amm

A Liquidity Provider in an Ether Pool (2/2)

- Daily Ether prices from 1 January 2024
  - 100 invested half in Ether, half in USD
  - Pool with no fees; dashed: the LVR benchmark (the pool's Ether, rebalanced daily)
  - Constant 50/50 mix: weights reset daily (final value 123.3); ETH in the legend: Ether's ticker
- Ether 2 352 $\to$ 2 611 USD ($r = 1.11$): impermanent loss only -0.14%
  - Final values: holding 105.5, pool 105.4, LVR benchmark 122.4
- Realised LVR 5.83% a year against $\sigma^2/8 = 5.79\%$ ($\sigma = 68.1\%$); 95% block-bootstrap CI $[4.88, 6.98]\%$
- Impermanent loss vanishes when the price returns
  - LVR accumulates with every move: fees must earn about 5.79% a year to match the benchmark
- Why the LVR benchmark beats holding here: with R_T the LVR benchmark, V_T^{LP} the pool and V_T^{H} holding, $R_T - V_T^{\mathrm{H}} = \mathrm{LVR}_T + (V_T^{\mathrm{LP}} - V_T^{\mathrm{H}})$, here $122.4 - 105.5 \approx 17.0 + (105.4 - 105.5)$
  - The price ended near its start ($r = 1.11$), so the impermanent-loss term is tiny; after a strong trend ($r = 4$: IL $= -20.00\%$ of V_T^{H}) it can outweigh LVR, and holding can come out ahead

MFM_ch16_amm

IBIT Follows Bitcoin



- Bitcoin and IBIT on ETF trading days, both 100 on 11 Jan 2024
- A **spot ETF** holds the coins with a custodian and trades in US hours; BlackRock's IBIT charges a sponsor fee of 0.25% a year BlackRock (2026)
- **Authorised participants** create or redeem shares against coins, which keeps the price within a cost band around the net asset value (NAV: the fund's holdings at market prices, per share), as for a stablecoin

MFM_ch16_etf

Measuring How Well an ETF Tracks (1/2): the Measures

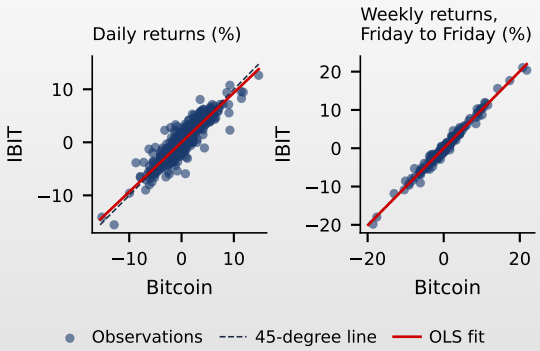
- Regression on the ETF's trading days: $r_t^{\mathrm{ETF}} = a + \beta\, r_t^{\mathrm{coin}} + u_t$
  - r_t^{ETF}, r_t^{coin}: log returns of the ETF and of its coin; a: intercept; β: how much the ETF moves per 1% move of the coin; u_t: error
  - R^2: share of the ETF's return variance explained by the coin, between 0 and 1; a perfect tracker has $\beta = 1$, $R^2 = 1$
- **Tracking error** $\mathrm{TE} = \mathrm{sd}(r_t^{\mathrm{ETF}} - r_t^{\mathrm{coin}})\sqrt{A}$: annualised volatility of the return gap, 0 for a perfect tracker
  - sd: sample standard deviation; A: return periods per year ($A = 252$ daily, 52 weekly)
- **Fee drift**: slope of $\ln(P_t^{\mathrm{ETF}}/P_t^{\mathrm{coin}})$ on time in years
  - P_t^{ETF}, P_t^{coin}: the two prices; the fund sells coins to pay its fee, so the ratio falls by about the fee each year

Measuring How Well an ETF Tracks (2/2): Results

	$\hat{\beta}$ daily	R^2 daily	$\hat{\beta}$ weekly	R^2 weekly	TE daily	TE weekly
IBIT / Bitcoin	0.93	0.82	1.01	0.98	20.8	6.2
ETHA / Ether	0.94	0.83	0.99	0.98	29.9	9.0

- TE in % a year; weekly data show a near-perfect tracker; fee drift -0.24% a year (IBIT) and -0.24% (ETHA), close to the 0.25% fee
- Daily data understate the fit: the coin and the ETF close at different times (next slide)
- IBIT trades 2.2 billion USD a day on average (maximum 10.3 billion on 5 Feb 2026)

Daily vs Weekly Returns and the Dimson Beta



- IBIT against Bitcoin: daily returns on common days (left), Friday-to-Friday returns (right); R^2 0.82 vs 0.98
- The coin's daily close and the New York close are hours apart
  - Part of day t's move reaches the ETF on $t \pm 1$: the daily OLS (ordinary least squares) beta is biased towards zero
  - Newey–West fixes the SE, not the bias
- The remedy, a regression with leads and lags: next slide

MFM_ch16_etf

The Dimson Beta: Leads and Lags

- Dimson (1979): $r_t^{\text{ETF}} = a + \sum_{j=-1}^{1} b_j \, r_{t+j}^{\text{coin}} + u_t$
  - b_j: response to the coin's return on day $t + j$ ($j = -1$: the day before, $+1$: the day after)
  - Sum beta $\hat{\beta}_D = \hat{b}_{-1} + \hat{b}_0 + \hat{b}_1$ collects the whole response
- Which lag should carry the delay? If the coin's daily close comes after the 16:00 New York close, the ETF's day-t return already contains the last hours of the coin's day $t - 1$: expect $b_{-1} > 0$ and $b_1 \approx 0$
  - IBIT: $\hat{b}_{-1} = 0.05$, $\hat{b}_0 = 0.94$, $\hat{b}_1 = 0.01$, the pattern of a coin close after New York; the close time is set by the data provider and is worth checking before reading a lead or a lag
- IBIT: $\hat{\beta}_D = 1.00$ (SE 0.023) vs OLS 0.94 on the same 671 days
  - The tracking table's 0.93 is the same-day OLS beta on all 673 ETF days, while the lead and the lag cost the Dimson sample two days: 0.9348 vs 0.9353, a gap of 0.0005 that rounding turns into 0.93 vs 0.94
  - ETHA: 0.99 (SE 0.022) vs 0.94: within one SE of 1, as the weekly betas
  - Where prices form is a separate question Hasbrouck (1995)

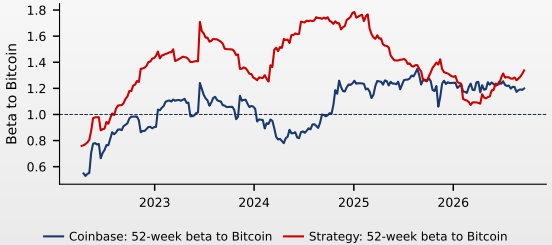
Crypto Equities: Coinbase and Strategy

- **Coinbase**: a US crypto exchange, revenue tied to trading volume; **Strategy** (formerly MicroStrategy): a software company holding Bitcoin as its main treasury asset
- Weekly two-factor regression, Newey–West SE, April 2021 – September 2026:
  $r_t^{\text{stock}} = a + \beta_{\text{BTC}}\, r_t^{\text{BTC}} + \beta_{\text{SPX}}\, r_t^{\text{S\&P}} + u_t$
  - r_t^{stock}, r_t^{BTC}, $r_t^{\text{S\&P}}$: weekly log returns of the stock, Bitcoin and the S&P 500
  - β_{BTC}, β_{SPX}: the stock's sensitivity to each, holding the other fixed

	$\hat{\beta}_{\text{BTC}}$	SE	$\hat{\beta}_{\text{SPX}}$	SE	R^2	σ_a (%)
Coinbase	0.79	0.09	2.05	0.22	0.52	84
Strategy	1.18	0.11	1.02	0.24	0.63	89

- Strategy: 1.18 times Bitcoin plus 1.02 times the S&P 500, a leveraged Bitcoin position with extra equity risk
- Coinbase: $\hat{\beta}_{\text{SPX}} = 2.05$, a high-beta US stock whose business depends on crypto activity, more an equity than a Bitcoin proxy

Beta to Bitcoin Over Time



- Rolling 52-week beta of weekly returns on Bitcoin; Strategy's ranged from 0.76 to 1.79; both stocks are more volatile than Bitcoin (55% a year on the same weeks)
- When did the beta change? A break test at an unknown date: Seminar 16, B3, B4

MFM_ch16_equities

Testing a Factor Model (1/2): the GRS Test

- Time-series regression for each test asset $i = 1, \ldots, N$: $r^e_{it} = \alpha_i + \beta_i' f_t + \varepsilon_{it}$, $t = 1, \ldots, T$
  - r^e_{it}: excess return of asset i in week t (over the Treasury bill); f_t: vector of the K traded factor returns; β_i: the asset's loadings; ε_{it}: residual
  - α_i: the mean excess return the factors leave unexplained; a correct model has $H_0 : \alpha_1 = \cdots = \alpha_N = 0$
- Gibbons, Ross & Shanken (1989): $\mathrm{GRS} = \dfrac{T-N-K}{N} \dfrac{\hat{\alpha}' \hat{\Sigma}^{-1} \hat{\alpha}}{1 + \bar{f}' \hat{\Omega}^{-1} \bar{f}} \sim F(N, T-N-K)$ under H_0
  - $\hat{\alpha}$: vector of the N estimated alphas; $\hat{\Sigma}$: residual covariance; $\bar{f}$, $\hat{\Omega}$: mean and covariance of the factors
  - Numerator: the alphas, weighted by the precision of the residuals; denominator: one plus the squared Sharpe ratio of the factors
  - A large GRS (small p-value from the F distribution) rejects zero alphas
- F is exact only for i.i.d. errors with the Normal distribution: with crypto tails, add a wild bootstrap under H_0 (▸ Appendix)

Testing a Factor Model (2/2): Fama–MacBeth

- Premium of a factor: Fama–MacBeth (FM), two passes
  - Pass 1: the betas $\hat{\beta}_i$ from the time-series regressions
  - Pass 2: each week, regress the N excess returns on the $\hat{\beta}_i$ across assets; the slope $\hat{\lambda}_t$ is that week's reward per unit of beta
  - $\hat{\lambda}$ = mean of the $\hat{\lambda}_t$: the factor premium; $\mathrm{SE}_{FM} = \mathrm{sd}(\hat{\lambda}_t)/\sqrt{T}$
- The FM standard error ignores that $\hat{\beta}_i$ is estimated: Shanken (1992) correction
  - Weak factors (small spread of β across assets) over-reject Kleibergen (2009)
- GRS judges the model, FM the premium: a priced factor in an incomplete model can fail GRS; formulas and robust versions in the ▸ Appendix

Worked Result: Eight Coins and the Crypto Market Factor

- Test assets: Bitcoin, Ethereum, XRP, BNB (the BNB Chain token), Cardano, Dogecoin, Litecoin, Chainlink; weekly excess returns over the 4-week bill (series DTB4WK in FRED, the St. Louis Fed database), $T = 450$; factor: the weekly excess return of the five-coin index in Section 3
  - CSMB (size: small minus big coins) and CMOM (momentum: past winners minus losers) of Liu, Tsyvinski & Wu (2022) need long–short portfolios over a broad cross-section with circulating market values: not buildable here
- GRS $= 2.16$: $p = 0.029$ from $F(8, 450 - 9)$ ($N + K = 8 + 1 = 9$), $p = 0.018$ from the wild bootstrap: zero alphas rejected at 5%
  - Largest alphas: Dogecoin +99% a year ($t = 1.23$), BNB +51% ($t = 1.97$)
  - No individual alpha is clearly significant: the largest $|t|$, BNB's 1.97, sits at the 5% edge of 1.96, the next is Chainlink's 1.69; the rejection is joint
- Fama–MacBeth: $\hat{\lambda} = 163\%$ a year, SE 147 (FM) and 156 (Shanken); the mean excess return of the factor was 39% ($t = 1.81$)
  - Betas range only from 0.90 to 1.40: eight coins cannot identify a premium
- Survivorship (coins chosen because they are large today) may contribute to the rejection; GRS does not identify its cause

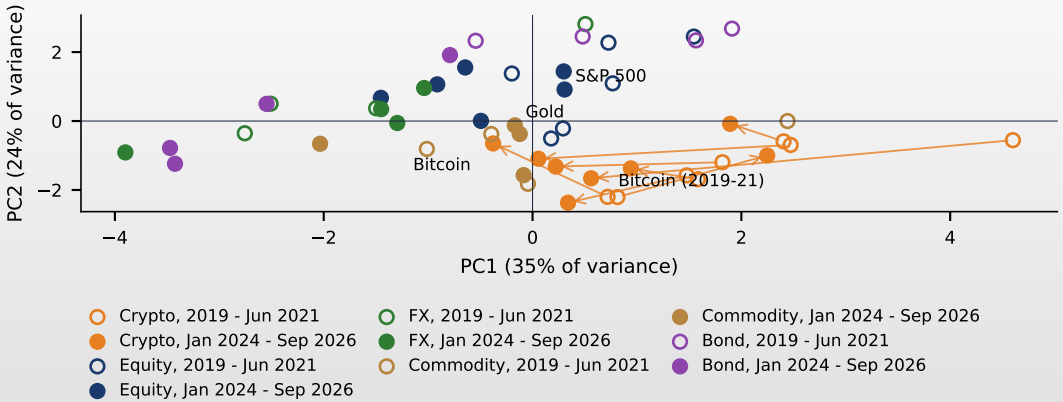
A Separate Asset Class? The Evidence So Far

- Pele et al. (2023) compare cryptos with stocks, exchange rates, commodities, bonds and real-estate indices
  - Dimension reduction: most of the variation is explained by **tail**, **memory** and **moment** factors of log returns
  - Classification methods place cryptos in a separate class, mainly because of the tail factor
  - The tail factor is a feature-space factor (76% of the variance), loading on quantiles, conditional tail expectations, variance and the α-stable tail and scale parameters (Pele et al. (2023), Sect. 3.2): it mixes scale with shape and is not a difference in Hill indices
  - Over time cryptos became more alike and more distinct from other assets
- Hedge or safe haven? A hedge is uncorrelated or negatively correlated on average; a safe haven on the worst market days Baur & Lucey (2010)
- Early evidence: Bitcoin a diversifier, a hedge only in some periods Bouri et al. (2017); tested with post-ETF data in Seminar 16, C2

A Small Replication: Features and Method

- 26 assets: 8 cryptos, 6 equity indices (including BET, the main index of the Bucharest Stock Exchange), 4 exchange rates, 4 commodities, 4 bond ETFs; two windows
  - W1: January 2019 – June 2021; W2: 11 January 2024 – 18 September 2026 (after the spot ETFs)
- Eight features of daily log returns, each asset on its own calendar
  - Moments: $\ln \sigma_a$ (log annualised volatility), skewness S, $\ln(1+K)$ (K: excess kurtosis)
  - Memory: $\hat{\rho}_1(r)$, $\hat{\rho}_1(r^2)$, lag-1 autocorrelations of the returns and of the squared returns
  - Tails: Hill indices of losses and gains (5% largest); maximum drawdown
- **Principal component analysis** (PCA) of the standardised features (Chapter 3): directions of maximum variance
  - Distance between class centroids in the full 8-dimensional standardised space

Cryptos Among Other Asset Classes (1/2)



MFM_ch16_alt_assets

Cryptos Among Other Asset Classes (2/2)

- First two principal components, PC1 and PC2 (35% and 24% of the variance)
  - Hollow: W1 (2019 – June 2021), filled: W2 (January 2024 – September 2026); FX in the legend: currencies
  - Arrows: each crypto
- PC1 measures volatility, kurtosis, drawdown depth and tail heaviness (loadings in the Appendix)
  - Cryptos sit apart in both windows, as in Pele et al. (2023)
- Point estimates move towards the others: distance to equities 3.03 → 2.63, to commodities 2.26 → 1.86; Bitcoin's volatility 77% → 47%
- Cryptos became more alike (spread around their centre 2.13 → 1.43)
  - Median Hill index of the cryptos 2.32 → 3.26; no interval is computed for this change: the next slide tests the distances on the map

MFM_ch16_alt_assets

Is the Move Significant? A Block Bootstrap of the Map

- In each window, draw the same calendar days for all 26 assets (moving blocks of 28 days: four weeks, 20 trading days)
  - Each asset keeps its own returns on those days, so cross-asset dependence is preserved
  - Recompute the eight features, the common standardisation and the centroid distances
  - 500 replications
  - Statistic: change $\Delta = d_{W2} - d_{W1}$ in the distance between the crypto centroid and each class centroid
- Equities: $\Delta = -0.39$, 95% CI $[-1.04, 0.59]$; commodities: -0.41 $[-1.38, 0.47]$; both contain 0
  - Currencies: -0.89 $[-1.67, 0.18]$; bonds: -0.62 $[-1.55, 0.92]$
  - Spread of the cryptos around their centre: -0.70 $[-1.09, 0.31]$
- Reading: all five indicative intervals contain 0; the map shows cryptos apart, but it does not establish a reliable post-ETF move towards the other classes
- Caveat: $\ln(1 + K)$ and $\hat{\rho}_1(r^2)$ have population targets only if $Er^4 < \infty$, doubtful here: they enter as sample descriptors, so the map is descriptive and these intervals are indicative

Case Study: Three Factors for Crypto Returns

- Liu, Tsyvinski & Wu (2022), *Journal of Finance*: a landmark study that applies the equity factor-model toolkit (portfolio sorts, factor regressions, GRS) to the whole cross-section of coins
  - Question: do a few crypto-specific factors price the expected returns of all coins, as market, size and momentum do for stocks?
- It builds on Liu & Tsyvinski (2021): aggregate crypto returns barely load on stock, currency or metal factors, but are predicted by crypto momentum and attention
  - The step: from the time series of three coins to the cross-section of 1,827 coins
- Why it matters: a factor model is the benchmark for every claimed crypto anomaly and for judging crypto funds
  - If coins are segmented from equities, stock factors are the wrong benchmark
- Link to this chapter: the market-value index of Section 3; the crypto market factor CMKT and the CSMB and CMOM factors (next slide), the GRS test of Section 7; data-quality threats in the ▸ Appendix

Case Study: Building Crypto Factors from 1,827 Coins

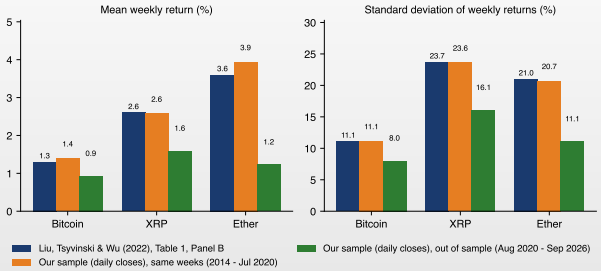
- Data: daily prices, volumes and market values of active and defunct coins from an aggregator of over 200 exchanges (Liu, Tsyvinski & Wu (2022), Sect. 2); market value of at least 1 million USD: 1,827 coins, first week of 2014 – July 2020
  - Weekly returns from daily closes: 52 weeks a year, the last one of eight days (nine in leap years); CMKT: value-weighted return of all coins minus the one-month bill rate
- Anomalies: 24 characteristics from the equity literature that need only price, volume and market value (Table 2); each week, value-weighted quintile portfolios and a top-minus-bottom spread
  - CSMB: bottom 30% minus top 30% of coins by market value; CMOM $= \frac{1}{2}(\mathrm{SH} + \mathrm{BH}) - \frac{1}{2}(\mathrm{SL} + \mathrm{BL})$ from 2×3 sorts on size and past three-week return (S/B: small/big coins; H/L: high/low past three-week return; Liu, Tsyvinski & Wu (2022), Sect. 4.1, Eq. 1)
- Tests: the ten significant spreads regressed on one-, two- and three-factor models (Tables 8–9); joint GRS tests on all 24 spreads (Table 11)
  - The one-factor model, CMKT alone, is the crypto CAPM (capital asset pricing model)
  - Multiple testing: *p*-values adjusted for the many strategies tried (the error rate used: ▸ Appendix)

Case Study: Size and Momentum Explain the Anomalies

Result	Numbers in the paper	Reading
Weekly returns (Table 1)	mean, s.d.: market 1.3%, 11.2%; Bitcoin 1.3%, 11.1%; XRP 2.6%, 23.7%; Ether 3.6%, 21.0%	the market behaves like Bitcoin
Size (Table 3)	top minus bottom quintile -5.8% a week ($t = -2.45$)	premium in the smallest coins
Momentum (Table 4)	three-week winners minus losers 3.1% a week ($t = 2.65$)	winners keep winning
Crypto CAPM (Table 8)	all ten alphas significant; momentum $\alpha = 3.0\%$ ($t = 2.55$), $R^2 = 0.001$	the market alone fails
Three factors (Table 9)	no alpha significant; momentum $\alpha = 1.1\%$ ($t = 1.41$), CMOM loading 1.12, $R^2 = 0.58$	size and momentum explain the ten anomalies
Joint test (Table 11)	GRS on 24 spreads: $p = 0.033$ (Fama–French six factors), $p = 0.194$ (crypto three)	stock factors are the wrong benchmark

- Momentum is a large-coin effect: 3.2% a week above the median size, insignificant below (Table IA.10; IA: the paper's Internet Appendix), the opposite of stocks

Case Study: Table 1 Rebuilt on the Course Data

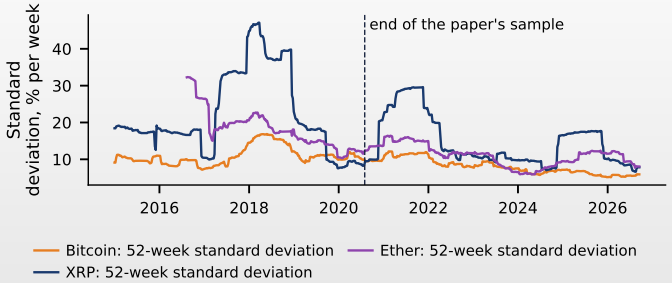


- Largest gap, the Ether mean: 0.26 i.i.d. standard errors (SE 1.29 points)
  - This checks the data and the calendar, not the factors
- Out of sample (319 weeks): means 0.9, 1.6, 1.2%
  - S.d. 8.0, 16.1, 11.1%: Ether's volatility roughly halves

- Weekly simple returns on the paper's calendar
  - Blue: Table 1, Panel B of Liu, Tsyvinski & Wu (2022)
  - Orange: our sample (daily closes), same weeks
  - Green: August 2020 – 16 Sep 2026

MFM_ch16_ltw

Case Study: Volatility Before and After the Paper's Sample



- Standard deviation of weekly returns over the last 52 weeks; dashed: end of July 2020, the end of the paper's sample; peaks: XRP 47.1% (Mar 2018), Bitcoin 16.9% (May 2018) a week
- Last window: Bitcoin 5.9%, XRP 7.8%, Ether 8.0% a week; volatility moves in regimes: factor tests across 2020 need HAC errors and a break test at an unknown date (Seminar 16, C3)

MFM_ch16_ltw

Replicating Crypto Factors

- **Replicate** with the course data: daily closes of the nine non-stable coins (Bitcoin, Ether, XRP, BNB, Cardano, Solana, Dogecoin, Litecoin, Chainlink); one-month Treasury bill rate from the Kenneth French Data Library, the bill of the paper's CMKT (the GRS test of Section 7 uses the 4-week bill DTB4WK instead)
  - Table 1, Panel B for Bitcoin, XRP and Ether: previous slides; next, the median, skewness and kurtosis (kurtosis as a sample descriptor only, since $\alpha < 4$ is plausible), and the CMKT proxy: the five-coin market-value index of Section 3
  - Data limit: CSMB and CMOM need breakpoints over hundreds of coins with market values, dead coins included; the nearest published design is the top-20 tercile of Table IA.13, but our nine coins are today's survivors

Limitations and Out-of-Sample Tests

- **Limits** stated by the authors (Section 6.2)
  - A short sample, 2014 – July 2020, while investors were still learning about the technology: only the split-half check (same signs)
  - Most coins cannot be shorted; for the top 20 coins, after trading costs, about 60% of the returns of long–short strategies that short Bitcoin remain (Table IA.16)
  - With prices from three reputable USD exchanges, the market-value size spread is no longer significant (Table IA.15)
- **Open extension**, out of sample, August 2020 – September 2026 (the 2022 collapses, the spot ETFs)
  - Does three-week momentum among large coins survive?
  - Is CMKT still segmented from equities?
  - Crypto-CAPM alphas; the segmentation regression of footnote 16 of Liu, Tsyvinski & Wu (2022) on the global five factors Fama & French (2015) and five exchange rates; a break test at an unknown date (Seminar 16, C3)

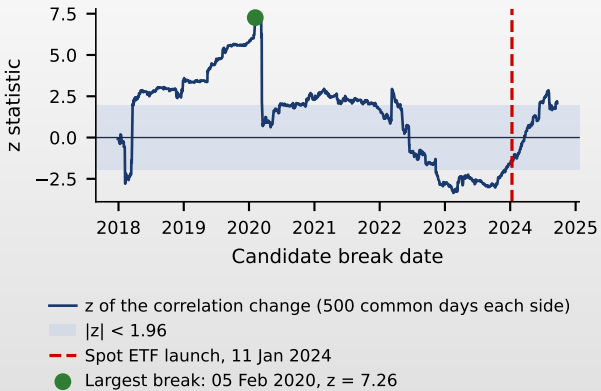
An Open Question and the Discovery Loop

- Did the US spot Bitcoin ETFs (since 11 January 2024) integrate Bitcoin with equities?
  - In the mean?
  - In the tails?
  - In price discovery?
- A live question: Lee & Lee (2025) study the shifted links after the approval; Kang & Lee (2026) title their answer "institutionalization without integration"
- Formal version, with ρ_{pre}, ρ_{post} the correlations before and after the launch: $H_0 : \rho_{\text{post}} = \rho_{\text{pre}}$ against $H_1 : \rho_{\text{post}} > \rho_{\text{pre}}$; the same for the lower-tail dependence λ_L of Chapter 6, the probability of an extreme Bitcoin loss given an extreme equity loss
  - Falsified if the data choose a break far from January 2024, or if placebo dates give changes as large
- Identification: the launch is an event, not an experiment; it overlaps with the rate cycle, the April 2024 halving and flows that respond to prices
- The discovery loop in one pass, the assistant drafts and the researcher verifies
  - literature (resolve every DOI, the digital object identifier of a paper), competing hypotheses, code that first replicates $\hat{\rho} = 0.39$ and 0.38, a pre-registered robustness plan, a referee-style critique; reusable prompts in the Appendix

What the Researcher Must Check

- **Calendars**: returns computed on each calendar and then matched pair Bitcoin's Sunday-to-Monday move with the S&P 500's Friday-to-Monday move; align prices on common trading days first, then compute returns
- **Closing times and volatility**: non-synchronous closes bias daily correlations towards zero, high volatility raises them Forbes & Rigobon (2002): check weekly returns, the lead–lag correlation and the adjusted ρ^*
- **Survivorship and supply**: use the universe as it was on each date and the circulating supply, as in the CRIX rule
- **Invented references and numbers**: every DOI resolves on Crossref with the same title; every number is regenerated by code in the repository

Mini-Case: Did the ETFs Break the Correlation? (1/2)



- Assistant's claim: "the spot ETFs integrated Bitcoin with equities"
  - $\hat{\rho} = 0.39$ in 2020–2023 ($n = 1\,006$), 0.38 since 11 January 2024 ($n = 674$)
- Difference -0.01, 95% CI $[-0.12, 0.13]$ (circular block bootstrap, blocks of 20 days)

ai_discovery_case.py

Mini-Case: Did the ETFs Break the Correlation? (2/2)

- Chart: for each candidate date τ, the statistic $z(\tau)$
  - $z(\tau) = [\operatorname{atanh} \hat{\rho}_{\text{after}} - \operatorname{atanh} \hat{\rho}_{\text{before}}] / \sqrt{2/(h-3)}$
  - Correlations over $h = 500$ days on each side of τ
  - atanh (Fisher) makes each estimate roughly Normal with variance $1/(h-3)$, so $z \approx N(0,1)$ without a change
- Largest $z(\tau)$ on 5 Feb 2020 ($z = 7.26$)
  - At the ETF date $z = -1.48$ and 82% of candidate break dates give a larger $|z|$: the scan points to early 2020, where the sup-Wald test places the break
  - Tails and price discovery remain open
- The ± 1.96 band of the chart is pointwise: valid for one date fixed in advance, not for the largest of more than a thousand overlapping dates
  - The maximum needs its own critical value, simulated or bootstrapped, as for $\sup W$; with heavy tails and volatility clustering the variance of $\operatorname{atanh} \hat{\rho}$ is typically larger than $1/(h-3)$, the i.i.d. Normal value
  - So the scan is descriptive; the inference is the sup-Wald test

Your Turn: a Project Seed

- **Question**: did institutionalisation change Bitcoin's **tail** dependence and **price discovery** with US equities?
- **Replicate first**: $\hat{\rho} = 0.39$ and 0.38 from this chapter, then one headline estimate of Lee & Lee (2025) or Kang & Lee (2026) on their own sample
  - Extend: lower-tail dependence from copulas (Chapter 6), weekly returns, IBIT–Bitcoin lead–lag in US trading hours
- **Deliverables**: a pre-registered plan committed before the results, a notebook that regenerates every number, effect sizes with CIs and the full list of specifications
- **AI log**: prompts that shaped the design in `AI_USE.md`; each wrong claim (an invented DOI, returns on mismatched calendars) and the test that caught it in `AI_ERRORS.md` (templates)

Key Takeaways

- Crypto returns: 3–7 times the volatility of shares, drawdowns of 80–96%, heavy tails in all six assets (Hill near 2.7), calmer weekends; the link with equities broke in March 2020, not at the ETF launch
- Indices: a Laspeyres-type index with a divisor; CRIX picks its size by an AIC on tracking residuals; value weights make crypto indices mostly Bitcoin
- Stablecoins: the peg is an arbitrage band that holds only while redemption works; a threshold AR fits it, tested with a bootstrap, not with χ^2
- DeFi: $xy = k$ prices a swap as fee plus price impact; providing liquidity is selling volatility, at $\sigma^2/8$ a year
- ETFs track one for one once closing times are handled; cryptos remain a separate class in distribution, a risk asset in co-movement

Check Your Understanding, and What Comes Next

Retrieval questions

- Why is Bitcoin's sample kurtosis not reported?
- What replaces it?
- Why can a threshold or a break date not be tested with χ^2 critical values?
- When does a stablecoin's arbitrage band stop holding the peg?
- Why does a liquidity provider fall behind the rebalancing benchmark even if the price returns to p_0?
- Why does a daily regression understate how well IBIT tracks Bitcoin?

Next chapters

- Chapter 17, bubbles and crashes: explosive-root tests and log-periodic power laws, crypto booms as a main case
- Chapter 18, systemic risk: spillovers between crypto, banks and equity markets
- Revisit
  - VaR and ES for crypto (Chapter 7, Pele & Mazurencu-Marinescu-Pele (2019b))
  - Liquidity on crypto exchanges (Chapter 10, Agakishiev, Härdle & Pele (2025))
- Reading
  - Härdle, Harvey & Reule (2020), Franke, Härdle & Hafner (2019), Schär (2021)

Tail Index, Inference Conditions (1/2)

- Regularly varying tail: $P(L > x) = x^{-\alpha}\ell(x)$, with ℓ slowly varying ($\ell(tx)/\ell(x) \to 1$ as $x \to \infty$, for every $t > 0$): the exact power law $Cx^{-\alpha}$ up to a slowly changing factor
  - n: number of losses; $k \to \infty$ and $k/n \to 0$: more tail observations, but a vanishing share of the sample
  - Second-order condition $\sqrt{k}\,A(n/k) \to 0$: $A(\cdot)$ measures how fast the tail approaches an exact power law; the condition makes the bias negligible
- Then $\sqrt{k}(\hat{\alpha}_H - \alpha) \Rightarrow N(0, \alpha^2)$, which gives $\mathrm{SE} = \hat{\alpha}_H/\sqrt{k}$
  - $\Rightarrow$: convergence in distribution; $N(0, \alpha^2)$: the Normal distribution with mean 0 and variance α^2
  - The 5% share is a finite-sample tuning choice, not implied by the theory
- With volatility clustering the variance depends on extremal dependence and is typically larger Drees (2000); a block bootstrap needs further tail and block-length conditions
  - Choose k from a Hill plot, $\hat{\alpha}_H$ against k, in a region where the curve is flat Drees, de Haan & Resnick (2000)
  - Bitcoin: $\hat{\alpha}_H = 2.73$ ($k = 78$), SE 0.31; with $k = 2.5\%$ or 10%: 3.48 and 2.59: the choice of k matters as much as the sampling error

Back

Tail Index, Inference Conditions (2/2)

- Block bootstrap intervals: up to 4.48 for the S&P 500; bootstrap CIs for the sample kurtosis are invalid when $\alpha < 4$ (bootstrap failure with infinite variance, Athreya (1987))
  - Dogecoin: sample kurtosis 106.4, driven by a few social-media jumps, with $\hat{\alpha}_H = 2.83$: it estimates nothing stable
- Equal tails, $H_0 : \alpha_1 = \alpha_2$: $z = (\hat{\alpha}_1 - \hat{\alpha}_2)/\sqrt{\mathrm{SE}_1^2 + \mathrm{SE}_2^2}$
  - Indices 1 and 2: the two assets; i.i.d. SE, independent estimates; under H_0, z is approximately $N(0,1)$: $|z| > 1.96$ rejects at 5%
  - Bitcoin vs S&P 500: $z = -0.28$; Dogecoin vs S&P 500: $z = -0.07$
- Overlapping CIs are not a test of equality, and non-rejection is not equivalence
  - Equivalence needs a dependence-aware interval for the difference, robust to the choice of k

Back

The sup-Wald Test, Null Limit and Critical Values

- Its limit is $\sup_r \|B(r) - rB(1)\|^2/[r(1-r)]$, with $r = \tau/T$ the break fraction and B a standard Brownian motion with one component per parameter
  - The sup runs over $r \in [0.15, 0.85]$, the trimming of the main slide; $B(r) - rB(1)$ is a Brownian bridge, zero at $r = 0$ and $r = 1$
  - Critical values simulated from this limit: 11.55 at 5%, 15.18 at 1%; they do not depend on T, so the same values serve every sample
- Why the trimming: near $r = 0$ or $r = 1$ one regime has very few days, $r(1-r) \to 0$ and the supremum over all of $(0, 1)$ would be infinite

Back

The Published CRIX Rule in Detail

- Total market index TMI (all coins) and candidate indices CRIX(k, β): the k_1 largest coins at weight 1, the next $s = k - k_1$ with free weights β (Trimborn & Härdle (2018), Sect. 3, eqs. (11)–(12) in Trimborn & Härdle (2018, author version))
  - $\hat{\varepsilon}_t(k, \beta)$: the tracking residual of the main slide; $\hat{\beta}$ minimises $\|\hat{\varepsilon}(k, \beta)\|^2$, the sum of squared tracking residuals
- $\mathrm{AIC}(k)$ of the main slide (eq. (20)); L evaluates the residuals under a kernel density estimated from the base residuals $\hat{\varepsilon}(k_1, 1)$
  - Epanechnikov kernel, a parabola $K(u) = \frac{3}{4\sqrt{5}}(1 - u^2/5)$ for $|u| \leq \sqrt{5}$ (variance 1); Sheather–Jones plug-in bandwidth, a data-driven choice of the smoothing width
  - Under Normal errors it reduces to $T \ln \hat{\sigma}_k^2 + 2s$, with T days and $\hat{\sigma}_k^2$ the residual variance (Trimborn & Härdle (2018, author version), Appendix, eq. (35))
- ECRIX: $k_1 = 1$, steps of one coin, stop at the first rise of the AIC (Trimborn & Härdle (2018, author version), Sect. 5, rule (26)); EFCRIX: global minimum (rule (27))

Back

Peg Deviations, Summary Statistics

	s.d. (bp)	$\lvert d_t \rvert > 50$ (%)	lowest close (bp)	lowest low (bp)
Tether	7.3	0.14	−41	−515
USD Coin	7.5	0.14	−285	−1 226
Dai	10.8	0.57	−261	−1 030

- Daily close, 1 January 2021 – 18 September 2026; $d_t = 10^4(P_t - 1)$ in basis points; lowest low from the daily low price

Back

Tokenisation of Real-World Assets

- **Tokenised yield-bearing instruments**: claims on Treasury bills and cash issued as tokens on a blockchain; the legal form differs
  - Fund shares: BlackRock `BUIDL` 2.38, Circle `USYC` 2.55; a tokenised secured note issued by a special-purpose vehicle: Ondo `USDY` 2.22 (Ondo) (billion USD, 18 Sep 2026, DefiLlama)
- Unlike a payment stablecoin, they pay the Treasury yield to the holder
- Benefits: settlement 24/7, use as collateral in DeFi; risks: smart-contract bugs, legal ownership, liquidity in stress Aramonte, Huang & Schrimpf (2021)
- Token economics and adoption: Cong, Li & Wang (2021)

Back

Decentralised Finance in Five Layers

- Schär (2021) describes DeFi as a stack of layers
  - **Settlement**: the blockchain and its native coin (Ether)
  - **Assets**: tokens issued on it (stablecoins, governance tokens)
  - **Protocols**: exchanges, lending, derivatives as smart contracts
  - **Applications** and **aggregators**: user interfaces that route trades across protocols
- Promise: open access, transparency, composability (protocols plug into each other)
- Risks: high leverage, liquidity mismatches, built-in interconnectedness, no shock absorber Aramonte, Huang & Schrimpf (2021)
  - "Decentralisation illusion": governance tends to concentrate power Aramonte, Huang & Schrimpf (2021)

Back

Deriving LVR $= \sigma^2/8$

- Pool value at market price P after arbitrage: $V(P) = Px + y = 2\sqrt{kP}$, concave: $V'(P) = \sqrt{k/P} = x$, $V''(P) = -\frac{1}{2}\sqrt{k}\,P^{-3/2}$
  - V', V'': first and second derivatives in P; $V'(P)$ is the pool's Ether holding, $V'' < 0$ its negative gamma
  - Rebalancing portfolio, value R_t: holds the same $x_t = V'(P_t)$ Ether at every instant, so $dR_t = V'(P_t)\,dP_t$
- Price: geometric Brownian motion $dP_t = \mu P_t\,dt + \sigma P_t\,dW_t$; Itô's lemma: $dV_t = V'\,dP_t + \frac{1}{2}V''\,\sigma^2 P_t^2\,dt$
  - μ: drift; σ: volatility; W_t: standard Brownian motion; $V_t = V(P_t)$
  - Subtracting: $d(R_t - V_t) = -\frac{1}{2}V''(P_t)\,\sigma^2 P_t^2\,dt = \frac{1}{4}\sqrt{kP_t}\,\sigma^2\,dt = \dfrac{\sigma^2}{8}\,V_t\,dt$
- LVR is the pool's negative gamma: a short-volatility position that loses $\sigma^2/8$ of its value per unit of time, whatever the drift Milionis et al. (2022)
- Assumptions: continuous geometric Brownian motion, frictionless arbitrage, no fees, continuous blocks
  - With fees and discrete blocks, arbitrageurs trade only when the price leaves the fee band, which lowers the loss Milionis, Moallemi & Roughgarden (2025)

Back

The Liquidity-Provider Benchmarks in Detail

- Realised LVR rate $\sum_t (\Delta R_t - \Delta V_t)/V_{t-1}$ per year; cumulative LVR 17.0 points over 2.7 years
  - ΔR_t, ΔV_t: daily changes in the values of the rebalancing portfolio and of the pool; V_{t-1}: the pool value the day before
- The match with theory is not a test: each day's loss is $g_t/2 - (\sqrt{1+g_t} - 1) \approx g_t^2/8$, with g_t the daily simple return of Ether, so the realised LVR is the realised variance over 8 (bootstrap correlation of the two: 0.9981)
- 2.7 years pin the realised rate down only to about ± 1 point: CI $[4.88, 6.98]\%$
- The constant 50/50 mix (final value 123.3) is a different benchmark with the same instantaneous rate: its log growth exceeds the pool's by $\sigma^2/8$, here 5.79% a year

Back

Crypto Factors, the Literature

- Value from on-chain adoption and segmentation from traditional markets Cong, Karolyi, Tang & Zhao (2026); mispricing versus risk compensation Babiak & Bianchi (2025)
- Fundamentals: network value follows Metcalfe's law in the long run, not in the short run Pele & Mazurencu-Marinescu-Pele (2019a); adoption and token value reinforce each other Cong, Li & Wang (2021)
- Two different questions, tested separately on the next slides
  - Model adequacy: are all test-asset alphas zero for the chosen factors? (time series, GRS)
  - Risk premium: is a factor's premium nonzero? (a traded factor: its mean excess return; cross-section: Fama–MacBeth λ)

Back

Time-Series Test, GRS

- Model, H_0: $\alpha_1 = \cdots = \alpha_N = 0$ and the GRS statistic Gibbons, Ross & Shanken (1989): main slide *Testing a Factor Model (1/2)*
  - $\hat{\Sigma}$, $\hat{\Omega}$: residual and factor covariances, both divided by T
- The F distribution is exact only if the errors, conditional on the factors, are i.i.d. with the Normal distribution; crypto returns have $\hat{\alpha}_H$ near 3 and volatility clustering, so these assumptions are doubtful for the residuals
  - Robust alternatives: a GMM (generalized method of moments) Wald test with HAC covariance, or a wild bootstrap under H_0: $r_t^{e*} = \hat{\beta}' f_t + \hat{\varepsilon}_t \eta_t$
  - r_t^{e*}: artificial excess returns with zero alphas; $\hat{\varepsilon}_t$: the residual vector of week t; $\eta_t = \pm 1$ with probability 1/2 (Rademacher), one draw for all assets (keeps the cross-correlation)
  - Valid if the residuals are serially uncorrelated (martingale differences), otherwise a block or dependent wild bootstrap
- $\hat{\Sigma}$ must be invertible and precise: N much smaller than T; weekly crypto data since 2018 give $T \approx 450$

Back

Cross-Section (1/2), Fama–MacBeth and Shanken

- Two passes: $\hat{\beta}_i$ from time series; then each week $r^e_{it} = \gamma_{\mathbf{0}t} + \hat{\beta}_i' \lambda_t + e_{it}$ across the N assets
  - $\gamma_{\mathbf{0}t}$: intercept, the excess return of a zero-beta asset; λ_t: the reward per unit of beta in week t; e_{it}: pricing error
  - $\hat{\lambda}$ and SE_{FM} as on the main slide; SE_{FM} is too small because $\hat{\beta}_i$ is estimated (errors in variables)
- Shanken (1992), for $\hat{\theta} = (\hat{\gamma}_{\mathbf{0}}, \hat{\lambda}')'$ and $X = [1, \hat{\beta}]$: $\widehat{\mathrm{Var}}(\hat{\theta}) = \frac{1}{T}\left[(X'X)^{-\mathbf{1}} X' \hat{\mathbf{\Sigma}} X (X'X)^{-\mathbf{1}} (1 + \hat{\lambda}' \hat{\mathbf{\Sigma}}_f^{-\mathbf{1}} \hat{\lambda}) + \hat{\mathbf{\Sigma}}_f^*\right]$
  - 1: a column of ones; $\hat{\beta}$: the $N \times K$ matrix of betas; $\hat{\mathbf{\Sigma}}$: residual covariance; $\hat{\mathbf{\Sigma}}_f$: factor covariance; $\hat{\mathbf{\Sigma}}_f^* = \begin{pmatrix} 0 & 0 \\ 0 & \hat{\mathbf{\Sigma}}_f \end{pmatrix}$
  - $\widehat{\mathrm{Var}}(\hat{\lambda})$ is its lower $K \times K$ block; the residual part is multiplied by $1 + \hat{\lambda}' \hat{\mathbf{\Sigma}}_f^{-\mathbf{1}} \hat{\lambda} > 1$, the factor part $\hat{\mathbf{\Sigma}}_f^*$ is added once
  - Assumptions: residuals serially uncorrelated and homoskedastic conditional on the factors

Back

Cross-Section (2/2), Weak Factors and Multiple Testing

- Weak or useless factors (small spread of β, e.g. attention proxies): λ is weakly identified and FM t-statistics over-reject
  - Use identification-robust tests of risk premia Kleibergen (2009)
- Multiple testing: dozens of crypto "anomalies" have been tried; a new factor should clear $|t| > 3$ Harvey, Liu & Zhu (2016), or control the false discovery rate
  - $t = \hat{\lambda}/\mathrm{SE}$; the usual 1.96 cut-off controls the error of one test, not of dozens
  - False discovery rate: the expected share of false findings among the rejected hypotheses
  - In the case study, Liu, Tsyvinski & Wu (2022): p-values adjusted under a generalised familywise error rate (k-FWER, probability of k or more false rejections, Lehmann & Romano (2005)), weaker than Holm's FWER (probability of at least one false rejection)

Back

Data-Quality Threats to Identification

- **Wash trading**: fake volume on unregulated exchanges inflates volume- and liquidity-based factors Cong, Li, Tang & Yang (2023)
  - Filter the exchanges and volumes using the rules of the cited paper before building any factor
- **Which price?** Large, recurrent gaps between exchanges, much larger across countries than within them Makarov & Schoar (2020)
- **Survivorship**: today's large coins are the winners; dead coins are missing from most samples Będowska-Sójka, Wójcik & Pele (2026)
- Remedy: fix the investable universe ex ante (exchange list, volume filter), keep delisted coins, report results by sub-universe

Back

The Asset-Class Map, PC1 Loadings

Feature (standardised)	PC1 loading
$\ln \sigma_a$, log annualised volatility	0.44
$\ln(1+K)$, kurtosis	0.52
Hill index of losses	−0.48
Hill index of gains	−0.34
Maximum drawdown (negative number)	−0.42
Skewness S	0.14
$\hat{\rho}_1(r^2)$	0.08
$\hat{\rho}_1(r)$	−0.05

- A loading is the weight of a standardised feature in PC1; the eight squared loadings sum to 1, and the sign is chosen so that PC1 rises with volatility
  - High PC1: high volatility and kurtosis, low Hill indices (heavy tails), a very negative maximum drawdown (deep falls); skewness and the two autocorrelations hardly enter
  - The PC2 loadings are not reported in the course output

Back

Common Pitfalls

- Reporting sample kurtosis as a population estimate when $\hat{\alpha}_H < 4$: if $\alpha < 4$ it estimates nothing and its bootstrap CI is invalid; assess first whether $Er^4 < \infty$ is plausible
- Testing a threshold or a break date with χ^2 critical values: the nuisance parameter is not identified under H_0
- Comparing daily closes recorded at different times and concluding that an ETF tracks badly
- Reading TVL growth as adoption without separating the valuation of the tokens from net deposits
- Treating a stablecoin as risk-free cash: the peg is only as good as redemption and reserves
- Total instead of circulating supply in a market-value index; survivorship: studying only today's large coins ignores the many that failed Będowska-Sójka, Wójcik & Pele (2026)

Back

The Discovery Loop, Reusable Prompts (1/2)

- **Literature**: map what is known, then verify every reference
  - "`List 2023-2026 papers that test for a change in Bitcoin-equity co-movement at the US spot ETF approval: DOI, sample, method, sign of the effect.`"
  - Resolve each DOI on Crossref, open the paper, find the number it is cited for; search with Semantic Scholar, Elicit or Google Scholar
- **Hypotheses**: competing mechanisms, each with an implication the others do not share
  - "`Give three competing mechanisms for the post-2020 Bitcoin-S&P 500 correlation and, for each, one testable implication that separates it from the others.`"
  - E.g. ETF flows predict an effect concentrated in US trading hours; a common monetary factor predicts co-movement with rate surprises already before 2024
- **Data and code**: on the course data, replicating a known number first
  - "`Take the daily closes of Bitcoin and the S&P 500 from the course data, align the prices on common trading days, then compute log returns and their 2020-2023 correlation.`"
  - Replicate first: this chapter reports 0.39 (2020–2023) and 0.38 (since the ETFs); no new result before the code hits both

Back

The Discovery Loop, Reusable Prompts (2/2)

- **Robustness**, fixed before the post-ETF results are seen
  - "`Draft a pre-registered robustness plan: break-date scan, placebo launch dates, weekly returns, tail dependence, block bootstrap; fix every tuning choice now.`"
  - The plan fixes windows, block length and the family of tests: the number of specifications is known and the FWER can be controlled
- **Critique**: ask the assistant to attack the result
  - "`Act as a referee: give the three strongest reasons this change in correlation could be spurious, and the test that would settle each.`"
  - Expect: non-synchronous closes, correlation that rises mechanically with volatility Forbes & Rigobon (2002), confounding events
- **Report**: effect size with a CI, every specification tried, what failed
  - The prompts that shaped the design go into `AI_USE.md`

Back

Further Study on Quantinar

Advanced courses on quantinar.com

- Are Cryptos becoming alternative Assets?
- Introduction to Blockchain and Cryptocurrencies
- Comparison of Cryptocurrency Indices
- Cryptocurrency as an Asset Class
- Smart Contracts
- Stabilising Onchain Stablecoin Risks
- Risk Premia in the Bitcoin Market

Code and data

- Quantlets: github.com/danpele/MFM/tree/main/Quantlets/Ch_16

Back

Acronyms used in this material (1/2)

- **AI**: Artificial Intelligence
- **AIC**: Akaike Information Criterion
- **AR**: AutoRegressive model
- **BET**: Bucharest Exchange Trading (index)
- **BNB**: BNB (native token of BNB Chain, formerly Binance Coin)
- **BSC**: BNB Smart Chain
- **CAPM**: Capital Asset Pricing Model
- **CI**: Confidence Interval
- **CMKT**: Crypto Market factor (value-weighted market excess return)
- **CMOM**: Crypto Momentum factor (past winners minus losers)
- **COVID-19**: COronaVIrus Disease 2019
- **CRIX**: CRypto IndeX
- **CSMB**: Crypto Size factor (small minus big coins)
- **DAI**: Dai (stablecoin of the Maker/Sky protocol)
- **DTB4WK**: 4-Week Treasury Bill Secondary Market Rate (FRED series code)
- **ECRIX**: Exact CRIX (number of constituents chosen in steps of one)
- **EFCRIX**: Exact Full CRIX (global minimum of the AIC)
- **ES**: Expected Shortfall
- **ETF**: Exchange-Traded Fund
- **ETHA**: iShares Ethereum Trust ETF (ticker)
- **EU**: European Union
- **FAQ**: Frequently Asked Questions
- **FDIC**: Federal Deposit Insurance Corporation
- **FM**: Fama–MacBeth (regressions)
- **FRED**: Federal Reserve Economic Data
- **FWER**: Family-Wise Error Rate
- **FX**: Foreign eXchange

Acronyms used in this material (2/2)

- **GENIUS**: Guiding and Establishing National Innovation for U.S. Stablecoins (Act)
- **GMM**: Generalized Method of Moments
- **GRS**: Gibbons–Ross–Shanken (test)
- **HAC**: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- **IA**: Internet Appendix (of the cited paper)
- **LP**: Liquidity Provider
- **LUNA**: Luna (native token of the Terra blockchain)
- **LVR**: Loss-Versus-Rebalancing
- **MAR33**: Basel Framework, chapter MAR33 (internal models approach: capital requirements calculation)
- **MDD**: Maximum DrawDown
- **NAV**: Net Asset Value
- **OLS**: Ordinary Least Squares
- **PC1**: first Principal Component
- **PC2**: second Principal Component
- **PCA**: Principal Component Analysis
- **SE**: Standard Error
- **SSR**: Sum of Squared Residuals
- **SVB**: Silicon Valley Bank
- **TAR**: Threshold AutoRegression
- **TE**: Tracking Error
- **TMI**: Total Market Index (all coins, CRIX methodology)
- **TVL**: Total Value Locked
- **US**: United States
- **USDC**: USD Coin (stablecoin issued by Circle)
- **UST**: TerraUSD (algorithmic stablecoin of the Terra blockchain)
- **XAU**: gold, ISO 4217 code (one troy ounce)
- **XRP**: XRP (the native token of the XRP Ledger)

References (1/5)

- Adams, H., Zinsmeister, N., Robinson, D. (2020). Uniswap v2 Core. Technical white paper.
- Agakishiev, I., Härdle, W.K., Pele, D.T. (2025). Cryptocurrency liquidity forecasting. In *Handbook of Blockchain Analytics*, Springer, 41–60.
- Andrews, D.W.K. (1993). Tests for parameter instability and structural change with unknown change point. *Econometrica*, 61(4), 821–856.
- Angeris, G., Kao, H.-T., Chiang, R., Noyes, C., Chitra, T. (2019). An analysis of Uniswap markets. arXiv:1911.03380.
- Aramonte, S., Huang, W., Schrimpf, A. (2021). DeFi risks and the decentralisation illusion. *BIS Quarterly Review*, December 2021.
- Athreya, K.B. (1987). Bootstrap of the mean in the infinite variance case. *Annals of Statistics*, 15(2), 724–731.
- Babiak, M., Bianchi, D. (2025). Mispricing and risk compensation in cryptocurrency returns. *Journal of Financial and Quantitative Analysis*, 1–27.
- Bai, J. (1997). Estimation of a change point in multiple regression models. *Review of Economics and Statistics*, 79(4), 551–563.
- Bai, J., Perron, P. (1998). Estimating and testing linear models with multiple structural changes. *Econometrica*, 66(1), 47–78.
- Balke, N.S., Fomby, T.B. (1997). Threshold cointegration. *International Economic Review*, 38(3), 627–645.
- Basel Committee on Banking Supervision (2023). MAR33: Internal models approach, capital requirements calculation. Basel Framework, Bank for International Settlements.
- Baur, D.G., Lucey, B.M. (2010). Is gold a hedge or a safe haven? An analysis of stocks, bonds and gold. *Financial Review*, 45(2), 217–229.

References (2/5)

- Będowska-Sójka, B., Wójcik, P., Pele, D.T. (2026). Early warning systems for cryptocurrency markets: predicting 'zombie' assets using machine learning. *North American Journal of Economics and Finance*, 81, 102543.
- Bitcoin.org (2026). Frequently asked questions. Web page.
- BlackRock (2026). iShares Bitcoin Trust ETF (IBIT). Product page.
- Bouri, E., Molnár, P., Azzi, G., Roubaud, D., Hagfors, L.I. (2017). On the hedge and safe haven properties of Bitcoin: is it really more than a diversifier?. *Finance Research Letters*, 20, 192–198.
- Buterin, V. (2014). Ethereum: a next-generation smart contract and decentralized application platform. White paper.
- Cong, L.W., Karolyi, G.A., Tang, K., Zhao, W. (2026). Crypto value, factor pricing, and market segmentation. *Management Science*, forthcoming.
- Cong, L.W., Li, X., Tang, K., Yang, Y. (2023). Crypto wash trading. *Management Science*, 69(11), 6427–6454.
- Cong, L.W., Li, Y., Wang, N. (2021). Tokenomics: dynamic adoption and valuation. *Review of Financial Studies*, 34(3), 1105–1155.
- Dimson, E. (1979). Risk measurement when shares are subject to infrequent trading. *Journal of Financial Economics*, 7(2), 197–226.
- Diop, P.O., Chevallier, J., Sanhaji, B. (2024). Collapse of Silicon Valley Bank and USDC depegging: a machine learning experiment. *FinTech*, 3(4), 569–590.
- Drees, H. (2000). Weighted approximations of tail processes for β-mixing random variables. *Annals of Applied Probability*, 10(4), 1274–1301.
- Drees, H., de Haan, L., Resnick, S. (2000). How to make a Hill plot. *Annals of Statistics*, 28(1), 254–274.

References (3/5)

- Ethereum.org (2026). The Merge. Web page.
- Fama, E.F., French, K.R. (2015). A five-factor asset pricing model. *Journal of Financial Economics*, 116(1), 1–22.
- FDIC (2023). FDIC creates a Deposit Insurance National Bank of Santa Clara to protect insured depositors of Silicon Valley Bank. Press release, 10 March 2023.
- Forbes, K.J., Rigobon, R. (2002). No contagion, only interdependence: measuring stock market comovements. *Journal of Finance*, 57(5), 2223–2261.
- Franke, J., Härdle, W.K., Hafner, C.M. (2019). Statistics of Financial Markets: An Introduction, 5th ed. Springer.
- Găman, Ş., Pele, D.T., Härdle, W.K., Będowska-Sójka, B. (2026). White paper on cryptocurrency indices. Bucharest University of Economic Studies.
- GENIUS Act (2025). Guiding and Establishing National Innovation for U.S. Stablecoins Act. Public Law 119-27.
- Gibbons, M.R., Ross, S.A., Shanken, J. (1989). A test of the efficiency of a given portfolio. *Econometrica*, 57(5), 1121–1152.
- Griffin, J.M., Shams, A. (2020). Is Bitcoin really untethered?. *Journal of Finance*, 75(4), 1913–1964.
- Hansen, B.E. (1996). Inference when a nuisance parameter is not identified under the null hypothesis. *Econometrica*, 64(2), 413–430.
- Härdle, W.K., Harvey, C.R., Reule, R.C.G. (2020). Understanding cryptocurrencies. *Journal of Financial Econometrics*, 18(2), 181–208.
- Harvey, C.R., Liu, Y., Zhu, H. (2016). . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1), 5–68.

References (4/5)

- Hasbrouck, J. (1995). One security, many markets: determining the contributions to price discovery. *Journal of Finance*, 50(4), 1175–1199.
- Hill, B.M. (1975). A simple general approach to inference about the tail of a distribution. *Annals of Statistics*, 3(5), 1163–1174.
- Kang, H., Lee, S.-G. (2026). Institutionalization without integration: Bitcoin after the spot ETF. *International Review of Economics & Finance*, 110, 105531.
- Kleibergen, F. (2009). Tests of risk premia in linear factor models. *Journal of Econometrics*, 149(2), 149–173.
- Lee, S., Lee, J. (2025). Shifting dynamics: Bitcoin spot ETF approval and Bitcoin's relationships with financial markets. *Applied Economics Letters*, 1–7.
- Lehmann, E.L., Romano, J.P. (2005). Generalizations of the familywise error rate. *Annals of Statistics*, 33(3), 1138–1154.
- Lin, M.-B., Pele, D.T., Ren, R. (2025). Understanding blockchain technology. In *Handbook of Blockchain Analytics*, Springer, 3–18.
- Liu, J., Makarov, I., Schoar, A. (2023). Anatomy of a run: the Terra Luna crash. NBER Working Paper 31160.
- Liu, Y., Tsyvinski, A. (2021). Risks and returns of cryptocurrency. *Review of Financial Studies*, 34(6), 2689–2727.
- Liu, Y., Tsyvinski, A., Wu, X. (2022). Common risk factors in cryptocurrency. *Journal of Finance*, 77(2), 1133–1177.
- Lyons, R.K., Viswanath-Natraj, G. (2023). What keeps stablecoins stable?. *Journal of International Money and Finance*, 131, 102777.
- Ma, Y., Zeng, Y., Zhang, A.L. (2025). Stablecoin runs and the centralization of arbitrage. NBER Working Paper 33882.

References (5/5)

- Makarov, I., Schoar, A. (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics*, 135(2), 293–319.
- MiCA (2023). Regulation (EU) 2023/1114 on markets in crypto-assets. European Securities and Markets Authority.
- Milionis, J., Moallemi, C.C., Roughgarden, T. (2025). Automated market making and arbitrage profits in the presence of fees. In *Financial Cryptography and Data Security 2024*, Lecture Notes in Computer Science, Springer, 159–171.
- Milionis, J., Moallemi, C.C., Roughgarden, T., Zhang, A.L. (2022). Automated market making and loss-versus-rebalancing. arXiv:2208.06046.
- Nakamoto, S. (2008). Bitcoin: a peer-to-peer electronic cash system. White paper.
- Pele, D.T., Mazurencu-Marinescu-Pele, M. (2019a). Metcalfe's law and log-period power laws in the cryptocurrencies market. *Economics*, 13(1), 2019-29.
- Pele, D.T., Mazurencu-Marinescu-Pele, M. (2019b). Using high-frequency entropy to forecast Bitcoin's daily value at risk. *Entropy*, 21(2), 102.
- Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). Are cryptos becoming alternative assets?. *European Journal of Finance*, 29(10), 1064–1105.
- Schär, F. (2021). Decentralized finance: on blockchain- and smart contract-based financial markets. *Federal Reserve Bank of St. Louis Review*, 103(2), 153–174.
- Shanken, J. (1992). On the estimation of beta-pricing models. *Review of Financial Studies*, 5(1), 1–33.
- Treasury, Federal Reserve, FDIC (2023). Joint statement by Treasury, Federal Reserve, and FDIC. Press release, 12 March 2023.
- Trimborn, S., Härdle, W.K. (2018). CRIX an index for cryptocurrencies. *Journal of Empirical Finance*, 49, 107–122; author version arXiv:2009.09782.