



Modelling Financial Markets

Chapter 20 (special chapter): Path Signatures and Realised Volatility — a Research Case Study



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AAPL: Apple Inc. (Nasdaq ticker)
- ▣ ACM: Association for Computing Machinery
- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model)
- ▣ ARIMA: AutoRegressive Integrated Moving Average (model)
- ▣ AUD: Australian Dollar (ISO 4217 code)
- ▣ BH: Benjamini–Hochberg (procedure)
- ▣ BIC: Bayesian Information Criterion
- ▣ BPQ: Bollerslev–Patton–Quaedvlieg (insanity filter of the HARQ paper)
- ▣ CAD: Canadian Dollar (ISO 4217 code)
- ▣ CL: Crude oil (WTI light sweet) futures, NYMEX ticker
- ▣ COVID-19: COronaVirus Disease 2019
- ▣ DE: Deutschland (jurisdiction of the CC licence) (Germany (jurisdiction of the CC licence))
- ▣ DM: Diebold–Mariano (test)
- ▣ ES: E-mini S&P 500 futures (CME ticker)
- ▣ EURUSD: Euro / US dollar exchange rate (FX pair code)
- ▣ FDR: False Discovery Rate
- ▣ FWER: Family-Wise Error Rate
- ▣ FX: Foreign eXchange
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GC: Gold futures, COMEX ticker
- ▣ GW: Giacomini–White (test)
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HAR: Heterogeneous AutoRegressive (model)
- ▣ HARQ: HAR model with Realised Quarticity
- ▣ JPM: JPMorgan Chase & Co. (NYSE ticker)

Acronyms used in this material (2/2)

- ▣ KDD: Knowledge Discovery and Data Mining (ACM SIGKDD conference)
- ▣ LARS: Least-Angle RegreSsion (algorithm)
- ▣ LASSO: Least Absolute Shrinkage and Selection Operator
- ▣ MCS: Model Confidence Set
- ▣ MSE: Mean Squared Error
- ▣ NG: Natural Gas futures, NYMEX ticker
- ▣ OLS: Ordinary Least Squares
- ▣ PDE: Partial Differential Equation
- ▣ PRDS: Positive Regression Dependence on a Subset (of p-values)
- ▣ QLIKE: Quasi-LIKElihood (loss function)
- ▣ RQ: Realised Quarticity
- ▣ RV: Realised Variance
- ▣ SHAR: Semivariance HAR model
- ▣ TSLA: Tesla, Inc. (Nasdaq ticker)
- ▣ VOLARE: VOLatility Archive for Realized Estimates (database)
- ▣ XOM: Exxon Mobil Corporation (NYSE ticker)

Outline

Motivation

Realised Measures: What We Forecast

The VOLARE Database

Benchmarks: the HAR Family

Path Signatures

The Paper: Gu et al. (2024)

Our Transfer Design (Pre-Registered)

Results

Ablations and Robustness

Limitations and Discussion

AI for Scientific Discovery


Conclusions

References

What This Special Chapter Does

A research case study, start to finish

- Take a method published for a different problem: freight-rate forecasting at Amazon [Gu et al. \(2024\)](#)
- Transfer it to a core problem of this course: forecasting daily realised variance
- Fix the design before looking at the out-of-sample results, then report whatever comes out
- Data: 50 assets from the VOLARE database (VOLatility Archive for Realized Estimates) [Cipollini et al. \(2026\)](#)



Wall Street and the New York Stock Exchange, 25 March 2020

Photo: [Billie Grace Ward \(2020\)](#); CC BY 2.0; [Wikimedia Commons](#)

Why Forecast Realised Variance?

- Volatility is an input to almost every risk measure in this course
 - ▶ Tail-risk measures (Chapter 7), option prices (Chapter 12), portfolio weights (Chapter 4)
- Intraday data make daily variance almost observable: realised variance (RV, Chapter 9)
 - ▶ So volatility forecasting becomes a regression problem with a noisy but model-free target
- The benchmark is hard to beat: the HAR model [Corsi \(2009\)](#) (Chapter 9)
 - ▶ Three averages of past RV; linear; estimated by OLS in a second
- Any new method must beat HAR out of sample, on many assets, with a proper test

The Research Question

- Gu et al. describe a recent window of a series as a **path** and summarise it by its **signature**
 - ▶ Signature features enter a sparse linear regression (two-step LASSO)
 - ▶ Past periods whose path looks like today's get more weight (a signature kernel)
- They report a more than fivefold accuracy gain over an industry forecast of freight rates
 - ▶ Question: does the same recipe improve on HAR-type models for realised variance?
- Two separate ingredients to test: the signature features and the kernel weights

Realised Variance and Quadratic Variation

- Log price p_t with $dp_t = \mu_t dt + \sigma_t dW_t + dJ_t$; day t split into M intraday returns $r_{t,i}$
 - ▶ $RV_t = \sum_{i=1}^M r_{t,i}^2 \xrightarrow{M \rightarrow \infty} \int_{t-1}^t \sigma_s^2 ds + \sum_{t-1 < s \leq t} (\Delta J_s)^2$ Barndorff-Nielsen & Shephard (2002)
- Without jumps and noise, on an equally spaced grid, the limit is the integrated variance IV_t ; the error is mixed Normal
 - ▶ $\sqrt{M}(RV_t - IV_t) \Rightarrow MN(0, 2IQ_t)$, $IQ_t = \int_{t-1}^t \sigma_s^4 ds$ (integrated quarticity)
- Microstructure noise makes very fine sampling biased; five-minute returns are the standard compromise (Chapter 9)
- Chapter 9 derives these results; here we use them as given

The Realised Measures in VOLARE

- Semivariances [Barndorff-Nielsen, Kinnebrock & Shephard \(2010\)](#): $RV_t^+ = \sum_i r_{t,i}^2 1\{r_{t,i} > 0\}$,
 $RV_t^- = \sum_i r_{t,i}^2 1\{r_{t,i} < 0\}$
 - ▶ They add up: $RV_t = RV_t^+ + RV_t^-$ (holds in VOLARE to machine precision)
- Realised quarticity: $RQ_t = \frac{M}{3} \sum_i r_{t,i}^4 \rightarrow IQ_t$, the scale of the measurement error of RV
 - ▶ The limit needs a continuous, noise-free price and an equally spaced grid; a jump ΔJ adds about $M(\Delta J)^4/3$, so raw RQ is jump-sensitive
- Bipower variation (robust to jumps), realised kernel (robust to noise) [Barndorff-Nielsen et al. \(2008\)](#), range-based estimators
- Target of this chapter: five-minute RV (rv5)
 - ▶ Across 400 assets, little beats five-minute RV as a proxy of daily variance [Liu, Patton & Sheppard \(2015\)](#)
 - ▶ Its semivariances and quarticity are on the same grid, so SHAR and HARQ use consistent inputs

Targets and Losses

- Horizon h : $RV_{t+1:t+h} = \frac{1}{h} \sum_{j=1}^h RV_{t+j}$, $h \in \{1, 5, 22\}$; log models forecast $\ln RV_{t+1:t+h}$
 - ▶ Back to levels with $F_t = \exp(\hat{y}_t + \hat{\sigma}^2/2)$, $\hat{\sigma}^2$ = residual variance: exact if the log error is conditionally Normal with constant variance, otherwise an approximation
 - ▶ For $h > 1$ targets overlap: this induces serial dependence (through lag $h - 1$ by construction, possibly longer), so every test uses a HAC variance
- A loss is **robust** if, for a conditionally unbiased proxy ($\mathbb{E}_{t-1}[RV_t] = \sigma_t^2$), the *expected*-loss ranking with RV_t equals that with the true variance [Patton \(2011\)](#)
 - ▶ $MSE_t = (RV_t - F_t)^2$, $QLIKE_t = RV_t/F_t - \ln(RV_t/F_t) - 1 \geq 0$ (zero only if $F_t = RV_t$)
- QLIKE depends on the ratio RV_t/F_t : the same proportional error has the same loss on a calm day as on a stormy day
- MSE is dominated by a few crisis days; QLIKE is our primary loss, MSE the secondary one

Recap: Realised Measures

- ▣ RV estimates the daily quadratic variation; its error scales with the quarticity
- ▣ Semivariances split RV by the sign of the intraday return; RQ measures how noisy RV is
- ▣ Target: five-minute RV averaged over $h = 1, 5, 22$ days; primary loss QLIKE, secondary MSE

VOLARE: an Open Archive of Realised Measures

- Built at the Universities of Messina and Florence from tick data [Cipollini et al. \(2026\)](#)
 - ▶ Asset-specific cleaning, trading calendars and timestamp precision; the same estimators for every asset class
- Free for research at volare.unime.it; not redistributed with this course
 - ▶ Download the realised-variance files and point the notebooks to your local folder
- Citation requested by the authors: Cipollini, Cruciani, Gallo, Insana, Otranto & Spagnolo (2026), arXiv:2602.19732



University of Messina, the home of VOLARE


Photo: [Superchilum \(2013\)](#); CC BY-SA 3.0; [Wikimedia Commons](#)

The Panel: 50 Assets in Three Classes

Class	Assets	Sample	Observations	Median volatility
American large-cap stocks	40	2 Jan 2015 – 30 Jan 2026	111 440	17.1%
Futures (C, CL, ES, GC, NG)	5	28 Sep 2009 – 30 Jan 2026	21 059	19.7%
FX (AUD, EUR, GBP, CAD, JPY vs USD)	5	25 Sep 2009 – 30 Jan 2026	21 203	7.3%

- Stocks: 40 members of the Dow Jones 30 and S&P 100 (AAPL to XOM); futures: corn, crude oil, E-mini S&P 500, gold, natural gas
- Median volatility: $\sqrt{252 \text{ median}(\text{RV}_t)}$ per asset, then the median over the class; weekends are absent in all three files
- Variables used: `rv5`, `rsp5`, `rsn5`, `rq5`, and the daily open and close prices

Realised Volatility Across Asset Classes



□ Annualised volatility $\sqrt{252 \overline{RV}_t}$, five-day mean, log scale; two assets per class

Reading the Volatility Chart

- All classes share March 2020; crude oil (CL) peaks in April 2020, when the front contract traded below zero
 - ▶ A good forecaster must adapt within days, not months
- Typical levels differ substantially: FX around 7.3%, stocks around 17.1% a year
- On a log scale volatility looks stationary and persistent: the long memory that HAR approximates
- Spikes decay within weeks: the shape of the recent path, not only its level, carries information

Recap: The VOLARE Database

- ▣ 50 assets, three classes, 2009–2026 (stocks from 2015), standardised realised measures
- ▣ Licensed for research: code and results are published, the data are downloaded by each user
- ▣ Common stress episodes (March 2020, April 2025) are natural test periods

HAR: Three Averages of the Past

- HAR (Heterogeneous AutoRegressive) model [Corsi \(2009\)](#)

- ▶ $RV_{t+1:t+h} = \beta_0 + \beta_d RV_t + \beta_w RV_t^{(w)} + \beta_m RV_t^{(m)} + \varepsilon_{t+h}$
- ▶ $RV_t^{(w)} = \frac{1}{5} \sum_{j=0}^4 RV_{t-j}$, $RV_t^{(m)} = \frac{1}{22} \sum_{j=0}^{21} RV_{t-j}$

- HAR is a restricted AR(22): the 22 lag coefficients take only three distinct values

- ▶ $\phi_1 = \beta_d + \beta_w/5 + \beta_m/22$, $\phi_{2..5} = \beta_w/5 + \beta_m/22$, $\phi_{6..22} = \beta_m/22$

- Traders at daily, weekly and monthly frequencies: a cascade that mimics long memory with three parameters

- log-HAR: $\ln RV_{t+1:t+h}$ on $\ln RV_t$, $\ln RV_t^{(w)}$, $\ln RV_t^{(m)}$; no negative forecasts, less driven by spikes

- ▶ Logs of averages are not averages of logs, so this log-HAR is not an exact restricted AR(22) in $\ln RV$ (Seminar 20, A5)

Two Refinements: HARQ and SHAR

- RV is a generated regressor: $RV_t = IV_t + \eta_t$, $\text{Var}(\eta_t) \approx 2IQ_t/M$; errors-in-variables attenuate β_d on noisy days
 - ▶ HARQ [Bollerslev, Patton & Quaedvlieg \(2016\)](#):
 $RV_{t+1:t+h} = \beta_0 + (\beta_d + \beta_Q(RQ_t^{1/2} - \overline{RQ^{1/2}}))RV_t + \beta_w RV_t^{(w)} + \beta_m RV_t^{(m)} + \varepsilon_{t+h}$, expected $\beta_Q < 0$
 - ▶ BPQ filter: a forecast outside the in-sample range of the target is replaced by the in-sample mean (“insanity filter”)
- SHAR [Patton & Sheppard \(2015\)](#): split the daily term by the sign of the intraday returns
 - ▶ $RV_{t+1:t+h} = \beta_0 + \beta^+ RV_t^+ + \beta^- RV_t^- + \beta_w RV_t^{(w)} + \beta_m RV_t^{(m)} + \varepsilon_{t+h}$
 - ▶ Patton and Sheppard find $\beta^- > \beta^+$: an intraday version of the leverage effect (Chapter 5)
- The signature can capture a related effect from daily data: the Lévy area between returns and log RV (next section)

Seven Models in One Table

Model	Target	Weights	Regressors
HAR	RV	equal	$RV_t, RV_t^{(w)}, RV_t^{(m)}$
log-HAR	$\ln RV$	equal	the same, in logs
HARQ	RV	equal	$HAR + RV_t (RQ_t^{1/2} - \overline{RQ^{1/2}})$
SHAR	RV	equal	$RV_t^+, RV_t^-, RV_t^{(w)}, RV_t^{(m)}$
log-HAR-K	$\ln RV$	signature kernel	log-HAR regressors
Sig-L	$\ln RV$	equal	log-HAR + 39 signature terms, two-step LASSO
Sig-LK	$\ln RV$	signature kernel	as Sig-L: the method of Gu et al.


- The last four form a 2×2 ablation: signature features (yes/no) $\times$ kernel weights (yes/no)
- Log-HAR is the natural baseline for the signature models: same target, and its regressors are among the candidates; but the LASSO may also drop log-HAR terms, so Sig-L vs log-HAR compares two procedures, not only two feature sets

Recap: Benchmarks

- HAR is a three-parameter restricted AR(22); log-HAR uses the logs of the three averages
- HARQ corrects for the measurement error of RV; SHAR separates good and bad volatility
- The candidate set of the signature models contains log-HAR; a gain can come from path-shape information or from a different selection of the log-HAR terms

Where Signatures Come From

- **Chen (1957)**: iterated integrals of a path form an algebra and determine it up to reparametrisation
 - ▶ Kuo-Tsai Chen, geometry of loop spaces
- **Lyons (1998)**: rough path theory; the signature is the natural feature set for controlled differential equations
 - ▶ A response to a path depends on it (to first order) only through its signature
- **Machine learning**: [Levin, Lyons & Ni \(2013\)](#), [Chevyrev & Kormilitzin \(2016\)](#), [Lyons & McLeod \(2022\)](#); finance: signature-based asset price models [Cuchiero, Gazzani & Svaluto-Ferro \(2023\)](#)



Terry Lyons (left), Oberwolfach, 2024

Photo: [Katrin Schmid \(2024\)](#); CC BY-SA 2.0 DE; [Wikimedia Commons](#)

A Window of Data Is a Path

- Take the last L days of $d - 1$ series and connect the points linearly
 - ▶ $X : [0, 1] \rightarrow \mathbb{R}^d$, $X_{u_i} = (u_i, \ln \text{RV}_{t-L+1+i}, \sum_{j \leq t-L+1+i} r_j)$, $u_i = i/(L - 1)$
- **Time augmentation:** the first channel is time itself
 - ▶ Without it the signature cannot tell a fast move from a slow one (reparametrisation invariance)
- A regression on the window sees 22 lags; a path summary sees their order, speed and co-movement
 - ▶ Example: RV rose while prices fell, or RV rose first and prices fell later

The Signature: Definition

- For a path $X : [a, b] \rightarrow \mathbb{R}^d$ of bounded variation and a word $(i_1, \dots, i_k)$ with letters in $\{1, \dots, d\}$
 - ▶ $S(X)_{a,b}^{i_1 \dots i_k} = \int_{a < u_1 < \dots < u_k < b} dX_{u_1}^{i_1} \dots dX_{u_k}^{i_k}$ (iterated integral)
- The signature is the whole sequence of levels
 - ▶ $\text{Sig}(X) = (1, S^{(1)}, S^{(2)}, S^{(3)}, \dots)$, $S^{(k)} \in (\mathbb{R}^d)^{\otimes k}$ has d^k terms
- Level 1: $S^i = X_b^i - X_a^i$, the total increments
 - ▶ Level 2: $S^{ij} = \int_a^b (X_u^i - X_a^i) dX_u^j$; level 3 adds one more integral

Linear Segments and Chen's Identity

- A linear segment with increment $\Delta \in \mathbb{R}^d$ has the tensor exponential as signature
 - ▶ $\text{Sig}(\text{segment}) = \exp_{\otimes}(\Delta) = (1, \Delta, \frac{1}{2!}\Delta^{\otimes 2}, \frac{1}{3!}\Delta^{\otimes 3}, \dots)$
- **Chen's identity** [Chen \(1957\)](#): the signature of a concatenation is the tensor product
 - ▶ $\text{Sig}(X * Y) = \text{Sig}(X) \otimes \text{Sig}(Y), \quad (A \otimes B)^{(k)} = \sum_{p=0}^k A^{(k-p)} \otimes B^{(p)}$
- A piecewise-linear path with m segments: $\text{Sig} = \exp_{\otimes}(\Delta_1) \otimes \dots \otimes \exp_{\otimes}(\Delta_m)$
 - ▶ This is the whole algorithm: m updates, each costing $O(d^N)$ at depth N (a few lines of numpy)
 - ▶ Rolling windows: Chen also gives $\text{Sig}(X_{[s,t]}) = \text{Sig}(X_{[a,s]})^{-1} \otimes \text{Sig}(X_{[a,t]})$

Worked Example: a Two-Segment Path

Path


- ▣ $X : (0, 0) \rightarrow (1, 0) \rightarrow (1, 2); \Delta_1 = (1, 0), \Delta_2 = (0, 2)$
- ▣ Level 1: $S^{(1)} = \Delta_1 + \Delta_2 = (1, 2)$
- ▣ Level 2 by Chen: $S^{ij} = \frac{1}{2}\Delta_1^i\Delta_1^j + \Delta_1^i\Delta_2^j + \frac{1}{2}\Delta_2^i\Delta_2^j$

Result

- ▣ $S^{11} = \frac{1}{2}, S^{12} = 0 + 1 \cdot 2 + 0 = 2$
- ▣ $S^{21} = 0, S^{22} = \frac{1}{2} \cdot 2^2 = 2$
- ▣ Check: $S^{11} = (\Delta X^1)^2/2$ and $S^{22} = (\Delta X^2)^2/2$ (shuffle identity)
- ▣ $S^{12} \neq S^{21}$: the order matters (first right, then up)

The Lévy Area: Order as a Number


- Antisymmetric part of level 2
 - ▶ $A^{ij} = \frac{1}{2} (S^{ij} - S^{ji}) =$ signed area between the path and its chord
- Symmetric part: $\frac{1}{2} (S^{ij} + S^{ji}) = \frac{1}{2} \Delta X^i \Delta X^j$, no new information
 - ▶ Example above: $A^{12} = (2 - 0)/2 = 1$, the area of the triangle under the path
- Channels (return, log RV): a price fall followed by a volatility rise gives $A < 0$, but so does a rise followed by a volatility fall
 - ▶ A lead-lag statistic: leverage-type sequences (Chapter 5) are one contributor to a negative area; the sign alone does not identify them



Paul Lévy (1886–1971)

Photo: Konrad Jacobs; CC BY-SA 2.0 DE;
[Wikimedia Commons](#)

Two Windows of JPM as Paths



- Path (cumulative open-to-close return, $\log RV$) over 22 days; the dashed chord joins start and end; the shaded region gives the Lévy area

Reading the Two Paths

- Calm window (to 15 Nov 2019): log RV moves by -0.58, the path wanders, Lévy area -0.011
 - ▶ Rises and falls of RV are not linked to the direction of prices
- Stress window (to 16 Mar 2020): log RV rises by +4.32 (RV multiplied by about 75)
 - ▶ Intraday returns first fall by about 10% and RV climbs as they fall: Lévy area -0.247
 - ▶ The net intraday return over the window is only +0.5%: the crash happened largely overnight, outside this channel
- Two contrasting shapes: the higher levels record the within-window order of moves that level 1 summarises away
 - ▶ What the input omits (the overnight returns) no level can recover: the open-to-close channel never saw it


Why Signatures Work: Three Theorems

- **Uniqueness** [Hambly & Lyons \(2010\)](#): the full $\text{Sig}(X)$ determines X up to tree-like equivalence and translation; with a strictly increasing time channel, up to translation only
 - ▶ Injective with time augmentation and a fixed (or separately supplied) starting point; a finite truncation Sig^N is not injective
- **Universality (linearisation)**: for a compact set K (1-variation topology) of time-augmented paths with a common starting point and $f : K \rightarrow \mathbb{R}$ continuous, for every $\varepsilon > 0$ there is a linear functional ℓ with
 - ▶ $\sup_{X \in K} |f(X) - \langle \ell, \text{Sig}(X) \rangle| < \varepsilon$ (Stone–Weierstrass, via the shuffle identity)
- **Factorial decay**: $\|S^{(k)}\| \leq \|X\|_{1\text{-var}}^k / k!$
 - ▶ Truncating at a small depth N loses little for paths of moderate length

The Shuffle Identity: Nonlinearity for Free

- Products of signature terms are linear combinations of higher terms
 - ▶ $S^i S^j = S^{ij} + S^{ji}$, $S^{ij} S^k = S^{ijk} + S^{ikj} + S^{kij}$
- So a linear model in $\text{Sig}^N(X)$ already contains polynomials of the increments
 - ▶ This is why a LASSO on signatures can approximate nonlinear functions of the path
- The price: the number of features grows as d^N
 - ▶ Sparsity (LASSO) is needed, as in Gu et al.

How Many Features? The Truncation Dimension



□ Levels 1..N of a d -dimensional path: $M(d, N) = \sum_{k=1}^N d^k = \frac{d^{N+1} - d}{d - 1}$ terms

Choosing d and N

- Our path: $d = 3$ channels (time, log RV, cumulative return), depth $N = 3$: $3 + 9 + 27 = 39$ terms
 - ▶ Three of them are constant (words with only the time letter), so 36 vary
- With $d = 5$ and $N = 4$ we would already have 780 terms for 1000 observations: the LASSO would select from noise
- Gu et al. also use the path of the target alone (with time), plus external factors outside the signature
 - ▶ The paper does not disclose its N and L ; we fix $N = 3$, $L = 22$ ex ante and check $N = 2$ as robustness

Two Invariances to Keep in Mind

- ▣ **Translation invariance:** $\text{Sig}(X + c) = \text{Sig}(X)$ — only increments enter
 - ▶ The signature does not see the **level** of volatility, only the shape of its recent path
 - ▶ So the level must enter separately: the log-HAR terms play the role of the external factors x_T of Gu et al.
- ▣ **Reparametrisation invariance:** running the path faster or slower leaves Sig unchanged
 - ▶ Removed by the time channel; basepoint augmentation (prepending a segment from 0 to the original starting point) would remove the first invariance [Morrill et al. \(2020\)](#)
 - ▶ Translating the path to start at 0 does not: it is itself a translation (Seminar 20, C1 uses the prepended segment)

The Signature Kernel

- Kernel = inner product of feature maps [Király & Oberhauser \(2019\)](#)
 - ▶ $k_N(X, Y) = \langle \text{Sig}^N(X), \text{Sig}^N(Y) \rangle = 1 + \sum_{k=1}^N \sum_{|w|=k} S(X)^w S(Y)^w$
- Distance induced by the kernel (Gu et al., Eq. 9)
 - ▶ $d_{\text{Sig}}(X, Y) = k(X, X) - 2k(X, Y) + k(Y, Y) = \|\text{Sig}^N(X) - \text{Sig}^N(Y)\|^2$
- Gu et al. and this chapter use this **truncated** kernel, with N as for the features
 - ▶ The untruncated kernel solves a Goursat PDE (partial differential equation) [Salvi et al. \(2021\)](#); the paper cites it but does not use it
- Channels are rescaled before the kernel (next section): otherwise the channel with the largest units dominates

Recap: Path Signatures

- Signature = all iterated integrals of a path; computed exactly for piecewise-linear paths with Chen's identity
- Level 2 holds the Lévy areas: who moved first; the shuffle identity turns products into linear terms
- Universality: on compact sets of time-augmented paths with a common start, linear models on signatures approximate any continuous function of the path
- Translation invariance: the level of volatility must be supplied by other regressors

The Problem at Amazon

- Weekly North American truckload marketplace rates, national and five regions, 2018–2022 [Gu et al. \(2024\)](#)
 - ▶ Forecast up to 12 weeks ahead, for procurement and pricing decisions
- More than 100 candidate factors (prices, housing, drilling, logistics indices, employment, weather), screened to about 40
 - ▶ Deployed in Amazon trucking operations since November 2022
- Published at KDD 2024 (ACM Conference on Knowledge Discovery and Data Mining) [Gu et al. \(2024\)](#)



A container truck on an American interstate, 2025

Photo: [DanTD \(2025\)](#); CC BY 4.0; [Wikimedia Commons](#)

The Method, Step 1: Features

- At week τ : regression features = external factors x_τ and the signature of the time-augmented target path
 - ▶ $\xi_\tau = [x_\tau, \text{Sig}^N(\hat{y}_{\tau-l:\tau})]$, $\hat{y}_{\tau-l:\tau} = \{(t_i, y_{t_i})\}$ (Section 4.1, Eq. 14)
- Hyperparameters: lookback l , depth N , horizon Δt , temperature γ
 - ▶ Their values are not published (confidentiality)
- One direct regression per horizon: $y_{\tau+\Delta t}$ on ξ_τ (no iteration of one-step forecasts)

The Method, Step 2: Two-Step LASSO

- ▣ Selection by a weighted LASSO (Least Absolute Shrinkage and Selection Operator) [Tibshirani \(1996\)](#) (Eq. 14)
 - ▶ $\hat{\theta}_L = \arg \min_{\theta} \sum_{\tau} w_{\tau} (y_{\tau+\Delta t} - \xi_{\tau}^{\top} \theta)^2 + \lambda \|\theta\|_1$
- ▣ Refit by weighted OLS on the selected support (Eq. 15)
 - ▶ $\hat{\theta} = \arg \min_{\text{supp}(\theta)=\text{supp}(\hat{\theta}_L)} \sum_{\tau} w_{\tau} (y_{\tau+\Delta t} - \xi_{\tau}^{\top} \theta)^2$
- ▣ Why refit: LASSO shrinks the kept coefficients towards zero; OLS after selection removes that bias [Belloni & Chernozhukov \(2013\)](#)
 - ▶ The paper calls it “adaptive” because of the weights w_{τ} , not in the sense of the adaptive LASSO of [Zou \(2006\)](#)
- ▣ How λ is chosen is not reported

The Method, Step 3: Kernel Weights

- ▣ Kernel path: the *pairs* $z_\tau = (x_\tau, y_{\tau+\Delta t})$ of predictors and later target, not the features ξ_τ (Algorithm 1)
 - ▶ Distance to the last window with an observed target (Eq. 11): $\delta_\tau = d_{\text{Sig}, N}(z_{\tau-l:\tau}, z_{t-\Delta t-l:t-\Delta t})$
- ▣ Softmax weights with temperature $\gamma \geq 0$ (Eq. 12)
 - ▶ $w_\tau = \frac{\exp(-\gamma \delta_\tau)}{\sum_s \exp(-\gamma \delta_s)}$
 - ▶ $\gamma = 0$: ordinary (equal-weight) regression; $\gamma \rightarrow \infty$: only the most similar past window
- ▣ A nearest-neighbour idea inside a regression: learn from the past periods that looked like today
 - ▶ Cost of the idea: fewer effective observations, $n_{\text{eff}} = 1 / \sum_\tau w_\tau^2$ (Kish)

What the Paper Reports

- ▣ National rates, three months ahead, April–November 2021 (Table 1): relative error of about 1–2% vs 13–27% for the industry forecast
 - ▶ Relative error = $|y - \hat{y}|/|y|$; “the prediction accuracy is improved by more than fivefold”
- ▣ Horizons 1–12 weeks (Table 2): 1.06–2.60% nationally, 1.05–5.69% by region
 - ▶ At 12 weeks alone: 2.38% nationally, 2.74–5.65% by region
- ▣ Kernel ablation (Table 3, six weekly values, 31 October–5 December 2021, one region): errors up to 9.85% without the kernel vs up to 5.15% with it
 - ▶ An estimated annual saving of about 50 million dollars for Amazon

Reading the Paper Critically

- ▣ **One benchmark:** an unnamed commercial forecast; no ARIMA, no LASSO without signatures, no random walk
 - ▶ A fivefold gain over a weak benchmark says little about the method itself
- ▣ **Short evaluation:** a few months of 2021–2022, a few series; no standard errors, no test of equal accuracy
- ▣ **Undisclosed tuning:** I , N , γ , λ and how they were chosen; the kernel ablation covers six weeks
- ▣ None of this makes the method wrong; it makes the evidence thin, which is why we test it elsewhere
 - ▶ Replication before extension is the rule of the course project (Chapter 19)

What Transfers to Volatility, and What Does Not

Transfers

- ▣ Signature of the time-augmented recent path of the target
- ▣ Two-step LASSO with weights, direct regression per horizon
- ▣ Softmax weights $\exp(-\gamma \delta_\tau)$ with the truncated signature kernel

Must be adapted

- ▣ Daily, noisy target instead of weekly rates: logs, HAC inference, robust losses
- ▣ External factors x_τ replaced by the log-HAR terms (they carry the level)
- ▣ Undisclosed choices (l, N, γ, λ) fixed by us, ex ante
- ▣ Kernel: our path (time, ln RV, return) ending at τ vs the window ending at t , updated every 5 days, instead of the pairs $(x_\tau, y_{\tau+\Delta t})$ ending at $t - \Delta t$: an adaptation
- ▣ Strong benchmarks and formal tests instead of one commercial forecast

Recap: The Paper

- ▣ Features: external factors plus the signature of the target path; estimation: weighted LASSO, then OLS on the support
- ▣ Weights: softmax of minus the signature distance to the last window with an observed target; Sig-LK adapts the kernel path
- ▣ Reported gains are large but measured against one benchmark, over a short period, without inference

Hypotheses

- ▣ **H1 (features):** signature terms add information beyond log-HAR
 - ▶ Test: Sig-L vs log-HAR, same target, same equal weights
- ▣ **H2 (weights):** kernel weights improve a given regression
 - ▶ Tests: Sig-LK vs Sig-L and log-HAR-K vs log-HAR
- ▣ **H3 (method):** the full method of Gu et al. beats the HAR family
 - ▶ Tests: Sig-LK vs HAR and vs log-HAR; model confidence set of all seven models
- ▣ **H4 (stress):** any gain is larger in stress periods (COVID-19 crash, April 2025), where similar past windows should matter most

Features and Scaling

- ▣ Regressors at day τ for the signature models: 3 log-HAR terms + 39 signature terms
 - ▶ $\xi_\tau = [\ln \text{RV}_\tau, \ln \text{RV}_\tau^{(w)}, \ln \text{RV}_\tau^{(m)}, \text{Sig}^3(X_{\tau-21:\tau})]$, path $X = (\text{time}, \ln \text{RV}, \text{cumulative } r)$, 22 points
- ▣ Non-time channel c divided by $s_c \sqrt{L-1}$, s_c = standard deviation of its daily increments in the training window up to the origin t
 - ▶ Time is not rescaled (constant increments, $s = 0$): it runs from 0 to 1; a word $(i_1..i_k)$ is rescaled by $\prod_j c_{i_j}$, so signatures are computed once
- ▣ Before the LASSO every column is standardised with weighted moments of the training window; constant columns (words of time only) are dropped
- ▣ Raw features at τ use data up to τ ; scales and standardisation use data up to the origin t : the no-look-ahead cutoff is t


Estimation: Choosing λ Without Look-Ahead

- ▣ Inside each training window: the whole weighted LASSO path by LARS (least-angle regression)
 - ▶ For every support on the path: weighted OLS refit (step 2 of Gu et al.)
- ▣ Pick the support that minimises a BIC-type criterion (Bayesian information criterion) with the Kish effective sample size, our heuristic
 - ▶ $\text{BIC}(S) = n_{\text{eff}} \ln \hat{s}_S^2 + (|S| + 1) \ln n_{\text{eff}}, \quad n_{\text{eff}} = 1 / \sum_{\tau} w_{\tau}^2$
- ▣ Why not cross-validation: 50 assets $\times$ 3 horizons $\times$ hundreds of re-estimations; BIC uses only the window
 - ▶ Kish size corrects for unequal weights, not for serial dependence (overlapping targets, $h > 1$): not a likelihood-based BIC for these data

Kernel Temperature: One Choice per Class, on a Validation Block

- Scale-free temperature: $\gamma = c / \text{median}_\tau(\delta_\tau)$ in every window
 - ▶ $c = 1$: a window at the median distance gets weight e^{-1} relative to an identical one
- Grid $c \in \{0.5, 1, 2, 4\}$, chosen per horizon and asset class on the first 250 forecast days of each asset, pooled within the class
 - ▶ Criterion: mean over the class of $\text{QLIKE}(c)$ divided by the asset's average over the grid
 - ▶ Evaluation of every asset starts after the last validation target of its class (a common calendar cutoff), so no tuning target postdates an evaluated origin
- Chosen c for Sig-LK ($h = 1, 5, 22$): stocks 0.5, 0.5, 0.5; futures 2.0, 1.0, 0.5; FX 0.5, 2.0, 0.5
 - ▶ log-HAR-K: stocks 0.5, 0.5, 1.0; futures 2.0, 1.0, 4.0; FX 4.0, 2.0, 4.0
- For the 40 stocks the flattest value wins at every horizon for Sig-LK; for futures and FX the choice changes with the horizon: a flat, noisy criterion

Rolling Estimation, Step by Step



- Re-estimate every 5 trading days on the last 1000 complete pairs; between re-estimations the coefficients are fixed but the regressors are updated daily
- Kernel distances are recomputed at each re-estimation, against the window ending at the origin t

Evaluation Protocol

- Per asset and horizon: mean QLIKE and MSE; DM (Diebold–Mariano) test with Newey–West variance, bandwidth $\max(h, \lceil T^{1/3} \rceil)$ [Diebold & Mariano \(1995\)](#)
 - ▶ Rolling estimation, so the test is read as the unconditional GW (Giacomini–White) test; the conditional GW test uses instruments $(1, d_{t-h})$ [Giacomini & White \(2006\)](#)
- MCS (model confidence set) at 90%, $T_{\max}$ statistic, stationary bootstrap, 1000 replications [Hansen, Lunde & Nason \(2011\)](#)
 - ▶ Block length $\max(h, 10)$ days
- Panel: DM test on the cross-sectional mean loss differential (one series per date)
 - ▶ Counting wins over 50 assets: Holm (FWER) and BH (FDR) corrections [Benjamini & Hochberg \(1995\)](#)
- Subperiods: 20 Feb–30 Apr 2020, 2–30 Apr 2025, all other days; the paper's relative error as a descriptive statistic

A Disclosed Deviation

- Pre-registered: the BPQ insanity filter on all models
 - ▶ A pipeline test on JPM (one asset, before any panel result) showed it replacing correct crash forecasts by the calm in-sample mean
 - ▶ On 12 March 2020 the realised next-day RV of JPM was about 40 times the filtered HAR forecast
- Changed before the panel run: primary results with positivity only (a forecast ≤ 0 becomes the in-sample mean)
 - ▶ The BPQ filter is kept as a robustness check and reported in full
- Rule: a deviation is allowed, a hidden deviation is not; state what was seen when it was changed

Recap: Design


- ▣ Seven models, three horizons, 50 assets, rolling 1000-day window, re-estimation every 5 days
- ▣ All undisclosed choices of the paper fixed ex ante; one temperature chosen on a validation block
- ▣ Primary loss QLIKE; DM/GW per asset, panel DM, MCS, Holm and BH for win counts

What Was Evaluated

Class	Evaluation starts	Ends	Days per asset (median)
Stocks	27 Jan 2020	29 Jan 2026	1 499
Futures	25 Sep 2014	29 Jan 2026	2 919
FX	27 Aug 2014	29 Jan 2026	2 952

- 265 929 out-of-sample forecasts per model ($50 \text{ assets} \times 3 \text{ horizons}$)
- Stocks start in 2015, so their evaluation begins in January 2020: the COVID-19 crash is in every stock sample
 - ▶ Futures and FX start in 2009 and are evaluated from 2014
- Computation: 18.5 core-hours; 87 minutes on 14 parallel workers; signatures and LASSO in numpy and scikit-learn

QLIKE Relative to HAR, Asset by Asset



□ Boxes: interquartile range over the 50 assets; line: median; diamond: mean; whiskers: 1.5 interquartile ranges

Reading the Headline Results

- $h = 1$: log-HAR, log-HAR-K and Sig-L are close (0.883, 0.885, 0.886); Sig-LK is behind (0.911)
 - ▶ Relative to log-HAR, Sig-LK is +4.5% on average: no gain from 36 extra features
- $h = 5$ and $h = 22$: the signature models fall behind log-HAR (+28.0% and +92.9% for Sig-LK)
 - ▶ Medians (0.997, 0.968) are much better than means: a few assets and days drive the loss
- HARQ and SHAR lose to HAR in levels here; without the BPQ filter HARQ produces unstable forecasts (a robustness slide returns to this)
- First verdict: H3 is not supported in the sense that matters: the method of Gu et al. does not beat the simplest log model (a near-tie is not a proof of equal accuracy)

Asset by Asset: Wins and Losses, Corrected


Comparison (QLIKE)	$h = 1$		$h = 5$		$h = 22$	
Sig-LK vs log-HAR	2/0/0	5/0/0	0/0/0	15/0/2	0/0/0	27/1/16
Sig-L vs log-HAR	3/1/1	5/0/0	0/0/0	9/0/0	1/0/0	21/0/0
log-HAR-K vs log-HAR	3/0/0	3/0/0	3/0/0	1/0/0	5/0/0	6/0/0
Sig-LK vs Sig-L	2/0/0	5/0/0	0/0/0	5/0/0	2/0/0	10/0/0
Sig-LK vs HAR	12/5/6	0/0/0	7/1/1	0/0/0	8/2/4	1/0/0
log-HAR vs HAR	21/5/17	0/0/0	23/13/20	0/0/0	42/25/42	0/0/0

- Cell: significantly better on raw/Holm/BH assets | significantly worse on raw/Holm/BH assets; one-sided DM at 5%, family of 50 assets
- Holm controls the probability of any false win (FWER) under any dependence; BH controls the expected share of false wins (FDR)
 - The BH guarantee needs independent or positively dependent (PRDS) p -values; cross-asset shocks make this plausible, not certain [Benjamini & Yekutieli \(2001\)](#)

Reading the Win Counts

- Signatures (H1): Sig-L beats log-HAR on 3 of 50 assets at $h = 1$ without correction, 1 after Holm
 - ▶ Under the null of no gain we expect about 2.5 raw wins by chance: this is what we see
- Weights (H2): log-HAR-K vs log-HAR, 3/3/5 raw wins, none after Holm
 - ▶ Sig-LK vs Sig-L: no Holm-significant difference at any horizon
- Against HAR at $h = 1$, Sig-LK wins on 12 assets (5 after Holm); log-HAR alone wins on 21 (5 after Holm, 17 after BH)
 - ▶ Lesson: choose the benchmark that isolates the idea; “beats HAR” would have credited signatures with the gain of the log transform

DM Statistics for Every Asset and Horizon




- Diebold–Mariano t -statistics of the QLIKE differential; blue: Sig-LK more accurate, red: less accurate

Reading the DM Map

- Left panel (vs log-HAR): mostly pale at $h = 1$, red at $h = 5$ and $h = 22$
 - ▶ The blue exceptions at $h = 1$: ES and TSLA ($|t| > 1.645$) and EURUSD; five assets are significantly red, before any correction
- Right panel (vs HAR): blue at $h = 1$ for most assets, because Sig-LK inherits the log target
 - ▶ At $h = 22$ the colour depends on the asset: the gain of the log target is offset by unstable forecasts
- One map shows why a single-asset case study can mislead: pick ES at $h = 1$ and signatures look like a clear success

Model Confidence Sets



□ Share of the 50 assets in which each model belongs to the 90% MCS under QLIKE

Reading the Model Confidence Sets


- log-HAR is in the MCS for 100%, 98% and 98% of the assets; Sig-LK for 100%, 94%, 88%
 - ▶ The MCS rarely separates the four log models: the data cannot distinguish them asset by asset
- HAR, HARQ and SHAR are excluded for roughly a fifth to a third of the assets: the level target is the clearest loser
- Under MSE almost every model stays in the set: MSE is dominated by a few crisis days and has little power [Patton \(2011\)](#)
- Low power is itself a result: small differences in average loss need very long samples [Pele & Mazurencu-Marinescu-Pele \(2026\)](#)

Panel Tests: Pooling the Evidence

Comparison: DM t (one-sided p)	$h = 1$	$h = 5$	$h = 22$
Sig-LK vs log-HAR	+1.45 (0.927)	+1.68 (0.953)	+1.80 (0.964)
Sig-L vs log-HAR	-1.91 (0.028)	+1.48 (0.931)	+1.51 (0.935)
log-HAR-K vs log-HAR	+1.89 (0.971)	+0.40 (0.654)	+1.92 (0.973)
Sig-LK vs Sig-L	+1.74 (0.959)	+0.98 (0.837)	-0.77 (0.221)
log-HAR vs HAR	-3.09 (<0.001)	-3.95 (<0.001)	-4.50 (<0.001)

- ▣ Series: the mean QLIKE differential over all assets trading that day; negative t = first model better
- ▣ Pooled, only the log transform is clearly significant; Sig-L vs log-HAR at $h = 1$ is borderline ($p = 0.028$), driven by futures
 - For the 40 stocks alone the sign flips: $t = +1.72$ against Sig-L
- ▣ Conditional GW test, Sig-LK vs log-HAR: 0, 0 and 0 Holm rejections out of 50

Inside the COVID-19 Crash: JPM, Next Day




□ Annualised volatility $\sqrt{252F_t}$ of the forecasts and $\sqrt{252RV_{t+1}}$ of the realised target, February–May 2020

Reading the Crash Forecasts

- All models lag the first jump (late February): no model sees a shock before it happens
 - ▶ Sig-LK reacts slightly faster than log-HAR in mid-March, when the path has already turned
- HAR in levels overshoots and then turns negative in late March: the positivity rule replaces it by the in-sample mean (the flat segment)
 - ▶ A level model with large weekly and monthly terms reacts badly after a spike
- After the peak, log-HAR and Sig-LK decay at similar speeds; the extra signature terms change little at $h = 1$

Stress Periods vs Normal Times



□ Mean over assets of QLIKE relative to HAR within each subperiod (fixed before the evaluation)

Reading the Subperiods (H4)


- $h = 1$, COVID-19 crash: log-HAR 0.70, Sig-LK 0.73, HARQ 1.89 (relative to HAR)
 - ▶ April 2025: log-HAR 0.88, Sig-LK 0.92; all other days: 0.94 and 0.97
- $h = 22$, COVID-19 crash: Sig-LK 6.35 and Sig-L 8.30 times the loss of HAR
 - ▶ Exactly where similar past windows were supposed to help, the signature models fail most
- H4 not supported by these results (descriptive comparison, no formal interaction test): no stress-period advantage at $h = 1$, a large disadvantage at longer horizons

The Paper's Metric on Our Data

Median relative error $ RV - F /RV$ (%)	$h = 1$	$h = 5$	$h = 22$
HAR	57.2	48.9	51.7
log-HAR	50.4	41.9	38.7
HARQ	55.0	49.0	50.7
SHAR	57.9	49.2	51.8
log-HAR-K	49.7	41.0	38.7
Sig-L	50.1	42.6	42.6
Sig-LK	50.6	43.6	43.7

- Median over assets of the mean relative error: about 40–60%, against 1–2% for freight rates in Gu et al.
- Daily RV is a noisy, heavy-tailed target: the evaluated models have much larger relative errors than those reported for freight rates, without establishing a lower bound on attainable error
 - ▶ Sig-LK forecasts above ten times the in-sample maximum: 8, 29, 36 ($h = 1, 5, 22$), on 10 assets at $h = 5$; log-HAR: none

Kernel Weights at the Origin 16 March 2020




- Weights of Sig-LK ($h = 1$) computed at the origin 16 March 2020 over the training days (24 Mar 2016 – 13 Mar 2020), times the window length: 1 = equal weighting

What the Kernel Picks in a Crash

- Weights stay close to equal: between 0.61 and 1.43 times $1/n$; effective sample size 975 of 1000
 - ▶ The temperature chosen on the stock validation block ($c = 0.5$) is the flattest in the grid
- The highest weights go to the last days before the origin, when volatility was already rising
 - ▶ The past stress window of the late-2018 sell-off (October 2018 – January 2019) gets *less* weight; overall, weights are almost uncorrelated with log volatility (+0.04)
- Why: the signature is translation-invariant; a spike-and-decay window has a different shape from a rising one, whatever the level

Does the Kernel Ever Favour Stress?




- Weight that Sig-LK ($h = 1$) puts on the 10% most volatile days of its training window, averaged by class, three-month mean

Reading the Weight Share

- For stocks the share averages 6.7%, between 4.4% and 9.9% (the maximum on 17 Mar 2020)
 - ▶ Always below the 10% of equal weighting: the kernel systematically *down-weights* stress days
- Futures and FX show the same pattern, with peaks during the March 2020 turmoil
- This is the mechanism behind the null result for H2: the weights do not implement “learn from past crises”

What the LASSO Selects



- Share of re-estimations in which a regressor is in the BIC-selected support of Sig-LK (all assets and horizons); t = time, lrv = ln RV, ret = cumulative return


Reading the Selection Frequencies

- The three log-HAR terms are kept in 99%, 94% and 93% of the re-estimations
 - ▶ The LASSO rediscovers HAR: the daily level and the weekly and monthly averages carry most of the signal
- The most frequent signature term is (t, r, r) , in 27%; all frequent signature terms are of level 3
 - ▶ Average support size: 8.7 regressors for Sig-L, 6.7 for Sig-LK
- Level-3 terms are cubic in the increments: in a crash window they take values far outside the training range, and $\exp(\cdot)$ amplifies the error

Recap: Results

- H1 (signatures add information): not supported; no gain at one day, a loss at a week and a month
- H2 (kernel weights help): not supported; weights are almost flat and down-weight past stress
- H3 (the method beats the HAR family): only HAR in levels, and only through the log target
- H4 (gains in stress): not supported; the signature models fail most in the COVID-19 crash at long horizons

The 2×2 Ablation




- Mean over the 50 assets of QLIKE relative to log-HAR: features (signature or not) crossed with weights (kernel or equal)

Reading the Ablation

- Weights alone (log-HAR-K): 1.002, 1.002, 1.004: no effect at any horizon
 - ▶ Consistent with nearly flat weights (average effective sample size 816 of 1000)
- Features alone (Sig-L): 1.004, 1.285, 2.238
 - ▶ Adding the weights (Sig-LK: 1.045, 1.280, 1.929) damps the one-month failure but does not remove it, and worsens $h = 1$ through crude oil (CL: 2.71)
- Contrast with Gu et al. (Table 3): their kernel lowered the error in six weekly values of 2021 for one region; on 50 assets and six to twelve years it does not

Robustness of Sig-LK



- Sig-LK relative to HAR (mean over 50 assets) in the main design and in three pre-registered variants; the line marks log-HAR

Reading the Robustness Checks

- Depth 2: 0.922 and 0.820 at $h = 5, 22$ (main: 1.104, 1.334); path without returns: 0.878 and 0.770
 - ▶ The long-horizon failures come from the level-3 terms that involve the return channel
- Even the stable variants only match log-HAR (0.871, 0.765): they do not beat it
 - ▶ So the null result for H1 does not depend on depth or channels; the instability does
- A simpler specification is more robust, but here simpler means “closer to log-HAR”

The Insanity Filter, Revisited

QLIKE relative to HAR, BPQ filter on all models	$h = 1$	$h = 5$	$h = 22$
HAR	1.000	1.000	1.000
log-HAR	0.763	0.654	0.778
HARQ	0.995	0.912	1.014
SHAR	1.014	1.003	1.000
log-HAR-K	0.772	0.659	0.779
Sig-L	0.871	0.819	0.904
Sig-LK	0.890	0.867	0.930

- The filter rescues HARQ relative to HAR (0.995 at $h = 1$ instead of 1.140), as BPQ intended
- Sig-LK relative to log-HAR, both filtered: 1.238, 1.424, 1.209 (unfiltered 1.045, 1.280, 1.929): the gap widens at $h = 1, 5$ and narrows at $h = 22$
- Filtered forecasts, $h = 1$: HAR 180, HARQ 140, Sig-LK 143, log-HAR 29 (of 89 893): the filter hits the signature models far more often than log-HAR

By Asset Class

Sig-L relative to log-HAR (QLIKE)	$h = 1$	$h = 5$	$h = 22$
Stocks (40)	1.009	1.332	2.505
Futures (5)	0.965	1.084	1.206
FX (5)	0.996	1.111	1.134

- The only gain: futures at $h = 1$ (-3.5%) and, marginally, FX
 - Five futures and one horizon: a candidate for a new, pre-registered study, not a finding of this one
- Stocks, 80% of the panel, show no gain: the pooled result is mostly a stock result

Limitations of This Study

- One design out of many: $L = 22$, $N = 3$, three channels, BIC, a four-point grid for c
 - ▶ A null result for this design is not a null result for signatures in general
- Truncated kernel only; the PDE kernel [Salvi et al. \(2021\)](#), log-signatures or a basepoint (Seminar 20, Part C) were not run in the main design
- No external factors (implied volatility, macro news), which were central in Gu et al.
- Realized GARCH [Hansen, Huang & Shek \(2012\)](#) and machine-learning benchmarks (Chapter 13) are outside the comparison
- Stocks are evaluated over six years only, with the COVID-19 crash at the start of every sample

Open Directions

- ▣ A level-aware kernel: basepoint augmentation or adding the level of $\ln RV$ to the distance (reference analysis in Seminar 20, Part C)
- ▣ Signatures of the intraday path (5-minute returns within the day) instead of the daily path
- ▣ Log-signatures (no redundant symmetric terms) and bounded ($\tanh$) channels to avoid explosive level-3 terms [Morrill et al. \(2020\)](#)
- ▣ Cross-asset paths (for example, S&P 500 futures and a stock) where Lévy areas measure lead-lag across markets
- ▣ Futures at $h = 1$: replicate on a longer futures panel, with this design fixed in advance

An Open Question

When does a method transfer across domains?

- Methods travel faster than evidence: a KDD paper on freight rates becomes a candidate for volatility, credit, energy
 - ▶ Which properties of the source problem made the method work, and does the target share them?
- Finance adds a hard filter: 164 teams on the same data disagree more than their standard errors suggest [Menkveld et al. \(2024\)](#)
- Formal version: a transfer succeeds if the gain over the strongest simple benchmark survives a pre-registered set $\mathcal{S}$ of designs, assets and horizons, with the family corrected
 - ▶ In this chapter: $\mathcal{S} = 50 \text{ assets} \times 3 \text{ horizons} \times 4 \text{ variants}$; Holm/BH were applied per comparison and horizon (families of 50 assets), variants reported descriptively; no joint correction over all 600 cells; the gain did not survive

The Discovery Loop, Step by Step (1/2)

- **Read:** extract exactly what the paper did
 - ▶ “From Gu et al. (2024, KDD), list every equation of the estimator, every hyperparameter and its reported value, and the benchmark of each table. Mark what is not disclosed.”
 - ▶ Check each item against the PDF: equation numbers, table numbers, the kernel actually used
- **Map:** what transfers, what must be adapted
 - ▶ “Our target is daily realised variance. For each component of the method, say whether it transfers, what must change, and which benchmark isolates it.”
- **Pre-register:** hypotheses, grid, tests, correction, committed before the first result
 - ▶ The assistant drafts; the team decides every number and signs the commit

The Discovery Loop, Step by Step (2/2)

- **Build and test the pieces:** signatures against identities, not against another library
 - ▶ “Write unit tests for a depth-3 signature: level 1 equals the increment, the symmetric part of level 2 equals half the outer product, Chen's identity on a split path.”
- **Run the grid, then attack it**
 - ▶ “Here are all 450 runs (50 assets $\times$ 3 horizons $\times$ 3 specifications: main, depth 2, no return channel). Which conclusions survive Holm over 50 assets? Which single asset or subperiod would make a misleading headline?”
- **Explain the null:** find the mechanism, not an excuse
 - ▶ Here: flat weights (translation invariance) and explosive level-3 terms (exponential link), both checked with data

What the Researcher Must Check

- **The method is the paper's:** kernel, weights, refit, as in the cited equations
 - ▶ Test: every design slide cites an equation, table or section of the paper; undisclosed choices are labelled as ours
- **No look-ahead:** λ , scales, c and features use only data available at the origin
 - ▶ Test: change every value after the origin (for example, multiply future RV by 10); the forecast made at the origin must not change
- **The benchmark isolates the idea**
 - ▶ Test: the 2×2 ablation is reported, not only the full model against HAR
- **Losses and tests are the right way round**
 - ▶ Test: $QLIKE(RV, RV) = 0$ and $QLIKE(RV, F) > 0$ otherwise; the DM variance is HAC with bandwidth at least h

Mini-Case: Three Headlines From the Same Runs

- “Signatures beat HAR on 12 of 50 assets” (one-day horizon, raw DM)
 - ▶ True, but log-HAR alone beats HAR on 21: the gain belongs to the log transform
- “Signatures improve E-mini S&P 500 volatility forecasts” (ES, $h = 1$, DM $t = -2.91$ against log-HAR)
 - ▶ One cell out of 150 asset-horizon cells: after Holm, 0 of 50 at $h = 1$ and none at $h = 5, 22$
- The headline of the pre-registered evaluation: “no gain over log-HAR; worse at longer horizons; kernel weights inert”
 - ▶ An AI assistant asked to “find where the method works” will produce the first two; the plan protects the third

Your Turn

- Pick one open direction of this chapter (level-aware kernel, intraday signatures, log-signatures, futures only)
 - ▶ Write H_0/H_1 , the benchmark that isolates it, and the result that would make you drop it
- Use an assistant to list what could go wrong in your design, then write one unit test per item
- Commit the pre-registration before the first out-of-sample number; run the grid; report the distribution over $\mathcal{S}$
- Log every AI error you catch in `AI_ERRORS.md` (Seminar 20, C2 gives typical ones)

Key Takeaways

- ▣ Signatures turn a window into a feature vector with a theory behind it: Chen's identity, uniqueness, universality
- ▣ Transferred to realised variance, the method of Gu et al. does not beat log-HAR: no gain at one day, a loss at a week and a month
- ▣ Kernel weights stay almost flat and down-weight past stress: the translation-invariant kernel sees shape, not level
- ▣ The strongest effect in the study is the oldest one: model the log of RV
- ▣ A pre-registered null result, with ablations and a mechanism, is a finding

Further Study

Reading

- ▣ Signatures: Chevyrev & Kormilitzin (2016); Lyons & McLeod (2022); realised measures: Franke, Härdle & Hafner (2019) and Chapter 9

Advanced courses on quantinar.com

- ▣  Statistics of Financial Markets
- ▣  Applied Time Series Analysis with Python
- ▣  Machine Learning in Financial Risk

Code and data

- ▣ Quantlets: github.com/danpele/MFM/tree/main/Quantlets/Ch_20; data: VOLARE, free for research at volare.unime.it

References (1/2)

- ▣ Barndorff-Nielsen, O.E., Hansen, P.R., Lunde, A., Shephard, N. (2008). Designing realized kernels to measure the ex post variation of equity prices in the presence of noise. *Econometrica*, 76(6), 1481–1536.
- ▣ Barndorff-Nielsen, O.E., Kinnebrock, S., Shephard, N. (2010). Measuring downside risk: realised semivariance. In *Volatility and Time Series Econometrics*, Oxford University Press, 117–136.
- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2002). Econometric analysis of realized volatility and its use in estimating stochastic volatility models. *Journal of the Royal Statistical Society B*, 64(2), 253–280.
- ▣ Belloni, A., Chernozhukov, V. (2013). Least squares after model selection in high-dimensional sparse models. *Bernoulli*, 19(2), 521–547.
- ▣ Benjamini, Y., Hochberg, Y. (1995). Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society B*, 57(1), 289–300.
- ▣ Benjamini, Y., Yekutieli, D. (2001). The control of the false discovery rate in multiple testing under dependency. *The Annals of Statistics*, 29(4), 1165–1188.
- ▣ Bollerslev, T., Patton, A.J., Quaedvlieg, R. (2016). Exploiting the errors: a simple approach for improved volatility forecasting. *Journal of Econometrics*, 192(1), 1–18.
- ▣ Chen, K.-T. (1957). Integration of paths, geometric invariants and a generalized Baker–Hausdorff formula. *Annals of Mathematics*, 65(1), 163–178.
- ▣ Chevyrev, I., Kormilitzin, A. (2016). A primer on the signature method in machine learning. [arXiv:1603.03788](https://arxiv.org/abs/1603.03788).
- ▣ Cipollini, F., Cruciani, G., Gallo, G.M., Insana, A., Otranto, E., Spagnolo, F. (2026). VOLatility Archive for Realized Estimates (VOLARE). [arXiv:2602.19732](https://arxiv.org/abs/2602.19732).
- ▣ Corsi, F. (2009). A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Cuchiero, C., Gazzani, G., Svaluto-Ferro, S. (2023). Signature-based models: theory and calibration. *SIAM Journal on Financial Mathematics*, 14(3), 910–957.
- ▣ Diebold, F.X., Mariano, R.S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- ▣ Giacomini, R., White, H. (2006). Tests of conditional predictive ability. *Econometrica*, 74(6), 1545–1578.
- ▣ Gu, H., Guo, X., Jacobs, T.L., Kaminsky, P., Li, X. (2024). Transportation marketplace rate forecast using signature transform. *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 4997–5005; [arXiv:2401.04857](https://arxiv.org/abs/2401.04857).

References (2/2)

- Hambly, B., Lyons, T. (2010). Uniqueness for the signature of a path of bounded variation and the reduced path group. *Annals of Mathematics*, 171(1), 109–167.
- Hansen, P.R., Huang, Z., Shek, H.H. (2012). Realized GARCH: a joint model for returns and realized measures of volatility. *Journal of Applied Econometrics*, 27(6), 877–906.
- Hansen, P.R., Lunde, A., Nason, J.M. (2011). The model confidence set. *Econometrica*, 79(2), 453–497.
- Király, F.J., Oberhauser, H. (2019). Kernels for sequentially ordered data. *Journal of Machine Learning Research*, 20(31), 1–45.
- Levin, D., Lyons, T., Ni, H. (2013). Learning from the past, predicting the statistics for the future, learning an evolving system. arXiv:1309.0260.
- Liu, L.Y., Patton, A.J., Sheppard, K. (2015). Does anything beat 5-minute RV? A comparison of realized measures across multiple asset classes. *Journal of Econometrics*, 187(1), 293–311.
- Lyons, T. (1998). Differential equations driven by rough signals. *Revista Matemática Iberoamericana*, 14(2), 215–310.
- Lyons, T., McLeod, A.D. (2022). Signature methods in machine learning. arXiv:2206.14674.
- Menkveld, A.J., Dreber, A., Holzmeister, F., Huber, J., Johannesson, M., Kirchler, M., et al. (2024). Nonstandard errors. *Journal of Finance*, 79(3), 2339–2390.
- Morrill, J., Fermanian, A., Kidger, P., Lyons, T. (2020). A generalised signature method for multivariate time series feature extraction. arXiv:2006.00873.
- Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- Patton, A.J., Sheppard, K. (2015). Good volatility, bad volatility: signed jumps and the persistence of volatility. *Review of Economics and Statistics*, 97(3), 683–697.
- Pele, D.T., Mazurencu-Marinescu-Pele, M. (2026). Finite-sample precision limits for expected shortfall forecast comparisons. *Mathematics*, 14(13), 2316.
- Salvi, C., Cass, T., Foster, J., Lyons, T., Yang, W. (2021). The signature kernel is the solution of a Goursat PDE. *SIAM Journal on Mathematics of Data Science*, 3(3), 873–899.
- Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society B*, 58(1), 267–288.
- Zou, H. (2006). The adaptive lasso and its oracle properties. *Journal of the American Statistical Association*, 101(476), 1418–1429.