



Modelling Financial Markets

Seminar 2: Market Efficiency — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ **AI:** Artificial Intelligence
- ▣ **AMH:** Adaptive Markets Hypothesis
- ▣ **AR:** AutoRegressive (model); in event studies: Abnormal Return
- ▣ **BDS:** Brock–Dechert–Scheinkman (test)
- ▣ **BET:** Bucharest Exchange Trading (index)
- ▣ **BET-TR:** BET Total Return (index)
- ▣ **BIST:** Borsa İstanbul (BIST 100 index) — *the Istanbul Stock Exchange (BIST 100 index)*
- ▣ **BUX:** Budapesti Értéktőzsde index — *the Budapest Stock Exchange index*
- ▣ **BVB:** Bursa de Valori București — *the Bucharest Stock Exchange*
- ▣ **CD:** Chow–Denning (multiple variance-ratio test)
- ▣ **CI:** Confidence Interval
- ▣ **COVID-19:** COronaVirus Disease 2019
- ▣ **CV:** Cross-Validation
- ▣ **DFA:** Detrended Fluctuation Analysis
- ▣ **ETF:** Exchange-Traded Fund
- ▣ **FRED:** Federal Reserve Economic Data
- ▣ **FTSE:** Financial Times Stock Exchange (FTSE Russell, index provider)

Acronyms used in this material (2/2)

- ▣ **FWER**: Family-Wise Error Rate
- ▣ **GARCH**: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ **HAC**: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ **IPC**: Índice de Precios y Cotizaciones — *the Mexican Stock Exchange index*
- ▣ **NW**: Newey–West (standard errors)
- ▣ **OLS**: Ordinary Least Squares
- ▣ **PX**: Prague Stock Exchange index (index Burzy cenných papírů Praha)
- ▣ **RW1**: Random Walk 1 (i.i.d. increments)
- ▣ **RW2**: Random Walk 2 (independent increments)
- ▣ **RW3**: Random Walk 3 (uncorrelated increments)
- ▣ **SD**: Standard Deviation
- ▣ **SE**: Standard Error
- ▣ **TOM**: Turn Of the Month
- ▣ **US**: United States
- ▣ **VR**: Variance Ratio
- ▣ **WIG20**: Warszawski Indeks Giełdowy 20 — *the Warsaw Stock Exchange index of the 20 largest companies*

Today's Question and Seminar Structure

- Question of the seminar
 - ▶ **which weak-form rejections of Chapter 2 survive robust inference, several horizons and several tests?**
- Outputs to produce
 - ▶ the answers to the A exercises, on the worksheets
 - ▶ one table and one figure per B exercise
 - ▶ the audit of the AI answer in C1
 - ▶ one interpretation sentence per exercise
- Notebook for today: [open the seminar notebook in Colab](#)
 - ▶ each task names its notebook section
- Exercises, types and order: the Exercise Map (next slide)
- Seminar: Prof. Daniel Traian Pele, danpele@ase.ro; Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Model
A1	what does $\widehat{VR}(2)$ compute, line by line?	Solved	–
A2, A3	does an implied $H(20) \neq 0.5$ prove long memory?	Solved, Proposed	A2
B1	do the S&P 500 and BET reject RW3 with robust errors?	Solved	–
B2	is the S&P 500 reversal a crisis effect?	Proposed	B1
C1	what is wrong in an AI variance-ratio answer?	Proposed	A1, B1
B3	is the turn-of-the-month effect real?	Solved	–
B4	which of 15 calendar effects survive Holm?	Proposed	B3
After the lecture	B5 [Solved], B6 [Proposed]: rolling DFA and rolling variance ratios with bootstrap bands	homework	B1
Extensions	A4–A6, B7–B9, C2, C3 (optional, after the seminar)	Solved or Proposed (marked on each)	on each task

- ▣ **[Solved]**: full solution in the slides and in the notebook, a model to copy; **[Proposed]**: you solve it, the solution is discussed in class
- ▣ Part A: derivations on paper; Part B: computation and inference on real data, each ending with an interpretation question; Part C: open questions and AI critique

Data Used

Series	Price used	Calendar	Used in
S&P 500	close	US trading days	A3, B1–B5, C1
BET	close	BVB trading days	A2, B1, B3, B4, B7, C2
BET-TR and 8 emerging indices	close (WIG20, BUX, PX, BIST 100, Bovespa, IPC, Nifty 50, Shanghai)	own calendar each	B7, B4, C2
Euro Stoxx 50	close	own calendar	B4
Bitcoin	close (USD)	every day	B6
Banca Transilvania, BRD	adjusted close (dividends, splits)	common days with BET	B9
S&P 500 monthly	Shiller S&P Composite, S&P 500 closes, SPY dividends; 3-month T-bill (FRED, TB3MS)	month ends	B8
49 industry portfolios	Kenneth French Data Library , value-weighted, monthly	month ends	C3

- Daily closing prices of the course data, one series per asset, from 2000 (Bitcoin from 2011) to 18 Sep 2026; daily log returns (Notation, 1/3) on each series' own calendar
- Cleaning rule for indices: weekdays only; a close equal to the previous close is removed (usually a filled holiday)

Setup: Run and Verify

```
# notebook sections "Data" and
# "Efficiency tests" define
#   log_returns(name, start=None)
#   variance_ratio(r, q) -> (VR, z, z*)
#   chow_denning(r), runs_test(r), dfa_hurst(r)
sp = log_returns('sp500')    # S&P 500
bet = log_returns('bet')     # BET
print(len(sp), sp.index[0].date(),
      sp.index[-1].date(), len(bet))
print(sp.head(3).round(5))
```

- Expected: 6 714 S&P 500 returns from 2000-01-04 to 2026-09-18; 6 670 BET returns

date	close	r_t
2000-01-03	1455.22	—
2000-01-04	1399.42	−0.03910
2000-01-05	1402.11	+0.00192
2000-01-06	1403.45	+0.00096

- If your count or first value differs, check the series, the price type (adjusted or unadjusted close) and the cleaning rule

Notation Used in Every Exercise (1/3): Returns and Sample Moments

- **Daily log return** of a series with closing prices P_t

$$r_t = \ln P_t - \ln P_{t-1}, \quad t = 1, \dots, T$$

- ▶ T : the number of returns
- ▶ q : a horizon in days; $r_t + r_{t-1} + \dots + r_{t-q+1}$ is the log return over q days

- **Sample mean, deviations and variance**

$$\bar{r} = \hat{\mu} = \frac{1}{T} \sum_{t=1}^T r_t, \quad e_t = r_t - \bar{r}, \quad \hat{\sigma}^2 = \frac{1}{T-1} \sum_{t=1}^T e_t^2$$

- ▶ $\hat{\cdot}$ (hat): an estimate computed from the sample; without a hat: the population value
- ▶ $\mathbb{E}[\cdot]$, $\text{Var}(\cdot)$, $\text{Cov}(\cdot, \cdot)$: expectation, variance and covariance of the population

- **i.i.d.:** independent and identically distributed

- ▶ the benchmark of the classical tests
- ▶ daily returns are not i.i.d.: calm years alternate with bursts such as 2008–2009 and 2020 (volatility clustering, [Chapter 1](#); chart: The Data at a Glance)

Notation Used in Every Exercise (2/3): Autocorrelation

- Autocorrelation of order k : the correlation between today's return and the return k days earlier

$$\rho_k = \frac{\text{Cov}(r_t, r_{t-k})}{\text{Var}(r_t)}, \quad \hat{\rho}_k = \frac{\sum_{t=k+1}^T e_t e_{t-k}}{\sum_{t=1}^T e_t^2}$$

- k : the lag in days; $\rho_k \in [-1, 1]$; $\rho_k = 0$ for every $k \geq 1$: no linear predictability from the past
- Standard error of $\hat{\rho}_k$ when $\rho_k = 0$
 - i.i.d. returns: $\text{se}(\hat{\rho}_k) = 1/\sqrt{T}$, hence the band $\pm 1.96/\sqrt{T}$
 - robust (volatility clustering allowed, [Chapter 1](#)): $\text{se}(\hat{\rho}_k) = \sqrt{\hat{\tau}_k}$
- Robust variance of $\hat{\rho}_k$ and the robust t statistic

$$\hat{\tau}_k = \frac{\sum_{t=k+1}^T e_t^2 e_{t-k}^2}{\left(\sum_{t=1}^T e_t^2\right)^2}, \quad \hat{\delta}_k = T \hat{\tau}_k, \quad t_k = \frac{\hat{\rho}_k}{\sqrt{\hat{\tau}_k}}$$

- $\hat{\delta}_k$: how many times the robust variance exceeds the i.i.d. one
 - $\hat{\delta}_k \approx 1$ for i.i.d. returns; $\hat{\delta}_k > 1$ when large days follow large days


Notation Used in Every Exercise (3/3): Tests and Decisions

- A **test** compares a statistic with its distribution under the null hypothesis H_0
 - ▶ $z \approx N(0, 1)$ under H_0 : the statistic has (approximately) the standard Normal distribution
 - ▶ Φ : the distribution function of $N(0, 1)$; $\Phi(1.96) = 0.975$
- **p-value** of a two-sided test: the probability, under H_0 , of a statistic at least as extreme

$$p = 2(1 - \Phi(|z|))$$

- ▶ reject H_0 at level $\alpha = 5\%$ when $p < 0.05$, that is when $|z| > 1.96$
 - ▶ the level α is the probability of a false rejection (type I error) that we accept
- Interval of 95%: estimate $\pm 1.96 \times$ standard error
 - ▶ it contains 0 exactly when the two-sided test does not reject at 5%
- Not rejecting is not proof of H_0
 - ▶ a test with low power (few observations, noisy data) seldom rejects even when H_0 is false


The Data at a Glance



What You Need for Today (1/23): Weak-Form Efficiency

- ▣ Details: [Chapter 2](#)
- ▣ **Efficient market Fama (1970)**: prices reflect the available information
 - ▶ **weak form**: the information is the past prices and returns
 - ▶ **semi-strong form**: all public information (announcements, news; [B9](#))
 - ▶ **strong form**: also private information
- ▣ Testable consequence of the weak form
 - ▶ past returns do not allow risk-adjusted profits after costs
 - ▶ in particular, $\mathbb{E}[r_t \mid r_{t-1}, r_{t-2}, \dots]$ does not depend on the past: returns are unpredictable from their own history
- ▣ Sign of the dependence
 - ▶ $\rho_k > 0$: persistence (momentum); $\rho_k < 0$: reversal (mean reversion)
 - ▶ a significant $\hat{\rho}_k$ is not yet a profitable rule: costs, risk and stability decide

What You Need for Today (2/23): Three Random Walks




- **Random walk of the log price** [Campbell, Lo & MacKinlay \(1997\)](#)

$$\ln P_t = \mu + \ln P_{t-1} + \varepsilon_t, \quad \text{so } r_t = \mu + \varepsilon_t$$

- ▶ μ : the drift, the expected daily log return; ε_t : the shock of day t , with mean 0
- ▶ RW1: ε_t i.i.d.; RW2: independent, the variance may change (middle panel)
- ▶ RW3: only uncorrelated, $\text{Cov}(\varepsilon_t, \varepsilon_{t-k}) = 0$, while ε_t^2 may be autocorrelated (GARCH, right)

- All three panels have $\hat{\rho}_1 \approx 0$; the right one has volatility clustering, the relevant null for daily returns is therefore RW3

What You Need for Today (3/23): Two Bands for the Autocorrelation



- i.i.d. band $\pm 1.96/\sqrt{T}$
- robust band $\pm 1.96\sqrt{\hat{t}_k}$
- $\hat{\rho}_k$ of a GARCH series

- ▣ Simulated GARCH returns, $T = 2500$: $\rho_k = 0$ at every lag
 - ▶ i.i.d. band $\pm 1.96/\sqrt{T} = \pm 0.039$ (shaded)
 - ▶ robust band $\pm 1.96\sqrt{\hat{t}_k}$, on average ± 0.061 (dashed)
- ▣ Reading the chart
 - ▶ 6 of 20 lags leave the i.i.d. band: false rejections, since the true ρ_k are all 0
 - ▶ no lag leaves the robust band
 - ▶ here $\hat{\delta}_k \approx 2.5$: the robust standard error is $\sqrt{2.5} \approx 1.6$ times larger
- ▣ Why
 - ▶ $\text{Var}(e_t e_{t-k})$ is larger than σ^4 when e_t^2 and e_{t-k}^2 are positively correlated




What You Need for Today (4/23): The Variance Ratio

- Uncorrelated returns with constant variance σ^2 : the variance of a q -day sum is $q\sigma^2$
 - ▶ variances add up only when the covariances are zero
- **Variance ratio** Lo & MacKinlay (1988): the variance of the q -day return relative to q one-day variances

$$VR(q) = \frac{\text{Var}(r_t + r_{t-1} + \dots + r_{t-q+1})}{q \text{Var}(r_t)}$$

- ▶ $VR(q) = 1$ under RW1, RW2 and RW3; $VR(1) = 1$ always
- ▶ $VR(q) > 1$: persistence, q -day moves larger than q independent days
- ▶ $VR(q) < 1$: reversal, part of each move is undone in the following days
- Case $q = 2$
$$\text{Var}(r_t + r_{t-1}) = 2\sigma^2 + 2\text{Cov}(r_t, r_{t-1}) \Rightarrow VR(2) = 1 + \rho_1$$
 - ▶ example: $\rho_1 = -0.1$ gives $VR(2) = 0.9$, two-day sums vary 10% less than two separate days

What You Need for Today (5/23): The Variance Ratio as a Weighted Sum



- Expanding the variance of the q -day sum: lag k appears in $2(q - k)$ covariances

$$VR(q) = 1 + 2 \sum_{k=1}^{q-1} \left(1 - \frac{k}{q}\right) \rho_k$$

- Left: the weight $2(1 - k/q)$ falls linearly with the lag; right: $VR(q)$ for $AR(1)$ returns, $\rho_k = \phi^k$
 - example: $\rho_1 = -0.1$, all other $\rho_k = 0$: $VR(5) = 1 + 2(1 - 1/5)(-0.1) = 0.84$

What You Need for Today (6/23): The Lo–MacKinlay Estimator

- One-day variance: $\hat{\sigma}_a^2 = \hat{\sigma}^2$, the sample variance of the Notation slides (1/3), with $\hat{\mu} = \bar{r}$
 - ▶ the subscript a marks the one-day estimator
- Variance of the **overlapping** q -day sums, divided by q

$$\hat{\sigma}_c^2(q) = \frac{1}{m} \sum_{t=q}^T \left(\sum_{j=0}^{q-1} r_{t-j} - q\hat{\mu} \right)^2, \quad m = q(T - q + 1) \left(1 - \frac{q}{T} \right)$$

- ▶ overlapping: a new q -day sum starts every day, $T - q + 1$ sums in total
- ▶ $q\hat{\mu}$: the expected q -day return; m : a divisor that removes the small-sample bias
- The estimator
$$\widehat{VR}(q) = \frac{\hat{\sigma}_c^2(q)}{\hat{\sigma}_a^2}$$
 - ▶ in [A1](#) the simple version (divisors T and $T - 1$, no bias correction) gives the same idea by hand

What You Need for Today (7/23): Testing $VR(q) = 1$

- Under RW1 (i.i.d.): the classical statistic

$$z(q) = \frac{\widehat{VR}(q) - 1}{\sqrt{2(2q-1)(q-1)/(3qT)}} \approx N(0, 1)$$

- the denominator is the i.i.d. standard error of $\widehat{VR}(q)$; for $q = 2$ it equals $1/\sqrt{T}$

- Under RW3 (volatility clustering allowed): the robust statistic


$$z^*(q) = \frac{\sqrt{T} (\widehat{VR}(q) - 1)}{\sqrt{\hat{\theta}(q)}}, \quad \hat{\theta}(q) = \sum_{k=1}^{q-1} \left(\frac{2(q-k)}{q} \right)^2 \hat{\delta}_k$$

- $\hat{\theta}(q)$: the robust variance of $\sqrt{T} \widehat{VR}(q)$; the weights are those of ρ_k in $VR(q)$, squared
- $\hat{\delta}_k = T \hat{\tau}_k$ from the Notation slides; $|z^*(q)| > 1.96$ rejects RW3 at 5%

- Example: $T = 2\,500$, $\widehat{VR}(2) = 0.95$

- $z(2) = -0.05 \times 50 = -2.5$: a rejection at 5%
- with $\hat{\delta}_1 = 4$: $z^*(2) = -2.5/\sqrt{4} = -1.25$, no rejection

What You Need for Today (8/23): Null Distributions of z and z^*



- $z(2)$: rejects 18.4%
- $z^*(2)$: rejects 6.2%
- $N(0, 1)$

- 1 000 simulated GARCH samples, $T = 2\,500$: no autocorrelation, so H_0 of RW3 is true
 - ▶ $z(2)$ is too dispersed (standard deviation 1.55, not 1)
 - ▶ $|z(2)| > 1.96$ in 18.4% of the samples instead of 5%
 - ▶ $z^*(2)$ follows $N(0, 1)$: 6.2% rejections
- Lesson for B1
 - ▶ with volatility clustering a large $|z|$ alone is not evidence: compare $|z^*|$ with 1.96
 - ▶ $z^* = z / \sqrt{\hat{\delta}_1}$ for $q = 2$: the effective sample is $\hat{\delta}_1$ times smaller


What You Need for Today (9/23): Several Horizons at Once: Chow–Denning

- Testing four horizons separately at 5% rejects *some* horizon more often than 5%
 - ▶ the more horizons, the more chances of a false rejection (multiple testing, [What You Need \(17/23\)](#))
- **Chow–Denning** statistic [Chow & Denning \(1993\)](#): the largest robust statistic

$$CD = \max_{q \in \{2, 5, 10, 20\}} |z^*(q)|$$


- ▶ critical value c from $(2\Phi(c) - 1)^4 = 0.95$, i.e. $c = 2.49$ (not 1.96)
 - ▶ p-value: $p = 1 - (2\Phi(CD) - 1)^4$
 - ▶ the exponent 4 is the number of horizons: the bound treats them as independent, so the test is slightly conservative
- Reading
 - ▶ $CD > 2.49$: RW3 is rejected jointly at 5%; the horizon with the largest $|z^*|$ shows where

What You Need for Today (10/23): The Hurst Exponent



- **Hurst exponent** $H \in (0, 1)$, from the hydrology of the Nile [Hurst \(1951\)](#): how fast the spread of a cumulated series grows with the horizon
 - ▶ self-similar increments (fractional Gaussian noise, fGn): $\text{Var}(r_t + \dots + r_{t-q+1}) = q^{2H} \text{Var}(r_t)$, so $VR(q) = q^{2H-1}$
 - ▶ $H = 0.5$: no memory; $H > 0.5$: smooth, trending paths; $H < 0.5$: anti-persistent, rough paths (left)
- Implied exponent $H(q)$: the H that solves $VR(q) = q^{2H-1}$ at a single horizon q ; [A2](#) derives its formula

What You Need for Today (11/23): Short and Long Memory



- AR(1), $\phi = 0.4$: $\rho_k = \phi^k$
- fGn, $H = 0.75$: $\rho_k \approx H(2H - 1)k^{2H-2}$

□ Short memory, AR(1) (first-order autoregression)

$$r_t = \phi r_{t-1} + u_t, \quad \rho_k = \phi^k$$



- ▶ ϕ : the lag-1 autocorrelation, $|\phi| < 1$; u_t : an i.i.d. shock
- ▶ geometric decay: a straight line on the log scale (blue)
- ▶ $\text{VR}(q) \rightarrow (1 + \phi)/(1 - \phi)$, bounded, so $H(q) \rightarrow 0.5$ as q grows

□ Long memory, fGn with $H > 0.5$

$$\rho_k = \frac{1}{2} (|k+1|^{2H} - 2|k|^{2H} + |k-1|^{2H})$$

- ▶ hyperbolic decay $\rho_k \approx H(2H - 1)k^{2H-2}$ (red): at $k = 20$, 0.084 against 10^{-8} for AR(1)
- One horizon q cannot tell the two apart: compare $H(q)$ over many q (A2, A3)


What You Need for Today (12/23): Detrended Fluctuation Analysis (DFA)



- | | |
|---|--|
| — Profile y_k | — i.i.d.: slope $\hat{\alpha}_{DFA} = 0.53$ |
| — Local linear trend | — fGn, $H = 0.75$: slope $\hat{\alpha}_{DFA} = 0.77$ |

- **DFA Peng et al. (1994):** estimates H from how the detrended spread of the profile grows with the segment length
 - ▶ profile $y_k = \sum_{i \leq k} (r_i - \bar{r})$; cut it into segments of length n and fit a line in each (left)
 - ▶ $F(n)$: root mean square of the residuals around the local lines; $F(n) \propto n^{\alpha_{DFA}}$
 - ▶ $\hat{\alpha}_{DFA}$: OLS slope of $\ln F(n)$ on $\ln n$ (right); 0.5 without memory, above 0.5 with long memory


What You Need for Today (13/23): The Sign-Flip (Wild) Bootstrap



□ **Wild bootstrap** Kim (2006): new series $e_t^* = e_t \eta_t$, with $\eta_t = \pm 1$ with probability 1/2, independent

- ▶ $|e_t^*| = |e_t|$: volatility clustering is kept; random signs destroy any autocorrelation, so H_0 holds in every copy
- ▶ repeat B times (here 999), compute the statistic each time; the 2.5% and 97.5% quantiles form the null band
- ▶ here the band for $\hat{\rho}_1$ is $[-0.072; 0.074]$, wider than the i.i.d. band ± 0.062

What You Need for Today (14/23): Rolling Windows and Adaptive Markets



band $\pm 1.96/\sqrt{W}$ if $\rho_1 = 0$ $--$ true ρ_1 $\bullet$ $\hat{\rho}_1$ on the last $W = 250$ days

- **Rolling window:** re-estimate on the last W days, moving the window forward by a fixed step
 - ▶ simulated: $\rho_1 = 0.25$ for 1000 days, then fading to 0; $W = 250$, step 21 days
 - ▶ overlapping windows share most data: neighbouring estimates are correlated, so counts are descriptive
- **Adaptive markets hypothesis (AMH)** [Lo \(2004\)](#): efficiency changes over time with competition and learning
 - ▶ a predictable pattern attracts traders until it fades, as in the simulation after day 1000

What You Need for Today (15/23): Calendar Effects as a Regression

- **Turn of the month (TOM)** [Ariel \(1987\)](#): the last trading day of a month and the first three of the next

$$r_t = a + b \text{TOM}_t + u_t$$

- ▶ $\text{TOM}_t = 1$ on a turn-of-the-month day, 0 otherwise (a dummy variable); u_t : the residual
- ▶ OLS (ordinary least squares) chooses $\hat{a}$, $\hat{b}$ to minimise $\sum_t \hat{u}_t^2$

- With a 0/1 regressor OLS compares two means

$$\hat{a} = \bar{r}_{\text{other}}, \quad \hat{b} = \bar{r}_{\text{TOM}} - \bar{r}_{\text{other}}$$

- ▶ $H_0: b = 0$, no TOM effect; $t = \hat{b} / \text{se}(\hat{b})$

- The other effects of [B4](#)

- ▶ **January**: monthly returns on a January dummy, same regression
- ▶ **day of the week**: five dummies; a Wald test of equal means, $W = (\mathbf{R}\hat{\beta})'(\mathbf{R}\hat{\mathbf{V}}\mathbf{R}')^{-1}\mathbf{R}\hat{\beta} \approx \chi_4^2$
- ▶ $\mathbf{R}\hat{\beta}$: the four differences from Monday's mean; $\hat{\mathbf{V}}$: their HAC covariance (next slide)


What You Need for Today (16/23): Newey–West (HAC) Standard Errors

- **HAC** (heteroskedasticity- and autocorrelation-consistent) [Newey & West \(1987\)](#): the variance of $\hat{\beta}$ when the residuals change variance and are correlated over L lags

$$\widehat{\text{Var}}(\hat{\beta}) = (\mathbf{X}'\mathbf{X})^{-1} \hat{\mathbf{S}} (\mathbf{X}'\mathbf{X})^{-1}, \quad \hat{\mathbf{S}} = \sum_{j=-L}^L \left(1 - \frac{|j|}{L+1}\right) \sum_t \mathbf{x}_t \hat{u}_t \hat{u}_{t-j} \mathbf{x}_{t-j}'$$

- ▶ $\mathbf{X}$: the $T \times 2$ matrix with rows $\mathbf{x}_t' = (1, \text{TOM}_t)$; $\hat{u}_t$: the OLS residuals
- ▶ weights $1 - |j|/(L+1)$ (Bartlett): 1 at $j = 0$, falling linearly to 0 after lag L
- ▶ $L = 0$: the White (heteroskedasticity-robust) errors; $L = 5$ in [B3](#)
- Classical OLS errors
 - ▶ $\widehat{\text{Var}}(\hat{\beta}) = \hat{\sigma}_u^2 (\mathbf{X}'\mathbf{X})^{-1}$: constant variance, no autocorrelation
 - ▶ HAC > OLS when turbulent days cluster or $\rho_1 > 0$; the point estimate is the same, only the uncertainty changes

What You Need for Today (17/23): Multiple Testing: FWER, Bonferroni and Holm



- **FWER** (family-wise error rate): the probability of at least one false rejection among M tests; independent tests at 5%: $1 - 0.95^M$ (left), $M = 10$: 40.1%

- ▶ **Bonferroni**: reject test i if $p_i < \alpha/M$; p_i : its p-value, α : the target FWER
- ▶ **Holm** [Holm \(1979\)](#): sort $p_{(1)} \leq \dots \leq p_{(M)}$, reject while $p_{(j)} < \alpha/(M - j + 1)$, stop at the first failure
- ▶ both keep $\text{FWER} \leq \alpha$; Holm rejects at least as much (right: 4 rejections against 2)


What You Need for Today (18/23): The Runs Test (Extension Pack 1)

- **Run:** a maximal block of consecutive returns with the same sign [Wald & Wolfowitz \(1940\)](#)
 - ▶ n_1 : positive days, n_2 : negative days, $n = n_1 + n_2$; R : the number of runs = sign changes + 1
- Under H_0 (independent signs), conditional on n_1 and n_2

$$\mathbb{E}[R] = \frac{2n_1n_2}{n} + 1, \quad \text{Var}[R] = \frac{2n_1n_2(2n_1n_2 - n)}{n^2(n-1)}, \quad z = \frac{R - \mathbb{E}[R]}{\sqrt{\text{Var}[R]}} \approx N(0, 1)$$

- ▶ too few runs ($z < 0$): signs persist; too many runs ($z > 0$): signs alternate
 - ▶ for small n use the exact distribution of R given n_1, n_2 instead of the Normal approximation
- Example: ++--+-: $n_1 = 3$, $n_2 = 3$, $R = 4$
 - ▶ $\mathbb{E}[R] = 2 \cdot 9/6 + 1 = 4$: exactly as many runs as expected

What You Need for Today (19/23): Predictive Regressions and the Kendall Bias (Pack 3)



OLS $\hat{\rho}$, 20 000 samples

true $\rho = 0.95$

mean of $\hat{\rho} = 0.859$

Kendall: $\rho - (1 + 3\rho)/T = 0.873$

- **Predictive regression:** next period's return on a predictor x_t known today

$$r_{t+1} = \alpha + \beta x_t + u_{t+1}, \quad x_{t+1} = \theta + \rho x_t + v_{t+1}$$

- ▶ x_t : e.g. the log dividend–price ratio
 $dp_t = \ln D_t - \ln P_t$, very persistent (ρ near 1)
- ▶ u, v : shocks; a price rise lowers x , so $\text{corr}(u, v) < 0$

- **Kendall bias:** OLS underestimates a persistent ρ in short samples

$$\mathbb{E}[\hat{\rho} - \rho] \approx -\frac{1 + 3\rho}{T}$$

- ▶ chart: $\rho = 0.95$, $T = 50$; the mean of $\hat{\rho}$ is 0.859, below the truth
- ▶ when u and v are correlated this bias passes to $\hat{\beta}$
 Stambaugh (1999); A6 derives how

What You Need for Today (20/23): Out-of-Sample Forecasts (Pack 3)

- Recursive forecast: at each t estimate on data up to t only, forecast $\hat{r}_{t+1}$; the benchmark is the historical mean $\bar{r}_t$
- **Out-of-sample R^2** Welch & Goyal (2008)

$$R_{OS}^2 = 1 - \frac{\sum_t (r_{t+1} - \hat{r}_{t+1})^2}{\sum_t (r_{t+1} - \bar{r}_t)^2}$$

- ▶ $R_{OS}^2 > 0$: the model beats the historical mean; it can be negative
- **Clark–West** Clark & West (2007): for nested models (the mean is the model with $\beta = 0$)

$$f_{t+1} = (r_{t+1} - \bar{r}_t)^2 - [(r_{t+1} - \hat{r}_{t+1})^2 - (\bar{r}_t - \hat{r}_{t+1})^2]$$

- ▶ t statistic of the mean of f_{t+1} (HAC), one-sided: reject equal accuracy when $t > 1.645$
- ▶ the last term removes the noise of estimating β , which Diebold–Mariano (the plain difference of squared errors) penalises
- **Campbell–Thompson** restrictions Campbell & Thompson (2008)
 - ▶ set $\hat{\beta} = 0$ when it has the wrong sign and replace a negative forecast by 0

What You Need for Today (21/23): Event Studies (Pack 4)

- ▣ **Market model** on the estimation window (T_0 days before the event)

$$r_{i,t} = a_i + b_i r_{m,t} + \varepsilon_{i,t}, \quad AR_{i,t} = r_{i,t} - \hat{a}_i - \hat{b}_i r_{m,t}$$

- ▶ $r_{m,t}$: the market (BET) return; AR : the abnormal return of an event-window day; $\hat{\sigma}_\varepsilon$: the residual SD

- ▣ Cumulated abnormal return and its variance

$$CAR[\tau_1, \tau_2] = \sum_{t=\tau_1}^{\tau_2} AR_{i,t}, \quad \text{Var}(AR_t) = \hat{\sigma}_\varepsilon^2 (1 + \mathbf{x}_t' (\mathbf{X}_0' \mathbf{X}_0)^{-1} \mathbf{x}_t)$$

- ▶ $\mathbf{x}_t = (1, r_{m,t})'$; $\mathbf{X}_0$: the rows $\mathbf{x}_t'$ of the estimation window
- ▶ the bracket (prediction error) exceeds 1 when $r_{m,t}$ is far from its usual values, e.g. a market crash on day 0

- ▣ Several firms on the same date: their residuals are correlated (average correlation $\bar{r}$)

- ▶ Kolari–Pynnönen [Kolari & Pynnönen \(2010\)](#): multiply the cross-sectional t by $\sqrt{(1 - \bar{r}) / (1 + (N - 1)\bar{r})}$, N : number of firms

What You Need for Today (22/23): Block Bootstrap and a Control Group (Pack 5)

- **Moving-block bootstrap** [Künsch \(1989\)](#): resample blocks of ℓ consecutive days, not single days
 - ▶ block starts drawn uniformly from $1, \dots, T - \ell + 1$; blocks are glued until the series has length T
 - ▶ the dependence inside each block (autocorrelation, volatility clustering) is kept
 - ▶ percentile interval: the 2.5% and 97.5% quantiles of the B bootstrap estimates
- **Before/after contrast** for an event (a reclassification)

$$\Delta_i = \hat{\rho}_{1,i}^{\text{after}} - \hat{\rho}_{1,i}^{\text{before}}$$

- ▶ i : a market; other events in the same years (a pandemic, new listings) also change Δ_i : confounders

- **Control group**: markets without the event

$$\Delta_{\text{treated}} - \overline{\Delta}_{\text{control}}$$

- ▶ valid under *parallel trends*: without the event the treated market would have changed like the controls

What You Need for Today (23/23): Many Predictors: Ridge and Cross-Validation (Pack 5)

- With J predictors close to the number of observations, OLS fits noise: high in-sample R^2 , poor forecasts
- **Ridge regression**: penalise large coefficients

$$\hat{h}_\lambda = \arg \min_h \sum_t (y_t - x'_t h)^2 + \lambda \|h\|^2 = (X'X + \lambda I)^{-1} X'y$$

- ▶ $\lambda \geq 0$: the penalty; $\lambda = 0$ is OLS, $\lambda \rightarrow \infty$ shrinks every coefficient to 0
- ▶ predictors standardised first (mean 0, SD 1), so that the penalty treats them alike
- **Cross-validation**: choose λ by the prediction error on data left out of the estimation
 - ▶ leave-one-year-out: estimate without year y , predict year y , repeat for every year, average the errors
- **In-sample against out-of-sample** [Martin & Nagel \(2022\)](#)
 - ▶ with learning, in-sample tests reject efficiency although nothing is predictable in real time; only the out-of-sample result is implementable

A1 [Solved]: The Variance Ratio, Line by Line — Task

▣ Questions

- ▶ what does $\widehat{VR}(2)$ compute?
- ▶ how does it relate to $\hat{\rho}_1$?

- ▣ **Data:** eight daily returns in %, in time order: 1.0; -0.5; 0.8; 0.3; -1.2; 0.6; 0.9; -0.4; no other data; simple estimator, sums of squares divided by the number of terms

▣ Tasks

1. Compute the mean $\hat{\mu}$, the deviations $e_t = r_t - \hat{\mu}$ and the one-day variance $\hat{\sigma}_1^2 = \frac{1}{8} \sum_t e_t^2$.
2. Compute the lag-1 autocorrelation $\hat{\rho}_1 = \sum_{t=2}^8 e_t e_{t-1} / \sum_{t=1}^8 e_t^2$.
3. Form the seven overlapping two-day sums $s_t = r_{t-1} + r_t$ and compute $\hat{\sigma}_2^2 = \frac{1}{7} \sum_t (s_t - 2\hat{\mu})^2$.
4. Compute $\widehat{VR}(2) = \hat{\sigma}_2^2 / (2\hat{\sigma}_1^2)$ and compare it with $1 + \hat{\rho}_1$.

- ▣ **Report:** a worksheet with columns t , r_t , e_t , e_t^2 , $e_t e_{t-1}$, s_t , $(s_t - 2\hat{\mu})^2$ and one sentence on the sign of $\widehat{VR}(2) - 1$

- ▣ Notebook: section [A1](#)


A1 [Solved]: Solution (1/3) — Worksheet

t	r_t	e_t	e_t^2	$e_t e_{t-1}$	s_t	$(s_t - 2\hat{\mu})^2$
1	+1.0	+0.8125	0.6602	—	—	—
2	−0.5	−0.6875	0.4727	−0.5586	+0.5	0.0156
3	+0.8	+0.6125	0.3752	−0.4211	+0.3	0.0056
4	+0.3	+0.1125	0.0127	+0.0689	+1.1	0.5256
5	−1.2	−1.3875	1.9252	−0.1561	−0.9	1.6256
6	+0.6	+0.4125	0.1702	−0.5723	−0.6	0.9506
7	+0.9	+0.7125	0.5077	+0.2939	+1.5	1.2656
8	−0.4	−0.5875	0.3452	−0.4186	+0.5	0.0156
sum			4.4688	−1.7639		4.4044


$$\hat{\mu} = 1.5/8 = 0.1875; \hat{\sigma}_1^2 = 4.4688/8 = 0.5586; \hat{\rho}_1 = -1.7639/4.4688 = -0.395$$

$$\hat{\sigma}_2^2 = 4.4044/7 = 0.6292; \widehat{VR}(2) = 0.6292/(2 \times 0.5586) = 0.563, \text{ against } 1 + \hat{\rho}_1 = 0.605$$

A1 [Solved]: Solution (2/3) — Chart and Interpretation



--- Mean $\hat{\mu} = 0.1875$



--- $2\hat{\mu} = 0.375$ ■ $s_t = r_{t-1} + r_t$

A1 [Solved]: Solution (3/3) — Chart and Interpretation



- $VR(2) = 1 + \rho_1$ holds for population moments; in a sample the two differ by end terms and divisors (7 sums, 8 squares): 0.563 vs 0.605 is a convention gap, not an error
- Interpretation: consecutive returns reverse, so two-day sums vary less than two single days: mean reversion. With $T = 8$ this is arithmetic, not evidence: no test is attempted

 MFM_ch2_seminar

A2 [Solved]: The Implied Hurst Exponent of BET — Task

- ▣ **Question:** does BET's $\widehat{VR}(20) > 1$ imply long memory?
- ▣ **Data:** BET, daily log returns 2000–2026 (closing prices): $\widehat{VR}(20) = 1.377$, $\hat{\rho}_1 = 0.109$ (Lecture 2)
- ▣ **Tasks**
 1. Take logarithms of the fractional-Gaussian-noise relation $VR(q) = q^{2H-1}$ ([What You Need \(10/23\)](#)) and solve for the implied $H(q)$.
 2. Compute the implied $H(20)$ for BET.
 3. Compute $VR(20)$ and $H(20)$ of a short-memory AR(1) with $\rho_k = 0.109^k$, using the weighted sum of [What You Need \(5/23\)](#).
- ▣ **Report:** the formula, the two $H(20)$ values with three decimals, and whether one horizon can identify long memory
- ▣ Notebook: section [A2](#)

A2 [Solved]: Solution (1/2)



— AR(1), $\rho_1 = 0.109$ (BET) — AR(1), $\rho_1 = -0.1$ ♦ BET, implied $H(20) = 0.553$ ♦ S&P 500, implied $H(20) = 0.453$

A2 [Solved]: Solution (2/2)

- (1) $\ln \text{VR}(q) = (2H - 1) \ln q$, so $H(q) = 0.5 + \ln \text{VR}(q)/(2 \ln q)$
- (2) $\ln 1.377 = 0.3199$, $2 \ln 20 = 5.991$: $H(20) = 0.5 + 0.3199/5.991 = 0.553$
- (3) AR(1): $\text{VR}(20) = 1.231$, $H(20) = 0.535$, and $H(250) = 0.520$ (red curve)
- Interpretation: short memory alone gives most of BET's $H(20)$, and its $H(q)$ drifts back to 0.5: one horizon cannot identify long memory

A3 [Proposed]: The Implied Hurst Exponent of the S&P 500 — Task

- ▣ **Question:** is the S&P 500 “anti-persistent”, or only short-run reversing?
- ▣ **Data:** S&P 500, daily log returns 2000–2026 (closing prices): $\widehat{VR}(20) = 0.753$; a short-memory benchmark AR(1) with $\rho_k = (-0.1)^k$
- ▣ **Tasks**
 1. Compute the implied $H(20)$ of the S&P 500 with the formula of [A2](#).
 2. Compute $VR(20)$ and $H(20)$ of the AR(1) benchmark.
 3. Compute $VR(250)$, $H(250)$ and the limit $VR(\infty) = (1 + \rho)/(1 - \rho)$ of the AR(1), and read the blue curve of the [A2](#) chart.
- ▣ **Report:** the five numbers and one sentence: does $H(20) < 0.5$ require long memory?
- ▣ **Model:** [A2](#) [Solved]
- ▣ Notebook: section [A3](#)

B1 [Solved]: Classical vs Robust Variance Ratio — Task

- **Question:** do the S&P 500 and BET reject the random walk once volatility clustering is allowed?
- **Data:** S&P 500 and BET, daily log returns 2000–2026, $T = 6\,714$ and $6\,670$; H_0 of $z(q)$: RW1 (i.i.d.); H_0 of $z^*(q)$: RW3 plus the conditions of [Lo & MacKinlay \(1988\)](#)
- **Tasks**
 1. For the S&P 500 and $q = 2$, compute $\hat{\sigma}_a^2$, $\hat{\sigma}_c^2(2)$, $\widehat{VR}(2)$, $\hat{\delta}_1$, $z(2)$ and $z^*(2)$ with the Lo–MacKinlay formulas ([What You Need \(6/23 and 7/23\)](#)).
 2. For both markets and $q = 2, 5, 10, 20$, compute $\widehat{VR}(q)$, $z(q)$, $z^*(q)$ and the two-sided p -values with `variance_ratio(r, q)`.
 3. Compute the Chow–Denning statistic $\max_q |z^*(q)|$ and decide at 5% with the critical value 2.49; report its p -value ([What You Need \(9/23\)](#)).
- **Report:** a table with one row per market and q ($\widehat{VR}$, z , z^* , p), the CD decision, and one sentence on why $|z| > |z^*|$
- Notebook: section [B1](#)

B1 [Solved]: Solution (1/3) — S&P 500, $q = 2$, Step by Step


- ▣ Returns in %: $T = 6\,714$, $\hat{\mu} = 0.0247$; one-day variance $\hat{\sigma}_a^2 = \frac{1}{T-1} \sum_t e_t^2 = 1.4714$
- ▣ 6 713 overlapping sums $r_t + r_{t-1} - 2\hat{\mu}$; divisor $m = q(T - q + 1)(1 - q/T) = 13\,422$; $\hat{\sigma}_c^2(2) = 1.3256$
- ▣ $\widehat{VR}(2) = 1.3256/1.4714 = 0.9009$
- ▣ Classical: $SD = \sqrt{2(2q-1)(q-1)/(3qT)} = \sqrt{1/T} = 0.0122$, so $z(2) = (-0.0991)/0.0122 = -8.12$
- ▣ Robust: $\hat{\delta}_1 = T \sum_t e_t^2 e_{t-1}^2 / (\sum_t e_t^2)^2 = 4.966$ (it would be about 1 under i.i.d. returns); for $q = 2$, $\hat{\theta}(2) = \hat{\delta}_1$
 - ▶ $z^*(2) = \sqrt{T} (\widehat{VR}(2) - 1) / \sqrt{\hat{\theta}(2)} = 81.94 \times (-0.0991) / \sqrt{4.966} = -3.64$, $p = 0.0003$
- ▣ The robust statistic is $1/\sqrt{\hat{\delta}_1}$ times the classical one: volatility clustering makes the effective sample about 4.966 times smaller

B1 [Solved]: Solution (2/3) — Both Markets, Four Horizons

	$z(2)$	$z^*(2)$	$z(5)$	$z^*(5)$	$z(10)$	$z^*(10)$	$z(20)$	$z^*(20)$
S&P 500	-8.12	-3.64	-6.37	-2.77	-5.65	-2.50	-4.06	-1.86
BET	8.83	3.84	6.44	2.98	5.44	2.73	6.19	3.31

- ▣ $\widehat{VR}(q)$ for $q = 2, 5, 10, 20$: S&P 500 0.901, 0.830, 0.767, 0.753; BET 1.108, 1.173, 1.225, 1.377
- ▣ Chow–Denning: S&P 500 3.64 ($p = 0.0011$), BET 3.84 ($p = 0.0005$): both reject RW3 jointly

B1 [Solved]: Solution (3/3) — Chart and Interpretation



- Interpretation: reversal in the S&P 500, persistence in BET; the robust evidence is two to three times weaker, and a rejection is not a trading rule (costs, risk, stability: [B2](#))

B2 [Proposed]: Is the S&P 500 Reversal a Crisis Effect? — Task

- **Question:** does the S&P 500 reversal of B1 survive without the two crisis periods?
- **Data**
 - ▶ S&P 500, daily log returns 2000–2026, $T = 6\,714$; crisis days, fixed in advance: 1 Sep 2008 – 30 Jun 2009 and 15 Feb – 30 Jun 2020
 - ▶ Removing a period leaves contiguous segments. A lag pair (r_t, r_{t-1}) with $t = 1$ Jul 2009 and $t - 1 = 29$ Aug 2008 would join two days nine months apart: it is forbidden; so is any 5-day sum that spans a gap
- **Tasks**
 1. On all days compute $\hat{\rho}_1$, the i.i.d. $t = \hat{\rho}_1 \sqrt{T}$, the robust $t = \hat{\rho}_1 / \sqrt{\hat{\tau}_1}$ and $z^*(5)$.
 2. Remove the crisis days and repeat within the retained segments; report the kept returns, the valid lag pairs and the valid 5-day sums.
 3. Give robust 95% intervals for ρ_1 in both samples and decide at 5%.
- **Report**
 - ▶ a two-row table and one sentence
 - ▶ by how much does the estimate fall?
 - ▶ is removing crisis days legitimate?
- **Model:** B1 [Solved]
- **Notebook:** section B2

C1 [Proposed]: Critique an AI Answer — the Answer

- **Prompt** sent by a student to an AI assistant: “Test whether the S&P 500 (daily closes, 2000–2026) follows a random walk with the variance ratio. Give Python code and interpret the result”

```
1 r = np.log(sp).diff().dropna()          # sp: S&P 500 closes, as in B1
2 T = len(r)
3 def vr(q):
4     rq = r.rolling(q).sum().dropna()    # overlapping q-day returns
5     return rq.var() / (q * r.var())
6 q = 5
7 z = (vr(q) - 1) / np.sqrt(2 * (2*q - 1) * (q - 1) / (3 * q * T))
8 p = 2 * stats.norm.sf(abs(z))
```

- **The AI's interpretation:** (a) “ $VR(2) = 0.901$, $VR(5) = 0.830$, $z(5) = -6.37$, $p < 0.001$: the random walk is rejected”; (b) “since $VR(2) = 1 + 2\rho_1$, $\rho_1 \approx -0.050$ ”; (c) “ $VR < 1$ means momentum: buy the S&P 500 after up days”

C1 [Proposed]: Critique an AI Answer — Task

- **Question:** which claims of the AI answer does its code actually support?
- **Data:** the code and interpretation on the previous slide; the same S&P 500 returns as [B1](#)
- **Tasks**
 1. Audit four points in a table (line number, claim, correct or wrong, how you check): the estimator convention, the null and its standard error, the VR identity, the economic interpretation.
 2. Say which reported numbers the code computes (it computes only $q = 5$).
 3. Give two numerical checks: $\hat{\rho}_1$ computed directly, and $z^*(5)$ from `variance_ratio`.
 4. Write the corrected interpretation in two sentences.
- **Report:** the audit table, the two checks, the corrected interpretation; log each error as in `AI_ERRORS.md`
- **Model:** [A1](#), [B1](#) [Solved]
- **Notebook:** section [B1](#)

B3 [Solved]: The Turn-of-the-Month Effect with HAC Errors — Task

- **Question:** do returns on turn-of-the-month days exceed those on other days?
- **Data**
 - ▶ S&P 500 and BET, daily log returns 2000–2026 in basis points, $r_t = 10\,000 \ln(P_t/P_{t-1})$, each on its own trading calendar; the incomplete last month (September 2026) is dropped
 - ▶ TOM_t : the turn-of-the-month dummy of [What You Need \(15/23\)](#); example: 31 Jul and 3, 4, 5 Aug 2026 are TOM days
- **Tasks**
 1. Build TOM_t from the trading calendar (rank of the day within its month, from the start and from the end).
 2. Estimate $r_t = a + b TOM_t + u_t$ by OLS; $\hat{a}$ is the mean on other days, $\hat{a} + \hat{b}$ the mean on TOM days.
 3. Compute the OLS and the Newey–West HAC standard errors of $\hat{b}$ with 5 lags [Newey & West \(1987\)](#), both t -statistics and the HAC p -value.
 4. Convert $\hat{b}$ into an approximate extra log return per month (4 TOM days), before costs.
- **Report:** a table (share of TOM days, $\hat{a}$, $\hat{b}$, both SEs and t), a 95% interval chart, one sentence on why OLS and HAC differ
- **Notebook:** section [B3](#)

B3 [Solved]: Solution (1/3) — Design Matrix

date	r_t (bp)	1	TOM _t
2026-07-27	+1.6	1	0
2026-07-28	+21.0	1	0
2026-07-29	−152.8	1	0
2026-07-30	+164.7	1	0
2026-07-31	+69.8	1	1
2026-08-03	+146.8	1	1
2026-08-04	+177.4	1	1
2026-08-05	−16.8	1	1
2026-08-06	−17.6	1	0
2026-08-07	+61.7	1	0

- S&P 500 rows around the July/August 2026 boundary: the regressors are a column of ones and the TOM dummy
- OLS: $\hat{b} = \bar{r}_{\text{TOM}} - \bar{r}_{\text{other}}$; S&P 500 $5.1 - 1.9 = 3.2$ bp; BET $20.4 - 3.1 = 17.3$ bp
- HAC: the Newey–West variance of [What You Need \(16/23\)](#) with $L = 5$, i.e. weights $1 - |j|/6$
- In code: `sm.OLS(r, X).fit(cov_type='HAC', cov_kwds={'maxlags': 5})`


 MFM_ch2_calendar_anomalies

B3 [Solved]: Solution (2/3) — Estimates

	TOM share	$\hat{a}$ (bp)	$\hat{b}$ (bp)	SE_{OLS}	t_{OLS}	SE_{HAC}	t_{HAC} (p)
S&P 500	19.1%	1.9	3.2	3.8	0.86	3.5	0.93 (0.35)
BET	19.2%	3.1	17.3	4.3	3.99	4.8	3.62 (0.0003)

- BET: HAC 95% interval [8.0, 26.7] bp; the HAC SE exceeds the OLS SE (clustering and $\rho_1 > 0$), t falls from 3.99 to 3.62, still a rejection
- S&P 500: [-3.6, 10.1] bp contains 0
- $4\hat{b} \approx 69$ bp (BET): an approximate extra monthly log return relative to the other-day mean, before costs, not a strategy return

B3 [Solved]: Solution (3/3) — Chart and Interpretation



- Interpretation: one effect in one market was picked among many: correct for the effects tried (B4) before looking for a mechanism (salary and pension inflows)

B4 [Proposed]: Fifteen Calendar Tests and Holm — Task (1/2)

- **Question:** which of fifteen calendar effects survive a correction for multiple testing?
- **Data:** the 15 p -values in the table below: 5 markets $\times$ 3 effects (day of week: Wald test of equal means, HAC 5 lags; January: monthly returns, HAC 3 lags; TOM: as in B3), 2000–2026

Market	Effect	p	Market	Effect	p
S&P 500 (US)	day of week (equal means)	0.826	Euro Stoxx 50	turn of the month	0.227
S&P 500 (US)	January	0.426	WIG20 (Poland)	day of week (equal means)	0.226
S&P 500 (US)	turn of the month	0.354	WIG20 (Poland)	January	0.686
BET (Romania)	day of week (equal means)	0.305	WIG20 (Poland)	turn of the month	0.084
BET (Romania)	January	0.137	BUX (Hungary)	day of week (equal means)	0.933
BET (Romania)	turn of the month	0.0003	BUX (Hungary)	January	0.024
Euro Stoxx 50	day of week (equal means)	0.632	BUX (Hungary)	turn of the month	0.0003
Euro Stoxx 50	January	0.774			

B4 [Proposed]: Fifteen Calendar Tests and Holm — Task (2/2)

□ Tasks

1. Count the raw rejections at 5% and compute the FWER of 15 independent tests.
2. Apply Bonferroni: reject if $p_i < 0.05/15$.
3. Apply Holm: sort the p -values, compare $p_{(j)}$ with $0.05/(16 - j)$ and stop at the first failure.

□ **Report:** the sorted table with ranks, thresholds and three decision columns (raw, Bonferroni, Holm), and one sentence on BUX January


□ **Model:** B3 [Solved]

□ **Notebook:** section B4

B5 [Solved]: Rolling DFA on the S&P 500 — Task


- **Question:** does the S&P 500 DFA exponent ever leave a null band that allows volatility clustering?
- **Data**
 - ▶ S&P 500, daily log returns 2000–2026; windows of 500 returns moved by 21 days (296 windows); function `dfa_hurst(x)`
 - ▶ DFA grid: 18 box lengths spaced evenly in $\log n$ from 10 to about 125 (a quarter of the window; 124 after rounding), rounded to integers, duplicates removed; the remainder $T - kn$ of the profile is discarded; random seed = 42 + index of the last day of the window
- **Tasks**
 1. For the last window (2024-09-20 to 2026-09-18) build the profile, detrend it in boxes and fit $\log F(n)$ on $\log n$.
 2. Compute α_{DFA} in every window.
 3. In each window build a sign-flip band ([What You Need \(13/23\)](#)) from 199 series $e_t \eta_t$; keep the 2.5% and 97.5% quantiles of α_{DFA} .
 4. Compute the share of windows outside their own band and outside one fixed i.i.d. Student- t_4 band (300 simulations of length 500).
- **Report:** the one-window fit, the rolling chart, both shares, and one sentence on what the band assumes
- **Notebook:** section [B5](#)

B5 [Solved]: Solution (1/3) — One Window, Step by Step



— Profile $y_k = \sum_{i \leq k} (r_i - \bar{r})$

— Linear trend in each box of $n = 50$ days




● $\log F(n)$, 18 box sizes from 10 to 124

— OLS slope $\hat{\alpha}_{DFA} = 0.457$

— Slope 0.5 (no memory), same start

- Left: the profile y_k of the last window, detrended box by box as in [What You Need \(12/23\)](#)
- 18 box lengths from 10 to 124; the OLS slope of $\log F(n)$ on $\log n$ is $\hat{\alpha}_{DFA} = 0.457$ (right), close to 0.5

B5 [Solved]: Solution (2/3) — All Windows



B5 [Solved]: Solution (3/3) — All Windows

- Outside their own band: 2.0% of 296 windows; outside the fixed Student- t_4 band $[0.378, 0.614]$: 6.4%
- Mean $\hat{\alpha}_{DFA} = 0.450$; minimum 0.341 (window ending 2022-03-15), maximum 0.598 (2021-02-12); the band is widest (0.500) in the window ending 2022-02-11
- Interpretation: the band is a pointwise benchmark (signs symmetric and independent given $|e_t|$), not a confidence interval; overlapping windows are correlated, so the shares describe, they do not test stability; no memory beyond what chance would produce

B6 [Proposed]: Adaptive Efficiency of Bitcoin — Task

- **Question:** has Bitcoin moved towards efficiency since its early years, as the AMH [Lo \(2004\)](#) would allow?
- **Data:** Bitcoin (USD) closing prices, daily log returns 7 days a week from 1 Jan 2011 to 18 Sep 2026 (earlier than the lecture's 2014 start, to include the first bubbles); windows of 365 days moved by 30 days; H_0 : RW3 in every window
- **Tasks**
 1. Compute the robust $z^*(5)$ in each window.
 2. In each window build a sign-flip band for $z^*(5)$ with 999 replications (random seed = 42 + index of the last day of the window) and keep its 2.5% and 97.5% quantiles.
 3. Count the windows outside their band and outside ± 1.96 and compare with $0.05 \times$ the number of windows.
 4. Compare with [Urquhart \(2016\)](#): 2010–2016, a battery of classical tests (Ljung–Box, runs, Bartels, automatic VR, BDS, R/S).
- **Report:** the chart, the flagged windows with z^* and band, both counts, and one sentence: what does “no detected trend” mean?
- **Model:** [B5](#), [B1](#) [Solved]
- Notebook: section [B6](#)

Extension Packs: What Each One Needs

Pack	Exercises	Prerequisites
1. Theory	A4 (null distribution), A5 [Solved] (runs), A6 (Stambaugh bias)	A1, B1
2. Many markets	B7	B1
3. Predictability	B8	A6, B3
4. Event study	B9	Seminar 0, B7; B3
5. Project seeds	C2, C3	B1, B5; the lecture case study

- Each exercise has its own notebook section; do the packs in order

A4 [Proposed]: The Variance Ratio and Its Null Distribution — Task

▣ Questions

- ▶ why is $\sqrt{T}(\widehat{\text{VR}}(q) - 1)$ approximately Normal with variance $2(2q - 1)(q - 1)/(3q)$?
- ▶ when does this fail?

- ▣ **Data:** r_t covariance-stationary, variance γ_0 , autocovariances $\gamma_k = \text{Cov}(r_t, r_{t-k})$, $\rho_k = \gamma_k/\gamma_0$; under RW1 $\sqrt{T}(\hat{\rho}_1, \dots, \hat{\rho}_{q-1})$ tends to $N(0, I_{q-1})$, I_{q-1} the identity matrix (independent, unit variances); S&P 500: $T = 6\,714$, $\widehat{\text{VR}}(2) = 0.901$

▣ Tasks

1. For $q = 3$, write $\text{Var}(r_t + r_{t-1} + r_{t-2})$ as the sum of the nine covariances and count how many times γ_0 , γ_1 , γ_2 appear.
2. Generalise: show that $\text{VR}(q) = 1 + \sum_{k=1}^{q-1} a_k \rho_k$ with $a_k = 2(1 - k/q)$.
3. Under RW1, with $a = (a_1, \dots, a_{q-1})'$, compute $a'la = \sum_k a_k^2$ and simplify it to $2(2q - 1)(q - 1)/(3q)$; evaluate it for $q = 2, 5, 20$.
4. Compute $z(2)$ for the S&P 500 and explain why I must be replaced by $\text{diag}(\delta_k)$, the diagonal matrix of the robust variances δ_k , under RW3.

- ▣ **Report:** the counting, the variance in two columns (asymptotic, of $\sqrt{T}(\widehat{\text{VR}} - 1)$; finite-sample, of $\widehat{\text{VR}}$: divide by T) and $z(2)$

- ▣ **Model:** [A1](#) [Solved]


- ▣ **Notebook:** section [A4](#)

A5 [Solved]: The Runs Test, Step by Step — Task

- ▣ **Question:** are 20 signs consistent with independence?
- ▣ **Data:** signs of 20 consecutive daily returns, no zeros: ++-+++--+-+---+---+; under H_0 the signs are independent with a constant probability of a positive day [Wald & Wolfowitz \(1940\)](#)
- ▣ **Tasks**
 1. Number the signs 1 to 20, mark every sign change and count n_1 , n_2 and the runs R .
 2. Compute $\mathbb{E}[R]$ and $\text{Var}[R]$ with the formulas of [What You Need \(18/23\)](#).
 3. Compute the approximate z and its two-sided p -value.
 4. Compare with the exact conditional distribution of R given n_1 and n_2 .
- ▣ **Report:** n_1 , n_2 , R , $\mathbb{E}[R]$, $\text{Var}[R]$, z , both p -values and what too few or too many runs would mean
- ▣ **Notebook:** section [A5](#)

A5 [Solved]: Solution (1/2)

Signs of 20 returns: $n_1 = 12$ plus, $n_2 = 8$ minus, $R = 11$ runs (black lines: run boundaries)



A5 [Solved]: Solution (2/2)

- (1) Runs: ++ | - | +++ | -- | + | - | ++ | - - - | + | - | +++, so $n_1 = 12$, $n_2 = 8$, $R = 11$ (10 sign changes)
- (2) $\mathbb{E}[R] = 2 \cdot 12 \cdot 8/20 + 1 = 10.60$; $\text{Var}[R] = 192 \cdot 172/(400 \cdot 19) = 4.345$; (3) approximate $z = 0.40/2.085 = 0.19$, $p = 0.85$
- (4) Exact: $P(R = 11) = 0.183$, $P(R \geq 11) = 0.520$, $P(R \leq 11) = 0.663$, so the exact two-sided p is 1. Interpretation: do not reject independence; too few runs would mean persistent signs, too many alternating signs

A6 [Proposed]: The Stambaugh Bias — Task

▣ **Question:** why does a persistent valuation ratio look like a significant predictor too often?

▣ **Data**

- ▶ System, notation and Kendall bias $\mathbb{E}[\hat{\rho} - \rho]$: [What You Need \(19/23\)](#), with x_t the log dividend–price ratio
- ▶ Write $u = \phi v + e$, $\phi = \sigma_{uv} / \sigma_v^2$, with $\sigma_{uv} = \text{Cov}(u, v)$, σ_u , σ_v the shock SDs and e the part of u unrelated to v [Stambaugh \(1999\)](#)
- ▶ Illustrative annual inputs (as fractions, not in %): $\rho = 0.98$, $T = 60$ years, $\text{corr}(u, v) = -0.9$, $\sigma_u = 0.17$, $\sigma_v = 0.14$

▣ **Tasks**

1. With demeaned $\tilde{x}_t = x_t - \bar{x}$, show $\hat{\beta} - \beta = \phi(\hat{\rho} - \rho) + \sum \tilde{x}_t e_{t+1} / \sum \tilde{x}_t^2$ and derive $\mathbb{E}[\hat{\beta} - \beta]$.
2. Compute ϕ , $\mathbb{E}[\hat{\rho} - \rho]$ and the bias of $\hat{\beta}$.
3. Compute $\text{SE}(\hat{\beta}) \approx \sigma_u / (\sigma_x \sqrt{T})$, $\sigma_x = \sigma_v / \sqrt{1 - \rho^2}$, and the bias in standard errors.

▣ **Report:** the derivation, ϕ , the bias, the SE, their ratio, and what a positive bias does to the false-rejection rate

▣ **Model:** [A1](#) [Solved]

▣ **Notebook:** section [A6](#)

B7 [Proposed]: BET-TR and Emerging Markets — Task

- **Question:** are emerging markets less efficient as a group?
- **Data:** BET-TR (from 2014) and WIG20, BUX, PX, BIST 100, Bovespa, IPC, Nifty 50, Shanghai, closing prices, 2000–2026, each on its own calendar
- **Tasks**
 1. Start with BET and BET-TR on the same dates (BET-TR's period): classical $z(2)$, robust $z^*(2)$, Chow–Denning.
 2. Repeat for the eight other indices and report three decisions separately: classical $q = 2$, robust $q = 2$, robust joint over four horizons.
 3. Before any group-wide claim, correct the joint p -values for the nine markets tested (Holm).
- **Report:** one table (one row per market), the matched BET comparison and one sentence: dividends or sample?
- **Model:** [B1](#) [Solved]
- **Notebook:** section [B7](#)

B8 [Proposed]: Predictive Regression and an Out-of-Sample Test — Task (1/2)

▣ **Question:** does the dividend–price ratio predict S&P 500 excess returns, in and out of sample?

▣ **Data**

- ▶ Monthly table 1934–2026 ($T = 1\,111$): P = Shiller S&P Composite monthly average price (to 2026-06), extended with S&P 500 month-end closes; D = trailing 12-month dividends (Shiller, interpolated from quarterly totals), extended with the dividends of the SPY fund (ETF on the S&P 500), inferred from the gap between its closing price and its dividend-adjusted price, spliced on the common period
- ▶ $r_{t+1} = \ln((P_{t+1} + D_{t+1}/12)/P_t) - \ln(1 + \text{TB3MS}_t/1\,200)$ (FRED, known at t); predictor $dp_t = \ln D_t - \ln P_t$, used one month later

B8 [Proposed]: Predictive Regression and an Out-of-Sample Test — Task (2/2)

▣ Tasks

1. Estimate $r_{t+1} = \alpha + \beta dp_t + u_{t+1}$ by OLS with Newey–West errors (12 lags); estimate the AR(1) of dp and $\text{corr}(\hat{u}, \hat{v})$; correct $\hat{\beta}$ for the Stambaugh bias (A6).
2. From January 1970, at each month re-estimate on data up to the previous month and forecast; the benchmark is the historical mean of the same data; compute R_{OS}^2 Welch & Goyal (2008) (pseudo out-of-sample: dividends are interpolated).
3. Compute the Clark–West t (HAC, 12 lags, one-sided) Clark & West (2007); impose the Campbell–Thompson restrictions (slope sign, non-negative forecast) Campbell & Thompson (2008).

▣ **Report:** $\hat{\beta}$, t , $\hat{\rho}$, the bias, R_{OS}^2 , the Clark–West t , and why Diebold–Mariano is the wrong test for nested models

▣ **Model:** A6, B3

▣ **Notebook:** section B8

B9 [Proposed]: The Bank Tax, TLV and BRD Jointly — Task (1/2)

- ▣ **Question:** how much of the bank-tax fall survives proper event-study inference?
- ▣ **Data**
 - ▶ Banca Transilvania (TLV) and BRD (dividend- and split-adjusted prices) and BET (closing prices), prices aligned on common trading days, then log returns
 - ▶ Timeline: tax announced on the evening of 18 Dec 2018; day 0 = 19 Dec; estimation window days $[-250, -11]$ ($T_0 = 240$); event window $[0, +5]$
 - ▶ X_0 (240×2) and X_e (6×2): the $[1, r_{m,t}]$ rows of the estimation and event windows; ι : a 6×1 vector of ones; I_6 : the 6×6 identity matrix; $N, \bar{r}$: as in [What You Need \(21/23\)](#)

B9 [Proposed]: The Bank Tax, TLV and BRD Jointly — Task (2/2)

▣ Tasks

1. For TLV fit the market model on $[-250, -11]$, compute AR_0 , the naive $t = AR_0 / \hat{\sigma}_\varepsilon$ and the t with the prediction-error variance; repeat for BRD.
2. For the equal-weighted portfolio of the two banks fit the market model to its own returns and compute its AR_0 and t ; compare with the t that assumes independent banks.
3. Compute $CAR[0, +5]$ and $CAR[+1, +5]$ with $\text{Var}(CAR) = \iota' \hat{\sigma}_\varepsilon^2 [\mathbf{I}_6 + X_e(X_0' X_0)^{-1} X_e'] \iota$ (adapted to the window used).
4. Compute the residual correlation $\bar{r}$ and the Kolari–Pynnönen factor $\sqrt{(1 - \bar{r}) / (1 + (N - 1)\bar{r})}$ [Kolari & Pynnönen \(2010\)](#).

▣ **Report:** one table (TLV, BRD, portfolio), the AR/CAR chart, and whether the reversal contradicts semi-strong efficiency; note that BET contains both banks

▣ **Model:** Seminar 0, B7; [B3](#)

▣ **Notebook:** section [B9](#)

C2 [Proposed]: BVB Efficiency after the 2020 FTSE Upgrade — Task (1/2)

- **Question:** did BET's lag-1 autocorrelation move towards zero after the upgrade, relative to WIG20 and BUX?
- **Data**
 - ▶ [FTSE Russell \(2019\)](#): announced 26 Sep 2019, effective 21 Sep 2020 (Frontier to Secondary Emerging); background: [Pele & Mazurencu-Marinescu \(2012\)](#)
 - ▶ BET, WIG20, BUX, closing prices; before: 21 Sep 2015 – 20 Sep 2020; after: 21 Sep 2020 – 21 Sep 2025; data after 21 Sep 2025 only in a separately reported extension

C2 [Proposed]: BVB Efficiency after the 2020 FTSE Upgrade — Task (2/2)

▣ Tasks

1. For each market compute $\hat{\rho}_1$ with robust SE and $z^*(5)$ in both windows.
2. Build a 95% interval for the change in $\hat{\rho}_1$ by a moving-block bootstrap, independent within each period (blocks of 20 days, 999 replications, seed 42) [Künsch \(1989\)](#).
3. Plot rolling 250-day $\hat{\rho}_1$ with robust intervals for the three markets, marking the announcement, the effective date and the end of the formal window.
4. List the confounders (COVID-19, the 2023 Hidroelectrica listing, liquidity) and say what a control-adjusted contrast would need.

▣ **Report:** a Quantlet folder and one slide: conclusion, main limitation, link with the AMH [Lo \(2004\)](#); judge movement towards zero, not the sign

▣ **Model:** [B1](#), [B5](#) [Solved]

▣ **Notebook:** section [C2](#)

C3 [Proposed]: Big Data and Efficiency, In and Out of Sample — Task (1/2)

- **Question:** with 238 predictors, is in-sample predictability of industry returns matched out of sample?
- **Data**
 - ▶ Lecture 2 case study [Martin & Nagel \(2022\)](#): rational learning makes in-sample tests reject although out-of-sample returns are unpredictable; replicate the design of Sec. 7 of their paper
 - ▶ 49 industry portfolios, [Kenneth French Data Library](#), value-weighted monthly returns, July 1926 to the last month; missing values −99.99 and −999 dropped month by month; returns divided by 100
 - ▶ Dimensions: in month t , X_t is $N_t \times 238$ (returns and squared returns at lags 2–120), N_t the industries with a complete 120-month history; h_t the 238×1 coefficients

C3 [Proposed]: Big Data and Efficiency, In and Out of Sample — Task (2/2)

▣ Tasks

1. Demean the returns and each predictor across industries every month, scale each predictor to unit cross-sectional SD, give each month equal weight.
2. Estimate pooled OLS h_t over the 240 months ending in t (data up to t only); compute $roos_{t+1} = \frac{1}{N} h_t' X_{t+1}' r_{t+1}$ and rs over the same window; state the first forecast month.
3. Average $roos$ over 10-year windows with descriptive ± 2 SE bands (as in Fig. 8 of the paper).
4. Fit a ridge from January 1971 with standardised predictors; choose the penalty on a log grid by leave-one-year-out CV (it may use later years: a validation, not a forecast); plot the coefficients by lag.

▣ **Report:** the chart of step 3, early and late means of rs and $roos$ with SE, the coefficient plot, and why 49 portfolios can differ from thousands of stocks

▣ **Model:** B3, B8

▣ **Notebook:** section C3 (lecture Quantlet MFM_ch2_big_data)

References (1/2)

- ▣ Ariel, R.A. (1987). [A monthly effect in stock returns](#). *Journal of Financial Economics*, 18(1), 161–174.
- ▣ Campbell, J.Y., Lo, A.W., MacKinlay, A.C. (1997). [The Econometrics of Financial Markets](#). Princeton University Press.
- ▣ Campbell, J.Y., Thompson, S.B. (2008). [Predicting excess stock returns out of sample: can anything beat the historical average?](#). *Review of Financial Studies*, 21(4), 1509–1531.
- ▣ Chow, K.V., Denning, K.C. (1993). [A simple multiple variance ratio test](#). *Journal of Econometrics*, 58(3), 385–401.
- ▣ Clark, T.E., West, K.D. (2007). [Approximately normal tests for equal predictive accuracy in nested models](#). *Journal of Econometrics*, 138(1), 291–311.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.
- ▣ French, K.R. (2026). [Kenneth R. French Data Library](#). Tuck School of Business, Dartmouth College.
- ▣ FTSE Russell (2019). [FTSE Russell announces results of country classification review](#) (press release, 26 September 2019).
- ▣ Harvey, C.R., Liu, Y., Zhu, H. (2016). [... and the cross-section of expected returns](#). *Review of Financial Studies*, 29(1), 5–68.
- ▣ Kim, J.H. (2006). [Wild bootstrapping variance ratio tests](#). *Economics Letters*, 92(1), 38–43.
- ▣ Kolari, J.W., Pynnönen, S. (2010). [Event study testing with cross-sectional correlation of abnormal returns](#). *Review of Financial Studies*, 23(11), 3996–4025.
- ▣ Künsch, H.R. (1989). [The jackknife and the bootstrap for general stationary observations](#). *Annals of Statistics*, 17(3), 1217–1241.

References (2/2)

- ▣ Lo, A.W. (1991). Long-term memory in stock market prices. *Econometrica*, 59(5), 1279–1313.
- ▣ Lo, A.W. (2004). The adaptive markets hypothesis. *Journal of Portfolio Management*, 30(5), 15–29.
- ▣ Lo, A.W., MacKinlay, A.C. (1988). Stock market prices do not follow random walks: evidence from a simple specification test. *Review of Financial Studies*, 1(1), 41–66.
- ▣ MacKinlay, A.C. (1997). Event studies in economics and finance. *Journal of Economic Literature*, 35(1), 13–39.
- ▣ Martin, I.W.R., Nagel, S. (2022). Market efficiency in the age of big data. *Journal of Financial Economics*, 145(1), 154–177.
- ▣ Newey, W.K., West, K.D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.
- ▣ Pele, D.T., Mazurencu-Marinescu, M. (2012). Modelling stock market crashes: the case of Bucharest Stock Exchange. *Procedia – Social and Behavioral Sciences*, 58, 533–542.
- ▣ Peng, C.-K., Buldyrev, S.V., Havlin, S., Simons, M., Stanley, H.E., Goldberger, A.L. (1994). Mosaic organization of DNA nucleotides. *Physical Review E*, 49(2), 1685–1689.
- ▣ Romano, J.P., Wolf, M. (2005). Stepwise multiple testing as formalized data snooping. *Econometrica*, 73(4), 1237–1282.
- ▣ Stambaugh, R.F. (1999). Predictive regressions. *Journal of Financial Economics*, 54(3), 375–421.
- ▣ Urquhart, A. (2016). The inefficiency of Bitcoin. *Economics Letters*, 148, 80–82.
- ▣ Wald, A., Wolfowitz, J. (1940). On a test whether two samples are from the same population. *Annals of Mathematical Statistics*, 11(2), 147–162.
- ▣ Welch, I., Goyal, A. (2008). A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 21(4), 1455–1508.