



Modelling Financial Markets

Seminar 4: Portfolio Optimisation — Hands-on



Daniel Traian PELE
Antoaneta Amza

Bucharest University of Economic Studies
IDA Institute Digital Assets
Blockchain Research Center
AI4EFin Artificial Intelligence for Energy Finance
Romanian Academy, Institute for Economic Forecasting
MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AAPL: Apple Inc. (Nasdaq ticker)
- ▣ AI: Artificial Intelligence
- ▣ AMZN: Amazon.com, Inc. (Nasdaq ticker)
- ▣ BAC: Bank of America Corporation (NYSE ticker)
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ BET-TR: BET Total Return (index)
- ▣ BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- ▣ CSCO: Cisco Systems, Inc. (Nasdaq ticker)
- ▣ ERC: Equal Risk Contribution portfolio
- ▣ ETF: Exchange-Traded Fund
- ▣ GLS: Generalised Least Squares
- ▣ GME: GameStop Corp. (NYSE ticker)
- ▣ GMV: Global Minimum-Variance portfolio
- ▣ GMV-LW: Global Minimum-Variance portfolio with Ledoit–Wolf covariance
- ▣ GMV-NL: Global Minimum-Variance portfolio with Nonlinear shrinkage covariance
- ▣ GOOGL: Alphabet Inc., Class A shares (Nasdaq ticker)
- ▣ GS: The Goldman Sachs Group, Inc. (NYSE ticker)
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HRP: Hierarchical Risk Parity
- ▣ HYG: iShares iBoxx High Yield Corporate Bond ETF
- ▣ IEF: iShares 7–10 Year Treasury Bond ETF
- ▣ JPM: JPMorgan Chase & Co. (NYSE ticker)
- ▣ LQD: iShares iBoxx Investment Grade Corporate Bond ETF
- ▣ MS: Morgan Stanley (NYSE ticker)
- ▣ MSFT: Microsoft Corporation (Nasdaq ticker)
- ▣ MSTR: Strategy Inc., formerly MicroStrategy (Nasdaq ticker)

Acronyms used in this material (2/2)

- MV: Mean–Variance tangency portfolio estimated from sample moments
- MV-LO: Mean–Variance tangency portfolio, Long-Only
- NVDA: NVIDIA Corporation (Nasdaq ticker)
- OLS: Ordinary Least Squares
- OOS: Out-Of-Sample
- SE: Standard Error
- US: United States
- WFC: Wells Fargo & Company (NYSE ticker)

Today's Questions and Route

- ▣ Questions of the seminar
 - ▶ can an optimised portfolio beat equal weights ($1/N$) out of sample?
 - ▶ how would we know that it does?
- ▣ Route
 - ▶ what you need for today: four primer slides (the lecture comes after this seminar)
 - ▶ A1, A3, A4: two-asset GMV and risk parity on paper
 - ▶ B1: the rolling backtest of four rules on nine sectors, then A7–A8: test the Sharpe differences it produces; B8: the break-even cost
 - ▶ C2: find the errors in an AI answer
 - ▶ after the lecture: B2 (does shrinkage lower risk?), B3 (ERC and HRP step by step); C1 is the take-home project
- ▣ What you produce: the A answers, one table and one figure per B exercise, the corrected C2 code, one interpretation sentence per exercise
- ▣ Notebook for today: [open the seminar notebook in Colab](#); each exercise has its own notebook section
- ▣ Seminar: Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Model
A1	where is the minimum-variance mix of SPY and TLT?	Solved	–
A3, A4	who carries the risk: ERC versus 60/40	Solved, Proposed	A4: A3
B1	do Markowitz rules beat 1/N on sectors?	Solved	–
A7, A8	is a Sharpe difference significant?	Solved, Proposed	A8: A7
B8	at which cost does MV-LO stop paying?	Proposed	B1
C2	what is wrong in an AI backtest?	Proposed	B1, A7
After the lecture	B2 [Solved], B3 [Solved]: covariance shrinkage, ERC and HRP	homework	B1
C1	a BVB portfolio in RON (take-home project)	Proposed	B1, B3
Extension packs	A2, A4 Extended, A5, A6, A9, B2 Extended, B4–B7, B9, C3 (optional)	Solved or Proposed	marked on each task

- **[Solved]**: full solution in the slides and in the notebook, a model to copy; **[Proposed]**: you solve it, the solution is discussed in class
- Part A: derivations on paper; Part B: backtests and inference on real data, each ending with an interpretation question; Part C: open questions and AI critique

Data Used

Universe	Assets (price used)	Currency	Monthly returns	Months
Sectors	XLB, XLE, XLF, XLI, XLK, XLP, XLU, XLV, XLY (adjusted close)	USD	Jan 1999 – Jul 2026	331
Multi-asset	SPY, EFA, EEM, TLT, IEF, LQD, HYG, GLD (adjusted close)	USD	May 2007 – Jul 2026	231
Combined	9 sectors + the 7 multi-asset ETFs other than SPY	USD	May 2007 – Jul 2026	231
Bitcoin	Bitcoin (BTC/USD, close)	USD	Sep 2010 – Jul 2026	B5 only
BVB	TLV, SNP, BRD, TGN, SNG, SNN, EL, TEL; the BET-TR index (close)	RON	Oct 2014 – Aug 2026	C1 only

- Daily prices of the course data, one series per asset (field used: close / adjusted close)
- Align the *prices* of a universe on common trading days, keep the last price of each month, then take simple returns in decimals ($0.01 = 1\%$)
- Risk-free rate: one-month Treasury bill, [Kenneth French Data Library](#), monthly decimals aligned on the month-end date; excess return $r_t^e = r_t - r_{f,t}$; USD/RON: [BNR](#) reference rate (C1)

Setup: Run and Verify

```
# notebook section "Setup" defines the helpers
r, rex, rf = us_monthly(SECTORS)
# r: monthly simple returns (T x 9)
# rex = r - rf: excess returns; rf: T-bill
print(r.shape, r.index[0], r.index[-1])
print(r.iloc[0][['XLB', 'XLK', 'XLU']], rf.iloc[0])
ret, to, W = backtest(r, rex)
# ret: OOS excess returns of each rule
# to: monthly turnover; W: weights per rule
```

- Expected output: 331 months, first return 1999-01-31, last 2026-07-31
- First row (January 1999): XLB -3.70%, XLK 15.90%, XLU -2.48%; T-bill 0.35% a month
- backtest uses months $t - 60, \dots, t - 1$ for the weights of month t : 271 out-of-sample months, 2004-01 to 2026-07
- Other helpers: `sr_diff_hac(r1, r2)` returns the annualised difference, its HAC standard error and p ; `sr_diff_boot(r1, r2, block)` the studentized bootstrap test; `lw_cc(X)` the shrunk covariance and δ
- If a count or a first value differs, stop and check the symbols, the price type (adjusted or close) and the date alignment

Notation Used in Every Exercise

- ▣ w : weights, $w^\top \mathbf{1} = 1$; $\boldsymbol{\mu}$, $\boldsymbol{\Sigma}$: mean vector and covariance matrix of excess returns; $\sigma_p^2 = w^\top \boldsymbol{\Sigma} w$
- ▣ GMV $w = \boldsymbol{\Sigma}^{-1} \mathbf{1} / \mathbf{1}^\top \boldsymbol{\Sigma}^{-1} \mathbf{1}$; tangency (MV) $w = \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} / \mathbf{1}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}$; MV-LO: tangency with $0 \leq w_i \leq 1$
- ▣ Risk contribution $RC_i = w_i (\boldsymbol{\Sigma} w)_i / \sigma_p$; risk share RC_i / σ_p ; ERC (Maillard, Roncalli & Teïletche (2010)): equal RC_i ; HRP (López de Prado (2016))
- ▣ Sharpe ratio $SR = \bar{r}^e / s(r^e)$, monthly, annualised $\times \sqrt{12}$; $\hat{\Delta} = \widehat{SR}_1 - \widehat{SR}_2$
- ▣ Weights after drift (moved by returns) $w_{i,t}^+ = w_{i,t} (1 + R_{i,t}) / \sum_j w_{j,t} (1 + R_{j,t})$; turnover $TO_{t+1} = \sum_i |w_{i,t+1} - w_{i,t}^+|$
- ▣ Net return $r_t^{\text{net}} = r_t - c TO_t$, c the cost per unit traded; 1 basis point (bp) = 10^{-4} , so 20 bp is $c = 0.002$
- ▣ Worked example: $\sigma_1 = 20\%$, $\sigma_2 = 10\%$, $\rho = 0$, $w = (0.5, 0.5)$: $\sigma_p^2 = 0.25(0.04) + 0.25(0.01) = 0.0125$, $\sigma_p = 11.2\%$

What You Need for Today (1/4): GMV, Tangency and 1/N

- This seminar comes **before** the lecture: these four slides give only what the in-class tasks use; depth in [Chapter 4](#), the lecture after this seminar
- Portfolio variance $\sigma_p^2 = \mathbf{w}^\top \Sigma \mathbf{w}$ [Markowitz \(1952\)](#); four rules for the weights, all with $\mathbf{w}^\top \mathbf{1} = 1$
 - ▶ **GMV** (global minimum variance): $\min_{\mathbf{w}} \mathbf{w}^\top \Sigma \mathbf{w}$, solution $\Sigma^{-1} \mathbf{1} / (\mathbf{1}^\top \Sigma^{-1} \mathbf{1})$; needs no expected returns
 - ▶ **MV**: the tangency portfolio of [Chapter 3](#), $\Sigma^{-1} \boldsymbol{\mu} / (\mathbf{1}^\top \Sigma^{-1} \boldsymbol{\mu})$; **MV-LO**: the same maximum Sharpe ratio with $0 \leq w_i \leq 1$, solved numerically
 - ▶ **1/N**: equal weights, nothing estimated [DeMiguel, Garlappi & Uppal \(2009\)](#)
- Example: two uncorrelated assets, $\sigma_1 = 20\%$, $\sigma_2 = 10\%$: $\Sigma^{-1} \mathbf{1} = (25, 100)$, GMV $\mathbf{w} = (0.2, 0.8)$, $\sigma_{\text{GMV}} = 1/\sqrt{125} = 8.94\%$, below both volatilities
- Why 1/N is a hard benchmark: $\boldsymbol{\mu}$ is estimated with large errors, and Σ^{-1} amplifies them into extreme weights

What You Need for Today (2/4): Risk Contributions and ERC

- Euler decomposition: $\sigma_p = \sum_i RC_i$ with $RC_i = w_i(\Sigma w)_i / \sigma_p$, the contribution of asset i to the volatility; the risk shares RC_i / σ_p add up to 100%
- Example ($\sigma_1 = 20\%$, $\sigma_2 = 10\%$, $\rho = 0$, $w = (0.5, 0.5)$): $\Sigma w = (0.020, 0.005)$, $\sigma_p^2 = 0.0125$
 - ▶ shares $0.5 \times 0.020 / 0.0125 = 80\%$ and $0.5 \times 0.005 / 0.0125 = 20\%$: half of the money, four fifths of the risk
- **ERC** (equal risk contribution, risk parity) [Maillard, Roncalli & Teïletche \(2010\)](#): choose w so that all RC_i are equal; with $\rho = 0$ and two assets, $w_i \propto 1/\sigma_i$, here $(1/3, 2/3)$
- Capital weights and risk weights are different things: a “balanced” portfolio in money can be concentrated in risk (A4)

What You Need for Today (3/4): An Honest Backtest

- ▣ **Rolling window**: the weights held in month t use only months $t - 60, \dots, t - 1$; any estimate that uses month t is look-ahead
- ▣ **Drift**: during the month the weights move with returns, $w_i^+ = w_i(1 + R_i) / \sum_j w_j(1 + R_j)$; **turnover** is the trade needed to return to the new target, $TO = \sum_i |w_{i,\text{new}} - w_i^+|$
 - ▶ example, $1/N$ with returns $+10\%$ and -10% : $w^+ = (0.55, 0.45)$, $TO = 0.05 + 0.05 = 0.10$; at 20 bp per unit traded: $0.002 \times 0.10 = 2$ bp of wealth
- ▣ Net return $r_t - c TO_t$; **maximum drawdown**: the largest fall of wealth from a previous peak
- ▣ Sharpe ratio of monthly excess returns, annualised by $\sqrt{12}$ (not $\sqrt{252}$, which is for daily data)

What You Need for Today (4/4): Is a Sharpe Difference Significant?

- ▣ $\hat{\Delta} = \widehat{SR}_1 - \widehat{SR}_2$ on the same months; the two estimates are correlated, so the test uses the pair
- ▣ **Delta method:** if $\sqrt{T}(\hat{\theta} - \theta) \rightarrow N(0, \Psi)$, then $\sqrt{T}(f(\hat{\theta}) - f(\theta)) \rightarrow N(0, \nabla f^\top \Psi \nabla f)$ for a smooth f
- ▣ i.i.d. returns with the Normal distribution [Jobson & Korkie \(1981\)](#), corrected in [Ledoit & Wolf \(2008\)](#):
 $T \text{Var}(\hat{\Delta}) \rightarrow 2 - 2\rho + \frac{1}{2}(SR_1^2 + SR_2^2 - 2SR_1SR_2\rho^2)$
 - ▶ example: $SR_1 = SR_2 = 0.2$ a month, $\rho = 0.5$, $T = 240$: bracket 1.03, SE $\sqrt{1.03/240} = 0.066$ a month, 0.23 a year: gaps below about 0.45 a year are not detectable in 20 years
- ▣ Robust versions [Ledoit & Wolf \(2008\)](#): HAC standard errors (heavy tails, autocorrelation) and a studentized block bootstrap; helpers `sr_diff_hac`, `sr_diff_boot`
- ▣ After the lecture: covariance shrinkage (B2) and hierarchical risk parity (B3)

A1 [Solved]: The GMV Portfolio — Task

▣ Questions

- ▶ which mix of SPY and TLT has the lowest variance?
- ▶ what does it earn?

▣ **Assumptions:** N risky assets, known positive definite Σ , short sales allowed, $w^\top \mathbf{1} = 1$

▣ **Inputs** (monthly total returns August 2002 – July 2026, $T = 288$, annualised; the numbers below suffice)


	mean	volatility	variance
SPY	11.73%	14.72%	0.02166
TLT	4.32%	13.49%	0.01820
correlation $\rho = -0.100$, covariance $\sigma_{12} = -0.00199$			

1. Write the Lagrangian of $\min_w w^\top \Sigma w$ s.t. $w^\top \mathbf{1} = 1$ and derive $w_{\text{GMV}} = \Sigma^{-1} \mathbf{1} / A$, $A = \mathbf{1}^\top \Sigma^{-1} \mathbf{1}$.
2. Show that the minimum variance is $1/A$.
3. Invert the 2×2 matrix and obtain $w^* = (\sigma_2^2 - \sigma_{12}) / (\sigma_1^2 + \sigma_2^2 - 2\sigma_{12})$.
4. Compute A , w^* , σ_{GMV} and the mean of the mix; **report** the four numbers and one sentence on why σ_{GMV} is below both volatilities.

A1 [Solved]: Solution (1/2) — Derivation

- ▣ Step 1: $\mathcal{L} = \mathbf{w}^\top \Sigma \mathbf{w} - \lambda(\mathbf{w}^\top \mathbf{1} - 1)$; first-order condition $2\Sigma \mathbf{w} - \lambda \mathbf{1} = 0$, so $\mathbf{w} = \frac{\lambda}{2} \Sigma^{-1} \mathbf{1}$
- ▣ Imposing $\mathbf{1}^\top \mathbf{w} = 1$: $\frac{\lambda}{2} A = 1$, hence $\mathbf{w}_{\text{GMV}} = \Sigma^{-1} \mathbf{1} / A$; the objective is strictly convex, so the minimum is unique
- ▣ Step 2: $\sigma_{\text{GMV}}^2 = \mathbf{1}^\top \Sigma^{-1} \Sigma \Sigma^{-1} \mathbf{1} / A^2 = A / A^2 = 1 / A$
- ▣ Step 3: $\Sigma^{-1} = \frac{1}{\det \Sigma} \begin{pmatrix} \sigma_2^2 & -\sigma_{12} \\ -\sigma_{12} & \sigma_1^2 \end{pmatrix}$, so $\Sigma^{-1} \mathbf{1} = \frac{1}{\det \Sigma} (\sigma_2^2 - \sigma_{12}, \sigma_1^2 - \sigma_{12})^\top$
- ▣ Dividing the first element by the sum of both: $w^* = (\sigma_2^2 - \sigma_{12}) / (\sigma_1^2 + \sigma_2^2 - 2\sigma_{12})$; $\det \Sigma$ cancels

A1 [Solved]: Solution (2/2) — Numbers and the Minimum



Step 4, numbers:

- ▶ $\det \Sigma = 0.02166 \times 0.01820 - 0.00199^2 = 0.000390$; $\Sigma^{-1} \mathbf{1} = (0.02019, 0.02366)/0.000390 = (51.7, 60.6)$
- ▶ $A = 112.4$; $w^* = 51.7/112.4 = 0.461$ in SPY, 0.539 in TLT; $\sigma_{\text{GMV}} = 1/\sqrt{A} = 9.43\%$; mean $0.461 \times 11.73\% + 0.539 \times 4.32\% = 7.73\%$

- Interpretation: with $\rho < 0$ the two risks partly cancel; the mix is 36% less volatile than SPY and keeps two-thirds of its mean; at $\rho = +0.5$ the minimum is 12.2%

A3 [Solved]: ERC with Two Assets — Task

- ▣ **Question:** which SPY–IEF mix gives both assets the same risk contribution?
- ▣ **Inputs:** SPY and IEF (7–10 year Treasuries), monthly August 2002 – July 2026, annualised


$$\Sigma = \begin{pmatrix} 0.02166 & -0.00098 \\ -0.00098 & 0.00441 \end{pmatrix}, \quad \sigma_1 = 14.72\%, \quad \sigma_2 = 6.64\%, \quad \rho = -0.101$$

- ▣ Three different quantities: $RC_i\sigma_p = w_i(\Sigma w)_i$ (contribution to the variance), RC_i (to the volatility), RC_i/σ_p (the share, in %)
- ▣ **Tasks**
 1. Write $RC_1\sigma_p$ and $RC_2\sigma_p$ explicitly for two assets, long-only weights $w_1 + w_2 = 1$.
 2. Prove that $RC_1 = RC_2$ if and only if $w_1\sigma_1 = w_2\sigma_2$, for any $\rho \in (-1, 1)$.
 3. Compute the ERC weight of SPY, then σ_{ERC} , then the two risk shares.
- ▣ **Report:** the proof, w_1 , σ_{ERC} and the two shares

A3 [Solved]: Solution (1/2) — Proof and Numbers

- (1) $RC_1\sigma_p = w_1^2\sigma_1^2 + w_1w_2\sigma_{12}$, $RC_2\sigma_p = w_2^2\sigma_2^2 + w_1w_2\sigma_{12}$
- (2) The shaded cross terms are identical and cancel: $RC_1 = RC_2 \iff w_1^2\sigma_1^2 = w_2^2\sigma_2^2 \iff w_1\sigma_1 = w_2\sigma_2$; ρ drops out of the weights, not out of σ_p
- (3) $w_1 = \sigma_1^{-1}/(\sigma_1^{-1} + \sigma_2^{-1}) = 6.79/(6.79 + 15.06) = 0.311$;
 $\sigma_p^2 = 0.311^2(0.02166) + 0.689^2(0.00441) + 2(0.311)(0.689)(-0.00098) = 0.00377$, $\sigma_{ERC} = 6.14\%$; shares 50%, 50%
- Interpretation: equal risk needs more than twice as much of the bond ETF, which is less than half as volatile

A3 [Solved]: Solution (2/2) — Weights, Risk Shares and Correlation



- Left: capital weights against shares of variance; right: the ERC weight of SPY is the same for every ρ , the ERC volatility is not (6.14% at the actual ρ)

A4 [Proposed]: Who Carries the Risk in 60/40? — Task

- ▣ **Model:** A3 [Solved]
- ▣ **Question:** a “balanced” 60/40 portfolio holds 40% in bonds; how much of its risk do the bonds carry?
- ▣ **Inputs:** Σ of A3; $w = (0.6, 0.4)$ in SPY and IEF
- ▣ **Tasks**
 1. Compute Σw row by row: $(\Sigma w)_1 = 0.6\sigma_1^2 + 0.4\sigma_{12}$ and $(\Sigma w)_2 = 0.6\sigma_{12} + 0.4\sigma_2^2$.
 2. Compute $\sigma_p^2 = w^\top (\Sigma w)$, σ_p and the shares $w_i(\Sigma w)_i / \sigma_p^2$.
 3. Compare σ_p with the ERC volatility of A3 and draw the capital and risk shares of both portfolios side by side.
- ▣ **Interpretation:** is 60/40 a balanced portfolio in risk terms?

B1 [Solved]: 1/N versus Markowitz on Sectors — Task


- ▣ **Question:** do estimated Markowitz rules beat 1/N out of sample on nine sector ETFs?
- ▣ **Data:** the sector panel of the setup; monthly excess returns January 1999 – July 2026
- ▣ **Tasks**
 1. Each month t estimate 1/N, MV (tangency), MV-LO (maximum Sharpe with $0 \leq w_i \leq 1$, solved numerically) and GMV on months $t - 60, \dots, t - 1$ only; hold them in month t : January 2004 – July 2026 ($T = 271$).
 2. Compute the first month step by step: the window, the weights, the January 2004 return, the drifted weights and the February 2004 turnover.
 3. Compute the annualised mean, volatility, Sharpe ratio, mean turnover and maximum drawdown of each rule.
 4. Test $H_0: SR_{\text{rule}} = SR_{1/N}$ with `sr_diff_hac` (Ledoit & Wolf (2008), Section 3.1; Bartlett kernel, 4 lags) and `sr_diff_boot` (studentized circular block bootstrap, 4999 draws; calibrated block lengths (given): MV 6, MV-LO 1, GMV 1), $\alpha = 5\%$.
- ▣ **Report:** the table below and a wealth chart

	Sharpe	Volatility	Turnover	Max. drawdown	$\hat{\Delta}$	HAC p / bootstrap p
1/N, MV, MV-LO, GMV	...	...	...	...	...	...

- ▣ **Interpretation:** MV-LO beats 1/N by a few hundredths: is that evidence of skill?

B1 [Solved]: Solution (1/4) — The First Out-of-Sample Month

Rolling 60-month estimation, one-month holding: 271 out-of-sample months




- Window 1999-01 – 2003-12 (60 months) gives the weights held in 2004-01; MV-LO holds XLB 81.8%, XLE 14.3%, XLF 3.9%, the others 0; MV has gross exposure 8.3
- January 2004 excess returns: $1/N \sum_i w_i r_i^e = 1.11\%$, MV-LO -3.24% , GMV 1.36% , MV 1.12% (T-bill 0.07%)
- Drift: XLB returned -4.4% , XLK 4.3% ; by the end of January the $1/N$ weights of 11.11% become 10.50% (XLB) and 11.45% (XLK)
- February 2004 turnover $\sum_i |w_{i,\text{Feb}} - w_{i,\text{Jan}}^+|$: $1/N$ 0.0191 , MV-LO 0.748 (the new window also moves the MV-LO weights)

B1 [Solved]: Solution (2/4) — Performance and Tests

	1/N	MV	MV-LO	GMV
Sharpe	0.66	0.26	0.72	0.71
Volatility	14.3%	69.0%	14.7%	11.9%
Turnover	0.02	3.86	0.23	0.23
Max. drawdown	-49%	-100%	-42%	-33%
$\hat{\Delta}$ vs 1/N	-	-0.40	+0.06	+0.05
HAC p	-	0.22	0.60	0.75
Bootstrap p	-	0.30	0.63	0.76


- MV: median gross exposure 4.9, turnover almost four times wealth a month; a monthly loss above 100% in 2009-03 ends its investable wealth: its Sharpe describes a hypothetical return series
- The other three rules land within 0.06 of 1/N; no p -value is below 0.05

B1 [Solved]: Solution (3/4) — Wealth, Drawdown, Exposure and Turnover



- Top: wealth of 1 USD in total returns and its drawdown; bottom: gross exposure $\sum_i |w_i|$ and turnover, both on log scales; MV-LO and 1/N have gross exposure 1 by construction

B1 [Solved]: Solution (4/4) — Inference



- MV-LO: $\hat{\Delta} = 0.063$, HAC standard error 0.120, HAC $p = 0.60$; studentized bootstrap (block 1) $p = 0.63$, 95% interval $[-0.19, 0.31]$ covers zero: do not reject
- Extrapolation (right): if the true gap equalled the estimate and the standard error fell like $1/\sqrt{\text{years}}$, $|z| = 1.96$ would need 311 years of data; a conditional significance calculation, not a power calculation
- Interpretation: the interval admits losses and gains up to 0.31; failure to reject does not establish equality

A7 [Solved]: The Variance of a Sharpe Difference — Task

- **Question:** is the GMV-1/N gap of B1 larger than sampling noise, under the simplest assumptions?
- **Inputs:** the B1 monthly **excess** returns of GMV and 1/N, $T = 271$; monthly $SR_{\text{GMV}} = 0.2051$, $SR_{1/N} = 0.1908$, correlation $\rho = 0.694$
- **Assumption:** i.i.d. bivariate Normal returns; Φ is the standard Normal cumulative distribution function
- **Tasks**
 1. Write the 4×4 asymptotic covariance matrix of $\sqrt{T}(\hat{\mu}_1 - \mu_1, \hat{\mu}_2 - \mu_2, \hat{\sigma}_1^2 - \sigma_1^2, \hat{\sigma}_2^2 - \sigma_2^2)$.
 2. Write the gradient of $f = \mu_1/\sigma_1 - \mu_2/\sigma_2$ and apply the delta method to obtain

$$T \text{Var}(\hat{\Delta}) \rightarrow 2 - 2\rho + \frac{1}{2}(SR_1^2 + SR_2^2 - 2SR_1SR_2\rho^2)$$
 3. Test $H_0: SR_{\text{GMV}} = SR_{1/N}$ with a two-sided z-test at 5%; compare with the HAC p -value 0.75.
- **Report:** the matrix, the formula, z , p and the decision

A7 [Solved]: Solution (1/2) — The Delta Method

- (1) Under Normality, sample means and variances are asymptotically independent, $\sigma_{12} = \rho\sigma_1\sigma_2$:

$$\Psi = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & 0 & 0 \\ \sigma_{12} & \sigma_2^2 & 0 & 0 \\ 0 & 0 & 2\sigma_1^4 & 2\sigma_{12}^2 \\ 0 & 0 & 2\sigma_{12}^2 & 2\sigma_2^4 \end{pmatrix}, \quad \nabla f = \begin{pmatrix} 1/\sigma_1 \\ -1/\sigma_2 \\ -SR_1/(2\sigma_1^2) \\ +SR_2/(2\sigma_2^2) \end{pmatrix}$$

- (2) $\nabla f^\top \Psi \nabla f$ splits into two blocks


► Means: $\sigma_1^2/\sigma_1^2 + \sigma_2^2/\sigma_2^2 - 2\sigma_{12}/(\sigma_1\sigma_2) = 2 - 2\rho$

► Variances: $\frac{SR_1^2}{4\sigma_1^4} 2\sigma_1^4 + \frac{SR_2^2}{4\sigma_2^4} 2\sigma_2^4 - 2 \frac{SR_1 SR_2}{4\sigma_1^2 \sigma_2^2} 2\sigma_{12}^2 = \frac{1}{2}(SR_1^2 + SR_2^2 - 2SR_1 SR_2 \rho^2)$

- Adding the blocks gives the formula (Jobson & Korkie (1981); corrected form in Ledoit & Wolf (2008))

A7 [Solved]: Solution (2/2) — The Test

- (3) Bracket: $2 - 2(0.694) + 0.5(0.2051^2 + 0.1908^2 - 2 \times 0.2051 \times 0.1908 \times 0.694^2)$
 $= 0.612 + 0.5(0.0421 + 0.0364 - 0.0377) = 0.6324$
- Standard error $\sqrt{0.6324/271} = 0.0483$;
 $z = (0.2051 - 0.1908)/0.0483 = 0.30$;
 $p = 2[1 - \Phi(0.30)] = 0.77$
- Decision: do not reject H_0 at 5%; the HAC test agrees ($p = 0.75$)
- Interpretation: the annualised gap, 0.71 against 0.66, is well inside sampling noise; 22 years cannot rank GMV and 1/N



A8 [Proposed]: ERC versus 1/N — Task

- ▣ **Model:** A7 [Solved]
- ▣ **Question:** ERC beats 1/N by less than GMV does; can the smaller gap be closer to significance?
- ▣ **Inputs:** B1 backtest with ERC added, $T = 271$; monthly $SR_{\text{ERC}} = 0.2018$, $SR_{1/N} = 0.1908$, correlation 0.9927; keep four decimals until the end (the bracket is small when ρ is near 1)
- ▣ Reference: $T \text{Var}(\hat{\Delta}) \rightarrow 2 - 2\rho + \frac{1}{2}(SR_1^2 + SR_2^2 - 2SR_1SR_2\rho^2)$
- ▣ **Tasks**
 1. Compute the bracket, the standard error, z and the two-sided p ; decide at 5%.
 2. Show that the bracket tends to $\frac{1}{2}(SR_1 - SR_2)^2$ as $\rho \rightarrow 1$.
 3. Fill a comparison table (gap, ρ , standard error, z) for A7 and A8 and plot both 95% intervals.
- ▣ **Interpretation**
 - ▶ why is a smaller difference here closer to significance?
 - ▶ what happens if the two return series are identical?

B8 [Proposed]: The Break-Even Cost — Task

- ▣ **Model:** B1 [Solved]
- ▣ **Question:** MV-LO beats 1/N gross on 16 ETFs; at which trading cost does the edge vanish?
- ▣ **Data:** the combined universe (16 ETFs), OOS May 2012 – July 2026 ($T = 171$); strategy 1 = MV-LO, strategy 2 = 1/N; turnover as in B1
- ▣ Net return $r_t - c TO_t$, c per unit traded in decimals: one month with $TO = 0.25$ at 20 bp costs $0.002 \times 0.25 = 0.0005$, i.e. 5 bp of wealth
- ▣ **Tasks**
 1. Compute the gross Sharpe ratios, the HAC p -value of their difference and the mean turnovers.
 2. On a grid of 0–200 bp (step 0.1 bp), recompute the mean **and** the standard deviation of both net return series and find the cost c^* where the net Sharpe ratios cross; plot both curves.
 3. Derive the fixed-volatility approximation $c^* \approx \Delta SR / [12(TO_1/\sigma_1 - TO_2/\sigma_2)]$ (annual σ) and compare it with the grid value.
- ▣ **Interpretation:** retail costs on BVB or US ETFs: is MV-LO worth implementing?

C2 [Proposed]: Critique an AI Answer — Task

- ▣ **Model:** B1 and A7 [Solved]. **Prompt:** “prices holds daily adjusted closes of the nine sector ETFs of B1 (December 1998 – July 2026). Write Python code for an out-of-sample GMV backtest with monthly rebalancing and give me its annualised Sharpe ratio.”

```
1 import numpy as np
2 r = prices.resample('ME').last().pct_change().dropna() # monthly simple returns
3 S = r.cov().values # covariance matrix
4 w = np.linalg.solve(S, np.ones(len(S))) # Sigma^{-1} 1
5 w = w / w.sum() # GMV weights, sum to one
6 port = r.dot(w) # portfolio return
7 sr = port.mean() / port.std() * np.sqrt(252)
8 print(f"Annualised Sharpe ratio: {sr:.2f}")
```

- ▣ **Tasks** (rf: the monthly T-bill of the setup, decimals, aligned on month-end dates)
1. Find the three intended errors; name the line, explain why it is wrong and how you would detect it (a date check, a number from B1 or A7, a definition). Further qualified remarks are welcome.
 2. Write the corrected code and report its Sharpe ratio; check the date lag, missing values, weight sums and the annualisation.
 3. Log each error in the format of AI_ERRORS.md: request, answer, error, how found, correction.
- ▣ **Interpretation**
- ▶ how far is the AI's number from a defensible one?
 - ▶ which error matters most?


B2 [Solved]: Does Shrinkage Lower GMV Risk? — Task

- ▣ **Question:** does the Ledoit–Wolf estimator lower the out-of-sample risk of GMV on sectors?
- ▣ **Data:** the B1 backtest; GMV with the sample covariance versus GMV-LW, the constant-correlation target of [Ledoit & Wolf \(2004\)](#) (helper `lw_cc(X)`: input a 60×9 window of excess returns, output $\hat{\Sigma}_{\text{LW}}$ and δ)
- ▣ **Tasks**
 1. On the last window, follow the pipeline: $S \rightarrow \text{target } F (f_{ij} = \bar{r} \sqrt{s_{ii}s_{jj}}) \rightarrow \delta \rightarrow \hat{\Sigma}_{\text{LW}} = \delta F + (1 - \delta)S \rightarrow \text{GMV weights}$; report $\bar{r}$, δ , both condition numbers and the gross exposures.
 2. Compute the out-of-sample volatilities of GMV, GMV-LW and $1/N$.
 3. Compute $\text{Var}(\text{GMV-LW}) / \text{Var}(\text{GMV})$ with a 95% percentile interval from a circular block bootstrap ([Künsch \(1989\)](#); blocks of 6 months, 5000 draws); both series use the **same** resampled months.
 4. Compare the mean monthly turnover of the two rules.
- ▣ **Interpretation**
 - ▶ why is the gain small with 9 assets and 60 months?
 - ▶ when would it be larger?

B2 [Solved]: Solution (1/2) — From S to Weights

- ▣ Last window 2021-08 – 2026-07, $T = 60$, $N = 9$
- ▣ Step 1: average sample correlation $\bar{r} = 0.511$; e.g. XLK–XLU 0.31 is replaced by $\bar{r}$ in F, the variances are kept
- ▣ Step 2: $\delta = 0.24$: $\hat{\Sigma}_{LW}$ is 24% target, 76% sample
- ▣ Step 3: condition number $51.6 \rightarrow 32.5$: the smallest eigenvalues are pulled up
- ▣ Step 4: GMV weights; gross exposure $\sum_i |w_i|$ 2.64 (sample) $\rightarrow$ 1.91 (shrunk); XLP 66.8% $\rightarrow$ 57.8%, XLK 10.1% $\rightarrow$ 6.4%
- ▣ One paired resample: draw 46 random start months, take 6 consecutive months from each (wrapping around), cut at $T = 271$, and compute the variance of both strategies on those same months

B2 [Solved]: Solution (2/2) — Risk and Turnover



- Volatility: GMV 11.90%, GMV-LW 11.80%, 1/N 14.28%: both GMV rules cut risk by about a sixth
- Variance ratio 0.984, interval [0.918, 1.062] contains 1: no significant reduction; turnover 0.234 $\rightarrow$ 0.137 (−41%): the benefit here is stability, not risk
- Interpretation: with $N/T = 0.15$ the sample matrix is already decent; the gain grows with N/T and with a target close to the truth

B3 [Solved]: ERC and HRP on Sectors — Task

▣ Questions

- ▶ how do the two risk-based rules allocate?
- ▶ do they beat $1/N$?

▣ Part 1, computing the weights (last window August 2021 – July 2026, sample covariance)

1. Apply HRP step by step on the four-asset example of the next slide: distances, tree, order, bisection.
2. Compute the nine ERC weights ($w_{\text{erc}}(S)$, the convex problem of [Maillard, Roncalli & Teiletche \(2010\)](#)) and HRP weights ($w_{\text{hrp}}(S)$: $d_{ij} = \sqrt{(1 - \rho_{ij})/2}$, Euclidean distance between the columns of D , single linkage, recursive bisection of the ordered list).
3. Compute the risk shares of $1/N$, ERC and HRP.


▣ Part 2, backtest (January 2004 – July 2026, rolling 60 months, as in B1)

4. Test the Sharpe ratios of ERC and HRP against $1/N$ with the HAC test at 5%; report turnover and the correlation with $1/N$.


▣ Interpretation: ERC and $1/N$ have correlation 0.993: is ERC a different strategy or a small tilt?

B3 [Solved]: Solution (1/3) — HRP on Four Assets

- Volatilities 10%, 12%, 20%, 25%; correlations A–B 0.8, C–D 0.7, all cross pairs 0.2
- Step 1: $d_{AB} = \sqrt{(1 - 0.8)/2} = 0.316$, $d_{CD} = 0.387$, cross pairs 0.632; distances between columns of D: $\tilde{d}_{AB} = 0.447$, $\tilde{d}_{CD} = 0.548$; single linkage merges A–B, then C–D, then the two clusters
- Step 2: order A, B, C, D; split into $\{A, B\}$ and $\{C, D\}$
- Step 3: inverse-variance weights inside each half, (0.59, 0.41) and (0.61, 0.39); cluster variances $V_1 = 0.01055$, $V_2 = 0.04105$; $\alpha = 1 - V_1/(V_1 + V_2) = 0.796$ to $\{A, B\}$
- Then split each pair again by inverse variance: A 47.0%, B 32.6%, C 12.5%, D 8.0%



B3 [Solved]: Solution (2/3) — Nine Sectors, Last Window



- Tree order XLE, XLK, XLY, XLU, XLV, XLP, XLF, XLB, XLI; ERC: XLP 14.8%, XLV 13.9% ... XLB 8.6%; HRP: XLP 18.4%, XLU 17.1% ... XLB 4.5%; 1/N risk shares from 7.8% (XLP) to 13.8% (XLB)

B3 [Solved]: Solution (3/3) — Backtest

- OOS January 2004 – July 2026, rolling 60 months, HAC test against $1/N$
- Backtest: ERC Sharpe 0.70 vs 0.66, $p = 0.09$; HRP 0.70, $p = 0.32$; turnover $1/N$ 0.024, ERC 0.028, HRP 0.082; correlation with $1/N$: ERC 0.993, HRP 0.979
- Interpretation: ERC is a small, cheap tilt towards low-volatility sectors; HRP a larger tilt that trades three times as much; neither rejects H_0

After the Lecture: What and When

- These two tasks use tools that [Chapter 4](#) introduces: work on them after the lecture, before the next seminar
 - ▶ B2 [Solved]: Ledoit–Wolf shrinkage of the covariance matrix and the risk of GMV
 - ▶ B3 [Solved]: ERC and HRP (hierarchical risk parity) weights, step by step, and their backtest
- The take-home project C1 uses both: start it after these two tasks

C1 [Proposed]: A BVB Portfolio in RON — the Take-Home Project

- ▣ **Model:** B1, B3 and A7 [Solved]. **Question:** does diversification beyond the BVB pay for a Romanian investor?
- ▣ **Data and currency:** TLV, SNP, BRD, TGN, SNG, SNN, EL, TEL and the BET-TR index (October 2014 – August 2026); SPY, EFA, IEF, GLD converted to RON on common days, $P_{\text{RON}} = P_{\text{USD}} \times \text{USD/RON}$ (BNR); e.g. 2017-10-02: SPY $219.84 \text{ USD} \times 3.9086 = 859.25 \text{ RON}$
- ▣ **Base case:** monthly returns in RON, 36-month window, first portfolio October 2017, $R_f = 0$
- ▣ **Tasks**
 1. Backtest 1/N, GMV-LW, ERC and HRP on the eight stocks, then on the stocks plus the four foreign ETFs.
 2. Compare each with BET-TR (never BET): Sharpe ratio, volatility, maximum drawdown; plot wealth and drawdown.
 3. Test each global rule against BET-TR and against its BVB-only version (HAC), with a Holm adjustment over this family of eight tests.
 4. Sensitivity: a RON rate of 5% a year, 0.05/12 a month, subtracted from all returns.
- ▣ **Deliverables:** a Quantlet folder in the team repository, a table of the eight tests (raw and Holm p), one slide with the conclusion and its main limitation (currency risk, costs, taxes)

Extension Packs: What Each One Needs

Pack	Exercises	Prerequisites
1. Risk parity theory	A4 Extended	A3, A4
2. Black–Litterman	A5 [Solved], A6, B7	lecture, Black–Litterman
3. Bias and regression inference	A9, A2, B9	A1, A7, B1
4. Covariance estimation	B2 Extended [Solved], B4	B2
5. Beyond the sectors	B5, B6	B1, B3
6. Research projects	C3 (C1 is the base homework)	B1, A7, the lecture case study

- Every pack has its own notebook section; do them in order within a pack
- Further study:  Hierarchical Clustering,  Minimum Spanning Tree

A4 Extended [Proposed]: Euler and the Uniqueness of ERC — Task

▣ **Model:** A3 and A4

▣ **Tasks**

1. Show that $\sigma_p(cw) = c \sigma_p(w)$ for $c > 0$ and deduce $\sum_i RC_i = \sigma_p$ from Euler's theorem for functions homogeneous of degree one.
2. For Σ positive definite, let $f(y) = \frac{1}{2}y^\top \Sigma y - \frac{1}{N} \sum_i \ln y_i$ on $y > 0$. Compute its Hessian and show that f is strictly convex.
3. Show that $f \rightarrow +\infty$ as any $y_i \rightarrow 0$ and as $\|y\| \rightarrow \infty$; conclude that a minimiser exists and is unique.
4. Write the first-order condition and show that it equalises the risk contributions, before and after the normalisation $w = y/1^\top y$.


▣ **Report:** the two proofs in at most ten lines

A5 [Solved]: Black–Litterman with One Asset — Task

- ▣ **Question:** how much should a precise view move an uncertain prior mean?
- ▣ **Inputs:** one asset, return volatility $\sigma = 16\%$, $\tau = 0.05$; prior for the **expected** return $\mu \sim N(\pi, \tau\sigma^2)$ with $\pi = 5\%$; view $q = 8\%$ with standard deviation 2%
 - ▶ Both variances measure uncertainty about the mean μ , not the volatility of returns; Normality is assumed for the posterior density
- ▣ **Tasks**
 1. Compute the prior variance $\tau\sigma^2$ and the two precisions (inverse variances).
 2. Compute μ_{BL} as the precision-weighted mean, the weight of the view and the posterior standard deviation; plot the three densities.
 3. Generalise: stack $\pi = \mu + e_1$ and $q = P\mu + e_2$ as $y = Z\mu + e$ and derive the GLS estimator (Theil & Goldberger (1961)); show that $P = I$, $\Omega = \tau\Sigma$ gives $(\pi + q)/2$.

A5 [Solved]: Solution (1/2) — One Asset

- (1) Prior variance $0.05 \times 0.16^2 = 0.00128$, precision 781; view variance $0.02^2 = 0.0004$, precision 2500
- (2) $\mu_{BL} = (781 \times 5\% + 2500 \times 8\%) / 3281 = 7.29\%$; view weight $2500 / 3281 = 0.76$; posterior standard deviation $1 / \sqrt{3281} = 1.75\%$
- The view is three times as precise as the prior, so it gets three-quarters of the weight; the posterior is more precise than either source



A5 [Solved]: Solution (2/2) — Black–Litterman as GLS

- (3) Dimensions: μ is $N \times 1$, q is $K \times 1$, P is $K \times N$

$$\underbrace{\begin{pmatrix} \pi \\ q \end{pmatrix}}_{y \ ((N+K) \times 1)} = \underbrace{\begin{pmatrix} I_N \\ P \end{pmatrix}}_Z \mu + e, \quad V = \text{Var}(e) = \begin{pmatrix} \tau \Sigma & 0 \\ 0 & \Omega \end{pmatrix}$$

- GLS: $\hat{\mu} = (Z^\top V^{-1} Z)^{-1} Z^\top V^{-1} y$; one block multiplication: $Z^\top V^{-1} = ((\tau \Sigma)^{-1}, P^\top \Omega^{-1})$, so

$$\mu_{\text{BL}} = [(\tau \Sigma)^{-1} + P^\top \Omega^{-1} P]^{-1} [(\tau \Sigma)^{-1} \pi + P^\top \Omega^{-1} q]$$

- Covariance $[(\tau \Sigma)^{-1} + P^\top \Omega^{-1} P]^{-1}$: the Black–Litterman posterior of the lecture
- $P = I$, $\Omega = \tau \Sigma$: $[2(\tau \Sigma)^{-1}]^{-1} (\tau \Sigma)^{-1} (\pi + q) = (\pi + q)/2$: equal precisions, equal weights

A6 [Proposed]: The Sector View of the Lecture — Task

- ▣ **Model:** A5 [Solved]
- ▣ **Inputs:** nine sectors in the order XLB, XLE, XLF, XLI, XLK, XLP, XLU, XLV, XLY; monthly excess returns August 2016 – July 2026, Σ annualised; benchmark $1/N$, $\delta_{RA} = 2.5$, $\tau = 0.05$
- ▣ View XLK minus XLU = 3% a year: $P = (0, 0, 0, 0, +1, 0, -1, 0, 0)$; implied $P\pi = 2.22\%$; spread variance $P\Sigma P^\top = 0.0477$ (enough: the full Σ is only needed for a check)
- ▣ **Tasks**
 1. Write the prior distribution of $P\mu$ implied by $\mu \sim N(\pi, \tau\Sigma)$.
 2. With $\Omega = \tau P\Sigma P^\top$ (He & Litterman (2002)), show that the posterior view is the simple average of prior and view, and compute it.
 3. Show that $\Delta w = P^\top(q - P\pi)/(2\delta_{RA}P\Sigma P^\top)$; compute the scalar, the new XLK and XLU weights, and check that the tilt sums to zero.
- ▣ **Interpretation:** how does the tilt change if the confidence in the view halves?

B7 [Proposed]: Black–Litterman on a Multi-Asset Set — Task

- ▣ **Model:** A5 and A6
- ▣ **Inputs:** SPY, EFA, EEM, TLT, IEF, LQD, HYG, GLD (this order), monthly excess returns August 2016 – July 2026 ($T = 120$), Σ annualised; benchmark $1/N$, $\delta_{RA} = 2.5$, $\tau = 0.05$; view GLD minus TLT = 4% a year, $P = (0, 0, 0, -1, 0, 0, 0, +1)$, $\Omega = \tau P \Sigma P^T$
- ▣ Unconstrained static allocation: a negative weight is a short position
- ▣ **Tasks**
 1. Compute $\pi = \delta_{RA} \Sigma w_b$ and the implied GLD minus TLT.
 2. Compute the spread variance, the surprise $q - P\pi$, the scalar tilt of A6 and the new weights; plot them against the benchmark.
 3. Compare with the sample mean of GLD minus TLT, with a ten times lower confidence ($\Omega \times 10$) and with a long-only version.
- ▣ **Interpretation:** the weights move far although the view is modest: why?

A9 [Proposed]: The In-Sample Sharpe Bias — Task

- ▣ **Model:** A7 [Solved] (for the moments), lecture on estimation risk
- ▣ **Assumptions:** i.i.d. $N(\mu, \Sigma)$ returns; maximum-likelihood $\hat{\mu}$, $\hat{\Sigma}$ (divisor T); given: $\hat{\mu} \sim N(\mu, \Sigma/T)$, $\mathbb{E}[\hat{\Sigma}^{-1}] = \frac{T}{T-N-2}\Sigma^{-1}$ for $T > N + 2$, $\hat{\mu}$ independent of $\hat{\Sigma}$, and $\mathbb{E}[x^\top A x] = \text{tr}(A \text{Var}(x)) + \mathbb{E}[x]^\top A \mathbb{E}[x]$
- ▣ **Tasks**
 1. Compute $\mathbb{E}[\hat{\theta}^2 \mid \hat{\mu}]$ for $\hat{\theta}^2 = \hat{\mu}^\top \hat{\Sigma}^{-1} \hat{\mu}$.
 2. Take the expectation over $\hat{\mu}$ and conclude $\mathbb{E}[\hat{\theta}^2] = (T\theta^2 + N)/(T - N - 2)$ (Kan & Zhou (2007)).
 3. Evaluate for $N = 9$, $T = 60$, $\theta = 0.66/\sqrt{12}$ monthly; report $\sqrt{12 \mathbb{E}[\hat{\theta}^2]}$, the annualised root mean square of the estimate (not its mean), and compare with a simulation.
 4. Give the unbiased estimator of θ^2 implied by the formula.

A2 [Proposed]: Tangency Weights by Regression — Task

- ▣ **Model:** A1 [Solved] (matrix algebra); lecture Appendix on [Britten-Jones \(1999\)](#)
- ▣ **Setting:** X is the $T \times N$ matrix of excess returns, rows $r_t^\top$; 1_T the response (a column of ones), 1_N the normalisation vector; $\hat{\mu} = X^\top 1_T / T$, $\hat{\Sigma} = X^\top X / T - \hat{\mu} \hat{\mu}^\top$, invertible
- ▣ **Tasks**
 1. Show $X^\top X / T = \hat{\Sigma} + \hat{\mu} \hat{\mu}^\top$ and write the OLS estimator of $1 = b^\top r_t + u_t$ (no constant) in terms of $\hat{\mu}$, $\hat{\Sigma}$.
 2. Apply Sherman–Morrison, $(A + uu^\top)^{-1} = A^{-1} - A^{-1}uu^\top A^{-1} / (1 + u^\top A^{-1}u)$, to show $\hat{b} = \hat{\Sigma}^{-1} \hat{\mu} / (1 + \hat{\theta}^2)$.
 3. Show that $\hat{b} / 1_N^\top \hat{b}$ are the sample tangency weights when $1_N^\top \hat{\Sigma}^{-1} \hat{\mu} > 0$.
 4. Optional: for $N = 3$ write the restriction matrix of “tangency = $1/N$ ”, then generalise.

B9 [Proposed]: Is the Tangency Portfolio 1/N? — Task


- ▣ **Model:** A2 and B1
- ▣ **Data:** nine sectors, monthly excess returns January 1999 – July 2026 ($T = 331$) and the halves January 1999 – September 2012, October 2012 – July 2026
- ▣ **Tests of H_0 :** $Rb = 0$, $R = [I_8, -1_8]$, 8 restrictions
 - ▶ Exact: $F = (R\hat{b})^\top [R(X^\top X)^{-1}R^\top]^{-1}(R\hat{b})/(8s^2) \sim F(8, T - 9)$, $s^2 = \hat{u}^\top \hat{u}/(T - 9)$
 - ▶ HAC Wald: $W = (R\hat{b})^\top (R\hat{V}R^\top)^{-1}R\hat{b} \rightarrow \chi_8^2$, $\hat{V} = T(X^\top X)^{-1}\hat{\Omega}(X^\top X)^{-1}$, $\hat{\Omega}$ Newey–West of $r_t \hat{u}_t$, Bartlett, $\lfloor 4(T/100)^{2/9} \rfloor$ lags
- ▣ **Tasks**
 1. Fill a 3×2 table of statistics and p -values (three samples, two tests); plot the coefficients with OLS and HAC intervals.
 2. Simulate 2000 samples of $T = 166$ from a Normal model with $\mu \propto \Sigma 1$ (so H_0 holds, Σ of the second half) and report the rejection rate of each test at 5%.
- ▣ **Interpretation**
 - ▶ when the tests disagree, which do you trust?
 - ▶ does non-rejection mean the tangency portfolio is 1/N?

B2 Extended [Solved]: Eigenvalues of the Last Window — Task

- ▣ **Question:** how many eigenvalues of the sector correlation matrix are distinguishable from noise?
- ▣ **Data:** last window of B2, August 2021 – July 2026, $N = 9$, $T = 60$
- ▣ **Tasks**
 1. Compute the nine eigenvalues of the sample correlation matrix.
 2. Compute the Marchenko–Pastur band $[(1 - \sqrt{c})^2, (1 + \sqrt{c})^2]$, $c = N/T$, and count the eigenvalues above and below it.
 3. Repeat with the noise variance rescaled to $\sigma^2 = 1 - \lambda_1/N$ (Laloux et al. (1999)).
- ▣ **Interpretation:** why does GMV care most about the eigenvalues *below* the band?

B2 Extended [Solved]: Solution

- (1)
5.25, 1.35, 0.84, 0.50, 0.32, 0.29, 0.21, 0.14, 0.11
- (2) $c = 0.15$, band $[0.38, 1.92]$: 1 above (the market mode), 5 below
- (3) $\sigma^2 = 0.42$, band $[0.16, 0.80]$: 3 above, 2 below
- Interpretation: GMV loads on the smallest eigenvalues, the near-riskless combinations; if they are underestimated, so is GMV risk; shrinkage pulls them up



B4 [Proposed]: Shrinkage on 16 ETFs — Task

- ▣ **Model:** B2 [Solved]
- ▣ **Data:** combined universe (nine sectors plus EFA, EEM, TLT, IEF, LQD, HYG, GLD), common months May 2007 – July 2026; 60-month window, OOS May 2012 – July 2026 ($T = 171$)
- ▣ **Tasks** (output: one table with rows GMV, GMV-LW, GMV-NL, 1/N and columns volatility, turnover, variance ratio with interval)
 1. Out-of-sample volatility of GMV, GMV-LW and 1/N; the variance ratio GMV-LW/GMV with the paired block bootstrap of B2; mean turnovers.
 2. Mean δ and the condition numbers of the sample and shrunk matrices on the last window.
 3. Extension: nonlinear shrinkage ([Ledoit & Wolf \(2020\)](#)); helper `nl_shrink(X)`, the GMV-NL backtest and its variance ratios to GMV and GMV-LW.
- ▣ **Interpretation:** is constant correlation a sensible target for stocks and Treasuries together? Distinguish a significant change from a point estimate

B5 [Proposed]: Risk Parity with Bitcoin — Task

- ▣ **Model:** B3 [Solved]; literature: [Petukhina, Trimborn, Härdle & Elendner \(2021\)](#), [Pele et al. \(2023\)](#)
- ▣ **Data:** the eight multi-asset ETFs with and without Bitcoin (BTC/USD), prices aligned on common trading days (weekends dropped), month-end returns from August 2010; both universes on the **same** months
- ▣ **Tasks**
 1. Descriptive, August 2010 – July 2026: Bitcoin's volatility and its risk share in 1/N of nine assets.
 2. Out of sample, 60-month window, August 2015 – July 2026 ($T = 132$): backtest 1/N, ERC and HRP with and without Bitcoin; Sharpe ratios, volatilities, drawdowns; HAC test of each with-minus-without difference.
 3. Plot Bitcoin's weight in ERC and HRP and its risk share in 1/N over time.
 4. Treat the three tests as one family: report raw and Holm-adjusted p -values.
- ▣ **Interpretation**
 - ▶ would you expect the past Bitcoin premium to persist?
 - ▶ what about the ex-post choice of Bitcoin?

B6 [Proposed]: HRP versus GMV on 16 ETFs — Task

- ▣ **Model:** B1 [Solved]
- ▣ **Data:** the combined universe and backtest of B4 (OOS May 2012 – July 2026, $T = 171$)
- ▣ **Tasks**
 1. Report Sharpe ratio, volatility and turnover of HRP and GMV, gross and net of 50 bp per unit of turnover; plot the gross and net points.
 2. Primary test on **gross** excess returns: $H_0: SR_{\text{HRP}} = SR_{\text{GMV}}$ with the HAC test and the studentized bootstrap (block 1), at 5% and 10%; report the bootstrap interval.
 3. Fill a decision table (test $\times$ level).
- ▣ **Interpretation**
 - ▶ GMV wins on volatility, HRP on Sharpe ratio: which criterion fits which investor?
 - ▶ what would a comparison at equal volatility require?

C3 [Proposed]: The Implementable Frontier on Our Data — Question

- ▣ **Question:** does the cost-aware ranking of [Jensen, Kelly, Malamud & Pedersen \(2026\)](#) (Table 2; lecture case study) hold on a small universe?
- ▣ **A multiweek project**, an alternative to C1; starter code: `jkmp_replication` in the Quantlet `MFM_ch4_implementable_frontier`
- ▣ **Notation**
 - ▶ π_t : risky weights; trade $\pi_t - g_t\pi_{t-1}$, with g_t the drift of last month's weights; cost $\frac{W_t}{2}(\pi_t - g_t\pi_{t-1})^\top \Lambda(\pi_t - g_t\pi_{t-1})$, $\Lambda_i = 0.2/V_i$, V_i the 6-month average daily dollar volume
 - ▶ Utility per month: net excess return $-\frac{\gamma}{2} \times$ its variance, $\gamma = 10$; the paper's rule trades part of the way (m) towards an aim portfolio learned from signals
- ▣ **Data:** JPM, BAC, C (Citigroup), GS, MS, WFC, NVDA, AAPL, MSFT, AMZN, GOOGL, GME, CSCO, MSTR and the nine sector ETFs, prices aligned on common days (from 19 August 2004); T-bill of the setup
- ▣ Signals at each month end, ranked into $[0, 1]$: $\ln(P_t/P_{t-1})$, $\ln(P_{t-1}/P_{t-6})$, $\ln(P_{t-1}/P_{t-12})$, 1-year volatility of daily returns, log 6-month dollar volume

C3 [Proposed]: The Implementable Frontier on Our Data — Tasks

- **Timeline:** signals and covariances use only past data; covariance windows expand from August 2004 until they reach 10 years (August 2014); hyperparameters chosen each year on validation years from 2011 only; first test weights end of January 2014, test returns February 2014 – August 2026
- **Tasks**
 1. Baseline first: build Portfolio-ML with one fixed hyperparameter set; check one month's cost step by step (W_t , trade vector, Λ).
 2. Tune on the grid of the paper (Table 1) and report, for Portfolio-ML, 1/N, minimum variance, Markowitz-ML and Static-ML: return, volatility, gross and net Sharpe, cost, utility, turnover, leverage (Table 2 layout).
 3. Posterior probability of a higher utility, $\Phi(\bar{d}/(s_d/\sqrt{T}))$ for the monthly utility differences d_t (Normal, flat prior, serial dependence ignored).
 4. Extension: wealth \$1, 10, 100 billion $\times \gamma \in \{1, 5, 10, 20, 100\}$; plot the implementable frontiers with 1/N's trading costs computed at each wealth level.
- **Report:** two tables, the frontier figure, a Quantlet folder; separate what follows the paper from any simplification you made
- **Interpretation**
 - ▶ if Portfolio-ML has a lower net Sharpe ratio than 1/N but a higher utility, which criterion matches the investor with $\gamma = 10$?
 - ▶ for which γ would the ranking reverse?

References (1/2)

- ▣ Black, F., Litterman, R. (1992). [Global portfolio optimization](#). *Financial Analysts Journal*, 48(5), 28–43.
- ▣ Britten-Jones, M. (1999). [The sampling error in estimates of mean-variance efficient portfolio weights](#). *Journal of Finance*, 54(2), 655–671.
- ▣ DeMiguel, V., Garlappi, L., Uppal, R. (2009). [Optimal versus naive diversification: how inefficient is the 1/N portfolio strategy?](#). *Review of Financial Studies*, 22(5), 1915–1953.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.
- ▣ He, G., Litterman, R. (2002). [The intuition behind Black-Litterman model portfolios](#). SSRN Working Paper 334304.
- ▣ Jensen, T.I., Kelly, B., Malamud, S., Pedersen, L.H. (2026). [Machine learning and the implementable efficient frontier](#). *Review of Financial Studies*, 39(10), 3035–3078.
- ▣ Jobson, J.D., Korkie, B.M. (1981). [Performance hypothesis testing with the Sharpe and Treynor measures](#). *Journal of Finance*, 36(4), 889–908.
- ▣ Kan, R., Zhou, G. (2007). [Optimal portfolio choice with parameter uncertainty](#). *Journal of Financial and Quantitative Analysis*, 42(3), 621–656.
- ▣ Laloux, L., Cizeau, P., Bouchaud, J.-P., Potters, M. (1999). [Noise dressing of financial correlation matrices](#). *Physical Review Letters*, 83(7), 1467–1470.
- ▣ Künsch, H.R. (1989). [The jackknife and the bootstrap for general stationary observations](#). *Annals of Statistics*, 17(3), 1217–1241.
- ▣ Ledoit, O., Wolf, M. (2004). [Honey, I shrunk the sample covariance matrix](#). *Journal of Portfolio Management*, 30(4), 110–119.
- ▣ Ledoit, O., Wolf, M. (2008). [Robust performance hypothesis testing with the Sharpe ratio](#). *Journal of Empirical Finance*, 15(5), 850–859.

References (2/2)

- Ledoit, O., Wolf, M. (2017). **Nonlinear shrinkage of the covariance matrix for portfolio selection: Markowitz meets Goldilocks.** *Review of Financial Studies*, 30(12), 4349–4388.
- Ledoit, O., Wolf, M. (2020). **Analytical nonlinear shrinkage of large-dimensional covariance matrices.** *Annals of Statistics*, 48(5), 3043–3065.
- López de Prado, M. (2016). **Building diversified portfolios that outperform out of sample.** *Journal of Portfolio Management*, 42(4), 59–69.
- Maillard, S., Roncalli, T., Teïletche, J. (2010). **The properties of equally weighted risk contribution portfolios.** *Journal of Portfolio Management*, 36(4), 60–70.
- Markowitz, H. (1952). **Portfolio selection.** *Journal of Finance*, 7(1), 77–91.
- Pele, D.T., Wesselhöfft, N., Härdle, W.K., Kolossiatis, M., Yatracos, Y.G. (2023). **Are cryptos becoming alternative assets?** *The European Journal of Finance*, 29(10), 1064–1105.
- Petukhina, A., Trimborn, S., Härdle, W.K., Elendner, H. (2021). **Investing with cryptocurrencies – evaluating their potential for portfolio allocation strategies.** *Quantitative Finance*, 21(11), 1825–1853.
- Theil, H., Goldberger, A.S. (1961). **On pure and mixed statistical estimation in economics.** *International Economic Review*, 2(1), 65–78.
- Data: Kenneth R. French Data Library; BNR reference exchange rates, www.bnr.ro.