



Modelling Financial Markets

Seminar 6: Multivariate Volatility and Dependence — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material

- ▣ AI: Artificial Intelligence
- ▣ AIC: Akaike Information Criterion
- ▣ AR: AutoRegressive (model); in event studies: Abnormal Return
- ▣ BAC: Bank of America Corporation (NYSE ticker)
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ CCC: Constant Conditional Correlation (model)
- ▣ CML: Canonical Maximum Likelihood
- ▣ COVID-19: COronaVirus Disease 2019
- ▣ COVOL: COmmon VOLatility (global COVOL: the size of the volatility shock common to all assets)
- ▣ DCC: Dynamic Conditional Correlation
- ▣ ETF: Exchange-Traded Fund
- ▣ FTSE: Financial Times Stock Exchange (FTSE Russell, index provider)
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GAS: Generalised Autoregressive Score (model)
- ▣ JPM: JPMorgan Chase & Co. (NYSE ticker)
- ▣ LQD: iShares iBoxx Investment Grade Corporate Bond ETF
- ▣ LR: Likelihood Ratio (test)
- ▣ PC1: first Principal Component
- ▣ US: United States
- ▣ USO: United States Oil Fund (ETF ticker)
- ▣ VIX: CBOE Volatility IndeX

Today's Question and Required Outputs

- Question of the seminar: **when do two markets move together, and how sure can we be?**
- Route
 - ▶ A1–A3 on paper: tail dependence of the t copula, copula parameters from Kendall's τ , the Forbes–Rigobon adjustment
 - ▶ B1: a DCC model for SPY and TLT; B2: Gaussian or t copula for the S&P 500 and the Euro Stoxx 50
 - ▶ B4: contagion or interdependence in 2008 and 2020; C2: find the errors in an AI answer; B3 and C1 at home
- Required outputs (checklist on the last core slide): the A answers, one table and one chart per B exercise, the corrected C2 lines, one interpretation sentence per exercise
- Notebook for today: [open the seminar notebook in Colab](#); each task names its notebook section
- Seminar: Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Worked example to follow
A1	do extremes arrive together in the t copula?	Solved	–
A2	which copula parameter matches Kendall's τ ?	Solved	–
A3	how much of a crisis correlation is volatility?	Solved	–
B1	how does the SPY–TLT correlation move?	Solved	–
B2	Gaussian or t copula for two equity indices?	Solved	–
B4	contagion in 2008 and 2020?	Proposed	A3
C2	what is wrong in an AI answer?	Proposed	A1, A2
B3, C1	joint losses; integration after 2020 (at home)	Proposed	B2; B1, B4
Extensions	A2, A3, B1, B2 extended; A4–A6, B5, B6, C3 (optional)	both	on each task

- **[Solved]**: full solution in the slides and in the notebook, a model to copy; **[Proposed]**: you solve it, the solution is discussed in class
- Part A: derivations on paper; Part B: estimation and inference on real data, each ending with an interpretation; Part C: AI critique and an open question

Data Used

Exercise	Series, price used	Frequency, sample	Rows
B1, A5	SPY and TLT, adjusted closes	daily, 31 July 2002 – 18 September 2026	6 073
A2, B2, B6	S&P 500 and Euro Stoxx 50, closes	weekly (Friday), 14 January 2000 – 18 September 2026	1 393
B3	BET and Euro Stoxx 50, closes	daily, 5 January 2010 – 18 September 2026	4 154
B4, A3 ext.	S&P 500, Euro Stoxx 50 and BET, closes	2-day, 5 January 2006 – 30 December 2021	3 880
B5	Bitcoin (USD) and S&P 500, closes	daily, 18 September 2014 – 18 September 2026	3 018
C1	BET and Euro Stoxx 50, closes	weekly, 15 January 2010 – 18 September 2026	871

- Rows: return observations after aligning the prices on common trading days (the rule for every joint analysis); log returns; weekly returns from the last common close of each week
- Daily prices of the course data, one series per asset: ETFs use the adjusted close, indices and Bitcoin the close; C3 uses 20 US-listed ETFs and the VIX index (see the C3 slide)

Setup: Run and Verify

```
# notebook section "Data" defines
# joint_returns(names, start) and
# weekly_returns(names, start):
# prices aligned on common days, then log returns
R = joint_returns(['spy', 'tlt'])
print(len(R), R.index[0].date(),
      R.index[-1].date())
print((100 * R).head(3).round(3))
W = weekly_returns(['sp500', 'stoxx'],
                  start='2000-01-01')
print(len(W))
```

- Expected: 6 073 rows from 31 July 2002 to 18 September 2026, then

date	SPY (%)	TLT (%)
31 July 2002	0.242	1.231
1 August 2002	-2.646	0.568
2 August 2002	-2.267	1.019


- and 1 393 weekly pairs
- If a count differs: check the series, the price used (close or adjusted close) and that prices were aligned before computing returns

Notation Used in Every Exercise

- $U = F_1(X_1)$, $V = F_2(X_2)$: the probability-integral transforms; pseudo-observations $\hat{u}_i = \text{rank}(x_i)/(n+1)$ (Genest, Ghoudi & Rivest (1995))
- t_ν : the distribution function of Student t with ν degrees of freedom; t_ν^{-1} its quantile function
- Tail dependence: $\lambda_L = \lim_{q \downarrow 0} C(q, q)/q$; at a finite q , $C(q, q)/q = P(V \leq q \mid U \leq q)$ is a probability, not the limit
- Kendall's τ : Gaussian and t $\tau = \frac{2}{\pi} \arcsin \rho$; Clayton $\theta/(\theta+2)$; Gumbel $1 - 1/\theta$
- DCC (Engle (2002)): $Q_t = (1 - a - b)\bar{Q} + a \varepsilon_{t-1} \varepsilon_{t-1}^\top + b Q_{t-1}$, $R_t = \text{diag}(Q_t)^{-1/2} Q_t \text{diag}(Q_t)^{-1/2}$, $\mathbb{E}(\varepsilon_t \varepsilon_t^\top \mid \mathcal{F}_{t-1}) = R_t$
- Forbes–Rigobon (Forbes & Rigobon (2002)): $\delta = \sigma_{x,c}^2 / \sigma_{x,calm}^2 - 1$, adjusted crisis correlation $\rho^* = \rho_c / \sqrt{1 + \delta(1 - \rho_c^2)}$

A1 [Solved]: Tail Dependence of the t Copula — Task

- ▣ **Question:** in the chart, as q falls the Gaussian probability heads to 0 while the t curves level off; what happens as $q \rightarrow 0$?
- ▣ **Inputs:** $\rho = 0.5$, $\nu \in \{3, 4, 10, 30\}$; fact: for a bivariate t, given $X_1 = x$, $(X_2 - \rho x) / \sqrt{(\nu + x^2)(1 - \rho^2)/(\nu + 1)} \sim t_{\nu+1}$
- ▣ **Tasks**
 1. Use L'Hôpital's rule and exchangeability to show that $\lambda_L = 2 \lim_{q \downarrow 0} P(V \leq q \mid U = q)$.
 2. Substitute the conditional t distribution, let $x = t_\nu^{-1}(q) \rightarrow -\infty$ and derive $\lambda = 2 t_{\nu+1} \left(-\sqrt{(\nu + 1)(1 - \rho)/(1 + \rho)} \right)$.
 3. Repeat the argument for the Gaussian copula and show that $\lambda = 0$ for $|\rho| < 1$.
 4. Evaluate λ for each ν , and $C(q, q)/q$ at $q = 5\%$ and 0.1% for the Gaussian and the t copula with $\nu = 4$, same ρ (notebook helper `joint_tail_prob`).
- ▣ **Report:** the derivation as a chain of equalities and a table of λ and of the four probabilities



- ▣ Exact $P(V \leq q \mid U \leq q)$ from the bivariate distribution functions; dotted: the limits λ ; notebook: section A1

A1 [Solved]: Solution (1/2) — Derivation

- (1) $C(q, q) \rightarrow 0$ and $q \rightarrow 0$, so by L'Hôpital $\lambda_L = \lim_{q \rightarrow 0} \frac{d}{dq} C(q, q) = \lim_{q \rightarrow 0} [\partial_1 C(q, q) + \partial_2 C(q, q)]$, with $\partial_1 C(u, v) = P(V \leq v \mid U = u)$; exchangeability makes the two terms equal: $\lambda_L = 2 \lim_{q \rightarrow 0} P(V \leq q \mid U = q)$
- (2) With $x = t_\nu^{-1}(q)$: $P(V \leq q \mid U = q) = P(X_2 \leq x \mid X_1 = x) = t_{\nu+1} \left(\frac{(1-\rho)x}{\sqrt{(\nu+x^2)(1-\rho^2)/(\nu+1)}} \right)$
- As $x \rightarrow -\infty$, $x/\sqrt{\nu+x^2} \rightarrow -1$; since $(1-\rho)/\sqrt{1-\rho^2} = \sqrt{(1-\rho)/(1+\rho)}$, the argument tends to $-\sqrt{(\nu+1)(1-\rho)/(1+\rho)}$, which gives the formula
- (3) Gaussian: $P(V \leq q \mid U = q) = \Phi(x\sqrt{(1-\rho)/(1+\rho)})$ with $x = \Phi^{-1}(q) \rightarrow -\infty$, so the limit is 0
- Why: in the t case the conditional scale grows with $|x|$, so the standardised argument stays bounded; in the Gaussian case the scale is fixed and the argument diverges

A1 [Solved]: Solution (2/2) — Numbers

ν	$a_\nu = \sqrt{(\nu+1)/3}$	$\lambda = 2t_{\nu+1}(-a_\nu)$
3	1.155	0.313
4	1.291	0.253
10	1.915	0.082
30	3.215	0.003


- $\nu = 4$ step by step: $t_5(-1.291) = 0.1266$, so $\lambda = 2 \times 0.1266 = 0.253$
- Tail dependence vanishes as ν grows

	$q = 5\%$	$q = 0.1\%$
Gaussian, $\rho = 0.5$	0.244	0.054
t, $\rho = 0.5, \nu = 4$	0.339	0.263

- $C(q, q)/q = P(V \leq q \mid U \leq q)$ with the same ρ in both models
- Interpretation: at 5% the two models look alike, but at 0.1% the Gaussian probability is a fraction of the t one; a model judged on moderate losses cannot reveal joint crashes

A2 [Solved]: Kendall's Tau and Copula Parameters — Task

- **Question:** which copula parameter reproduces the rank dependence of two equity indices, and does it fix the tail risk?
- **Inputs:** weekly S&P 500 and Euro Stoxx 50 log returns (Data Used), $n = 1\,393$; $\hat{\tau} = 0.5642$, $\hat{\rho}_S = 0.7462$, Pearson 0.781; pairs treated as i.i.d. (an unconditional weekly working model)
- **Tasks**
 1. Invert $\hat{\tau}$ for the Gaussian ($\rho = \sin(\pi\tau/2)$), Clayton ($\theta = 2\tau/(1-\tau)$), Gumbel ($\theta = 1/(1-\tau)$) and Frank copulas; for Frank solve $1 - \frac{4}{\theta}[1 - D_1(\theta)] = \hat{\tau}$ with $D_1(\theta) = \frac{1}{\theta} \int_0^\theta \frac{t}{e^t - 1} dt$ (helper theta_from_tau).
 2. Compute the tail dependence implied by each fitted copula.
 3. Compute the Spearman's $\rho_S = \frac{6}{\pi} \arcsin(\rho/2)$ implied by the Gaussian copula and compare it with $\hat{\rho}_S$.
- **Report:** a table of the four parameters and their λ_L , λ_U , and the gap of task 3



- Left: weekly returns; right: pseudo-observations, boxes at the 5% corners; notebook: section A2

A2 [Solved]: Solution

Copula	Parameter	λ_L	λ_U
Gaussian	$\rho = 0.775$	0	0
Clayton	$\theta = 2.589$	0.765	0
Gumbel	$\theta = 2.295$	0	0.647
Frank	$\theta = 7.044$	0	0

- Gaussian: $\sin(\pi \times 0.564/2) = 0.775$; Clayton: $2 \times 0.564/0.436 = 2.589$; Gumbel: $1/0.436 = 2.295$; Frank: numerical root
- Clayton $\lambda_L = 2^{-1/\theta}$, Gumbel $\lambda_U = 2 - 2^{1/\theta}$


- Task 3: $\frac{6}{\pi} \arcsin(0.775/2) = 0.760$ against $\hat{\rho}_S = 0.746$: a gap of 0.0134 (its precision: extension A2)
- The same $\hat{\tau}$ gives tail dependence between 0 and 0.77, depending only on the family chosen
- Interpretation: $\hat{\tau}$ identifies a parameter within a family, never the family; the joint tail must be estimated (B2), not read off τ

A3 [Solved]: How Much of a Crisis Correlation Is Volatility? — Task

- **Question:** a correlation rises from 0.40 in a calm year to 0.60 in a crisis while the source market's variance quadruples; is this contagion?
- **Model:** $y_t = \alpha + \beta x_t + \varepsilon_t$, $\text{Cov}(x_t, \varepsilon_t) = 0$, β and σ_ε^2 constant; in the crisis $\sigma_{x,c}^2 = (1 + \delta)\sigma_x^2$
 - ▶ Scaffold: $\text{Cov}(x, y) = \beta\sigma_x^2$, $\text{Var}(y) = \beta^2\sigma_x^2 + \sigma_\varepsilon^2$, hence $\rho^2 = \beta^2\sigma_x^2/(\beta^2\sigma_x^2 + \sigma_\varepsilon^2)$ and $\rho_c = \rho\sqrt{(1 + \delta)/(1 + \delta\rho^2)}$
- **Tasks**
 1. Compute δ and the adjusted crisis correlation $\rho^* = \rho_c/\sqrt{1 + \delta(1 - \rho_c^2)}$.
 2. Compute the crisis correlation we would measure if the true correlation stayed at 0.40.
 3. State whether the rise is evidence of contagion under the maintained assumptions, and name the assumption you trust least.
- **Report:** each number with its formula and substituted values, and one sentence of interpretation; notebook: section A3

A3 [Solved]: Solution


- (1) Variance ratio 4, so $\delta = 3$;
 $\rho^* = 0.60 / \sqrt{1 + 3(1 - 0.36)} = 0.60 / \sqrt{2.92} = 0.60 / 1.709 = 0.351$
- (2) $\rho_c = 0.40 \sqrt{4 / (1 + 3 \times 0.16)} = 0.40 \sqrt{4 / 1.48} = 0.40 \times 1.644 = 0.658$
- (3) The adjusted 0.351 is below the calm 0.40, and volatility alone would lift 0.40 to 0.658, above the observed 0.60: no evidence of an increase after this adjustment
- Least credible: no common shock; a global shock raises σ_ε too, and the adjustment then over-corrects (A4, lecture)



- The observed 0.60 lies below the curve of an unchanged 0.40 and on the curve of a true 0.351

B1 [Solved]: DCC for SPY and TLT — Task

- **Question:** how has the stock–bond correlation moved since 2002, and is it constant?
- **Data:** SPY and TLT adjusted closes, aligned on common days, log returns in %, $T = 6073$
- **Tasks**
 1. Fit a GARCH(1,1) with constant mean to each series, form $\varepsilon_{it} = (r_{it} - \hat{\mu}_i) / \hat{\sigma}_{it}$ and compare the autocorrelation of r_{it}^2 and of ε_{it}^2 .
 2. Set $\bar{Q}$ to the sample covariance of ε_t , $Q_1 = \bar{Q}$, compute R_2 from ε_1 , then maximise $\ell_C(a, b) = -\frac{1}{2} \sum_t (\ln |R_t| + \varepsilon_t^\top R_t^{-1} \varepsilon_t - \varepsilon_t^\top \varepsilon_t)$ subject to $a, b \geq 0$, $a + b < 1$.
 3. Plot $\rho_{12,t}$ against the 252-day rolling correlation and the target, and re-estimate on 2002–2019 and 2020–2026.
 4. Test CCC ($a = b = 0$) with the likelihood ratio, calibrated by a bootstrap of both steps under CCC (199 paths, a few minutes).




- **Report:** GARCH and DCC tables, R_2 , the path chart, the LR with its bootstrap p , one interpretation; notebook: section B1

B1 [Solved]: Solution (1/4) — Step 1, GARCH Filtering

	ω	α	β	ann. vol., mean
SPY	0.0283	0.127	0.848	16.0%
TLT	0.0057	0.049	0.943	13.6%

- Returns in %; volatility peaks: SPY 105% (17 March 2020), TLT 53% (24 March 2020)
- Checkpoint: lag-1 autocorrelation of squares, SPY 0.273 $\rightarrow$ -0.008 , TLT 0.373 $\rightarrow$ 0.021; mean over lags 1–20 SPY 0.233 $\rightarrow$ -0.002
- Filtering removes the volatility clustering: the correlation of ε_t is no longer driven by volatility regimes




B1 [Solved]: Solution (2/4) — Step 2, the DCC Estimates

- One recursion step: $\bar{Q} = \begin{pmatrix} 0.998 & -0.248 \\ -0.248 & 1.000 \end{pmatrix}$,
 $\varepsilon_1 = (0.080, 1.498)$ on 31 July 2002
- $Q_2 = (1 - a - b)\bar{Q} + a\varepsilon_1\varepsilon_1^\top + b\bar{Q}$: $Q_{2,11} = 0.947$,
 $Q_{2,22} = 1.065$, $Q_{2,12} = -0.2293$
- $R_{2,12} = -0.2293/\sqrt{0.947 \times 1.065} = -0.2283$: a positive TLT shock with a flat SPY day moves the correlation towards zero

	estimate	se (Hessian)
a	0.0522	0.0059
b	0.9335	0.0084
$a + b$	0.9857	—
$\bar{Q}_{12}$	-0.248	—

- Half-life of a shock to Q_t : $\ln 0.5 / \ln(a + b) = 48.3$ days, an approximation since $R_t \neq Q_t$
- Hessian se treats step 1 as known; both-step bootstrap: extension B1

B1 [Solved]: Solution (3/4) — Path and Subsamples




	2002–2019	2020–2026
a	0.0565	0.0341
b	0.9199	0.9527
$\bar{Q}_{12}$	-0.350	0.003
T	4 386	1 687

- Full sample: $\rho_{12,t}$ from -0.79 to 0.49 ; above the target on 92% of days since 2022
- Persistence is similar in both halves, but the target moves from about -0.35 to about 0
- Interpretation: a full-sample DCC pulls today's correlation towards a target that no longer applies

 MFM_ch6_dcc_spy_tlt

B1 [Solved]: Solution (4/4) — Is the Correlation Constant?

- $LR = 2(\ell_C^{DCC} - \ell_C^{CCC}) = 2(460.3 - 193.5) = 533.4$
- χ^2_2 is not the null law: under $a = 0$, b is not identified; the bootstrap under CCC
 1. decorrelate ε_t with the constant correlation matrix and resample the pairs;
 2. rebuild CCC innovations and GARCH returns with the fitted parameters;
 3. re-estimate both steps and the LR on each of 199 paths;
 4. $p = (1 + \#\{LR^{(b)} \geq LR\})/200$.
- Largest null value 15.7, 95% null quantile 6.6: $p = 0.005$, the smallest possible; reject constant correlation



- This bootstrap simulates under CCC to test dynamics; the both-step bootstrap of the extension simulates under the fitted DCC to measure uncertainty

B2 [Solved]: Gaussian or t Copula for the S&P 500 and the Euro Stoxx 50? — Task


- **Question:** do weekly joint crashes of US and euro-area equities need a copula with tail dependence?
- **Data and model:** the pseudo-observations of A2 ($n = 1\,393$); an unconditional i.i.d. weekly working model; canonical maximum likelihood (CML) on the ranks; the Gaussian copula is the t copula with $\nu = \infty$, a boundary point
- **Tasks**
 1. Fit the Gaussian and the t copula by CML and compare them by AIC.
 2. Compute the likelihood ratio for $H_0 : \nu = \infty$; its p -value comes from a parametric bootstrap under H_0 (simulate, re-rank, refit both; 499 draws; with all B2 bootstraps the notebook cell runs for about 15 minutes).
 3. Run the Rosenblatt goodness-of-fit test ([Genest, Rémillard & Beaudoin \(2009\)](#)) for both copulas (simulate from the fitted copula, re-rank, refit; 499 draws).
 4. Compare the empirical $C_n(q, q)/q$ with the fitted $C_\theta(q, q)/q$ at the same $q = 5\%$, and report the limit λ of the t copula separately.
- **Report:** a table ($\hat{\rho}$, $\hat{\nu}$, ℓ , AIC, fit p), the LR with its p , the finite-threshold chart; notebook: section B2

B2 [Solved]: Solution (1/2) — Fit and Tests

	$\hat{\rho}$	$\hat{\nu}$	ℓ	AIC	fit p
Gaussian	0.771	–	623.3	-1244.6	0.002
Student t	0.775	3.91	674.2	-1344.4	0.874

- (1) The t copula has the lower AIC by 99.8 points: the extra parameter ν is strongly supported
- (2) $LR = 2(674.2 - 623.3) = 101.8$; the largest of 499 null draws is 87.1, so $p = 0.002$, the smallest possible (1/500): reject the Gaussian copula
- χ_1^2 would be wrong: $\nu = \infty$ lies on the boundary, and the $\frac{1}{2}\chi_0^2 + \frac{1}{2}\chi_1^2$ mixture is not established for a rank-based pseudo-likelihood
- (3) Fit test: Gaussian $p = 0.002$, rejected at 5%; t $p = 0.874$, not rejected: absence of evidence against the t copula, not proof

B2 [Solved]: Solution (2/2) — The Tails at a Finite Threshold




- (4) At $q = 5\%$: empirical lower 0.517 (36 weeks in the corner), upper 0.488 (34 weeks)
- Fitted at the same q : Gaussian 0.462, t 0.544; the limit of the t copula is $\lambda = 0.467$, a different object
- The t curve tracks the empirical points down to small q ; the Gaussian curve falls away, as A1 predicts
- Interpretation: for weekly equity pairs a risk model needs tail dependence; the empirical upper corner at small q rests on a handful of weeks

B4 [Proposed]: Contagion in 2008 and 2020? — Task

- Question:** did the links between the S&P 500 and the Euro Stoxx 50 or BET strengthen in the two crises, once the higher US volatility is accounted for?
- Data:** 2-day log returns $\ln P_t - \ln P_{t-2}$ on common trading days; the S&P 500 is the source; a return belongs to the window of its end date

	calm	crisis
2008	14 Sep 2007 – 12 Sep 2008	15 Sep 2008 – 31 Mar 2009
2020	19 Feb 2019 – 19 Feb 2020	20 Feb – 30 Apr 2020
n (2-day)	242, 241	130, 47



B4 [Proposed]: Contagion in 2008 and 2020? — Task (continued)

▣ **Worked example to follow:** A3; **Tasks**

1. Work 2008 with BET fully: compute ρ_{calm} , ρ_c , δ and ρ^* , then repeat for the other three cases.
2. As a historical comparison, run the one-sided Fisher z tests of $H_0 : \rho_c \leq \rho_{calm}$ and $H_0 : \rho^* \leq \rho_{calm}$, with n half the number of overlapping 2-day returns.
3. Test $D = \rho^* - \rho_{calm}$ by bootstrap with δ estimated in each draw:
 - ▶ resample the joint vectors (S&P 500, target) within each window with a stationary bootstrap, mean block 10 days, 1 999 draws;
 - ▶ recompute $\hat{\delta}$, ρ^* and ρ_{calm} ; report the 90% interval of D and $p =$ the share of draws with $D \leq 0$.
4. Sensitivity: end the crisis windows on 31 December 2008 and on 31 March 2020.

▣ **Report:** a table by crisis and market (δ , ρ_{calm} , ρ_c , p_{raw} , ρ^* , p_{adj} , interval of D) and one chart; notebook: section B4

▣ **Interpretation:** why is δ so much larger in 2020, and which assumption is least credible for a pandemic?

C2 [Proposed]: Critique an AI Answer — the Answer

- **Prompt** sent to an AI assistant: “e holds the GARCH-standardised residuals of JPM and BAC (two US bank stocks). Calibrate a t copula with $\nu = 4$ from Kendall's τ and give me its tail dependence.”
- **Answer:** the code below and the conclusion

```
import numpy as np
from scipy import stats
tau = stats.kendalltau(e['JPM'], e['BAC'])[0]      # Kendall's tau
rho = 2 * np.sin(np.pi * tau / 6)                # copula correlation
u = e.rank() / len(e)                             # pseudo-observations
z = stats.t.ppf(u, df=4)                          # input of the t-copula likelihood
nu = 4
x = np.sqrt((nu + 1) * (1 - rho) / (1 + rho))
lam = 2 * stats.t.cdf(-x, df=nu)                  # tail dependence
```


- **Conclusion of the AI:** “With $\hat{\tau} = 0.50$: $\rho = 0.518$ and $\lambda_L = \lambda_U = 0.276$.”

C2 [Proposed]: Critique an AI Answer — Task

- ▣ **Worked example to follow:** A1 and A2 [Solved]; **inputs:** the code above with $\hat{\tau} = 0.50$ and $\nu = 4$; for the data check, the JPM–BAC residuals of the notebook (section C2)
- ▣ The answer contains **three** mistakes typical of copula modelling; everything else is correct
- ▣ **Tasks**
 1. Find the three wrong lines and explain in one sentence why each is wrong.
 2. For each, give the evidence that exposes it: a formula from A1 or A2, a check of the largest pseudo-observation, or a small simulation.
 3. Write the three corrected lines and recompute ρ and λ for $\hat{\tau} = 0.50$, $\nu = 4$.
 4. Say which errors change the displayed λ and which one only breaks the likelihood inputs.
- ▣ **Report:** a three-row table (error, evidence, correction) and the corrected ρ and λ ; log in `AI_ERRORS.md` the errors you actually found


B3 [Proposed]: Clayton or Gumbel for Joint Losses of BET and Euro Stoxx 50? — Task

- **Worked example to follow:** B2 and A2; **data:** daily BET and Euro Stoxx 50 closes from 5 January 2010, aligned on common days (almost synchronous closes), $T = 4\,154$; notebook: section B3
- Filter each series with GARCH(1,1); losses $L = -\varepsilon$ (minus the standardised returns) and their pseudo-observations: joint crashes sit in the *upper-right* corner
- **Tasks**
 1. Predict which family puts its tail dependence in the corner of joint crashes: Clayton or Gumbel on losses.
 2. Fit Clayton and Gumbel by CML (Gaussian and t as benchmarks); report AIC, λ_L , λ_U and the Rosenblatt fit p (200 draws; 100 for t).
 3. Compute the empirical corner probabilities at $q = 5\%$ and 10% , and compare them with the fitted curves.
 4. Explain what the 180° rotated (survival) Clayton on losses represents.



C1 [Proposed]: Did Romanian and Euro-Area Equities Become More Integrated after 2020? — Task

- ▣ **Worked example to follow:** B1, B4; **data:** weekly BET and Euro Stoxx 50 returns on common days, 2010–2019 ($n = 520$) against 2020–2026 ($n = 351$); notebook: section C1
- ▣ 2020 brought COVID-19 and the FTSE Russell upgrade of Romania (FTSE Russell (2019)): a descriptive period comparison, not a causal test of the upgrade
- ▣ **Tasks**
 1. Estimate the change in weekly correlation with a 95% stationary-bootstrap interval (Politis & Romano (1994), resample within each period, mean block 8 weeks, 2000 draws) and the one-sided p for H_0 : no increase, the share of draws ≤ 0 .
 2. Run two separate checks against the same baseline: the Forbes–Rigobon adjustment for the change in Euro Stoxx 50 volatility, and the correlation without 20 February – 30 April 2020.
 3. Optional: Kendall's τ , the beta $\text{Cov}(\text{BET}, \text{Stoxx})/\text{Var}(\text{Stoxx})$, the $q = 10\%$ lower-corner probability and a weekly DCC.



- ▣ Context from the lecture: 104-week rolling correlations
- ▣ **Report:** one table of all measures in both periods and a verdict with its caveats

Submission Checklist and Team Repository

Hand in before the next seminar

- ▣ A1–A3: derivations and the tables of numbers
- ▣ B1, B2: one table and one chart each, as in the solutions
- ▣ B4: the table by crisis and market and the chart of D
- ▣ C2: the three-row table and the corrected numbers
- ▣ B3 and C1: table, chart and verdict
- ▣ One interpretation sentence per exercise

Team repository

- ▣ Choose **one** dependence question for your project (for example C1, B5 or C3) rather than repeating every estimator
- ▣ A Quantlet folder: code, Metainfo.txt, charts
- ▣ Document your checks; log AI use in AI_USE.md and the errors actually found in AI_ERRORS.md
- ▣ Extension packs are optional

Extension Packs: What Each One Needs

Pack	Exercises	Prerequisites
1. Copula inference	A2 Extended [Solved], B2 Extended [Solved]	A2, B2
2. Contagion inference	A4 [Proposed], A3 Extended [Solved]	A3, B4
3. DCC dynamics	A5 [Proposed], B1 Extended [Solved], B5 [Proposed]	B1
4. Margins and dynamic copulas	A6 [Proposed], B6 [Proposed]	A2, B2
5. Project brief	C3 [Proposed]: global common volatility	B1, C1, lecture case study

- Each pack has its own notebook section; long bootstraps take a few minutes
- C3 is a project, not a seminar exercise: its reference outputs are partial, so it suits a team that wants a replication

A2 Extended [Solved]: How Precise Are Kendall's τ and the Copula Parameters? — Task

- ▣ **Inputs:** the pseudo-observations of A2, $n = 1\,393$, $\hat{\tau} = 0.5642$; empirical copula


$$C_n(u, v) = \frac{1}{n} \sum_i 1\{\hat{U}_i \leq u, \hat{V}_i \leq v\}$$

- ▣ **Tasks**

1. With the Hoeffding projection $W_i = 2C_n(\hat{U}_i, \hat{V}_i) - \hat{U}_i - \hat{V}_i$ (sample standard deviation 0.1239), show that $\text{se}(\hat{\tau}) = 4 \text{sd}(W)/\sqrt{n}$ and compute it.
 2. Propagate it to $\hat{\theta}_{\text{Clayton}}$ and $\hat{\rho}_{\text{Gauss}}$ by the delta method.
 3. Bootstrap the weekly pairs jointly (1 999 draws), recompute both rank statistics in each draw and report the 95% percentile interval of the gap $g(\hat{\tau}) - \hat{\rho}_S$ of A2.
- ▣ **Report:** three standard errors, the interval and a decision at 5%; the i.i.d. bootstrap assumes no serial dependence in the weekly pairs

A2 Extended [Solved]: Solution

- (1) $\hat{\tau}$ is a U-statistic with $E[h | Z_i] = 2W_i + 1$, so $\text{Var}(\hat{\tau}) \approx 16\text{Var}(W)/n$: $\text{se} = 4 \times 0.1239/\sqrt{1393} = 0.0133$
- (2) Clayton: $d\theta/d\tau = 2/(1-\tau)^2 = 10.53$, $\text{se } 0.140$; Gaussian: $\frac{\pi}{2} \cos(\pi\tau/2) = 0.993$, $\text{se } 0.0132$ (Genest & Rivest (1993))
- (3) Gap 0.0134, $\text{se } 0.0026$, 95% interval [0.0088, 0.0187]: zero is outside, so the gap is significant at 5%
- Why so precise: $\hat{\tau}$ and $\hat{\rho}_S$ move together across samples, so their difference has a much smaller se than each
- Interpretation: the pair $(\hat{\tau}, \hat{\rho}_S)$ already rejects the Gaussian copula, consistently with B2




B2 Extended [Solved]: Bootstrap Distributions and Sandwich Standard Errors — Task

▣ **Inputs:** the B2 fits; the ranks are estimated, so the inverse Hessian understates the uncertainty of CML estimates

▣ **Tasks**

1. Plot the null distribution of the LR of B2 against the observed value.
2. Build 90% intervals for ν and λ from 200 draws of the fitted t copula, re-ranked and refitted.
3. Compute the se of $\hat{\rho}$ and $\hat{\nu}$ from the inverse Hessian A^{-1}/n and from the [Genest, Ghoudi & Rivest \(1995\)](#) sandwich $A^{-1}\Sigma A^{-1}/n$, $\Sigma = \text{Var}[\dot{\ell}_{\theta} + W_1(U) + W_2(V)]$.

B2 Extended [Solved]: Solution



- (1) 499 draws under H_0 : 95% quantile 3.2, largest 87.1, against the observed 101.8
- (2) $\hat{\nu} = 3.91$, 90% interval [2.98, 4.77];
 $\hat{\lambda} = 0.467$, [0.423, 0.518]
- (3) Gaussian $\text{se}(\hat{\rho})$ 0.0086 versus 0.0138; t $\text{se}(\hat{\rho})$ 0.0115 versus 0.0131; $\text{se}(\hat{\nu})$ 0.56 versus 0.58
- Interpretation: ranking adds variance to $\hat{\rho}$; ν is driven by the tails, where ranking errors matter little

A4 [Proposed]: Deriving the Forbes–Rigobon Bias — Task


- ▣ **Worked example to follow:** A3; **model:** $y_t = \alpha + \beta x_t + \varepsilon_t$ with $\text{Cov}(x_t, \varepsilon_t) = 0$ maintained; β and σ_ε^2 equal in calm and crisis; only $\sigma_{x,c}^2 = (1 + \delta)\sigma_x^2$ changes
- ▣ **Tasks**
 1. Write ρ^2 as a function of $k = \beta^2 \sigma_x^2 / \sigma_\varepsilon^2$, using $\text{Cov}(x, y)$ and $\text{Var}(y)$.
 2. Derive $\rho_c = \rho \sqrt{(1 + \delta) / (1 + \delta \rho^2)}$ and show that $|\rho_c| > |\rho|$ for $0 < |\rho| < 1$.
 3. Invert the relation to obtain $\rho^* = \rho_c / \sqrt{1 + \delta(1 - \rho_c^2)}$.
 4. Suppose that in the crisis σ_ε^2 also rises, to $(1 + \gamma)\sigma_\varepsilon^2$ with $0 < \gamma < \delta$: compute the true crisis correlation and say in which direction ρ^* is then biased.
- ▣ **Report:** each step with one sentence of justification and the answer to task 4 in economic terms

A3 Extended [Solved]: The Forbes–Rigobon Adjustment with an Estimated δ — Task

- ▣ **Inputs** (2008 windows of B4): $\hat{\delta} = 6.54$; effective $n_c = 65$, $n_0 = 121$, half the overlapping 2-day returns (an approximation)
- ▣ **Tasks**
 1. Under i.i.d. Normality, $(1 + \hat{\delta})/(1 + \delta) \sim F(n_c - 1, n_0 - 1)$: derive $\text{Var}(\hat{\delta})$ to first order.
 2. Write the gradient (g_r, g_δ) of ρ^* in $(\hat{\rho}_c, \hat{\delta})$ and their covariance (same crisis window).
 3. Compute $\text{se}(\rho^*)$ with δ fixed and random for both markets, and the z of $\rho^* - \hat{\rho}_{calm}$ including the uncertainty of $\hat{\rho}_{calm}$.
- ▣ These Normal, effective-sample formulas are illustrative; the stationary bootstrap of B4 is the primary inference
- ▣ **Report:** $\text{se}(\hat{\delta})$, the two se of ρ^* , their percentage difference, the two z values

A3 Extended [Solved]: Solution

- (1) $\text{Var} \ln(1 + \hat{\delta}) \approx \frac{2}{n_c - 1} + \frac{2}{n_0 - 1}$, so
 $\text{se}(\hat{\delta}) = 7.54 \sqrt{2(1/64 + 1/120)} = 1.65$
- (2) $k = 1 + \delta(1 - \rho_c^2)$, $g_r = (1 + \delta)k^{-3/2}$,
 $g_\delta = -\rho_c(1 - \rho_c^2)k^{-3/2}/2$; $\text{Cov}(\hat{\rho}_c, \hat{\delta}) \approx (1 + \delta)\rho_c(1 - \rho_c^2)/n_c$
- $\text{Var}(\rho^*) \approx g_r^2 \text{Var} \hat{\rho}_c + g_\delta^2 \text{Var} \hat{\delta} + 2g_r g_\delta \text{Cov}(\hat{\rho}_c, \hat{\delta})$
- (3) Euro Stoxx 50: $\text{se } 0.055 \rightarrow 0.051$ (−7%), $z = -2.44$; BET:
 $0.049 \rightarrow 0.047$ (−3%), $z = -0.19$
- A random δ need not widen the interval: a volatile crisis window raises $\hat{\delta}$ and $\hat{\rho}_c$ together, and their errors partly cancel; still no evidence of an increase



A5 [Proposed]: DCC Algebra and a Shock — Task

- ▣ **Worked example to follow:** B1; target $\bar{Q} = \begin{pmatrix} 1 & -0.248 \\ -0.248 & 1 \end{pmatrix}$, $a = 0.0522$, $b = 0.9335$, $Q_{t-1} = \bar{Q}$, a shock $\varepsilon_{t-1} = (-2.0, 1.0)$; notation: $Q_t^* = \text{diag}(Q_t)$, a time-varying matrix, not the fixed target

- ▣ **Tasks**

1. Show by induction that Q_t is positive definite for every t when $Q_0 = \bar{Q}$ is.
 2. Compute Q_t and $R_{t,12}$ after the shock, and the approximate half-life $\ln 0.5 / \ln(a + b)$.
 3. Show that $\mathbb{E}[Q_t] = \bar{Q}$ requires $\mathbb{E}[R_t] = \bar{Q}$, which fails in general; for the corrected DCC of [Aielli \(2013\)](#), with $Q_t^{*1/2} \varepsilon_{t-1} \varepsilon_{t-1}^\top Q_t^{*1/2}$ in place of $\varepsilon_{t-1} \varepsilon_{t-1}^\top$, show $\mathbb{E}[Q_t] = \bar{Q}^*$.
- ▣ **Report:** the proofs, $R_{t,12}$ and the half-life; is the half-life a rule for how often to rebalance a hedge?


B1 Extended [Solved]: When Did the Correlation Break, and How Precise Is DCC? — Task

▣ **Inputs:** the standardised residuals and the DCC fit of B1

▣ **Tasks**

1. Compute the [Wied, Krämer & Dehling \(2012\)](#) statistic $Q_T = \max_j \frac{j}{\sqrt{T\hat{D}}} |\hat{\rho}_j - \hat{\rho}_T|$ ($\hat{\rho}_j$: correlation of the first j residual pairs, $\hat{D}$: Bartlett long-run se, bandwidth $\lfloor T^{1/4} \rfloor$), test at 5% (1.358) and date the break.
2. Explain why a Fisher test at January 2022, a date read off the lecture chart, is not valid.
3. Simulate 199 paths under the fitted DCC (resampled decorrelated residuals); compute the se of (a, b) re-estimating step 2 only and both steps ([Newey & McFadden \(1994\)](#)).

B1 Extended [Solved]: Solution



se of	$\hat{a}$	$\hat{b}$
Hessian	0.0059	0.0084
step 2 only	0.0051	0.0077
both steps	0.0052	0.0077

- (1) $Q_T = 4.25 > 1.358$ ($p < 0.001$); break 3 August 2020; residual correlation -0.361 before, 0.066 after; a second break on 9 July 2007
- The limit law is proved for observed series: on estimated GARCH residuals p is approximate
- (2) Fisher at January 2022 gives $z = 14.9$, but the date was chosen after seeing the data and Fisher assumes i.i.d. Normal pairs
- (3) Step 1 adds only 2% and 1%: $T \approx 6000$ pins down GARCH; the Hessian $se(\hat{b})$ is conservative here; in short samples the correction matters

B5 [Proposed]: Bitcoin and the S&P 500: DCC with Bootstrap Bands — Task

- ▣ **Worked example to follow:** B1; **data:** Bitcoin (USD) and the S&P 500 index (the index, not SPY: the lecture's SPY result is a separate benchmark), common trading days since 18 September 2014, $T = 3018$; weekend Bitcoin moves enter the Monday return
- ▣ **Tasks**
 1. Fit GARCH(1,1) and DCC(1,1); compute the average conditional correlation before and after 1 January 2020.
 2. Bootstrap the parameter uncertainty in five steps, 200 draws:
 - ▶ keep the GARCH residuals and the target $\bar{Q}$ fixed; decorrelate $\eta_t = L_t^{-1}\varepsilon_t$ (L_t the Cholesky factor of R_t); resample the η_t jointly;
 - ▶ simulate a DCC with $(\hat{a}, \hat{b})$, re-estimate (a, b) , then refilter the *observed* residuals with each bootstrap (a, b) .
 3. Draw the 90% pointwise band of ρ_t and a 90% interval for the change in the mean correlation after 2020.
- ▣ **Interpretation:** which uncertainty does the band ignore, and why is it not a prediction interval?

A6 [Proposed]: Attainable Correlations — Task

- ▣ **Worked example to follow:** A2; $X = e^{\sigma_1 Z_1}$, $Y = e^{\sigma_2 Z_2}$, $Z_i \sim N(0, 1)$: hypothetical lognormal variables with fixed margins and free dependence
- ▣ **Tools:** $M(u, v) = \min(u, v)$, $W(u, v) = \max(u + v - 1, 0)$, $W \leq C \leq M$; Hoeffding:
 $\text{Cov}(X, Y) = \iint [F(x, y) - F_1(x)F_2(y)] dx dy$; $\mathbb{E}(e^{aZ}) = e^{a^2/2}$
- ▣ **Tasks**
 1. Show that $\rho(X, Y)$ is largest for $Z_2 = Z_1$ and smallest for $Z_2 = -Z_1$.
 2. Compute in turn $\mathbb{E}X$, $\mathbb{E}Y$, $\mathbb{E}XY$, $\text{Var}X$, $\text{Var}Y$ and derive $\rho_{\max}$ and $\rho_{\min}$ in closed form.
 3. Evaluate them for $\sigma_1 = 1$, $\sigma_2 = 2$ and interpret $\rho = 0.3$ relative to this interval.

B6 [Proposed]: A Dynamic t Copula for the S&P 500 and the Euro Stoxx 50 — Task

- **Worked example to follow:** B2 and the lecture's dynamic t copula; data: the weekly pseudo-observations of A2 ($n = 1393$); benchmark: the static t copula of B2
- **GAS t copula** (Creal, Koopman & Lucas (2013), unit scaling): $\rho_t = \tanh f_t$, $f_{t+1} = \omega + A s_t + \beta_{GAS} f_t$, $f_1 = \omega / (1 - \beta_{GAS})$; with x, y the t_ν quantiles and $Q = x^2 + y^2 - 2\rho xy$:

$$\ln c = \text{const} - \frac{1}{2} \ln(1 - \rho^2) - \frac{\nu+2}{2} \ln\left(1 + \frac{Q}{\nu(1-\rho^2)}\right) + \frac{\nu+1}{2} \sum \ln(1 + x^2/\nu)$$
- **Tasks**
 1. Derive $s_t = (1 - \rho_t^2) \partial \ln c / \partial \rho$ by the chain rule and check it numerically against a finite difference.
 2. Estimate $(\omega, A, \beta_{GAS}, \nu)$ under $0 \leq A \leq 1$, $0 \leq \beta_{GAS} < 1$, $\nu > 2$, starting from the static fit.
 3. Compute the LR against the static t copula and its p by a bootstrap under the static copula ($B_{boot} = 199$, simulate, re-rank, refit both); plot ρ_t against the static ρ .

C3 [Proposed]: Global Common Volatility, a Project Brief — Task

- ▣ **Worked example to follow:** B1, C1 and the lecture case study ([Engle & Campos-Martins \(2023\)](#)): $e_{it}^2 = (s_i x_t + 1 - s_i) \varepsilon_{it}^2$, x_t the squared global COVOL, s_i the loadings
- ▣ **Data:** adjusted closes of 20 US-listed ETFs from the course market data (11 sector ETFs; Russell 1000 Growth and Value; Russell 2000; developed and emerging markets outside the US; gold; long Treasuries; corporate bonds; crude oil): XLB, XLC, XLE, XLF, XLI, XLK, XLP, XLRE, XLU, XLV, XLY, IWF, IWD, IWM, EFA, EEM, GLD, TLT, LQD, USO, each from its first day to 18 September 2026 (unbalanced panel); the VIX index, daily closes
- ▣ **Tasks**
 1. Filter: daily log returns; AR(1) mean with PC1 of the returns as factor (state how you form PC1 on the unbalanced panel); GARCH(1,1); keep the standardised residuals of the ETFs and of PC1.
 2. Equal loadings: $\hat{x}_t = \frac{1}{N} \sum_i e_{it}^2$ over the series observed that day; list its 20 largest days before and after 1 March 2021 with mean return and an event label.
 3. Compare the average correlation of the residuals with that of their squares; then, as a team project, estimate unequal loadings by the alternating algorithm (Section 5.3 of the paper), the test ξ and the VIX regression.
- ▣ **Report:** the two top-20 tables, the correlation comparison and, for the project, loadings, ξ and the regression with standard errors; interpretation: why can a day with a positive mean return rank high?

References (1/3)

- ▣ Aielli, G.P. (2013). [Dynamic conditional correlation: on properties and estimation](#). *Journal of Business & Economic Statistics*, 31(3), 282–299.
- ▣ Creal, D., Koopman, S.J., Lucas, A. (2013). [Generalized autoregressive score models with applications](#). *Journal of Applied Econometrics*, 28(5), 777–795.
- ▣ Demarta, S., McNeil, A.J. (2005). [The t copula and related copulas](#). *International Statistical Review*, 73(1), 111–129.
- ▣ Engle, R.F. (2002). [Dynamic conditional correlation: a simple class of multivariate generalized autoregressive conditional heteroskedasticity models](#). *Journal of Business & Economic Statistics*, 20(3), 339–350.
- ▣ Engle, R.F., Campos-Martins, S. (2023). [What are the events that shake our world? Measuring and hedging global COVOL](#). *Journal of Financial Economics*, 147(1), 221–242.
- ▣ Engle, R.F., Ledoit, O., Wolf, M. (2019). [Large dynamic covariance matrices](#). *Journal of Business & Economic Statistics*, 37(2), 363–375.

References (2/3)

- ▣ Forbes, K.J., Rigobon, R. (2002). No contagion, only interdependence: measuring stock market comovements. *Journal of Finance*, 57(5), 2223–2261.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- ▣ FTSE Russell (2019). *FTSE Russell announces results of country classification review* (press release, 26 September 2019).
- ▣ Genest, C., Ghoudi, K., Rivest, L.-P. (1995). A semiparametric estimation procedure of dependence parameters in multivariate families of distributions. *Biometrika*, 82(3), 543–552.
- ▣ Genest, C., Rivest, L.-P. (1993). Statistical inference procedures for bivariate Archimedean copulas. *Journal of the American Statistical Association*, 88(423), 1034–1043.
- ▣ Genest, C., Rémillard, B., Beaudoin, D. (2009). Goodness-of-fit tests for copulas: a review and a power study. *Insurance: Mathematics and Economics*, 44(2), 199–213.

References (3/3)

- ▣ Longin, F., Solnik, B. (2001). [Extreme correlation of international equity markets](#). *Journal of Finance*, 56(2), 649–676.
- ▣ Nelsen, R.B. (2006). [An Introduction to Copulas](#), 2nd ed. Springer.
- ▣ Newey, W.K., McFadden, D. (1994). [Large sample estimation and hypothesis testing](#). In R.F. Engle, D.L. McFadden (eds.), *Handbook of Econometrics*, Vol. 4, 2111–2245. Elsevier.
- ▣ Politis, D.N., Romano, J.P. (1994). [The stationary bootstrap](#). *Journal of the American Statistical Association*, 89(428), 1303–1313.
- ▣ Wied, D., Krämer, W., Dehling, H. (2012). [Testing for a change in correlation at an unknown point in time using an extended functional delta method](#). *Econometric Theory*, 28(3), 570–589.