



Modelling Financial Markets

Seminar 7: Value-at-Risk and Expected Shortfall — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material

- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model)
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ BNR: Banca Națională a României (the National Bank of Romania)
- ▣ BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- ▣ CF: Cornish–Fisher (expansion)
- ▣ CI: Confidence Interval
- ▣ COVID-19: COronaVirus Disease 2019
- ▣ DM: Diebold–Mariano (test)
- ▣ ES: Expected Shortfall
- ▣ EVT: Extreme Value Theory
- ▣ FHS: Filtered Historical Simulation
- ▣ FRTB: Fundamental Review of the Trading Book
- ▣ FZ0: Fissler–Ziegel loss, zero-homogeneous version
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GPD: Generalised Pareto Distribution
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HS: Historical Simulation
- ▣ LGD: Loss Given Default
- ▣ QML: Quasi-Maximum Likelihood
- ▣ RMSFE: Root Mean Squared Forecast Error
- ▣ SAV: Symmetric Absolute Value (CAViaR specification)
- ▣ SE: Standard Error
- ▣ US: United States

Today's Question and Required Outputs

- Question of the seminar: **how large are tomorrow's tail losses, and how precisely can we tell?**
- Route
 - ▶ A1–A4: VaR and ES on paper, from the Normal distribution to a discrete default loss (30 min)
 - ▶ B1, B6: historical VaR and ES on BET with bootstrap intervals; B2: an extreme-value tail (45 min)
 - ▶ B3: tomorrow's filtered VaR for Bitcoin; A5, B5: who carries the portfolio risk (40 min)
 - ▶ C2: find the errors in an AI answer (15 min)
- Required outputs (checklist on the last core slide): the A answers, one table and one figure per B exercise, the corrected C2 code, one interpretation sentence per exercise
- Notebook: [open the seminar notebook in Colab](#); each task names its section
- Seminar: Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Model
A1, A2	VaR and ES under the Normal and the Student- t distributions	Solved	–
A3	VaR and ES of a discrete default loss	Solved	–
A4	when VaR punishes diversification	Proposed	A3
B1	historical VaR and ES on BET, with bootstrap intervals	Solved	–
B6	how precise is an ES estimate?	Proposed	B1
B2	a GPD tail for BET	Solved	–
B3	tomorrow's filtered VaR for Bitcoin	Proposed	B1
A5	risk contributions of two assets	Solved	–
B5	risk contributions of a four-asset portfolio	Proposed	A5
C2	what is wrong in an AI answer?	Proposed	B1
Extension packs	A2, A4, A5, B1, B2, B3, B5 Extended; A6–A8, B4, B7; projects C1, C3 (optional, after the seminar)	both	named on each task

- ▣ **[Solved]**: full solution in the slides and the notebook, a model to copy; **[Proposed]**: you solve it, the solution is discussed in class
- ▣ Part A: derivations on paper; Part B: estimation and inference on real data, each ending with an interpretation question; Part C: AI critique and open projects

Data Manifest

Series	Symbol, field	Sample	Calendar	Used in
BET	BET, close	2000–2026	weekdays	B1, B2, B6, C2
S&P 500	GSPC.INDX, close	2000–2026	US trading days	B6
Bitcoin	BTC-USD.CC, close	2014–2026	every day	B3
SPY, TLT, GLD	SPY.US, TLT.US, GLD.US, adjusted close	2014–2026	common days with Bitcoin	B5
Gold; EUR/RON	XAUUSD.FOREX, close; BNR reference rate	2000–2026; 2005–2026	weekdays; Romanian working days	B4
BVB blue chips	TLV.RO, SNP.RO, . . . , adjusted close	2016–2026	BVB trading days	C1

- Market files: `data/market/<SYMBOL>.csv`, columns `date`, `open`, `high`, `low`, `close`, `adjusted_close`, `volume`; all samples end on 18 September 2026
- EUR/RON: the official BNR reference rate, not the EODHD series `EURRON.FOREX` (it contains erroneous quotes, Seminar 0, B6)
- Two or more assets: join the *prices* on common days first, then compute returns

Setup: Run and Verify

```
# notebook section "Setup"
import numpy as np, pandas as pd
URL = ('https://raw.githubusercontent.com/'
       'danpele/MFM/main/data/market/')
px = pd.read_csv(URL + 'BET.csv', index_col='date',
                 parse_dates=True)['close']
px = px.loc['2000-01-01':'2026-09-18']
px = px[px.index.dayofweek < 5]
r = 100 * np.log(px).diff().dropna()    # % log returns
L = -r                                  # losses
print(px.head(3)); print(len(r), r.index[-1])
```

- Expected output: $n = 6\,670$ returns, last date 2026-09-18

date	close	r_t (%)
2000-01-05	460.67	–
2000-01-06	484.04	4.95
2000-01-07	526.34	8.38

- Last close: 32757.81 on 2026-09-18
- If your count or first value differs, check the symbol, the field and the weekday filter
- Other series: change only the symbol and the field (Data Manifest)

Notation and Conventions Used in Every Exercise

- ▣ Return X (in %), loss $L = -X$; VaR and ES are **positive** numbers when they are losses
- ▣ Level α = tail probability: VaR 1%, ES 2.5%; $\text{VaR}_\alpha(X) = -\inf\{x : P(X \leq x) > \alpha\} = -q_\alpha(X)$,
 $\text{ES}_\alpha(X) = -\frac{1}{\alpha} \int_0^\alpha q_u(X) du$
- ▣ Empirical convention (all B exercises)
 - ▶ $\widehat{\text{VaR}}_\alpha = \text{np.quantile}(L, 1 - \alpha)$: linear interpolation at position $1 + (n - 1)(1 - \alpha)$ of the sorted losses
 - ▶ $\widehat{\text{ES}}_\alpha = \text{mean of the } k \text{ losses } \geq \widehat{\text{VaR}}_\alpha, k \approx n\alpha$: an approximation of the exact ES, equal to it when $n\alpha$ is an integer
- ▣ Units: log losses in % of the position for one asset; simple returns for portfolios (they add across assets); percentage points (pp) for standard errors and interval widths
- ▣ Symbols that look alike: α = tail probability, α_1 = GARCH coefficient; ξ = GPD shape, $\hat{\xi}_\alpha$ = residual quantile (B3); w = money exposures (A5) or portfolio weights (B5)

A1 [Solved]: Normal VaR and ES — Task

- ▣ **Question:** what are VaR and ES of a position whose daily return is Normal?
- ▣ **Data**
 - ▶ $X \sim N(\mu, \sigma^2)$, daily log return in % of the position, $\mu = 0.04\%$, $\sigma = 1.2\%$; loss $L = -X$ with density f
 - ▶ $z_\alpha = \Phi^{-1}(\alpha) < 0$; Φ , φ : distribution and density functions of the standard Normal distribution
- ▣ **Tasks**
 1. Derive $ES_\alpha = -\mu + \sigma\varphi(z_\alpha)/\alpha$ from $ES_\alpha = -\mathbb{E}[X \mid X \leq q_\alpha(X)]$.
 2. Compute z_α and $\varphi(z_\alpha)$ at $\alpha = 1\%$ and $\alpha = 2.5\%$.
 3. Compute VaR and ES at both levels, in % of the position.
- ▣ **Report:** a four-cell table (VaR, ES $\times$ 1%, 2.5%)
- ▣ **Interpretation:** which of the four numbers are almost equal, and why does that matter for regulation?
- ▣ Notebook: section A1

A1 [Solved]: Solution

- **1. Derivation.** Write $X = \mu + \sigma Z$; $X \leq q_\alpha$ is $Z \leq z_\alpha$, of probability α

► Since $\varphi'(u) = -u\varphi(u)$:

$$\int_{-\infty}^z u\varphi(u) du = [-\varphi(u)]_{-\infty}^z = -\varphi(z)$$

► $\mathbb{E}[X \mid X \leq q_\alpha] = \mu + \frac{\sigma}{\alpha} \int_{-\infty}^{z_\alpha} u\varphi(u) du = \mu - \sigma \frac{\varphi(z_\alpha)}{\alpha};$

change the sign

- **2–3.** $\text{VaR}_\alpha = -(\mu + \sigma z_\alpha)$, e.g.
 $-(0.04 + 1.2 \cdot (-2.326)) = 2.75$;
 $\text{ES}_{1\%} = -0.04 + 1.2 \cdot 2.665 = 3.16$

α	z_α	$\varphi(z_\alpha)$	φ/α	VaR_α	ES_α
1%	-2.326	0.0267	2.665	2.75	3.16
2.5%	-1.960	0.0584	2.338	2.31	2.77

- VaR and ES in % of the position; on 1 million: VaR 1% $\approx 27\,516$, ES 1% $\approx 31\,583$
- Interpretation: ES 2.5% (2.77%) $\approx$ VaR 1% (2.75%): under the Normal distribution the switch to ES 2.5% keeps the capital, which is why FRTB chose this level

A2 [Solved]: Student- t ES — Task

▣ **Question:** how much do heavy tails change VaR and ES when the standard deviation is the same as in A1?

▣ **Data**

► $X = m + sT$, $T \sim t_\nu$ with density g_ν , $t_\alpha = t_\nu^{-1}(\alpha) < 0$; $\nu = 4$, $m = 0.04\%$, standard deviation 1.2%

► Tail integral (proof in Extension Pack 1): $\int_{-\infty}^t x g_\nu(x) dx = -g_\nu(t) \frac{\nu + t^2}{\nu - 1}$ for $\nu > 1$

▣ **Tasks**

1. Use the tail integral to write $ES_\alpha = -m + s \frac{g_\nu(t_\alpha)}{\alpha} \frac{\nu + t_\alpha^2}{\nu - 1}$.
2. Find the scale s that gives a standard deviation of 1.2%, using $\text{Var}(T) = \nu/(\nu - 2)$ (finite only for $\nu > 2$).
3. Compute VaR and ES at 1% and 2.5% and put them next to A1.

▣ **Report:** the A1 table with a Student- t column

▣ **Interpretation:** at which level do the two models differ most, and why does ES see the difference first?


▣ Common mistake: using $\sigma = 1.2$ instead of s as the scale

A2 [Solved]: Solution

- 1. Divide the tail integral by α and rescale with $X = m + sT$, as in A1
- 2. $s^2 \cdot 4/2 = 1.2^2$, so $s = 1.2\sqrt{2/4} = 0.849$
- 3. At 1%: $t_{1\%} = -3.747$, $g_4(t_{1\%}) = 0.0087$, ES factor $\frac{0.0087}{0.01} \cdot \frac{4 + 14.04}{3} = 5.221$
 - $\text{VaR}_{1\%} = -(0.04 + 0.849 \cdot (-3.747)) = 3.14$; $\text{ES}_{1\%} = -0.04 + 0.849 \cdot 5.221 = 4.39$

	$\text{VaR}_{2.5\%}$	$\text{ES}_{2.5\%}$	$\text{VaR}_{1\%}$	$\text{ES}_{1\%}$
Normal	2.31	2.77	2.75	3.16
t_4	2.32	3.35	3.14	4.39

A2 [Solved]: Solution — Chart



- Interpretation: VaR 2.5% is the same in both models; VaR 1% and both ES are much higher under t_4 : heavy tails show only deep in the tail, and ES sees them first

A3 [Solved]: VaR and ES of a Discrete Loss — Task

- **Question:** what are VaR and ES when the loss has an atom (a default)?
- **Data:** a bond loses 60 monetary units (loss given default, LGD) with probability $p = 3\%$ and 0 otherwise
- **Tasks**
 1. Compute VaR 5% and VaR 2.5% from $\text{VaR}_\alpha = -\inf\{x : P(X \leq x) > \alpha\}$.
 2. Compute ES 5% and ES 2.5% with $\text{ES}_\alpha = \frac{1}{\alpha} (\mathbb{E}[L1\{L > \text{VaR}_\alpha\}] + \text{VaR}_\alpha(\alpha - P(L > \text{VaR}_\alpha)))$ Acerbi & Tasche (2002).
 3. Explain why $\text{ES}_{5\%}$ differs from $\mathbb{E}[L \mid L \geq \text{VaR}_{5\%}]$.
- **Report:** four numbers and one sentence


Loss L	0	60
$P(L)$	97%	3%
$P(L \geq \cdot)$	100%	3%

- The worst α of outcomes must be filled with probability mass; at $\alpha = 5\%$ the 3% default mass is not enough

A3 [Solved]: Solution

- 1. $P(X \leq -60) = 3\% \leq 5\%$, $P(X \leq 0) = 100\% > 5\%$: $q_{5\%} = 0$, $\text{VaR}_{5\%} = 0$; $P(X \leq -60) = 3\% > 2.5\%$: $\text{VaR}_{2.5\%} = 60$
- 2. $\text{ES}_{5\%} = \frac{60 \cdot 0.03 + 0 \cdot (0.05 - 0.03)}{0.05} = 36$; $\text{ES}_{2.5\%} = \frac{0 + 60 \cdot (0.025 - 0)}{0.025} = 60$
- 3. $\mathbb{E}[L \mid L \geq 0] = 1.8$ averages over **all** outcomes; ES 5% uses the 3% default mass plus 2% of the zero-loss atom (chart)
- Interpretation: this fractional mass is the exact version of the tail average used in B1 with $k \approx n\alpha$ losses

A3 [Solved]: Solution — Chart




A4 [Proposed]: When VaR Punishes Diversification — Task

- **Question:** can a VaR limit make a diversified portfolio look riskier than its parts?
- **Data:** two independent bonds A and B; each loses 100 with probability $p = 0.7\%$ (the lecture used 0.9%) and 0 otherwise; portfolio A + B [Artzner et al. \(1999\)](#)
- **Tasks**
 1. Write the distribution of the portfolio loss: the values 0, 100, 200 and their probabilities.
 2. Compute VaR 1% and ES 1% of one bond and of the portfolio with the discrete formula of A3.
 3. Show that VaR 1% is not subadditive here while ES 1% is.
 4. Repeat steps 2–3 at $\alpha = 1.5\%$.
- **Report:** a table of VaR and ES for A, B and A + B at both levels
- **Interpretation:** what would a VaR-based limit reward?
- **Model:** A3

B1 [Solved]: Historical VaR and ES on BET — Task

- **Question:** how large are the tail losses of the Romanian stock market, and how precise are the estimates?
- **Data**
 - ▶ BET index, data/market/BET.csv, field close; closes 2000-01-05 – 2026-09-18, weekdays only; $n = 6\,670$ daily log returns in %, $r_t = 100 (\ln P_t - \ln P_{t-1})$, loss $L_t = -r_t$
 - ▶ Two samples: the full period and the last 500 returns (from 2024-09-12, shaded in the chart)
- **Tasks**
 1. Compute historical VaR and ES at 1% and 2.5% on both samples, and count the losses beyond VaR 1%.
 2. Build a 95% percentile bootstrap interval for VaR 1% with 2 000 i.i.d. resamples, seed 42 [Efron \(1979\)](#).
 3. Repeat step 2 with a moving-block bootstrap: 20-day blocks drawn with replacement, concatenated and cut to length n [Künsch \(1989\)](#).
- **Report:** a table of the eight VaR/ES numbers and the two intervals per sample
- **Interpretation:** which estimate would you give a risk committee, and why is the block interval wider on the full sample?

B1 [Solved]: Historical VaR and ES on BET — Data



□ Look first: the largest losses (October 2008: 10.45%; March 2020: 10.08%) come in clusters


B1 [Solved]: Solution (1/3) — Step by Step

- **1. Last 500 days.** Sort the losses; VaR 1% sits at position $1 + 499 \cdot 0.99 = 495.01$
 - ▶ $\widehat{\text{VaR}}_{1\%} = L_{(495)} + 0.01 (L_{(496)} - L_{(495)}) = 2.710 + 0.01 (2.925 - 2.710) = 2.712\%$
 - ▶ The five losses beyond it: 3.68, 3.35, 3.12, 3.03, 2.93: their mean is $\widehat{\text{ES}}_{1\%} = 3.22\%$

	n	$\text{VaR}_{1\%}$	$\text{ES}_{1\%}$	$\text{VaR}_{2.5\%}$	$\text{ES}_{2.5\%}$	beyond VaR 1%
2000–2026	6 670	4.13	6.28	2.80	4.57	67
Last 500 days	500	2.71	3.22	2.05	2.69	5


- The 500-day ES 1% averages only five numbers: report it with that count
- Both are unconditional summaries: the full sample describes 26 years, the last 500 days a calm period; tomorrow's risk needs a conditional forecast (B3)

B1 [Solved]: Solution (2/3) — The Tail and the Bootstrap



- Left: BET losses above 1.5% with historical VaR 1% and ES 1%; right: the 2000 i.i.d. bootstrap values of VaR 1% for each sample
- Each resample recomputes both the quantile and which losses lie beyond it; the 2.5% and 97.5% percentiles of the 2000 values give the interval
- The two bootstrap distributions do not overlap: the recent market is calmer than the 26-year average

B1 [Solved]: Solution (3/3) — Intervals



Dot: historical VaR 1%; bar: 95% interval; number: interval width in percentage points

- 2. i.i.d. bootstrap: full [3.85, 4.46], last 500 days [2.11, 3.12]
- 3. Block bootstrap: [3.61, 4.80] and [2.02, 3.03]

- Interpretation: exceedances cluster in 2008 and 2020, so blocks that keep a crisis together widen the full-sample interval by about half (in B1 Extended: asymptotic and HAC); the calm last 500 days show no such effect

B6 [Proposed]: How Precise Is an ES Estimate? — Task

- ▣ **Question:** does a 500-day historical ES 2.5% pin the tail risk down to within ± 0.3 pp [Pele & Mazurencu-Marinescu-Pele \(2026\)](#)?
- ▣ **Data**
 - ▶ S&P 500 (GSPC.INDX, close) and BET, daily log losses in %, 2000–2026; samples: full and last 500 days
 - ▶ Use the resampling function of B1 (2 000 resamples, seed 42, 20-day blocks) with ES 2.5% as the statistic
- ▣ **Tasks**
 1. Compute ES 2.5%, the tail count $n\alpha$ and the number k of losses that the tail mean actually averages.
 2. Draw one i.i.d. and one block resample of the S&P 500 last 500 days, and recompute VaR 2.5%, the tail membership and ES on each.
 3. Build the i.i.d. and the block 95% intervals for all four samples and compare their widths with 0.6 pp (± 0.3).
- ▣ **Report:** a table ($n\alpha$, k , ES, two intervals) and an interval chart
- ▣ **Interpretation:** which samples meet the ± 0.3 pp target, and what governs the width?
- ▣ **Model:** B1

B2 [Solved]: A GPD Tail for BET — Task

- ▣ **Question:** what does an extreme-value tail add beyond historical simulation, deep in the tail?
- ▣ **Data**
 - ▶ BET index, data/market/BET.csv, field close; closes 2000-01-05 – 2026-09-18, weekdays only; $n = 6\,670$ daily log returns in %, $r_t = 100 (\ln P_t - \ln P_{t-1})$, loss $L_t = -r_t$
 - ▶ Peaks over threshold: excesses $Y = L - u$ above a high threshold u follow a generalised Pareto distribution (GPD) with shape ξ and scale β [McNeil, Frey & Embrechts \(2015\)](#); n = sample size, N_u = number of exceedances
- ▣ **Tasks**
 1. Set u = the empirical 95% quantile of L and form the excesses $L - u$ of the losses above it.
 2. Fit the GPD by maximum likelihood with location fixed at 0: `scipy.stats.genpareto.fit(excesses, floc=0)`.
 3. Compute the i.i.d. standard errors $\text{SE}(\hat{\xi}) = (1 + \hat{\xi})/\sqrt{N_u}$ and $\text{SE}(\hat{\beta}) = \hat{\beta}\sqrt{2(1 + \hat{\xi})/N_u}$.
 4. Compute GPD VaR and ES at 1%, 0.5% and 0.1% with $\text{VaR}_\alpha = u + \frac{\beta}{\xi} [(\frac{n}{N_u}\alpha)^{-\xi} - 1]$, $\text{ES}_\alpha = \frac{\text{VaR}_\alpha + \beta - \xi u}{1 - \xi}$, and compare with historical simulation.
 5. Refit for thresholds leaving 10% down to 1.5% of the losses above them and plot $\hat{\xi}$ with 95% bands.
- ▣ **Report:** the fit with standard errors, a six-row VaR/ES table with tail counts, and the stability chart
- ▣ **Interpretation:** is the BET tail heavier than the Normal distribution, and how sure are we?


B2 [Solved]: Solution (1/2) — Fit and Table

- 1–2. $u = 1.93\%$ leaves $N_u = 334$ of $n = 6\,670$ losses above it; $\hat{\xi} = 0.209$, $\hat{\beta} = 1.20$
- 3. $SE(\hat{\xi}) = (1 + 0.209)/\sqrt{334} = 0.066$, $SE(\hat{\beta}) = 0.10$; $\hat{\xi}/SE = 3.2$
- 4. VaR 1%: $N_u/n = 0.050$, $(0.01/0.050)^{-0.209} = 1.401$
 - ▶ $VaR_{1\%} = 1.93 + \frac{1.20}{0.209}(1.401 - 1) = 4.22\%$

	VaR GPD	VaR hist.	ES GPD	ES hist.
1%	4.22	4.13	6.34	6.28
0.5%	5.47	5.38	7.92	7.86
0.1%	9.18	9.50	12.61	11.25

- Losses beyond historical VaR: 67 at 1%, 34 at 0.5%, 7 at 0.1%
- The two methods agree down to 0.5%; at 0.1% historical ES averages 7 numbers, while the GPD uses all 334 excesses through a fitted tail

B2 [Solved]: Solution (2/2) — Diagnostics



- Left: the GPD follows the empirical tail, the Normal tail falls far below it; right: $\hat{\xi}$ ranges from 0.15 to 0.30 across thresholds, while VaR 0.1% moves only between 9.04% and 9.51%
- Interpretation: the tail is heavier than the Normal one, but the i.i.d. SE assumes independent exceedances; BET losses cluster, so the honest uncertainty is larger (B2 Extended: extremal index)

B3 [Proposed]: Tomorrow's Filtered VaR for Bitcoin — Task

□ **Question:** Bitcoin's VaR 1% for 19 September 2026: filtered or historical?

□ **Data**

- ▶ BTC-USD.CC, close, 17 September 2014 – 18 September 2026, 7 days a week; daily log returns in %
- ▶ AR(1)-GARCH(1,1) by quasi-maximum likelihood (QML): $r_t = \mu_t + \sigma_t z_t$ Barone-Adesi et al. (1999)

□ **Tasks**


1. Estimate the model on the full sample and forecast $\hat{\mu}_{t+1}$, $\hat{\sigma}_{t+1}$.
2. Compute $\hat{\xi}_{1\%} = -q_{1\%}(\hat{z})$ from the standardised residuals and the filtered VaR 1% and ES 1%.
3. Compare with historical VaR and ES on the last 500 days and on the full sample.
4. With the forecasts of the timeline (right), count the days of the last two years with a loss above the FHS VaR and above a rolling 500-day HS VaR.

□ **Report:** three VaR/ES pairs, two exceedance counts

□ **Interpretation:** why do the three VaRs differ, and which method misses more often?

□ **Model:** B1

Forecast timeline (step 4)



□ No forecast uses information from its own day or later: a full-sample threshold is not an out-of-sample forecast

A5 [Solved]: Risk Contributions of Two Assets — Task

□ **Question:** does each asset carry risk in proportion to the money invested in it?

□ **Data**

► Money exposures $w = (600\,000, 400\,000)$ EUR in equity and bonds; daily simple returns R in decimals, $\sigma_1 = 0.012$, $\sigma_2 = 0.009$, $\rho = -0.2$

► $\Sigma = \begin{pmatrix} 0.000144 & -0.0000216 \\ -0.0000216 & 0.000081 \end{pmatrix}$; jointly Normal returns, mean ignored, $\alpha = 1\%$, $z_{99\%} = 2.326$

□ **Tasks**

1. Compute $\sigma_p = \sqrt{w^\top \Sigma w}$ in EUR and the portfolio VaR $1\% = z_{99\%} \sigma_p$.
2. Compute Σw , the marginal VaR $z_{99\%}(\Sigma w)_i / \sigma_p$ per EUR and the component VaR $C_i = w_i \times \text{marginal VaR}$; check $C_1 + C_2 = \text{VaR}$ [Tasche \(2008\)](#).
3. Compute the diversification benefit: the sum of stand-alone VaRs minus the portfolio VaR.
4. Under the Normal distribution, rescale each C_i by $\frac{\varphi(z_{1\%})/0.01}{z_{99\%}}$ to get the ES 1% contributions.


□ **Report:** a table of C_i , C_i^{ES} and shares, and the diversification benefit in EUR

□ **Interpretation:** do the bonds carry risk in proportion to their 40% money weight?

A5 [Solved]: Solution

- 1. $\sigma_p^2 = 7\,200^2 + 3\,600^2 + 2(-0.2)(7\,200)(3\,600) = 54\,432\,000$, $\sigma_p = 7\,378$ EUR, $\text{VaR}_{1\%} = 2.326 \cdot 7\,378 = 17\,163$ EUR
- 2. $\Sigma w = (77.76, 19.44)$; marginal VaR per EUR (0.0245, 0.0061); $C_1 = 14\,711$ (85.7%), $C_2 = 2\,452$ (14.3%), sum 17 163 EUR
- 3. Stand-alone $16\,750 + 8\,375 = 25\,125$ EUR vs 17 163 EUR
- 4. Factor 1.146: $C_1^{ES} = 16\,854$, $C_2^{ES} = 2\,809$, total 19 663 EUR = ES 1%; same shares as VaR under the Normal distribution

A5 [Solved]: Solution — Chart



□ Interpretation: the bonds hold 40% of the money but 14.3% of the risk: the negative correlation makes them a partial hedge

B5 [Proposed]: Risk Contributions of a Four-Asset Portfolio — Task

- **Question:** in an equally weighted portfolio of SPY, TLT, GLD and Bitcoin, who carries the risk?
- **Data**
 - ▶ SPY.US, TLT.US, GLD.US (adjusted close) and BTC-USD.CC (close), 18 September 2014 – 18 September 2026; prices joined on the 3017 common days, then daily simple returns R in %
 - ▶ Fractional weights $w_i = 25\%$, rebalanced daily; Σ = sample covariance of R (in $\%^2$); Normal VaR 1% with the mean ignored; historical ES 2.5%
- **Tasks**
 1. Compute the Normal VaR 1%, the marginal VaR and the component VaR of each asset, as in A5.
 2. Compute the historical ES 2.5% of the portfolio loss $L_p = -w^\top R$ and the contributions $\mathbb{E}[-w_i R_i \mid L_p \geq \text{VaR}_{2.5\%}]$ as averages over the tail days.
 3. Compute the undiversified VaR $\sum_i z_{99\%} w_i \sigma_i$.
- **Report:** a table of weights, contributions and shares, and a chart of money weights vs risk shares
- **Interpretation:** is equal money equal risk? (comparing Normal VaR 1% with historical ES 2.5% changes both the model and the level)
- **Model:** A5

C2 [Proposed]: Critique an AI Answer — the Answer

- **Prompt:** "px holds daily BET closes, 2000–2026. Give me the one-day historical VaR 1% and ES 2.5% as positive log losses in %, and the 10-day VaR 1%."
- **Answer:** the code below, which runs without an error message

```
1  import numpy as np
2  r = 100 * np.log(px).diff().dropna()           # daily log returns, %
3  loss = -r                                       # log loss, %
4  var1 = np.percentile(r, 99)                   # VaR 1%
5  es25 = loss[loss >= np.percentile(loss, 99)].mean() # ES 2.5%
6  var10 = 10 * var1                             # 10-day VaR 1%
```

- **Conclusion of the AI:** "VaR 1% = 3.95%, ES 2.5% = 6.28%, 10-day VaR 1% = 39.46%."

C2 [Proposed]: Critique an AI Answer — Task

- ▣ **Question:** which lines of the AI answer are wrong, and what are the correct numbers?
- ▣ **Data:** the BET data of B1; the answer contains **three** mistakes typical for VaR and ES, everything else is correct
- ▣ **Tasks**
 1. For each mistake, name the line number and explain in one sentence why it is wrong, using $\text{VaR}_\alpha = -q_\alpha(r) = q_{1-\alpha}(L)$.
 2. Say how you would detect it: a count of observations, a number from B1, or a slide of the lecture.
 3. Write the corrected lines, including the 10-day VaR from overlapping 10-day losses, and recompute the three numbers.
 4. Check the corrected ES 2.5%: how many losses does it average, and how many should it?
- ▣ **Report:** an error table (line, error, detection, correction) and the corrected numbers; log each error you actually found in `AI_ERRORS.md`
- ▣ **Interpretation:** why does the code run without an error message although the answer is wrong?
- ▣ **Model:** B1

Submission Checklist and Exit Check

Hand in before the next seminar

- A1–A4: the tables and the derivations
- B1, B6, B2, B3, B5: one table and one figure each, as in the solutions
- C2: the error table, the corrected code and numbers
- One interpretation sentence per exercise
- Log AI use in AI_USE.md and every error you actually caught in AI_ERRORS.md (no quota)

Exit check: can you say why

- a 500-day ES 1% should be reported with the number 5?
- $\sqrt{10}$ scaling is too low for BET?
- a 25% position can carry 80% of the risk?
- portfolio losses need simple, not log returns?
- Team project: choose one of C1 or C3 (Extension Pack 4), with the common rubric

Extension Packs: What Each One Needs

Pack	Exercises	Prerequisites
1. Tail algebra	A2 Extended [Solved], A4 Extended, A6, A7, A5 Extended [Solved], A8	A1–A5
2. Inference	B1 Extended [Solved], B3 Extended, B4, B5 Extended	B1, B3, B5
3. Extremes and dynamics	B2 Extended [Solved], B7	B2, B3
4. Projects	C1 (ES capital, BVB vs S&P 500), C3 (deep quantile regression); rubric	the lecture case study, B1–B3

- Each exercise names its notebook section and its model; do them in order within a pack

A2 Extended [Solved]: The Student- t Tail Integral — Task and Solution

- **Question:** where does the Student- t ES factor of A2 come from?
- **Data:** $g_\nu(x) = c_\nu(1 + x^2/\nu)^{-(\nu+1)/2}$
- **Tasks**
 1. Differentiate $(1 + x^2/\nu)^{-(\nu+1)/2}$.
 2. Deduce $\int_{-\infty}^t x g_\nu(x) dx = -g_\nu(t) \frac{\nu+t^2}{\nu-1}$ and state where $\nu > 1$ is needed.
- **Report:** the two-line proof

Solution

- 1. $\frac{d}{dx}(1 + x^2/\nu)^{-(\nu+1)/2} = -\frac{\nu+1}{\nu} x (1 + x^2/\nu)^{-(\nu+1)/2}$, so $x g_\nu(x) = -\frac{\nu c_\nu}{\nu-1} \frac{d}{dx}(1 + x^2/\nu)^{-(\nu+1)/2}$
- 2. Integrating from $-\infty$ to t : $-\frac{\nu}{\nu-1} g_\nu(t)(1 + t^2/\nu) = -g_\nu(t) \frac{\nu+t^2}{\nu-1}$; the term at $-\infty$ vanishes only if $\nu > 1$
- ES needs $\nu > 1$; matching a variance needs $\nu > 2$

A4 Extended [Proposed]: For Which Levels Does VaR Fail? — Task

- ▣ **Question:** for which tail probabilities α is VaR not subadditive in A4?
- ▣ **Data:** the A4 bonds, $p = 0.7\%$, and the definition $\text{VaR}_\alpha = -\inf\{x : P(X \leq x) > \alpha\}$
- ▣ **Tasks**
 1. Write $\text{VaR}_\alpha(A)$ and $\text{VaR}_\alpha(A + B)$ as step functions of α .
 2. Find the interval of α where $\text{VaR}_\alpha(A + B) > 2\text{VaR}_\alpha(A)$ and explain both endpoints.
- ▣ **Report:** the interval and one sentence per endpoint
- ▣ **Model:** A4

A6 [Proposed]: When Is Cornish–Fisher a Valid Quantile? — Task

□ **Question:** for which excess kurtosis is the Cornish–Fisher map still a quantile function?

□ **Data**

▶ $\mu = 0$, $\sigma = 1.2\%$, skewness $S = -0.5$, excess kurtosis K [Cornish & Fisher \(1938\)](#)

▶ $\tilde{z}(z) = z + \frac{(z^2-1)S}{6} + \frac{(z^3-3z)K}{24} - \frac{(2z^3-5z)S^2}{36}$, $\text{VaR}_\alpha = -\sigma\tilde{z}(z_\alpha)$

□ **Tasks**

1. With $K = 3$, compute the three correction terms at $z_{1\%}$, $\tilde{z}_{1\%}$ and the Cornish–Fisher VaR 1%.
2. Write $d\tilde{z}/dz = Az^2 + Bz + C$ and locate the vertex of this parabola.
3. Find the largest K for which $\tilde{z}$ is increasing on $[z_{0.1\%}, 0]$, and decide whether $K = 10$ is admissible.

□ **Report:** the VaR, the bound on K and a plot of $\tilde{z}$ and $d\tilde{z}/dz$ for $K = 3$ and $K = 10$

□ **Interpretation:** does a valid quantile map also give an accurate VaR?

□ **Model:** A1

A7 [Proposed]: ES as a Minimum, and Subadditivity — Task

- **Question:** why can ES never reward splitting a portfolio?
- **Data**
 - ▶ Integrable loss L ; $(x)_+ = \max(x, 0)$; objective $G_\alpha(v) = v + \mathbb{E}[(L - v)_+]/\alpha$ Rockafellar & Uryasev (2000)
 - ▶ Numbers: the A4 portfolio (two bonds, $p = 0.7\%$) at $\alpha = 1\%$
- **Tasks**
 1. For the A4 portfolio, compute the slopes of G on $(0, 100)$ and $(100, 200)$, the minimiser and the minimum.
 2. Show that the right derivative of G is $1 - P(L > v)/\alpha$ and deduce that VaR_α minimises G , with minimum ES_α .
 3. Deduce $\text{ES}_\alpha(L_1 + L_2) \leq \text{ES}_\alpha(L_1) + \text{ES}_\alpha(L_2)$ from $(a + b)_+ \leq a_+ + b_+$.
- **Report:** the three numbers of step 1 and two short proofs
- **Model:** A3, A4

A5 Extended [Solved]: ES Contributions from the Minimum — Task and Solution

- **Question:** why are the historical ES contributions of B5 averages over the tail days?
- **Data:** $ES_\alpha(w) = \min_v G(w, v)$,
 $G(w, v) = v + \mathbb{E}[(L_p - v)_+]/\alpha$, $L_p = -w^\top R$ (A7)
- **Tasks**
 1. Use the envelope argument at $v^* = \text{VaR}_\alpha$ to compute $\partial ES_\alpha / \partial w_i$.
 2. Show that $w_i \partial ES_\alpha / \partial w_i = \mathbb{E}[-w_i R_i \mid L_p \geq \text{VaR}_\alpha]$ for a continuous L_p .
- **Report:** a short proof
- **Model:** A7

Solution

- $\partial G / \partial v = 0$ at v^* , so
 $\partial ES_\alpha / \partial w_i = \partial G / \partial w_i(w, v^*) = \mathbb{E}[-R_i \mathbf{1}\{L_p > v^*\}]/\alpha$
- For continuous L_p , $P(L_p > v^*) = \alpha$: this is $\mathbb{E}[-R_i \mid L_p \geq \text{VaR}_\alpha]$; summing $w_i \times$ it gives ES_α (Euler)
- Historical estimate: average $-w_i R_i$ over the tail days (B5); under the Normal distribution it equals the VaR rescaling of A5

A8 [Proposed]: The Sampling Variance of Historical ES — Task

- ▣ **Question:** how imprecise is a 500-day historical ES 2.5%, even for Normal losses?
- ▣ **Data**
 - ▶ n i.i.d. losses, continuous, positive density at VaR_α , $\mathbb{E}[L^2] < \infty$; $Y = (L - \text{VaR}_\alpha)_+$
 - ▶ Influence function $\text{IF}(x) = \text{VaR}_\alpha + (x - \text{VaR}_\alpha)_+ / \alpha - \text{ES}_\alpha$ [Chen \(2008\)](#); $\text{Var}(\widehat{\text{ES}}_\alpha) \approx \text{Var}(\text{IF}(L)) / n$
- ▣ **Tasks**
 1. Show that $\text{Var}(\text{IF}(L)) = \text{Var}(Y) / \alpha^2$ and explain, with A7, why estimating VaR adds no first-order term.
 2. For $L \sim N(0, 1)$, $\alpha = 2.5\%$, integrate $\mathbb{E}[Y] = \int_v^\infty (x - v)\varphi(x) dx$ and $\mathbb{E}[Y^2]$ in closed form.
 3. Compute the SE of $\widehat{\text{ES}}_{2.5\%}$ for $n = 500$ and $\sigma = 1.2\%$, and check it by simulation.
- ▣ **Report:** the two moments, the SE in pp and the simulated standard deviation
- ▣ **Interpretation:** why is the S&P 500 value (B6) much larger?
- ▣ **Model:** A1, A7


B1 Extended [Solved]: Asymptotic and HAC Intervals for VaR — Task

- ▣ **Question:** do the asymptotic formulas agree with the bootstrap of B1?
- ▣ **Data**
 - ▶ i.i.d.: $SE = \sqrt{\alpha(1-\alpha)/n} / \hat{f}(\widehat{VaR})$, $\hat{f}$ = Gaussian kernel density of L (Silverman bandwidth)
 - ▶ HAC: replace $\alpha(1-\alpha)$ by the Newey–West long-run variance of $l_t = 1\{L_t > \widehat{VaR}\}$ (Bartlett kernel, lag $\lfloor 4(n/100)^{2/9} \rfloor$)
- ▣ **Tasks**
 1. Normal benchmark (from A1): with $f(VaR_{1\%}) = \varphi(z_{1\%})/1.2$, compute the SE for $n = 500$ and $n = 6\,670$.
 2. For BET, compute $\hat{f}$, the i.i.d. and HAC SEs and the two 95% intervals on both samples.
- ▣ **Report:** the four intervals of VaR 1% next to the two bootstrap intervals of B1
- ▣ **Interpretation:** why are the BET intervals wider than the Normal benchmark?

B1 Extended [Solved]: Solution

- ▣ 1. $f = 0.0267/1.2 = 0.0222$: SE 0.20 pp ($n = 500$, width 0.79), 0.055 pp ($n = 6\,670$, width 0.22)
- ▣ 2. BET full: $\hat{f} = 0.0076$, SE 0.16 (i.i.d.), 0.23 (HAC); intervals $[3.81, 4.44]$, $[3.69, 4.57]$
- ▣ Last 500 days: $\hat{f} = 0.0163$, SE 0.27 and 0.27
- ▣ Interpretation: the BET tail is heavier than the Normal one, so the density at VaR is lower and the quantile less precise; HAC and blocks agree that clustering widens the full-sample interval

B1 Extended [Solved]: Solution — Chart



Dot: historical VaR 1%; bar: 95% interval; number: interval width in percentage points

B3 Extended [Proposed]: A Bootstrap Interval for Tomorrow's FHS VaR — Task

- ▣ **Question:** how uncertain is the filtered VaR of B3?
- ▣ **Data:** the B3 model and data; residual bootstrap [Christoffersen & Gonçalves \(2005\)](#) with 999 paths, seed 42, routine `fhs_bootstrap` of the Quantlet
- ▣ **Tasks**
 1. For each path, simulate an AR(1)-GARCH(1,1) series of the sample length with resampled $\hat{z}$ and re-estimate $\theta^* = (c^*, \phi^*, \omega^*, \alpha_1^*, \beta_1^*)$.
 2. Compute μ_{t+1}^* and σ_{t+1}^* by running the **original** Bitcoin returns through θ^* , and $\hat{\xi}_{1\%}^*$ from the path's own residuals.
 3. Report the 90% percentile interval of $\text{VaR}^* = -\mu_{t+1}^* + \sigma_{t+1}^* \hat{\xi}_{1\%}^*$ and the intervals of its two sources.
- ▣ **Report:** the histogram of the 999 forecasts with the interval
- ▣ **Interpretation:** does the interval include the 500-day historical VaR, and is that a test?
- ▣ **Model:** B3

B4 [Proposed]: Cornish–Fisher vs Empirical Quantiles — Task

- ▣ **Question:** does the Cornish–Fisher correction fix the Normal VaR on real data?
- ▣ **Data**
 - ▶ BET (compulsory): the B1 data; extension: GSPC.INDX, BTC-USD.CC, XAUUSD.FOREX (close, weekdays for gold) and the BNR EUR/RON reference rate (2005–2026), all to 18 September 2026
 - ▶ Moments of the returns X : mean, standard deviation with `ddof=1`, `scipy.stats.skew` and `scipy.stats.kurtosis` (excess, biased)
- ▣ **Tasks**
 1. Compute S and K for BET.
 2. Compute VaR 1% and VaR 2.5% by the Normal distribution, Cornish–Fisher (A6) and historical simulation.
 3. Build 95% i.i.d. bootstrap intervals (1 000 resamples) for the two historical VaRs and check whether the Cornish–Fisher values fall inside.
 4. Separately, check whether K exceeds the monotonicity bound of A6.
- ▣ **Report:** a small table per asset and a chart of the three methods with the bootstrap intervals
- ▣ **Interpretation:** why does Cornish–Fisher work better at 2.5% than at 1%?
- ▣ **Model:** A1, A6, B1

B5 Extended [Proposed]: Risk-Parity Weights — Task

- ▣ **Question:** which long-only, fully invested weights give each asset of B5 the same share of Normal VaR?
- ▣ **Data:** the B5 returns and covariance Σ
- ▣ **Tasks**
 1. Solve $w_i(\Sigma w)_i = w_j(\Sigma w)_j$ for all i, j , with $w \geq 0$ and $\sum_i w_i = 1$, numerically.
 2. Compare with the inverse-volatility weights $w_i \propto 1/\sigma_i$.
 3. At the risk-parity weights, compute the Normal VaR 1%, the historical ES 2.5% and the ES shares.
- ▣ **Report:** the weights, VaR, ES and shares
- ▣ **Interpretation:** does equal VaR risk mean equal ES risk?
- ▣ **Model:** B5


B2 Extended [Solved]: Clustered Exceedances and the Extremal Index — Task

- ▣ **Question:** how many independent pieces of tail information do the 334 BET exceedances of B2 contain?
- ▣ **Data**
 - ▶ Exceedance times $S_1 < \dots < S_N$ of u (B2), gaps $T_i = S_{i+1} - S_i$
 - ▶ Intervals estimator [Ferro & Segers \(2003\)](#): $\hat{\theta} = \min\left\{1, \frac{2(\sum_i (T_i - 1))^2}{(N-1) \sum_i (T_i - 1)(T_i - 2)}\right\}$ if $\max T_i > 2$
- ▣ **Tasks**
 1. Compute $\hat{\theta}$ and the target number of clusters $C = \lfloor \hat{\theta}N \rfloor + 1$.
 2. Split the exceedances at the $C - 1$ largest gaps; if the $(C - 1)$ -th and C -th largest gaps are equal, reduce C until they differ.
 3. Fit the GPD to the cluster maxima and compare $\hat{\xi}/\text{SE}$ with B2.
- ▣ **Report:** $\hat{\theta}$, the number of clusters, the declustered fit and a timeline chart
- ▣ **Interpretation:** is the i.i.d. SE of B2 credible?

B2 Extended [Solved]: Solution

- ▣ 1. $\hat{\theta} = 0.32$, $C = 108$
- ▣ 2. Tie at the cut ($T = 12$ days): 100 clusters of 3.3 exceedances on average; largest cluster 33, 47 isolated exceedances
- ▣ 3. $\hat{\xi} = 0.269$ (SE 0.127) vs 0.209 (SE 0.066); $\hat{\xi}/\text{SE}$ falls from 3.2 to 2.1; at the VaR 1% threshold $\hat{\theta} = 0.40$
- ▣ Interpretation: the tail stays heavy, but the i.i.d. SE is not credible; report the declustered fit or filter with GARCH first (conditional EVT)

B2 Extended [Solved]: Solution — Chart



- Exceedances (clusters alternate colour)
- Cluster maximum
- Empirical mean excess
- GPD VaR 0.1%

B7 [Proposed]: CAViaR on BET — Task

- ▣ **Question:** does modelling the BET quantile directly beat filtered and historical simulation out of sample?
- ▣ **Data**
 - ▶ BET index, data/market/BET.csv, field close; closes 2000-01-05 – 2026-09-18, weekdays only; $n = 6\,670$ daily log returns in %, $r_t = 100(\ln P_t - \ln P_{t-1})$, loss $L_t = -r_t$
 - ▶ SAV CAViaR [Engle & Manganelli \(2004\)](#): $\text{VaR}_t = \beta_0 + \beta_1 \text{VaR}_{t-1} + \beta_2 |r_{t-1}|$, $|\beta_1| < 1$, $\text{VaR}_1 = -\hat{q}_1\%$ of the first 300 returns; routines `caviar_path`, `caviar_fit` (Quantlet): 10^4 starts in $(0, 1)^3$, seed 42, the best 10 refined
- ▣ **Tasks**
 1. Estimate β on 2000–2015 by minimising the pinball loss $\sum_t (\alpha - 1\{r_t < -\text{VaR}_t\})(r_t + \text{VaR}_t)$ at $\alpha = 1\%$.
 2. Compute one pinball-loss term step by step for the first out-of-sample day and check that the forecast for day t uses only r_{t-1} .
 3. With β fixed, forecast VaR for 2016–2026 and compare exceedances and mean pinball loss with FHS (refitted each January) and rolling 500-day HS.
 4. For 20 February – 30 April 2020 count exceedances and report the peak VaR of each method.
- ▣ **Report:** the estimates, a comparison table and a chart with a COVID-19 zoom
- ▣ **Interpretation:** can the difference be attributed to the model class when CAViaR is fitted once and FHS every year?
- ▣ **Model:** B3; lecture, CAViaR on the S&P 500

C1 [Open]: ES Capital, BVB vs S&P 500 — Task (1/2)

- ▣ **Question:** which portfolio needs more ES capital per unit invested: Bucharest blue chips or the S&P 500?
- ▣ **Data**
 - ▶ BVB basket: TLV.R0, SNP.R0, BRD.R0, TGN.R0, SNG.R0, SNN.R0, EL.R0, TEL.R0 (adjusted close), equal weights, rebalanced daily; days on which no stock traded are removed
 - ▶ S&P 500 portfolio: SPY.US (adjusted close); both from 2016-09-20 to 18 September 2026; daily simple returns R_t in % in local currency
- ▣ **Tasks**
 1. Construction: compute the one-day historical ES 2.5% of each portfolio (loss $-R_t$) and its tail count.
 2. Horizon: compute the 10-day ES 2.5% by $\sqrt{10}$ scaling, from overlapping compounded losses $100[1 - \prod_{i=1}^{10}(1 + R_{t+i}/100)]$, and by filtered Monte Carlo (AR(1)-GARCH(1,1) on log returns, today's state, 100 000 paths, seed 42, each path converted to a simple loss).
- ▣ **Report:** continued on the next slide
- ▣ **Model:** B1, B3; lecture, horizon and Basel sections

C1 [Open]: ES Capital, BVB vs S&P 500 — Task (2/2)


▣ Tasks (continued)

3. Stress: find the 250 consecutive days with the largest one-day historical ES 2.5%, plot the rolling ES with the window shaded, and report the stressed ES at 1 day and, by $\sqrt{10}$, at 10 days.
 4. Capital: capital per 1 000 000 invested = $1\,000\,000 \times 1.5 \times \text{ES}_{10}^{\text{stress}} (\%) / 100$, for a 10-day and a 20-day liquidity horizon (a further factor $\sqrt{2}$); amounts in RON for BVB and in USD for SPY, ratios comparable.
 5. Uncertainty: 95% moving-block bootstrap intervals (20-day blocks, 2 000 resamples) for the one-day ES and for the stressed ES with the window held fixed.
- ▣ **Report:** a table of all numbers with intervals and a one-page argument: which portfolio needs more capital, and how do thin trading and currency change the answer?
- ▣ A simplified illustration inspired by FRTB [Basel Committee \(2019\)](#), not the full capital calculation of [Basel Committee, MAR33](#)


C3 [Open]: Deep Quantile Regression for VaR — Task (1/2) — Minimum Version

- **Question:** does a small neural network beat linear quantile regression for the S&P 500 VaR 1%, as in the lecture case study [Chronopoulos, Raftapostolos & Kapetanios \(2024\)](#)?
- **Data**
 - ▶ GSPC.INDX, 1990-01-03 – 31 August 2020 (the paper starts in 1985); unscaled daily log returns r_t ; tuning: first 5 726 days, evaluation: last 2 000 (from 2012-09-19)
 - ▶ One predictor set: GARCH(1,1) σ_t (estimated on the tuning sample only); $\alpha = 1\%$
- **Tasks**
 1. Build training pairs (x_t, r_{t+10}) whose targets all lie in the tuning sample, as in the chart.
 2. Fit a linear quantile regression and one network (one hidden ReLU layer of 5 units, ridge penalty 0.01, Adam, early stopping) by the pinball loss; scale inputs within each training fold only.
 3. Compare mean pinball losses on the evaluation sample with a Diebold–Mariano test (HAC, 10 lags) [Diebold & Mariano \(1995\)](#), over 5 random seeds.
- **Report:** a table of mean pinball losses, DM statistics and exceedance counts, per seed

C3 [Open]: Deep Quantile Regression for VaR — Data



Alignment of one training pair



— Tuning sample (5,726 days)

— CV test day +/- 5 excluded days (20 random)

— Evaluation sample (last 2,000 days)

- Choose the target explicitly: the daily return r_{t+10} , not the cumulative 10-day return

C3 [Open]: Deep Quantile Regression for VaR — Task (2/2) — Extensions

- ▣ **Replication extension:** the four predictor sets and the cross-validation grid of the paper (Appendix of the lecture); relative RMSFE only to reproduce Table 6, since RMSFE is minimised by the conditional mean, not the quantile
- ▣ **Rolling extension:** re-estimate once a year on a rolling window and forecast September 2020 – 2026 for the S&P 500 and the BET
- ▣ **Joint extension:** two outputs (v_t, e_t) trained on the FZ0 loss [Patton, Ziegel & Chen \(2019\)](#), with $e_t = -\exp(a_t)$ and $v_t = e_t + \exp(b_t)$ so that $e_t < 0$ and $e_t \leq v_t$
- ▣ **Interpretation** (one page): can a gain be attributed to non-linearity when architecture, penalty and tuning change together?

Project Choice and Common Rubric

Choose one

- C1: ES capital for a portfolio of your choice (at least three assets not covered in class), with the four-step structure of C1
- C3: the minimum version, then one extension
- Deliver a Quantlet folder (code, `Metainfo.txt`, charts) and three slides: numbers, interpretation, limitations

Rubric

- Reproducibility: data loaded by code, fixed seeds, every number from a script (30%)
- Correct conventions: sign, level, horizon, return type (20%)
- Uncertainty: tail counts and intervals for every tail estimate (25%)
- Interpretation and limits, including logged AI use and the AI errors actually found (25%)

References (1/3)

- ▣ Acerbi, C., Tasche, D. (2002). [On the coherence of expected shortfall](#). *Journal of Banking & Finance*, 26(7), 1487–1503.
- ▣ Artzner, P., Delbaen, F., Eber, J.-M., Heath, D. (1999). [Coherent measures of risk](#). *Mathematical Finance*, 9(3), 203–228.
- ▣ Barone-Adesi, G., Giannopoulos, K., Vosper, L. (1999). [VaR without correlations for portfolios of derivative securities](#). *Journal of Futures Markets*, 19(5), 583–602.
- ▣ Basel Committee on Banking Supervision (2019). [Minimum capital requirements for market risk](#). Bank for International Settlements.
- ▣ Chen, S.X. (2008). [Nonparametric estimation of expected shortfall](#). *Journal of Financial Econometrics*, 6(1), 87–107.
- ▣ Christoffersen, P., Gonçalves, S. (2005). [Estimation risk in financial risk management](#). *Journal of Risk*, 7(3), 1–28.
- ▣ Chronopoulos, I., Raftapostolos, A., Kapetanios, G. (2024). [Forecasting Value-at-Risk using deep neural network quantile regression](#). *Journal of Financial Econometrics*, 22(3), 636–669.
- ▣ Cornish, E.A., Fisher, R.A. (1938). [Moments and cumulants in the specification of distributions](#). *Revue de l'Institut International de Statistique*, 5(4), 307–320.
- ▣ Diebold, F.X., Mariano, R.S. (1995). [Comparing predictive accuracy](#). *Journal of Business & Economic Statistics*, 13(3), 253–263.

References (2/3)

- ▣ Efron, B. (1979). [Bootstrap methods: another look at the jackknife](#). *Annals of Statistics*, 7(1), 1–26.
- ▣ Engle, R.F., Manganelli, S. (2004). [CAViaR: conditional autoregressive value at risk by regression quantiles](#). *Journal of Business & Economic Statistics*, 22(4), 367–381.
- ▣ Ferro, C.A.T., Segers, J. (2003). [Inference for clusters of extreme values](#). *Journal of the Royal Statistical Society: Series B*, 65(2), 545–556.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.
- ▣ Harvey, D., Leybourne, S., Newbold, P. (1997). [Testing the equality of prediction mean squared errors](#). *International Journal of Forecasting*, 13(2), 281–291.
- ▣ Künsch, H.R. (1989). [The jackknife and the bootstrap for general stationary observations](#). *Annals of Statistics*, 17(3), 1217–1241.
- ▣ McNeil, A.J., Frey, R. (2000). [Estimation of tail-related risk measures for heteroscedastic financial time series: an extreme value approach](#). *Journal of Empirical Finance*, 7(3–4), 271–300.
- ▣ McNeil, A.J., Frey, R., Embrechts, P. (2015). [Quantitative Risk Management: Concepts, Techniques and Tools](#), revised ed. Princeton University Press.
- ▣ Patton, A.J., Ziegel, J.F., Chen, R. (2019). [Dynamic semiparametric models for expected shortfall \(and Value-at-Risk\)](#). *Journal of Econometrics*, 211(2), 388–413.

References (3/3)

- ▣ Pele, D.T., Mazurencu-Marinescu-Pele, M. (2026). Finite-sample precision limits for expected shortfall forecast comparisons. *Mathematics*, 14(13), 2316.
- ▣ Rockafellar, R.T., Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2(3), 21–41.
- ▣ Tasche, D. (2008). Capital allocation to business units and sub-portfolios: the Euler principle. arXiv:0708.2542.