



Modelling Financial Markets

Seminar 9: Realised Volatility — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ **AI:** Artificial Intelligence
- ▣ **AR:** AutoRegressive (model); in event studies: Abnormal Return
- ▣ **BH:** Benjamini–Hochberg (procedure)
- ▣ **BV:** Bipower Variation
- ▣ **BVB:** Bursa de Valori București — *the Bucharest Stock Exchange*
- ▣ **CI:** Confidence Interval
- ▣ **DM:** Diebold–Mariano (test)
- ▣ **EWMA:** Exponentially Weighted Moving Average
- ▣ **FDR:** False Discovery Rate
- ▣ **FRED:** Federal Reserve Economic Data
- ▣ **GARCH:** Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ **HAC:** Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ **HAR:** Heterogeneous AutoRegressive (model)
- ▣ **HAR-CJ:** HAR model with Continuous and Jump components
- ▣ **HARQ:** HAR model with Realised Quarticity
- ▣ **IQ:** Integrated Quarticity

Acronyms used in this material (2/2)

- ▣ **JB**: Jarque–Bera (test)
- ▣ **MCS**: Model Confidence Set
- ▣ **ML**: Machine Learning
- ▣ **MSE**: Mean Squared Error
- ▣ **MZ**: Mincer–Zarnowitz (regression)
- ▣ **NW**: Newey–West (standard errors)
- ▣ **OLS**: Ordinary Least Squares
- ▣ **QLIKE**: Quasi-LIKElihood (loss function)
- ▣ **QQ**: Quantile–Quantile (plot)
- ▣ **RF**: Random Forest
- ▣ **RQ**: Realised Quarticity
- ▣ **RV**: Realised Variance
- ▣ **SE**: Standard Error
- ▣ **SHAR**: Semivariance HAR model
- ▣ **TQ**: Tripower Quarticity
- ▣ **TSRV**: Two-Scale Realised Variance
- ▣ **UTC**: Coordinated Universal Time (Temps Universel Coordonné)
- ▣ **VIX**: CBOE Volatility Index

Today's Questions and Route

- ▣ Questions of the seminar
 - ▶ **how precisely does one day of intraday prices measure volatility?**
 - ▶ **does the measurement improve a forecast?**
- ▣ Route
 - ▶ Primer, what you need today from [Chapter 9](#) (the lecture comes after this seminar): RV, noise, jumps, HAR, forecast evaluation
 - ▶ Part I, measurement and diagnosis: A1 precision of RV, A3 noise and sampling, B3 sampling grids on SPY, B1 jump days in SPY, A2 one jump in six returns
 - ▶ Part II, models and forecasts: A5 HAR as a restricted AR(22), B5 log-HAR with robust inference, A6 losses, the forecast-evaluation workflow, A8 accuracy versus calibration, C2 an AI answer
- ▣ What you produce: the A answers, one figure and one table per B exercise, the C2 correction table, one interpretation sentence per exercise
- ▣ Notebook for today: [open the seminar notebook in Colab](#); each task names its notebook section; background: [Chapter 9](#) and [Franke, Härdle & Hafner \(2019\)](#)
- ▣ Seminar: Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Worked example
A1	how precise is one day's RV?	Solved	–
A3	how often should a noisy price be sampled?	Solved	–
B3	does SPY show a noise bias at 5 minutes?	Solved	A3
B1	which SPY days contain a jump?	Solved	–
A2	can six returns reveal a jump?	Proposed	B1, A1
A5	what does HAR impose on an AR(22)?	Solved	–
B5	how persistent is SPY's log variance?	Solved	A5
A6	which loss ranks variance forecasts correctly?	Solved	–
A8	is GARCH less accurate, biased, or both?	Proposed	A5, A6
C2	what is wrong in an AI answer?	Proposed	A3, B5
Extension packs	A2, A3, A5, A6, B1, B3, B5 extended; A4, A7, B2, B4, B6–B9, C1, C3 (optional, after the seminar)	Solved and Proposed	named on each task

- ☐ **[Solved]:** full solution in the slides and in the notebook, a model to copy; **[Proposed]:** you solve it, the solution is discussed in class
- ☐ Part A: derivations on paper; Part B: computation and inference on SPY and Bitcoin, each ending with an interpretation question; Part C: open questions and AI critique

Data Used

Series	Content	Sample
SPY, 5-minute bars	prices of the regular session, UTC time stamps, each bar labelled by its start	2020-10-12 – 2026-09-18, 1 489 days
Bitcoin (USD), 5-minute bars	24-hour UTC days	659 of 747 days kept
SPY and Bitcoin, daily prices	adjusted closes (SPY) and closes (Bitcoin), for GARCH	since 2000 (SPY), 2014 (Bitcoin)
SPY daily variances (derived)	intraday RV, overnight return and total variance V_t ($\%^2$, $\%$)	one row per SPY day
One-day forecasts (derived)	log-HAR, HAR, GARCH(1,1)- t , EWMA, yesterday's RV (column RW) and the target V_t	from 3 Jan 2022 (SPY), 1 Mar 2025 (Bitcoin)


- SPY: UTC converted to New York time; bars 09:30–15:55; missing bars inside a day carry the last price; 78 returns a day, 43 on the 12 half-days
- Bitcoin: days with fewer than 95% of the 288 bars are dropped; short gaps carry the last price
- Units: returns $r = 100 \Delta \ln P$ in $\%$, variances in $\%^2$; annualise with 252 (SPY) or 365 (Bitcoin) days; all data end on 18 September 2026

Setup: Run and Verify

```
# notebook section "Data" defines
# spy_intraday, btc_intraday,
# intraday_returns, spy_overnight
P = spy_intraday()           # day x 79 prices
R = intraday_returns(P)     # day x 78 returns, %
print(P.shape)
print(R.notna().sum(axis=1).value_counts())
Pb = btc_intraday()         # day x 289 prices, UTC
print(Pb.shape, Pb.index[0], Pb.index[-1])
```

- Expected output
 - ▶ (1489, 79): SPY days, from 2020-10-12 to 2026-09-18
 - ▶ 1 477 days with 78 returns, 12 with 43
 - ▶ (659, 289): Bitcoin, 2024-09-01 to 2026-09-17
- Column 0 of P is the opening price of the 09:30 bar; columns 1–78 are the closes of the bars
- If your counts differ, stop and check the time zone, the session filter and the cut-off date

From Bars to Returns: One SPY Day



- 7 April 2025: opening price $P_0 = 489.19$, then bar closes $P_1 = 487.60$, $P_2 = 485.88$, $\dots$, $P_{78} = 504.54$: 79 prices give 78 returns; differencing the 78 closes alone would give 77
- $r_1 = 100 \ln(P_1/P_0) = -0.33\%$, $r_2 = -0.35\%$; $RV = \sum_i r_i^2 = 45.1\%^2$; the overnight return -3.27% is outside the bars, total variance $V = RV + r_{on}^2 = 55.8\%^2$
- Overnight return: the adjusted close-to-close return minus the open-to-close return, so dividends fall in the overnight part

Notation Used in Every Exercise

- ▣ Day t has M intraday returns $r_{t,i}$ (%); $RV_t = \sum_i r_{t,i}^2$; realised quarticity $RQ_t = \frac{M}{3} \sum_i r_{t,i}^4$
- ▣ Bipower variation [Barndorff-Nielsen & Shephard \(2004\)](#) $BV_t = \frac{\pi}{2} \frac{M}{M-1} \sum_{i \geq 2} |r_{t,i}| |r_{t,i-1}|$; tripower quarticity $TQ_t = M \mu_{4/3}^{-3} \frac{M}{M-2} \sum_{i \geq 3} |r_{t,i}|^{4/3} |r_{t,i-1}|^{4/3} |r_{t,i-2}|^{4/3}$, a jump-robust estimate of IQ_t , $\mu_{4/3} = \mathbb{E}|Z|^{4/3}$, $Z \sim N(0, 1)$
- ▣ IV_t : integrated variance, IQ_t : integrated quarticity; ω^2 : variance of the microstructure noise
- ▣ Total daily variance $V_t = RV_t + r_{on,t}^2$ for SPY; $V_t = RV_t$ for Bitcoin (24 hours)
- ▣ Forecast F_t of V_t , made at the close of day $t-1$; $QLIKE(V, F) = V/F - \ln(V/F) - 1$; $MSE = (V - F)^2$, on variance

Realised Variance: Volatility Measured after the Fact

- ▣ From inferring to **measuring** daily variance Andersen et al. (2001)
 - ▶ GARCH (Chapter 5) infers today's variance from daily returns through a model
 - ▶ with intraday prices we add up squared 5-minute returns
- ▣ Integrated variance: the variance the day actually delivered
 - ▶ $IV_t = \int_{t-1}^t \sigma_s^2 ds$, σ_s : the spot volatility inside the day
- ▣ Realised variance: the sum of squared intraday returns
 - ▶ $RV_t = \sum_{i=1}^M r_{t,i}^2 \rightarrow IV_t$ as $M \rightarrow \infty$ (continuous price, no noise)
 - ▶ $r_{t,i}$: the i -th intraday return of day t ; SPY: $M = 78$ five-minute returns a session
- ▣ Measurement error: $RV_t - IV_t \approx N(0, 2IQ_t/M)$
 - ▶ $IQ_t = \int \sigma_s^4 ds$: integrated quarticity, estimated by RQ_t (notation slide)
 - ▶ constant volatility: relative standard error $\sqrt{2/M} = \sqrt{2/78} = 0.16$, so one day's RV is a noisy measurement
- ▣ Tiny example: returns 0.10, -0.20, 0.10, 0.20 (%)
 - ▶ $RV = 0.01 + 0.04 + 0.01 + 0.04 = 0.10\%$
 - ▶ daily volatility $\sqrt{0.10} = 0.32\%$, annualised $\sqrt{252 \times 0.10} = 5.0\%$
- ▣ Depth: Chapter 9, after the seminar

Two Things That Distort RV: Noise and Jumps

- **Microstructure noise:** observed log price = efficient price + error
 - ▶ $p_i^{\text{obs}} = p_i^* + \varepsilon_i$, $\text{Var}(\varepsilon_i) = \omega^2$ (bid-ask bounce, price ticks)
 - ▶ each return carries two noise terms: $\mathbb{E}[\text{RV}^{(n)}] = \text{IV} + 2n\omega^2$ with n returns
 - ▶ the bias grows with n ; the variance of the sampling error falls like $1/n$
 - ▶ hence an optimal sampling frequency [Bandi & Russell \(2008\)](#); in practice 5 minutes is hard to beat [Liu, Patton & Sheppard \(2015\)](#)
 - ▶ signature plot: mean RV against the sampling interval, flat when the noise bias is negligible
- **Jumps:** RV also collects the squared jumps, $\text{RV}_t \rightarrow \text{IV}_t + \sum J^2$
 - ▶ bipower variation [Barndorff-Nielsen & Shephard \(2004\)](#): $\text{BV}_t = \frac{\pi}{2} \frac{M}{M-1} \sum_{i \geq 2} |r_{t,i}| |r_{t,i-1}| \rightarrow \text{IV}_t$
 - ▶ a jump enters BV only multiplied by a small neighbouring return
 - ▶ $\pi/2 = 1/(\mathbb{E}|Z|)^2$, $Z \sim N(0, 1)$: rescales products of absolute Normal returns
- Tiny example, $M = 4$, returns 0.1, 0.1, 0.1, 1.0 (%)
 - ▶ $\text{RV} = 1.03$; $\text{BV} = \frac{\pi}{2} \cdot \frac{4}{3} (0.01 + 0.01 + 0.10) = 0.25$; jump part $\text{RV} - \text{BV} = 0.78$
- Depth: [Chapter 9](#), after the seminar

Testing for a Jump: the Ratio Statistic

- H_0 : no jump on day t
 - ▶ the relative jump measure $(RV_t - BV_t)/RV_t$ is near 0 under H_0 and large on a jump day
- Ratio statistic [Huang & Tauchen \(2005\)](#)
 - ▶ $z_t = \frac{(RV_t - BV_t)/RV_t}{\sqrt{\theta M^{-1} \max(1, TQ_t/BV_t^2)}}, \theta = \pi^2/4 + \pi - 5 = 0.609$
 - ▶ TQ_t : tripower quarticity (notation slide), estimates IQ_t without being inflated by a jump
 - ▶ TQ_t/BV_t^2 : how uneven volatility was within the day, floored at 1 (its constant-volatility value)
- Under H_0 , $z_t \approx N(0, 1)$; one-sided test
 - ▶ reject when $z_t > 1.645, 2.326$ or 3.09 at 5%, 1% or 0.1%
 - ▶ $p_t = 1 - \Phi(z_t)$, Φ : the standard Normal distribution function
- Many days, many tests (multiple testing, [Chapter 3](#))
 - ▶ with N days tested, about αN rejections appear by chance alone: compare the count with αN
- Depth: [Chapter 9](#), after the seminar

HAR: a Cascade of Three Horizons

- ▣ Stylised fact ([Chapter 1](#)): volatility has long memory
 - ▶ the autocorrelation of $|r_t|$ decays slowly; a GARCH(1,1) forecast reverts at an exponential rate ([Chapter 5](#))
- ▣ HAR [Corsi \(2009\)](#): three averages of past RV
 - ▶ $RV_{t+1} = \beta_0 + \beta_d RV_t + \beta_w RV_t^{(5)} + \beta_m RV_t^{(22)} + \varepsilon_{t+1}$
 - ▶ $RV_t^{(5)}, RV_t^{(22)}$: means over the last 5 and 22 days; β_0 : constant; ε_{t+1} : forecast error
 - ▶ $\beta_d, \beta_w, \beta_m$: weights on yesterday, the last week and the last month
- ▣ Intuition: traders with daily, weekly and monthly horizons
 - ▶ three slopes reproduce a slowly decaying memory
 - ▶ log-HAR uses $y_t = \ln V_t$ in place of RV_t : positive forecasts, residuals closer to the Normal distribution
- ▣ Estimated by OLS with Newey–West (HAC) standard errors ([Chapter 3](#))
 - ▶ the errors are autocorrelated and heteroskedastic
- ▣ Tiny example: $\hat{\beta} = (0.1, 0.4, 0.3, 0.2)$, $RV_t = 2$, $RV_t^{(5)} = 1.5$, $RV_t^{(22)} = 1$
 - ▶ forecast $0.1 + 0.8 + 0.45 + 0.2 = 1.55\%^2$; persistence $\beta_d + \beta_w + \beta_m = 0.9$
- ▣ Depth: [Chapter 9](#), after the seminar

Judging Variance Forecasts: the Tools of Chapter 5

- The true variance is never observed: forecasts are scored against a proxy
 - ▶ here $V_t = \text{RV}_t + r_{\text{on},t}^2$ for SPY ($r_{\text{on},t}$: overnight return), noisy but conditionally unbiased
- Robust losses [Patton \(2011\)](#): rank forecasts as the true variance would
 - ▶ variance MSE $(V - F)^2$ and QLIKE(V, F) = $V/F - \ln(V/F) - 1$; F : the forecast
 - ▶ QLIKE penalises under-prediction more
- Accuracy: Diebold–Mariano [Diebold & Mariano \(1995\)](#)
 - ▶ $t = \bar{d}/\text{se}_{\text{HAC}}(\bar{d})$ on $d_t = L(V_t, F_t^A) - L(V_t, F_t^B)$ for models A and B
 - ▶ L : QLIKE or MSE; $\bar{d}$: mean of d_t ; $t < -1.96$: A significantly better
- Calibration: Mincer–Zarnowitz regression $V_t = a + bF_t + u_t$
 - ▶ a : intercept, b : slope, u_t : error; unbiased forecast: $a = 0$ and $b = 1$, tested jointly (Wald)
 - ▶ $b < 1$: forecasts too extreme
- Out of sample
 - ▶ at each origin $t - 1$ the model is re-estimated on data up to $t - 1$ only, then F_t is compared with V_t
- Depth: [Chapter 9](#), after the seminar


A1 [Solved]: Precision of Realised Variance — Task

- ▣ **Question:** how precisely does one day of 5-minute returns measure integrated variance?
- ▣ **Inputs:** constant volatility, M intervals of length $\Delta = 1/M$, $r_i = \sigma\sqrt{\Delta} Z_i$ with Z_i i.i.d. $N(0, 1)$, $IV = \sigma^2$; use $\mathbb{E}Z^2 = 1$, $\text{Var}(Z^2) = 2$, $\mathbb{E}Z^4 = 3$; χ_M^2 : chi-square with M degrees of freedom (a sum of M squared i.i.d. $N(0, 1)$); no data file
- ▣ **Tasks**
 1. Show that $RV = \sigma^2 \chi_M^2 / M$ and derive $\mathbb{E}[RV]$ and $\text{Var}(RV)$.
 2. Show that $RQ = \frac{M}{3} \sum_i r_i^4$ is unbiased for $IV = \sigma^4$ and write the estimated standard error of RV .
 3. Use the delta method for $\ln RV$ and write the 95% interval for IV built on the log scale.
 4. Compute the interval factors $e^{\pm 1.96\sqrt{2/M}}$ for $M = 78$ and $M = 6$, and the relative width for $M = 78$.
- ▣ **Report:** the formulas, the two pairs of factors and one sentence on how precise one day's RV is
- ▣ **Notebook:** section A1

A1 [Solved]: Solution (1/2) — Derivation

- ▣ $RV = \sum_i \sigma^2 \Delta Z_i^2 = \sigma^2 \frac{1}{M} \sum_i Z_i^2 = \sigma^2 \chi_M^2 / M$, so $\mathbb{E}[RV] = \sigma^2$ (unbiased) and $\text{Var}(RV) = \sigma^4 \cdot 2M/M^2 = 2\sigma^4/M$
- ▣ $\mathbb{E}[r_i^4] = \sigma^4 \Delta^2 \mathbb{E}Z^4 = 3\sigma^4/M^2$, so $\mathbb{E}[RQ] = \frac{M}{3} \cdot M \cdot 3\sigma^4/M^2 = \sigma^4$; hence $\widehat{\text{Var}}(RV) = 2RQ/M$ and $\text{se}(RV) = \sqrt{\frac{2}{3} \sum_i r_i^4}$
- ▣ Delta method with $g(x) = \ln x$, $g'(x) = 1/x$: $\text{se}(\ln RV) \approx \text{se}(RV)/RV = \sqrt{2RQ/M}/RV$
- ▣ 95% interval: $RV \exp(\pm 1.96 \text{se}(\ln RV))$, always positive and asymmetric in levels; with constant volatility $RQ/RV^2 \approx 1$, so the factors are $e^{\pm 1.96\sqrt{2/M}}$
- ▣ For volatility $\sqrt{RV}$, $g'(x) = 1/(2\sqrt{x})$ halves the relative standard error: the interval for volatility is narrower than the one for variance

A1 [Solved]: Solution (2/2) — Interval Factors and Coverage



- $M = 78$: $\sqrt{2/78} = 0.160$, factors [0.73, 1.37]: IV lies between 0.73 RV and 1.37 RV, a total relative width of 0.64
- $M = 6$ (30-minute returns): factors [0.32, 3.10]: almost uninformative
- Simulated coverage of the nominal 95% interval with RQ (200 000 days): 94.2% for $M = 78$, 86.7% for $M = 6$: the 95% is an asymptotic statement
- With stochastic volatility the same studentised limit holds asymptotically, for a continuous price and negligible noise, because the convergence is stable (Chapter 9)
- **Interpretation:** one day of 5-minute data pins IV down to roughly $-27\%/+37\%$


A3 [Solved]: Noise Bias and the Optimal Frequency — Task

- ▣ **Question:** how often should we sample a price that contains noise?
- ▣ **Inputs:** $p_i^{\text{obs}} = p_i^* + \varepsilon_i$, ε_i i.i.d. $(0, \omega^2)$ independent of p^* ; zero drift, constant volatility; $IV = 1\%^2$, $IQ = 1$, $\omega = 0.01\%$, so $\omega^2 = 0.0001\%^2$; session of 23 400 seconds
- ▣ **Tasks**
 1. Write the observed return, expand $\mathbb{E}[r_i^2]$ and derive $\mathbb{E}[RV^{(n)}] = IV + 2n\omega^2$ for n returns.
 2. Derive $\text{cov}(r_i, r_{i-1}) = -\omega^2$.
 3. From $\text{MSE}(n) \approx 2IQ/n + 4n^2\omega^4$ [Bandi & Russell \(2008\)](#), derive the MSE-optimal n^* and its sampling interval.
 4. Fill a table for 5 minutes, 1 minute, 1 second and n^* : n , bias $2n\omega^2$, $\mathbb{E}[RV]$, MSE.
- ▣ **Report:** the derivations, the table and which interval you would use in practice, with the reason
- ▣ **Notebook:** section A3

A3 [Solved]: Solution (1/2) — Derivation

- $r_i = p_i^{\text{obs}} - p_{i-1}^{\text{obs}} = r_i^* + \varepsilon_i - \varepsilon_{i-1}$
- $\mathbb{E}[r_i^2] = \mathbb{E}[r_i^{*2}] + \mathbb{E}[\varepsilon_i^2] + \mathbb{E}[\varepsilon_{i-1}^2] = \mathbb{E}[r_i^{*2}] + 2\omega^2$ (cross terms have mean 0 by independence); summing over n returns: $\mathbb{E}[\text{RV}^{(n)}] = \text{IV} + 2n\omega^2$
- Adjacent returns share only ε_{i-1} , with opposite signs: $\text{cov}(r_i, r_{i-1}) = \text{cov}(-\varepsilon_{i-1}, \varepsilon_{i-1}) = -\omega^2$
- $\frac{d}{dn} \text{MSE} = -2\text{IQ}/n^2 + 8n\omega^4 = 0$ gives $n^3 = \text{IQ}/(4\omega^4)$, $n^* = (\text{IQ}/(4\omega^4))^{1/3}$; the second derivative is positive: a minimum
- Numbers: $n^* = (1/(4 \times 10^{-8}))^{1/3} = 292$, one return every $23\,400/292 = 80$ seconds

A3 [Solved]: Solution (2/2) — Table and Chart



Interval	n	$2n\omega^2$	$E[RV]$	MSE
5 minutes	78	0.0156	1.0156	0.0259
1 minute	390	0.0780	1.0780	0.0112
1 second	23 400	4.6800	5.6800	21.9025
n^* , every 80 s	292	0.0585	1.0585	0.0103

- $MSE = 2/n + 4n^2 \times 10^{-8}$ in %⁴; at one second RV is almost six times IV
- In this model 1 minute is close to the optimum; 5 minutes has 2.5 times the minimal MSE
- **Interpretation:** the usual 5-minute choice rests on arguments outside this model
 - ▶ ω is uncertain; real noise is not i.i.d.
 - ▶ 5-minute RV is hard to beat empirically [Liu, Patton & Sheppard \(2015\)](#)

B3 [Solved]: Which Sampling Frequency? — Task

▣ **Worked example to follow:** A3 [Solved]

▣ **Questions**

- ▶ does SPY show a measurable noise bias at 5 minutes?
- ▶ which grid gives the most useful daily measure?

▣ **Data:** the SPY price matrix P of the setup (1 489 days); grid of $k \times 5$ minutes, $k = 1, 3, 6, 13$ (5, 15, 30, 65 minutes): prices at columns $o, o + k, o + 2k, \dots$ plus the open and the close, for offsets $o = 0, \dots, k - 1$ (helper `sparse_points`)

▣ **Tasks**

1. Compute daily RV on one grid ($o = 0$) and the average over the k shifted grids, and report $\sqrt{252 \overline{RV}}$ for each interval.
2. For each interval compute the ratio of mean one-grid RV to mean 5-minute RV and a 95% moving-block bootstrap interval [Künsch \(1989\)](#): resample 20-day blocks of the paired daily series, 2 000 times, and recompute the ratio of means.
3. Compute the correlation of today's log RV (one grid and averaged grids) with tomorrow's 5-minute log RV.

▣ **Report:** one table (volatility, ratio [CI], correlation per interval) and one chart

▣ **Interpretation**

- ▶ is there a measurable noise bias at 5 minutes?
- ▶ why a block bootstrap and not an i.i.d. one?
- ▶ which grid is most useful for forecasting?

▣ **Notebook:** section B3


B3 [Solved]: Solution (1/2) — Procedure and Table

- Procedure: for each day take the grid prices, compute $RV^{(o)} = \sum (100 \Delta \ln P)^2$ on each offset o , average over o ; for the bootstrap, draw block starts at random, glue 20-day blocks of the paired series $(RV_t^{(k)}, RV_t^{(s)})$ to the original length, take $\sum RV^{(k)} / \sum RV^{(s)}$, and read the 2.5% and 97.5% percentiles

Interval	Vol. (one grid)	Vol. (avg. grids)	Ratio [95% CI]	Corr. next day	Corr. (avg.)
5 min	12.93	12.93	1.000 [1.000, 1.000]	0.73	0.73
15 min	12.88	12.79	0.992 [0.954, 1.035]	0.71	0.72
30 min	12.75	12.77	0.973 [0.933, 1.015]	0.70	0.72
65 min	12.76	12.88	0.974 [0.913, 1.037]	0.64	0.70

- Block bootstrap: daily RVs are strongly autocorrelated; i.i.d. resampling would destroy this dependence and make the interval too narrow

B3 [Solved]: Solution (2/2) — Chart and Answer



- All intervals contain 1: the coarse grids detect no difference from 5-minute RV; the signature is flat
- The predictive correlation falls from 0.73 (5 minutes) to 0.64 (65 minutes, one grid): sparse RV is noisier; averaging the shifted grids recovers part of it (0.70)
- **Interpretation:** 5-minute RV is the most useful measure here; in the kernel comparison (extension pack) the realised kernel averages 0.952 of the 5-minute RV (95% CI [0.910, 0.987]): a small but significant difference

B1 [Solved]: Jump Days in SPY — Task

□ Questions

- ▶ on which SPY days does the ratio test reject?
- ▶ how many rejections would chance alone produce?

- **Data:** the SPY return matrix R (1 489 days; $M_t = 78$, or 43 on half-days); ratio test [Huang & Tauchen \(2005\)](#), formulas on the notation slide; helper `jump_test`

□ Tasks

1. For every day compute RV , BV , TQ , z_t and the one-sided $p_t = 1 - \Phi(z_t)$; work 27 August 2021 step by step.
2. Count the days that reject at 5%, 1% and 0.1% and compare each count with αN .
3. On the 0.1% days find the interval with the largest $|r_{t,i}|$ and plot the histogram of its end time.
4. Compute the jump share $\sum_t J_t / \sum_t RV_t$, with $J_t = RV_t - BV_t$ on 0.1% days and 0 otherwise.

- **Report:** the worked day, a table of counts against αN , the histogram and the jump share

□ Interpretation

- ▶ are there more rejections than the level alone would produce?
- ▶ what could explain them?

- Notebook: section B1

B1 [Solved]: Solution (1/3) — One Day Step by Step


- ▣ 27 August 2021, $M = 78$: $RV = 0.1354$, $BV = 0.0687$, $TQ = 0.00539$ ($\%^2$, $\%^2$, $\%^4$)
- ▣ $(RV - BV)/RV = 0.493$; $TQ/BV^2 = 1.143 > 1$, so the floor does not bind
- ▣ $z = 0.493/\sqrt{0.609 \times 1.143/78} = 5.22$, $p = 9.0 \times 10^{-8}$: reject at 0.1%
- ▣ The largest return, +0.24% in the interval ending 10:05, is 44% of RV: 10:00–10:05 is when the Federal Reserve chair's Jackson Hole speech was released
- ▣ The same four steps, vectorised over days, are `jump_test(R)`

B1 [Solved]: Solution (2/3) — Counts and Jump Share

Level	Rejections	αN
5%	235	74.5
1%	110	14.9
0.1%	35	1.5

- z_t has mean 0.50 and standard deviation 1.14 instead of 0 and 1: far more rejections than the level implies
- Jump share of total RV with 0.1% days: 2.8%
- Explanations: real jumps, and the intraday U-shape, which biases BV down (extension pack: FDR control and the periodicity correction)

B1 [Solved]: Solution (3/3) — When Do the Largest Returns Occur?



- Time of the largest $|r|$ on flagged days: it peaks after the open; attributing a flag to news needs separate evidence

A2 [Proposed]: Six Returns and One Jump — Task


- ▣ **Worked example to follow:** B1 [Solved] (one day step by step) and A1
- ▣ **Question:** can six 5-minute returns reveal a jump?
- ▣ **Inputs:** the returns in the table (%); the last one is a news jump;
 $TQ = 0.0097$ (supplied); $\theta = \pi^2/4 + \pi - 5 = 0.609$; critical value at 0.1%: 3.09
- ▣ **Tasks**
 1. Complete the table and compute RV and $BV = \frac{\pi}{2} \frac{M}{M-1} \sum_{i \geq 2} |r_i r_{i-1}|$ with $M = 6$.
 2. Compute the jump part $RV - BV$ and its share of RV.
 3. Compute TQ/BV^2 , apply the floor $\max(1, TQ/BV^2)$ and compute
$$z = \frac{(RV - BV)/RV}{\sqrt{\theta M - 1} \max(1, TQ/BV^2)}.$$
 4. Decide at 0.1% and compare z with its largest possible value $\sqrt{M/\theta}$.
- ▣ **Report:** the completed table, RV, BV, z , the decision and one sentence of interpretation

i	r_i	r_i^2	$ r_i r_{i-1} $
1	0.08	0.0064	–
2	–0.15		
3	0.03		
4	0.10		
5	–0.04		
6	–0.90		

A5 [Solved]: HAR as a Restricted AR(22) — Task

- **Question:** which restrictions does HAR impose on an AR(22) for log variance?
- **Inputs:** log-HAR [Corsi \(2009\)](#) $y_{t+1} = \beta_0 + \beta_d y_t + \beta_w \bar{y}_t^{(5)} + \beta_m \bar{y}_t^{(22)} + \varepsilon_{t+1}$, $y_t = \ln V_t$, $\bar{y}_t^{(5)}$ the mean of $y_t, \dots, y_{t-4}$, $\bar{y}_t^{(22)}$ the mean of $y_t, \dots, y_{t-21}$; rounded SPY estimates $\hat{\beta} = (-0.08, 0.36, 0.34, 0.17)$
- **Tasks**
 1. Expand the two means and write the implied AR coefficients $\phi_1, \phi_2 = \dots = \phi_5$ and $\phi_6 = \dots = \phi_{22}$.
 2. With $\theta = (\phi_0, \phi_1, \dots, \phi_{22})$, write two rows of the matrix R in $R\theta = 0$ and count the restrictions.
 3. Compute $\phi_1, \phi_{2..5}, \phi_{6..22}$ and the persistence $\sum_k \phi_k$ for $\hat{\beta}$, and draw the 22 weights.
- **Report:** the three weight formulas, two rows of R , the numbers and what the persistence implies for shocks
- Notebook: section A5

A5 [Solved]: Solution



- Lag 1 is in all three components: $\phi_1 = \beta_d + \beta_w/5 + \beta_m/22$; lags 2–5 in the weekly and monthly ones: $\phi_k = \beta_w/5 + \beta_m/22$; lags 6–22 only in the monthly one: $\phi_k = \beta_m/22$
- Rows of R (column 0: intercept): $(0, 0, 1, -1, 0, \dots, 0)$ for $\phi_2 = \phi_3$; $(0, \dots, 0, 1, -1, 0, \dots)$ at columns 6, 7 for $\phi_6 = \phi_7$
- Count: 3 equalities among $\phi_2, \dots, \phi_5$ plus 16 among $\phi_6, \dots, \phi_{22}$: 19 rows, R is 19×23 (22 lags minus 3 free slopes)
- $\phi_1 = 0.36 + 0.34/5 + 0.17/22 = 0.436$; $\phi_{2..5} = 0.076$; $\phi_{6..22} = 0.008$; $\sum_k \phi_k = \beta_d + \beta_w + \beta_m = 0.87$
- **Interpretation:** persistence $0.87 < 1$: shocks to log variance decay slowly but die out; the data test of the 19 restrictions is in the extension pack

B5 [Solved]: log-HAR for SPY with Robust Inference — Task

- ▣ **Worked example to follow:** A5 [Solved]
- ▣ **Questions**
 - ▶ how persistent is SPY's log total variance?
 - ▶ do robust standard errors change any conclusion?
- ▣ **Data:** SPY total daily variance from the setup, $V_t = RV_t + r_{\text{on},t}^2$, 1 489 days; regressors dated $t - 1$: $\ln V_{t-1}$ and the means of $\ln V$ over days $t - 5..t - 1$ and $t - 22..t - 1$ (mean of logs); the first 22 days are lost
- ▣ **Tasks**
 1. Build the design matrix and check your row for 8 April 2025 against the solution.
 2. Estimate log-HAR by OLS; report the usual and the Newey–West (HAC) standard errors with 7 lags, rule $\lfloor 4(T/100)^{2/9} \rfloor$.
 3. Test $\beta_w = \beta_m$ with a Wald test using the Newey–West covariance.
 4. Run Ljung–Box tests with 22 lags on $\ln V_t$ and on the residuals, plot both ACFs, and compare R^2 with an AR(1).
- ▣ **Report:** the coefficient table with both standard errors, the tests with p -values and the chart
- ▣ **Interpretation**
 - ▶ which component matters most?
 - ▶ do robust errors change a conclusion?
 - ▶ has HAR captured the persistence?
- ▣ **Notebook:** section B5


B5 [Solved]: Solution (1/2) — Design Row and Estimates

- Row for 8 April 2025: $y_t = \ln V_t = 2.976$; regressors $\ln V_{t-1} = 4.022$ (7 April), weekly mean 2.068, monthly mean 0.612
- Fitted value $\hat{y}_t = 2.177$: the model expects a partial return towards the monthly level
- Usable observations: $1\,489 - 22 = 1\,467$; $R^2 = 0.470$ against 0.413 for an AR(1) ($\hat{\phi} = 0.643$); persistence 0.866

	$\hat{\beta}$	SE OLS	SE NW	t_{NW}
β_o	-0.079	0.025	0.027	-3.0
β_d	0.362	0.031	0.037	9.8
β_w	0.337	0.049	0.048	7.0
β_m	0.167	0.046	0.044	3.8

- Newey–West errors are up to 20% larger than the OLS ones; all slopes stay significant

B5 [Solved]: Solution (2/2) — Tests and Chart




- Wald $\beta_w = \beta_m$: 4.17, $p = 0.041$: rejected at 5%, the weekly weight is larger
- Ljung–Box(22): $\ln V_t$ 4 641, residuals 17.5 ($p = 0.73$): three coefficients absorb the persistence; first-order residual autocorrelation -0.014
- Interpretation:** the weekly and daily components dominate; robust errors do not change a conclusion; HAR leaves no significant autocorrelation

A6 [Solved]: Which Loss Ranks Variance Forecasts? — Task


- **Question:** which loss ranks variance forecasts correctly when the target is a noisy proxy?
- **Inputs:** proxy V_t conditionally unbiased: $\mathbb{E}[V_t | \mathcal{F}_{t-1}] = \mathbb{E}[IV_t | \mathcal{F}_{t-1}]$ [Patton \(2011\)](#);
QLIKE(V, F) = $V/F - \ln(V/F) - 1$; variance MSE $(V - F)^2$
- **Tasks**
 1. Minimise $\mathbb{E}[\text{QLIKE}(V_t, F) | \mathcal{F}_{t-1}]$ over F and show that the minimiser is $\mathbb{E}[IV_t | \mathcal{F}_{t-1}]$.
 2. For two forecasts F_A, F_B show that $\Delta L = V(1/F_A - 1/F_B) + \ln(F_A/F_B)$ is linear in V , so its conditional mean is the same with V_t or with IV_t : the ranking is preserved.
 3. With true variance 1, compute QLIKE and variance MSE for $F = 0.5$ and $F = 2$ and compare the two rankings.
- **Report:** the first-order condition, the ranking argument, four loss values and which loss you would use
- **Notebook:** section A6

A6 [Solved]: Solution




- $\partial_F \mathbb{E}[V/F - \ln V + \ln F - 1] = -\mathbb{E}V/F^2 + 1/F = 0$ gives $F = \mathbb{E}[V_t | \mathcal{F}_{t-1}] = \mathbb{E}[IV_t | \mathcal{F}_{t-1}]$; second derivative $1/F^2 > 0$
- $\text{QLIKE}(V, F_A) - \text{QLIKE}(V, F_B) = V(1/F_A - 1/F_B) + \ln(F_A/F_B)$: linear in V , and F_A, F_B are known at $t - 1$, so replacing V_t by the conditionally unbiased IV_t leaves $\mathbb{E}[\Delta L | \mathcal{F}_{t-1}]$ unchanged
- QLIKE: $F = 0.5$ gives $2 - \ln 2 - 1 = 0.307$, $F = 2$ gives $0.5 + \ln 2 - 1 = 0.193$; variance MSE: 0.25 and 1, the opposite ranking
- **Interpretation:** both losses are robust, but QLIKE penalises under-prediction more, which matters for risk; MSE on volatility is not robust (extension pack)

Forecast Evaluation: Origin, Target, Window



- At each origin: re-estimate on all data up to $t - 1$, forecast F_t , then observe V_t ; no value from day t enters the forecast
- SPY: 1 180 origins from 3 January 2022 to 18 September 2026; losses compared on common dates only
- One-day forecasts of log-HAR, HAR, GARCH(1,1)- t , EWMA and yesterday's RV (column RW), and the target V_t

One-Day Forecasts and Cumulative Losses



- Mean QLIKE: log-HAR 0.310, GARCH(1,1)-t 0.343; a rising cumulative line means log-HAR has the lower loss
- $d_t = \text{QLIKE}(\text{log-HAR}) - \text{QLIKE}(\text{GARCH})$ is the input of A8: is the gap more than chance?

A8 [Proposed]: Accuracy versus Calibration — Task

- ▣ **Worked example to follow:** A5 (Wald statistic) and A6 [Solved]
- ▣ **Question:** is GARCH less accurate than log-HAR, biased, or both?
- ▣ **Inputs:** SPY one-day forecasts, January 2022 – September 2026; $d_t = \text{QLIKE}(\log\text{-HAR}) - \text{QLIKE}(\text{GARCH})$ has mean $\bar{d} = -0.0327$, Newey–West SE 0.0142; MZ regression for GARCH, $V_t = a + bF_t + u_t$: $\hat{a} = 0.2418$ (SE 0.2310), $\hat{b} = 0.7046$ (SE 0.2382), $\text{corr}(\hat{a}, \hat{b}) = -0.9750$ (keep four decimals)
- ▣ **Tasks**
 1. Accuracy, H_0 : equal expected loss. Compute the Diebold–Mariano t , its two-sided p -value, and decide at 5%.
 2. Calibration, H_0 : $a = 0$, $b = 1$. Compute the separate t -statistics for $a = 0$ and $b = 1$.
 3. Build the 2×2 covariance matrix, compute the Wald statistic $W = \delta^\top \hat{V}^{-1} \delta$ (the distance of the estimates from $(0, 1)$, weighted by their covariance $\hat{V}$) with $\delta = (\hat{a}, \hat{b} - 1)$ and its χ^2_2 p -value.
 4. Draw the 95% confidence ellipse, mark $(0, 1)$ and explain when separate and joint tests disagree.
- ▣ **Report:** three statistics, three decisions and one sentence separating accuracy from calibration

C2 [Proposed]: Critique an AI Answer — Task

- ▣ **Worked example to follow:** A3 and B5 [Solved]; a code-reading exercise, no tick data are needed
- ▣ **Setting:** a student gives an AI assistant a pandas Series `px` of clean, time-sorted SPY trade prices in New York time, 09:30–16:00, several days, in USD (so returns are decimal, not %)
- ▣ **Tasks**
 1. The answer on the next slide contains three conceptual errors; everything else is correct under this setting. Find them.
 2. For each error give the evidence: a derivation, a check on the output, or a number from the lecture or this seminar.
 3. Give the correction: formula, code line and the corrected number where one is asked.
 4. Log each error you actually found in the format of `AI_ERRORS.md`: request, answer, error, how found, correction.
- ▣ **Report:** a table with one row per error (what, evidence, correction)
- ▣ Hints: A3 on $\mathbb{E}[RV^{(n)}]$ with noise; what happens between 16:00 and 09:30; plug in $RV_t = 1.0 \times 10^{-4}$ ($1\%^2$)

C2 [Proposed]: Critique an AI Answer — the Answer

- **Request:** “Compute the daily realised variance, the annualised volatility and the HAR regressors from px.”
- **Answer, text**
 - ▶ $RV_t = \sum_i r_{t,i}^2$ with $r_{t,i}$ the intraday log returns of day t ; use every tick: more returns always make RV_t more precise
 - ▶ RV_t is SPY's daily variance: compare it directly with squared close-to-close returns and with the GARCH variance
 - ▶ HAR: regress RV_{t+1} on RV_t and on its 5-day and 22-day averages [Corsi \(2009\)](#)

```
import numpy as np, pandas as pd
lp = np.log(px)
r = lp.groupby(lp.index.date).diff().dropna()
RV = (r**2).groupby(r.index.date).sum()
vol_ann = np.sqrt(RV) * 252
H = pd.DataFrame({'d': RV,
                  'w': RV.rolling(5).mean(),
                  'm': RV.rolling(22).mean(),
                  'y': RV.shift(-1)}).dropna()
```

Take-Home Choice

Take-home: choose one

- ▣ B2, jumps in Bitcoin (model: B1), or B6, forecasting the next week (model: B5, A8)
- ▣ Write it as a Quantlet folder (code, `Metainfo.txt`, the chart): the format the team project uses
- ▣ AI use is allowed; note the errors you actually catch, in the format of `AI_ERRORS.md`
- ▣ Extension packs are optional

Extension Packs: What Each One Needs

Pack	Exercises	Needs
1. Noise	A3 extended [Solved], A4, B3 extended [Solved]	A3, B3
2. Jumps	A2 extended, B1 extended [Solved], B2	A2, B1
3. Distributions	B4	B3 (block bootstrap)
4. HAR	A5 extended [Solved], B5 extended [Solved], A7, B7	A5, B5, B1
5. Forecasting	A6 extended [Solved], B6, B6 extended, B9	A6, A8, the evaluation workflow
6. Roughness	B8	B3, Chapter 9 (rough volatility)
7. Projects	C1, C3	B5, A6, B9, the lecture case study

- Each exercise names its notebook section; within a pack, work in order
- C1 and C3 can seed the team project

A3 Extended [Solved]: Estimating the Noise Variance — Task

▣ Questions

- ▶ how can ω^2 be estimated from the data?
- ▶ in which limit is each estimator consistent?

▣ **Inputs:** the model of A3, n returns per day

▣ Tasks

1. Show that $\mathbb{E}[\text{RV}^{\text{all}}/(2n)] = \omega^2 + \text{IV}/(2n)$ and state its limit as $n \rightarrow \infty$.
2. Propose a second estimator from the first-order autocovariance and state when it is consistent.

▣ Notebook: section A3


A3 Extended [Solved]: Solution

- $\mathbb{E}[\text{RV}^{\text{all}} \mid \sigma] = \text{IV} + 2n\omega^2$, so $\mathbb{E}[\text{RV}^{\text{all}}/(2n)] = \omega^2 + \text{IV}/(2n)$; a law of large numbers for the noise terms gives $\hat{\omega}^2 = \text{RV}^{\text{all}}/(2n) \xrightarrow{P} \omega^2$ as $n \rightarrow \infty$ [Bandi & Russell \(2008\)](#)
- $-\widehat{\text{cov}}(r_i, r_{i-1})$ estimates ω^2 (A3), consistent as the number of returns grows, provided the noise is i.i.d. and independent of the price
- In finite samples the first estimator is contaminated by $\text{IV}/(2n)$; A4 compares it with the oracle value

A4 [Proposed]: Two-Scale RV — Task

▣ **Worked example to follow:** A3 [Solved]

▣ **Inputs:** a synthetic illiquid stock with a stipulated 390-minute session: $IV = IQ = 1$, $\omega = 0.03\%$; one day: 1-minute returns ($n = 390$) give $RV^{\text{all}} = 1.70$; $K = 5$ shifted 5-minute grids give $RV^{\text{avg}} = 1.10$




1. Compute the noise bias $2n\omega^2$ at 5 and 1 minutes, and n^* of A3 with its interval.
2. With $\bar{n} = (n - K + 1)/K$, compute $TSRV = RV^{\text{avg}} - \frac{\bar{n}}{n} RV^{\text{all}}$, with and without the factor $1/(1 - \bar{n}/n)$.
3. Estimate ω^2 twice: the oracle $(RV^{\text{all}} - IV)/(2n)$ and the feasible $RV^{\text{all}}/(2n)$ of A3 extended.
4. Show that TSRV removes the noise bias to first order. Optional: with $\text{Var}(\text{error}) \approx aK/n + bn/K^2$, find the K that balances the terms and the rate of TSRV.

B3 Extended [Solved]: A Realised Kernel on 5-Minute Bars — Task

- ▣ **Question:** does a noise-robust estimator change SPY's 5-minute RV?
- ▣ **Data:** SPY returns R and prices P ; helper `realized_kernel` [Barndorff-Nielsen et al. \(2008\)](#): Parzen weights, bandwidth $H = \lceil 3.5134 \xi^{4/5} n^{3/5} \rceil$ per day ($\lceil \cdot \rceil$: round up), noise-to-signal ratio $\xi^2 = \hat{\omega}^2 / \widehat{IV}$, the course grid endpoints (open and close)
- ▣ **Tasks**
 1. Compute the kernel for every day and the ratio of its mean to mean RV, with the block bootstrap of B3.
 2. Plot the distribution of $\ln(RV^{(k)} / RV^{(5)})$ for 15, 30 and 65 minutes.
- ▣ **Notebook:** section B3

B3 Extended [Solved]: Solution



- Kernel / RV = 0.952, CI [0.910, 0.987] (median bandwidth $H = 7$): the kernel adds back small negative autocorrelations ($\hat{\rho}_1 = -0.008$); significant but economically minor
- Coarser grids scatter around the 5-minute value: sparse RV adds noise without removing a detectable bias


A2 Extended [Proposed]: Why BV Ignores a Jump — Task

- ▣ **Worked example to follow:** A1 [Solved]
- ▣ **Inputs:** one jump of size κ in an **interior** interval j ($1 < j < M$): $r_j = \kappa + r_j^c$; the other returns are diffusive, $O_p(\Delta^{1/2})$
- ▣ **Tasks**
 1. Show that to leading order the jump enters BV only through $|r_{j-1}| |\kappa| + |\kappa| |r_{j+1}|$, and that these products are $O_p(\Delta^{1/2})$.
 2. Conclude that $BV \rightarrow IV$ while $RV \rightarrow IV + \kappa^2$; how does a jump in the last interval (as in A2) differ?

B1 Extended [Solved]: Multiplicity and Intraday Seasonality — Task

- ▣ **Question:** how many of B1's rejections survive a false-discovery-rate control and a correction for the intraday U-shape?
- ▣ **Data:** the $N = 1\,489$ daily p -values of B1; helpers `bh_count`, `periodicity_sd`
- ▣ **Tasks**
 1. Benjamini–Hochberg [Benjamini & Hochberg \(1995\)](#): sort the p -values and reject the k smallest, k the largest index with $p_{(k)} \leq kq/N$; apply $q = 5\%$ and 1% , and Bonferroni ($p < 0.05/N$).
 2. Periodicity factor [Boudt, Croux & Laurent \(2011\)](#): scale returns by $\sqrt{\text{BV}_t/M_t}$, take the root mean square of each 5-minute slot, normalise so that the mean of f_i^2 is 1; divide $r_{t,i}$ by f_i and repeat B1(2) and step 1.
- ▣ **Interpretation:** BH controls the FDR for valid p -values that are independent or positively dependent [Benjamini & Yekutieli \(2001\)](#); which count would you use to build a jump regressor?

B1 Extended [Solved]: Solution



- BH at 5%: the largest k with $p_{(k)} \leq k \times 0.05/1489$ is $k = 42$ ($p_{(k)} = 0.00140$ against 0.00141): 42 jump days; at 1%: 9; Bonferroni: 7
- After dividing by f_i ($f = 1.78$ at the open, 0.73 at the minimum): 167 rejections at 5%, 20 at 0.1%, 17 under FDR 5%
- **Interpretation:** build the jump regressor from FDR-controlled, periodicity-corrected days

B2 [Proposed]: Jumps in Bitcoin — Task

- **Worked example to follow:** B1 [Solved]
- **Data:** Bitcoin price matrix P_b , 24-hour UTC days with at least 95% of the 288 bars, 659 days of 747; a day is labelled by its UTC date
- **Tasks**
 1. Apply the ratio test of B1 at 0.1% and build a daily table: date, M_t , RV, BV, z , weekend flag.
 2. For weekdays and weekends report the number of days, the number of rejections and their rate, and the mean BV/RV.
 3. Compute the jump share of total variation and a one-sided binomial test of the overall rate against 0.001.
- **Interpretation:** the binomial test assumes independent, correctly sized daily tests
 - ▶ what does a rejection establish?
 - ▶ what does it not establish?
 - ▶ which group rejects more often?
 - ▶ why might that group reject more often?

B4 [Proposed]: Returns Standardised by Realised Volatility — Task

- ▣ **Worked example to follow:** B3 [Solved] (block bootstrap)
- ▣ **Data:** SPY open-to-close returns (1 489 days) and Bitcoin 24-hour returns (659 days), each aligned with the same-day RV Andersen et al. (2001); one row: 7 April 2025, $r_t = 3.09\%$, $RV_t = 45.1$, $z_t = r_t / \sqrt{RV_t} = 0.46$
- ▣ **Tasks**
 1. Compute the Pearson kurtosis $\mathbb{E}(x - \bar{x})^4 / s^4$ of $u_t = r_t / \text{sd}(r)$ and of z_t , with 95% intervals from a paired moving-block bootstrap (20-day blocks, 2 000 resamples).
 2. Plot histograms and Normal QQ plots of u_t and z_t ; run Jarque–Bera on z_t ; compute the share of $|z_t| > 3$ (Normal: 0.27%).
- ▣ **Interpretation:** where do the heavy tails come from? Optional benchmark: for i.i.d. Gaussian intraday returns z_t has kurtosis $3M/(M + 2)$

A5 Extended [Solved]: The Wald Test of the HAR Restrictions — Task

▣ **Worked example to follow:** A5 [Solved]

▣ **Tasks**

1. Write the Wald statistic for $R\theta = 0$ with a Newey–West covariance $\hat{V}_{NW}$ and give its limiting distribution under HAR.
2. Explain why a HAC covariance is needed for the log-variance regression.

A5 Extended [Solved]: Solution

- ▣ $W = (R\hat{\theta})^\top (R\hat{V}_{\text{NW}}R^\top)^{-1}R\hat{\theta} \rightarrow \chi_{19}^2$ under the 19 HAR restrictions
- ▣ The errors of the log-variance regression are heteroskedastic and may be autocorrelated: the usual OLS covariance would misstate $\text{Var}(R\hat{\theta})$
- ▣ The data test is B5 extended, step 1

B5 Extended [Solved]: The AR(22) Test and HARQ — Task

- ▣ **Worked example to follow:** A5 extended, B5 [Solved]
- ▣ **Data:** as in B5; realised quarticity RQ_t from the intraday returns (it does not cover the overnight return)
- ▣ **Tasks**
 1. Estimate the unrestricted AR(22) of $\ln V_t$ and test the 19 HAR restrictions with a Newey–West Wald test.
 2. Estimate HARQ in levels [Bollerslev, Patton & Quaadvlieg \(2016\)](#),
$$V_{t+1} = \beta_0 + (\beta_d + \beta_Q \sqrt{RQ_t})V_t + \beta_w V_t^{(w)} + \beta_m V_t^{(m)} + \varepsilon$$
, on intraday RV and on total variance.
 3. Forecast from 3 January 2022 on an expanding window, replacing forecasts that fall outside the $[\min, \max]$ range of the in-sample values by the sample mean [Bollerslev, Patton & Quaadvlieg \(2016\)](#), and run DM against HAR in levels under QLIKE.
- ▣ **Notebook:** section B5

B5 Extended [Solved]: Solution

- ▣ (1) $W = 18.1$, $p = 0.51$ (χ^2_{19}); R^2 0.477 (AR(22)) vs 0.470 (HAR): the restrictions are not rejected; HAR saves 19 parameters
- ▣ (2) $\hat{\beta}_Q = -0.0088$ (SE 0.0042) on intraday RV, significant at 5%; -0.0022 (SE 0.0024) on total variance, not significant
- ▣ (3) QLIKE 0.344 (HARQ) vs 0.326 (HAR), DM $t = 1.23$, $p = 0.22$: no significant difference
- ▣ **Interpretation:** the sign predicted by A7 appears, but with 5-minute RV the measurement error is small and the overnight return, not covered by RQ, dominates the noise in total variance

A7 [Proposed]: Attenuation and the Case for HARQ — Task

- ▣ **Worked example to follow:** A5 [Solved], B5 extended
- ▣ **Inputs:** $RV_t = V_t^* + u_t$ (latent signal plus error), V_t^* the integrated variance; $\mathbb{E}[u_t | \sigma] = 0$, $\text{Var}(u_t | \sigma) = 2 IQ_t / M$ (asymptotic, A1); errors serially uncorrelated and uncorrelated with V^* and η ; $V_{t+1}^* = \beta_0 + \beta V_t^* + \eta_{t+1}$, $\beta > 0$
- ▣ **Tasks**
 1. Using $\text{plim } \hat{\beta} = \text{cov}(RV_{t+1}, RV_t) / \text{Var}(RV_t)$, derive the attenuation of the OLS slope.
 2. Illustration (signal variance $\sigma_{V^*}^2$ assumed not to depend on IQ_t): derive the best linear weight on RV_t given IQ_t and show that it falls with IQ_t .
 3. Write an instrumental-variables estimator with instrument RV_{t-1} and state its exclusion and relevance conditions.

B7 [Proposed]: Do Jumps Help? The HAR-CJ Model — Task

- ▣ **Worked example to follow:** B5 and B1 [Solved]
- ▣ **Data:** SPY as in B5; replication baseline: B1's 0.1% jump days, without corrections; $J_t = RV_t - BV_t$ on those days, 0 otherwise; $C_t = RV_t - J_t + r_{on,t}^2$ (all in %²) [Andersen, Bollerslev & Diebold \(2007\)](#)
- ▣ **Tasks**
 1. Regress $\ln V_{t+1}$ on $\ln C_t$, $\ln C_t^{(5)}$, $\ln C_t^{(22)}$, $\ln(1 + J_t)$, $\ln(1 + J_t^{(5)})$, $\ln(1 + J_t^{(22)})$, where (5), (22) are arithmetic means logged afterwards; report all coefficients with Newey–West SEs and test the three jump coefficients jointly.
 2. Forecast from 3 January 2022 on an expanding window and compare QLIKE with log-HAR on common dates by DM.
 3. Sensitivity: repeat with the FDR-controlled, periodicity-corrected jump days of B1 extended.
- ▣ **Interpretation:** HAR-CJ averages levels before logging, log-HAR averages logs: the two are not nested; why can a significant jump coefficient give no forecasting gain?

A6 Extended [Solved]: Why MSE on Volatility Is Not Robust — Task

- ▣ **Worked example to follow:** A6 [Solved]
- ▣ **Inputs:** constant IV and $RV = IV \chi_M^2 / M$; supplied moment: $\mathbb{E}\sqrt{\chi_M^2} = \sqrt{2} \Gamma(\frac{M+1}{2}) / \Gamma(\frac{M}{2})$, Γ : the gamma function
- ▣ **Tasks**
 1. Minimise $\mathbb{E}(\sqrt{RV} - \sqrt{F})^2$ (MSE on volatility) over F and compare the minimiser with $\mathbb{E}RV$ by Jensen's inequality.
 2. Compute F^*/IV for $M = 1$ (squared daily return), 6 and 78.

A6 Extended [Solved]: Solution

- $\mathbb{E}(\sqrt{RV} - \sqrt{F})^2 = \mathbb{E}RV - 2\sqrt{F}\mathbb{E}\sqrt{RV} + F$, minimal at $\sqrt{F} = \mathbb{E}\sqrt{RV}$; by Jensen $(\mathbb{E}\sqrt{RV})^2 < \mathbb{E}RV$: the optimal forecast is biased down
- $F^*/IV = \frac{2}{M} \left(\Gamma(\frac{M+1}{2}) / \Gamma(\frac{M}{2}) \right)^2 = 0.637$ ($M = 1$), 0.920 ($M = 6$), 0.994 ($M = 78$); plotted in A6
- **Interpretation:** against squared daily returns, MSE on volatility rewards a forecast 36% too low; with 5-minute RV the distortion almost disappears

B6 [Proposed]: Forecasting the Next Week — Task

- ▣ **Worked example to follow:** B5 [Solved], A8, the evaluation workflow
- ▣ **Data:** SPY total variance V_t (B5) and SPY adjusted closes since 2000 for GARCH; target $Y_t = V_t + \dots + V_{t+4}$, forecast at the close of $t - 1$; 1 176 origins from 3 January 2022
- ▣ **Tasks**
 1. Direct log-HAR: regress $\ln Y$ on the B5 regressors, re-estimated at every origin using only targets fully observed by $t - 1$ (the last four are not); forecast $\exp(\hat{y} + \hat{s}^2/2)$ [Duan \(1983\)](#).
 2. GARCH(1,1)- t re-estimated every 21 days: with σ_t^2 the one-step variance, ω , α , β the GARCH constant, reaction and persistence coefficients and $\bar{\sigma}^2 = \omega/(1 - \alpha - \beta)$ the long-run variance, forecast $\sum_{h=0}^4 [\bar{\sigma}^2 + (\alpha + \beta)^h (\sigma_t^2 - \bar{\sigma}^2)]$.
 3. Compute QLIKE and MSE, DM on QLIKE with HAC errors, both MZ coefficients with the joint test of $(a, b) = (0, 1)$, and plot the cumulative QLIKE difference.
- ▣ **Interpretation**
 - ▶ does the one-day advantage of log-HAR survive at a weekly horizon?
 - ▶ what do overlapping targets do to the DM test?

B6 Extended [Proposed]: A Conditional Giacomini–White Test — Task

▣ **Worked example to follow:** B6

▣ **Tasks**

1. Re-estimate both methods on fixed-length rolling windows (log-HAR: last 250 targets; GARCH: last 1 000 returns), as the theory of [Giacomini & White \(2006\)](#) requires.
2. With instruments $Z_t = (1, d_{t-5})$ and the loss difference d_{t+5} realised five days later, compute $n \bar{g}^\top \hat{\Omega}^{-1} \bar{g}$ with $g = Z_t d_{t+5}$ (n : number of origins, $\bar{g}$: the mean of g , $\hat{\Omega}$: its long-run covariance); under the null g is uncorrelated beyond lag 4, so $\hat{\Omega}$ sums the unweighted autocovariances of lags 0–4.

B9 [Proposed]: Which Models Survive? The Model Confidence Set — Task

- ▣ **Worked example to follow:** B5 extended, B6
- ▣ **Data:** SPY one-day forecasts (log-HAR, HAR, GARCH(1,1)-t, EWMA, yesterday's RV) plus HARQ (B5 extended) and SHAR [Patton & Sheppard \(2015\)](#): $\beta_d V_t$ replaced by $\beta^+ RS_t^+ + \beta^- RS_t^-$, $RS_t^\pm = \sum_i r_{t,i}^2 1\{r_{t,i} \gtrless 0\}$, overnight return counted as one more return; 1 180 common dates
- ▣ **Tasks**
 1. Build the date $\times$ model QLIKE matrix on common dates and check that every forecast is positive.
 2. Run the MCS [Hansen, Lunde & Nason \(2011\)](#) with the range statistic, 20-day blocks, 10 000 resamples, seed 42: `MCS(L, size=0.10, reps=10000, block_size=20, method='R', seed=42)`.
 3. Report the elimination order and every model's MCS p -value; compare the sets at 90% (size 0.10) and 75% (size 0.25).
- ▣ **Interpretation:** why is the MCS preferable to 21 pairwise DM tests?

B8 [Proposed]: How Rough Is Volatility? — Task

- ▣ **Worked example to follow:** B3 [Solved] (block bootstrap), [Chapter 9](#)
- ▣ **Data:** SPY daily series of four measures: total variance, intraday RV, BV, 30-minute subsampled RV; $x_t = \frac{1}{2} \ln(\text{measure}_t)$; lags $\Delta = 1, \dots, 30$ trading days [Gatheral, Jaisson & Rosenbaum \(2018\)](#)
- ▣ **Tasks**
 1. Compute $m(q, \Delta)$, the mean of $|x_{t+\Delta} - x_t|^q$ over all valid pairs, for $q \in \{0.5, 1, 1.5, 2, 3\}$; check $m(2, 1)$ for total variance.
 2. For each q regress $\ln m(q, \Delta)$ on $\ln \Delta$ with an intercept: slope ζ_q ; then $H = \text{slope of } \zeta_q \text{ on } q \text{ through the origin.}$
 3. 95% CI: moving-block bootstrap, 60-day blocks, 300 resamples, increments computed only within each block.
 4. Simulate a Gaussian random walk (increments $N(0, 0.1^2)$, seed 42), remove the line joining its endpoints (endpoint-detrended, still $H = 0.5$), estimate H , add i.i.d. $N(0, 0.25^2)$ noise, estimate again.
- ▣ **Interpretation**
 - ▶ is H below 0.5 for every measure?
 - ▶ what does measurement error do?
- ▣ One path is an illustration, not a size or power study

C1 [Proposed]: Is Bitcoin's Volatility More Predictable Than the S&P 500's? — Task

- **Data:** SPY and Bitcoin one-day forecasts, common calendar window 1 March 2025 – 18 September 2026 (567 calendar days: SPY 389 trading days, Bitcoin 422 days)
- **Tasks**
 1. Define three measures and compute them per asset: the calibration R^2 of $\ln V_t$ on $\ln F_t$; the QLIKE gain $1 - \overline{\text{QLIKE}}_{\text{log-HAR}} / \overline{\text{QLIKE}}_{\text{RV yest.}}$; the same gain over GARCH.
 2. Build a 95% CI for the R^2 difference (Bitcoin minus SPY): resample the same 28-calendar-day blocks for both assets (2 000 times), keeping each asset's own missing days.
 3. Add a dummy for a weekend target day to the Bitcoin log-HAR, forecast again and run DM under QLIKE with and without it.
- **Report:** the table, the CI, the DM result and a two-sentence answer; project extension: Ethereum or a BVB stock, and the entropy measure of [Pele & Mazurencu-Marinescu-Pele \(2019\)](#)

C3 [Proposed]: Does Machine Learning Beat HAR for SPY? — Project Brief (1/4)

- ▣ Context: Christensen, Siggaard & Veliyev (2023) (Chapter 9, case study)
 - ▶ random forests and neural networks beat HAR once predictors beyond past RV enter
 - ▶ the gain grows with the horizon
 - ▶ you test this on SPY, with the sample split and the model settings of the paper
- ▣ (a) Build the target series (Section 4 of the paper).
 - ▶ daily open-to-close RV of SPY from the 5-minute bars of the regular session (78 returns a day, as in B1)
 - ▶ three targets: RV_{t+1} , the average RV over the next 5 days and over the next 22 days
- ▣ (b) Build two information sets dated t .
 - ▶ $\mathcal{M}_{\text{HAR}}$: the daily, weekly and monthly RV
 - ▶ $\mathcal{M}_{\text{ALL}}$ adds the VIX, the squared daily log return of the Hang Seng and the 1-week SPY momentum
 - ▶ and the first differences of log dollar volume and of the 3-month T-bill rate (FRED, DTB3)
 - ▶ and the economic policy uncertainty index (FRED, USEPUINXD) and the Aruoba–Diebold–Scotti index (Philadelphia Fed)

C3 [Proposed]: Does Machine Learning Beat HAR for SPY? — Project Brief (2/4)

- (b, continued) Date each predictor by its release, not its reference date.
 - ▶ use only values published by the close of day t
 - ▶ the Aruoba–Diebold–Scotti index is revised: use the Philadelphia Fed's real-time vintages
 - ▶ disclose any series for which vintages are unavailable
 - ▶ state which predictors of the paper (Table 1) cannot be built for an index fund, and why: a data limitation, not a reason to change the method
- (c) Split the sample as in Appendix A.1 of the paper.
 - ▶ the first 1000 usable days for training, the next 200 for validation, the rest for the test
 - ▶ a row belongs to a subset only if its whole target window lies inside it: drop the last $h - 1$ rows before each boundary
 - ▶ report the usable sample size per horizon
 - ▶ standardise every predictor with the training mean and variance only

C3 [Proposed]: Does Machine Learning Beat HAR for SPY? — Project Brief (3/4)

- ▣ (d) Estimate the linear models.
 - ▶ HAR, log-HAR and OLS on all predictors, on a rolling window over training and validation merged
 - ▶ log-HAR forecast: $\exp(\hat{f} + \hat{\sigma}^2/2)$, $\hat{f}$: fitted log RV, $\hat{\sigma}^2$: residual variance
 - ▶ elastic net on a rolling training window, without merging; λ and the mixing weight chosen on the validation set (paper, Section 2)
 - ▶ HAR keeps its three RV regressors under both information sets
- ▣ (d, continued) Estimate the machine-learning models (paper, Table A.6).
 - ▶ random forest: 500 trees on the merged window, $\lfloor J/3 \rfloor$ candidate predictors per split (J : number of predictors), minimum node size 5
 - ▶ NN_2^{10} : layers of 4 and 2 neurons, leaky ReLU, Adam with learning rate 0.001, dropout 0.8
 - ▶ early stopping with patience 100; the average of the best 10 of 100 seeds; trained once on a fixed window
 - ▶ replace negative forecasts with the smallest in-sample RV

C3 [Proposed]: Does Machine Learning Beat HAR for SPY? — Project Brief (4/4)

- (e) Report the test-set comparison, as in B9.
 - ▶ a table of MSE and QLIKE relative to HAR, for the three horizons and both information sets
 - ▶ one-sided DM p -values against HAR and the 90% MCS
 - ▶ overlapping 5- and 22-day targets: HAC errors with a bandwidth above $h - 1$ (Bartlett weights down-weight lag $h - 1$) and a sensitivity check over bandwidths
- Interpretation
 - ▶ does the ranking of Tables 2, 3 and 7 of the paper survive on SPY?
 - ▶ does the ML advantage grow with the horizon?
 - ▶ does the ML advantage survive QLIKE?
 - ▶ with a test set of this length, what can a non-significant DM test tell you?
- Worked example to follow: B5, A6.
- Project extension
 - ▶ repeat on Bitcoin's 24-hour RV
 - ▶ compare across models the importance of RV and VIX from accumulated local effects (Figure 7 of the paper)

References (1/2)

- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X. (2007). [Roughing it up: including jump components in the measurement, modeling, and forecasting of return volatility](#). *Review of Economics and Statistics*, 89(4), 701–720.
- ▣ Andersen, T.G., Bollerslev, T., Diebold, F.X., Labys, P. (2001). [The distribution of realized exchange rate volatility](#). *Journal of the American Statistical Association*, 96(453), 42–55.
- ▣ Bandi, F.M., Russell, J.R. (2008). [Microstructure noise, realized variance, and optimal sampling](#). *Review of Economic Studies*, 75(2), 339–369.
- ▣ Barndorff-Nielsen, O.E., Hansen, P.R., Lunde, A., Shephard, N. (2008). [Designing realized kernels to measure the ex post variation of equity prices in the presence of noise](#). *Econometrica*, 76(6), 1481–1536.
- ▣ Barndorff-Nielsen, O.E., Shephard, N. (2004). [Power and bipower variation with stochastic volatility and jumps](#). *Journal of Financial Econometrics*, 2(1), 1–37.
- ▣ Benjamini, Y., Hochberg, Y. (1995). [Controlling the false discovery rate: a practical and powerful approach to multiple testing](#). *Journal of the Royal Statistical Society, Series B*, 57(1), 289–300.
- ▣ Benjamini, Y., Yekutieli, D. (2001). [The control of the false discovery rate in multiple testing under dependency](#). *The Annals of Statistics*, 29(4), 1165–1188.
- ▣ Bollerslev, T., Patton, A.J., Quaedvlieg, R. (2016). [Exploiting the errors: a simple approach for improved volatility forecasting](#). *Journal of Econometrics*, 192(1), 1–18.
- ▣ Boudt, K., Croux, C., Laurent, S. (2011). [Robust estimation of intraweek periodicity in volatility and jump detection](#). *Journal of Empirical Finance*, 18(2), 353–367.
- ▣ Christensen, K., Siggaard, M., Veliyev, B. (2023). [A machine learning approach to volatility forecasting](#). *Journal of Financial Econometrics*, 21(5), 1680–1727.
- ▣ Corsi, F. (2009). [A simple approximate long-memory model of realized volatility](#). *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Diebold, F.X., Mariano, R.S. (1995). [Comparing predictive accuracy](#). *Journal of Business & Economic Statistics*, 13(3), 253–263.

References (2/2)

- ▣ Duan, N. (1983). Smearing estimate: a nonparametric retransformation method. *Journal of the American Statistical Association*, 78(383), 605–610.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- ▣ Gatheral, J., Jaisson, T., Rosenbaum, M. (2018). Volatility is rough. *Quantitative Finance*, 18(6), 933–949.
- ▣ Giacomini, R., White, H. (2006). Tests of conditional predictive ability. *Econometrica*, 74(6), 1545–1578.
- ▣ Hansen, P.R., Lunde, A., Nason, J.M. (2011). The model confidence set. *Econometrica*, 79(2), 453–497.
- ▣ Huang, X., Tauchen, G. (2005). The relative contribution of jumps to total price variance. *Journal of Financial Econometrics*, 3(4), 456–499.
- ▣ Künsch, H.R. (1989). The jackknife and the bootstrap for general stationary observations. *Annals of Statistics*, 17(3), 1217–1241.
- ▣ Liu, L.Y., Patton, A.J., Sheppard, K. (2015). Does anything beat 5-minute RV? A comparison of realized measures across multiple asset classes. *Journal of Econometrics*, 187(1), 293–311.
- ▣ Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- ▣ Patton, A.J., Sheppard, K. (2015). Good volatility, bad volatility: signed jumps and the persistence of volatility. *Review of Economics and Statistics*, 97(3), 683–697.
- ▣ Pele, D.T., Mazurencu-Marinescu-Pele, M. (2019). Using high-frequency entropy to forecast Bitcoin's daily value at risk. *Entropy*, 21(2), 102.
- ▣ Zhang, L., Mykland, P.A., Aït-Sahalia, Y. (2005). A tale of two time scales: determining integrated volatility with noisy high-frequency data. *Journal of the American Statistical Association*, 100(472), 1394–1411.