



Modelling Financial Markets

Seminar 10: Market Microstructure — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AAPL: Apple Inc. (Nasdaq ticker)
- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model); in event studies: Abnormal Return
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ BET-TR: BET Total Return (index)
- ▣ BIC: Bayesian Information Criterion
- ▣ BNR: Banca Națională a României (the National Bank of Romania)
- ▣ BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- ▣ CBOE: Chicago Board Options Exchange
- ▣ CI: Confidence Interval
- ▣ EKOP: Easley–Kiefer–O'Hara–Paperman (model of informed trading)
- ▣ ETF: Exchange-Traded Fund
- ▣ GME: GameStop Corp. (NYSE ticker)
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ JPM: JPMorgan Chase & Co. (NYSE ticker)
- ▣ MA: Moving Average
- ▣ MSFT: Microsoft Corporation (Nasdaq ticker)
- ▣ MSTR: Strategy Inc., formerly MicroStrategy (Nasdaq ticker)
- ▣ NVDA: NVIDIA Corporation (Nasdaq ticker)
- ▣ OHLC: Open, High, Low, Close
- ▣ OLS: Ordinary Least Squares
- ▣ PIN: Probability of INformed trading
- ▣ SD: Standard Deviation
- ▣ SE: Standard Error
- ▣ US: United States
- ▣ UTC: Coordinated Universal Time (Temps Universel Coordonné)

Acronyms used in this material (2/2)

- **VECM: Vector Error-Correction Model**
- **VIX: CBOE Volatility IndeX**
- **XRP: XRP (the native token of the XRP Ledger)**



Today's Question and Required Outputs

- ▣ Question of the seminar: **how precisely can trading costs and liquidity be measured from public prices and volumes?**
- ▣ Route
 - ▶ Step 1: A1–A2, the Roll model on paper and its sampling error; B1, the intraday U-shape with a day bootstrap
 - ▶ Step 2: B2, spread estimators at two frequencies and a Monte Carlo check; B3, 25 assets; B4, illiquidity and the VIX with HAC errors; C2, critique of an AI answer
- ▣ Required outputs (checklist at the end of the core part, before the extension packs): the A answers, one figure and one table per B exercise, the corrected C2 code, one interpretation sentence per exercise
- ▣ Notebook: [open the seminar notebook in Colab](#); each task names its notebook section
- ▣ Seminar: Antoaneta Amza, PhD candidate, antoaneta.amza@csie.ase.ro

Exercise Map

Exercise	Question	Type	Worked example
A1	what do bid-ask bounces do to price changes?	Solved	–
A2	how noisy is one Roll estimate?	Proposed	A1
B1	is the intraday U-shape systematic?	Solved	–
B2	which frequency gives a credible spread?	Solved	–
B3	do BVB stocks look as liquid as US stocks?	Proposed	B2
B4	how strongly does illiquidity move with the VIX?	Solved	–
C2	what is wrong in an AI answer?	Proposed	A1, B2
Extension packs	A3–A10 (information models, execution, PIN), B5–B8, C1, C3 (optional, after the seminar)	both	named on each task

- **[Solved]:** full solution in the slides and the notebook, a model to copy; **[Proposed]:** you solve it, the solution is discussed in class
- Part A: derivations on paper; Part B: estimation and inference on real data, each ending with an interpretation question; Part C: open questions and AI critique

Data Used

Series	Content	Window	Units
SPY 5-minute bars	OHLC and volume; UTC time stamps converted to New York time, bars 09:30–15:55, only days with all 78 bars	12 Oct 2020 – 18 Sep 2026, 1 477 days	USD; volume in shares
Bitcoin 5-minute bars	OHLC, UTC time	31 Aug 2024 – 18 Sep 2026	USD
US, daily (7)	SPY, AAPL, MSFT, NVDA, JPM, GME, MSTR	19 Sep 2024 – 18 Sep 2026	USD; volume in shares
BVB, daily (11)	TLV, SNP, H2O, BRD, SNG, SNN, TGN, EL, DIGI, TEL and the BET ETF TVBETETF	same	lei; converted to USD at the BNR rate of the day
Crypto, daily (7)	BTC, ETH, SOL, XRP, Dogecoin, Cardano, Litecoin	same, 7 days a week	volume already in USD
VIX; BET-TR	VIX (CBOE implied-volatility index) and BET-TR (BET total-return index), daily closes	as in B4, B8	index points

- Daily series: open, high, low, close, adjusted close and volume; OHLC scaled by the ratio adjusted close / close (dividends and splits); traded value = split-adjusted close × volume; days with high = low or zero volume dropped
- BNR reference rate (RON per USD), official daily fixing; 1 bp = 0.01%

Setup: Run and Verify

```
# the data-loading functions are in the notebook'  
s Data section  
s = intraday_spy()          # SPY 5-minute bars  
print(len(s), s['date'].nunique())  
print(s.index[0], s.index[-1])  
b = intraday_btc()          # Bitcoin, UTC  
print(len(b), b['date'].nunique())  
r = returns('TLV', '2024-09-19')  
dv = dollar_volume('TLV', '2024-09-19')  
print(len(r), dv.median() / 1e6)
```

- Expected output
 - ▶ 115 206 SPY bars, 1477 days
 - ▶ first bar 12 Oct 2020 09:30, last 18 Sep 2026 15:55 (New York time)
 - ▶ 213 754 Bitcoin bars on 749 days (all bars loaded; the lecture's 208 596 counts only bars with a 5-minute return, i.e. whose previous bar is present)
 - ▶ 492 daily returns for Banca Transilvania; median traded value USD 3.90 million
- If a count differs, stop and check the series, the time zone and the start date


Notation Used in Every Exercise

- ▣ p_t trade (or bar close) price, m_t efficient price, $q_t = \pm 1$ trade sign, c half-spread, $s = 2c$ spread; relative spreads in bp
- ▣ $\gamma_0 = \text{Var}(\Delta p_t)$, $\gamma_1 = \text{Cov}(\Delta p_t, \Delta p_{t-1})$; T = number of price changes (78 closes give 77 changes); sample moments demeaned within the sample
- ▣ Intraday returns: log changes of the close within a day; the first bar runs from the open; $|r|$ in bp
- ▣ Amihud ratio $\text{ILLIQ} = \frac{1}{D} \sum_d |r_d| / \text{DVOL}_d$, in bp per USD 1 million [Amihud \(2002\)](#)
- ▣ **Day bootstrap**: resample whole days with replacement, recompute the statistic, take the 2.5% and 97.5% percentiles (1000 resamples, seed 42) [Efron \(1979\)](#)
- ▣ One worked line: 78 closes, $\hat{\gamma}_1 = -1.0 \cdot 10^{-8}$ in log units gives $\hat{s} = 2\sqrt{10^{-8}} = 2 \cdot 10^{-4} = 2 \text{ bp}$

A1 [Solved]: The Roll Model — Task

- **Question:** what does the bid-ask bounce do to the moments of price changes, and how do we recover the spread from them?
- **Data and assumptions**
 - ▶ Efficient price $m_t = m_{t-1} + \varepsilon_t$, ε_t i.i.d., mean 0, variance σ^2 ; trade price $p_t = m_t + c q_t$ [Roll \(1984\)](#)
 - ▶ Signs $q_t = \pm 1$, i.i.d. with probability 1/2, independent of ε
 - ▶ Numbers: a BVB stock at about 20 lei, daily $\text{Var}(\Delta p) = 0.0520$, $\text{Cov}(\Delta p_t, \Delta p_{t-1}) = -0.0081$ (lei²)
- **Tasks**
 1. Write Δp_t in terms of ε_t , q_t and q_{t-1} , then derive $\text{Var}(\Delta p_t)$ and $\text{Cov}(\Delta p_t, \Delta p_{t-1})$.
 2. From the two numbers, compute c , s in lei and in % of the price, σ^2 , the bounce share $2c^2/\text{Var}(\Delta p)$ and the lag-1 autocorrelation ρ_1 .
 3. Extension: let the signs be a stationary Markov chain with $\text{Corr}(q_t, q_{t-k}) = \rho^k$, $\text{Var}(q_t) = 1$, still independent of ε ; expand the four covariance terms, and find the limit of $2\sqrt{-\hat{\gamma}_1}$ for $\rho = 0.3$.
- **Report:** the two moments and a table: c , s (lei), s (%), σ^2 , bounce share, ρ_1 , and the extension limit
- **Interpretation:** does persistent order flow make the Roll spread too high or too low?

A1 [Solved]: Picture the Bounce Before the Algebra



- Simulated with the A1 numbers ($c = 0.09$ lei, $\sigma^2 = 0.0358$), 20 000 trades: the trade price jumps between $m_t + c$ and $m_t - c$ around the efficient price
- Sample lag-1 autocorrelation of Δp_t : -0.150, close to the theoretical value derived in the solution; no autocorrelation at longer lags; notebook: section A1

A1 [Solved]: Solution (1/2) — Moments and Numbers

- Step 1. $\Delta p_t = \varepsilon_t + c(q_t - q_{t-1})$: the efficient innovation plus the change of side
- Step 2. Independence: $\text{Var}(\Delta p_t) = \sigma^2 + c^2(1 + 1) = \sigma^2 + 2c^2$
- Step 3. $\text{Cov}(\Delta p_t, \Delta p_{t-1}) = c^2 \text{Cov}(q_t - q_{t-1}, q_{t-1} - q_{t-2})$; the only common term is q_{t-1} , with a minus sign: $-c^2$
- Step 4. $c = \sqrt{0.0081}$; $\sigma^2 = 0.0520 - 2c^2$; $\rho_1 = -0.0081/0.0520$
- Interpretation: about a third of the daily variance is bid-ask noise, so volatility from trade prices overstates the efficient volatility

Quantity	Value
c	0.090 lei
$s = 2c$	0.18 lei
$s/20$	0.9%
$2c^2$	0.0162
σ^2	0.0358
bounce share	31.2%
ρ_1	-0.156

A1 [Solved]: Solution (2/2) — Correlated Trade Signs

- ▣ Step 5. With $\text{Cov}(q_t, q_{t-k}) = \rho^k$, expand the four terms of $\text{Cov}(q_t - q_{t-1}, q_{t-1} - q_{t-2})$
- ▣ $\text{Cov}(q_t, q_{t-1}) - \text{Cov}(q_t, q_{t-2}) - \text{Var}(q_{t-1}) + \text{Cov}(q_{t-1}, q_{t-2}) = \rho - \rho^2 - 1 + \rho = -(1 - \rho)^2$
- ▣ Step 6. $\text{Cov}(\Delta p_t, \Delta p_{t-1}) = -c^2(1 - \rho)^2 = -0.0081 \cdot 0.49 = -0.0040$
- ▣ Step 7. The half-spread estimator converges to $c(1 - \rho) = 0.063$ lei instead of 0.090; the spread estimator to 0.126 instead of 0.18 lei: both 30% too low
- ▣ Interpretation: persistent order flow (orders split into several trades) makes reversals rarer, hides part of the bounce and biases the Roll spread **downwards**

A2 [Proposed]: Precision of the Roll Estimator — Task

- ▣ **Worked example to follow:** A1
- ▣ **Question:** how much does the sample size T matter for one Roll estimate, holding the population moments fixed?
- ▣ **Data and assumptions**
 - ▶ Population moments of A1: $\gamma_0 = 0.0520$, $\gamma_1 = -0.0081$, so $c = 0.090$, $s = 0.18$ lei; estimator $\hat{s} = 2\sqrt{-\hat{\gamma}_1}$ Harris (1990)
 - ▶ T = number of price changes; Gaussian approximation (Bartlett, MA(1)): $T \text{Var}(\hat{\gamma}_1) \rightarrow \gamma_0^2 + 3\gamma_1^2$
- ▣ **Tasks**
 1. Derive $g'(\gamma_1)$ for $g(\gamma_1) = 2\sqrt{-\gamma_1}$ and write $\text{se}(\hat{s})$ by the delta method.
 2. Fill the table below for $T = 78$ and $T = 250$ price changes.
 3. Approximate $P(\hat{\gamma}_1 > 0)$ with the Normal distribution and explain why binary signs and demeaning make the approximation rough.
- ▣ **Report:** the completed table
- ▣ **Interpretation:** how often would a true Roll market show a positive covariance in one sample of 78 changes?

T	$\text{se}(\hat{\gamma}_1)$	$\text{se}(\hat{s})$ (lei)	$\text{se}(\hat{s})/s$	$P(\hat{\gamma}_1 > 0)$
78	...	...	...	...
250	...	...	...	...


B1 [Solved]: The Intraday U-Shape of SPY — Task

- **Question:** are volatility and volume systematically higher at the open and the close, or is this driven by a few extreme days?
- **Data and assumptions**
 - ▶ SPY 5-minute bars, 1 477 full days (Data Used); $|r|$ in bp; bars are labelled by their start time
 - ▶ One row per day: $|r|$ of the first bar (09:30), the mean $|r|$ of the 31 midday bars starting 11:30 to 14:00 inclusive, and the share of the day's volume in the last six bars
- **Tasks**
 1. Compute the ratio of the cross-day means, $\overline{|r|}_{\text{open}}/\overline{|r|}_{\text{mid}}$, and the mean last-30-minute volume share.
 2. Build 95% day-bootstrap intervals for both statistics (1 000 resamples of complete rows, seed 42).
 3. Plot mean $|r|$ and the mean volume share by bar with pointwise day-bootstrap bands.
- **Report:** the two estimates with intervals and the two profiles
- **Interpretation:** why resample days rather than bars, and what does the pattern imply for a trader who must buy 1 million shares today?

B1 [Solved]: Solution (1/2) — Resampling Days, Step by Step

- ▣ Step 1. Point estimates: $\overline{r}|_{\text{open}} = 10.6$ bp, $\overline{r}|_{\text{mid}} = 5.1$ bp, ratio 2.06; last 30 minutes 18.6% of daily volume against $6/78 = 7.7\%$ if even
- ▣ Step 2. One resample: draw 1 477 day-rows with replacement (here 933 distinct days appear), keep each row whole, recompute
 - ▶ First resample: 10.7 bp and 5.0 bp, ratio 2.13
- ▣ Step 3. Repeat 1 000 times; the 2.5% and 97.5% percentiles: ratio [1.98, 2.15], volume share [18.3, 19.0]%
- ▣ Why whole days: the bars of one day share its volatility and news, so resampling days keeps that dependence; days are still treated as independent, although volatility persists

B1 [Solved]: Solution (2/2) — Volatility and Volume Profiles



- Both intervals are narrow and exclude the even benchmarks (ratio 1, 7.7%): a systematic pattern, not a few extreme days; the bands are pointwise, not a simultaneous band for the whole curve
- Interpretation: a large buyer finds most volume at the open and the close; the open is also the most volatile, so the best time also depends on spreads, impact, urgency and the risk of waiting


B2 [Solved]: Spread Estimators at Two Frequencies — Task (1/2)

- **Question:** which sampling frequency gives a credible bid-ask spread for SPY, the most traded ETF?
- **Data and assumptions**
 - ▶ Intraday: the 1477 days of SPY 5-minute bars; per day 78 closes give 77 log returns
 - ▶ Daily: SPY daily bars, 19 Sep 2024 – 18 Sep 2026 (501 days); one tick = USD 0.01
- **Recipes (log prices; all spreads relative)**
 - ▶ Roll: per day the demeaned lag-1 autocovariance of the 77 returns, averaged over days, then $2\sqrt{-\bar{\gamma}_1}$, set to 0 if $\bar{\gamma}_1 \geq 0$ [Roll \(1984\)](#)
 - ▶ Corwin–Schultz: the two-bar estimate for each pair of consecutive bars (negatives set to 0), averaged; daily bars add the overnight adjustment [Corwin & Schultz \(2012\)](#)
 - ▶ Abdi–Ranaldo: the mean of $4(c_t - \eta_t)(c_t - \eta_{t+1})$, then the square root, 0 if negative [Abdi & Ranaldo \(2017\)](#)

B2 [Solved]: Spread Estimators at Two Frequencies — Task (2/2)


- **Question:** which sampling frequency gives a credible bid-ask spread for SPY, the most traded ETF?
- **Tasks**
 1. Compute the three estimators from the 5-minute bars with 95% day-bootstrap intervals (resample precomputed daily moment rows, not prices).
 2. Repeat on the same two-year window as the daily bars, so that only the frequency differs.
 3. Compute the three estimators from daily bars with a moving-block bootstrap (blocks of 20 days, 1 000 resamples), and report how often each is truncated at 0.
 4. Compare every estimate with one tick relative to the price, a floor for the quoted (not the effective) spread.
 5. Simulate the Roll model with SPY's bar variance, 78 prices a day, 1 477 days per sample, 200 samples, for $c = \text{one tick}$ and $c = 0$; compute the share of days with $\hat{\gamma}_1 > 0$ and the pooled estimate [Harris \(1990\)](#).
 6. Correct the pooled 5-minute Roll estimate for the demeaning bias $\mathbb{E}[\hat{\gamma}_1] \approx \gamma_1 - (\gamma_0 + 2\gamma_1)/77$.
- **Report:** an interval plot of the estimates, the truncation shares, and the Monte Carlo results next to the observed ones
- **Interpretation:** which frequency gives a credible spread, what do the daily estimators measure, and is the share of positive covariances evidence against Roll?

B2 [Solved]: Solution (1/2) — Estimates and Intervals



- 5-minute (2020–2026): Roll 2.82 [2.09, 3.43], Corwin–Schultz 3.11 [3.03, 3.20], Abdi–Ranaldo 1.06 [0.00, 1.67] bp; same two years (496 days): 1.95, 2.83, 0.00 bp
- Daily: Roll 64.3 [0.0, 114.4], Corwin–Schultz 26.5 [21.3, 33.7], Abdi–Ranaldo 55.1 [0.0, 92.2] bp; truncated at 0 in 9.7% (Roll) and 3.4% (Abdi–Ranaldo) of resamples; 5-minute Abdi–Ranaldo in 11.2%
- One tick is 0.21 bp: daily estimates are two orders of magnitude above it, because the daily SD is 103.4 bp; even 5-minute Corwin–Schultz is 15 ticks

B2 [Solved]: Solution (2/2) — Monte Carlo under the Roll Model



- One tick: $c = 0.10$ bp against a bar SD of 9.2 bp; share of positive days 45.3% [43.0, 48.1], with $c = 0$ 45.3%: indistinguishable; observed 40.8%, below both simulated reference intervals
- Pooled Roll: 2.07 bp (one tick), 2.06 bp ($c = 0$), observed 2.82 bp; bias correction $\hat{\gamma}_1^c = (\bar{\gamma}_1 + \bar{\gamma}_0/77)/(1 - 2/77)$ gives 1.92 bp
- Interpretation: only intraday bars give credible magnitudes; positive covariances are expected under Roll, so they do not reject it; the remaining negative autocorrelation points to transient price pressure. These are simulated reference intervals, not bootstrap CIs


B3 [Proposed]: Spread Estimates Across 25 Assets — Task

- **Worked example to follow:** B2
- **Question:** do daily spread proxies rank BVB stocks as liquid as US stocks, and does the Amihud ratio agree?
- **Data and assumptions**
 - ▶ The 25 assets listed on the Data Used slide: 7 US, 10 BVB stocks plus the BET ETF (11), 7 crypto; daily bars, 19 Sep 2024 – 18 Sep 2026
 - ▶ Estimators as in B2 (daily recipes); volatility annualised with 252 days (stocks) or 365 (crypto); Amihud in bp per USD 1 million, crypto volume already in USD, BVB converted at the BNR rate
- **Tasks**
 1. Compute Corwin–Schultz and Abdi–Rinaldo for each asset, and the median of each estimator in each group.
 2. Count, per group, the assets for which Abdi–Rinaldo is set to 0.
 3. Compute the Spearman correlations of Corwin–Schultz with volatility, of Amihud with the median traded value, and of Corwin–Schultz with Amihud, and draw the three scatter plots.
- **Report:** a table with rows US, BVB, crypto and columns Corwin–Schultz median, Abdi–Rinaldo median, Abdi–Rinaldo zeros; the three correlations and plots
- **Interpretation:** which measure would you trust to compare markets, and what would you need to validate it?

B4 [Solved]: Illiquidity and the VIX — Task

- **Question:** by how much does SPY's intraday illiquidity rise when implied volatility rises, and how precise is that elasticity?
- **Data and assumptions**
 - ▶ $ILLIQ_d = \frac{1}{78} \sum_k 10^4 |r_{d,k}| / (DV_{d,k} / 10^8)$: bp per USD 100 million, from the 5-minute bars ($DV = \text{close} \times \text{volume}$); bars with zero value dropped
 - ▶ Aligned on trading days with the close of the VIX index: 1477 days, 2020-10-12 – 2026-09-18; mean ILLIQ 1.83
- **Tasks**
 1. Estimate $\ln ILLIQ_d = a + b \ln VIX_d + e_d$ by OLS and plot the fit.
 2. Plot the residual autocorrelations; report OLS and HAC (Newey–West, 20 lags) standard errors and the 95% HAC interval [Newey & West \(1987\)](#).
 3. Extension: add a linear trend t (in years from the first day) and re-estimate b .
 4. Extension: with $D_d = 1$ from 2023-09-27 (the middle of the sample, fixed before estimation), add $gD_d + hD_d \ln VIX_d$ and test $h = 0$ with a HAC Wald test at 5%.
- **Report:** a table of specifications ($\hat{b}$, SE, interval) and the coefficient plot
- **Interpretation:** why do the two standard errors differ, is the elasticity stable, and is the relation causal?

B4 [Solved]: Solution (1/2) — Fit and Residuals



- (1) $\hat{b} = 0.95$, $R^2 = 0.30$: a 10% higher VIX goes with about $0.95 \times 10\%$ higher illiquidity; (2) residual lag-1 autocorrelation 0.76, decaying slowly: OLS SE 0.038, HAC SE 0.121 (3.2 times larger), HAC CI [0.72, 1.19]

B4 [Solved]: Solution (2/2) — Specifications and the Break Test

Specification	$\hat{b}$	SE	95% interval
Baseline, OLS SE	0.95	0.038	—
Baseline, HAC SE	0.95	0.121	[0.72, 1.19]
With trend (HAC)	0.80	0.120	—
Before 2023-09-27 (HAC)	1.03	0.090	—
From 2023-09-27 (HAC)	0.60	0.238	—

- Trend -0.062 a year (SE 0.021): illiquidity falls over time, the VIX link remains; break: $\hat{h} = -0.43$, Wald $(\hat{h}/SE)^2 = (-0.43/0.25)^2 = 2.90$, $p = 0.089$: not rejected at 5%
- Interpretation: persistent residuals make OLS errors far too small; the elasticity is close to 1; non-rejection at 5% does not establish stability; not causal, since both respond to news

C2 [Proposed]: Critique an AI Answer — the Answer

- **Prompt:** “I have daily closes P (lei) and volumes V (number of shares, positive every day) for a BVB stock. Estimate the Roll spread and the Amihud illiquidity. What does a round trip cost me?”
- **Answer, text:** Roll: $\text{Cov}(\Delta p_t, \Delta p_{t-1}) = -c^2$ and $s = 2c$; if the covariance is positive, take its absolute value, the formula still gives the spread; Amihud: the average of $|r_d|$ divided by the volume of day d , per 1 million; result: $s = 0.05$ lei on a 20-lei stock, 0.25% of the price, so a round trip costs $2s$, 0.50%
- **Answer, code**

```
import numpy as np
dp = P.diff().dropna().values
g = np.cov(dp[1:], dp[:-1])[0, 1]
s_roll = 2 * np.sqrt(abs(g))
s_pct = 100 * s_roll / P.mean()
r = np.log(P).diff().dropna()
illiq = 1e6 * (r.abs() / V.loc[r.index]).mean()
```

C2 [Proposed]: Critique an AI Answer — Task

- ▣ **Worked example to follow:** A1, B2
- ▣ **Question:** which parts of the AI answer are wrong, and how would you detect each error before using the numbers?
- ▣ **Data and assumptions**
 - ▶ A conceptual audit: no data file is needed; P and V are hypothetical
 - ▶ Three checks to run first: the sign of the Roll covariance (A1), the units of the Amihud denominator (Notation), the bid, mid and ask of a round trip
- ▣ **Tasks**
 1. The answer contains three errors; identify each one in the text and, where applicable, in the code.
 2. For each error, state how you would detect it: a test, a published number or a formula from the lecture.
 3. Correct each error: give the formula, the code line and the corrected number where the information permits one; state the units of your Amihud ratio.
 4. Log the errors you actually found in `AI_ERRORS.md`: request, answer, error, how found, correction.
- ▣ **Report:** the three errors with detection and correction
- ▣ **Interpretation:** which error would change an investment decision?

Submission Checklist and Team Project

Hand in before the next seminar

- ☐ A1: the moments and the table; A2: the completed table
- ☐ B1, B2, B4: the figures and tables of the solutions, reproduced; B3: your table and the three plots
- ☐ C2: the three errors, detection, corrected code
- ☐ One interpretation sentence per exercise

Team project: choose one

- ☐ C1: is illiquidity priced on the BVB? C3: the implied price of the speed race for SPY; or: the three high-low estimators for three assets not covered here (documenting the data you use)
- ☐ Common rubric: question, data and conventions, method, uncertainty (intervals, robustness), limitations, reproducible Quantlet folder
- ☐ Declare AI use in AI_USE.md; log every AI error you actually found in AI_ERRORS.md

Extension Packs: What Each One Needs

Pack	Exercises	Prerequisites
1. Information models	A3 [Solved], A4, A5 [Solved], A6	lecture section <i>Information and Prices</i> , Bayes' rule, projections
2. Optimal execution	A7 [Solved], A8	lecture section <i>Price Impact and Optimal Execution</i>
3. Informed trading (PIN)	A9 [Solved], A10	Poisson likelihood, log-sum-exp
4. Illiquidity across markets	B5	B2, B3
5. Volume and price impact	B6 [Solved]	B1, instrumental-variable (IV) estimation
6. The Bitcoin clock	B7	B1
7. Price discovery	B8 [Solved]	cointegration, VECM
8. Projects	C1, C3	B4, B5; the lecture case study


- Each pack has its own notebook section; do them in order within a pack; C1 and C3 are possible team projects

A3 [Solved]: Glosten–Milgrom for Any Prior — Task

- **Question:** how wide is the spread when the market maker learns from order flow, and when is it widest?
- **Data and assumptions**
 - ▶ $V \in \{V_L, V_H\}$, $\Delta V = V_H - V_L$; $\mathcal{F}_t$ = the trades observed so far, $\theta = P(V = V_H \mid \mathcal{F}_t)$ [Glosten & Milgrom \(1985\)](#)
 - ▶ Likelihoods: $P(\text{buy} \mid V_H) = (1 + \mu)/2$, $P(\text{buy} \mid V_L) = (1 - \mu)/2$ (informed share μ , uninformed buy or sell with probability $1/2$)
 - ▶ Pre-trade quotes for the next trade: $a = \mathbb{E}[V \mid \mathcal{F}_t, \text{buy}]$, $b = \mathbb{E}[V \mid \mathcal{F}_t, \text{sell}]$; $V_L = 90$, $V_H = 110$, $\mu = 0.3$
- **Tasks**
 1. At $\theta = 0.5$, compute $P(V_H \mid \text{buy})$, the ask, the bid and the spread.
 2. After one buy, the posterior becomes the new θ : compute the next-trade quotes and spread.
 3. Derive the spread for any $0 < \theta < 1$ and show that it is largest at $\theta = 1/2$; check it at $\theta = 0.7$.
 4. Extension: show that $\mathbb{E}[\text{next trade price} \mid \mathcal{F}_t] = \mathbb{E}[V \mid \mathcal{F}_t]$ and that beliefs are a martingale.
- **Report:** the quotes, the spread formula and the martingale argument
- **Interpretation:** why is the spread largest when the market maker is most uncertain?

A3 [Solved]: Solution — Quotes and the General Spread

- Step 1 ($\theta = 0.5$).
 $P(V_H | \text{buy}) = 0.65 \cdot 0.5 / (0.65 \cdot 0.5 + 0.35 \cdot 0.5) = 0.6500$:
 $a = 90 + 20 \cdot 0.6500 = 103.00$, $b = 97.00$, spread $6.00 = \mu \Delta V$
- Step 2. After a buy $\theta = 0.6500$: $a = 105.50$, $b = 100.00$, spread 5.50
- Step 3. With $x = 2\theta - 1$: $P(V_H | \text{buy}) = \theta(1 + \mu)/(1 + \mu x)$,
 $P(V_H | \text{sell}) = \theta(1 - \mu)/(1 - \mu x)$
- Step 4. $a - b = \mu \Delta V (1 - x^2)/(1 - \mu^2 x^2) \leq \mu \Delta V$, equality only at $x = 0$; at $\theta = 0.7$: 5.11
- Step 5. $P(\text{buy})a + P(\text{sell})b = \mathbb{E}[V | \mathcal{F}_t]$ (total expectation); iterated expectations give the martingale
- Interpretation: at $\theta = 1/2$ one trade moves beliefs most, so adverse selection costs most



A4 [Proposed]: How Fast Does the Market Maker Learn? — Task

- ▣ **Worked example to follow:** A3
- ▣ **Question:** how many consecutive buys does it take before the bid exceeds 108, and how does this depend on μ ?
- ▣ **Data and assumptions**
 - ▶ A3 with $\theta_0 = 1/2$, $V_L = 90$, $V_H = 110$; true value V_H ; a **conditioned scenario**: every trade is a buy, not the expected learning time
 - ▶ Timeline: after k buys the belief is θ_k ; the bid quoted then conditions on a hypothetical next sell
- ▣ **Tasks**
 1. Show that each buy multiplies the odds $\theta/(1 - \theta)$ by $(1 + \mu)/(1 - \mu)$ and each sell divides them by the same factor.
 2. Show that the bid after k buys is $V_L + \Delta V \theta_{k-1}$.
 3. Show that bid > 108 requires odds above 9, derive $k_{\min} = \lfloor \ln 9 / \ln \frac{1+\mu}{1-\mu} \rfloor + 2$, and evaluate it for $\mu = 0.1, 0.3, 0.5$, checking the boundary case.
 4. Compare with the simulated paths of the lecture, where uninformed traders also sell.
- ▣ **Report:** the three values of $k_{\min}$ and a plot of the bid against k
- ▣ **Interpretation:** how does the speed of price discovery depend on the share of informed traders?

A5 [Solved]: Deriving the Kyle Equilibrium — Task

- **Question:** how aggressively does an insider trade, and how much does the market maker move the price per share?
- **Data and assumptions**
 - ▶ Timeline: (1) the insider sees $v \sim N(p_0, \sigma_v^2)$ and sends x ; (2) noise traders send $u \sim N(0, \sigma_u^2)$, independent; (3) the market maker sees only $y = x + u$ and sets $p = \mathbb{E}[v \mid y]$ [Kyle \(1985\)](#)
 - ▶ Units: v, p in USD; x, u, y in shares; β in shares per USD; λ in USD per share. Numbers: $\sigma_v = 2$, $\sigma_u = 10\,000$ and $20\,000$
- **Tasks**
 1. Insider best response: given $p = p_0 + \lambda y$, maximise the expected profit and write β in $x = \beta(v - p_0)$.
 2. Market-maker pricing: given $x = \beta(v - p_0)$, derive λ from the projection theorem.
 3. Solve the two equations; compute the expected profit, $\text{Var}(x)$ and $\text{Var}(v \mid y)$.
 4. Fill a two-column table ($\sigma_u = 10\,000$ and $20\,000$): β , λ , profit, price change for $y = 15\,000$.
- **Report:** the equilibrium formulas and the table
- **Interpretation:** why does more noise trading help the insider?

A5 [Solved]: Solution — Equilibrium and Numbers

- Step 1. $\mathbb{E}[(v - p_0 - \lambda x - \lambda u)x \mid v] = (v - p_0)x - \lambda x^2$; first-order condition: $x = (v - p_0)/(2\lambda)$, $\beta = 1/(2\lambda)$
- Step 2. (v, y) jointly Normal: $\lambda = \text{Cov}(v, y)/\text{Var}(y) = \beta\sigma_v^2/(\beta^2\sigma_v^2 + \sigma_u^2)$
- Step 3. With $2\lambda\beta = 1$: $\beta^2\sigma_v^2 = \sigma_u^2$, so $\beta = \sigma_u/\sigma_v$, $\lambda = \sigma_v/(2\sigma_u)$
- Step 4. Profit $\beta\sigma_v^2 - \lambda\beta^2\sigma_v^2 = \sigma_v\sigma_u/2$; $\text{Var}(x) = \sigma_u^2$; $\text{Var}(v \mid y) = \sigma_v^2/2$ in both economies
- Interpretation: λ measures informativeness per share; more noise makes the flow less informative, the insider trades more and earns more, and half the information still enters the price

	$\sigma_u = 10\,000$	$\sigma_u = 20\,000$
β	5 000	10 000
λ	0.0001	0.00005
profit (USD)	10 000	20 000
$\Delta p, y = 15\,000$	1.50	0.75

A6 [Proposed]: Kyle's Lambda as a Regression Slope — Task

- ▣ **Worked example to follow:** A5
- ▣ **Question:** what does an Amihud-type ratio estimate in a Kyle economy, and does the choice of volume denominator matter?
- ▣ **Data and assumptions**
 - ▶ A5 with $p_0 = 50$ USD, $\sigma_v = 2$, $\sigma_u = 10\,000$; simulation (seed 42): 200 000 independent auctions, $v = p_0 + \sigma_v z_1$, $u = \sigma_u z_2$, $x = \beta(v - p_0)$, $y = x + u$, $p = p_0 + \lambda y$
 - ▶ $|x| + |u|$ is a constructed gross-order proxy (the two orders counted separately), not exchange volume
- ▣ **Tasks**
 1. Show that the population OLS slopes of $p - p_0$ and of $v - p_0$ on y both equal λ , and give $\text{Var}(y)$.
 2. Show that $|r|/(p_0|y|)$, with $r = (p - p_0)/p_0$, equals λ/p_0^2 , and convert it to bp per USD 1 million (factor $10^4 \times 10^6$).
 3. Replace $|y|$ by $|x| + |u|$, derive the direction of the bias and measure it by simulation; repeat with $\sigma_u = 20\,000$.
- ▣ **Report:** the slopes, the ratio for both denominators and both noise levels
- ▣ **Interpretation:** what does the Amihud ratio measure in a Kyle world, and why is it biased with gross volume?

A7 [Solved]: Almgren–Chriss Execution — Task


- **Question:** how should a risk-averse seller split 1 million shares over five days, and what does the risk reduction cost?
- **Data and assumptions**
 - ▶ Parameters [Almgren & Chriss \(2001\)](#): $X = 10^6$ shares, $T = 5$ days, $N = 5$, $\tau = 1$; $\sigma = 0.95$ USD per day^{1/2}; temporary impact $\eta = 2.5 \cdot 10^{-6}$, permanent $\gamma = 2.5 \cdot 10^{-7}$ (USD per share per share); risk aversion $\lambda = 2 \cdot 10^{-6}$ per USD
 - ▶ x_k = holdings after trade k ($x_0 = X$, $x_5 = 0$), $n_k = x_{k-1} - x_k$; $\mathbb{E}[C] = \frac{1}{2}\gamma X^2 + \frac{\tilde{\eta}}{\tau} \sum_{k=1}^5 n_k^2$, $\text{Var}[C] = \sigma^2 \tau \sum_{k=1}^5 x_k^2$, $\tilde{\eta} = \eta - \gamma\tau/2$
- **Tasks**
 1. Compute $\tilde{\eta}$, $\tilde{\kappa}^2 = \lambda\sigma^2/\tilde{\eta}$ and κ from $\cosh(\kappa\tau) = 1 + \tilde{\kappa}^2\tau^2/2$.
 2. Fill a table of k , $x_k = X \sinh(\kappa(5-k))/\sinh(5\kappa)$ and n_k .
 3. Compute $\mathbb{E}[C]$ and $\text{sd}[C]$ and compare with selling evenly ($\lambda = 0$).
 4. Extension: minimise $\int_0^T (\eta\dot{x}^2 + \lambda\sigma^2 x^2) dt$ with $x(0) = X$, $x(T) = 0$ by Euler–Lagrange and compare κ_c with κ .
- **Report:** the schedule table, the cost–risk comparison and the holdings plot
- **Interpretation:** what does the trader buy with the extra expected cost?

A7 [Solved]: Solution (1/2) — Schedule, Cost and Risk

- Step 1. $\tilde{\eta} = 2.5 \cdot 10^{-6} - 1.25 \cdot 10^{-7} = 2.375 \cdot 10^{-6}$; $\tilde{\kappa}^2 = 0.76$;
 $\cosh \kappa = 1.38$, $\kappa = 0.8463$, $\sinh(5\kappa) = 34.40$
- Step 2. The table: 57% is sold on day 1
- Step 3. $\mathbb{E}[C] = 1.078$, $\text{sd}[C] = 0.449$ million USD; evenly: 0.600 and 1.041
- Interpretation: front-loading pays 0.48 million USD more expected impact to cut the SD of cost by more than half

k	x_k	n_k
0	1 000 000	—
1	428 599	571 401
2	182 933	245 666
3	76 296	106 637
4	27 643	48 652
5	0	27 643

A7 [Solved]: Solution (2/2) — Euler–Lagrange and the Frontier



□ $F = \eta \dot{x}^2 + \lambda \sigma^2 x^2$: $\frac{d}{dt} \partial F / \partial \dot{x} = \partial F / \partial x \Rightarrow \ddot{x} = \kappa_c^2 x$, $\kappa_c^2 = \lambda \sigma^2 / \eta = 0.722$; with $x(0) = X$, $x(T) = 0$:
 $x(t) = X \sinh(\kappa_c(T - t)) / \sinh(\kappa_c T)$, $x(1) = 427151$

- $\kappa_c = 0.8497$ against $\kappa = 0.8463$: close here because discretisation and $\tilde{\eta}$ partly offset ($\kappa T \approx 0.85$ is not small); dots: the discrete schedule, dashed: the continuous one


A8 [Proposed]: A More Risk-Averse Seller — Task

- ▣ **Worked example to follow:** A7
- ▣ **Question:** how does the schedule change when risk aversion is ten times higher?
- ▣ **Data and assumptions**
 - ▶ A7 parameters: $X = 10^6$, $T = N = 5$, $\sigma = 0.95$, $\eta = 2.5 \cdot 10^{-6}$, $\gamma = 2.5 \cdot 10^{-7}$; only λ changes, to $2 \cdot 10^{-5}$
- ▣ **Tasks**
 1. Compute $\tilde{\kappa}^2$, κ , κ_c and the holdings x_k .
 2. Compute $\mathbb{E}[C]$ and $\text{sd}[C]$ and the ratios new/A7, with their direction.
 3. Overlay the holdings of A7 and A8 on one chart.
 4. Optional: explain why κ and κ_c now differ much more.
- ▣ **Report:** one comparison table and the chart
- ▣ **Interpretation:** when would a trader choose this schedule?

A9 [Solved]: The PIN Likelihood, Step by Step — Task

- **Question:** how do we evaluate the likelihood of one day of buys and sells without numerical overflow, and what does it say about that day?
- **Data and assumptions**
 - ▶ EKOP [Easley et al. \(1996\)](#), three states: no news, prob. $1 - \alpha$, $(B, S) \sim \text{Poisson}(\varepsilon_b), \text{Poisson}(\varepsilon_s)$; bad news, $\alpha\delta$, sells rate $\mu + \varepsilon_s$; good news, $\alpha(1 - \delta)$, buys rate $\mu + \varepsilon_b$
 - ▶ $\alpha = 0.3$, $\delta = 0.5$, $\mu = 400$, $\varepsilon_b = \varepsilon_s = 1\,000$; one day with $B = 1\,450$, $S = 1\,000$
- **Tasks**
 1. Write the likelihood as a three-component mixture and compute $\text{PIN} = \alpha\mu / (\alpha\mu + \varepsilon_b + \varepsilon_s)$.
 2. Explain why evaluating the terms directly fails in double precision.
 3. Write each weighted log component, combine them as $m + \ln \sum e^{\ell_j - m}$ with $m = \max_j \ell_j$, and normalise the same components into posterior probabilities.
- **Report:** PIN, $\ln L$, the three posterior probabilities
- **Interpretation:** is the posterior probability of good news the same thing as PIN?

A9 [Solved]: Solution — Stable Log Components and Posteriors



- $\text{PIN} = 120/2120 = 0.0566$; directly, $e^{-1000} \rightarrow 0$ and $1000^{1450} \rightarrow \infty$: $0 \cdot \infty = \text{NaN}$ Lin & Ke (2011). Factor $K = e^{-\varepsilon_b - \varepsilon_s} (\mu + \varepsilon_b)^B (\mu + \varepsilon_s)^S / (B!S!)$, $x_b = \varepsilon_b / (\mu + \varepsilon_b)$, x_s likewise; $\ln B!$ by log-gamma
- Components: $\ell_1 = \ln(1 - \alpha) + B \ln x_b + S \ln x_s$, $\ell_2 = \ln \alpha \delta - \mu + B \ln x_b$, $\ell_3 = \ln \alpha(1 - \delta) - \mu + S \ln x_s$; $\ln L = \ln K + m + \ln \sum e^{\ell_j - m} = -11.71$; posterior $e^{\ell_j - m} / \sum e^{\ell_i - m}$: good news ≈ 1
- Interpretation: 450 extra buys against a Poisson SD of about 32 classify the day almost surely; $\text{PIN} = 0.0566$ is the unconditional expected share of informed trades, not this posterior

A10 [Proposed]: Reading Days with the PIN Model — Task

- ▣ **Worked example to follow:** A9
- ▣ **Question:** how confidently can known PIN parameters classify a day, and what changes when the parameters are unknown?
- ▣ **Data and assumptions**
 - ▶ The A9 parameters; two observed days (a) $B = 1\,000$, $S = 1\,420$; (b) $B = S = 1\,000$; (c) a hypothetical ambiguous day, $B = 1\,195$, $S = 1\,000$
- ▣ **Tasks**
 1. For each day, compute the stable log-likelihood and the three posterior probabilities.
 2. Compute PIN when μ doubles to 800.
 3. Conceptual: explain why estimating $(\alpha, \delta, \mu, \varepsilon_b, \varepsilon_s)$ from many days remains hard for very liquid stocks.
- ▣ **Report:** a table of days by posterior probabilities, and the new PIN
- ▣ **Interpretation:** can an order imbalance always be read as private information?


B5 [Proposed]: Amihud Illiquidity on Three Markets — Task

- **Worked example to follow:** B2, B3
- **Question:** how many times less liquid is the BVB than the US by the Amihud ratio, and how precise and stable is that multiple?
- **Data and assumptions**
 - ▶ The 25 assets of B3 (BVB includes the BET ETF); fixed group membership
 - ▶ Aggregation: average the daily ratios within each asset over 19 Sep 2024 – 18 Sep 2026, then take the group median
 - ▶ Years: 19 Sep 2024 – 18 Sep 2025 and 19 Sep 2025 – 18 Sep 2026
- **Tasks**
 1. Compute the three group medians and the ratio BVB median / US median.
 2. Build a 95% interval by resampling assets with replacement independently within each group (1000 resamples), recomputing the median ratio each time.
 3. Compute the Spearman correlation of the asset rankings in the two years.
- **Report:** the medians, the ratio with its interval, the rank correlation and log-scale plots
- **Interpretation:** what does the interval condition on, and is illiquidity a stable asset characteristic?

B6 [Solved]: Volume and Price Moves in SPY Bars — Task

- **Question:** how do absolute returns scale with relative volume in 5-minute bars, and can that slope be read as a price-impact law?
- **Data and assumptions**
 - ▶ SPY 5-minute bars with positive volume: 115 052 bars on 1 477 days (115 206 bars minus 154 without volume)
 - ▶ Relative volume $v = \text{bar volume} / \text{median volume at the same time of day}$; $z = |r|/\sigma_{\text{day}}$, σ_{day} the SD of that day's 5-minute returns
- **Tasks**
 1. Sort bars into 20 equal-size groups by v (bin edges fixed on the full sample), regress $\ln \bar{z}$ on $\ln \bar{v}$, and give a 95% day-bootstrap interval (300 resamples, same edges).
 2. At bar level, estimate $\ln(|r| + 10^{-6}) - \ln \sigma_{\text{day}} = a + b \ln v + e$ with SEs clustered by day, test $b = 0.5$, and repeat without the bars with $r = 0$.
 3. Extension: instrument $\ln v$ by the log relative volume of the same bar on the previous trading day; report the first stage and the slope estimated by two-stage least squares (2SLS), with clustered SE.
- **Report:** the grouped fit, the four slope estimates with intervals and the first stage
- **Interpretation:** is the relation concave, is it the square-root law [Tóth et al. \(2011\)](#), and is the exclusion restriction plausible?

B6 [Solved]: Solution — Fitted Relation and Slope Estimates



- Grouped slope 0.274 [0.262, 0.289]; bar-level 0.309 (SE 0.008), $t = (0.5 - 0.309)/0.008 \approx 24$: 0.5 rejected; 0.8% of bars have $r = 0$, without them 0.291 (SE 0.008)
- First stage slope 0.339 ($t = 32.8$): relevant; 2SLS slope 0.063 (SE 0.018); relevance does not establish exclusion: yesterday's volume at that time reflects persistent activity and scheduled news (lagged correlation of normalised $\ln |r|$ 0.043)
- Interpretation: concave but flatter than a square root; a bar mixes buyers and sellers, while the law concerns one trader's net metaorder, so slope 0.5 is not a direct test of it


B7 [Proposed]: The Bitcoin Clock — Task

- ▣ **Worked example to follow:** B1
- ▣ **Question:** does a market open 24/7 still keep the clock of traditional markets?
- ▣ **Data and assumptions**
 - ▶ Bitcoin 5-minute bars, UTC, 749 days (535 weekdays, 214 weekend days); returns spanning missing bars are discarded; coverage is incomplete: only 106 days have all 288 bars (minimum 26); all days are kept
 - ▶ Weighting: first the mean $|r|$ of each day in a window (hours $[h_1, h_2]$), then the mean over days; ratios are ratios of these day means
- ▣ **Tasks**
 1. On weekdays, compute the mean $|r|$ in 13:00–17:00 UTC and in 03:00–11:00 UTC and their ratio, with a 95% day-bootstrap interval.
 2. Compute the weekday / weekend ratio of mean $|r|$ with a 95% day-bootstrap interval (days resampled from the pooled sample).
 3. Plot the two hourly profiles with day-bootstrap bands.
- ▣ **Report:** the two ratios with intervals and the profiles
- ▣ **Interpretation:** what spread, depth and volume evidence would an execution decision need beyond volatility?

B8 [Solved]: Price Discovery: the BET ETF and BET-TR — Task


- **Question:** does the BET ETF or the BET-TR index lead price discovery, and how sure can we be with daily closes?
- **Data and assumptions**
 - ▶ Log closes $y_t = (\ln P_t^{\text{ETF}}, \ln P_t^{\text{idx}})'$ of Patria-TVBETETF (days with trades) and BET-TR, common days 2024-09-19 – 2026-09-18, 493 VECM observations; the ETF reinvests dividends; BET-TR is computed, not traded
 - ▶ VECM: $\Delta y_t = a_0 + \alpha z_{t-1} + \sum_{i=1}^p \Gamma_i \Delta y_{t-i} + e_t$, $z_{t-1} = \ln P_{t-1}^{\text{ETF}} - \ln P_{t-1}^{\text{idx}}$, $\text{Var}(e_t) = \Omega$; unrestricted constant; p lags of the *differences*, chosen by BIC over 1–10 [Johansen \(1991\)](#)
- **Tasks**
 1. Plot the rebased prices and the basis; run the Johansen trace test (rank one does not by itself establish $\beta = (1, -1)'$, which is imposed).
 2. Estimate the VECM and test $\alpha_{\text{ETF}} = 0$ and $\alpha_{\text{idx}} = 0$ at 5%.
 3. With $\alpha_{\perp} = (-\alpha_{\text{idx}}, \alpha_{\text{ETF}})'$ compute the component share $\alpha_{\perp} / \sum \alpha_{\perp}$ [Gonzalo & Granger \(1995\)](#) and the information share $([\alpha'_{\perp} F]_j)^2 / \alpha'_{\perp} \Omega \alpha_{\perp}$ for both Cholesky orderings F [Hasbrouck \(1995\)](#).
 4. Bootstrap: resample moving blocks of 20 regression rows $(\Delta y_t, 1, z_{t-1}, \Delta y_{t-1}, \dots)$, 1 000 times; then re-estimate α on each half.
- **Report:** the test, α with asymptotic and bootstrap intervals, the ordering bounds with endpoint intervals, the two halves
- **Interpretation:** which price leads, why are the bounds wide, and why can the component share leave $[0, 1]$?

B8 [Solved]: Solution (1/2) — Prices, Test and Adjustment



- BIC: 1 lag; trace 49.6 for $r = 0$ (5% critical value 15.49), 0.10 for $r \leq 1$ (3.84): one cointegrating vector
- $\alpha_{\text{ETF}} = -0.12$ ($t = -2.11$, rejected at 5%), $\alpha_{\text{idx}} = -0.03$ ($t = -0.63$); block bootstrap: ETF $[-0.29, 0.01]$, index $[-0.18, 0.09]$: the asymptotic test rejects, the bootstrap interval touches 0, so the evidence is weak

B8 [Solved]: Solution (2/2) — Shares, Bounds and Halves



- Index information share within [0.25, 0.98] (ordering bounds, not a CI); bootstrap CI of the lower bound [0.00, 0.64], of the upper [0.72, 1.00]; component share 1.39 [-0.14, 3.32]
- Halves: ETF $\alpha = -0.26$ ($t = -2.34$), then -0.09 ($t = -0.85$); the index is never significant ($t = -0.77, 0.47$)
- Interpretation: the ETF appears to adjust toward the index; residual correlation 0.93 makes the ordering decisive, hence wide bounds; both α negative push the component share outside [0, 1]

C1 [Proposed]: Does Illiquidity Carry a Premium on the BVB? — Task

- ▣ **Worked example to follow:** B4, B5
- ▣ **Question:** is illiquidity priced on the BVB as Amihud (2002) predicts: expected illiquidity raises future returns, an unexpected rise lowers current returns [Amihud \(2002\)](#)?
- ▣ **Data and assumptions**
 - ▶ Markets: BVB (return BET-TR; illiquidity from the ten blue chips), US (SPY; six US stocks), crypto (Bitcoin; six other coins); monthly log returns in %, 2016–2026
 - ▶ Index: for each asset and month the mean daily Amihud ratio (months with at least 10 trading days), then the cross-sectional mean of its log; shock = residual of an AR(1) of the index
- ▣ **Tasks**
 1. Milestone 1: build the index and its shock for each market.
 2. Milestone 2: estimate $r_t = a + b_1 L_{t-1} + b_2 \text{shock}_t + e_t$ with HAC SEs (6 lags).
 3. Milestone 3: correct b_1 for the Stambaugh bias [Stambaugh \(1999\)](#) with the reduced-bias estimator: $\hat{\phi}_c = \hat{\phi} + (1 + 3\hat{\phi})/T + 3(1 + 3\hat{\phi})/T^2$, residuals v_t^c from $\hat{\phi}_c$, regress r_t on L_{t-1} and v_t^c , SE including the uncertainty of $\hat{\phi}_c$ [Amihud & Hurvich \(2004\)](#).
 4. Milestone 4 (sorts): each January rank the blue chips with more than 150 trading days by the previous year's Amihud; hold equal-weighted illiquid and liquid halves for the year (monthly simple total returns); test the mean difference with HAC SEs.
- ▣ **Report:** a table per market, the reference charts and a discussion of power (few stocks, ten years) [Pele, Mazurencu-Marinescu & Nijkamp \(2013\)](#)
- ▣ **Interpretation:** which of Amihud's two predictions survives on each market?

C3 [Proposed]: The Price of the Speed Race in SPY — Task

- ▣ **Worked example to follow:** B1, B4
- ▣ **Question:** if the London latency-arbitrage tax applied to SPY, how large would the prize be, and how much would it move on stress days?
- ▣ **Data and assumptions**
 - ▶ Lecture case study [Aquilina, Budish & O'Neill \(2022\)](#): tax 0.419 bp (Table X); Table XV: col. 2 $\Pi_t = 0.4213 \cdot 10^{-4} V_t$; col. 6 $\Pi_t = 10^{-4}(0.3354 V_t + 0.0066 \sigma_t \bar{V})$; Table XIV scenarios 0.20–0.74 bp
 - ▶ Units: Π_t , V_t in USD; V_t = sum of close $\times$ volume over the 78 bars; $\sigma_t = 100 \sqrt{252 \sum r^2}$ in % (percentage points); $\bar{V}$ the mean of the year; SPY full days 2021–2025
 - ▶ Outputs are implications of transferred coefficients, not measured US race profits
- ▣ **Tasks**
 1. Daily inputs: compute V_t and σ_t ; report yearly means and the yearly correlation of V_t and σ_t with 95% intervals from an i.i.d. day bootstrap and a 20-day moving-block bootstrap; compare with the paper's 0.82.
 2. Annual extrapolation: compute each year's prize $\sum_t \Pi_t$ under col. 2, col. 6 and the two scenarios.
 3. Show that under col. 6 the annual tax is $0.3354 + 0.0066 \bar{\sigma}$ bp and find the $\bar{\sigma}$ above which col. 6 exceeds col. 2.
 4. Stress days: plot the tax at mean traded value, $0.3354 + 0.0066 \sigma_t$, and the realised ratio $\Pi_t/V_t = 0.3354 + 0.0066 \sigma_t \bar{V}/V_t$; work 7 April 2025 completely.
- ▣ **Report:** the annual table, the derivation and threshold, the daily chart and the worked day
- ▣ **Interpretation:** which assumptions carry a London 2015 tax over to SPY today, and can the data at hand sign their effect?

References (1/2)

- ▣ Abdi, F., Rinaldo, A. (2017). A simple estimation of bid-ask spreads from daily close, high, and low prices. *Review of Financial Studies*, 30(12), 4437–4480.
- ▣ Almgren, R., Chriss, N. (2001). Optimal execution of portfolio transactions. *Journal of Risk*, 3(2), 5–39.
- ▣ Amihud, Y. (2002). Illiquidity and stock returns: cross-section and time-series effects. *Journal of Financial Markets*, 5(1), 31–56.
- ▣ Amihud, Y., Hurvich, C.M. (2004). Predictive regressions: a reduced-bias estimation method. *Journal of Financial and Quantitative Analysis*, 39(4), 813–841.
- ▣ Aquilina, M., Budish, E., O'Neill, P. (2022). Quantifying the high-frequency trading “arms race”. *Quarterly Journal of Economics*, 137(1), 493–564.
- ▣ Bouchaud, J.-P., Bonart, J., Donier, J., Gould, M. (2018). *Trades, Quotes and Prices: Financial Markets Under the Microscope*. Cambridge University Press.
- ▣ Brauneis, A., Mestel, R., Riordan, R., Theissen, E. (2021). How to measure the liquidity of cryptocurrency markets? *Journal of Banking & Finance*, 124, 106041.
- ▣ Budish, E., Cramton, P., Shim, J. (2015). The high-frequency trading arms race: frequent batch auctions as a market design response. *Quarterly Journal of Economics*, 130(4), 1547–1621.
- ▣ Corwin, S.A., Schultz, P. (2012). A simple way to estimate bid-ask spreads from daily high and low prices. *Journal of Finance*, 67(2), 719–760.
- ▣ Easley, D., Kiefer, N.M., O'Hara, M., Paperman, J.B. (1996). Liquidity, information, and infrequently traded stocks. *Journal of Finance*, 51(4), 1405–1436.
- ▣ Efron, B. (1979). Bootstrap methods: another look at the jackknife. *Annals of Statistics*, 7(1), 1–26.

References (2/2)

- Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- Glosten, L.R., Milgrom, P.R. (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1), 71–100.
- Gonzalo, J., Granger, C.W.J. (1995). Estimation of common long-memory components in cointegrated systems. *Journal of Business & Economic Statistics*, 13(1), 27–35.
- Harris, L. (1990). Statistical properties of the Roll serial covariance bid/ask spread estimator. *Journal of Finance*, 45(2), 579–590.
- Hasbrouck, J. (1995). One security, many markets: determining the contributions to price discovery. *Journal of Finance*, 50(4), 1175–1199.
- Hasbrouck, J. (2007). *Empirical Market Microstructure*. Oxford University Press.
- Johansen, S. (1991). Estimation and hypothesis testing of cointegration vectors in Gaussian vector autoregressive models. *Econometrica*, 59(6), 1551–1580.
- Kyle, A.S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315–1335.
- Lin, H.-W.W., Ke, W.-C. (2011). A computing bias in estimating the probability of informed trading. *Journal of Financial Markets*, 14(4), 625–640.
- Newey, W.K., West, K.D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.
- Pele, D.T., Mazurencu-Marinescu, M., Nijkamp, P. (2013). Herding behaviour, bubbles and log periodic power laws in illiquid stock markets: a case study on the Bucharest Stock Exchange. SSRN Working Paper 2305945.
- Roll, R. (1984). A simple implicit measure of the effective bid-ask spread in an efficient market. *Journal of Finance*, 39(4), 1127–1139.
- Stambaugh, R.F. (1999). Predictive regressions. *Journal of Financial Economics*, 54(3), 375–421.
- Tóth, B., Lempérière, Y., Deremble, C., de Lataillade, J., Kockelkoren, J., Bouchaud, J.-P. (2011). Anomalous price impact and the critical nature of liquidity in financial markets. *Physical Review X*, 1(2), 021006.