



Modelling Financial Markets

Seminar 14: Deep Learning and Time-Series Foundation Models — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ ACF: AutoCorrelation Function
- ▣ ACI: Adaptive Conformal Inference
- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model)
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ BVB: Bursa de Valori București (the Bucharest Stock Exchange)
- ▣ CI: Confidence Interval
- ▣ CRPS: Continuous Ranked Probability Score
- ▣ CRSP: Center for Research in Security Prices
- ▣ CW: Clark–West (test for nested models)
- ▣ DM: Diebold–Mariano (test)
- ▣ DQ: Dynamic Quantile (test)
- ▣ ES: Expected Shortfall
- ▣ ETF: Exchange-Traded Fund
- ▣ FHS: Filtered Historical Simulation
- ▣ FRED: Federal Reserve Economic Data
- ▣ FZ0: Fissler–Ziegel loss, zero-homogeneous version
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GPD: Generalised Pareto Distribution
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HAR: Heterogeneous AutoRegressive (model)
- ▣ HS: Historical Simulation
- ▣ LR: Likelihood Ratio (test)
- ▣ LSTM: Long Short-Term Memory (network)
- ▣ MCS: Model Confidence Set
- ▣ MDE: Minimum Detectable Effect

Acronyms used in this material (2/2)

- ▣ **MSE: Mean Squared Error**
- ▣ **OLS: Ordinary Least Squares**
- ▣ **PIT: Probability Integral Transform**
- ▣ **QLIKE: Quasi-LIKElihood (loss function)**
- ▣ **QS: Quantile Score**
- ▣ **RFF: Random Fourier Features**
- ▣ **RV: Realised Variance**
- ▣ **SPA: Superior Predictive Ability (test)**
- ▣ **TS: Time Series (in the name TS-Arena)**
- ▣ **US: United States**
- ▣ **VAR: Vector AutoRegression**

Outline

Introduction

Key Formulas

Part A: Derivations

Part B: Inference on the Forecasts

Part C: Open Question and AI Critique

Wrap-up

References

Seminar Objectives

By the end of this seminar you can

- ▣ Derive vanishing gradients as VAR impulse responses and the LSTM cell state as an AR(1)
- ▣ Read VaR and ES from a quantile grid, with a GPD tail; compute the exact power of a backtest
- ▣ Test forecasts correctly: Clark–West, SPA, Romano–Wolf, DQ, PIT and quantile scores, conformal recalibration
- ▣ Detect seed effects, stale prices and look-ahead leakage

- ▣ Background: Lecture 14 and Lecture 8 (backtesting); textbook [Franke, Härdle & Hafner \(2019\)](#)

Structure

- ▣ **Part A:** ten derivations (A1–A10)
- ▣ **Part B:** eleven computational exercises with inference (B1–B11)
- ▣ **Part C:** an open question, a seed for the team project (C1), an AI answer to critique (C2), and a project seed from the lecture case study (C3)

How to Work in This Seminar

- **[Solved]** problems show the full worked solution; **[Proposed]** problems show only the task: their solutions are discussed in class
- Every proposed problem names its **model**: the solved problem that uses the same technique
 - ▶ Every Part B exercise ends with interpretation questions
- Always name the benchmark: zero for returns, log-HAR for volatility, FHS for tail risk
- A model that is "not significantly worse" is not "as good": ask how much power the test had

Solved and Proposed Problems

Solved: worked examples in this deck

- Part A: A1 (gradients as VAR impulse responses), A3 (mean scaling), A5 (VaR and ES from a grid), A8 (exact power of Kupiec), A10 (QLIKE elicits the mean)
- Part B: B1 (returns vs zero, Clark–West, SPA), B3 (realised volatility vs log-HAR), B5 (VaR 1% and ES 2.5%), B7 (seeds), B9 (PIT, Berkowitz, quantile score)

Proposed: you solve them

- Part A: A2 (model: A1), A4 (model: A3), A6 and A7 (model: A5), A9 (model: A8)
- Part B: B2 (model: B1), B4, B8, B11 (model: B3), B6 (model: B5), B10 (models: B5, B9)
- Part C: C1, open question and project seed; C2, critique of an AI answer (a foundation-model VaR backtest); C3, project seed from the lecture case study (the virtue of complexity in return prediction)

Setup: Python Environment and Data

-  **Open in Colab:** `chapter14_seminar_notebook.ipynb`
- Google Colab (recommended) or local Python ≥ 3.10 with `torch`, `chronos-forecasting`, `timesfm`, `arch`, `pandas`, `scipy`
- Series: S&P 500, Bitcoin, BET (daily closes); SPY 5-minute bars for realised variance; up to 18 September 2026
- Forecasts: the one-day-ahead forecasts of Lecture 14 (context of 512 days for the foundation models; windows of 500 or 1,000 days for the classical models)
- Returns r_t in %, losses $L_t = -r_t$; $\text{VaR}_\alpha = -q_\alpha$, $\text{ES}_\alpha = -\frac{1}{\alpha} \int_0^\alpha q_u du$

Key Formulas (1/2): Networks and Tokens

- LSTM (Hochreiter & Schmidhuber (1997)): $c_t = f_t \odot c_{t-1} + i_t \odot g_t$, $h_t = o_t \odot \tanh(c_t)$; gates $\sigma(Wx_t + Uh_{t-1} + b) \in (0, 1)$
- Parameters of one layer: $4(hd + h^2 + h)$, or $4(hd + h^2 + 2h)$ with two bias vectors (PyTorch)
- Memory along the cell state: $\partial c_T / \partial c_{T-k} = \prod_j f_{T-j}$; half-life $k_{1/2} = \ln 0.5 / \ln f$
- Attention (Vaswani et al. (2017)): $\text{softmax}(QK^\top / \sqrt{d_k}) V$
- Chronos tokens (Ansari et al. (2024)): $s = \frac{1}{\epsilon} \sum |x_t|$, $z_t = x_t / s$, bin index $j = \text{round}((z_t + 15) / \Delta)$, $\Delta = 30/4092$

Key Formulas (2/2): Evaluation

- Out-of-sample $R^2 = 1 - \sum_t (r_t - \hat{r}_t)^2 / \sum_t r_t^2$ (benchmark: zero forecast)
- Diebold–Mariano (Diebold & Mariano (1995)): $DM = \bar{d} / \sqrt{\widehat{LRV} / T}$, $d_t = L_t^A - L_t^B$, Newey–West long-run variance; negative: A better
- QLIKE (Patton (2011)): $RV_t / \hat{h}_t - \log(RV_t / \hat{h}_t) - 1$; HAR (Corsi (2009)): RV on its daily, 5-day and 22-day averages
- Kupiec (Kupiec (1995)): $LR_{uc} = -2[(T - x) \ln(1 - \alpha) + x \ln \alpha - (T - x) \ln(1 - \hat{\pi}) - x \ln \hat{\pi}] \sim \chi_1^2$
- FZ0 (Patton, Ziegel & Chen (2019)): $\frac{1}{\alpha \text{ES}} 1\{L > \text{VaR}\}(L - \text{VaR}) + \frac{\text{VaR}}{\text{ES}} + \log \text{ES} - 1$; MCS (Hansen, Lunde & Nason (2011)) at 90%

Key Formulas (3/3): Inference and Calibration

- Clark–West (Clark & West (2007)), nested benchmark: $\tilde{d}_t = e_{0,t}^2 - [e_{1,t}^2 - (\hat{y}_{0,t} - \hat{y}_{1,t})^2]$, one-sided HAC t
- Many models: SPA (Hansen (2005)), $\max_k \sqrt{T} \bar{d}_k / \hat{\omega}_k$ with a stationary bootstrap; Romano–Wolf stepwise adjusted p (Romano & Wolf (2005))
- DQ (Engle & Manganelli (2004)): regress $I_t - \alpha$ on $(1, I_{t-1}, \dots, I_{t-4}, \text{VaR}_t)$; $\hat{\beta}^\top X^\top X \hat{\beta} / (\alpha(1 - \alpha)) \sim \chi^2_6$
- PIT $u_t = \hat{F}_t(y_t)$ i.i.d. $U(0, 1)$ (Diebold, Gunther & Tay (1998)); Berkowitz LR on $\Phi^{-1}(u_t)$ (Berkowitz (2001)); QS $= \frac{2}{K} \sum_k \text{pinball}_{u_k}$, an equal-weight grid analogue of the CRPS (Gneiting & Raftery (2007))
- Adaptive conformal (Gibbs & Candès (2021)): $\alpha_{t+1} = \alpha_t + \gamma(\alpha - \text{err}_t)$; $\hat{Q} = +\infty$ when $\alpha_t \leq 0$; long-run breach rate $\rightarrow \alpha$ for any sequence

A1. Vanishing Gradients as VAR Impulse Responses [Solved]

Task

- **Context:** a simple recurrent network $h_t = \tanh(Wh_{t-1} + Ux_t + b)$ with two hidden units, trained by backpropagation through time; no data are needed
 - ▶ Given: $W = \begin{pmatrix} 0.5 & 0.8 \\ 0 & 0.5 \end{pmatrix}$; linearise at the fixed point $h^* = 0$ with $b = 0$ and $x = 0$
- (a) Derive the Jacobian $\partial h_T / \partial h_{T-k}$ with the chain rule and prove that $\|\partial h_T / \partial h_{T-k}\|_2 \leq \|W\|_2^k$
- (b) Show that the linearised network is a VAR(1) in h_t and that the Jacobian equals its impulse response at horizon k
- (c) Compute the spectral radius $\rho(W)$, the spectral norm $\|W\|_2$, the bound $\|W\|_2^k$ and the exact $\|W^k\|_2$ for $k = 5, 20, 50$
- **Report** the numbers of (c) and answer: is the norm bound or the spectral radius the right criterion for vanishing gradients?

A1. Vanishing Gradients as VAR Impulse Responses: solution (1/2)

Solution

- (a) One step: $\partial h_t / \partial h_{t-1} = D_t W$ with $D_t = \text{diag}(1 - h_t^2)$, since $\tanh'(z) = 1 - \tanh^2(z)$
 - ▶ Chain rule over k steps: $\partial h_T / \partial h_{T-k} = \prod_{j=0}^{k-1} D_{T-j} W$
 - ▶ $\|D_t\|_2 = \max_i |1 - h_{t,i}^2| \leq 1$ and $\|AB\|_2 \leq \|A\|_2 \|B\|_2$, hence the product is at most $\|W\|_2^k$
- (b) At $h^* = 0$: $\tanh'(0) = 1$, so $D_t = I$ and $h_t \approx Wh_{t-1} + Ux_t$
 - ▶ This is a VAR(1) with coefficient matrix W and innovation Ux_t
 - ▶ Its moving-average form $h_T = \sum_k W^k Ux_{T-k}$ shows that $\partial h_T / \partial h_{T-k} = W^k$ is the impulse response at horizon k

A1. Vanishing Gradients as VAR Impulse Responses: solution (2/2)

Solution

- (c) W is triangular: its eigenvalues are the diagonal entries, so $\rho(W) = 0.50$; the largest singular value gives $\|W\|_2 = 1.040$
 - ▶ Bound at $k = 20$: $1.040^{20} = 2.20$, which grows with k and says nothing
 - ▶ Exact: $W^k = 0.5^k \begin{pmatrix} 1 & 1.6k \\ 0 & 1 \end{pmatrix}$, so $\|W^5\|_2 = 0.254$, $\|W^{20}\|_2 = 3.1 \times 10^{-5}$, $\|W^{50}\|_2 = 7.1 \times 10^{-14}$
- Interpretation: the norm bound is sufficient, not necessary; the gradient vanishes at the rate $k 0.5^k$ because $\rho(W) < 1$ at this fixed point, where $D = I$
 - ▶ Gelfand's formula $\|W^k\|^{1/k} \rightarrow \rho(W)$: the same criterion as the stability of a VAR; away from $h^* = 0$ it is $\rho(DW)$ that matters, not $\rho(W)$

A2. The LSTM Cell State as an AR(1) [Proposed]

Solution discussed in class

Task

- **Context:** an LSTM cell with a constant forget gate $f = \sigma(b_f)$, a constant input gate i and i.i.d. candidates g_t with mean 0 and variance σ_g^2
 - ▶ Data for (c): S&P 500 daily log returns (GSPC.INDX, data/market) since January 2000; ACF of $|r_t|$: 0.289 (lag 1), 0.132 (lag 50), 0.059 (lag 150)
- (a) Show that $c_t = fc_{t-1} + ig_t$ is a stationary AR(1) for $0 < f < 1$; derive $\text{Var}(c_t)$ and $\text{Corr}(c_t, c_{t-k})$
- (b) Compute the half-life $\ln 0.5 / \ln f$ for $b_f = 1$ and for $f = 0.99$, the bias b_f that gives $f = 0.99$, and f^{250} for $b_f = 5$
- (c) Fit a geometric decay with its own amplitude, $\rho_k = \rho_1 f^{k-1}$ (an AR(1) signal observed with noise): choose f so that $f^{49} = \rho_{50} / \rho_1$; compute b_f , the half-life and the predicted $\rho_{150} = \rho_1 f^{149}$
- **Report** the numbers of (b) and (c); why can no fixed $f < 1$ reproduce long memory, and what does the network need instead?
- Model: A1

A3. Mean Scaling in Chronos: Who Moves the Scale? [Solved]

Task

- **Context:** the original Chronos (Ansari et al. (2024)) divides a context of C returns by $s = \frac{1}{C} \sum_t |x_t|$ and maps $z = x/s$ to 4,093 equally spaced bins in $[-15, 15]$, $\Delta = 30/4092$
 - ▶ Given for (b): a four-day context of returns in %, $x = (0.8, -1.2, 0.4, 2.0)$
- (a) The window moves one day: a new value M enters and x_{out} leaves. Derive s_{new} as a function of s_{old} , M , x_{out} and C , and the factor by which the scaled value of every unchanged day changes
- (b) For the four-day context compute s , the scaled values z and the bin indices $j = \text{round}((z + 15)/\Delta)$
- **Report** the formula of (a) and the table of z and j ; how does the effect of one crash on the scale depend on C ?

A3. Mean Scaling in Chronos: Who Moves the Scale?: solution

Solution

- (a) The sum of absolute values loses $|x_{\text{out}}|$ and gains $|M|$: $\sum_{\text{new}} = \sum_{\text{old}} - |x_{\text{out}}| + |M|$
 - ▶ Dividing by C : $s_{\text{new}} = s_{\text{old}} + (|M| - |x_{\text{out}}|)/C$
 - ▶ Unchanged day x : $z_{\text{new}} = x/s_{\text{new}} = z_{\text{old}} \cdot s_{\text{old}}/s_{\text{new}}$
- (b) $s = (0.8 + 1.2 + 0.4 + 2.0)/4 = 1.10$
 - ▶ $z = x/s = (+0.727; -1.091; +0.364; +1.818)$; $\Delta = 0.00733$
 - ▶ $j = (2145, 1897, 2096, 2294)$: all near the centre of the vocabulary, $j = 2046$ for $z = 0$
- Interpretation: a crash moves s by about $|M|/C$: a short context reacts strongly, a 512-day context hardly at all

A4. A Crash Day in the Context [Proposed]

Solution discussed in class

Task

- ▣ **Context:** the scaling of A3; S&P 500 daily log returns (GSPC.INDX, data/market) in a 512-day context
 - ▶ On 16 March 2020 the return -12.77% enters and a day of 0.51% leaves; before the crash $s_{\text{old}} = 0.741\%$
- ▣ (a) Replace the last value of the A3 context by -9.0 and recompute s , z , j and the factor $s_{\text{old}}/s_{\text{new}}$
- ▣ (b) Apply the formula of A3(a) to the S&P 500 case and compute s_{new} and the factor
- ▣ (c) Compute the scaled value of an ordinary 0.4% day before and after the crash, in both cases
- ▣ **Report** s , the factor and z of the 0.4% day; does a 512-day context make ordinary days look calm after a crash?
- ▣ Model: A3

A5. VaR 1% and ES 2.5% from a Quantile Grid [Solved]

Task

- ▣ **Context:** the last zero-shot Chronos-2 forecast for the S&P 500 (day after 18.09.2026; context of 512 daily returns, GSPC.INDX); grid levels 1%, 5%, 10%, ..., 99%
 - ▶ Given: $q_{0.01} = -2.704\%$, $q_{0.05} = -1.400\%$; convention $\text{VaR}_\alpha = -q_\alpha$, $\text{ES}_\alpha = -\frac{1}{\alpha} \int_0^\alpha q_u du$
- ▣ (a) Compute VaR 1%
- ▣ (b) Interpolate $q_{0.025}$ linearly between $q_{0.01}$ and $q_{0.05}$ and compute VaR 2.5%
- ▣ (c) Compute ES 2.5%, holding $q_u = q_{0.01}$ for $u < 0.01$ and using the linear q_u on $[0.01, 0.025]$
- ▣ **Report** VaR 1%, VaR 2.5%, ES 2.5% and the ratio ES 2.5%/VaR 2.5%; compare this ratio at the same tail level with its value for a zero-mean Normal distribution

A5. VaR 1% and ES 2.5% from a Quantile Grid: solution

Solution

- (a) $\text{VaR } 1\% = -q_{0.01} = 2.704\%$
- (b) $q_{0.025} = q_{0.01} + \frac{0.025-0.01}{0.05-0.01}(q_{0.05} - q_{0.01}) = -2.704 + 0.375(-1.400 + 2.704) = -2.215\%$
 - ▶ $\text{VaR } 2.5\% = -q_{0.025} = 2.215\%$
- (c) Split the integral: $\int_0^{0.01} q_u du = 0.01 q_{0.01}$; $\int_{0.01}^{0.025} q_u du = 0.015 (q_{0.01} + q_{0.025})/2$ (trapezoid)
 - ▶ $\text{ES } 2.5\% = -\frac{1}{0.025} \left[0.01 q_{0.01} + 0.015 \frac{q_{0.01} + q_{0.025}}{2} \right] = 2.557\%$
- Interpretation: $\text{ES } 2.5\% / \text{VaR } 2.5\% = 1.15$, below the zero-mean Normal value at the same level, $\varphi(1.96)/(0.025 \cdot 1.96) = 1.19$: the grid cuts the tail below 1% (A6)

A6. Clamping the Grid Tail; a GPD Tail [Proposed]

Solution discussed in class

Task

- **Context:** the grid of A5; S&P 500 daily losses $L = -r$ (GSPC.INDX) over the last 1,000 trading days to 18.09.2026
 - ▶ Peaks over threshold: u = empirical 95% quantile of L ; excesses $L - u$ follow a $\text{GPD}(\xi, \beta)$ with $\xi < 1$ (McNeil & Frey (2000))
- (a) With the quantile function exact on $[0.01, 0.025]$, prove that holding $q_u = q_{0.01}$ for $u < 0.01$ never raises ES 2.5%; then show that the linear interpolation of A5 on $[0.01, 0.025]$ can err in either direction
- (b) Derive $\text{ES}_\alpha = \frac{\text{VaR}_\alpha}{1-\xi} + \frac{\beta-\xi u}{1-\xi}$ from the mean excess function of the GPD
- (c) Fit ξ, β by maximum likelihood; compute $\text{VaR}_\alpha = u + \frac{\beta}{\xi} \left[(n\alpha/N_u)^{-\xi} - 1 \right]$ for $\alpha = 1\%$ and 2.5% , and ES 2.5%
- **Report** u , N_u , $\hat{\xi}$, $\hat{\beta}$, VaR and ES; compare ES/VaR with A5
- **Model:** A5

A7. A “VaR 1%” from Deciles [Proposed]

Solution discussed in class

Task

- **Context:** Chronos-Bolt (small), zero-shot, same day and context as A5; it returns only the deciles $q_{0.1}, \dots, q_{0.9}$
 - ▶ Given: $q_{0.1} = -0.691\%$, $q_{0.5} = 0.071\%$
- (a) State the value returned when VaR 1% is requested from a decile grid that is clamped at its lowest level
- (b) Assume a Normal predictive distribution with median $q_{0.5}$ and $\hat{\sigma} = (q_{0.5} - q_{0.1})/z_{0.9}$; compute VaR 1% $= -(q_{0.5} + \hat{\sigma}z_{0.01})$ and ES 2.5% $= -(q_{0.5} - \hat{\sigma}\varphi(z_{0.025})/0.025)$
- (c) Compare with Chronos-2 (A5) and assess the Normal-extrapolation assumption: which causes of the too-frequent breaches of raw foundation-model quantiles (Lecture 14, Test 3) can it address, and which test would settle it?
- **Report** the three numbers and the comparison
- **Model:** A5

A8. The Exact Power of the Kupiec Test [Solved]

Task

- **Context:** VaR 1% forecasts over the 220 trading days after the release of the weights (3.11.2025–18.09.2026); the number of breaches $x \sim \text{Bin}(T, \pi)$ with $T = 220$
 - ▶ Kupiec (Kupiec (1995)) LR test, two-sided, level 5%: reject if $\text{LR}_{uc} > \chi^2_{1;0.95} = 3.84$
- (a) Compute $\text{LR}_{uc}(x) = -2[(T - x) \ln 0.99 + x \ln 0.01 - (T - x) \ln(1 - x/T) - x \ln(x/T)]$ for $x = 0, 1, 2, \dots$ and find the rejection region
- (b) Compute the exact size $P_{\pi=1\%}(\text{reject})$
- (c) Compute the power $P_{\pi=1.67\%}(\text{reject})$; 1.67% is the breach rate of raw Chronos-2 on the S&P 500
- **Report** the rejection region, size and power; what does a non-rejection over 220 days tell a risk manager?

A8. The Exact Power of the Kupiec Test: solution (1/2)

Solution

- (a) $x = 0$: $LR = -2 \cdot 220 \ln 0.99 = 4.42 > 3.84$: reject (too few breaches)
 - ▶ $x = 5$: $LR = 2.65 < 3.84$; $x = 6$: $LR = 4.51 > 3.84$
 - ▶ Rejection region: $x = 0$ or $x \geq 6$; acceptance region: $1 \leq x \leq 5$
- (b) $\text{Size} = P_{1\%}(x = 0) + P_{1\%}(x \geq 6) = 0.134$, not 0.05: with a discrete statistic the nominal level cannot be hit

A8. The Exact Power of the Kupiec Test: solution (2/2)

Solution

- (c) Power = $P_{1.67\%}(x = 0) + P_{1.67\%}(x \geq 6) = 0.9833^{220} + 0.165 = 0.025 + 0.165 = 0.189$
 - Expected number of breaches: $220 \times 1\% = 2.2$ under H_0 , $220 \times 1.67\% = 3.7$ under the alternative
- Interpretation: four times in five a model with 67% too many breaches passes; over 220 days, “not rejected” is no evidence of correct coverage

A9. How Long Must the Post-Release Window Be? [Proposed]

Solution discussed in class

Task

- ▣ **Context:** the Kupiec test of A8 (two-sided, 5%) for VaR 1%; true breach rate $\pi = 1.67\%$ or 2%
- ▣ (a) Compute the exact power against $\pi = 1.67\%$ for $T = 1,000$ and $T = 2,693$ (the full sample of Lecture 14)
- ▣ (b) Find the smallest T with power $\geq 80\%$ and the T from which it stays above 80%; explain why power is not monotone in T
- ▣ (c) Repeat for $\pi = 2\%$: the power at $T = 220$ and the T for 80% power
- ▣ **Report** the numbers; how many years of post-release data does a clean backtest need?
- ▣ Model: A8

A10. Which Functional Does QLIKE Elicit? [Solved]

Task

- ▣ **Context:** SPY realised variance RV (sum of squared 5-minute returns), one-day-ahead forecasts from 11.10.2022 to 18.09.2026 (979 days); loss QLIKE $L(y, h) = y/h - \log(y/h) - 1$
 - ▶ Chronos-2 on log RV (context 512) returns quantiles q_u ; average ratio of $\exp(q_{0.5})$ to $\int_0^1 \exp(q_u) du$ over the days: 0.8048
- ▣ (a) Minimise $E[L(Y, h) \mid \mathcal{F}]$ over $h > 0$ and show that the minimiser is $E[Y \mid \mathcal{F}]$
- ▣ (b) Assume $\log Y \mid \mathcal{F} \sim N(\mu, s^2)$: derive the median, the mean, their ratio and the expected QLIKE excess of the median forecast
- ▣ (c) From the Chronos-2 ratio infer an effective common s and the QLIKE excess (a day-averaged ratio does not identify each day's s)
- ▣ **Report** the derivations and the two numbers; why is $\exp(\text{median})$ a biased volatility forecast?

A10. Which Functional Does QLIKE Elicit?: solution (1/2)

Solution

- (a) $E[L] = E[Y]/h - E[\log Y] + \log h - 1$; $\partial_h E[L] = -E[Y]/h^2 + 1/h = 0 \Rightarrow h^* = E[Y \mid \mathcal{F}]$
 - Second derivative $2E[Y]/h^3 - 1/h^2 = 1/h^{*2} > 0$ at h^* : a minimum; QLIKE is consistent for the mean (Gneiting (2011))
- (b) Median e^μ , mean $m = e^{\mu+s^2/2}$, ratio $e^{-s^2/2}$
 - With $h = me^{-s^2/2}$: $E[L(h)] - E[L(m)] = e^{s^2/2} - 1 - s^2/2 > 0$

A10. Which Functional Does QLIKE Elicit?: solution (2/2)

Solution

- (c) $s = \sqrt{-2 \ln(\text{ratio})} = \sqrt{-2 \ln 0.8048} = 0.659$; excess $e^{s^2/2} - 1 - s^2/2 = 1/\text{ratio} - 1 + \ln(\text{ratio}) = 0.0254$
 - Observed in Lecture 14: QLIKE 0.259 with the median against 0.231 with the mean, a gap of the same order
- Interpretation: monotone maps preserve quantiles, so $\exp(\text{median of } \log \text{RV})$ is the median of RV, about 20% below its mean; QLIKE rewards the mean

B1. Can Anything Beat Zero for the S&P 500? [Solved]

Task

- ▣ **Context:** S&P 500 daily log returns in % (GSPC.INDX closes, data/market); test period 04.01.2016–18.09.2026 (2 693 days); benchmark: the zero forecast
 - ▶ One-day-ahead forecasts of Lecture 14, using data up to the previous close: historical mean, AR(1) (rolling 1,000 days), LSTM (average of 5 seeds), Chronos-2, Chronos-Bolt and TimesFM-2.5 (median)
- ▣ (a) Compute the out-of-sample $R^2 = 1 - \sum_t (r_t - \hat{r}_t)^2 / \sum_t r_t^2$ of each model with a 95% circular block-bootstrap CI (blocks of 20 days, 1,000 draws)
- ▣ (b) Test each model against zero with the DM test on squared errors (Newey–West, two-sided, 5%) and compute the share of correct signs
- ▣ (c) Apply Clark–West (Clark & West (2007), one-sided) to the nested historical mean and AR(1); apply SPA (Hansen (2005)) and Romano–Wolf (Romano & Wolf (2005)) to the six models (stationary bootstrap, mean block 20 days, 2,000 draws)
- ▣ **Report** one table (R^2 , CI, DM, p , correct signs) and the Clark–West, SPA and Romano–Wolf p -values

B1. Can Anything Beat Zero for the S&P 500?: interpretation questions

Interpretation

- Is a positive R^2 of about 1% statistically meaningful here? Why is 55% correct signs not evidence of skill?
- Does anything beat zero once nesting and the six-model search are accounted for?

B1. Solution (1/3): Procedure

Solution

- Step 1: align forecasts and returns on the 2693 test days; losses $e_{0,t}^2 = r_t^2$ (zero) and $e_{1,t}^2 = (r_t - \hat{r}_t)^2$ (model)
- Step 2: $R^2 = 1 - \sum e_1^2 / \sum e_0^2$; bootstrap: draw blocks of 20 consecutive days with wrap-around, recompute R^2 1,000 times, take the 2.5% and 97.5% percentiles
- Step 3: $d_t = e_{1,t}^2 - e_{0,t}^2$, $DM = \bar{d} / \sqrt{\widehat{LRV} / T}$ with $\lfloor 4(T/100)^{2/9} \rfloor$ Bartlett lags; $p = 2[1 - \Phi(|DM|)]$; reject equal accuracy if $p < 0.05$
- Step 4: Clark–West $\tilde{d}_t = e_{0,t}^2 - (e_{1,t}^2 - \hat{r}_t^2)$, HAC t -statistic, $p = 1 - \Phi(t)$
- Step 5: SPA and Romano–Wolf: studentised gains t_k of $e_0^2 - e_{1,k}^2$; bootstrap the maximum of the recentred t_k ; Romano–Wolf removes rejected models step by step

B1. Solution (2/3): Results

Model	R^2 (%)	95% CI	DM [p]	Correct sign (%)
Historical mean	+0.15	[-0.02; +0.40]	-1.67 [0.095]	54.6
AR(1)	+0.94	[-2.65; +4.40]	-0.47 [0.639]	54.1
LSTM	+0.93	[+0.17; +1.59]	-1.70 [0.090]	53.5
Chronos-2	-0.26	[-0.89; +0.38]	+0.81 [0.418]	54.7
Chronos-Bolt	-1.57	[-2.73; -0.24]	+1.76 [0.078]	54.6
TimesFM-2.5	-3.35	[-4.62; -1.95]	+3.23 [0.001]	53.6

- LSTM: the CI [+0.17; +1.59] excludes zero, but DM = -1.70, $p = 0.090 > 0.05$: equal accuracy with zero is not rejected at 5%
- TimesFM-2.5 is significantly worse than zero; AR(1) has a huge CI, driven by 2020
- The S&P 500 rose on 54.6% of days: "always up" already gets that share of signs right

B1. Solution (3/3): Nesting and Many Models

- Zero is nested in the historical mean and in AR(1): under H_0 their estimation noise makes DM undersized
 - ▶ Clark–West: $\tilde{d}_t = r_t^2 - [(r_t - \hat{r}_t)^2 - \hat{r}_t^2]$, one-sided HAC t
 - ▶ With the zero benchmark $\tilde{d}_t = 2r_t\hat{r}_t$: CW tests the martingale-difference null $E[r_t\hat{r}_t] = 0$, valid for any forecast built from past data (Clark & West (2006)); DM, SPA and Romano–Wolf test equal MSE, a different hypothesis
 - ▶ One-sided p , DM vs CW: historical mean 0.048 vs 0.007; AR(1) 0.320 vs 0.055
- Six models tested against the same zero forecast: SPA $p = 0.130$ (H_0 : none beats zero)
 - ▶ Romano–Wolf adjusted p : historical mean 0.164, LSTM 0.161
- Reading: CW detects the small drift of the historical mean; once the six-model search is accounted for, nothing beats zero on the S&P 500

B2. Bitcoin and BET [Proposed]

Solution discussed in class

Task

- **Context:** as B1, for Bitcoin (BTC-USD.CC closes, calendar days, 2818 days from 01.01.2019) and the BET (official index, data/market/BET.csv, 2675 days from 04.01.2016)
 - ▶ Stale-price check: the BET-tracking ETF Patria-TVBETETF (adjusted close), joined with the BET on common days before computing returns
- (a) Repeat B1(a)–(b) for both markets
- (b) For the BET LSTM compute the one-sided DM and Clark–West p , Romano–Wolf within the six BET models, SPA for the BET, and the Holm-adjusted p over all 18 model–market tests
- (c) Compute the BET lag-1 autocorrelation, the correlation of the LSTM forecast with r_{t-1} (Lo & MacKinlay (1990)), and R^2 and Clark–West p of the LSTM signal applied to the ETF return of the same day and of the next day
- **Report** the tables of (a) and the numbers of (b)–(c). Model: B1

B2. Bitcoin and BET: interpretation questions

Interpretation

- Which series has predictable daily returns, and which model detects this predictability?
- Does the signal survive multiplicity and stale prices? Relate the answer to Chapter 2

B3. Realised Volatility: Better than log-HAR? [Solved]

Task

- ▣ **Context:** SPY realised variance $RV_t = \sum_{j=1}^{78} r_{t,j}^2$ from 5-minute bars (data/market/intraday/SPY.US_5m.csv), days with at least 70 bars; 979 one-day-ahead forecasts, 11.10.2022–18.09.2026
 - Models: random walk, HAR and log-HAR (Corsi (2009); rolling 500 days), GARCH-t, LSTM, Chronos-2, Chronos-Bolt, TimesFM-2.5 (grid mean $\int_0^1 \exp(q_u) du$)
- ▣ (a) Compute the mean QLIKE loss (Patton (2011)) of each model
- ▣ (b) Test each model against log-HAR with DM on QLIKE differences (Newey–West, two-sided, 5%)
- ▣ (c) Compute the 90% model confidence set (Hansen, Lunde & Nason (2011); $T_{\max}$ statistic, circular block bootstrap with blocks of 10 days, 1,000 draws); repeat (a)–(b) on the 218 days after 3 November 2025
- ▣ **Report** one table: QLIKE, DM [p], MCS p , QLIKE and DM after the release

Interpretation

- ▣ What does membership of the MCS tell you, and what does it not tell you?
- ▣ Why does the post-release window matter for foundation models only?

B3. Solution (1/2): Procedure

Solution

- Step 1: daily loss $L_{t,m} = \text{RV}_t / \hat{h}_{t,m} - \log(\text{RV}_t / \hat{h}_{t,m}) - 1$ for each model m ; average over the 979 days
- Step 2: $d_t = L_{t,m} - L_{t,\log \text{HAR}}$; DM with Newey–West variance; negative DM: model m better; reject equality if $p < 0.05$
- Step 3: MCS: for the surviving set compute $t_m = (\bar{L}_m - \bar{L}.) / \widehat{\text{se}}$, bootstrap $\max_m t_m$; if $p < 0.10$ eliminate the worst model and repeat; the MCS p of a model is the running maximum of these p -values
- Step 4: repeat steps 1–3 on the days from 3.11.2025, when all three sets of weights were public

B3. Solution (2/2): Results

Model	QLIKE	DM vs log-HAR [p]	MCS p	QLIKE (post)	DM (post)
Random walk	0.326	+4.53 [<0.001]	0.00	0.384	+2.81
HAR	0.247	+1.30 [0.195]	0.29	0.270	+1.40
log-HAR	0.234	–	0.95	0.238	–
GARCH-t	0.283	+3.43 [<0.001]	0.03	0.295	+2.35
LSTM	0.284	+1.68 [0.094]	0.29	0.233	-0.52
Chronos-2	0.231	-0.62 [0.537]	1.00	0.242	+0.57
Chronos-Bolt	0.232	-0.18 [0.857]	0.95	0.236	-0.39
TimesFM-2.5	0.233	-0.08 [0.938]	0.95	0.244	+0.60

- Chronos-2 has the lowest QLIKE but is not significantly better than log-HAR (DM -0.62, $p = 0.537$)
- MCS: random walk and GARCH-t are eliminated; the MCS says "not distinguishable", not "equally good"
- After 3.11.2025 the LSTM is best (0.233); no DM test is significant against log-HAR except for the random walk and GARCH-t

B4. Median or Mean? And How Much Context? [Proposed]

Solution discussed in class

Task

- **Context:** the Chronos-2 and TimesFM-2.5 quantile forecasts of log RV of B3 (same 979 days); Chronos-2 also with contexts of 32, 64, 128, 256 and 512 days on the same forecast days
- (a) For both models compute the QLIKE and the ratio mean forecast / mean RV of $\exp(q_{0.5})$ and of $\int_0^1 \exp(q_u) du$ (trapezoid on the grid)
- (b) Test median against mean for Chronos-2 with DM on the QLIKE differences (Newey–West, 5%)
- (c) Compute the Chronos-2 QLIKE for each context length and compare with log-HAR
- **Report** the table of (a), the DM statistic and the QLIKE by context. Model: B3

Interpretation

- Why is the median biased as a variance forecast (A10)?
- What does the context effect say about the memory of volatility?

B5. Backtesting the Tail Forecasts on the S&P 500 [Solved]

Task

- ▣ **Context:** S&P 500 daily log returns (GSPC.INDX), 2 693 days from 04.01.2016; losses $L_t = -r_t$; one-day VaR 1%, VaR 2.5% and ES 2.5%
 - ▶ Six models: HS, GARCH(1,1)- t and FHS (1,000-day window, re-estimated every 20 days), raw Chronos-2 quantiles, FHS with Chronos-2 or TimesFM volatility
- ▣ (a) Count the VaR 1% breaches $I_t = 1\{L_t > \text{VaR}_t\}$; apply the Kupiec test (Kupiec (1995); two-sided, 5%) and the Christoffersen independence test (first-order Markov chain, χ^2_1)
- ▣ (b) Compute the mean FZ0 loss (Patton, Ziegel & Chen (2019)) of (VaR 2.5%, ES 2.5%) and DM tests against FHS (Newey–West, 5%)
- ▣ (c) For FHS and the Chronos-2 hybrid, compute the DQ test (Engle & Manganelli (2004); constant, four lagged hits and VaR_t , χ^2_6)
- ▣ **Report** one table: breaches, rate, Kupiec p , independence p , FZ0, DM against FHS

B5. Backtesting the Tail Forecasts on the S&P 500: interpretation questions

Interpretation

- Why do the raw Chronos-2 quantiles fail while the hybrid passes the coverage test?
- After passing the coverage test, does the hybrid also pass the independence test and match FHS?

B5. Solution (1/2): Procedure

Solution

- ▣ Step 1: breaches $x = \sum_t I_t$, rate $\hat{\pi} = x/T$; LR_{uc} as in A8; reject correct coverage if $p < 0.05$
- ▣ Step 2: independence: count transitions n_{ij} of $I_{t-1} \rightarrow I_t$, $\hat{\pi}_{01} = n_{01}/(n_{00} + n_{01})$, $\hat{\pi}_{11} = n_{11}/(n_{10} + n_{11})$; LR of $\pi_{01} = \pi_{11}$, χ^2_1 (Christoffersen (1998))
- ▣ Step 3: $FZ0 = \frac{\mathbf{1}\{L > v\}(L-v)}{\alpha e} + \frac{v}{e} + \log e - 1$ with $v = \text{VaR } 2.5\%$, $e = \text{ES } 2.5\%$; DM on the differences against FHS
- ▣ Step 4: DQ: regress $I_t - 0.01$ on $(1, I_{t-1}, \dots, I_{t-4}, \text{VaR}_t)$; $DQ = \hat{\beta}^\top X^\top X \hat{\beta} / (0.01 \cdot 0.99)$; reject if $p < 0.05$

B5. Solution (2/2): Results (S&P 500, 2 693 days)

Model	Breaches	Rate (%)	p Kupiec	p indep.	FZ0	DM vs FHS
HS	37	1.37	0.065	0.001	1.309	+2.46
GARCH-t	44	1.63	0.002	0.205	1.026	+1.34
FHS	30	1.11	0.559	0.045	0.999	–
Chronos-2 quantiles	45	1.67	0.001	0.223	1.147	+3.16
FHS + Chronos-2 vol.	25	0.93	0.705	0.021	1.084	+2.65
FHS + TimesFM vol.	27	1.00	0.989	0.029	1.059	+1.89

- Raw Chronos-2: 1.67% breaches, Kupiec $p = 0.001$: its quantiles are too narrow
- Hybrid (Chronos-2): 0.93%, $p = 0.705$: the tail comes from the data, through FHS
- But breaches cluster (independence $p = 0.021$, DQ $p = 0.001$) and FZ0 is significantly worse than FHS (DM +2.65): right on average, too much capital

B6. Bitcoin, and VaR 10% for All Three Models [Proposed]

Solution discussed in class

Task

- **Context:** the six tail models of B5 for Bitcoin (BTC-USD.CC, 2818 calendar days from 01.01.2019); and the 10% quantiles of Chronos-2, Chronos-Bolt and TimesFM-2.5 against HS and FHS in the three markets
- (a) Repeat B5(a)–(b) for Bitcoin
- (b) For VaR 10% compute the breach rates and the Kupiec test with $\alpha = 10\%$ (two-sided, 5%) for the five models in the three markets
- **Report** the table of (a) and a 3×5 table of breach rates with Kupiec p . Model: B5

Interpretation

- Is the narrowness of foundation-model quantiles specific to the 1% level?

B7. The Seed Lottery [Solved]

Task

- **Context:** the return LSTM of Lecture 14: input the last 20 standardised returns, 16 hidden units, MSE loss, Adam, weight decay 10^{-4} , early stopping on the last 20% of up to 4,000 training days; trained once before each test period of B1–B2
 - ▶ Five networks identical except for the random seed (0–4)
- (a) Compute the out-of-sample R^2 against zero of each seed in the three markets
- (b) Compute the R^2 of the average of the five forecasts
- **Report** a 3×6 table of R^2 (%)

Interpretation

- What would you conclude from a single run on Bitcoin?
- Why is averaging seeds a form of regularisation?

B7. Solution (1/2): Procedure

Solution

- Step 1: for each seed s , $R_s^2 = 1 - \sum_t (r_t - \hat{r}_{t,s})^2 / \sum_t r_t^2$
- Step 2: average forecast $\bar{r}_t = \frac{1}{5} \sum_s \hat{r}_{t,s}$ and its R^2
- Step 3: compare the spread across seeds with the size of R^2 ; by convexity of squared error, the ensemble forecast has no greater squared error than the average squared error of its members

B7. Solution (2/2): Results

R^2 (%)	Seed 0	1	2	3	4	Average
S&P 500	+0.55	+0.37	+0.70	+0.54	+1.96	+0.93
Bitcoin	+0.11	+0.03	-0.79	-0.12	+0.24	+0.04
BET	+1.06	+1.01	+1.12	+1.22	+1.12	+1.15

- Bitcoin: seeds range from -0.79% to +0.24%: a single run could "show" skill or its absence
- BET: all five seeds agree; averaging reduces variance without adding bias, like bagging

B8. Leakage Through a Look-Ahead Regressor [Proposed]

Solution discussed in class

Task

- **Context:** SPY log RV as in B3; log-HAR on a rolling 500-day window with regressors $\log RV_{t-1}$ and the 5- and 22-day averages; forecast $\exp(\hat{y}_t + \hat{s}^2/2)$
- (a) Build a wrong version in which the 5- and 22-day averages include the forecast day t
- (b) Compute the QLIKE of the correct and the wrong log-HAR and of Chronos-2 on the same 979 days; test wrong against correct with DM (Newey–West, 5%)
- **Report** the three QLIKE values, the relative gain and the DM test. Model: B3

Interpretation

- How large is the gain from the look-ahead? Could any honest model get it?
- Where can the same error hide in a deep-learning pipeline?

B9. Calibration of the Whole Predictive Distribution [Solved]

Task

- ▣ **Context:** S&P 500 (GSPC.INDX), 2 693 days from 04.01.2016; the raw Chronos-2 grid of 21 quantiles (1%, 5%, 10%, ..., 99%); FHS and GARCH-t (1,000-day window, refit every 20 days) evaluated on the same grid
- ▣ (a) Compute the PIT $u_t = \hat{F}_t(r_t)$ (Diebold, Gunther & Tay (1998)) by linear interpolation between the grid quantiles, censored outside $[q_{1\%}, q_{99\%}]$; apply Pearson's χ^2 to the counts in the 22 grid intervals (21 degrees of freedom)
- ▣ (b) Apply the Berkowitz LR test (Berkowitz (2001)) of $\mu = 0$, $\sigma = 1$ to $z_t = \Phi^{-1}(u_t)$, with a Normal likelihood censored below 1% and above 99% (χ^2_2 , 5%)
- ▣ (c) Compute the quantile score $QS_t = \frac{2}{K} \sum_k (1\{r_t < q_{u_k}\} - u_k)(q_{u_k} - r_t)$ (an equally weighted quantile score; the CRPS would weight each level by its spacing, Gneiting & Raftery (2007)) and DM of Chronos-2 minus FHS, overall, in the centre (10%–90%) and in the tails
- ▣ **Report** the PIT histogram and a table: tail shares, $\hat{\mu}/\hat{\sigma}$, Berkowitz p , χ^2 p , QS

Interpretation


- ▣ Is the failure of the foundation model in the centre or in the tails?
- ▣ Why censor instead of extrapolating beyond the grid?

B9. Solution (1/2): Procedure

Solution

- Step 1: for each model, on its own 21-quantile grid: if $q_{1\%,t} \leq r_t \leq q_{99\%,t}$, u_t = linear interpolation of the levels between the two neighbouring quantiles; otherwise record only “below” or “above”
- Step 2: $\chi^2 = \sum_{j=1}^{22} (O_j - Tp_j)^2 / (Tp_j)$, p_j = width of interval j
- Step 3: $\ell(\mu, \sigma) = \sum_{\text{inside}} \log \frac{1}{\sigma} \varphi\left(\frac{z_t - \mu}{\sigma}\right) + n_L \log \Phi\left(\frac{c_L - \mu}{\sigma}\right) + n_H \log \left[1 - \Phi\left(\frac{c_H - \mu}{\sigma}\right)\right]$, $c_L = \Phi^{-1}(0.01)$, $c_H = \Phi^{-1}(0.99)$
- LR = $2[\ell(\hat{\mu}, \hat{\sigma}) - \ell(0, 1)] \sim \chi^2_2$; reject calibration if $p < 0.05$; $\hat{\sigma} > 1$ means outcomes more dispersed than forecast
- Step 4: QS on all 21 levels, on the 17 central levels and on the 4 tail levels; DM with Newey–West variance

B9. Solution (2/2): Results



Model	Below $q_1\%$ (%)	Above $q_{99}\%$ (%)	$\hat{\mu} / \hat{\sigma}$	Berkowitz p	$\chi^2 p$	QS
Chronos-2	1.67	1.08	-0.035 / 1.030	0.019	0.004	0.5101
FHS	1.11	1.34	+0.014 / 0.990	0.599	0.747	0.5062
GARCH-t	1.63	0.30	-0.041 / 0.976	0.023	<0.001	0.5063

- Chronos-2: too many returns below its 1% quantile, $\hat{\sigma} > 1$: Berkowitz and χ^2 reject; FHS passes both
- QS, Chronos-2 minus FHS: DM +3.06 (p 0.002); centre +2.51, tails +3.65: worse in both, most in the tails
- GARCH-t: Berkowitz (p 0.023) and χ^2 reject: 1.63% below $q_1\%$ but only 0.30% above $q_{99}\%$, a symmetric t misses the asymmetric tails; censoring avoids inventing a tail the grid does not have

B10. Conformal Recalibration of Chronos-2 [Proposed]

Solution discussed in class

Task

- ▣ **Context:** raw Chronos-2 quantiles of B5, in the three markets (S&P 500 and BET from 2016, Bitcoin from 2019); evaluation on the common days after the first 250 calibration days
- ▣ (a) Correct the 1% and 2.5% quantiles with adaptive conformal inference ([Gibbs & Candès \(2021\)](#)): scores $E_t = \hat{q}_{\alpha,t} - r_t$ on the last 250 days, corrected quantile $\hat{q}_{\alpha,t} - \hat{Q}_{1-\alpha_t}(E)$, update $\alpha_{t+1} = \alpha_t + \gamma(\alpha - \text{err}_t)$ with $\gamma = 0.001$; shift ES 2.5% by the VaR 2.5% correction
- ▣ (b) Repeat with $\gamma = 0$ (plain split conformal)
- ▣ (c) Backtest FHS, raw Chronos-2, the Chronos-2 hybrid and both conformal versions: breaches, conditional coverage, DQ ([Engle & Manganelli \(2004\)](#)), FZ0 and DM against FHS (5%)
- ▣ **Report** one table per market. Models: B5, B9

Interpretation

- ▣ Does conformal recalibration make a foundation model a usable VaR model without GARCH?
- ▣ Why can a model with correct coverage still lose the FZ0 comparison?

B11. A Formal Pre/Post-Release Test [Proposed]

Solution discussed in class

Task

- ▣ **Context:** SPY realised variance, Chronos-2 and log-HAR forecasts of B3 (979 days); common post-release cutoff: 3 November 2025, the first trading day after all three weights were public (Chronos-2: 30.10.2025); 218 days after
- ▣ (a) Compute $d_t = \text{QLIKE}(\text{Chronos-2}) - \text{QLIKE}(\text{log-HAR})$ and its means before and after the release
- ▣ (b) Regress d_t on a constant and $1\{t \geq 3.11.2025\}$ with Newey–West errors ($\lfloor 4(T/100)^{2/9} \rfloor$ lags) and test $\beta_1 = 0$ at 5%
- ▣ (c) Compute the minimum detectable effect of β_1 , $(z_{0.975} + z_{0.8}) \text{se}_{HAC}(\hat{\beta}_1)$, and its limit when the post-release window grows without bound (n_{pre} fixed); can the pre-release gap ever be detected?
- ▣ **Report** the means, $\hat{\beta}_1$, t , p , MDE. Model: B3

Interpretation

- ▣ Does a non-significant $\hat{\beta}_1$ show that Chronos-2 did not memorise the test period?
- ▣ What design would give the test power?

C1. Can a Zero-Shot Foundation Model Replace GARCH for the BET? [Proposed]

Solution discussed in class

Task

- ▣ **Context:** a risk manager proposes Chronos-2 for the daily VaR 1% and ES 2.5% of a BET position; data: official BET index (data/market/BET.csv), daily log returns, 2 675 days from 04.01.2016 to 18.09.2026; post-release window from 3.11.2025 (217 days)
- ▣ (a) Backtest VaR 1% of the six models of B5: breaches, Kupiec, Christoffersen independence and DQ (5%)
- ▣ (b) Compare (VaR 2.5%, ES 2.5%) with FZ0, DM against FHS and the 90% MCS
- ▣ (c) Classify ten non-overlapping 250-day blocks into Basel zones (green ≤ 4 , yellow 5–9, red ≥ 10 breaches); repeat (a) on the post-release window
- ▣ (d) Pre-register: submit your VaR forecasts for the next 20 trading days before they occur and backtest them afterwards
- ▣ **Report** a one-page recommendation with the tables of (a)–(c)

C1. Can a Zero-Shot Foundation Model Replace GARCH for the BET?: interpretation questions

Interpretation

- Accept, reject, or accept as a challenger model?
- Project extension: fine-tune a foundation model on Romanian stocks and re-run the test

C2: Critique an AI Answer [Proposed]

- A student asks an AI assistant: “Backtest zero-shot Chronos-2 one-day-ahead VaR 1% for the S&P 500 on the last 1500 days, with a 512-day context”
- The assistant answers with this code (spx: daily closes, imports omitted):

```
pipe = BaseChronosPipeline.from_pretrained('amazon/chronos-2')
r = 100 * np.log(spx).diff().dropna().to_numpy()      # daily returns, %
test = np.arange(len(r) - 1500, len(r))
q01 = np.empty(len(test))
for i, t in enumerate(test):                          # forecast of r[t]
    ctx = torch.tensor(r[t - 511:t + 1], dtype=torch.float32) # 512 days
    q, _ = pipe.predict_quantiles([ctx[None, :]], prediction_length=1,
                                  quantile_levels=[0.01])
    q01[i] = q[0][0, 0, 0].item()
x = int((r[test] < q01).sum())                          # breaches
print(f"VaR1%={q01.mean():.2f}%,{breaches}{x}")
```

- “Output: VaR 1% = −2.41%, breaches = 5, i.e. 0.33% against 15 expected: the model is conservative, so it passes the Kupiec test easily”

C2: Critique an AI Answer, Task [Proposed]

Solution discussed in class

Task

- ▣ **Context:** the code and conclusion of the previous slide; S&P 500 daily closes, last 1 500 days, zero-shot Chronos-2 with a 512-day context; the answer contains **three** errors, in the code and in the conclusion
- ▣ (a) For each error, say what is wrong and why
- ▣ (b) Say how you would detect it: an index check, the VaR convention of the lecture, the test statistic
- ▣ (c) Give the corrected line or conclusion, and compute the Kupiec statistic and its p -value for 5 breaches in 1 500 days
- ▣ Log the three errors as in `AI_ERRORS.md`: what was asked, what the AI answered, what was wrong, how you found out, the correct version

Interpretation

- ▣ Which element of $\mathbf{r}$ should be the last one in the context when forecasting $\mathbf{r}[t]$?
- ▣ Is “too few breaches” good news for a risk manager? For a regulator?
- ▣ The weights of Chronos-2 were published in autumn 2025: is the last 1 500 days a clean test for it?

C3. The Virtue of Complexity on Course Data [Proposed]

Solution discussed in class

Task

- Context: the case study of Lecture 14 (Kelly, Malamud & Zhou (2024)) times the US stock market with ridge regressions on up to 12 000 random Fourier features of the 15 predictors of Welch & Goyal (2008); you replicate its Table I on the course data, monthly, from 1993
 - ▶ Model: B1 (out-of-sample return forecasts against the zero forecast)
- (a) Build the monthly data: SPY total returns from month-end adjusted closes (data/market) minus the 3-month Treasury bill rate (FRED TB3MS); predictors tbl , lty (GS10), tms , dfy (BAA minus AAA), $infl$ (CPIAUCSL), $svar$ (sum of squared daily S&P 500 returns in the month) and the lagged excess return
- (b) Standardise as in Section V.A of the paper: returns by their trailing 12-month uncentred standard deviation, predictors by an expanding-window standard deviation, after 36 months of warm-up
- (c) Generate 12 000 random Fourier features with $\gamma = 2$ (eq. 20); for P from 2 to 12 000 and $\log_{10} z = -3, \dots, 3$, estimate ridge without intercept on rolling 12-month windows, with features scaled by their training-window standard deviation, and form the timing return $\hat{\beta}' S_t R_{t+1}$; average all statistics over 1 000 draws of the features
- (d) Estimate the linear kitchen sink on your seven predictors with $z = 0^+$ and $z = 10^3$ on the same windows
- Report the curves of $\|\hat{\beta}\|^2$, out-of-sample R^2 and Sharpe ratio against $c = P/T$ (as in Figures 7–8), and a table in the format of Table I for the two linear models and RFF with $c = 1\,000$, $z = 10^3$: R^2 , annualised Sharpe ratio with t , information ratio against SPY and against the linear model with t , maximum loss, skewness

C3. The Virtue of Complexity on Course Data: interpretation questions

Interpretation

- Your kitchen sink has $c = 7/12$, not $15/12$ as in the paper: is it still near the interpolation boundary, and what does the ridgeless curve of Figure 7 predict for its R^2 ?
- Can the Sharpe ratio be positive while $R^2 < 0$? Explain with the split of the forecast error into a correlation part and a scale part (footnote 5 of the paper)
- With about 30 years of test data, is an information ratio of 0.3 detectable at 5%? Use $t \approx \text{ratio} \times \sqrt{\text{years}}$
- Project extension: replace the random features by the internal representations (embeddings) of a pretrained foundation model and repeat (c)

Take-Home Assignment (Team Project)

- ▣ Choose one asset not covered in class (a BVB stock, a crypto-asset, an exchange rate or a commodity)
- ▣ Forecast its daily volatility or VaR 1% with one foundation model, zero-shot, and with a GARCH benchmark
 - ▶ Rolling origin, and the context window must end the day before the forecast
 - ▶ Report the period after the model's release date separately
- ▣ Compare the models with a DM test and, for VaR, with the coverage and independence tests
- ▣ Log every AI error you catch in `AI_ERRORS.md` (at least three per project)
- ▣ Deliver a Quantlet folder (code, Metainfo, charts) and three slides: numbers, interpretation, limitations

References (1/3)

- ▣ Ansari, A.F., Shchur, O., Küken, J., Auer, A., et al. (2025). Chronos-2: from univariate to universal forecasting. arXiv:2510.15821.
- ▣ Ansari, A.F., Stella, L., Turkmen, C., Zhang, X., et al. (2024). Chronos: learning the language of time series. *Transactions on Machine Learning Research*; arXiv:2403.07815.
- ▣ Berkowitz, J. (2001). Testing density forecasts, with applications to risk management. *Journal of Business & Economic Statistics*, 19(4), 465–474.
- ▣ Christoffersen, P.F. (1998). Evaluating interval forecasts. *International Economic Review*, 39(4), 841–862.
- ▣ Clark, T.E., West, K.D. (2006). Using out-of-sample mean squared prediction errors to test the martingale difference hypothesis. *Journal of Econometrics*, 135(1–2), 155–186.
- ▣ Clark, T.E., West, K.D. (2007). Approximately normal tests for equal predictive accuracy in nested models. *Journal of Econometrics*, 138(1), 291–311.
- ▣ Corsi, F. (2009). A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2), 174–196.
- ▣ Diebold, F.X., Gunther, T.A., Tay, A.S. (1998). Evaluating density forecasts with applications to financial risk management. *International Economic Review*, 39(4), 863–883.

References (2/3)

- ▣ Diebold, F.X., Mariano, R.S. (1995). [Comparing predictive accuracy](#). *Journal of Business & Economic Statistics*, 13(3), 253–263.
- ▣ Engle, R.F., Manganelli, S. (2004). [CAViaR: conditional autoregressive value at risk by regression quantiles](#). *Journal of Business & Economic Statistics*, 22(4), 367–381.
- ▣ Franke, J., Härdle, W.K., Hafner, C.M. (2019). [Statistics of Financial Markets: An Introduction](#), 5th ed. Springer.
- ▣ Gibbs, I., Candès, E. (2021). [Adaptive conformal inference under distribution shift](#). *NeurIPS*; arXiv:2106.00170.
- ▣ Gneiting, T. (2011). [Making and evaluating point forecasts](#). *Journal of the American Statistical Association*, 106(494), 746–762.
- ▣ Gneiting, T., Raftery, A.E. (2007). [Strictly proper scoring rules, prediction, and estimation](#). *Journal of the American Statistical Association*, 102(477), 359–378.
- ▣ Hansen, P.R. (2005). [A test for superior predictive ability](#). *Journal of Business & Economic Statistics*, 23(4), 365–380.
- ▣ Hansen, P.R., Lunde, A., Nason, J.M. (2011). [The model confidence set](#). *Econometrica*, 79(2), 453–497.

References (3/3)

- ▣ Hochreiter, S., Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735–1780.
- ▣ Kelly, B., Malamud, S., Zhou, K. (2024). The virtue of complexity in return prediction. *Journal of Finance*, 79(1), 459–503.
- ▣ Kupiec, P.H. (1995). Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2), 73–84.
- ▣ Lo, A.W., MacKinlay, A.C. (1990). An econometric analysis of nonsynchronous trading. *Journal of Econometrics*, 45(1–2), 181–211.
- ▣ McNeil, A.J., Frey, R. (2000). Estimation of tail-related risk measures for heteroscedastic financial time series: an extreme value approach. *Journal of Empirical Finance*, 7(3–4), 271–300.
- ▣ Pascanu, R., Mikolov, T., Bengio, Y. (2013). On the difficulty of training recurrent neural networks. *ICML*; arXiv:1211.5063.
- ▣ Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- ▣ Patton, A.J., Ziegel, J.F., Chen, R. (2019). Dynamic semiparametric models for expected shortfall (and value-at-risk). *Journal of Econometrics*, 211(2), 388–413.
- ▣ Romano, J.P., Wolf, M. (2005). Stepwise multiple testing as formalized data snooping. *Econometrica*, 73(4), 1237–1282.
- ▣ Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., et al. (2017). Attention is all you need. *NeurIPS*; arXiv:1706.03762.
- ▣ Welch, I., Goyal, A. (2008). A comprehensive look at the empirical performance of equity premium prediction. *Review of Financial Studies*, 21(4), 1455–1508.