



Modelling Financial Markets

Seminar 20: Path Signatures and Realised Volatility — Hands-on



Daniel Traian PELE

Drd. Antoaneta Amza, antoaneta.amza@csie.ase.ro

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material (1/2)

- ▣ AI: Artificial Intelligence
- ▣ AR: AutoRegressive (model)
- ▣ AUDUSD: Australian dollar / US dollar exchange rate (FX pair code)
- ▣ BH: Benjamini–Hochberg (procedure)
- ▣ BIC: Bayesian Information Criterion
- ▣ CI: Confidence Interval
- ▣ CL: Crude oil (WTI light sweet) futures, NYMEX ticker
- ▣ COVID-19: COronaVirus Disease 2019
- ▣ CVX: Chevron Corporation (NYSE ticker)
- ▣ DM: Diebold–Mariano (test)
- ▣ ES: E-mini S&P 500 futures (CME ticker)
- ▣ EURUSD: Euro / US dollar exchange rate (FX pair code)
- ▣ FDR: False Discovery Rate
- ▣ FWER: Family-Wise Error Rate
- ▣ FX: Foreign eXchange
- ▣ GBPUSD: British pound / US dollar exchange rate (FX pair code)
- ▣ GW: Giacomini–White (test)
- ▣ HAC: Heteroskedasticity and Autocorrelation Consistent (standard errors)
- ▣ HAR: Heterogeneous AutoRegressive (model)
- ▣ HARQ: HAR model with Realised Quarticity
- ▣ JPM: JPMorgan Chase & Co. (NYSE ticker)
- ▣ LASSO: Least Absolute Shrinkage and Selection Operator
- ▣ OLS: Ordinary Least Squares
- ▣ OPEC: Organization of the Petroleum Exporting Countries
- ▣ PRDS: Positive Regression Dependence on a Subset (of p-values)
- ▣ QLIKE: Quasi-LIKElihood (loss function)

Acronyms used in this material (2/2)

- ▣ RV: Realised Variance
- ▣ SHAR: Semivariance HAR model
- ▣ SHW: Sherwin-Williams Company (NYSE ticker)
- ▣ USDCAD: US dollar / Canadian dollar exchange rate (FX pair code)
- ▣ USDJPY: US dollar / Japanese yen exchange rate (FX pair code)
- ▣ VOLARE: VOLatility Archive for Realized Estimates (database)

Outline

Introduction

Part A: Derivations

Part B: Signatures and Volatility on VOLARE

Part C: Open Problems

Wrap-up

References

Seminar Objectives

By the end of this seminar you can

- Compute signatures of piecewise-linear paths step by step and use Chen's identity
- Read the Lévy area as a lead-lag statistic and test whether it predicts volatility
- Compare two volatility forecasts with DM, GW and a bootstrap interval
- Count wins over many assets with Holm and BH corrections

- Background: Lecture 20 and Chapter 9; data: VOLARE [Cipollini et al. \(2026\)](#), free for research at volare.unime.it

Structure

- **Part A:** derivations (A1–A10): signatures, Chen, Lévy area, HAR as AR(22), QLIKE, kernel weights
- **Part B:** VOLARE data with inference (B1–B8)
- **Part C:** a level-aware kernel as project seed (C1); critique of an AI answer (C2)

How to Work in This Seminar

- **[Solved]** problems show the full worked solution; **[Proposed]** problems show only the task: their solutions are discussed in class
- Every proposed problem names its **model**: the solved problem that uses the same technique
 - ▶ Every Part B exercise ends with an interpretation question
- Notation: RV_t five-minute realised variance, r_t open-to-close log return,
 $QLIKE(RV, F) = RV/F - \ln(RV/F) - 1$
- Code: [seminar notebook \(Colab\)](#), built on the Quantlet MFM_ch20_seminar; set the folder with your VOLARE files in the first cell

Solved and Proposed Problems

Solved: worked examples in this deck

- Part A: A1 (signature step by step), A3 (Chen and path reversal), A5 (HAR as AR(22)), A7 (QLIKE), A9 (kernel weights)
- Part B: B1 (Lévy area as predictor, JPM), B3 (HAR, SHAR, HARQ, JPM), B5 (Sig-LK vs log-HAR, JPM), B7 (Holm and BH)

Proposed: you solve them

- Part A: A2 (model A1), A4 (model A3), A6 (model A5), A8 (model A7), A10 (model A9)
- Part B: B2 (model B1, ES), B4 (model B3, CL), B6 (model B5, ES), B8 (model B7)
- Part C: C1, a level-aware kernel; C2, critique an AI answer (model: A1, A7, A8, B5)

A1. A Depth-2 Signature, Step by Step [Solved]

Task

- A series takes the values $y = (0, 2, 3, 3)$ on days $t = 0, 1, 2, 3$
 - Assumption: the time-augmented path $X_t = (t, y_t)$ is linear between consecutive days (channel 1 = time, channel 2 = y)
- (a) Write the three increments $\Delta_1, \Delta_2, \Delta_3 \in \mathbb{R}^2$ and compute level 1, $S^{(1)} = \sum_k \Delta_k$
- (b) Use Chen's identity to show that level 2 equals $S^{(2)} = \sum_k \frac{1}{2} \Delta_k \Delta_k^\top + \sum_{j < k} \Delta_j \Delta_k^\top$, then compute the 2×2 matrix
- (c) Check the result with the shuffle identity $S^i S^j = S^{ij} + S^{ji}$ for $i = j = 1$ and for $i = 1, j = 2$
- (d) Compute the Lévy area $A^{12} = (S^{12} - S^{21})/2$
- **Report:** $S^{(1)}$, $S^{(2)}$ and A^{12} . **Interpretation:** what does the sign of A^{12} say about when the series rose?

A1. Solution (1/2)

Solution

- (a) Increments: $\Delta_1 = (1, 2)$, $\Delta_2 = (1, 1)$, $\Delta_3 = (1, 0)$; level 1: $S^{(1)} = (3, 3)$, the total change in time and in y
- (b) One segment has $\text{Sig} = (1, \Delta, \frac{1}{2}\Delta\Delta^\top, \dots)$; Chen multiplies the segments, and level 2 of the product collects every square term once and every ordered pair $j < k$ once
- Squares: $\frac{1}{2} \begin{bmatrix} 1 & 2 \\ 2 & 4 \end{bmatrix} + \frac{1}{2} \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} + \frac{1}{2} \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} = \begin{bmatrix} 1.5 & 1.5 \\ 1.5 & 2.5 \end{bmatrix}$
- Ordered pairs: $\Delta_1\Delta_2^\top + \Delta_1\Delta_3^\top + \Delta_2\Delta_3^\top = \begin{bmatrix} 1 & 1 \\ 2 & 2 \end{bmatrix} + \begin{bmatrix} 1 & 0 \\ 2 & 0 \end{bmatrix} + \begin{bmatrix} 1 & 0 \\ 1 & 0 \end{bmatrix} = \begin{bmatrix} 3 & 1 \\ 5 & 2 \end{bmatrix}$
- Sum: $S^{(2)} = \begin{bmatrix} 4.5 & 2.5 \\ 6.5 & 4.5 \end{bmatrix}$

A1. Solution (2/2)

Solution

- (c) Shuffle, $i = j = 1$: $S^1 S^1 = 9 = 2S^{11} = 2 \cdot 4.5$; $i = 1, j = 2$: $S^1 S^2 = 9 = S^{12} + S^{21} = 2.5 + 6.5$: both hold
- (d) Lévy area: $A^{12} = (2.5 - 6.5)/2 = -2$
- Geometric check: the path lies above its chord $y = t$ (at $t = 1$: 2 vs 1; at $t = 2$: 3 vs 2); the enclosed polygon has area 2, traversed clockwise, hence the minus sign
- Interpretation: with channels (time, y), a negative area means that y rose *early* in the window; the same net change spread evenly would give $A^{12} = 0$

A2. Signature and Shuffle Identity [Proposed]

Solution discussed in class

Task

- Model: A1. A path in $\mathbb{R}^2$ through the points $(0, 0) \rightarrow (2, 1) \rightarrow (3, -1) \rightarrow (4, 2)$, linear between points
- (a) Write the increments and compute level 1
- (b) Compute level 2 with the formula of A1 (squares plus ordered pairs)
- (c) Verify $S^1 S^2 = S^{12} + S^{21}$ and compute the Lévy area A^{12}
- **Report:** $S^{(1)}, S^{(2)}, A^{12}$. **Interpretation:** which channel moved first, on balance?

A3. Chen's Identity and Path Reversal [Solved]

Task

- Path $X : (0, 0) \rightarrow (1, 0) \rightarrow (1, 2)$ from the lecture, with $S^{(1)} = (1, 2)$ and $S^{(2)} = \begin{bmatrix} 1/2 & 2 \\ 0 & 2 \end{bmatrix}$
- $\overleftarrow{X}$ is the same path run backwards: $(1, 2) \rightarrow (1, 0) \rightarrow (0, 0)$
- (a) Compute levels 1 and 2 of $\text{Sig}(\overleftarrow{X})$ with the A1 formula
- (b) Apply Chen's identity, $(A \otimes B)^{(2)} = A^{(2)} + A^{(1)}B^{(1)\top} + B^{(2)}$, to compute level 2 of $X * \overleftarrow{X}$ (there and back)
- (c) State what the result implies for the uniqueness of the signature and for time augmentation
- **Report:** the two signatures and level 2 of the concatenation. **Interpretation:** why do we add a time channel in practice?

A3. Solution

Solution

- (a) Reversed increments: $\overleftarrow{\Delta}_1 = (0, -2)$, $\overleftarrow{\Delta}_2 = (-1, 0)$; level 1: $\overleftarrow{S}^{(1)} = (-1, -2) = -S^{(1)}$
- Level 2: squares $\begin{bmatrix} 0 & 0 \\ 0 & 2 \end{bmatrix} + \begin{bmatrix} 1/2 & 0 \\ 0 & 0 \end{bmatrix}$ plus the pair $\overleftarrow{\Delta}_1 \overleftarrow{\Delta}_2^\top = \begin{bmatrix} 0 & 0 \\ 2 & 0 \end{bmatrix}$, so $\overleftarrow{S}^{(2)} = \begin{bmatrix} 1/2 & 0 \\ 2 & 2 \end{bmatrix}$
- (b) Chen: $\begin{bmatrix} 1/2 & 2 \\ 0 & 2 \end{bmatrix} + \begin{bmatrix} -1 & -2 \\ -2 & -4 \end{bmatrix} + \begin{bmatrix} 1/2 & 0 \\ 2 & 2 \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix}$; level 1 is $S^{(1)} - S^{(1)} = 0$ as well
- (c) There and back is “tree-like”: its signature is that of the constant path; the signature identifies a path only up to such pieces [Hambly & Lyons \(2010\)](#)
- Interpretation: a strictly increasing time channel cannot go back, so tree-like pieces are impossible and the full signature determines the path up to its starting point (translation): this is why we time-augment and supply the level separately

A4. Lévy Area of a Loop [Proposed]

Solution discussed in class

Task

- Model: A3. The unit square $(0, 0) \rightarrow (1, 0) \rightarrow (1, 1) \rightarrow (0, 1) \rightarrow (0, 0)$, four linear segments
- (a) Compute level 1 and the Lévy area A^{12} for the square run anticlockwise
- (b) Repeat for the square run clockwise, and for the square translated by $(5, 5)$
- (c) Explain why level 1 cannot distinguish the loop from a constant path while level 2 can
- **Report:** three values of $S^{(1)}$ and A^{12} . **Interpretation:** which invariance does the translation illustrate?

A5. HAR as a Restricted AR(22) [Solved]

Task

- Data: VOLARE Cipollini et al. (2026), JPM (JPMorgan Chase), daily five-minute realised variance rv5, 2 Jan 2015 – 30 Jan 2026, trading days only
- HAR for $\ln RV_{t+1}$ with averages of logs: $\ell_t = \ln RV_t$, $\ell_t^{(w)} = \frac{1}{5} \sum_{j=0}^4 \ell_{t-j}$, $\ell_t^{(m)} = \frac{1}{22} \sum_{j=0}^{21} \ell_{t-j}$ (the lecture's log-HAR uses logs of averages, which have no exact AR form)
- Full sample (2764 usable days), OLS, Newey–West (5 lags) s.e.: $\hat{\beta}_d = 0.408$ (0.027), $\hat{\beta}_w = 0.299$ (0.046), $\hat{\beta}_m = 0.196$ (0.040)
- (a) Substitute the definitions of $\ell_t^{(w)}$ and $\ell_t^{(m)}$ into the HAR equation and collect the coefficient of each lag ℓ_{t-j+1} , $j = 1, \dots, 22$
- (b) Compute $\phi_1, \phi_2 = \dots = \phi_5$ and $\phi_6 = \dots = \phi_{22}$ from the JPM estimates
- (c) Compute the sum $\sum_{j=1}^{22} \phi_j$
- **Report:** the three distinct values and the sum. **Interpretation:** what does the sum measure, and is the process stationary?

A5. Solution

Solution

- (a) $\beta_w \ell_t^{(w)} = \sum_{j=1}^5 (\beta_w/5) \ell_{t-j+1}$ and $\beta_m \ell_t^{(m)} = \sum_{j=1}^{22} (\beta_m/22) \ell_{t-j+1}$; with $\ln RV_t^{(w)} \neq \ell_t^{(w)}$ (Jensen) this step fails for the lecture's log-HAR
- Lag 1 collects all three terms, lags 2–5 the weekly and monthly ones, lags 6–22 only the monthly one: 22 coefficients, 3 free parameters
- (b) $\phi_1 = 0.408 + 0.299/5 + 0.196/22 = 0.476$
- $\phi_{2..5} = 0.299/5 + 0.196/22 = 0.0688$; $\phi_{6..22} = 0.196/22 = 0.0089$
- (c) $\sum_j \phi_j = \phi_1 + 4\phi_2 + 17\phi_6 = \beta_d + \beta_w + \beta_m = 0.903$
- Interpretation: the sum is the persistence of log RV; all $\phi_j \geq 0$ and $\sum_j \phi_j = 0.903 < 1$, so every root of $1 - \sum_j \phi_j z^j$ lies outside the unit circle: stationary, with slow (long-memory-like) decay

A6. Testing the HAR Restrictions [Proposed]

Solution discussed in class

Task

- Model: A5. The unrestricted AR(22) for $\ell_t = \ln RV_t$ has an intercept and 22 slope coefficients $\phi_1, \dots, \phi_{22}$; HAR with averages of logs (A5) is its restricted version
- Assumption: ϕ is the linear projection of ℓ_{t+1} on its 22 lags (errors orthogonal to the regressors), weakly dependent data with finite fourth moments; the errors may be heteroskedastic and serially correlated
- (a) Count the linear restrictions that HAR imposes on ϕ and write them as $R\phi = 0$
- (b) Write the Wald statistic with an HAC covariance and its distribution under H_0
- (c) State the 5% critical value
- **Interpretation:** if the test rejects in a long sample, should we drop HAR for forecasting?



A7. QLIKE: Minimiser and Asymmetry [Solved]

Task

- QLIKE(RV, F) = $RV/F - \ln(RV/F) - 1$ with realised proxy $RV > 0$ and forecast $F > 0$; $\mathbb{E}_t$ = expectation given the information at the forecast origin
- (a) Differentiate $\mathbb{E}_t[\text{QLIKE}(RV, F)]$ with respect to F and solve the first-order condition
- (b) Check the second-order condition
- (c) For $RV = 1$, compute QLIKE and the squared error for $F = 0.5$ and for $F = 1.5$
- **Report:** the minimiser and the four losses. **Interpretation:** which error does QLIKE punish more, and why does this suit risk management?

A7. Solution

Solution

- (a) $\mathbb{E}_t[\text{QLIKE}] = \mathbb{E}_t[\text{RV}]/F + \ln F - \mathbb{E}_t[\ln \text{RV}] - 1$; $\partial/\partial F = -\mathbb{E}_t[\text{RV}]/F^2 + 1/F = 0 \Rightarrow F^* = \mathbb{E}_t[\text{RV}]$
- (b) $\partial^2/\partial F^2 = 2\mathbb{E}_t[\text{RV}]/F^3 - 1/F^2$, which equals $1/F^{*2} > 0$ at F^* : a minimum; F^* depends only on $\mathbb{E}_t[\text{RV}]$, so with a conditionally unbiased proxy the *expected-loss* ranking is robust [Patton \(2011\)](#)
- (c) $\text{QLIKE}(1, 0.5) = 2 - \ln 2 - 1 = 0.307$; $\text{QLIKE}(1, 1.5) = 0.667 + \ln 1.5 - 1 = 0.072$; squared errors 0.25 and 0.25
- Interpretation: QLIKE punishes under-prediction of variance about four times more than the same over-prediction; under-predicted risk is the costly error (too little capital, breached VaR)

A8. QLIKE With Its Arguments Swapped [Proposed]

Solution discussed in class

Task

- Model: A7. A program computes $L(RV, F) = F/RV - \ln(F/RV) - 1$, i.e. QLIKE with the arguments swapped
- (a) Find the forecast $F^\dagger$ that minimises $\mathbb{E}_t[L(RV, F)]$
- (b) Compare $F^\dagger$ with $\mathbb{E}_t[RV]$ using Jensen's inequality
- (c) Say which models the swapped loss favours
- **Interpretation:** why can this bug pass elementary checks?

A9. Kernel Weights and Effective Sample Size [Solved]

Task

- ▣ Four past windows with signature distances $\delta = (0.2, 0.5, 1.0, 2.0)$ to the current window
- ▣ Weights of Gu et al. (Eq. 12): $w_\tau = e^{-\gamma\delta_\tau} / \sum_s e^{-\gamma\delta_s}$; Kish effective sample size $n_{\text{eff}} = 1 / \sum_\tau w_\tau^2$
- (a) Compute the four weights for $\gamma = 1$
- (b) Compute the four weights for $\gamma = 3$
- (c) Compute n_{eff} in both cases
- ▣ **Interpretation:** what is gained and what is lost when γ increases?

A9. Solution

Solution

- ▣ (a) $e^{-\delta} = (0.819, 0.607, 0.368, 0.135)$, sum 1.929: $w = (0.425, 0.315, 0.191, 0.070)$
- ▣ (b) $e^{-3\delta} = (0.549, 0.223, 0.050, 0.002)$, sum 0.824: $w = (0.666, 0.271, 0.060, 0.003)$
- ▣ (c) $n_{\text{eff}} = 3.12$ for $\gamma = 1$ and 1.92 for $\gamma = 3$ (out of 4)
- ▣ Interpretation: a higher temperature focuses on the most similar window (less bias from dissimilar periods) but halves the information (more variance); in the lecture $n_{\text{eff}} \approx 816$ of 1000: little is lost, and little changes

A10. Limits of the Temperature [Proposed]

Solution discussed in class

Task

- Model: A9. n past windows, distances $\delta_\tau > 0$, temperature $\gamma = c/\text{median}(\delta)$ as in the lecture
- (a) Find the limits of w and n_{eff} when $c \rightarrow 0$
- (b) Find the limits when $c \rightarrow \infty$ (assume a unique nearest window)
- (c) Show that multiplying all distances by $k > 0$ leaves the weights unchanged
- **Interpretation:** why is the median scaling needed to compare c across 50 assets?

B1. Does the Lévy Area Predict Volatility? (JPM) [Solved]

Task

- Data: VOLARE [Cipollini et al. \(2026\)](#), JPM (JPMorgan Chase), daily five-minute realised variance `rv5`, 2 Jan 2015 – 30 Jan 2026, trading days only; open and close prices for the open-to-close return r_t
- Target: $y_t = \ln \text{RV}_{t+1:t+22} - \ln \text{RV}_t^{(m)}$, the change of next-month log RV relative to the past month (overlapping monthly targets)
- (a) For every day, build the path (cumulative r , $\ln \text{RV}$) of the last 22 days and compute its Lévy area A_t from the depth-2 signature; standardise A_t
- (b) Regress y_t by OLS on a constant, the three log-HAR terms and A_t ; compute Newey–West standard errors with 22 lags
- (c) Test H_0 : coefficient of $A_t = 0$ at 5%, two-sided, and compare the R^2 with and without A_t
- **Report:** coefficient, HAC s.e., t , decision, both R^2 . **Interpretation:** why 22 lags, and is the gain economically relevant?

B1. Solution (1/2)

Solution

- (a) Procedure: signature of each 22-point window with channels (return, $\ln RV$), $A_t = (S^{12} - S^{21})/2$; negative in 46% of the windows (consistent with, but not identifying, leverage-type sequences)
- (b) 2743 usable days; coefficient of $A_t = +0.042$, HAC s.e. 0.036
- (c) $t = +0.042/0.036 = +1.18$; $|t| < 1.96$: do not reject H_0 at 5%
- R^2 : 28.0% without A_t , 28.6% with A_t

B1. Solution (2/2)

Solution

- Why 22 lags: consecutive targets share 21 of their 22 days, so the errors are serially dependent even under H_0 ; 22 lags is our bandwidth choice, and persistent RV can extend dependence further: check 44 lags too
- Relevance: half a point of R^2 next to the 28% explained by the log-HAR terms is negligible
- Conclusion: for JPM the order of moves between price and volatility adds no reliable information about next month's volatility

B2. The Lévy Area for E-mini S&P 500 Futures [Proposed]

Solution discussed in class

Task

- Model: B1. Data: VOLARE, ES (E-mini S&P 500 futures), rv5 and open/close prices, 28 Sep 2009 – 30 Jan 2026
- (a) Repeat steps (a)–(c) of B1 for ES, with the same target, regressors and 22 Newey–West lags
- (b) Test the coefficient of A_t at 1% and 5%, two-sided
- (c) Compare the R^2 gain with the one for JPM
- **Interpretation:** what does a positive coefficient say about windows with a negative area, and can the sign alone label them leverage episodes?

B3. HAR, SHAR and HARQ for JPM, With HAC Errors [Solved]

Task

- Data: VOLARE [Cipollini et al. \(2026\)](#), JPM (JPMorgan Chase), daily five-minute realised variance `rv5`, 2 Jan 2015 – 30 Jan 2026, trading days only; semivariances `rsp5`, `rsn5`; quarticity `rq5`
- Target: next-day RV in levels; full sample; Newey–West standard errors with 5 lags
- (a) Estimate HAR [Corsi \(2009\)](#) by OLS and report the three slopes with HAC standard errors
- (b) Estimate SHAR [Patton & Sheppard \(2015\)](#) and test $H_0 : \beta^- = \beta^+$ against $H_1 : \beta^- > \beta^+$ with a Wald t that uses the HAC covariance of the two estimates
- (c) Estimate HARQ [Bollerslev, Patton & Quaedvlieg \(2016\)](#) with $RQ_t^{1/2}$ demeaned and test $H_0 : \beta_Q = 0$ against $\beta_Q < 0$
- **Report:** estimates, t , p , decisions. **Interpretation:** do the two effects of the literature hold for this stock?

B3. Solution (1/2)

Solution

- (a) HAR (2 764 days): $\beta_d = 0.367$ (0.080), $\beta_w = 0.494$ (0.149), $\beta_m = -0.026$ (0.095)
- (b) SHAR: $\beta^+ = 0.401$ (0.252), $\beta^- = 0.324$ (0.236); $\hat{\beta}^- - \hat{\beta}^+ = -0.078$
- $t = (\hat{\beta}^- - \hat{\beta}^+) / \sqrt{V_{--} + V_{++} - 2V_{+-}} = -0.17$, one-sided $p = 0.566$: do not reject H_0
- (c) HARQ: $\hat{\beta}_Q = -35.062$ (15.904), $t = -2.20 < -1.645$: reject H_0 at 5% one-sided

B3. Solution (2/2)

Solution

- Reading HARQ: on a day with high quarticity (noisy RV) the effective daily coefficient $\beta_d + \beta_Q(RQ_t^{1/2} - \overline{RQ^{1/2}})$ falls: the model trusts that day's RV less
- Reading SHAR: the standard errors of β^+ and β^- are large because the two semivariances are highly correlated for a single stock
- Conclusion: the HARQ effect holds for JPM; the good/bad volatility asymmetry is not detectable for JPM alone (it is a panel result in Patton and Sheppard)

B4. HAR, SHAR and HARQ for Crude Oil [Proposed]

Solution discussed in class

Task

- Model: B3. Data: VOLARE, CL (crude oil futures), rv5, rsp5, rsn5, rq5, 28 Sep 2009 – 30 Jan 2026
- (a) Estimate SHAR with Newey–West (5 lags) and test $\beta^- > \beta^+$ as in B3(b)
- (b) Estimate HARQ and test $\beta_Q < 0$ as in B3(c)
- (c) Compare both results with JPM
- **Interpretation:** why might the asymmetry be stronger for a commodity than for a bank stock?

B5. Sig-LK vs log-HAR for JPM, One Day Ahead [Solved]

Task

- Data: VOLARE [Cipollini et al. \(2026\)](#), JPM (JPMorgan Chase), daily five-minute realised variance $rv5$, 2 Jan 2015 – 30 Jan 2026, trading days only
- Design of the lecture: rolling window of 1000 pairs, re-estimation every 5 days, path (time, $\ln RV$, cumulative r) over 22 days, depth 3, two-step LASSO with BIC, $c = 0.5$
- (a) Produce one-day forecasts of Sig-LK and log-HAR from 27 Jan 2020 to 29 Jan 2026 (1 511 days) and compute the mean QLIKE of each
- (b) Compute the ratio of Sig-LK mean QLIKE to log-HAR mean QLIKE and a 95% CI with the stationary bootstrap (mean block length 10, 1000 replications)
- (c) Run the DM test (HAC variance, bandwidth $\max(h, \lceil T^{1/3} \rceil)$) and the conditional GW test with instruments $(1, d_{t-1})$
- **Report:** losses, ratio and CI, t , p , decisions. **Interpretation:** is there any evidence that the signature adds value for JPM?

B5. Solution

Solution

- (a) Mean QLIKE: Sig-LK 0.1726, log-HAR 0.1727; on average 4.2 regressors selected, effective sample size 851 of 1000
- (b) Ratio 1.000; bootstrap 95% CI [0.956, 1.047] contains 1
- (c) DM: $d_t = \text{QLIKE}_t^{\text{Sig-LK}} - \text{QLIKE}_t^{\text{log-HAR}}$, $t = \bar{d} / \sqrt{\hat{\Omega}_{\text{NW}} / T} = -0.02$, one-sided $p = 0.493$: do not reject equal accuracy
- GW: Wald statistic $2.85 \sim \chi^2(2)$, $p = 0.241$: the instruments $(1, d_{t-1})$ give no significant evidence that the loss difference is predictable
- Interpretation: no evidence; the two forecasts are practically identical, and the CI rules out gains or losses larger than about 5%

B5. Cumulative Loss Difference (JPM, One Day Ahead)



- Upward segments: log-HAR better; downward segments: Sig-LK better; the largest moves sit in March 2020 and April 2025
- A handful of stress days decides the sign of the total: report the cumulative path, not only the mean

B6. Sig-LK vs log-HAR for ES, One Week Ahead [Proposed]

Solution discussed in class

Task

- ▣ Model: B5. Data: VOLARE, ES (E-mini S&P 500 futures), rv_5 and open/close prices, 28 Sep 2009 – 30 Jan 2026; horizon $h = 5$, $c = 1.0$ (futures validation block)
- (a) Produce weekly forecasts (average RV over the next 5 days) of Sig-LK and log-HAR with the design of B5
- (b) Compute the QLIKE ratio and its stationary-bootstrap 95% CI (mean block length 10)
- (c) Run the DM test (bandwidth at least $h = 5$) and the conditional GW test
- ▣ **Interpretation:** why is the CI so asymmetric, and what does it imply for the power of DM?

B7. Holm and BH, Step by Step [Solved]

Task

- From the lecture: one-sided DM p -values for “Sig-LK better than log-HAR” at $h = 1$ (QLIKE), one per asset; family = the 50 assets
- The five FX pairs, sorted: EURUSD 0.071, USDJPY 0.213, USDCAD 0.336, AUDUSD 0.718, GBPUSD 0.817
- (a) Apply Holm (FWER 5%) to the five FX p -values: compare $p_{(k)}$ with $0.05/(m - k + 1)$ and stop at the first non-rejection
- (b) Apply BH (FDR 5%) to the same five: find the largest k with $p_{(k)} \leq 0.05 k/m$
- (c) Repeat both procedures for all 50 assets and count the rejections
- **Report:** rejections raw, Holm, BH. **Interpretation:** what may we claim, and about how many assets?

B7. Solution

Solution

- (a) Holm, $m = 5$: $p_{(1)} = 0.071$ vs $0.05/5 = 0.01$: no rejection, stop; Holm-adjusted p : 0.354, 0.852, then 1
- (b) BH: thresholds 0.01, 0.02, 0.03, 0.04, 0.05; every $p_{(k)}$ is above its threshold: no rejection; the smallest BH-adjusted p is 0.354
- (c) 50 assets: 2 raw rejections; Holm: the smallest, ES with $p = 0.002$, is above $0.05/50 = 0.001$: stop, 0 rejections; BH: $p_{(1)} > 0.001$ and $p_{(2)} = 0.028 > 0.002$, no k qualifies: 0
- Interpretation: no asset-level improvement survives either correction; ES is at most a lead for a new, pre-registered study. Holm is valid under any dependence; BH needs independent or positively dependent (PRDS) p -values [Benjamini & Yekutieli \(2001\)](#)

B8. Multiple Testing at One Month Against HAR [Proposed]

Solution discussed in class

Task

- Model: B7. From the lecture: 50 one-sided DM p -values for “Sig-LK better than HAR” at $h = 22$ (QLIKE)
- (a) Count the rejections at 5% without correction
- (b) Apply Holm and BH as in B7 and count the rejections; name the asset with the smallest p -value
- (c) Compare with the pooled mean QLIKE ratio of Sig-LK relative to HAR at $h = 22$ (1.334)
- **Interpretation:** how can Sig-LK beat HAR on several assets and still lose on average?

C1. A Level-Aware Signature Kernel [Proposed]

Solution discussed in class

Task

- Context: the kernel of the lecture is translation-invariant, so it cannot recognise a high-volatility regime (Lecture 20, weights on 16 March 2020); data: VOLARE, the 40 American stocks, $rv5$ and open/close prices, 2015–2026
- Idea: prepend a **basepoint** [Morrill et al. \(2020\)](#): the path starts at the training mean of $\ln RV$ and jumps to the first value of the window; by Chen, $\text{Sig} = \exp_{\otimes}(v) \otimes \text{Sig}(X)$ with v the standardised level jump
- (a) Write a one-page pre-registration: 40 stocks, $h = 1$, the lecture design, $c = 0.5$, QLIKE, DM with Holm, the COVID-19 subperiod reported separately
- (b) Implement the basepoint kernel (use the Chen update of the notebook) with the same features as Sig-LK, and forecast
- (c) Compare with Sig-LK and log-HAR: mean QLIKE ratios (all days and COVID-19), wins over log-HAR raw and after Holm
- **Interpretation:** if seeing the level does not help, where would you look next (features, target, horizon)?

C2. Critique an AI Answer (1/2) [Proposed]

Solution discussed in class

- **Student prompt:** “rv holds the daily 5-minute RV of one stock. Build the signature LASSO of Gu et al. for next-day log RV and evaluate it with QLIKE on the last 30% of the sample.”

- **Answer of the AI assistant**

```
lrv = np.log(rv.values)
S = np.array([sig_path(lrv[t-21:t+1, None], depth=3)      # 22-day window
              for t in range(21, len(lrv) - 1)])
y = lrv[22:]                                              # next-day log RV
lam = LassoCV(cv=5).fit(S, y).alpha_                    # choose lambda
cut = int(0.7 * len(y))
m = Lasso(alpha=lam).fit(S[:cut], y[:cut])
f = np.exp(m.predict(S[cut:]))
qlike = lambda a, b: a / b - np.log(a / b) - 1
print(qlike(f, np.exp(y[cut:])).mean())
```

- “The signature LASSO has a lower QLIKE than HAR on the test set, so path signatures capture the shape of volatility and improve the forecast.”

C2. Critique an AI Answer (2/2) [Proposed]

Solution discussed in class

Task

- ▣ Context: the answer on the previous slide contains **at least three** errors that are typical of signature and volatility code
- ▣ Model: A1 and A3 (time augmentation), A7 and A8 (QLIKE), B5 (evaluation without look-ahead)
- (a) Find at least three errors; for each, name the line and explain in one sentence why it is wrong; then assess the exponential back-transform and whether the code is the two-step weighted method of Gu et al.
- (b) For each error, describe a check that would reveal it without knowing the answer (a unit test, a timing test, a known case)
- (c) Write the corrected code and compare it with log-HAR on the same split, with a DM test
- (d) Log each error as in [AI_ERRORS.md](#): what the AI answered, what was wrong, how you found out, the correct version
- ▣ **Interpretation**: does the conclusion of the assistant follow even if the code were correct?

Common Pitfalls

- ▣ Signatures of one-dimensional paths, or without the time channel
- ▣ Hyperparameters (λ , c , depth) tuned on data that are later used for evaluation
- ▣ QLIKE with swapped arguments; DM tests without HAC variance for $h > 1$
- ▣ Comparing only with HAR in levels, so that the log transform gets credited to the new method
- ▣ Reporting the best asset or horizon instead of the family, corrected with Holm or BH

Take-Home Assignment

- ▣ Solve the proposed problems A2, A4, A6, A8, A10 and B2, B4, B6, B8 in the [seminar notebook](#)
- ▣ C1: write a one-page pre-registration for your own variant of the level-aware kernel and run it on five assets
- ▣ C2: correct the AI code and compare it with log-HAR on the same split
 - ▶ Log every AI error you catch from now on in [AI_ERRORS.md](#), as in C2: what the AI answered, what was wrong, how you found out, the correct version

Further Study on Quantinar

Advanced courses on quantinar.com

- ☐  Statistics of Financial Markets
- ☐  Applied Time Series Analysis with Python

Code and data

- ☐ Seminar notebook (Colab); Quantlets: github.com/danpele/MFM/tree/main/Quantlets/Ch_20; data: volare.unime.it

References

- Benjamini, Y., Hochberg, Y. (1995). Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society B*, 57(1), 289–300.
- Benjamini, Y., Yekutieli, D. (2001). The control of the false discovery rate in multiple testing under dependency. *The Annals of Statistics*, 29(4), 1165–1188.
- Bollerslev, T., Patton, A.J., Quaedvlieg, R. (2016). Exploiting the errors: a simple approach for improved volatility forecasting. *Journal of Econometrics*, 192(1), 1–18.
- Chen, K.-T. (1957). Integration of paths, geometric invariants and a generalized Baker–Hausdorff formula. *Annals of Mathematics*, 65(1), 163–178.
- Chevyrev, I., Kormilitzin, A. (2016). A primer on the signature method in machine learning. [arXiv:1603.03788](https://arxiv.org/abs/1603.03788).
- Cipollini, F., Cruciani, G., Gallo, G.M., Insana, A., Otranto, E., Spagnolo, F. (2026). VOLatility Archive for Realized Estimates (VOLARE). [arXiv:2602.19732](https://arxiv.org/abs/2602.19732).
- Corsi, F. (2009). A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2), 174–196.
- Diebold, F.X., Mariano, R.S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263.
- Franke, J., Härdle, W.K., Hafner, C.M. (2019). *Statistics of Financial Markets: An Introduction*, 5th ed. Springer.
- Giacomini, R., White, H. (2006). Tests of conditional predictive ability. *Econometrica*, 74(6), 1545–1578.
- Gu, H., Guo, X., Jacobs, T.L., Kaminsky, P., Li, X. (2024). Transportation marketplace rate forecast using signature transform. *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 4997–5005.
- Hambly, B., Lyons, T. (2010). Uniqueness for the signature of a path of bounded variation and the reduced path group. *Annals of Mathematics*, 171(1), 109–167.
- Morrill, J., Fermanian, A., Kidger, P., Lyons, T. (2020). A generalised signature method for multivariate time series feature extraction. [arXiv:2006.00873](https://arxiv.org/abs/2006.00873).
- Patton, A.J. (2011). Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1), 246–256.
- Patton, A.J., Sheppard, K. (2015). Good volatility, bad volatility: signed jumps and the persistence of volatility. *Review of Economics and Statistics*, 97(3), 683–697.